

Balanced Trees

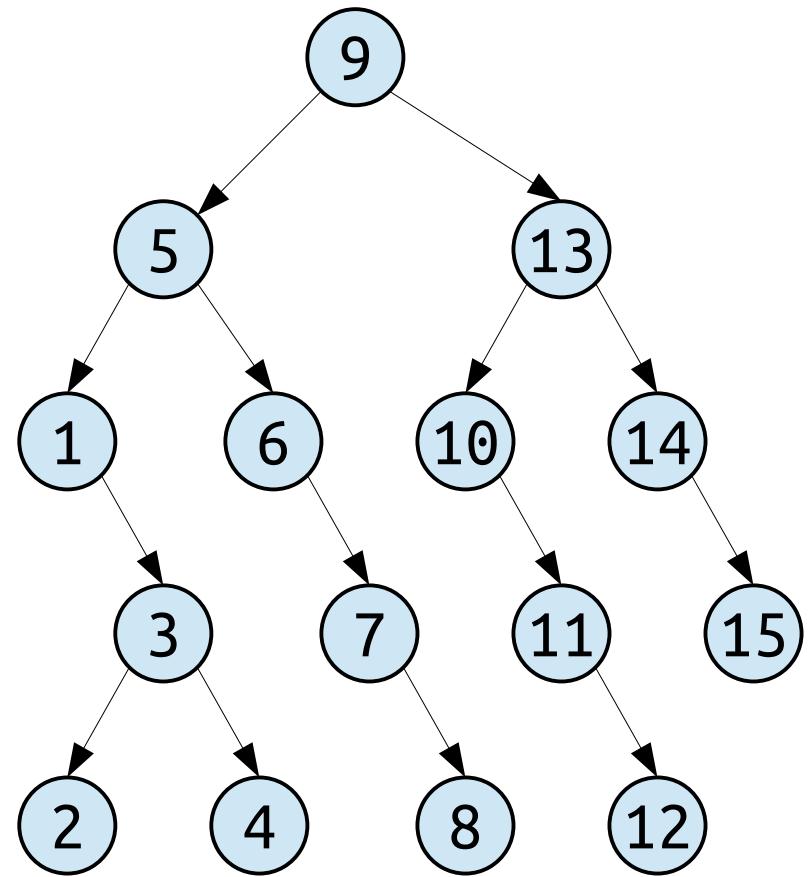
Outline for Today

- ***BST Review***
 - Refresher on basic BST concepts and runtimes.
- ***Overview of Red/Black Trees***
 - What we're building toward.
- ***B-Trees and 2-3-4 Trees***
 - Simple balanced trees, in depth.
- ***Intuiting Red/Black Trees***
 - A much better feel for red/black trees.

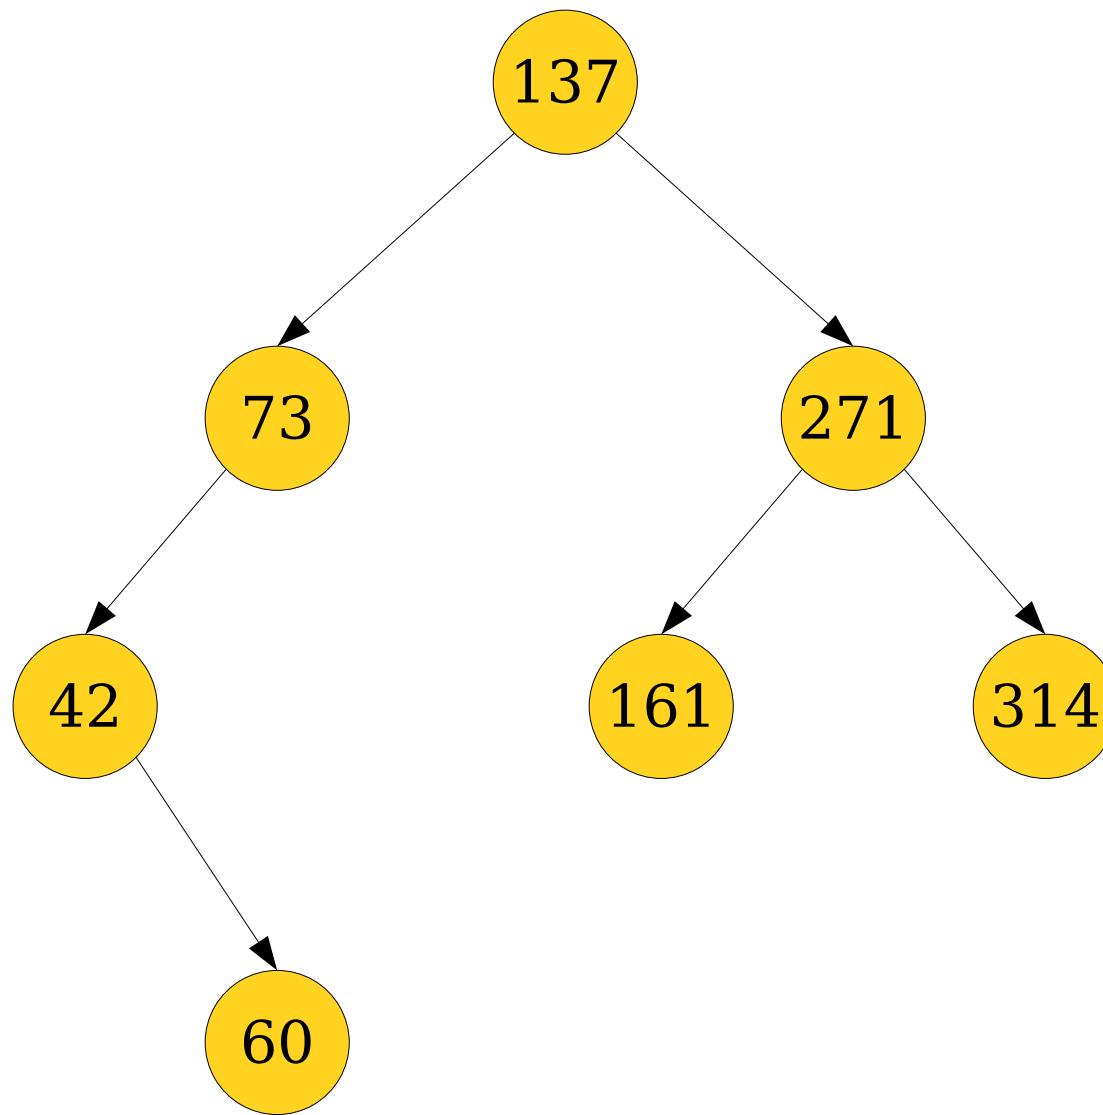
A Quick BST Review

Binary Search Trees

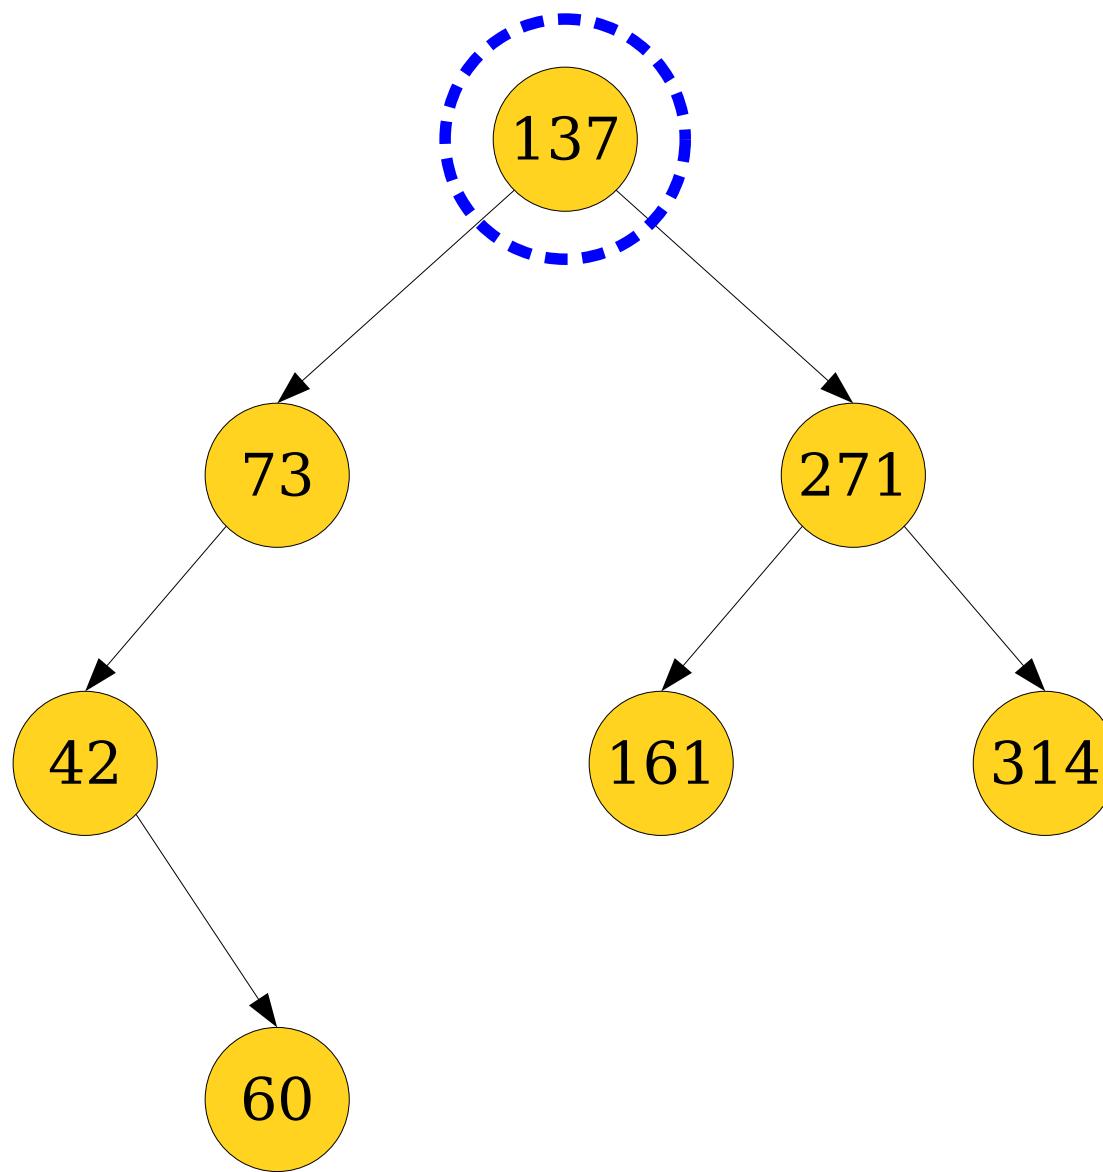
- A ***binary search tree*** is a binary tree with the following properties:
 - Each node in the BST stores a ***key***, and optionally, some auxiliary information.
 - The key of every node in a BST is strictly greater than all keys to its left and strictly smaller than all keys to its right.



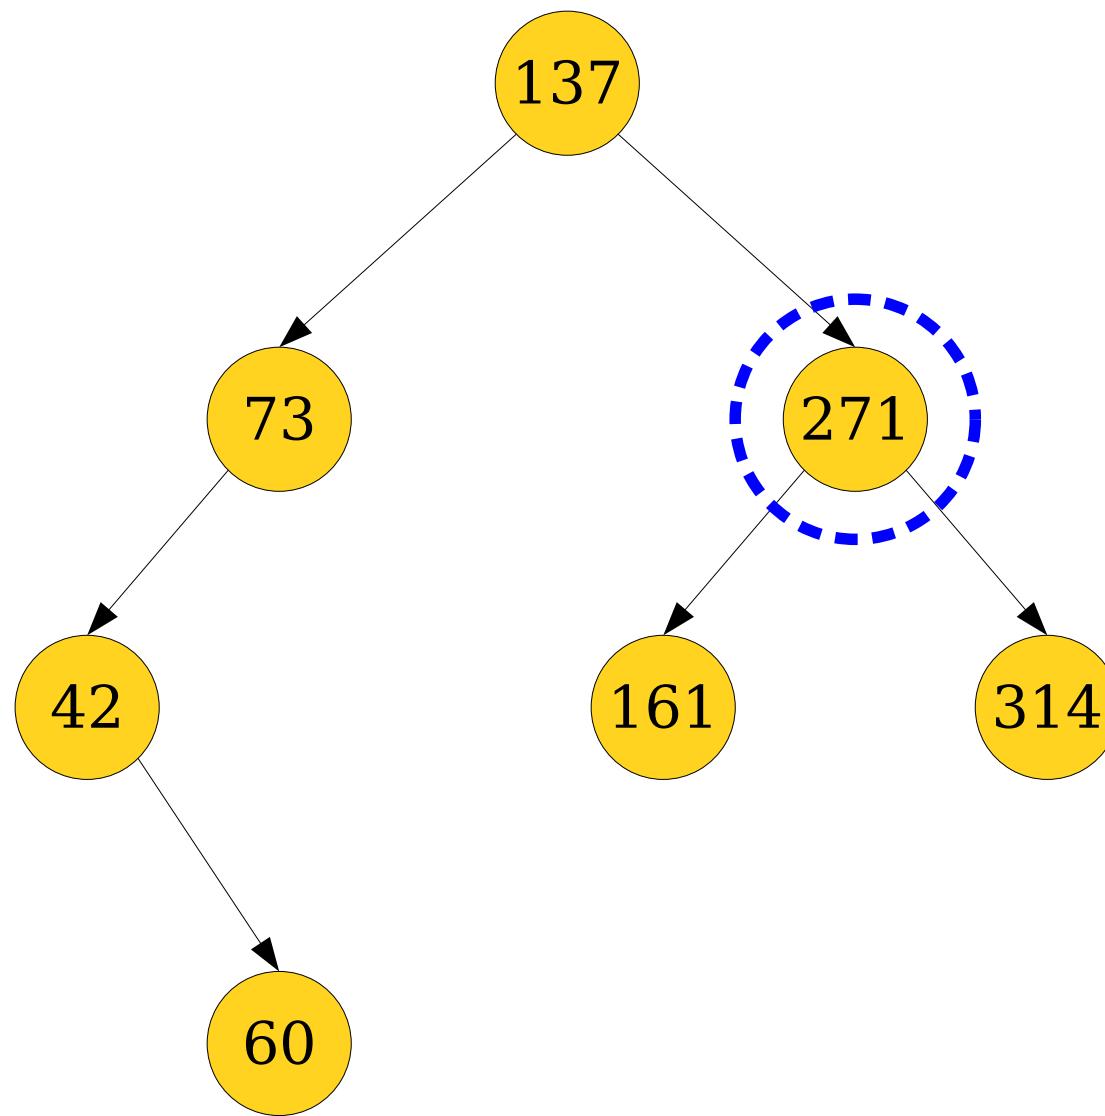
Searching a BST



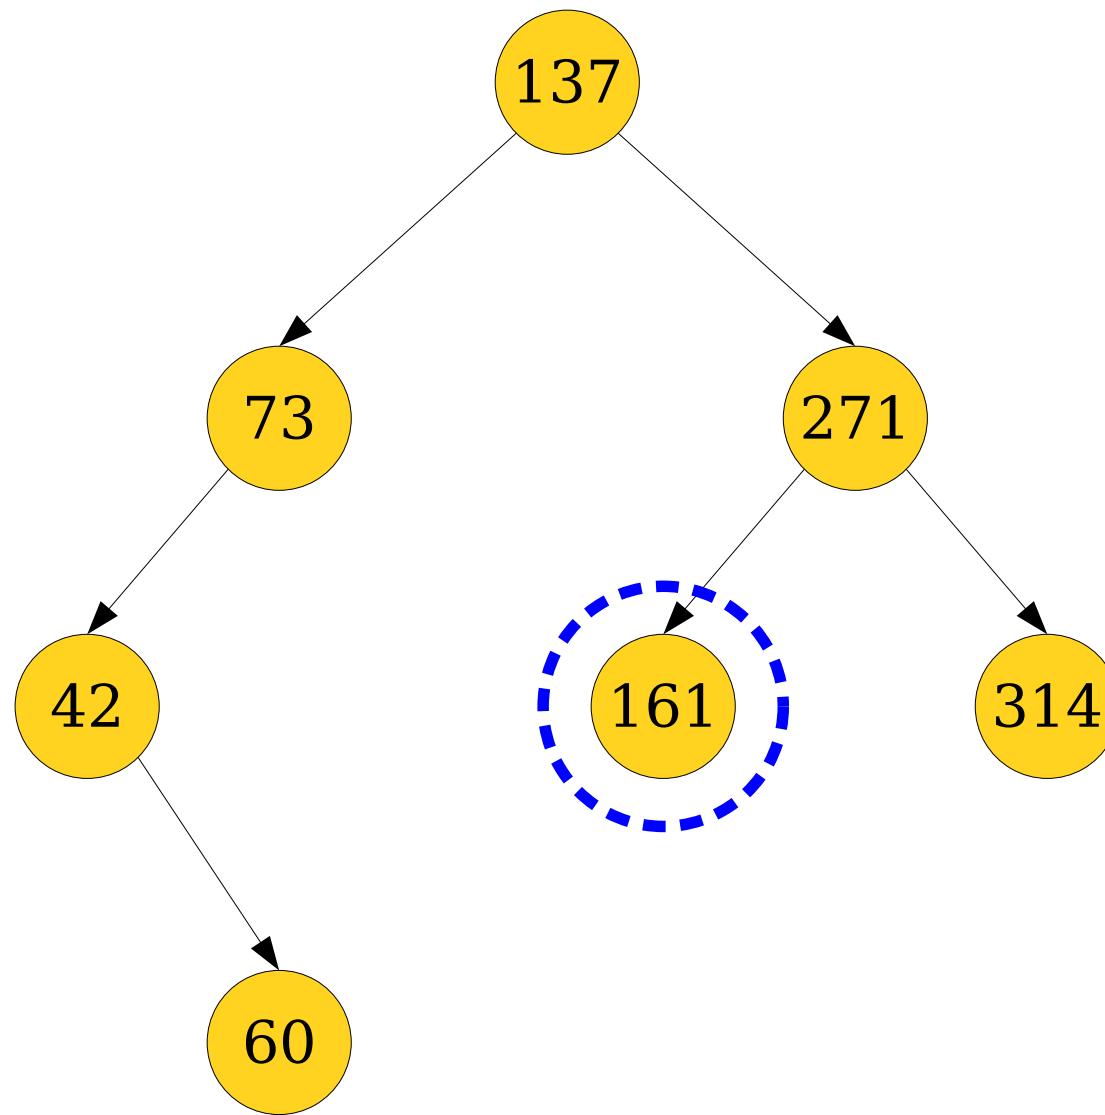
Searching a BST



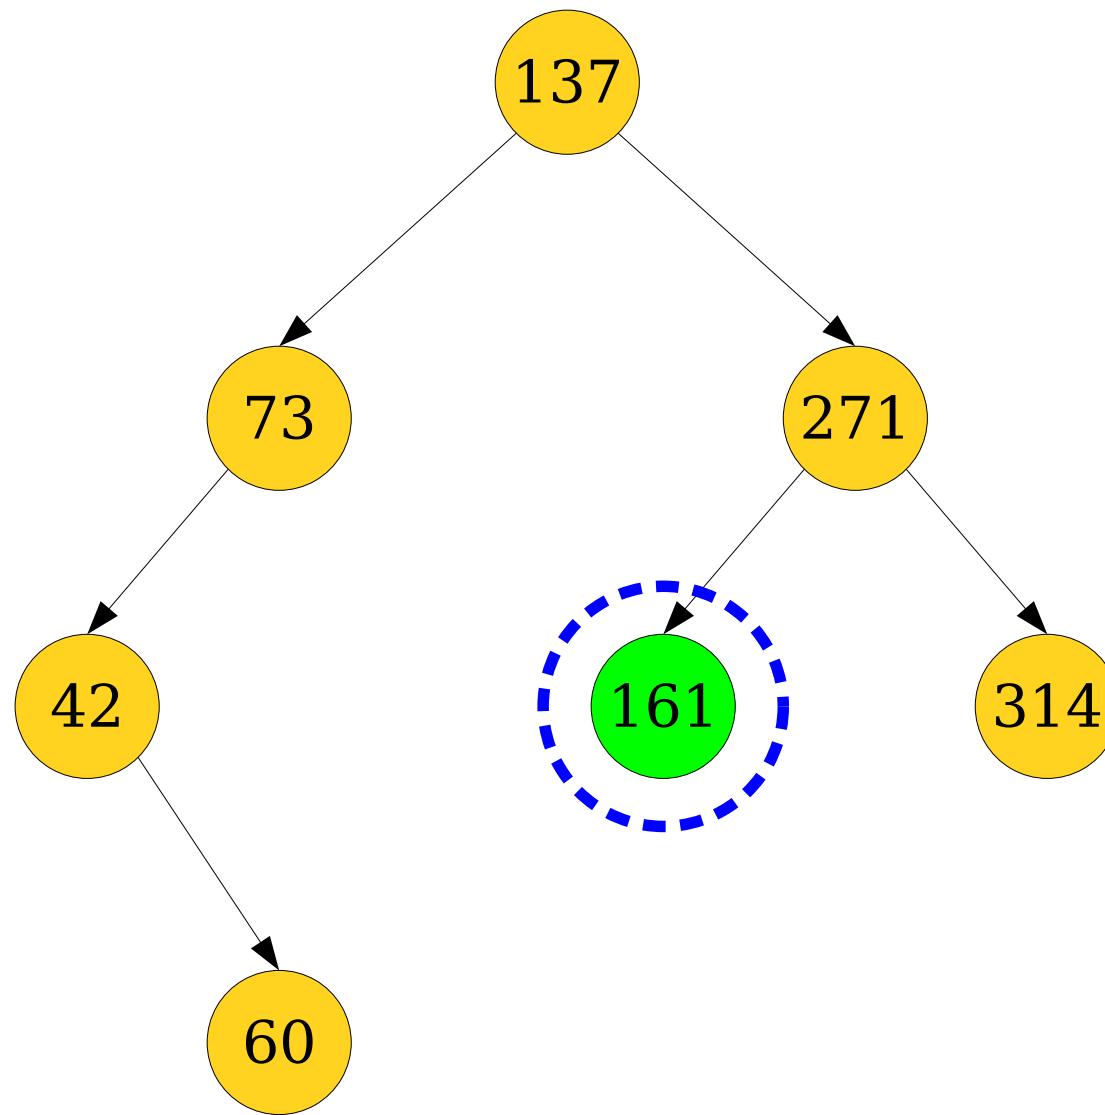
Searching a BST



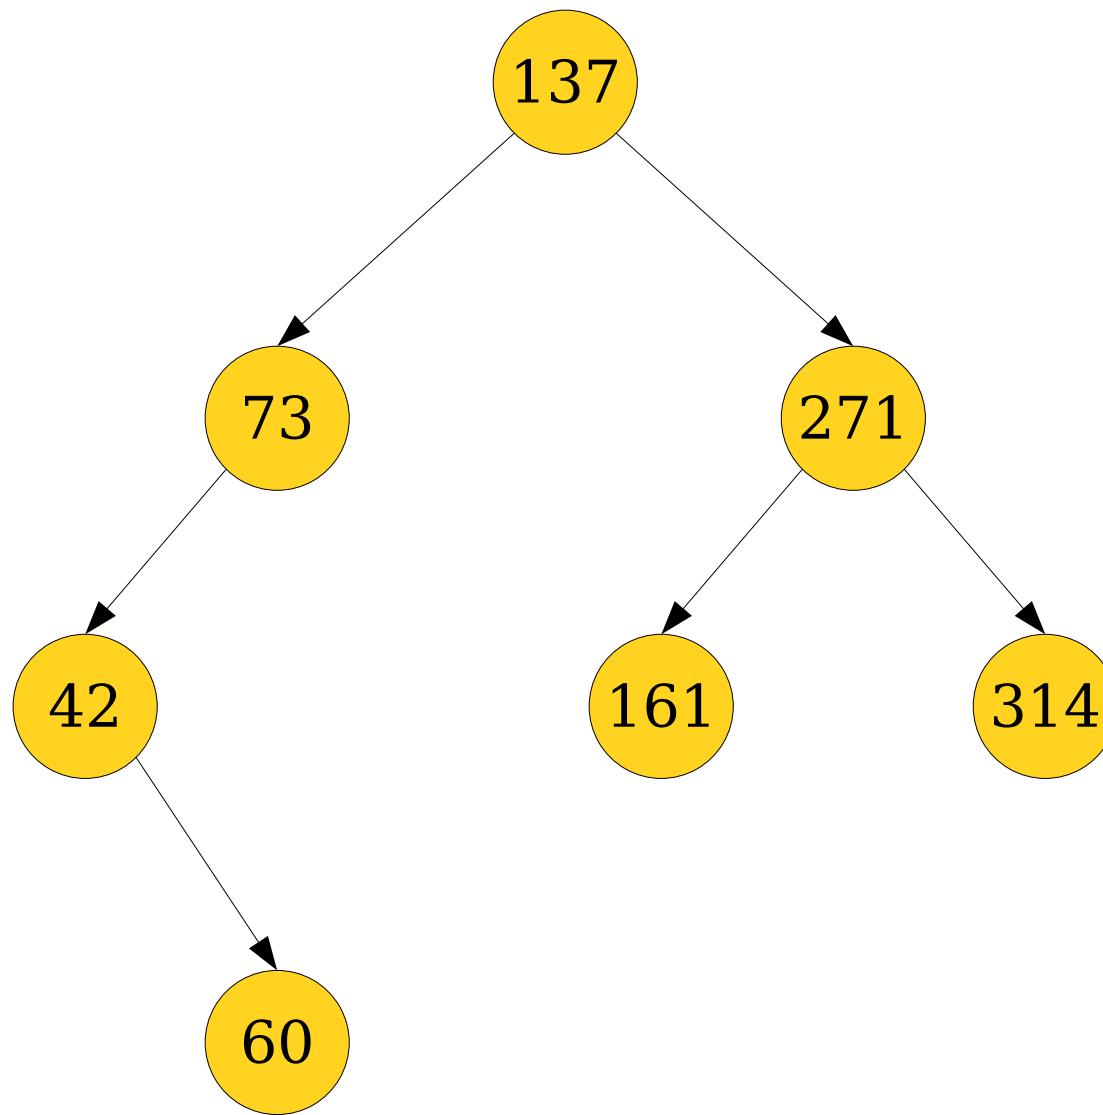
Searching a BST



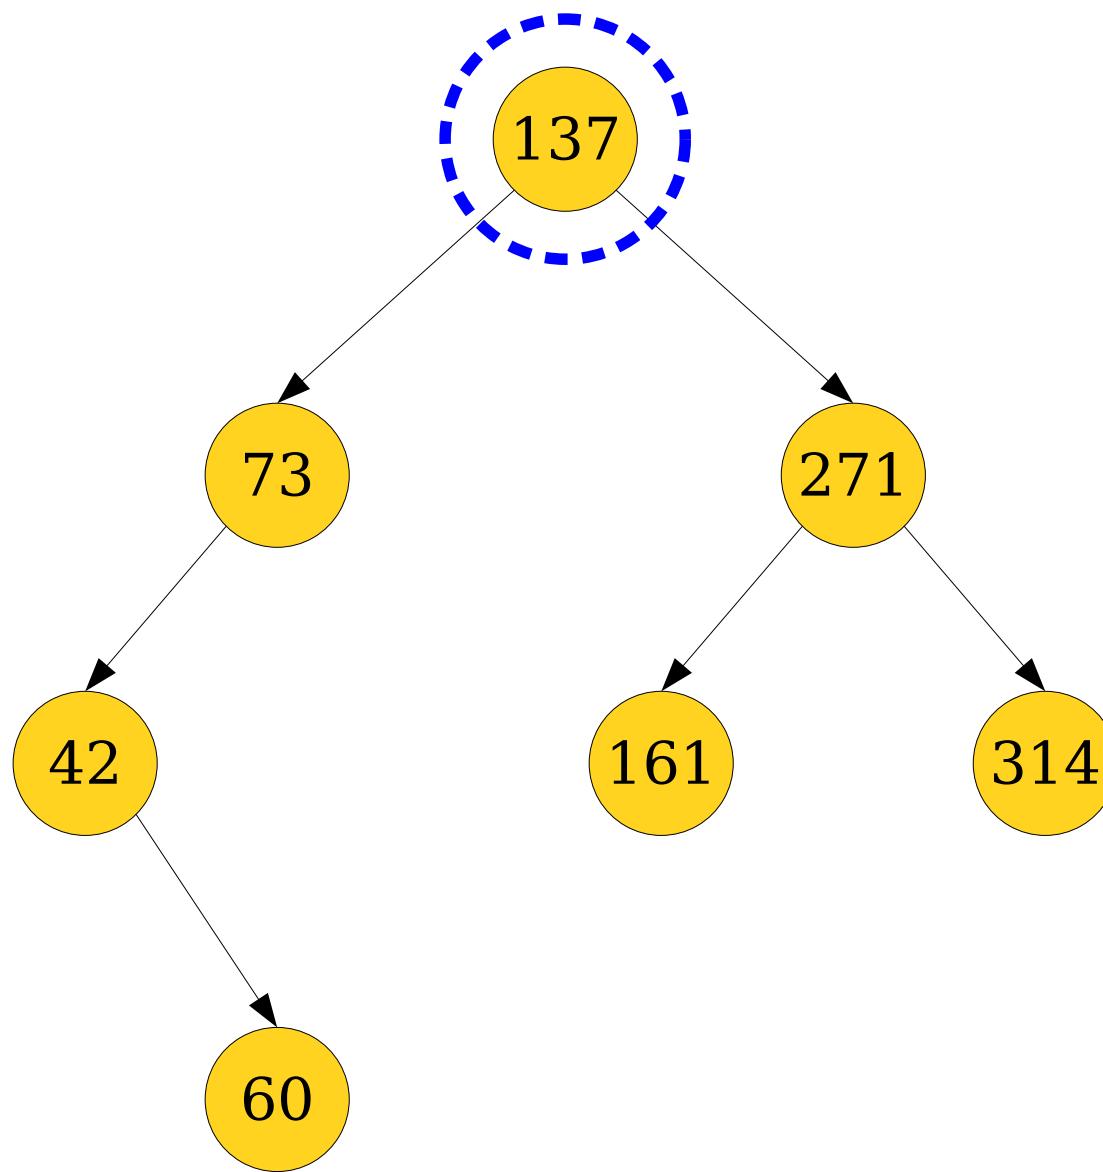
Searching a BST



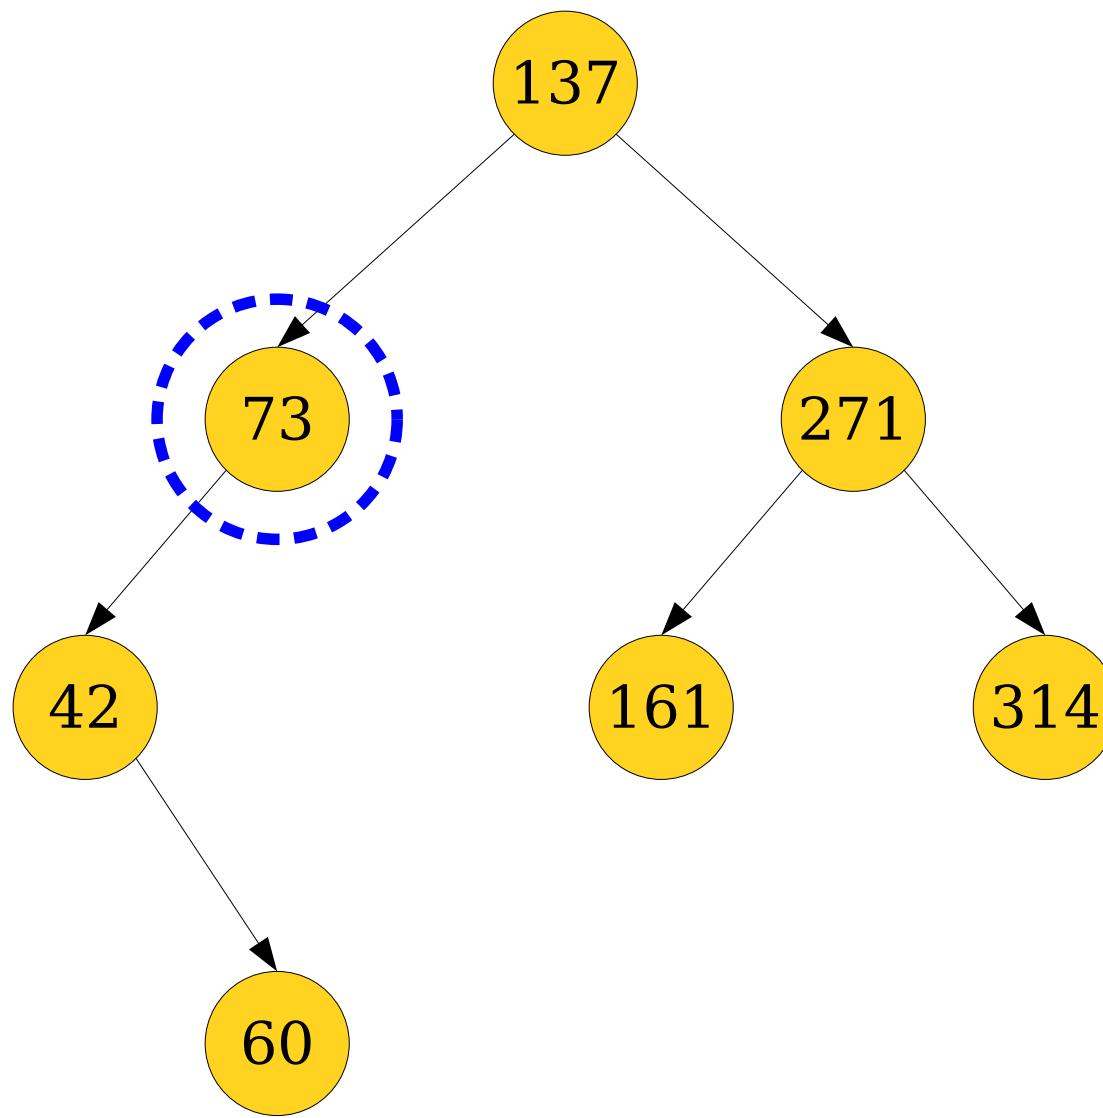
Searching a BST



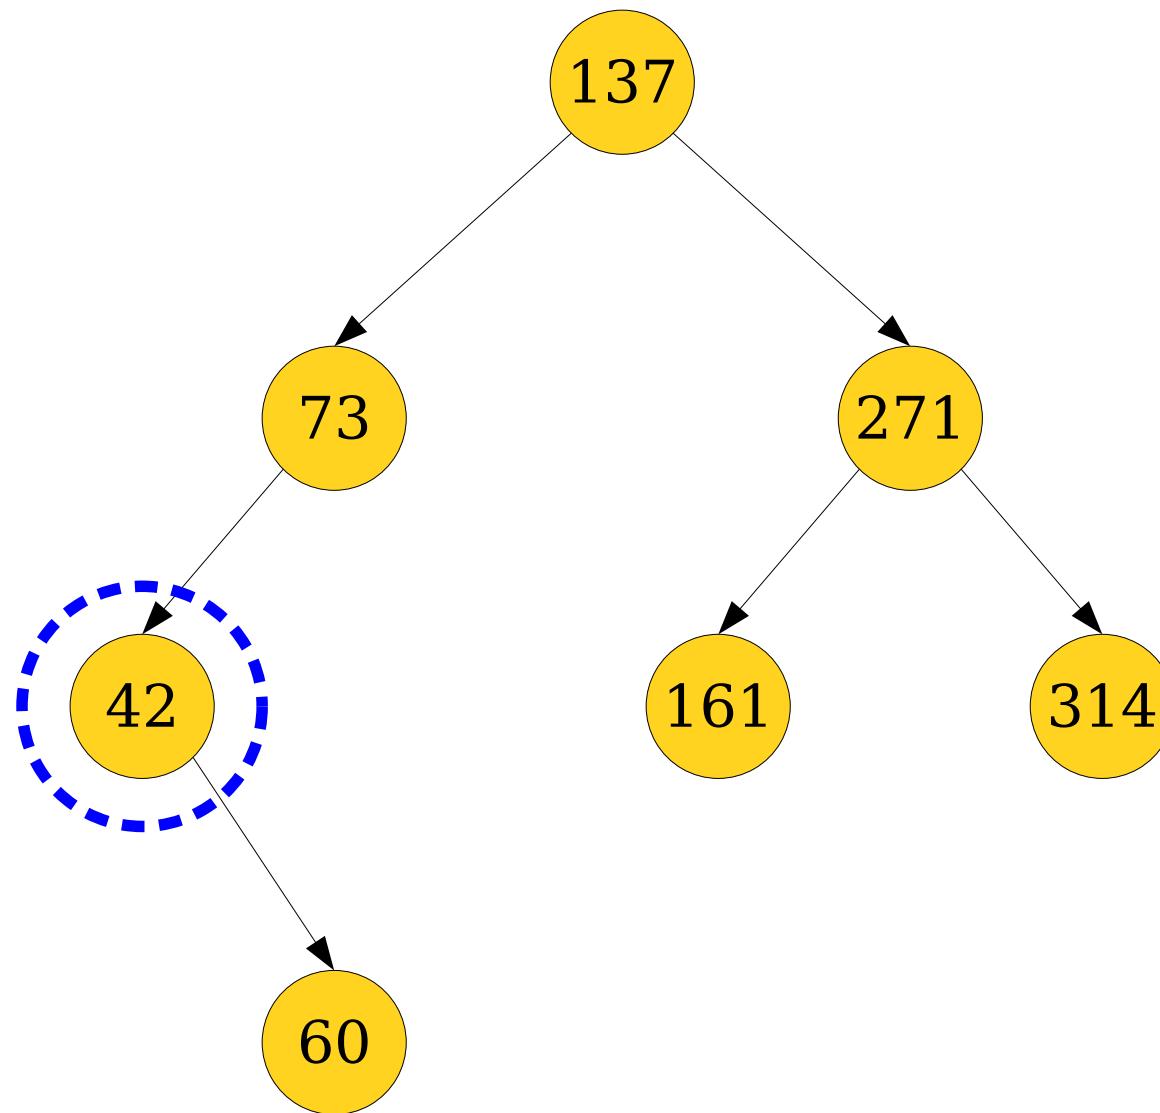
Searching a BST



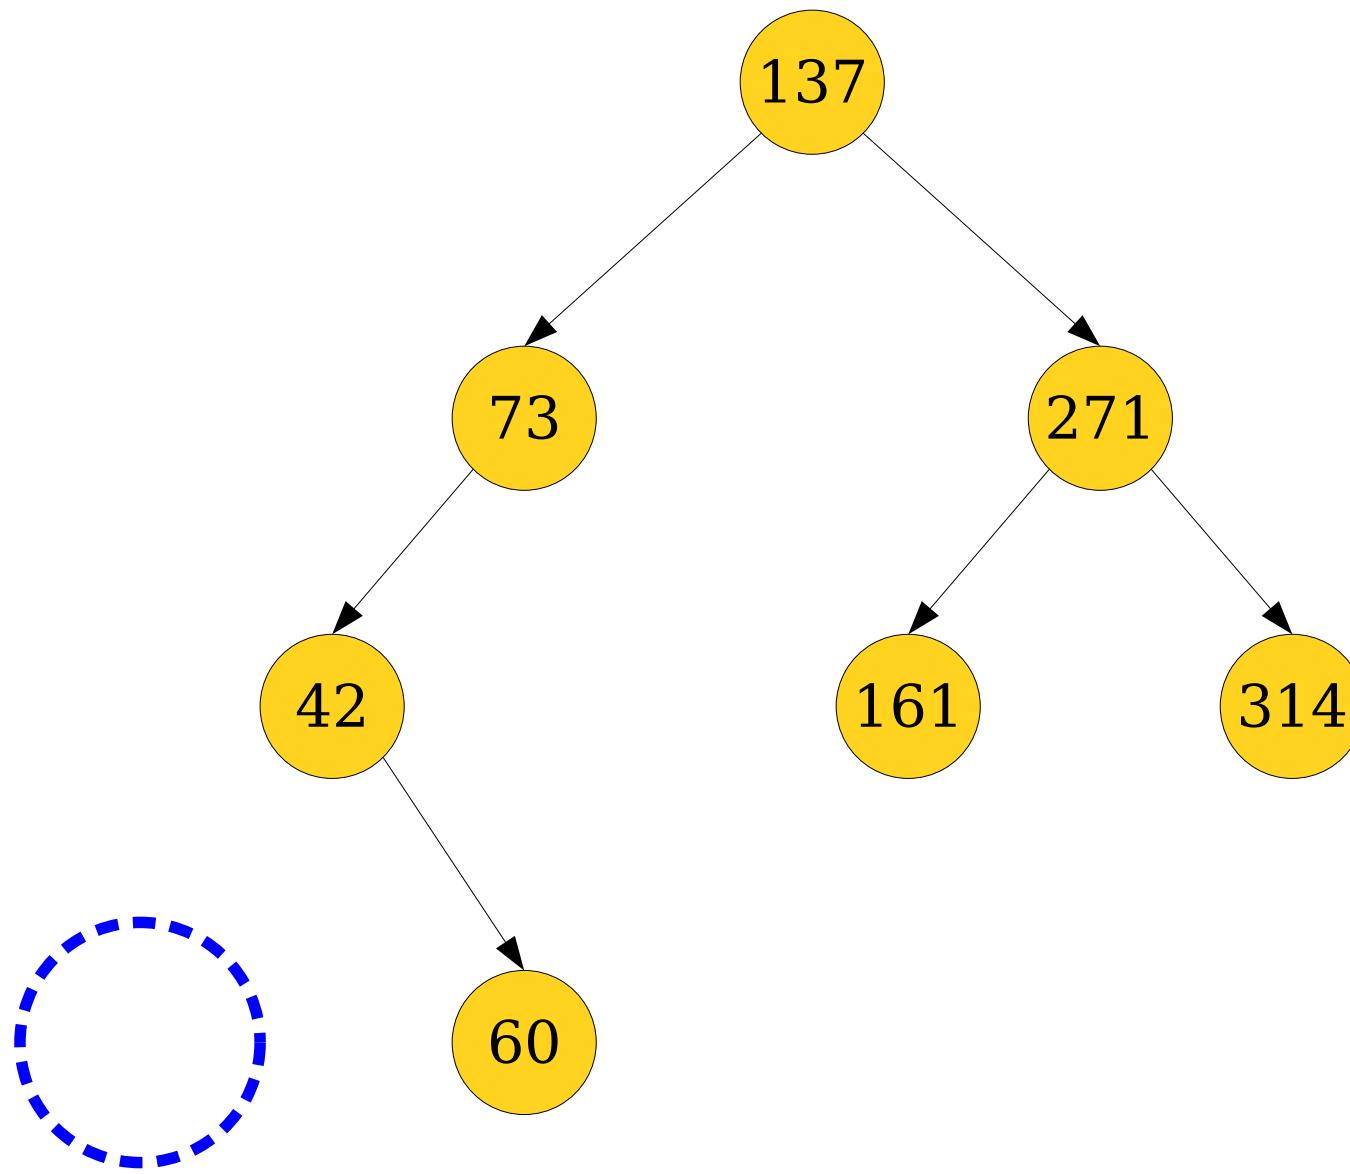
Searching a BST



Searching a BST

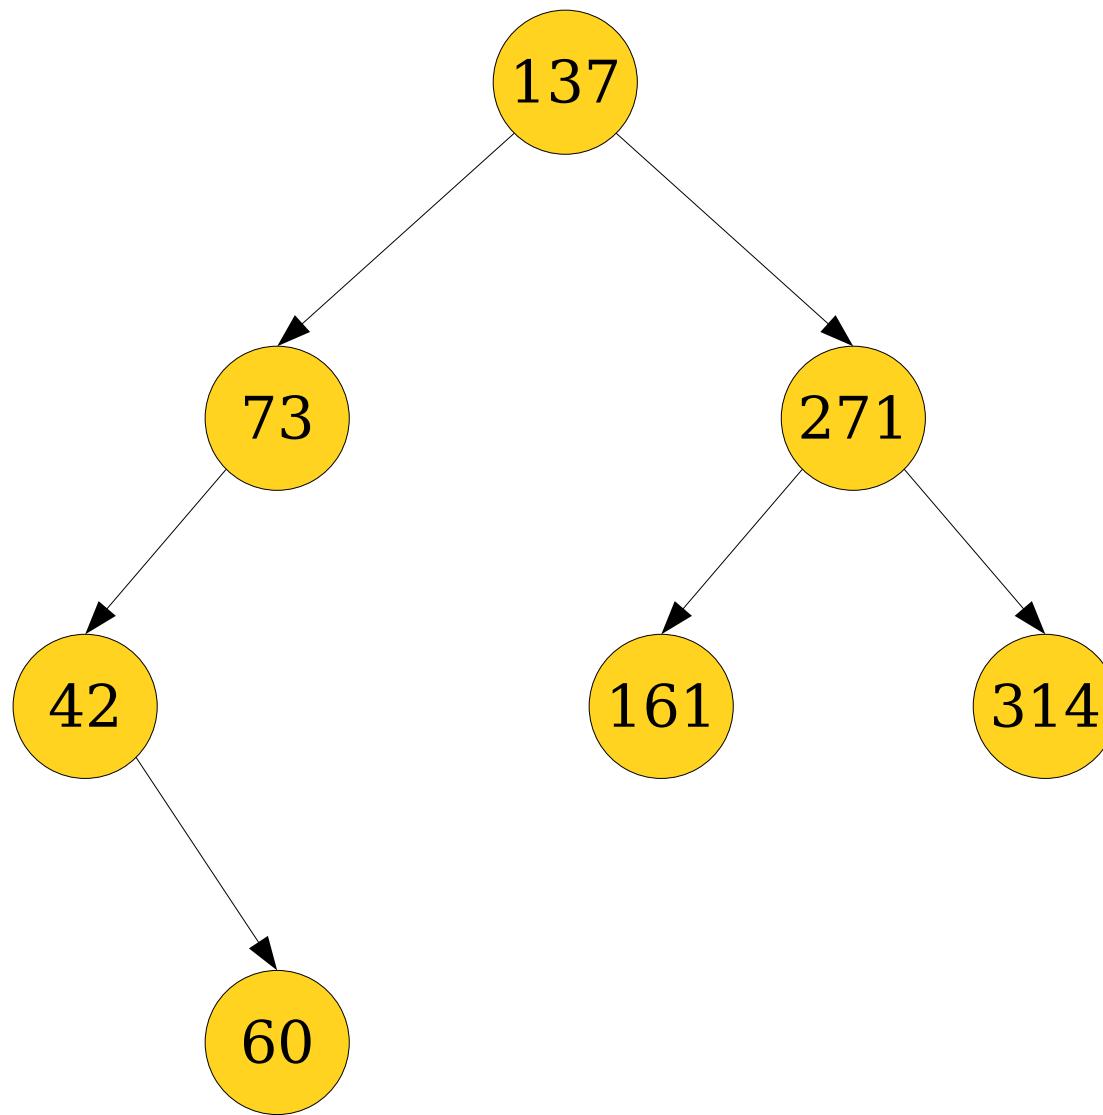


Searching a BST

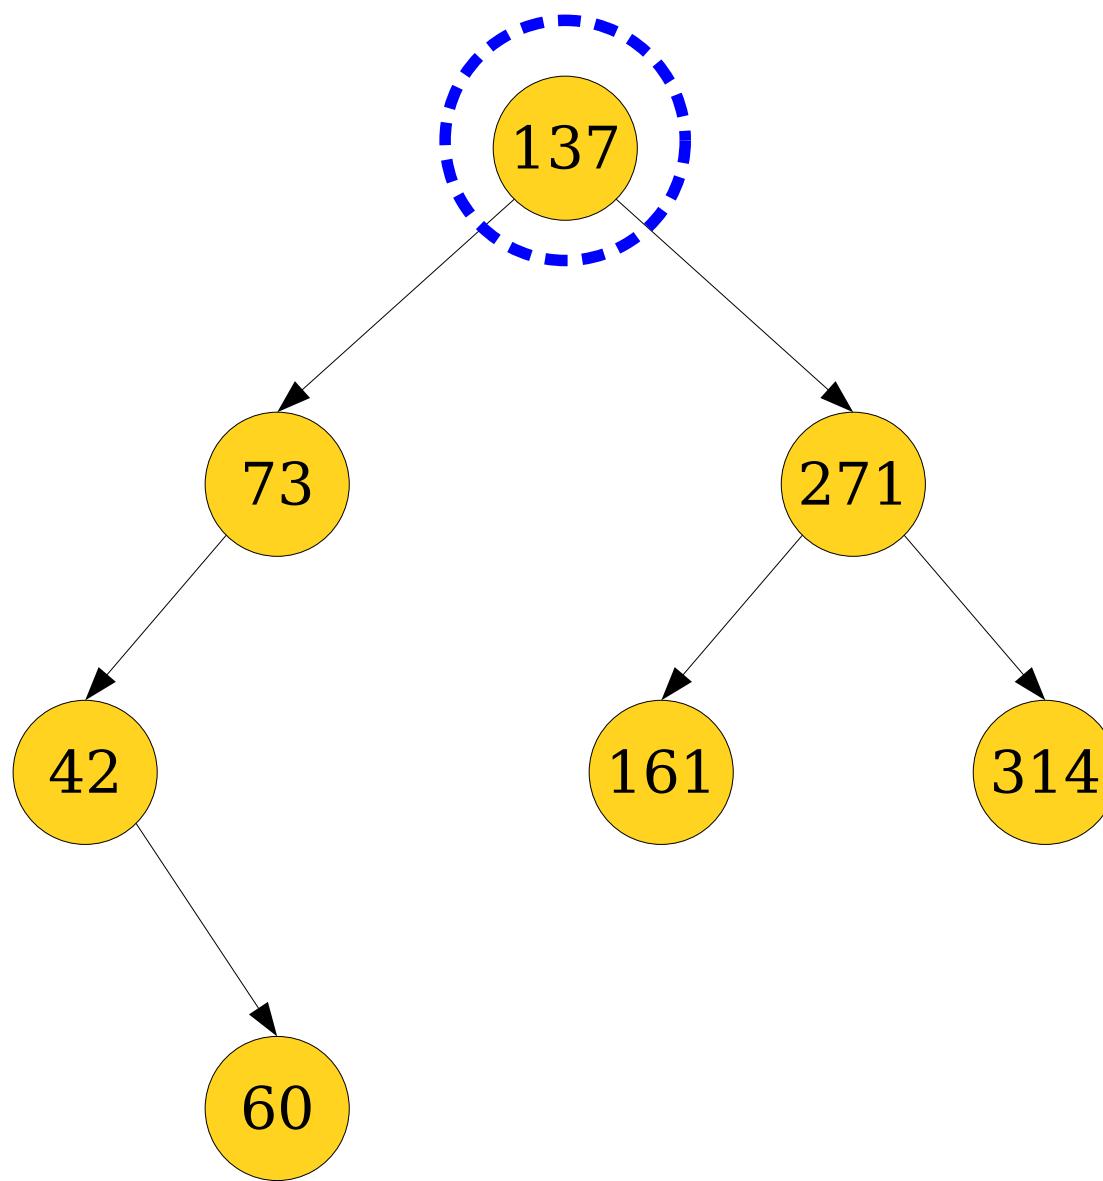


Inserting into a BST

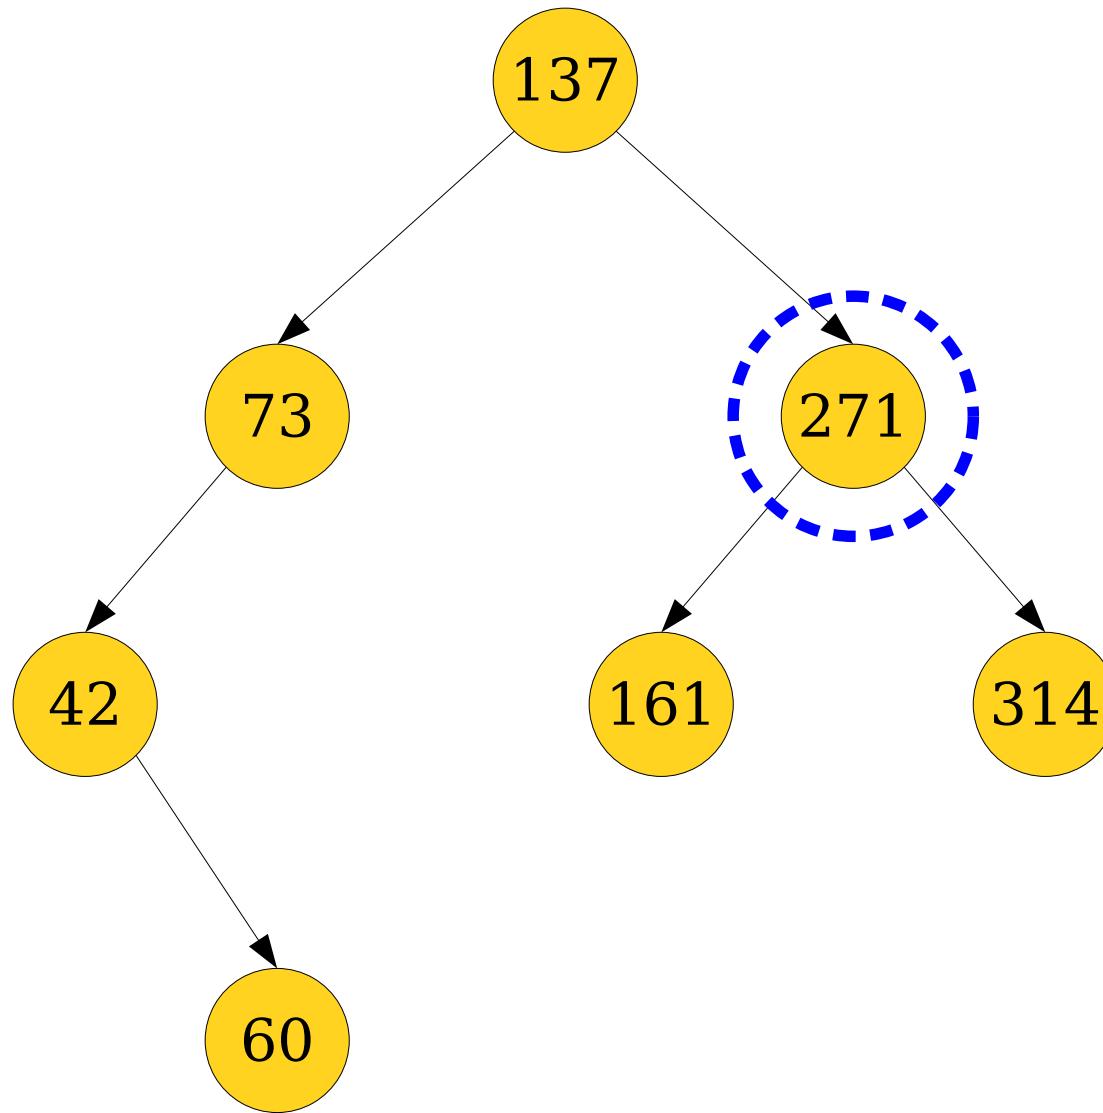
Inserting into a BST



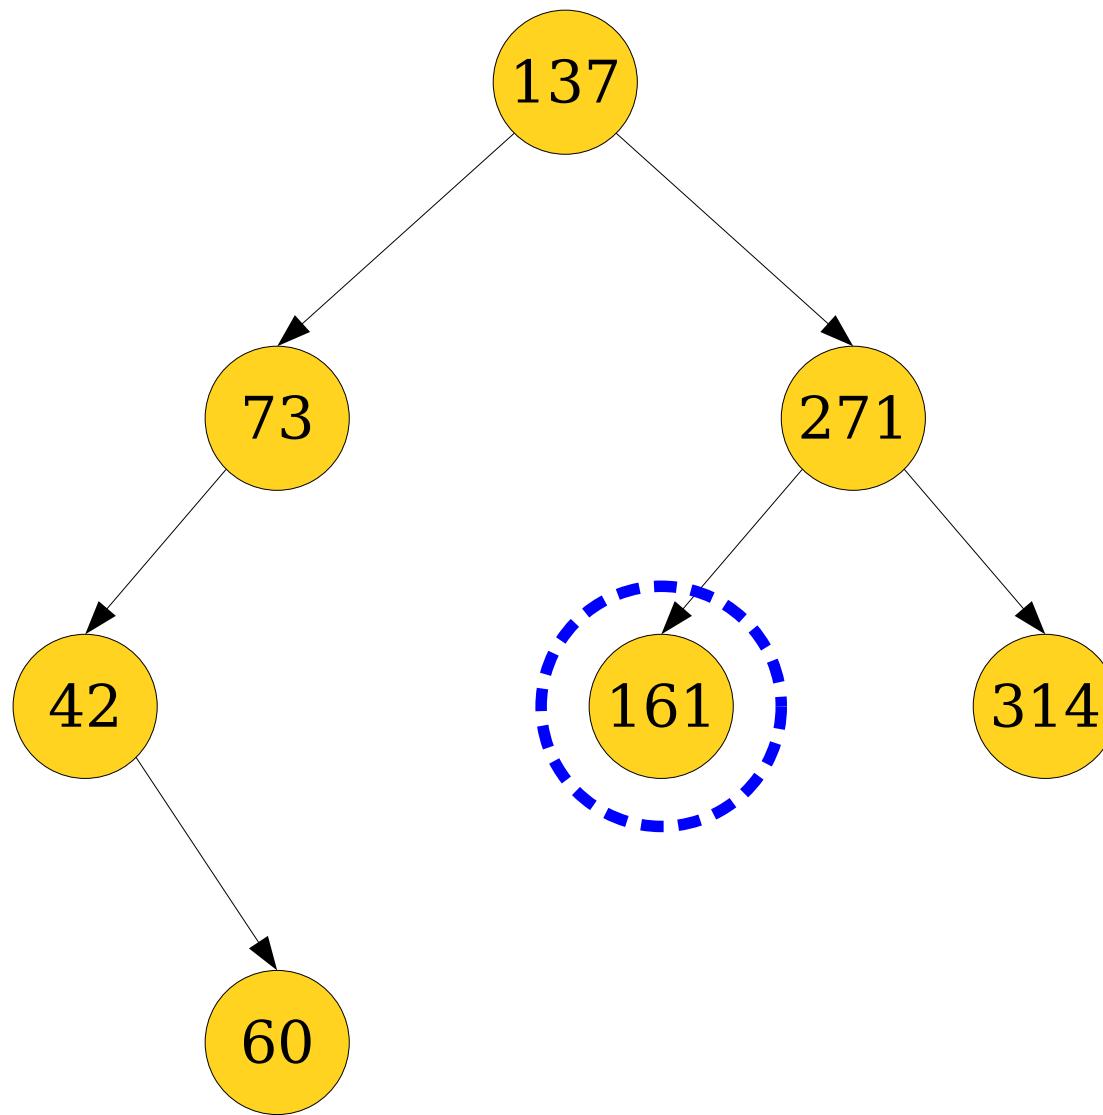
Inserting into a BST



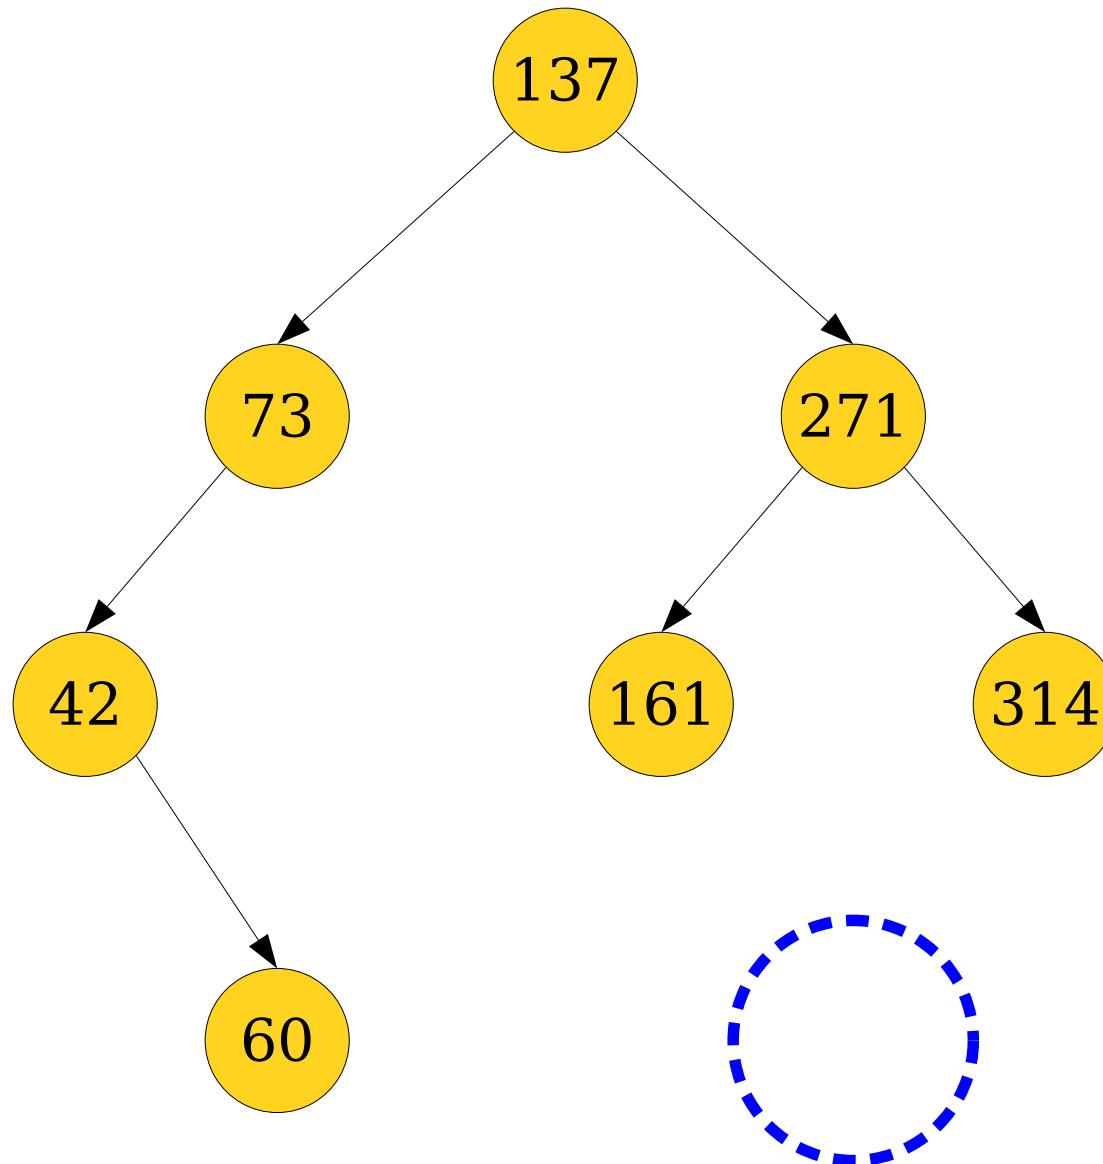
Inserting into a BST



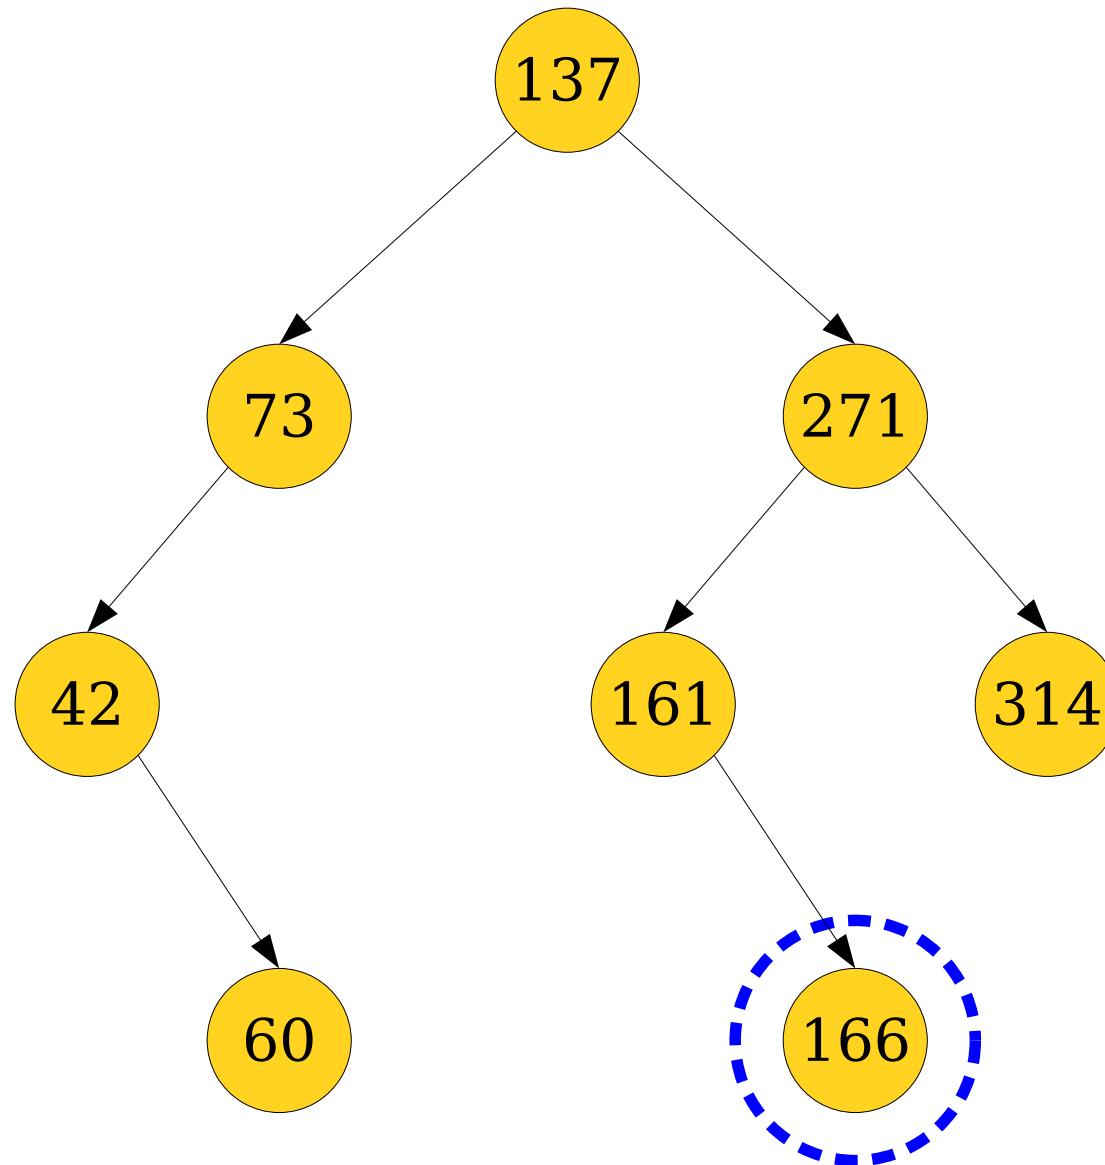
Inserting into a BST



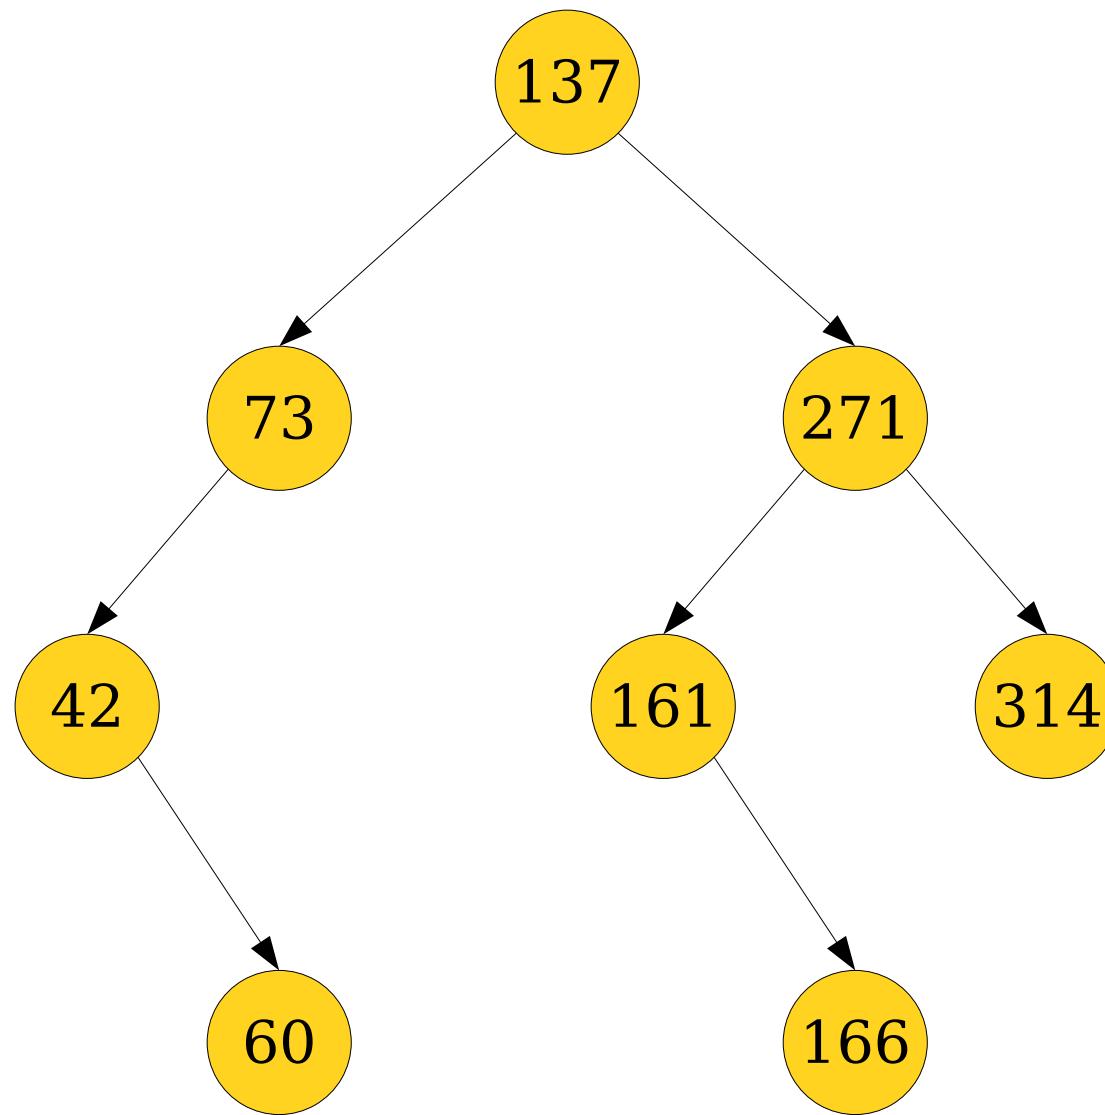
Inserting into a BST



Inserting into a BST

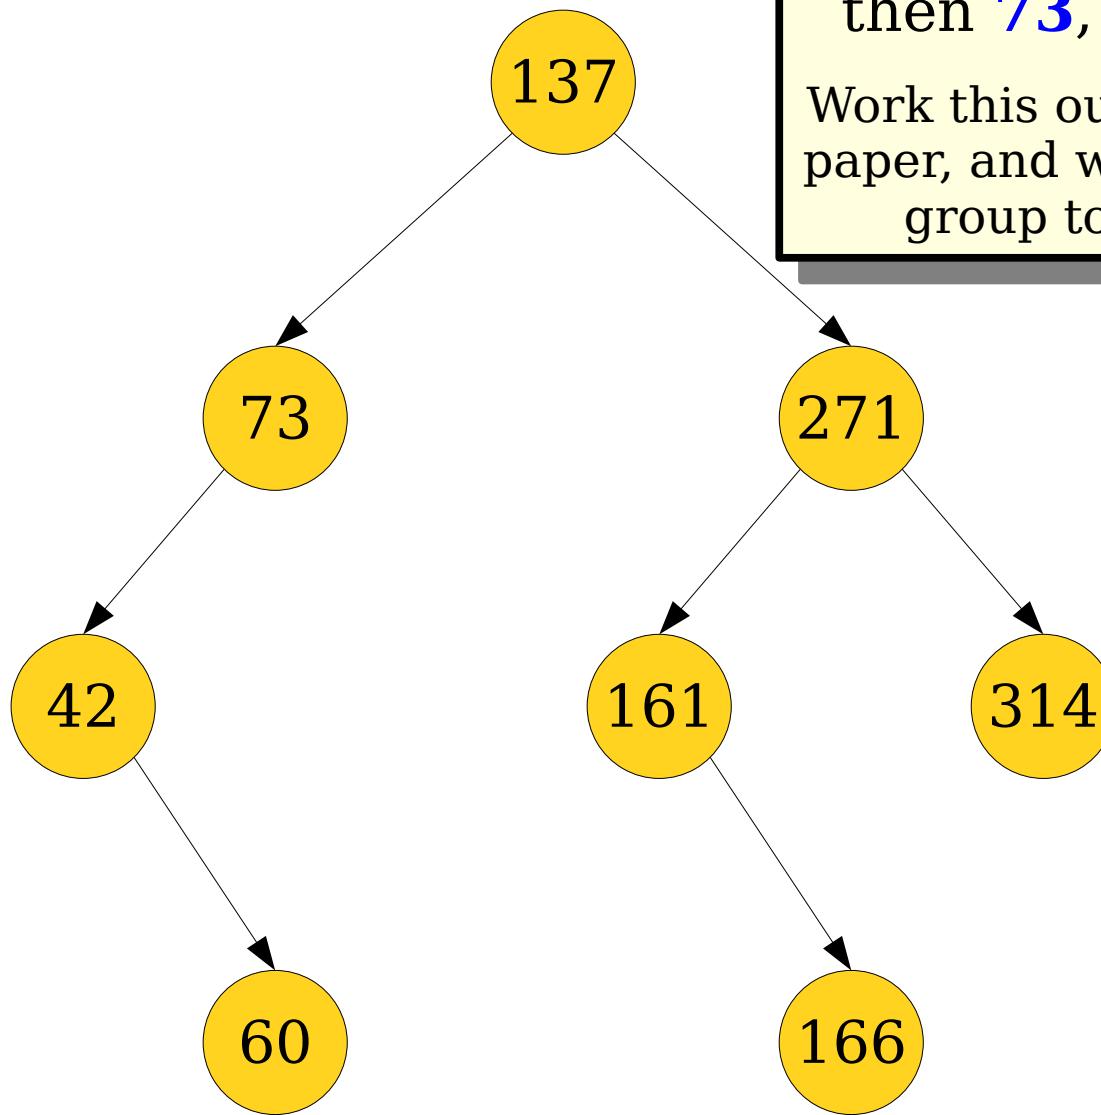


Inserting into a BST



Deleting from a BST

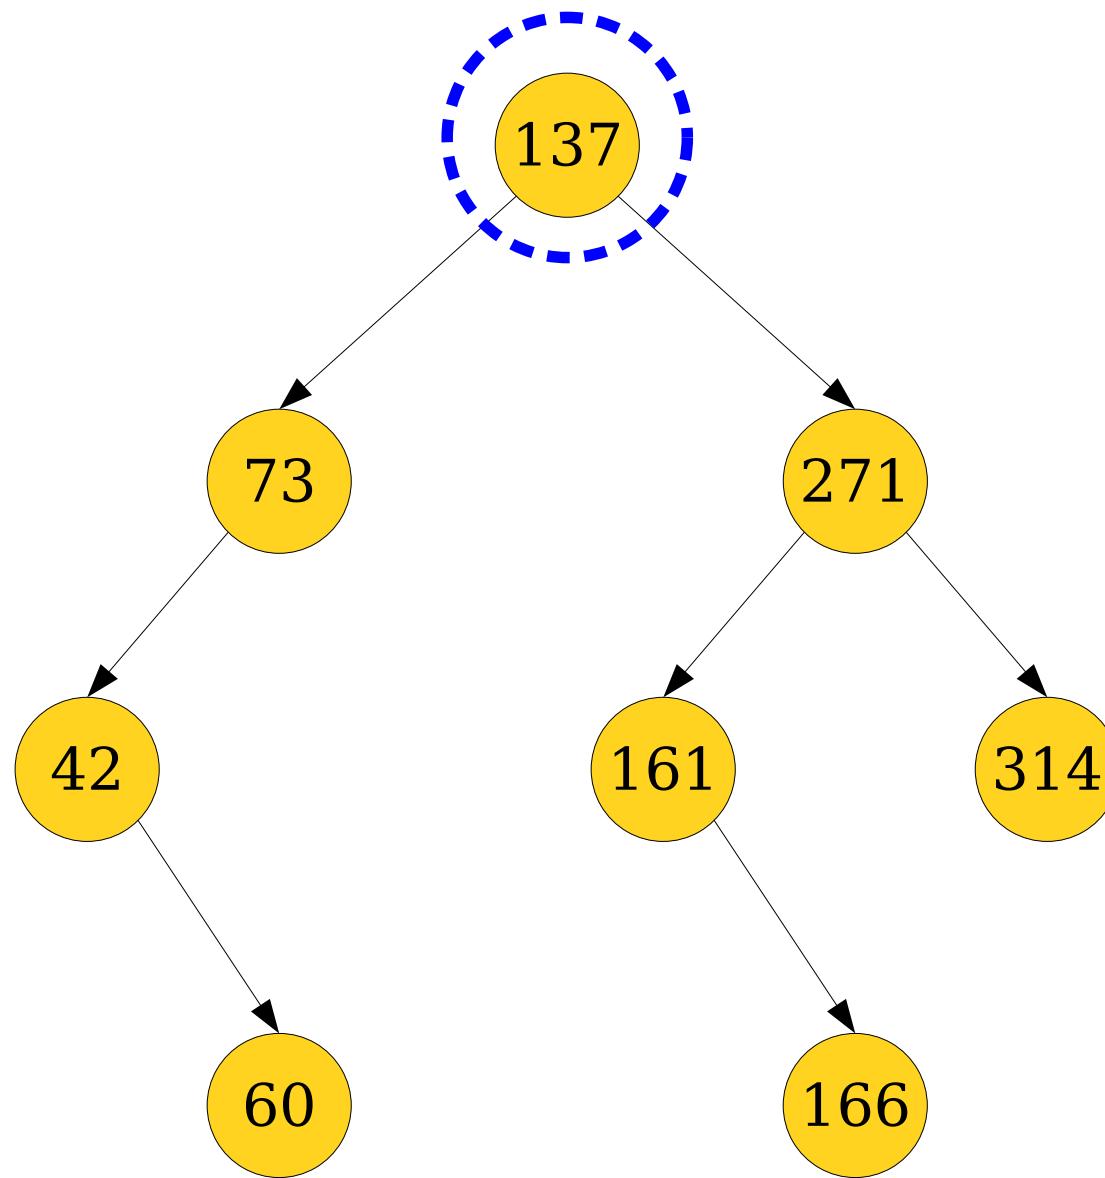
Deleting from a BST



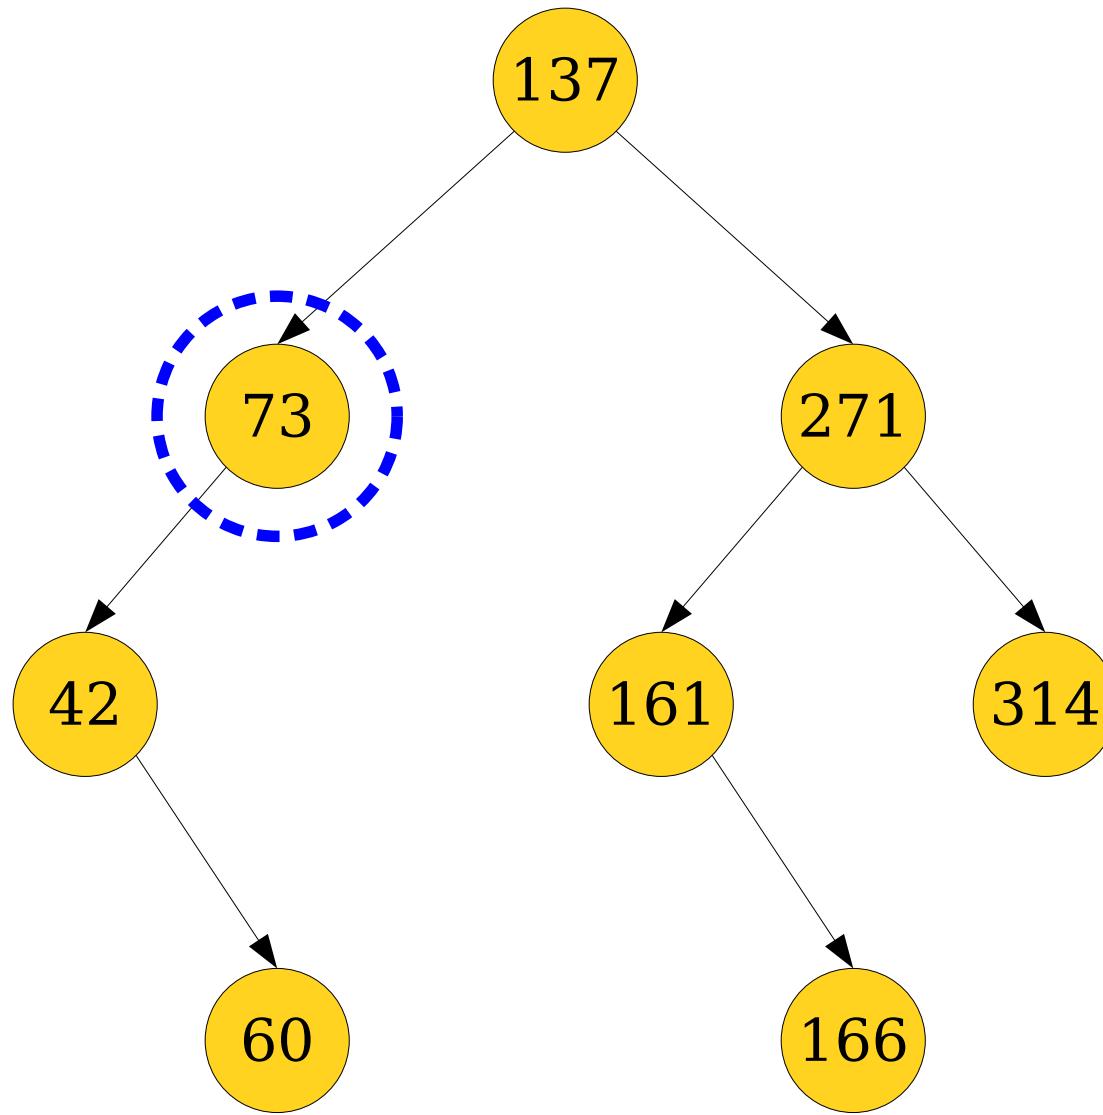
Delete **60** from this tree, then **73**, and then **137**.

Work this out with a pencil and paper, and we'll reconvene as a group to do it together.

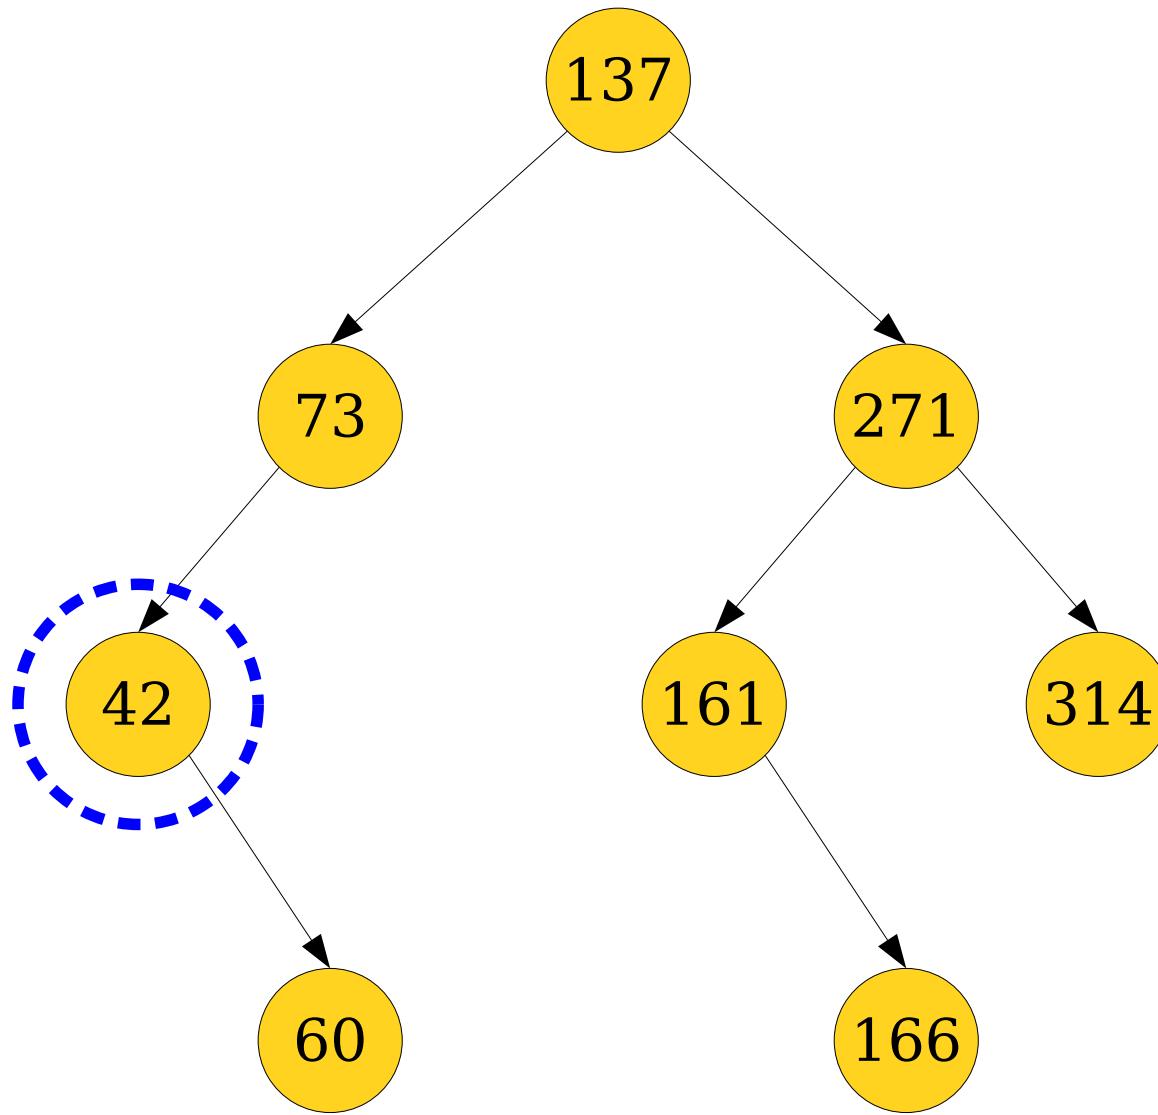
Deleting from a BST



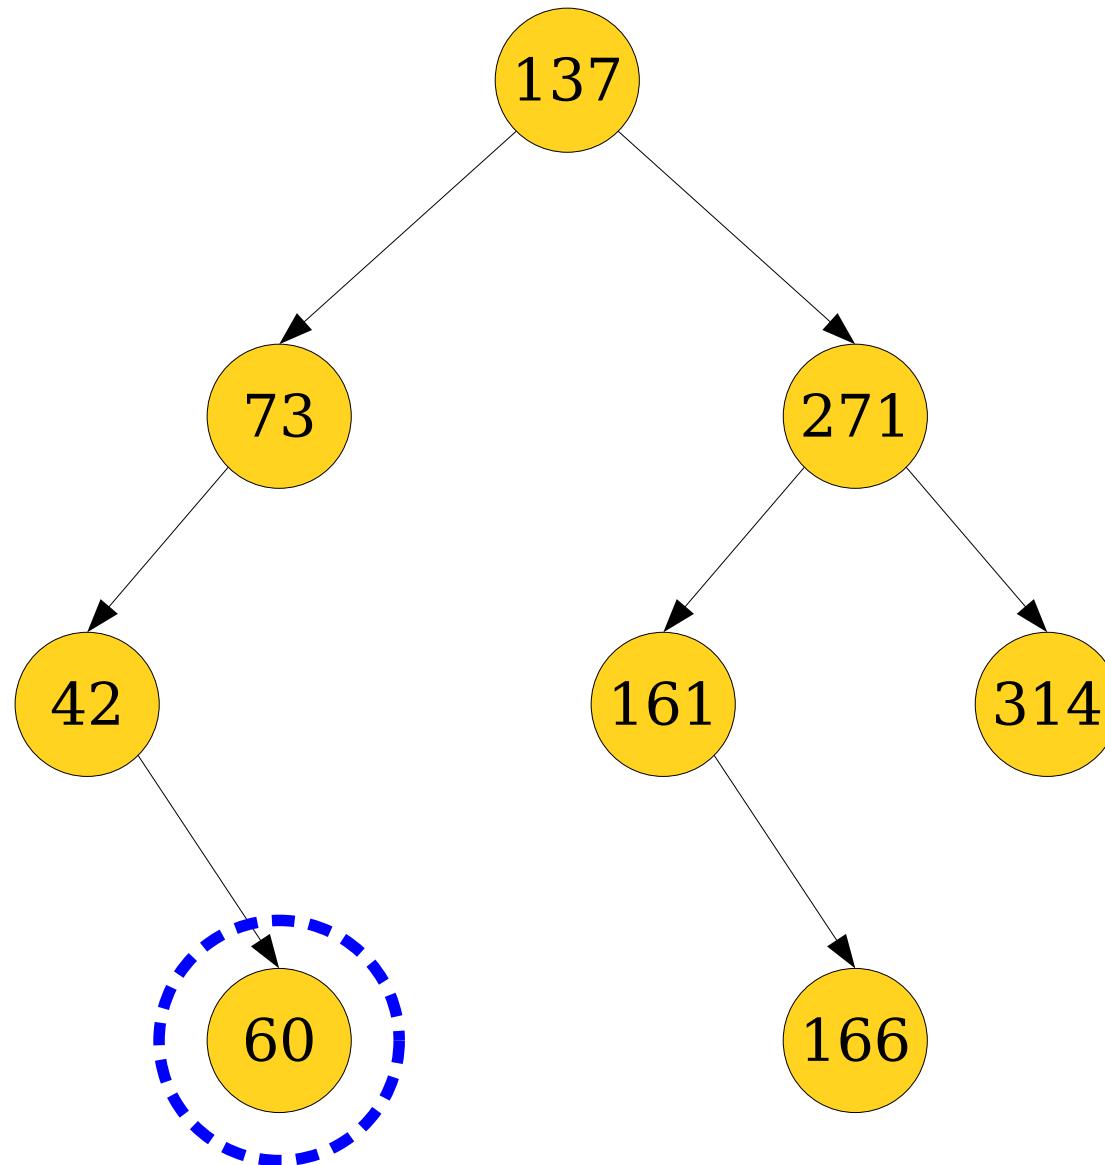
Deleting from a BST



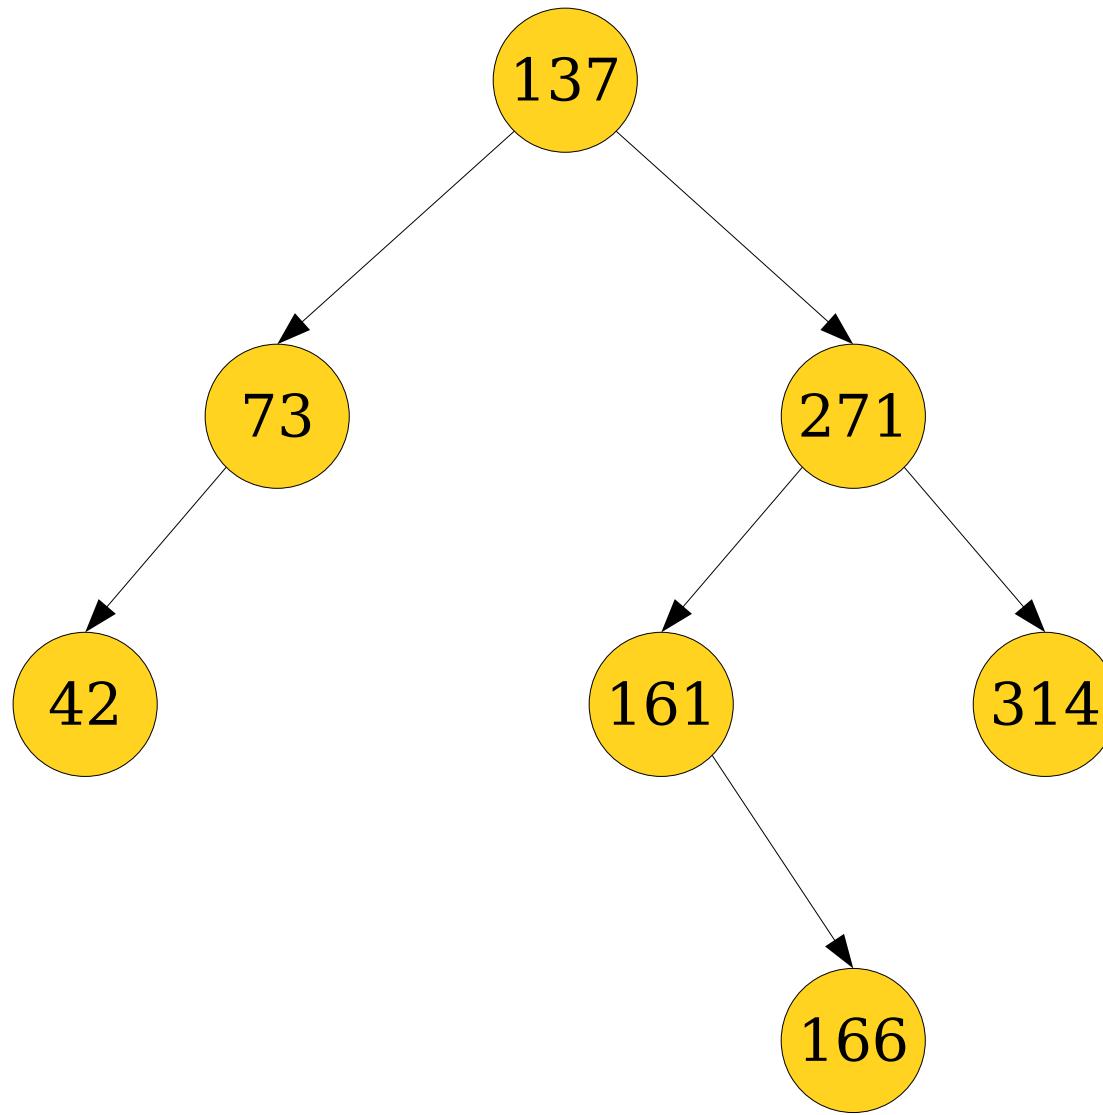
Deleting from a BST



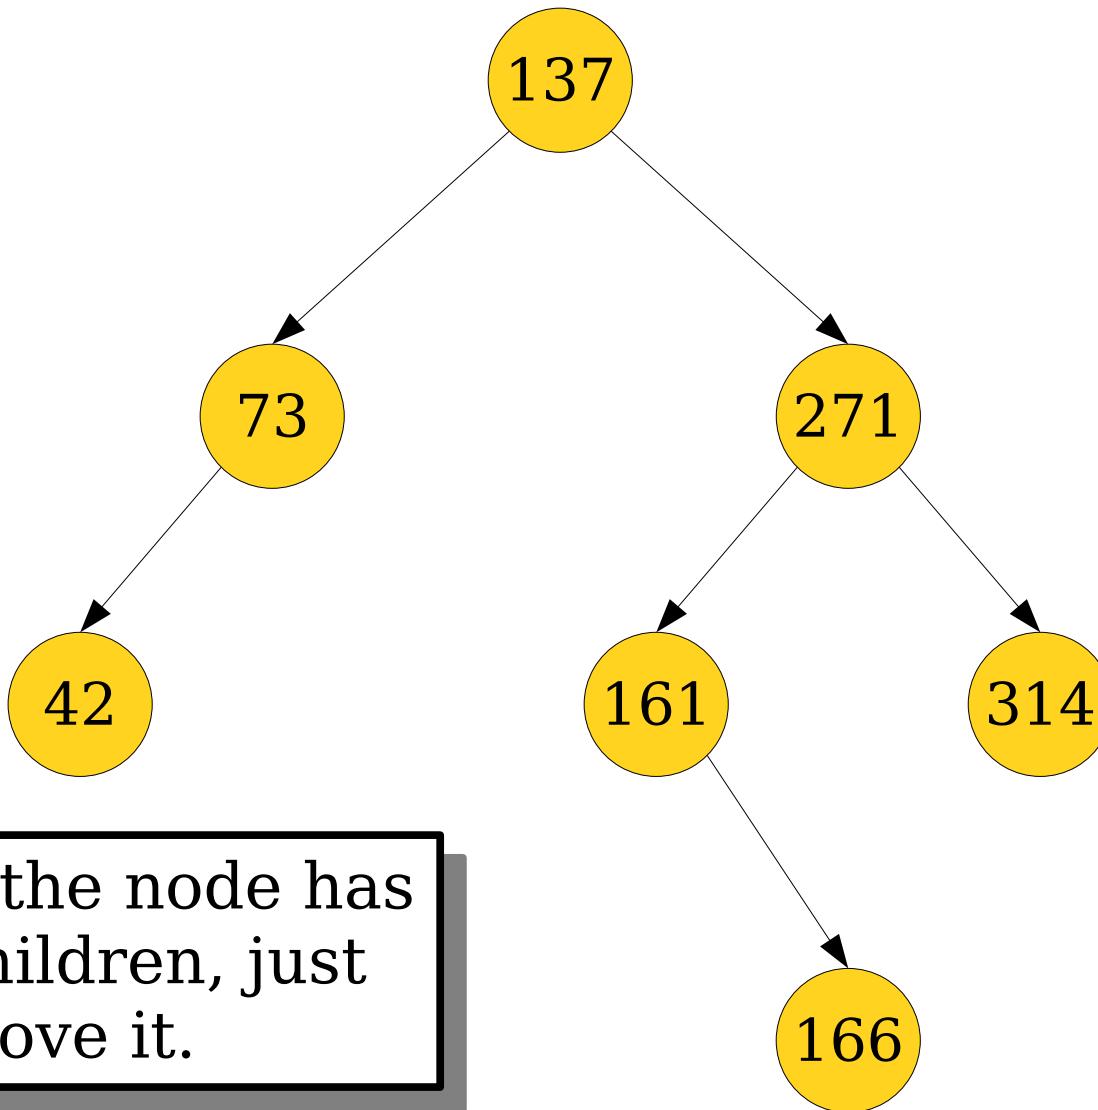
Deleting from a BST



Deleting from a BST

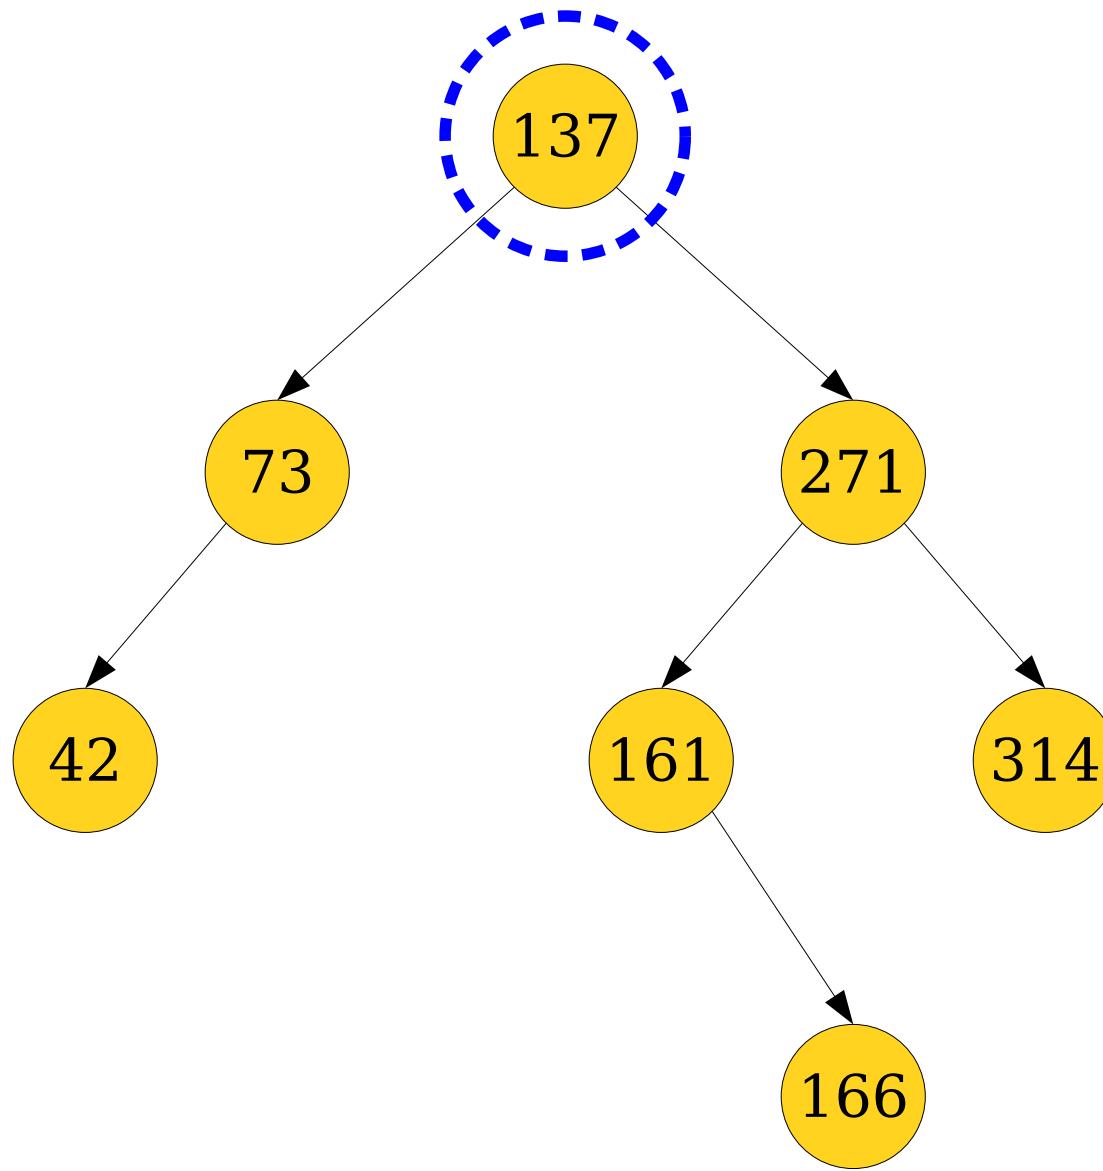


Deleting from a BST

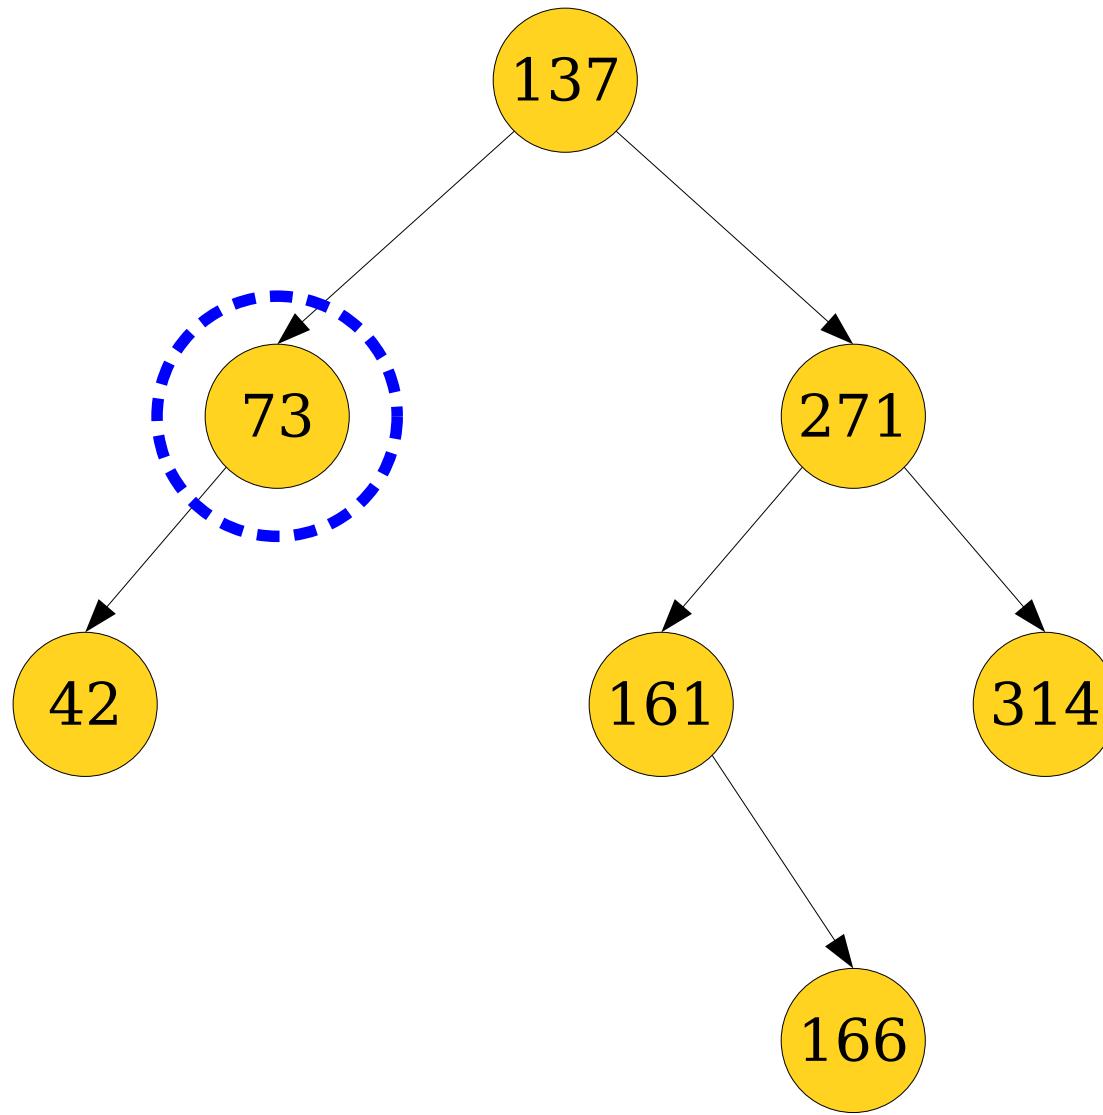


Case 0: If the node has just no children, just remove it.

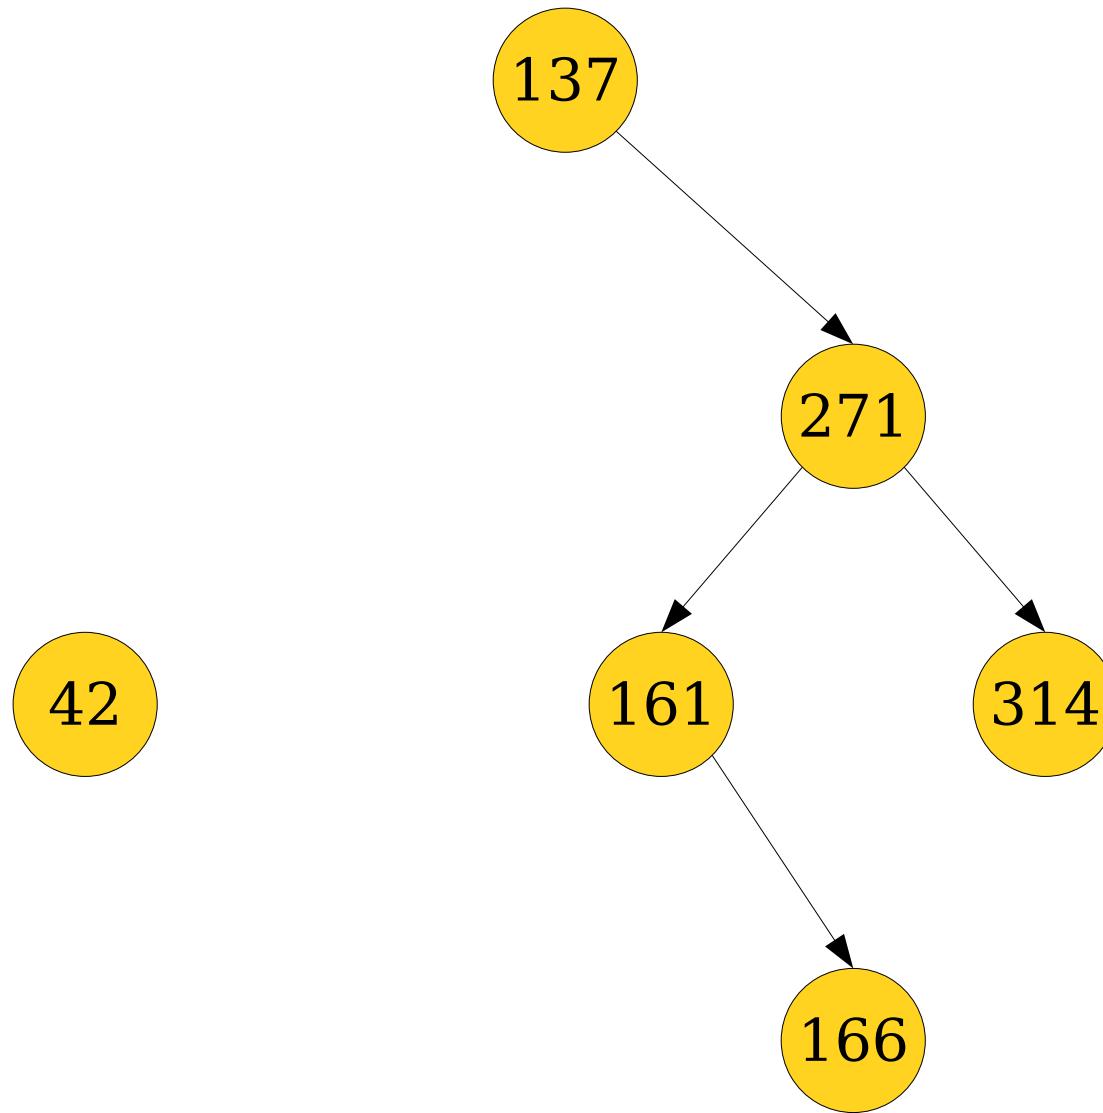
Deleting from a BST



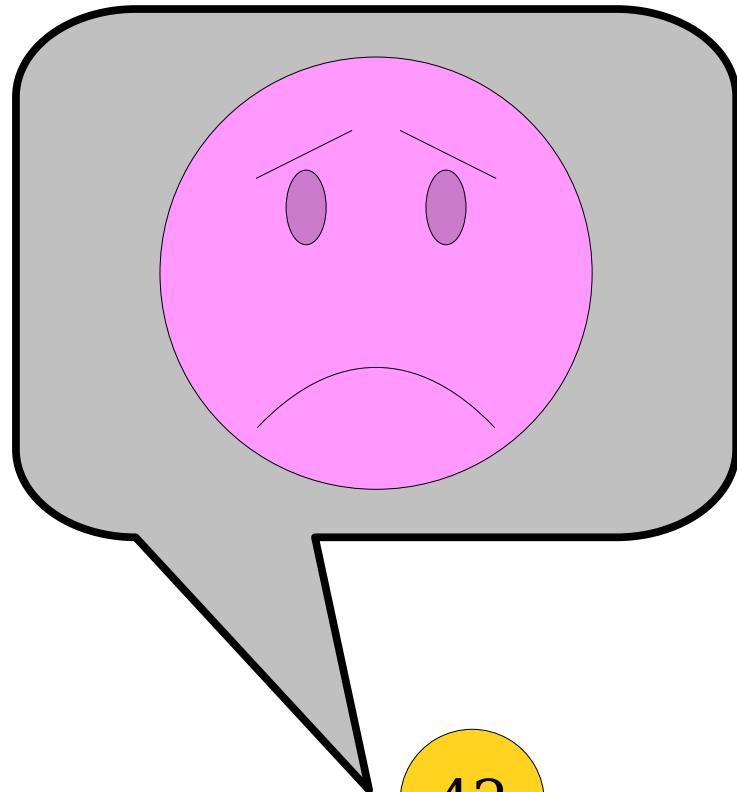
Deleting from a BST



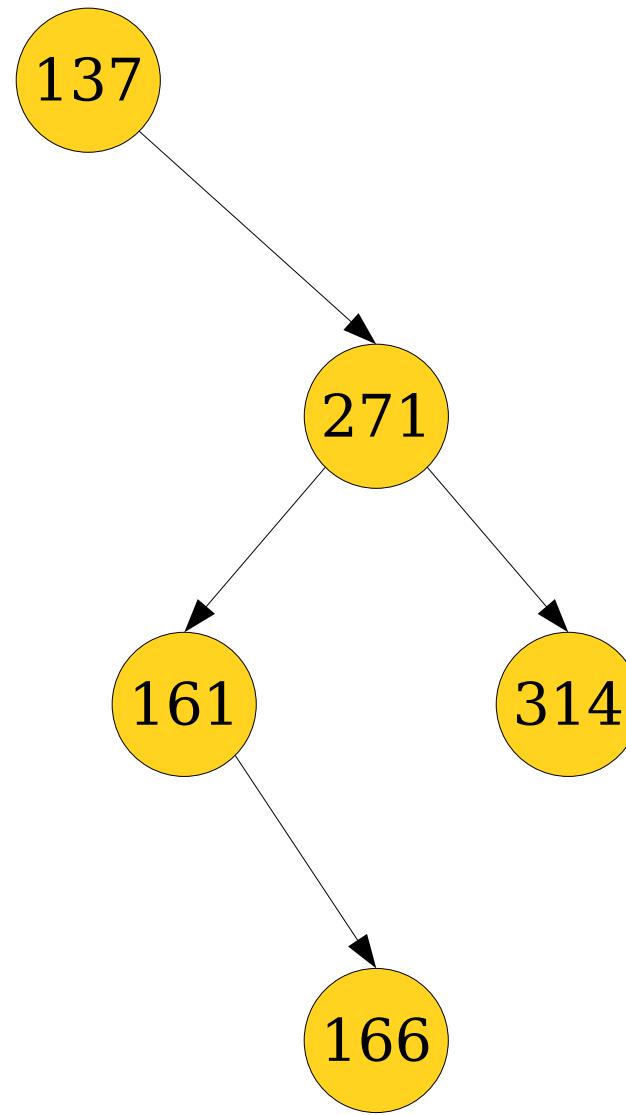
Deleting from a BST



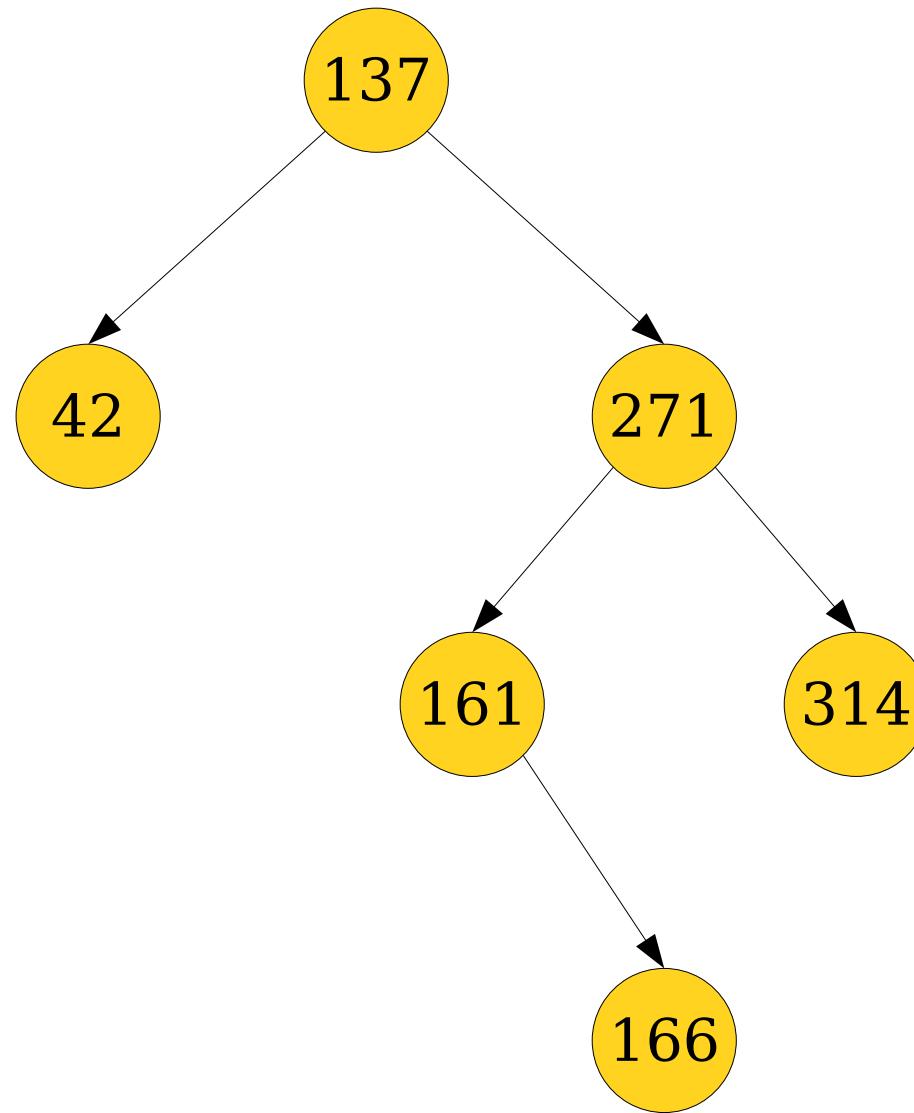
Deleting from a BST



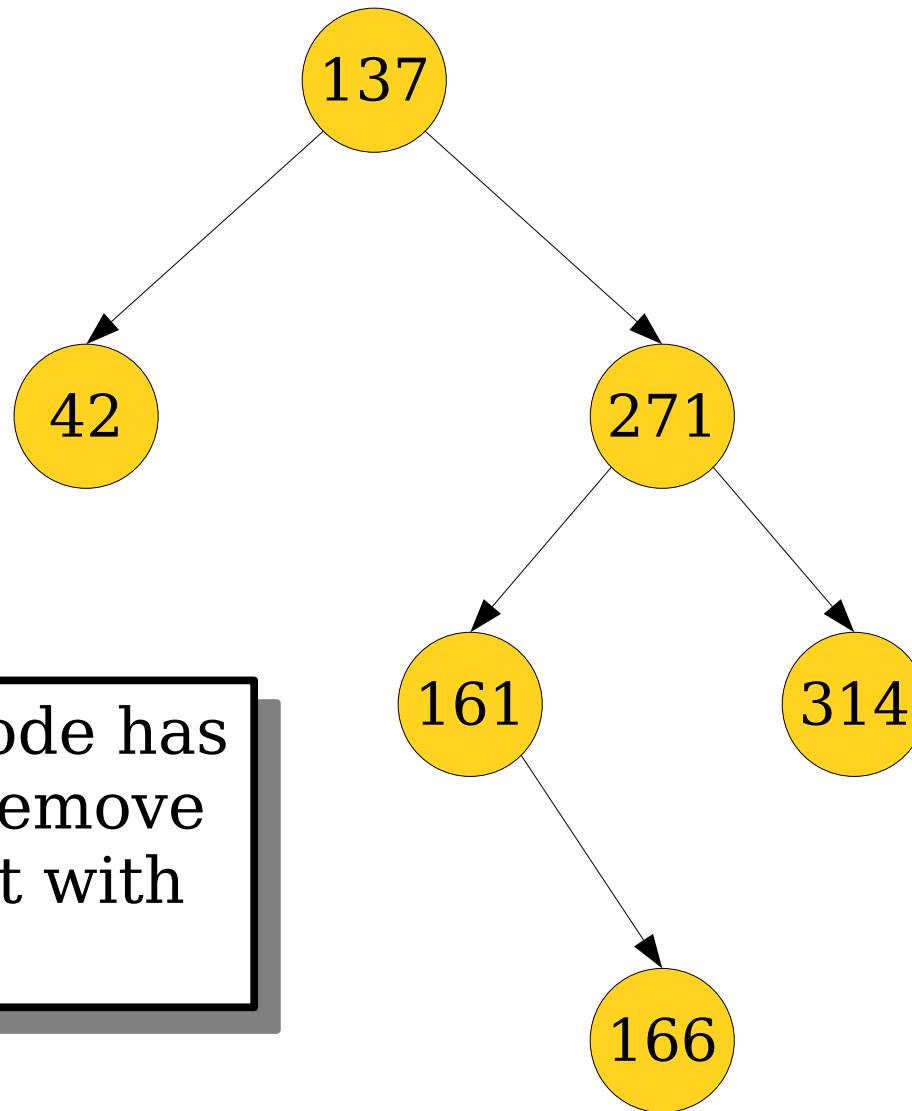
42



Deleting from a BST

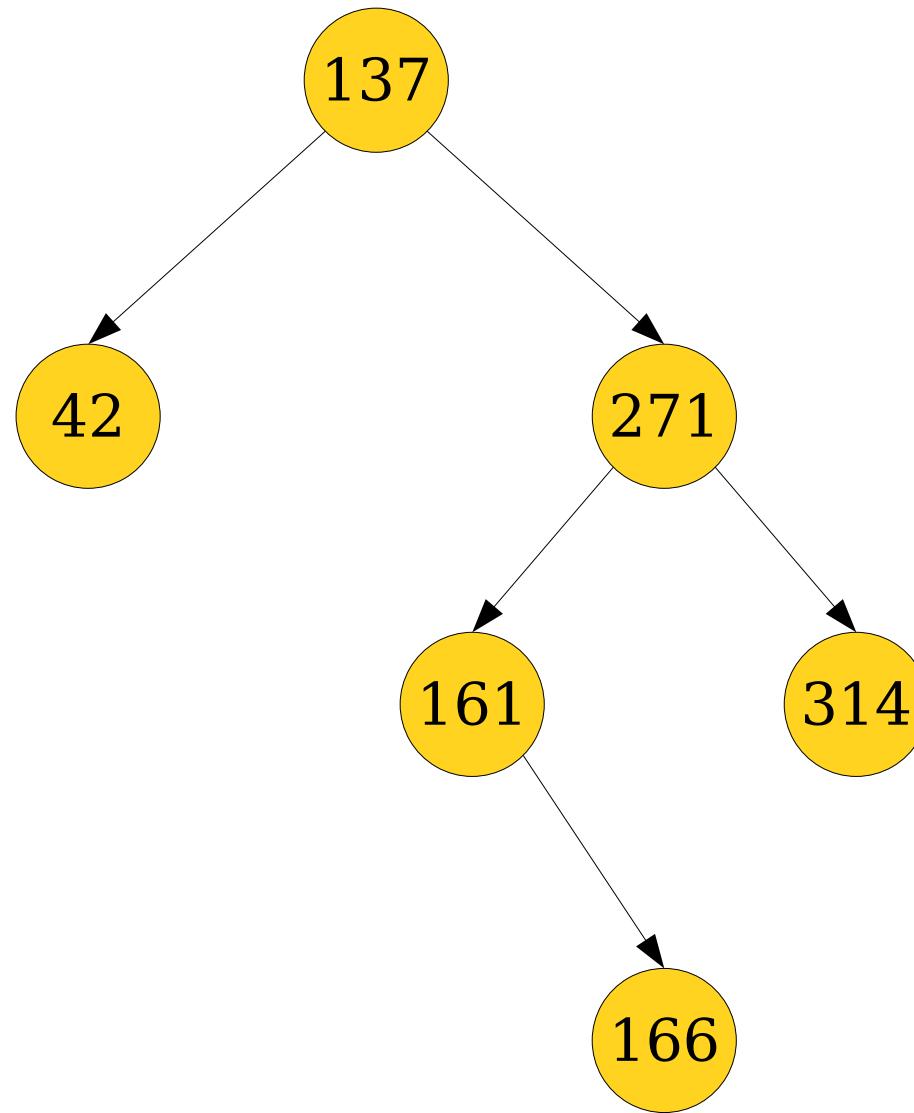


Deleting from a BST

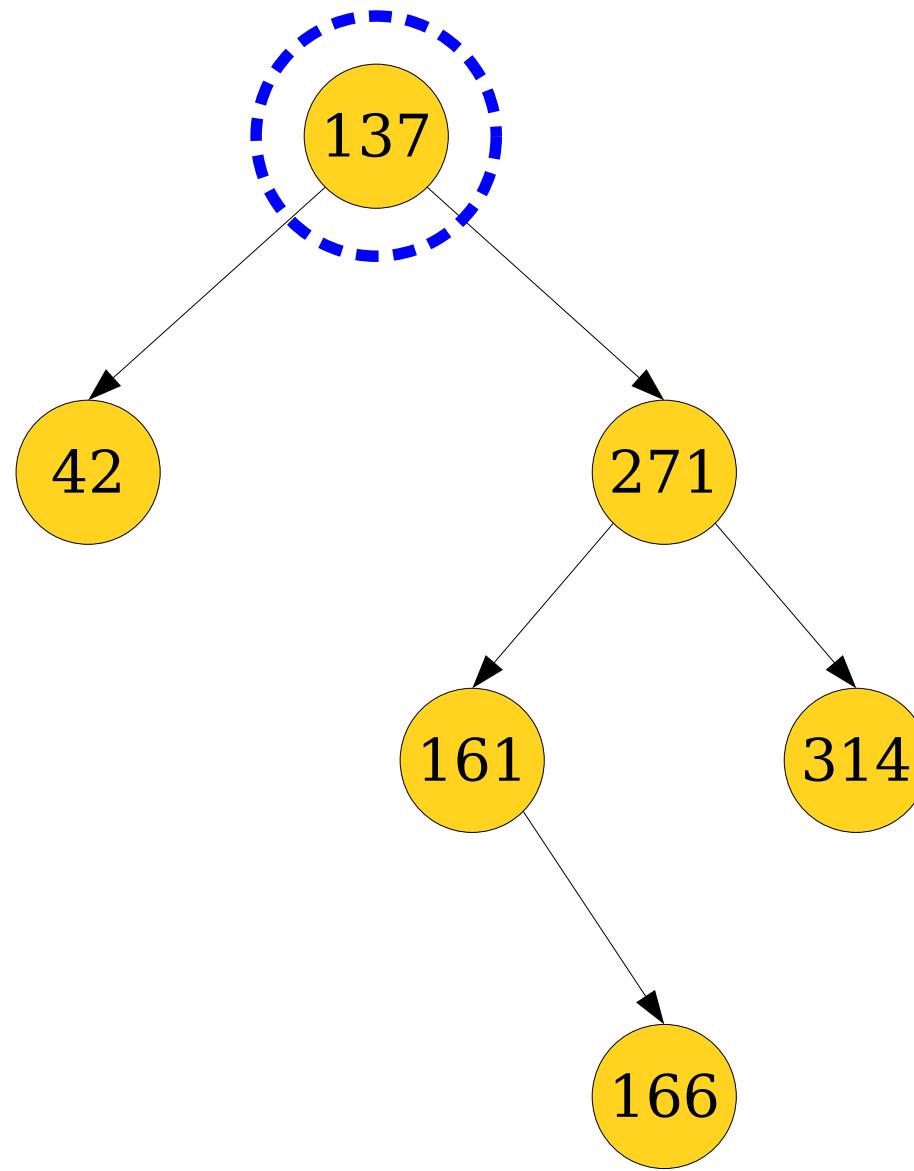


Case 1: If the node has just one child, remove it and replace it with its child.

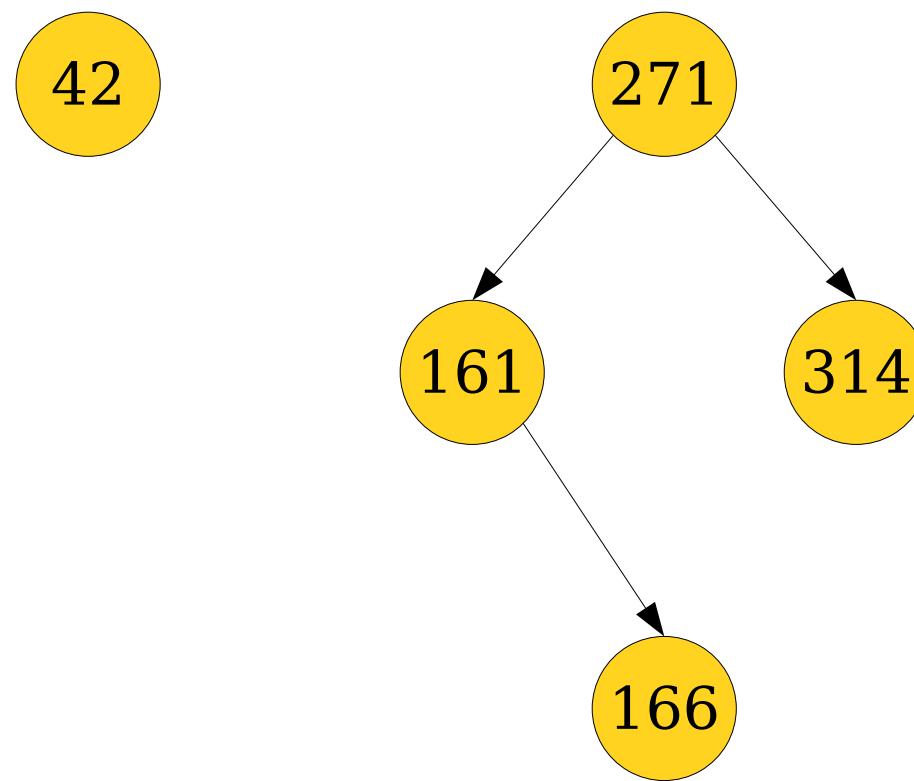
Deleting from a BST



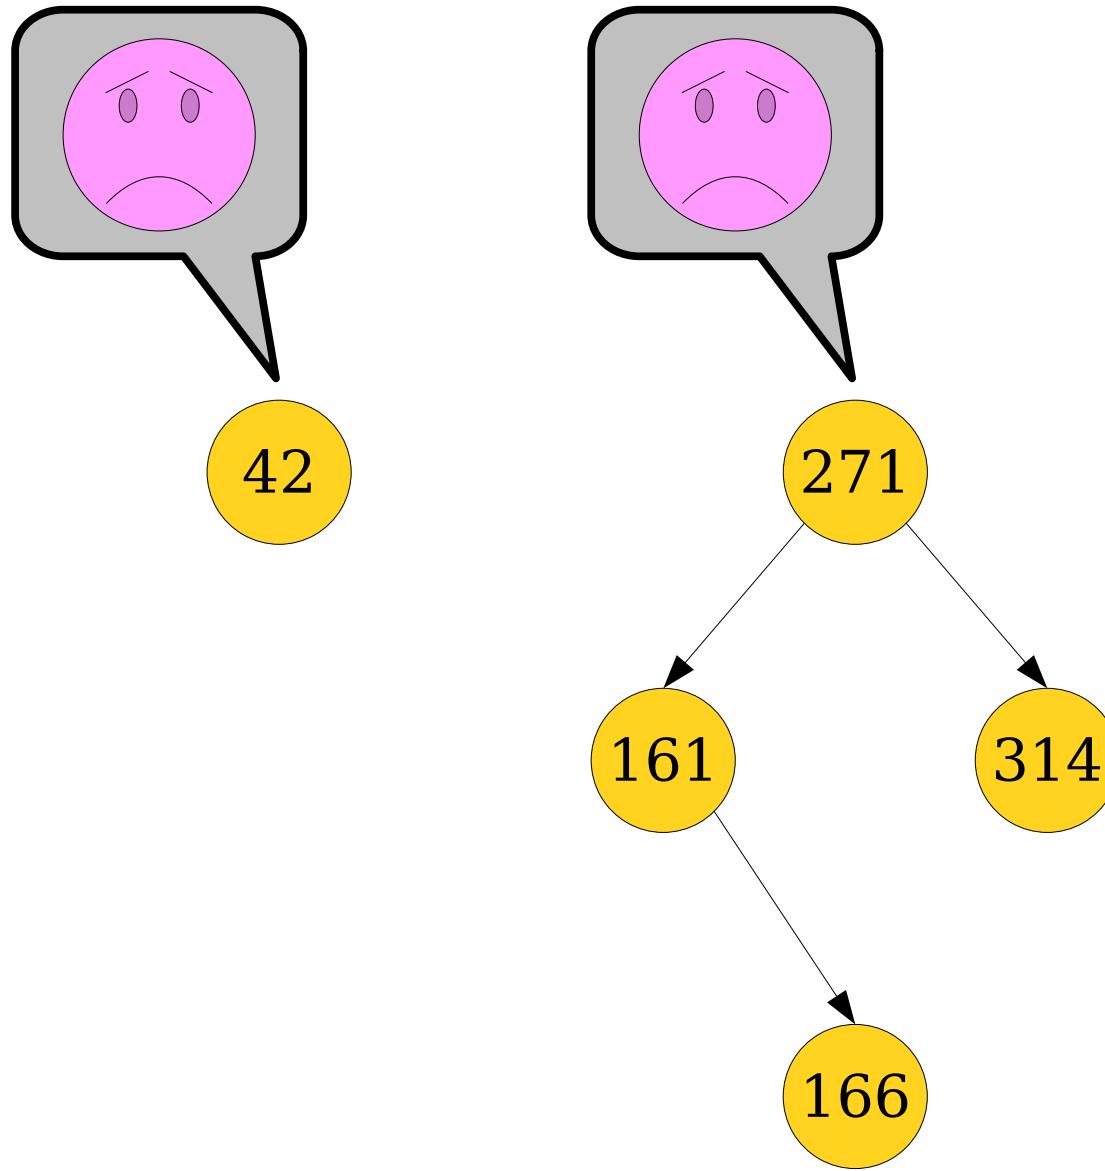
Deleting from a BST



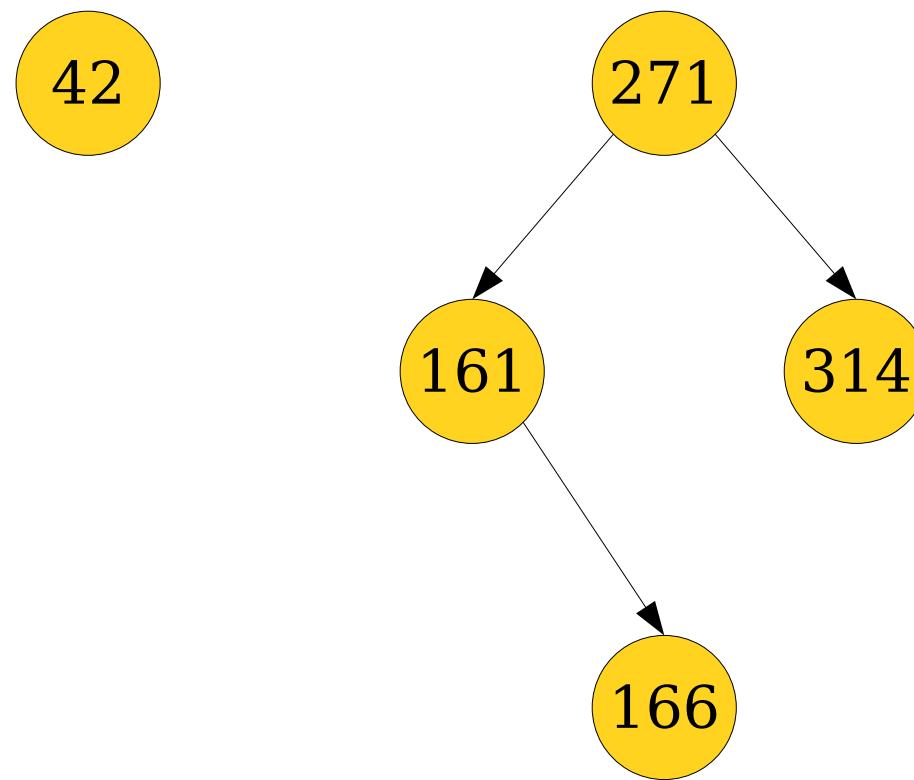
Deleting from a BST



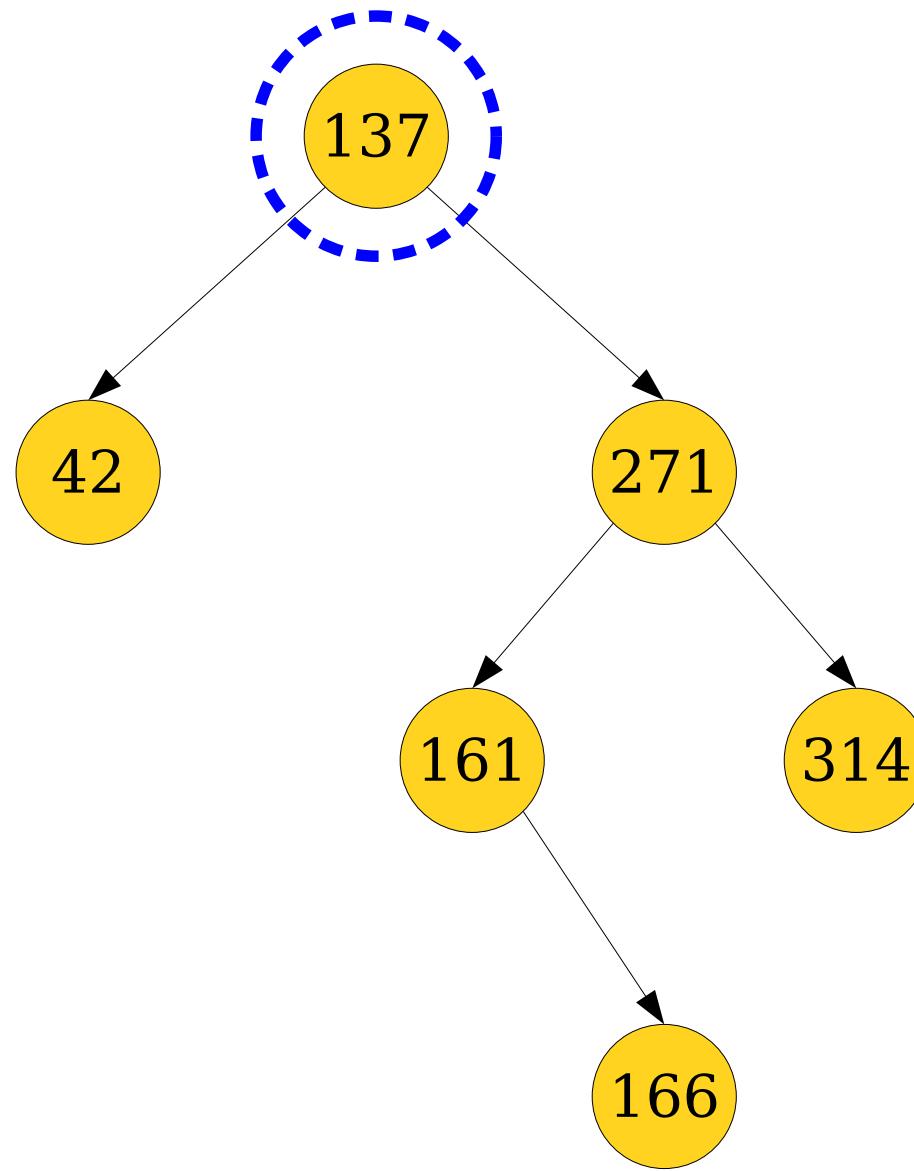
Deleting from a BST



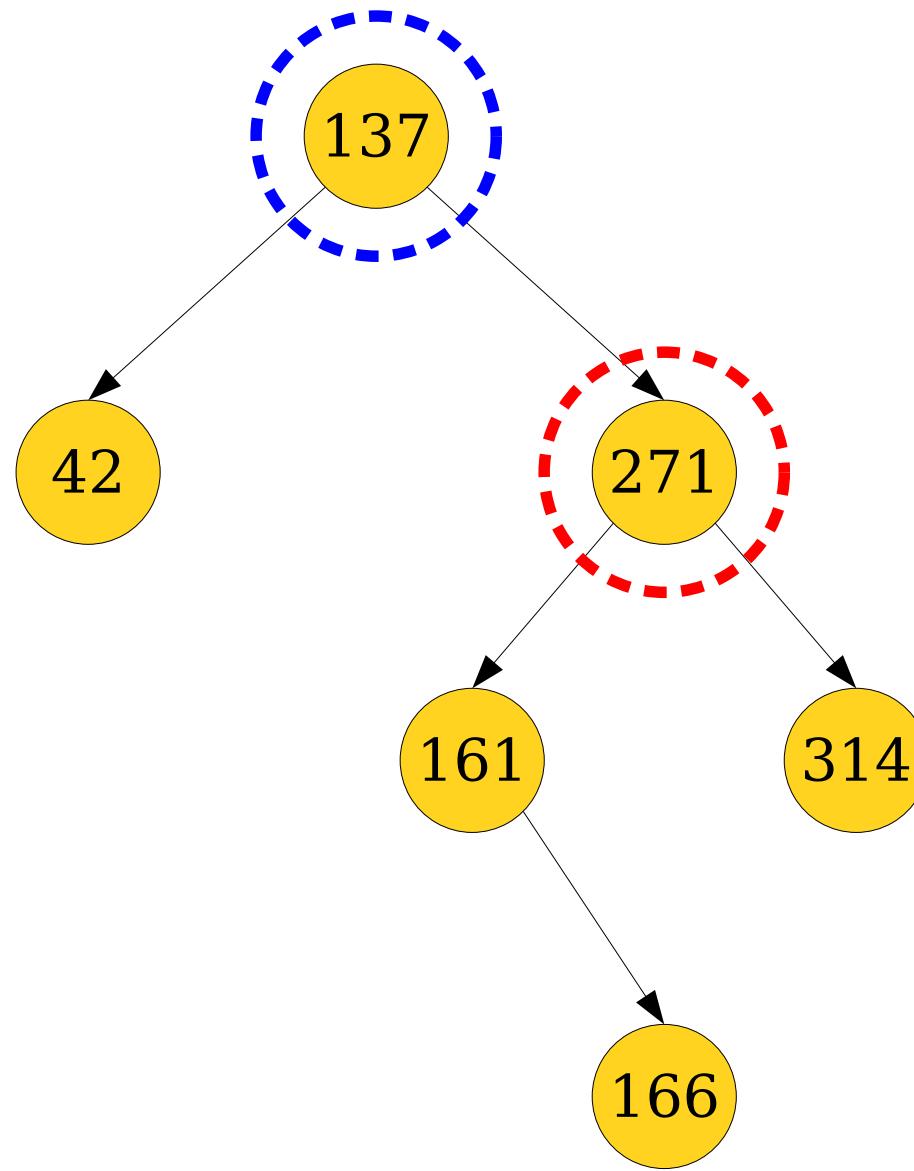
Deleting from a BST



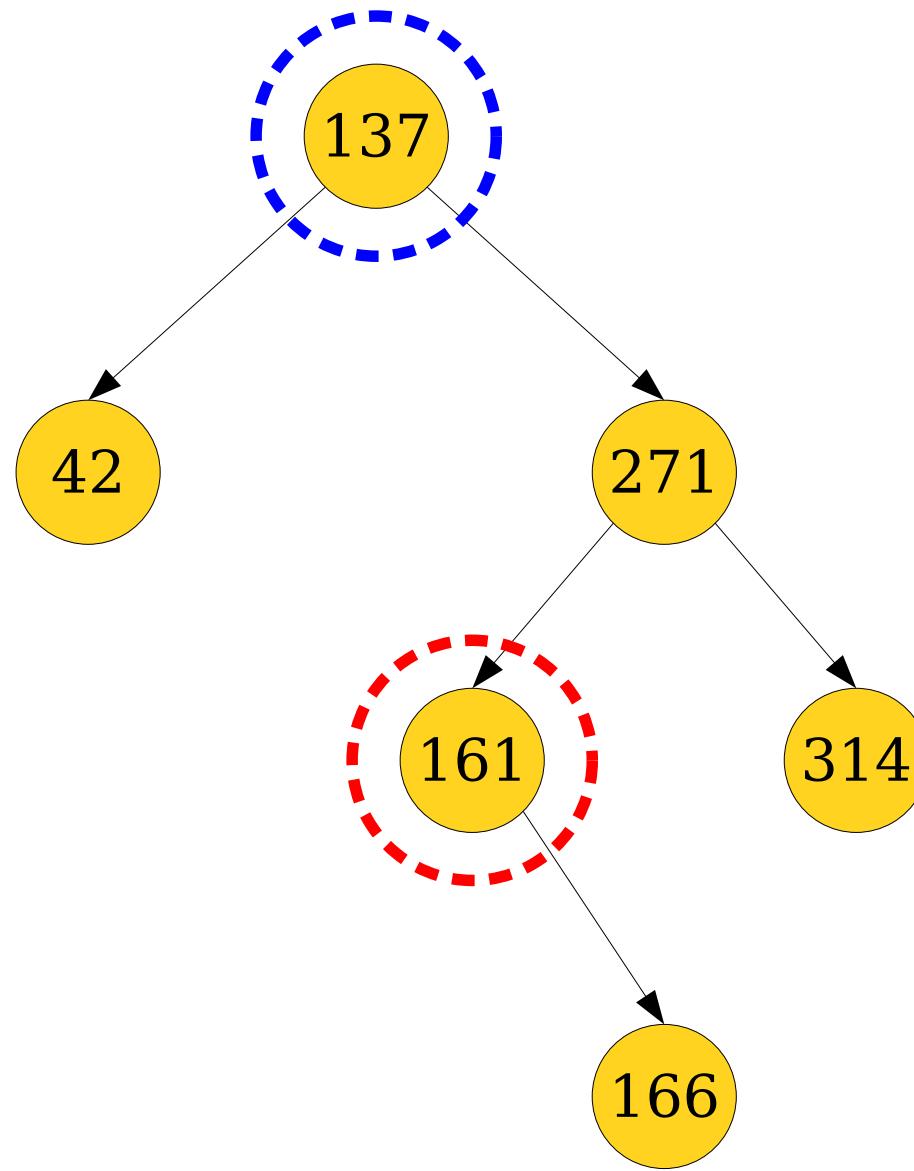
Deleting from a BST



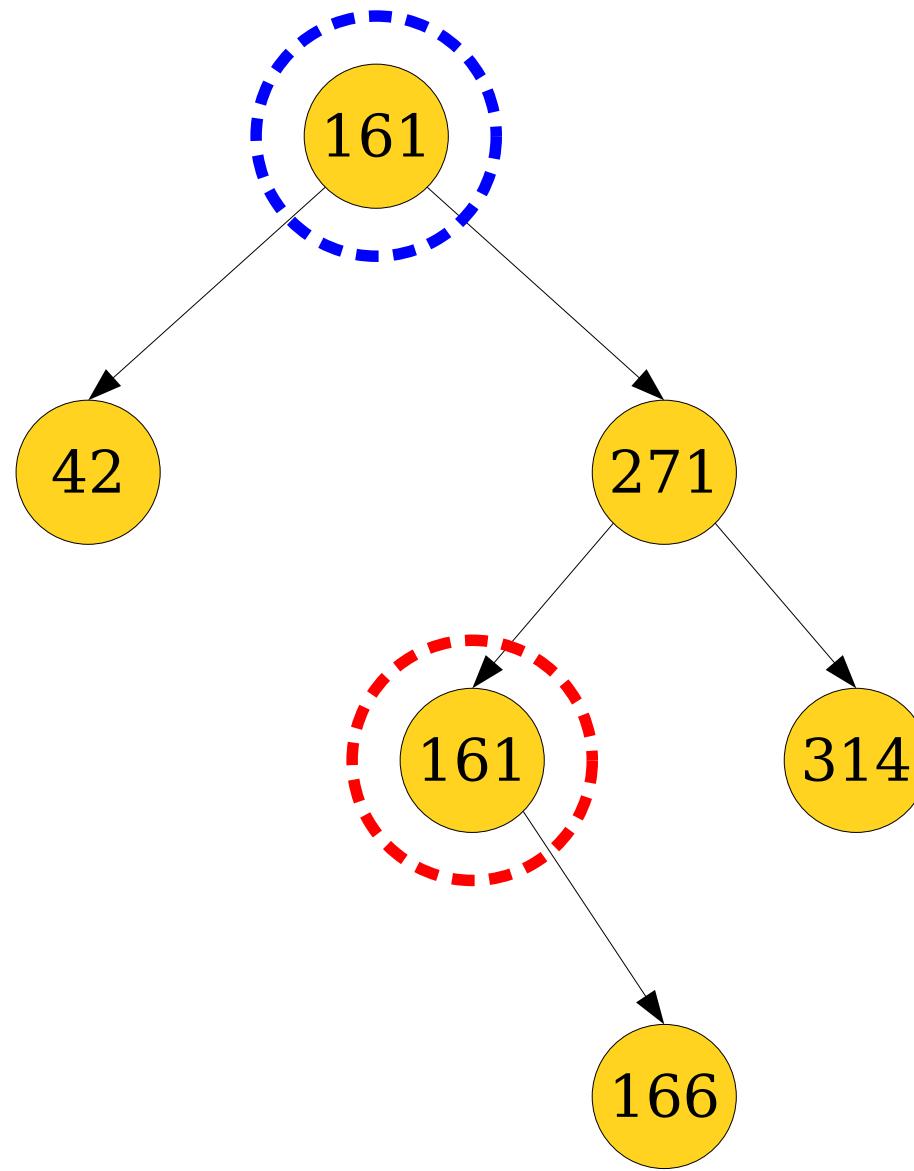
Deleting from a BST



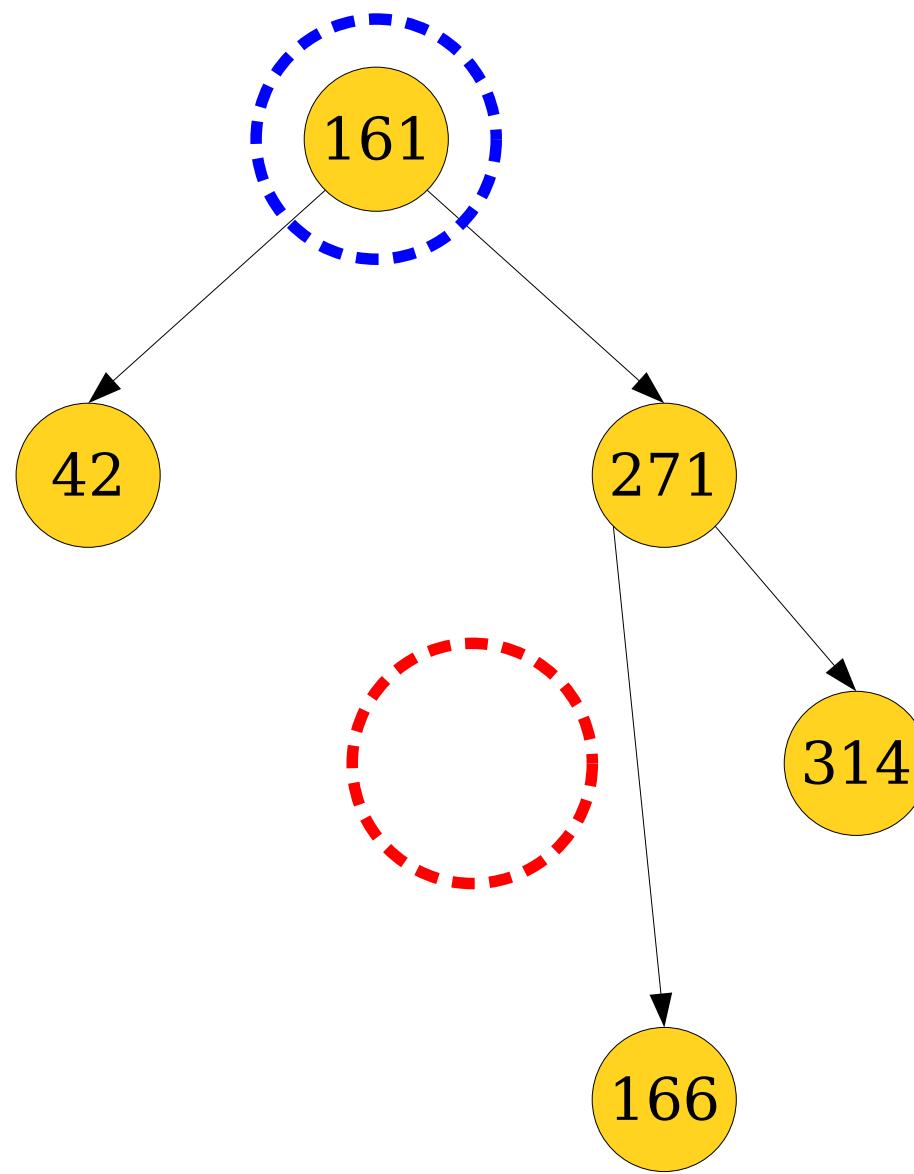
Deleting from a BST



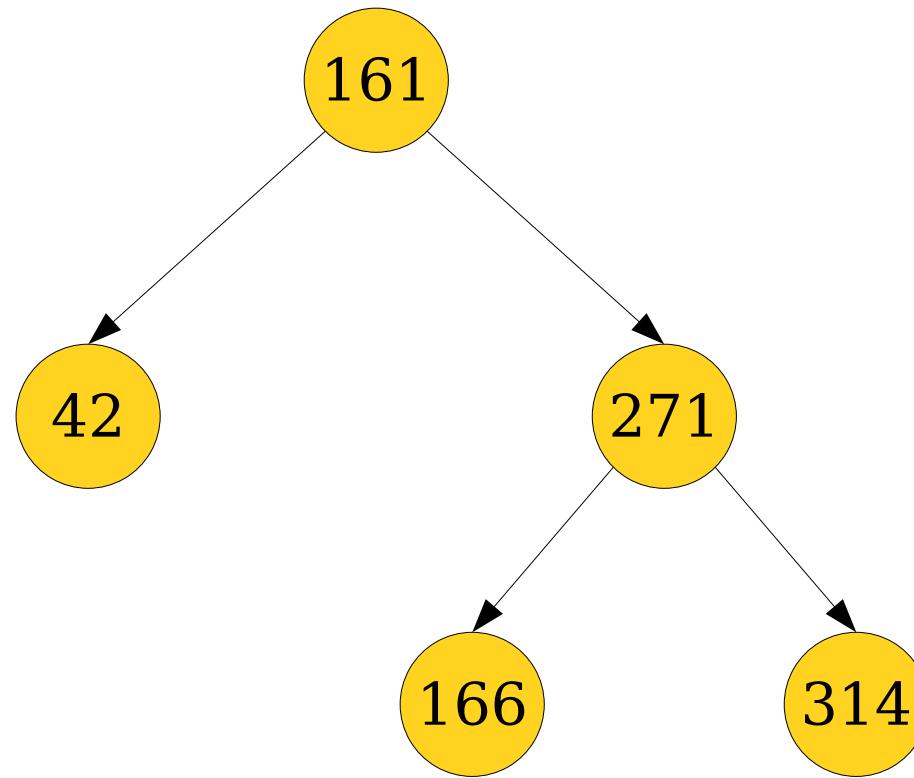
Deleting from a BST



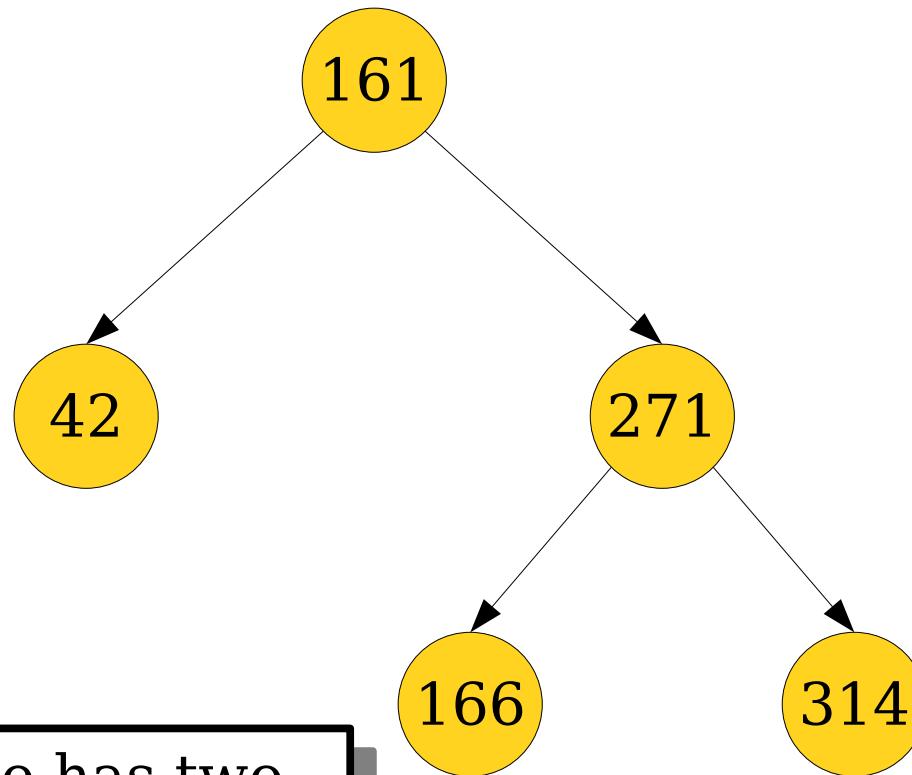
Deleting from a BST



Deleting from a BST



Deleting from a BST



Case 2: If the node has two children, find its inorder successor (which has zero or one child), replace the node's key with its successor's key, then delete its successor.

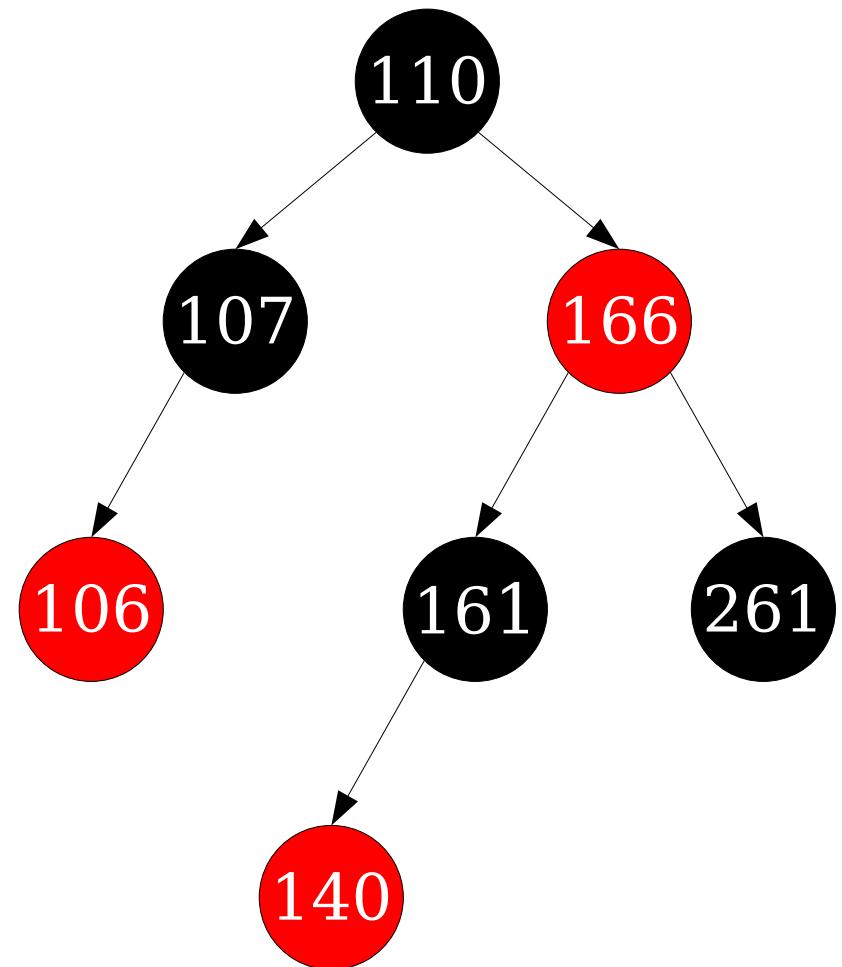
Runtime Analysis

- The time complexity of all these operations is $O(h)$, where h is the height of the tree.
 - That's the longest path we can take.
- In the best case, $h = O(\log n)$ and all operations take time $O(\log n)$.
- In the worst case, $h = \Theta(n)$ and some operations will take time $\Theta(n)$.
- ***Challenge:*** How do you efficiently keep the height of a tree low?

A Glimpse of Red/Black Trees

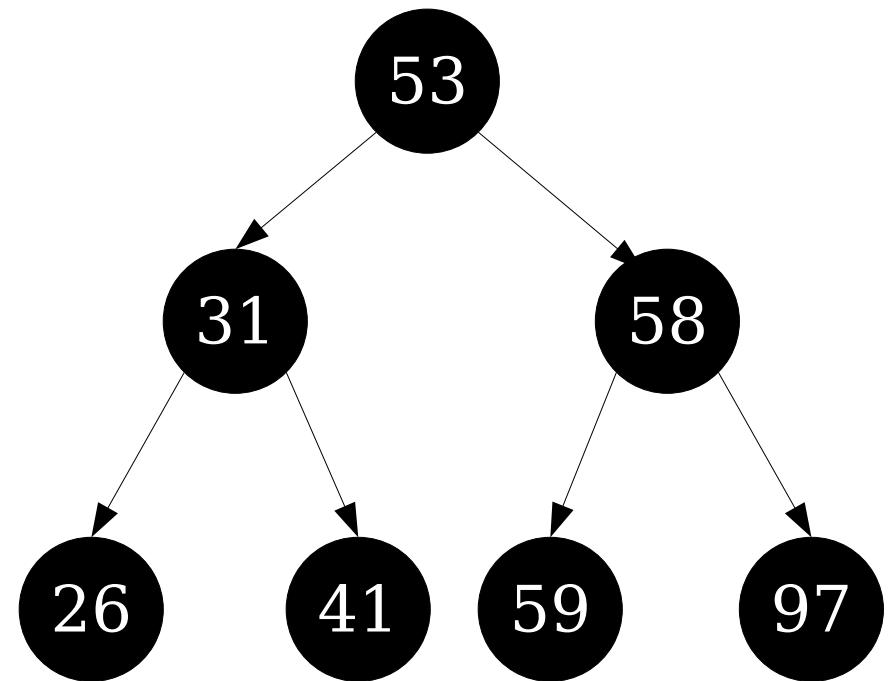
Red/Black Trees

- A ***red/black tree*** is a BST with the following properties:
 - Every node is either red or black.
 - The root is black.
 - No red node has a red child.
 - Every root-null path in the tree passes through the same number of black nodes.



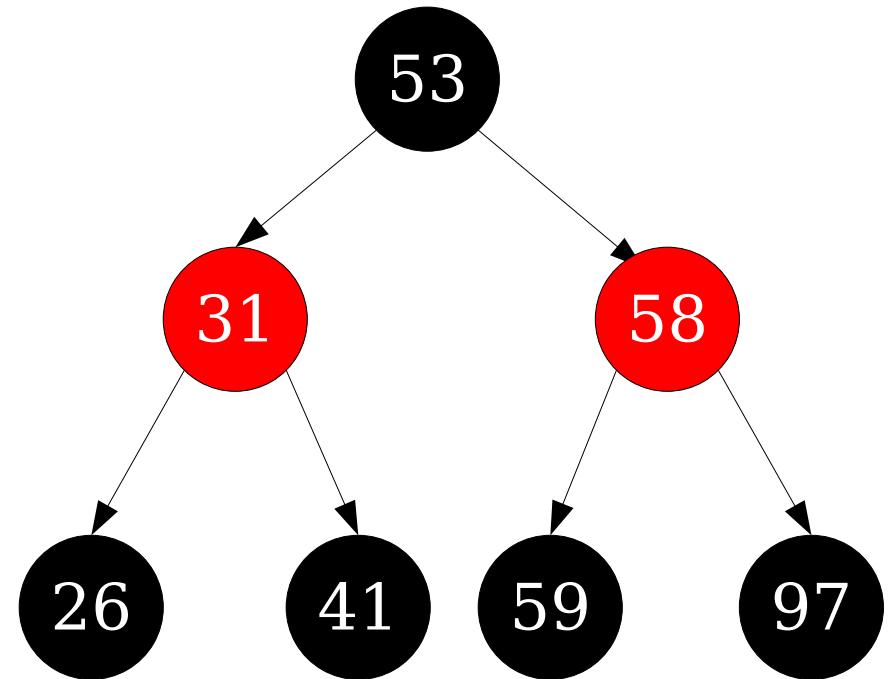
Red/Black Trees

- A ***red/black tree*** is a BST with the following properties:
 - Every node is either red or black.
 - The root is black.
 - No red node has a red child.
 - Every root-null path in the tree passes through the same number of black nodes.



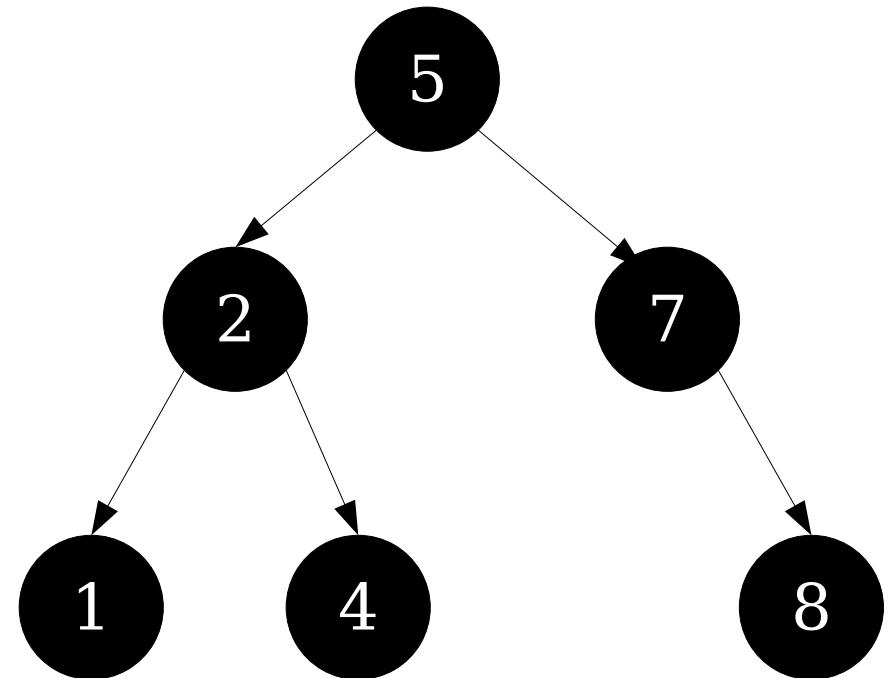
Red/Black Trees

- A ***red/black tree*** is a BST with the following properties:
 - Every node is either red or black.
 - The root is black.
 - No red node has a red child.
 - Every root-null path in the tree passes through the same number of black nodes.



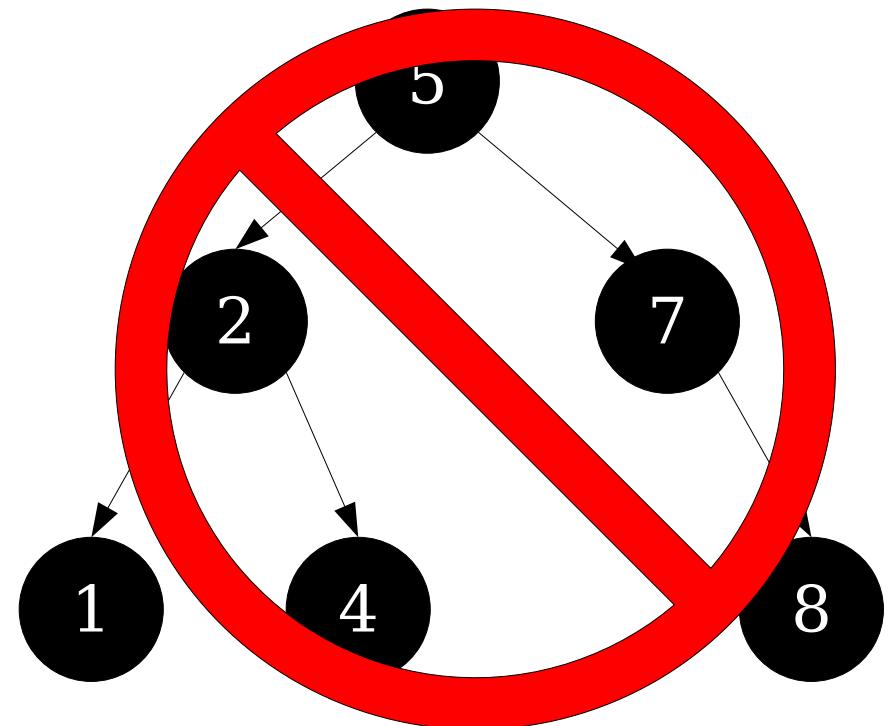
Red/Black Trees

- A ***red/black tree*** is a BST with the following properties:
 - Every node is either red or black.
 - The root is black.
 - No red node has a red child.
 - Every root-null path in the tree passes through the same number of black nodes.



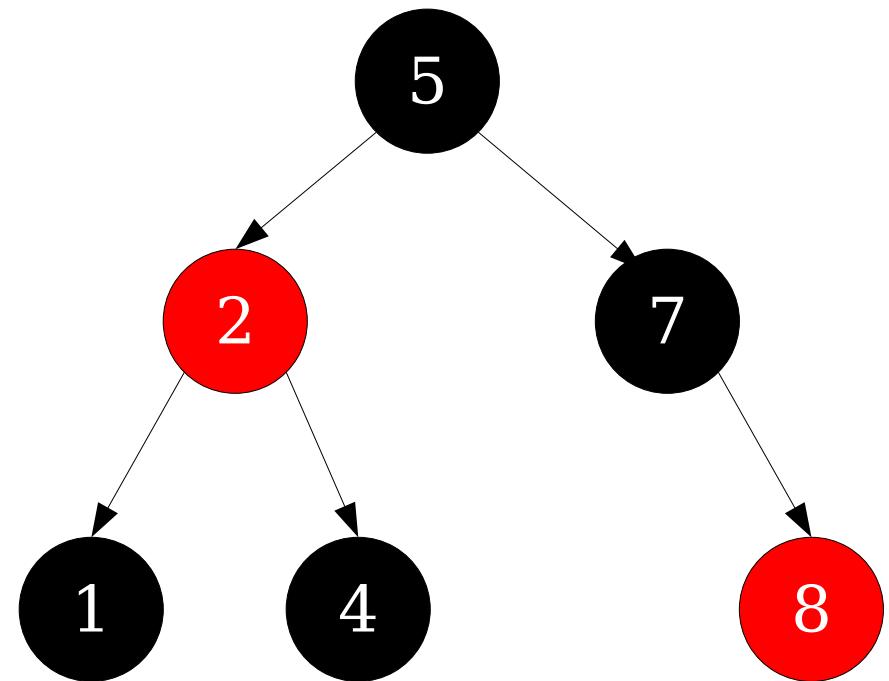
Red/Black Trees

- A ***red/black tree*** is a BST with the following properties:
 - Every node is either red or black.
 - The root is black.
 - No red node has a red child.
 - Every root-null path in the tree passes through the same number of black nodes.



Red/Black Trees

- A ***red/black tree*** is a BST with the following properties:
 - Every node is either red or black.
 - The root is black.
 - No red node has a red child.
 - Every root-null path in the tree passes through the same number of black nodes.



Red/Black Trees

- **Theorem:** Any red/black tree with n nodes has height $O(\log n)$.
 - We could prove this now, but there's a *much* simpler proof of this we'll see later on.
- Given a fixed red/black tree, lookups can be done in time $O(\log n)$.

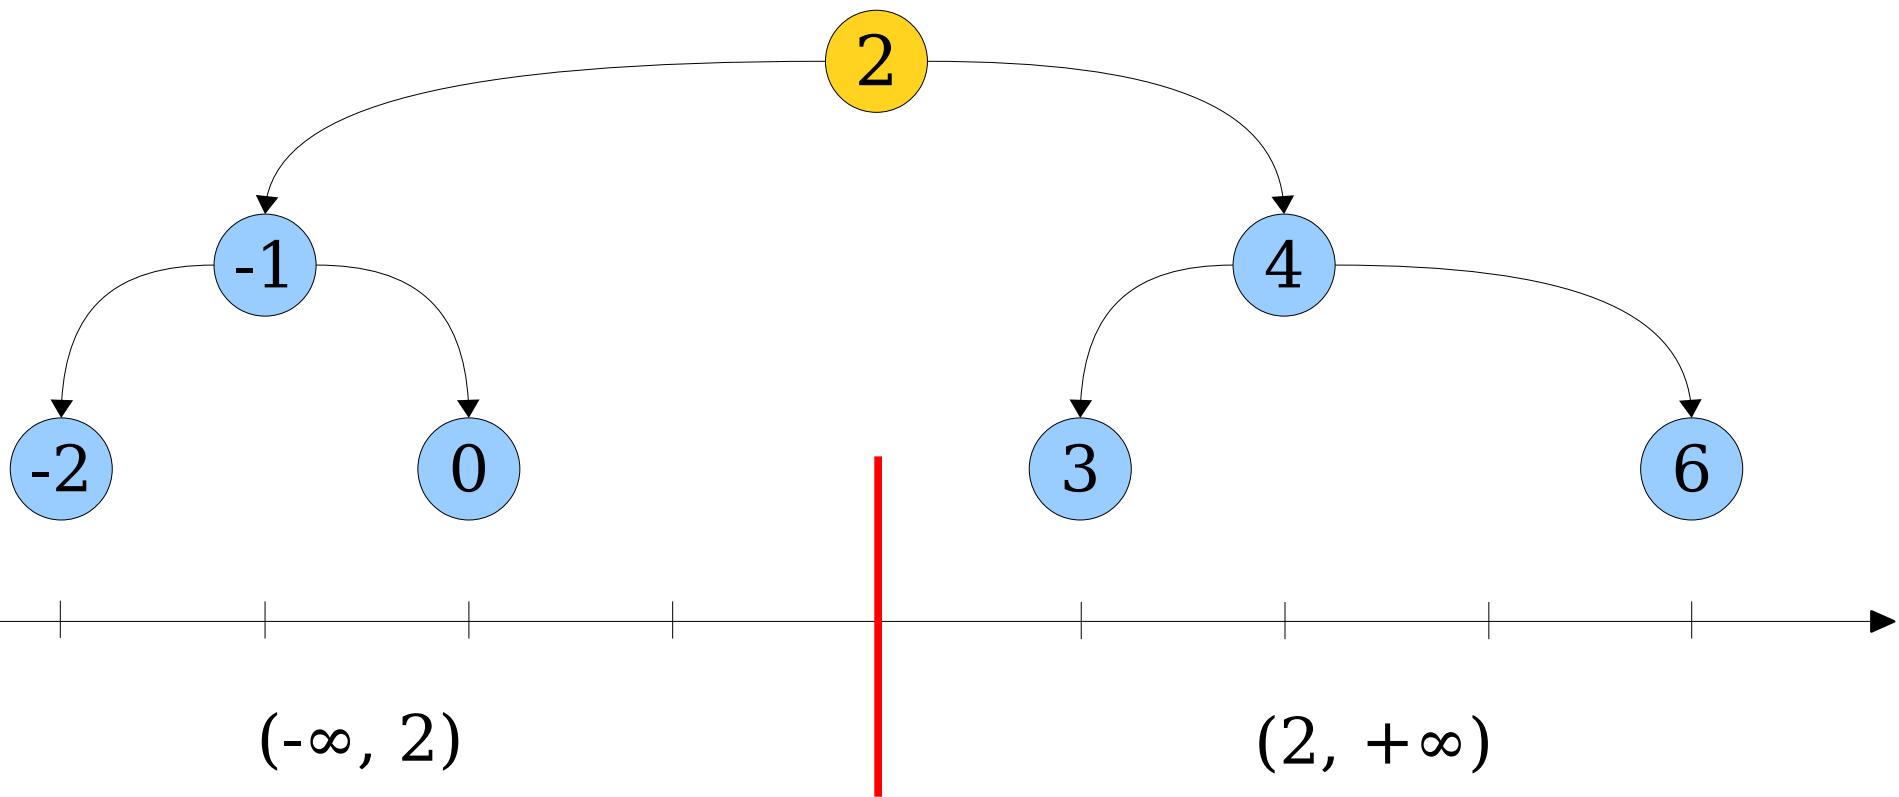
Fixing Up Red/Black Trees

- **The Good News:** After doing an insertion or deletion, we can locally modify a red/black tree in time $O(\log n)$ to fix up the red/black properties.
- **The Bad News:** There are a *lot* of cases to consider and they're not trivial.
- Some questions:
 - How do you memorize / remember all the rules for fixing up the tree?
 - How on earth did anyone come up with red/black trees in the first place?

B-Trees

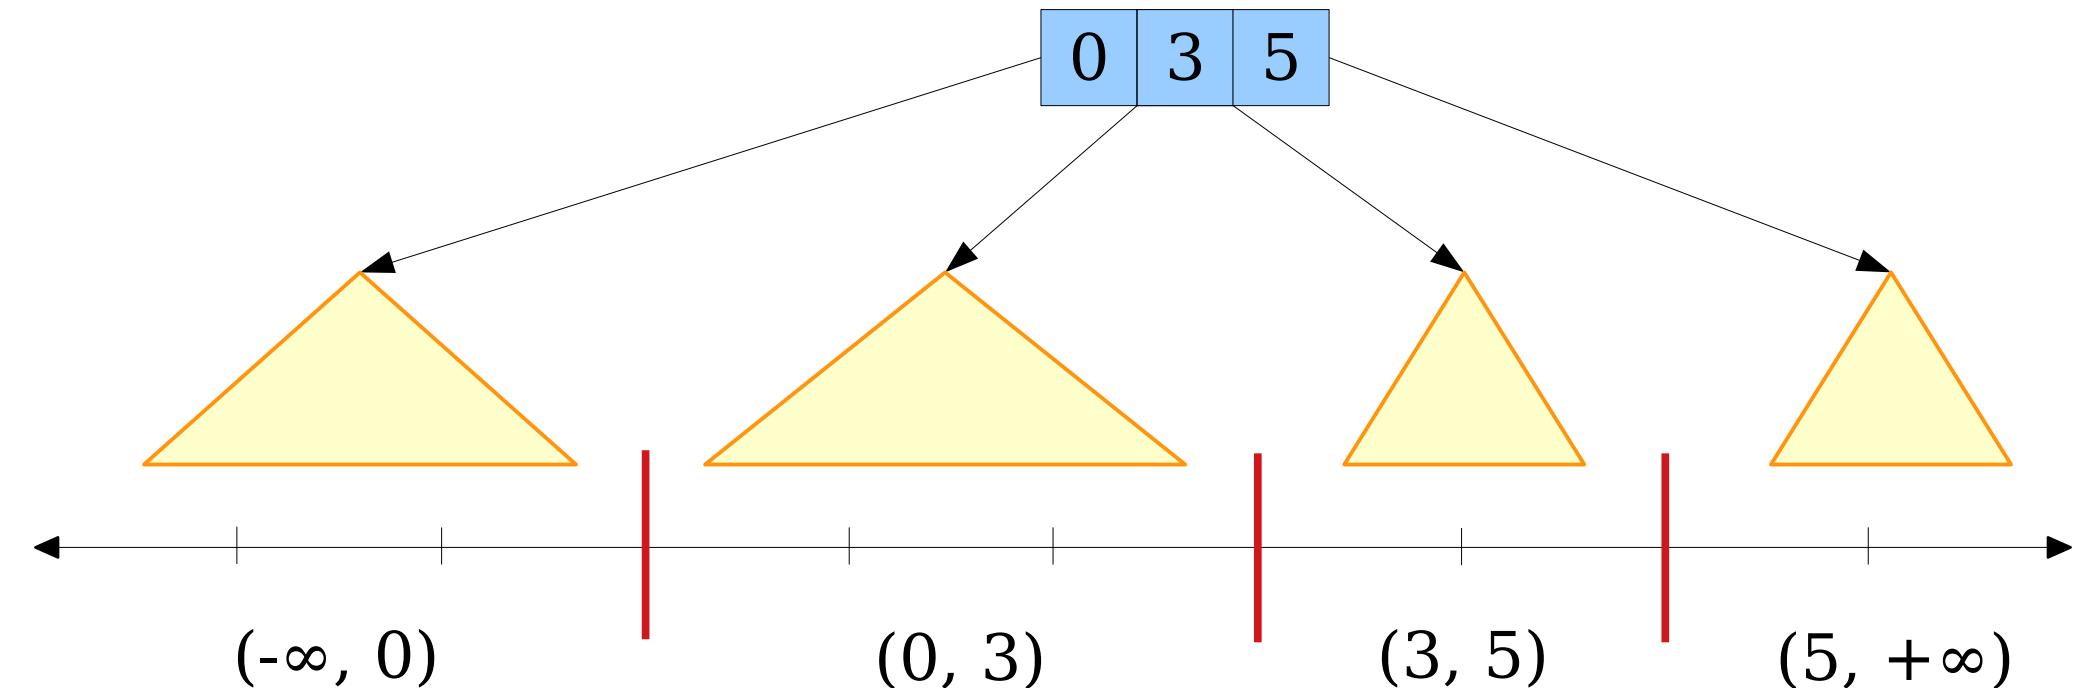
Generalizing BSTs

- In a binary search tree, each node stores a single key.
- That key splits the “key space” into two pieces, and each subtree stores the keys in those halves.



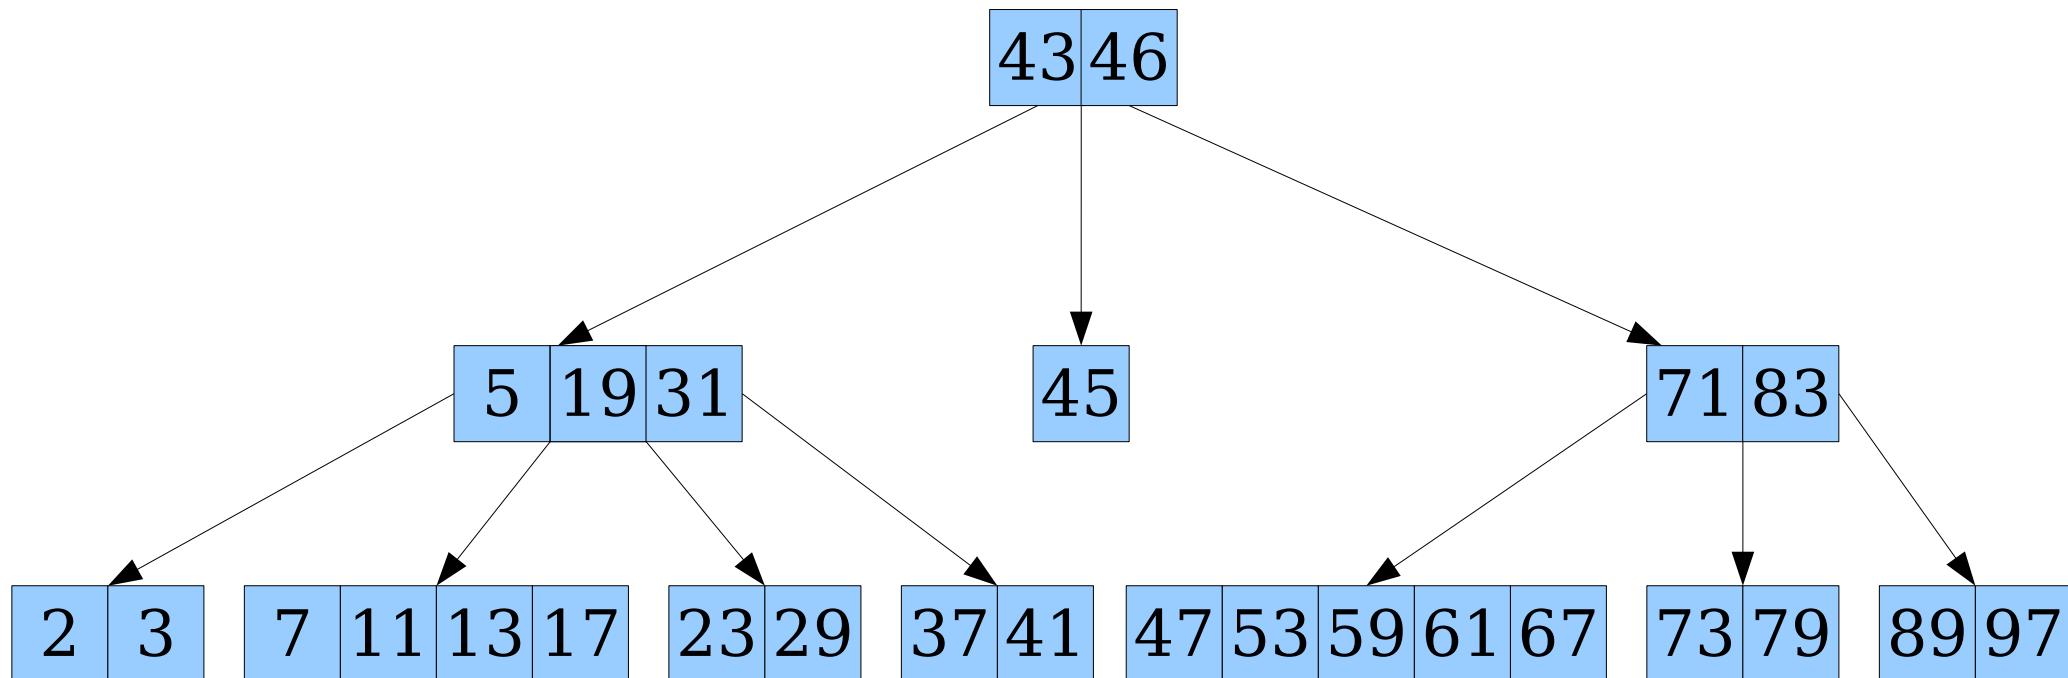
Generalizing BSTs

- In a ***multiway search tree***, each node stores an arbitrary number of keys in sorted order.
- A node with k keys splits the key space into $k+1$ regions, with subtrees for keys in each region.



Generalizing BSTs

- In a ***multiway search tree***, each node stores an arbitrary number of keys in sorted order.



- Surprisingly, it's a bit easier to build a balanced multiway tree than it is to build a balanced BST. Let's see how.

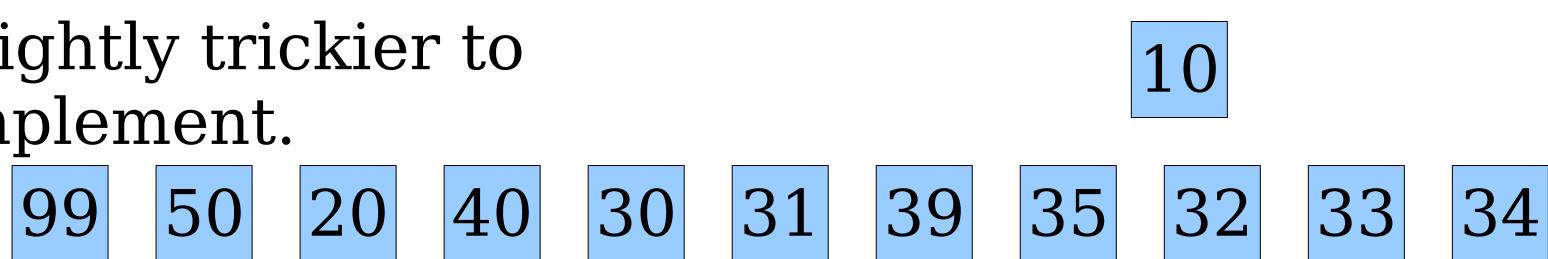
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.

10 99 50 20 40 30 31 39 35 32 33 34

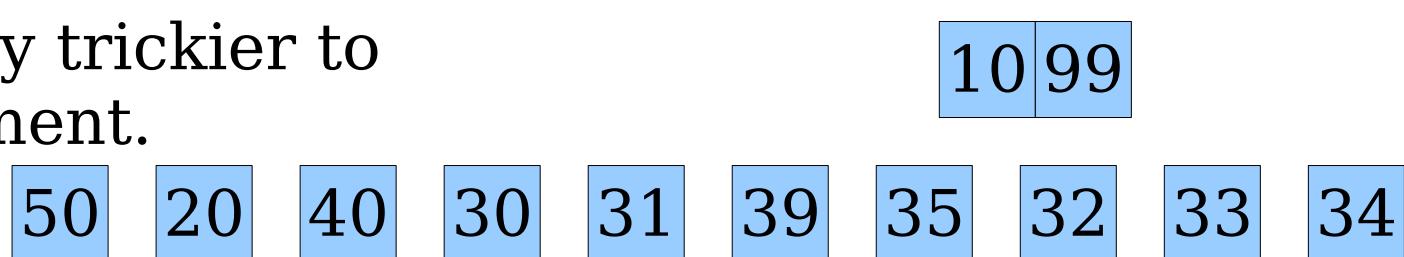
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



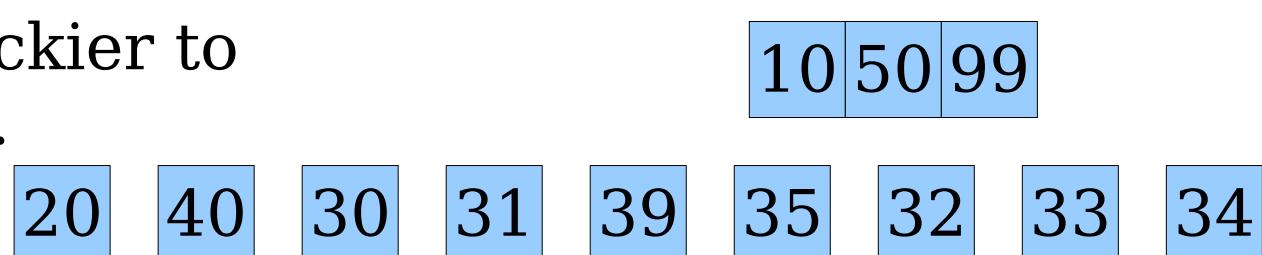
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



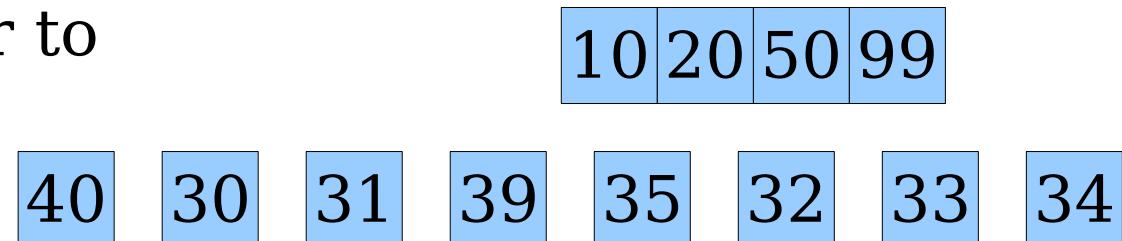
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



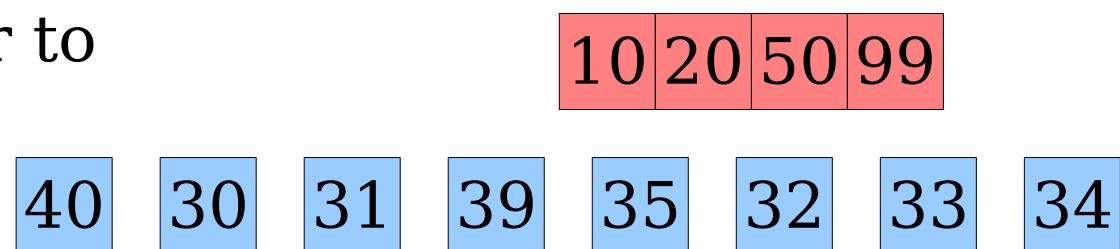
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



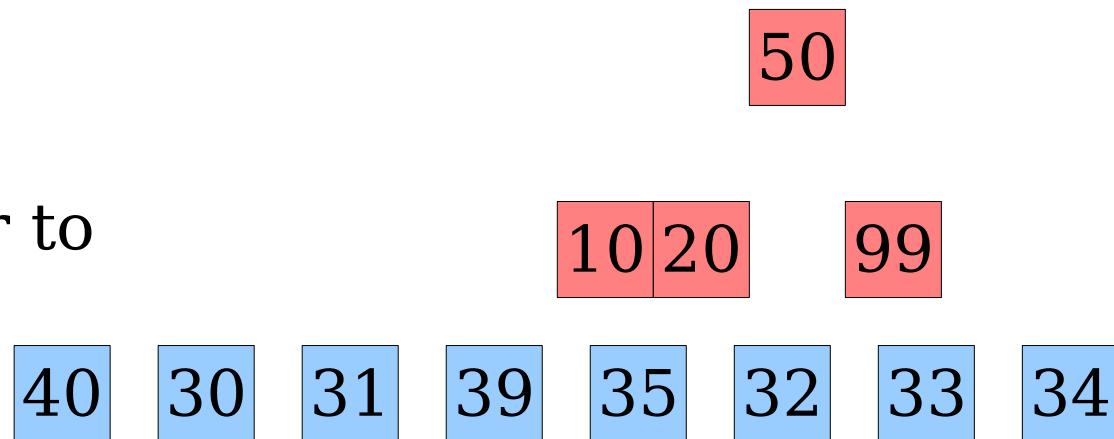
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



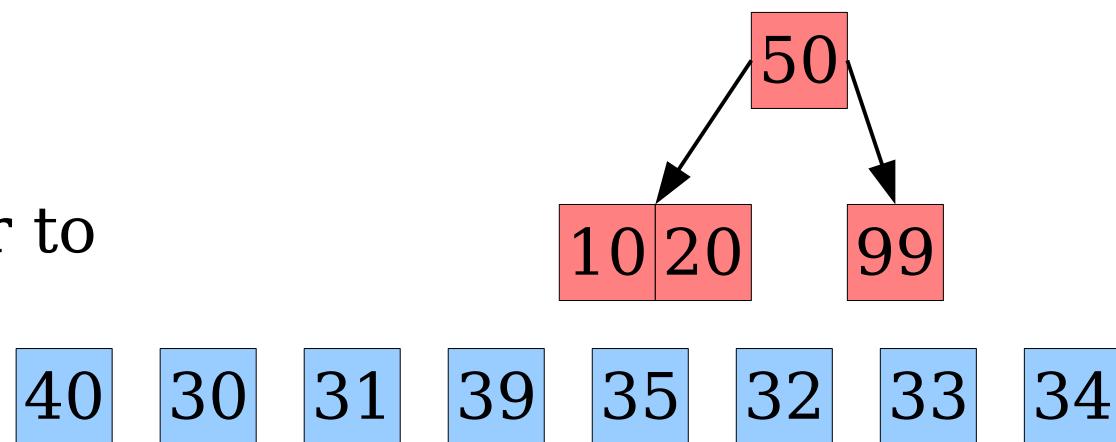
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



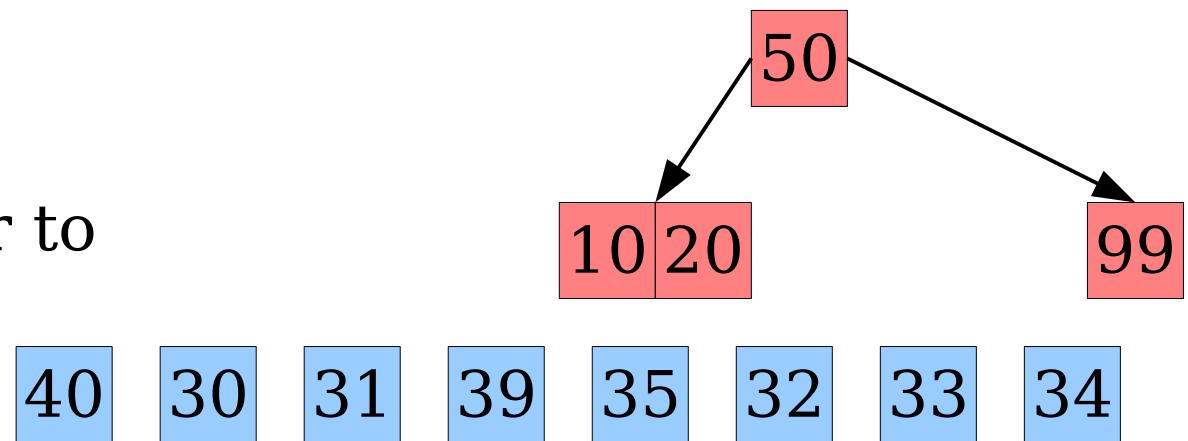
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



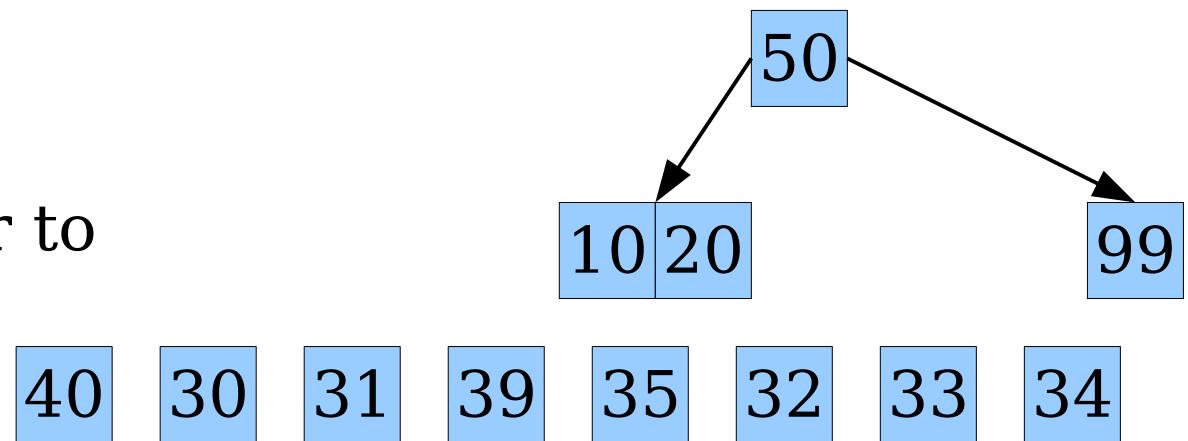
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



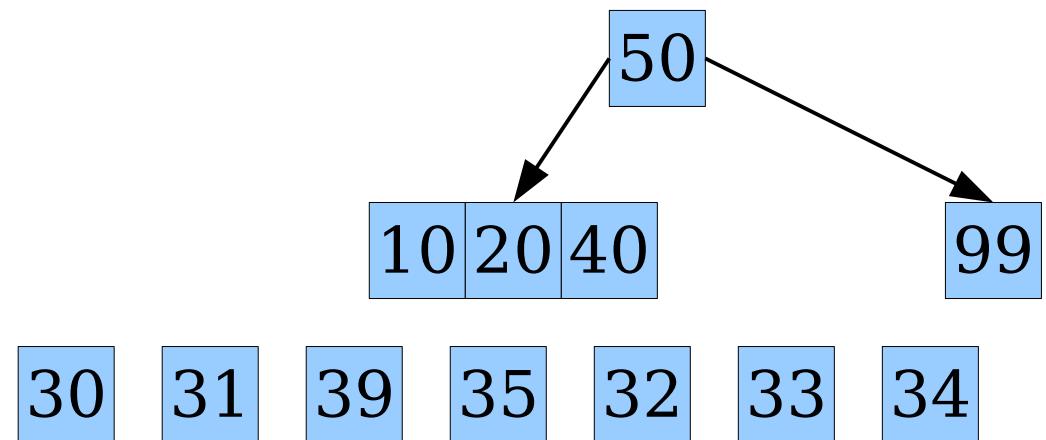
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



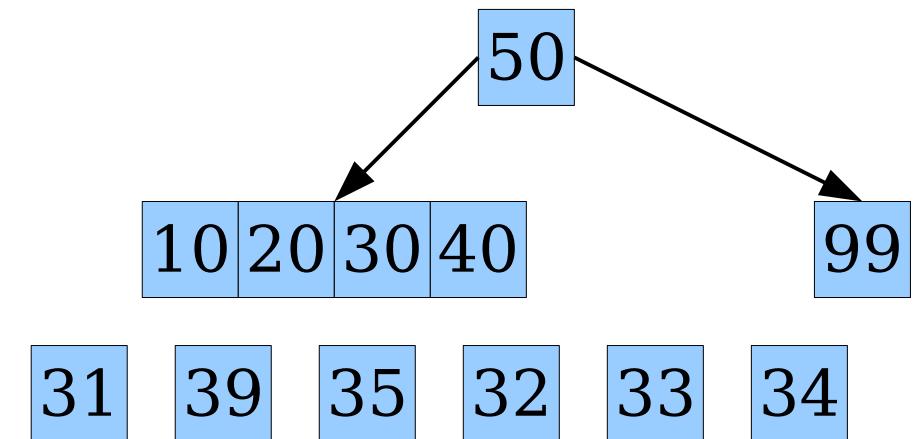
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



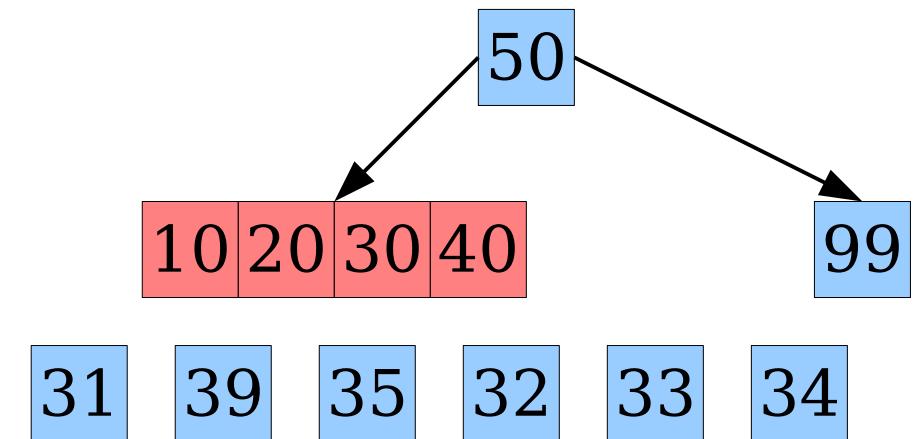
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



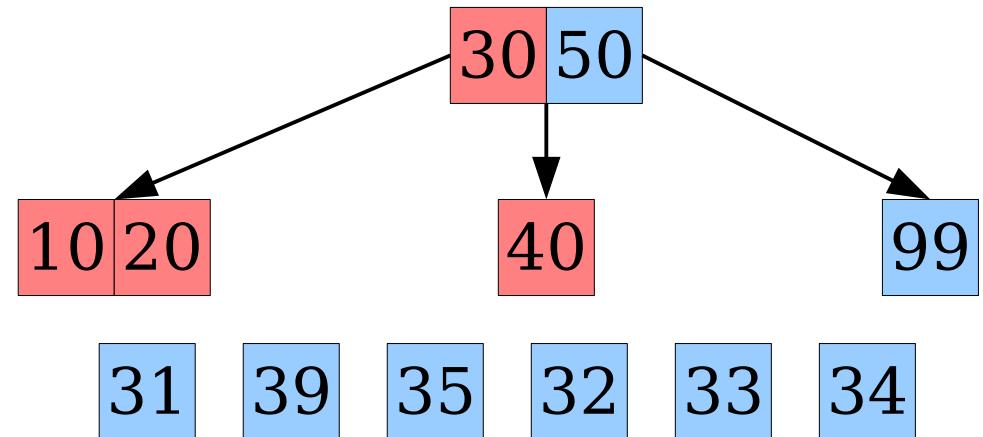
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



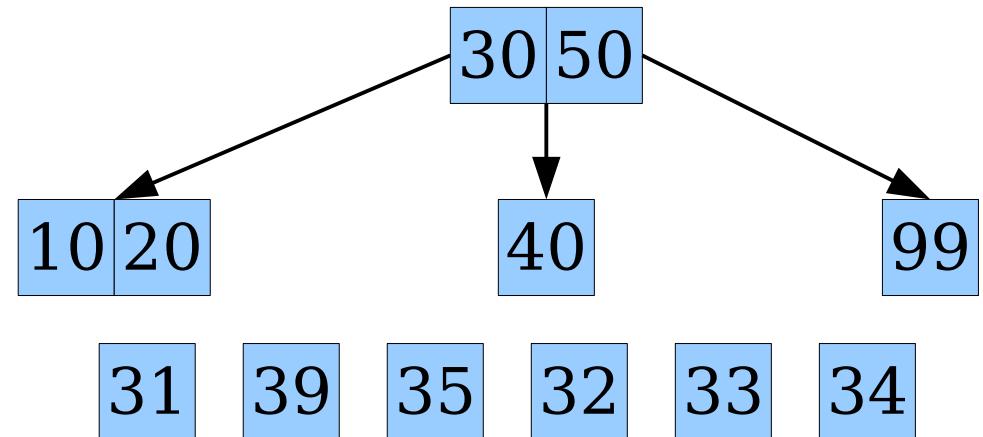
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



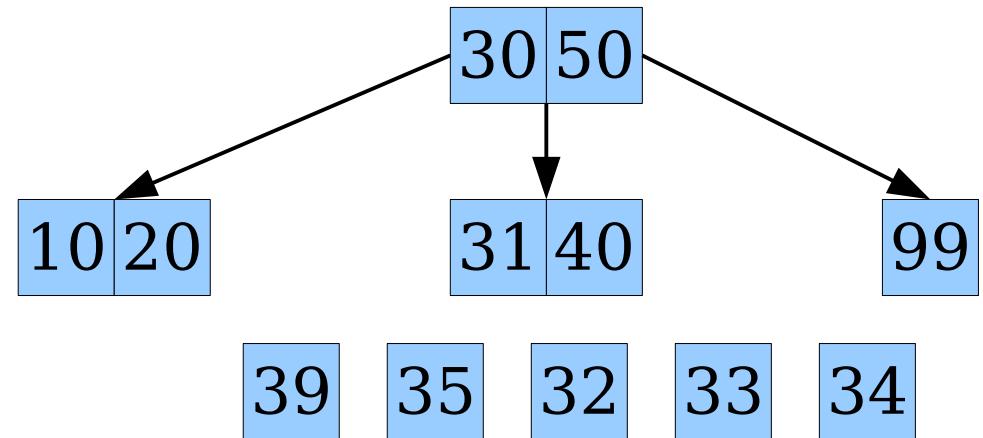
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



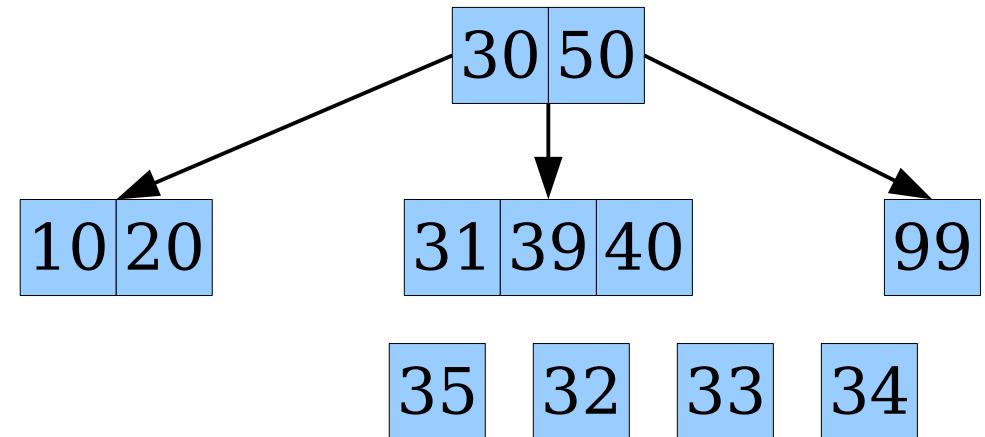
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



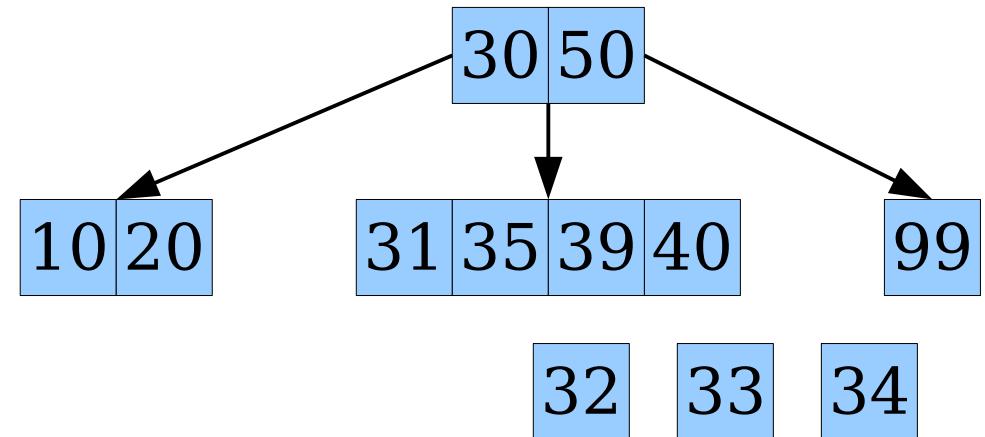
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



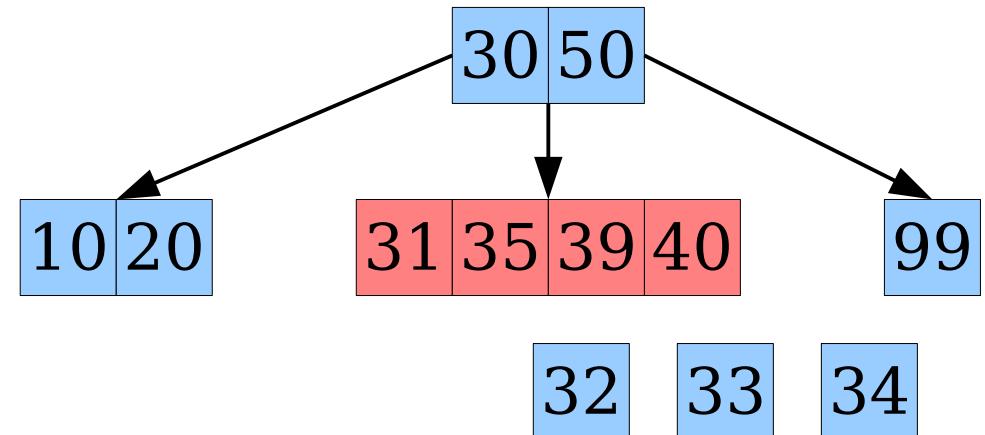
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



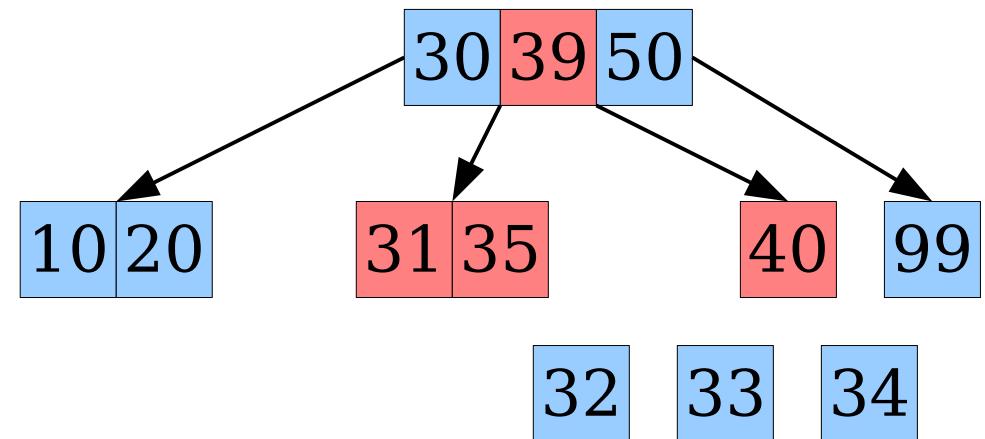
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



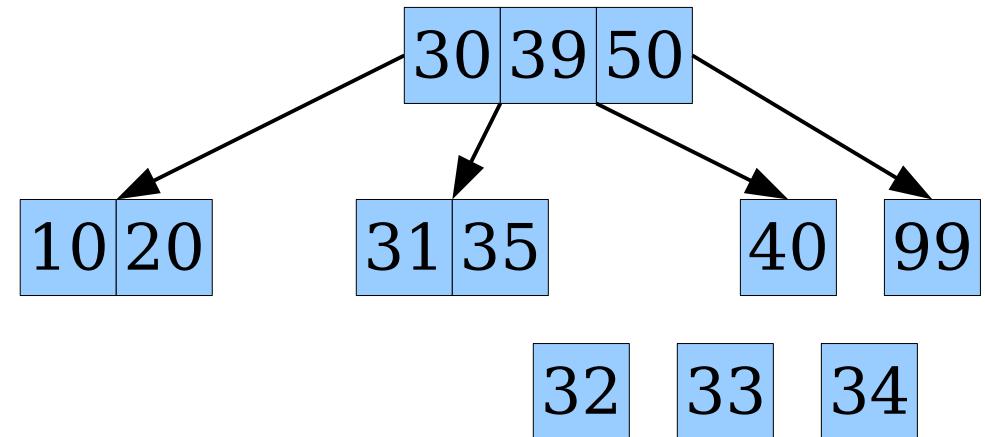
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



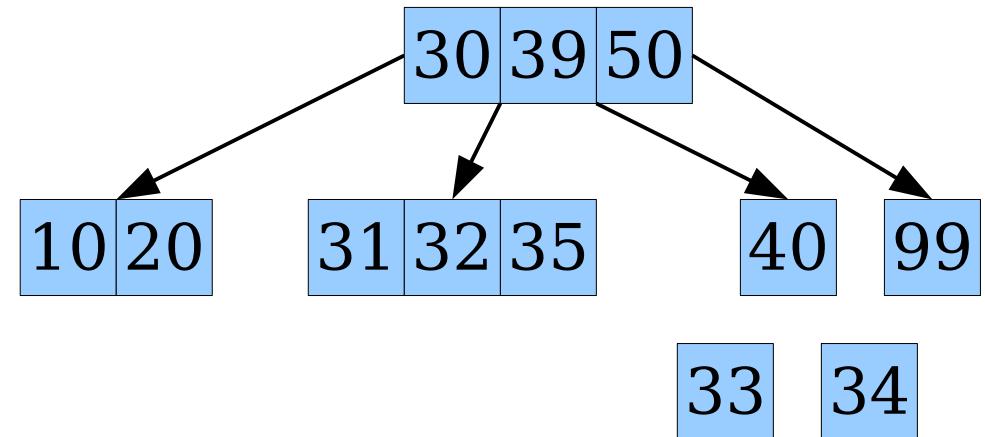
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



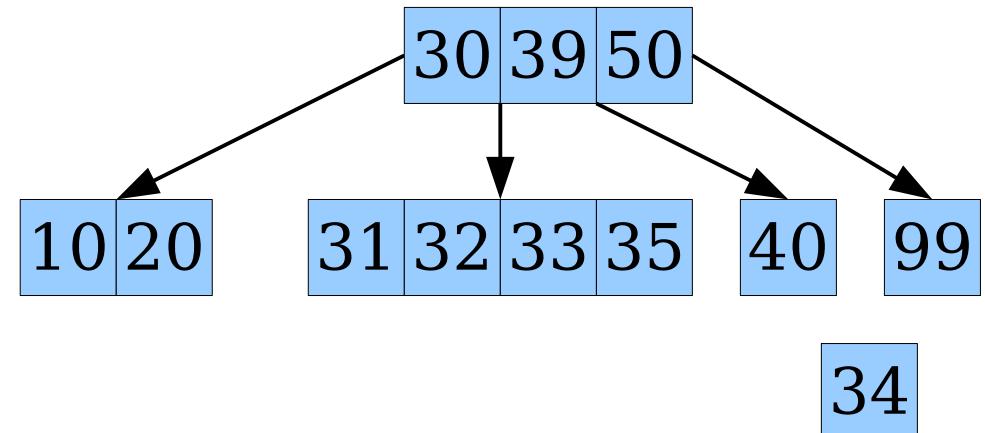
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



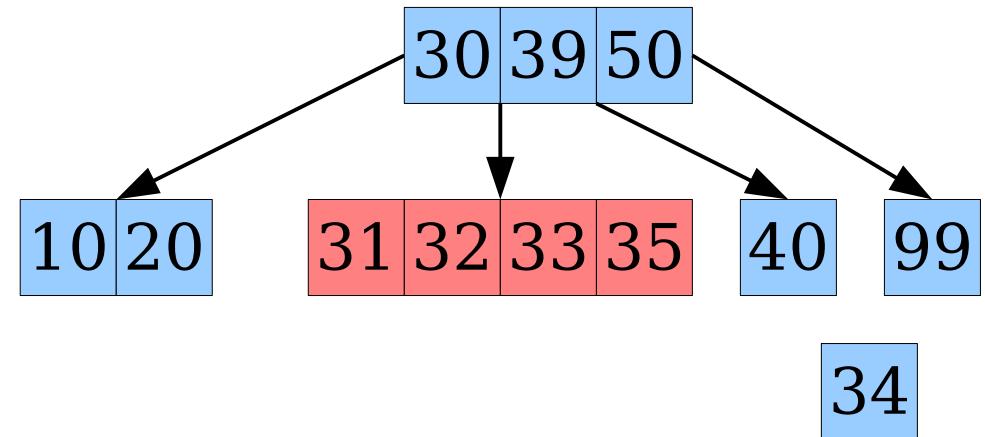
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



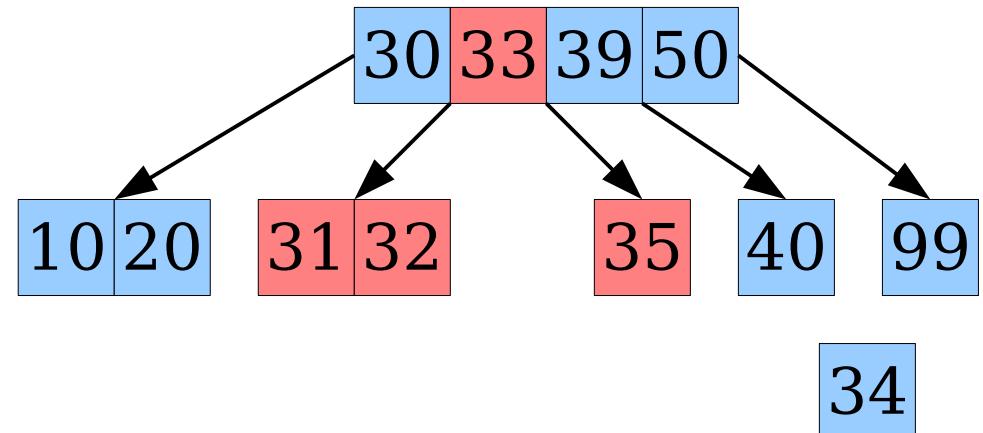
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



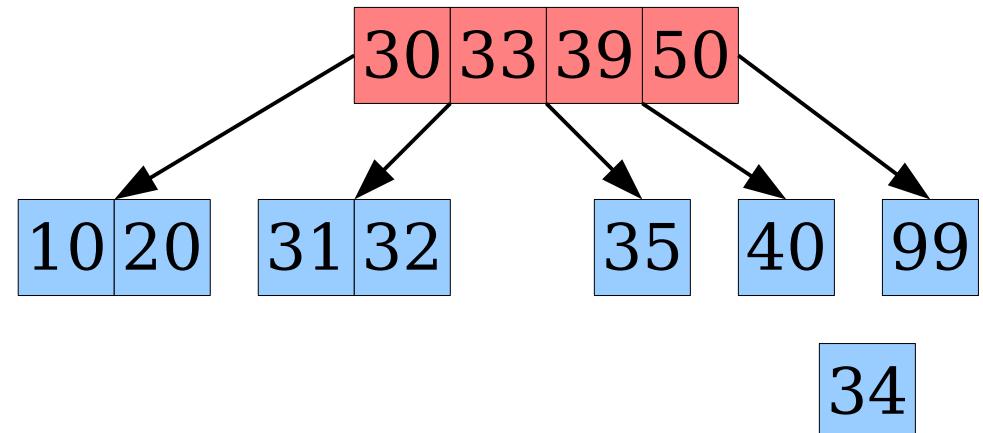
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



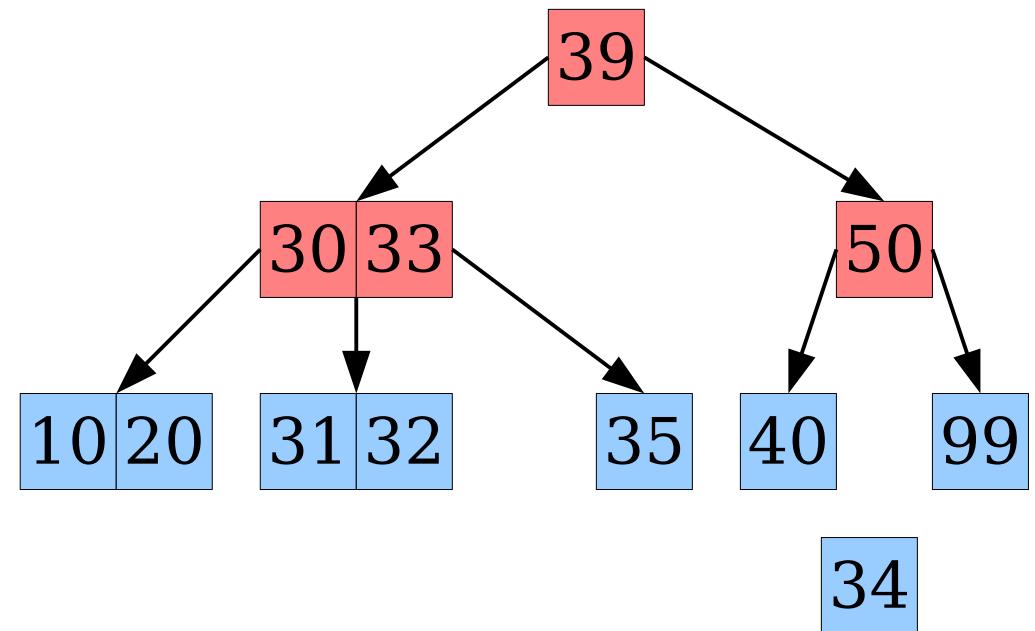
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



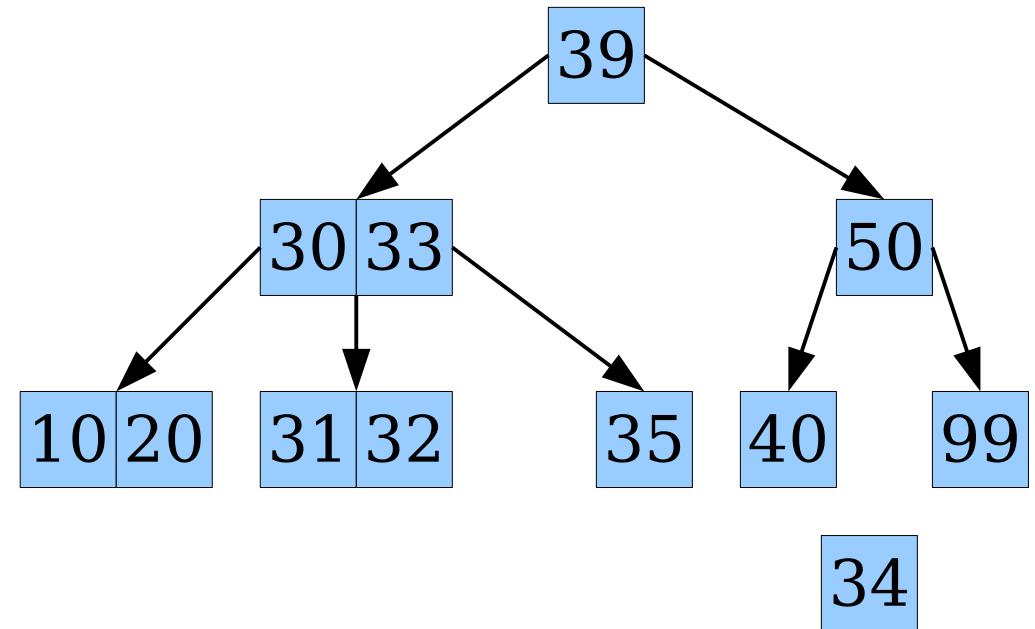
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



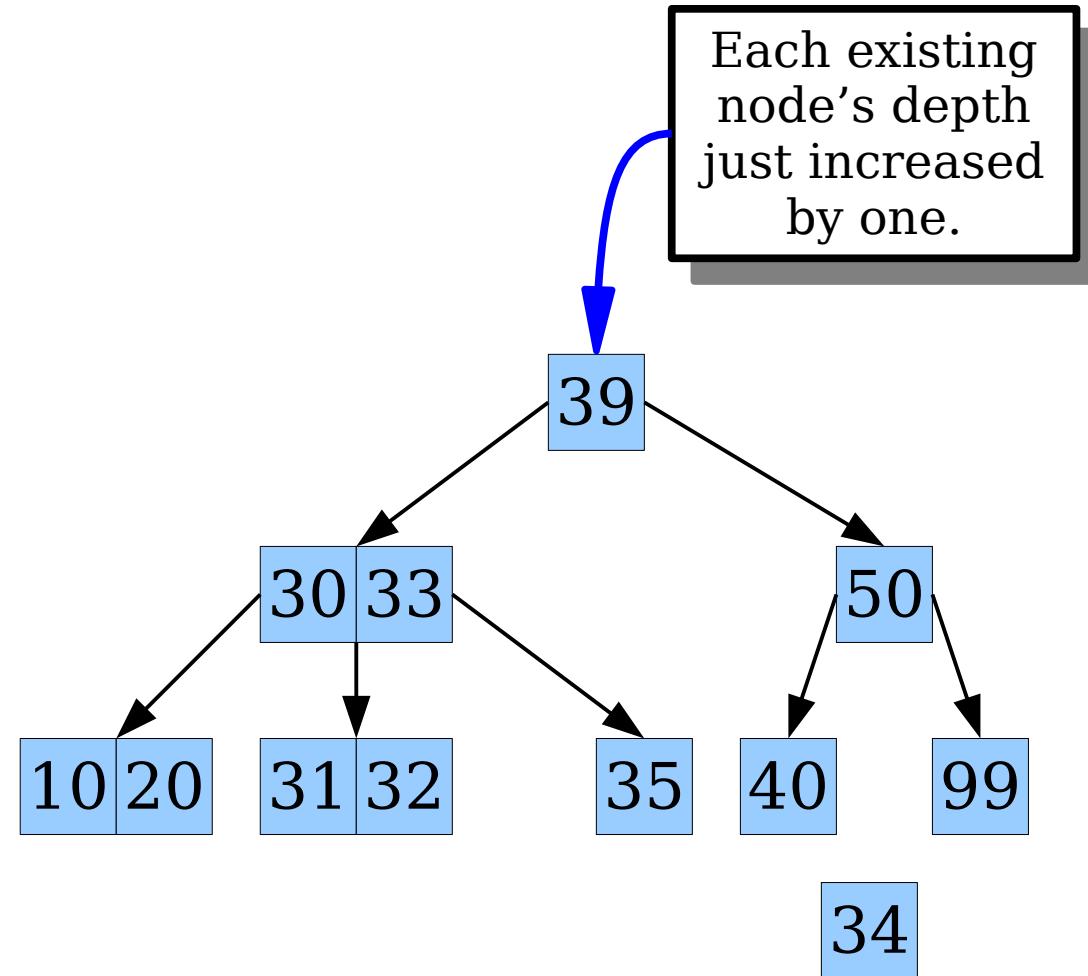
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



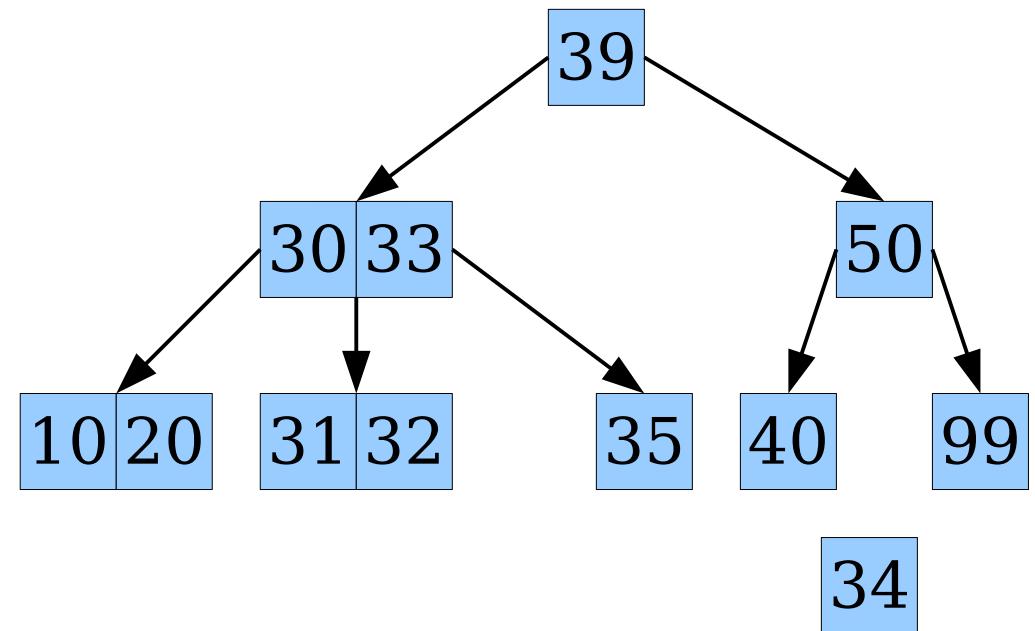
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



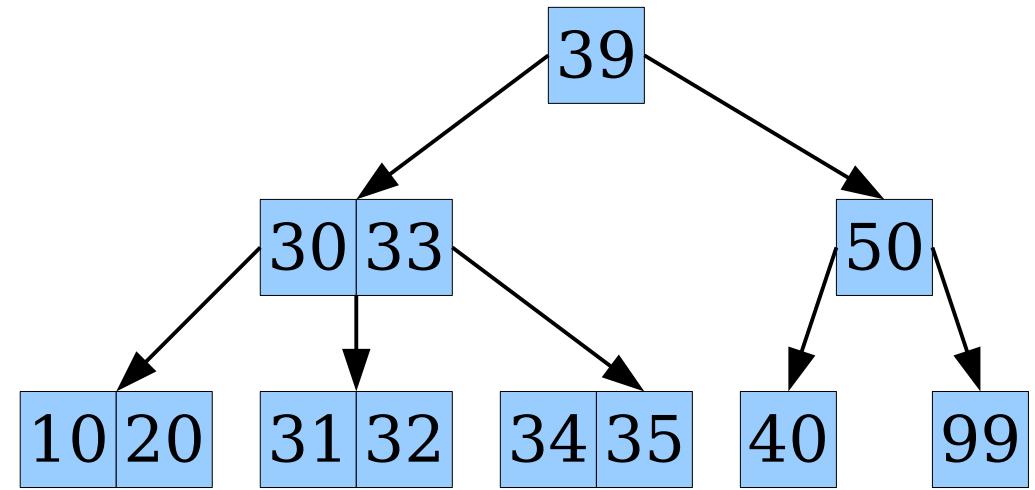
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



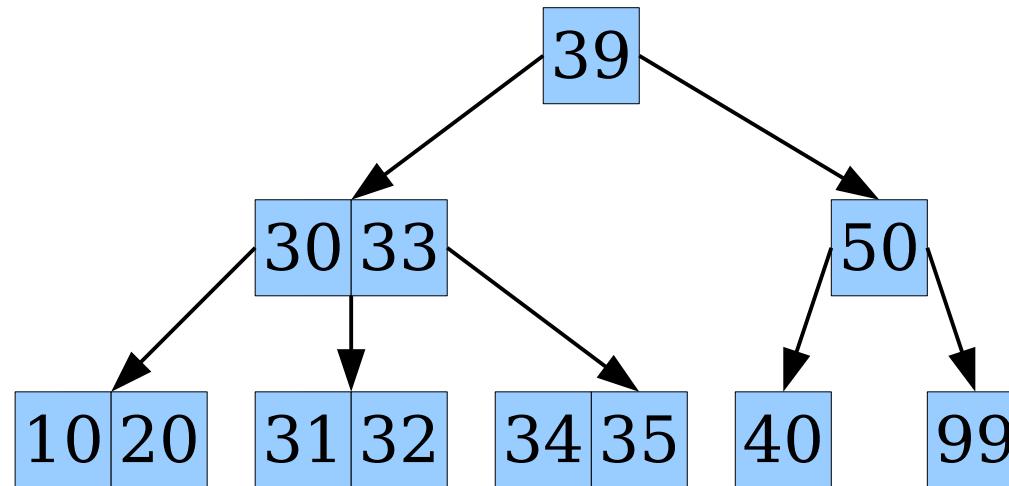
Balanced Multiway Trees

- **Option 1:** Push keys down into new nodes.
 - Simple to implement.
 - Can lead to tree imbalances.
- **Option 2:** Split big nodes, kicking keys higher up.
 - Keeps the tree balanced.
 - Slightly trickier to implement.



Balanced Multiway Trees

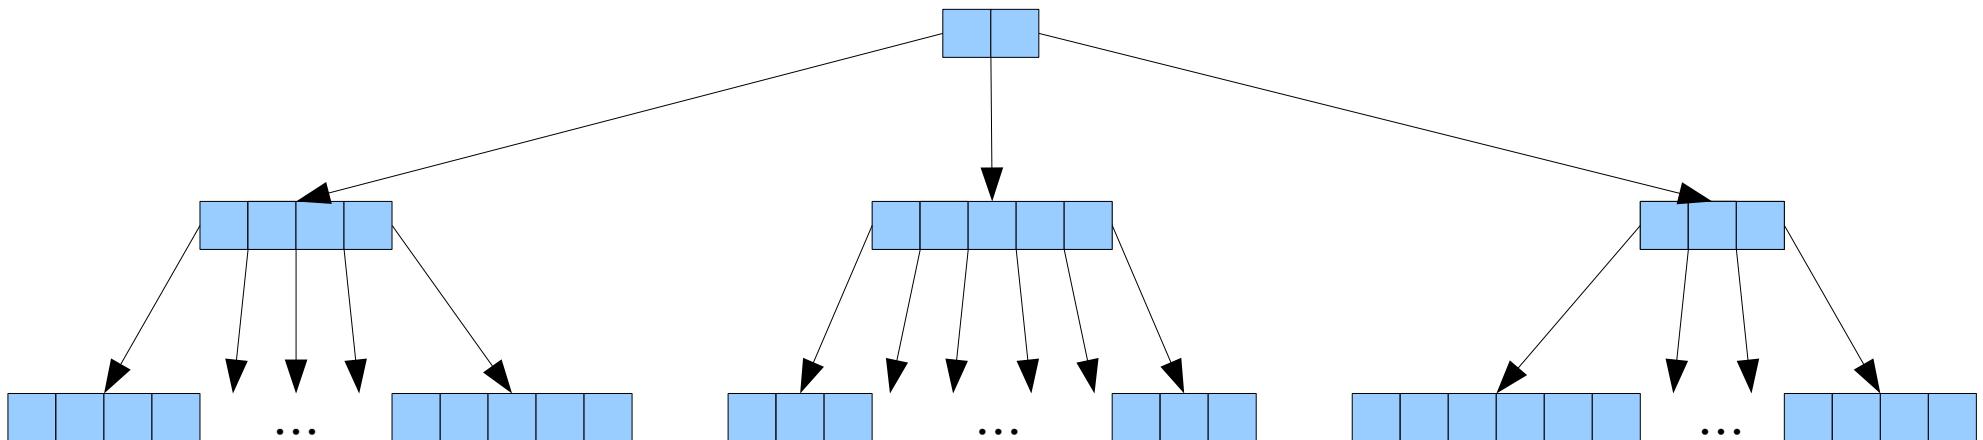
- **General idea:** Cap the maximum number of keys in a node. Add keys into leaves. Whenever a node gets too big, split it and kick one key higher up the tree.



- **Advantage 1:** The tree is always balanced.
- **Advantage 2:** Insertions and lookups are pretty fast.

B-Trees

- A ***B-tree of order b*** is a multiway search tree where
 - each node has between $b-1$ and $2b-1$ keys, except the root, which may only have between 1 and $2b-1$ keys;
 - each node is either a leaf or has one more child than key; and
 - all leaves are at the same depth.
- Different authors give different bounds on how many keys can be in each node. The ranges are often $[b-1, 2b-1]$ or $[b, 2b]$. For the purposes of today's lecture, we'll use the range $[b-1, 2b-1]$ for the key limits, just for simplicity.



Analyzing B-Trees

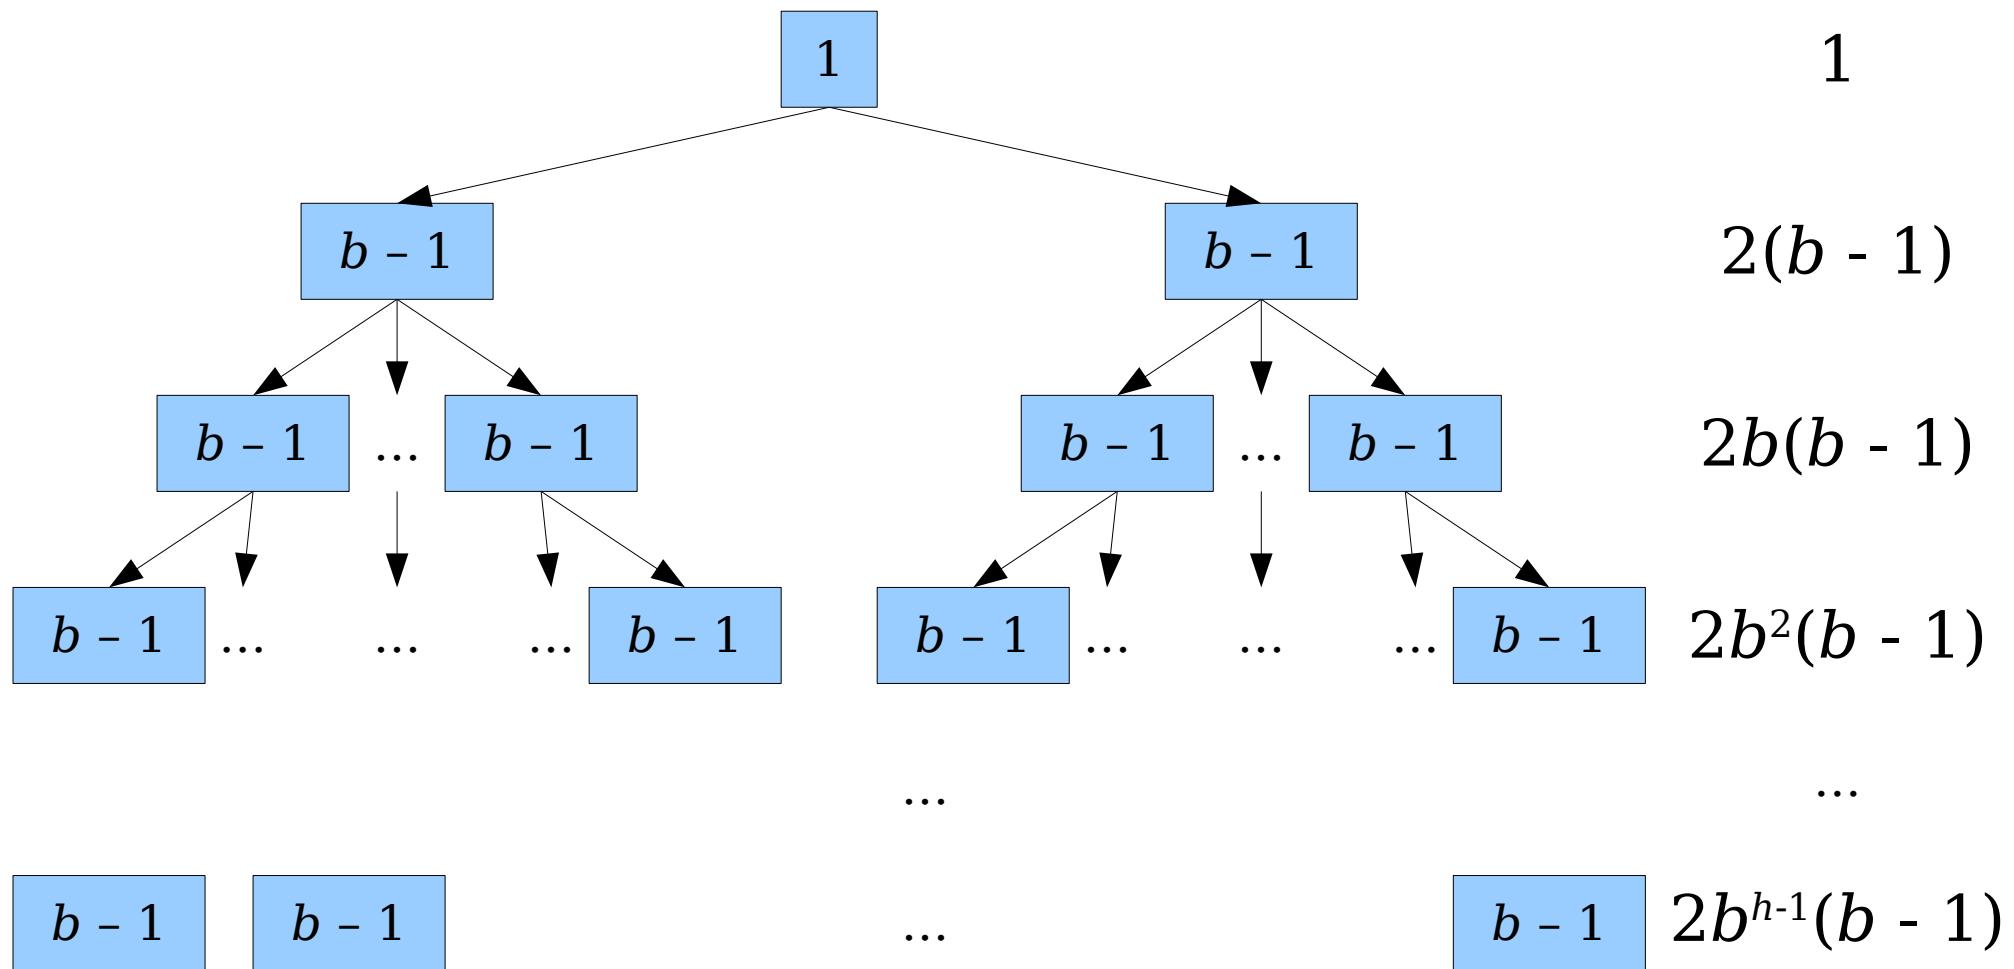
The Height of a B-Tree

- What is the maximum possible height of a B-tree of order b that holds n keys?

Intuition: The branching factor of the tree is at least b , so the number of keys per level grows exponentially in b . Therefore, we'd expect something along the lines of $O(\log_b n)$.

The Height of a B-Tree

- What is the maximum possible height of a B-tree of order b that holds n keys?



The Height of a B-Tree

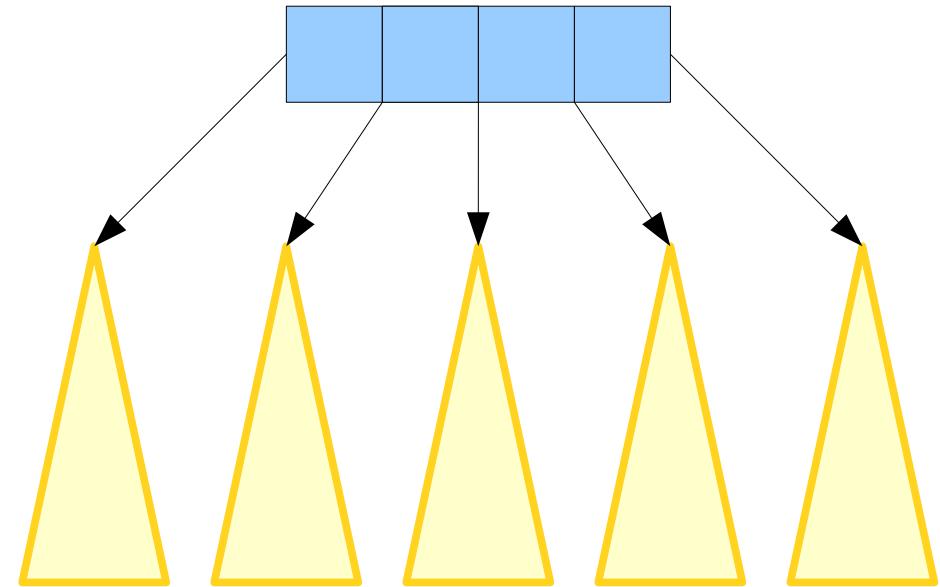
- **Theorem:** The maximum height of a B-tree of order b containing n keys is $O(\log_b n)$.
- **Proof:** Number of keys n in a B-tree of height h is guaranteed to be at least

$$\begin{aligned} & 1 + 2(\mathbf{b - 1}) + 2\mathbf{b}(\mathbf{b - 1}) + 2\mathbf{b}^2(\mathbf{b - 1}) + \dots + 2\mathbf{b}^{h-1}(\mathbf{b - 1}) \\ &= 1 + 2(\mathbf{b - 1})(1 + \mathbf{b} + \mathbf{b}^2 + \dots + \mathbf{b}^{h-1}) \\ &= 1 + 2(\mathbf{b - 1})((\mathbf{b}^h - 1) / (\mathbf{b - 1})) \\ &= 1 + 2(\mathbf{b}^h - 1) = 2\mathbf{b}^h - 1. \end{aligned}$$

Solving $n = 2b^h - 1$ yields $h = \log_b ((n + 1) / 2)$, so the height is $O(\log_b n)$. ■

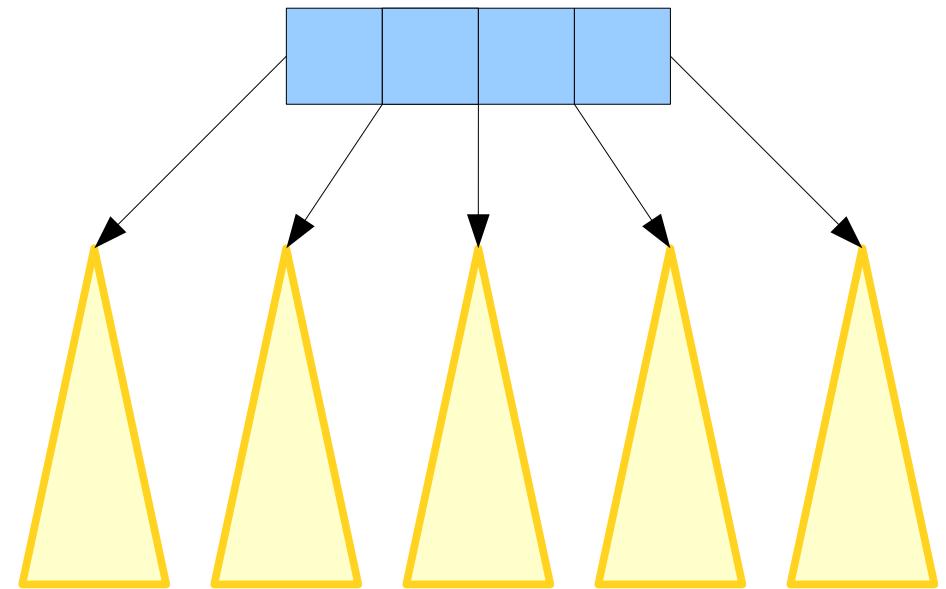
Analyzing Efficiency

- Suppose we have a B-tree of order b .
- What is the worst-case runtime of looking up a key in the B-tree?
- **Answer:** It depends on how we do the search!



Analyzing Efficiency

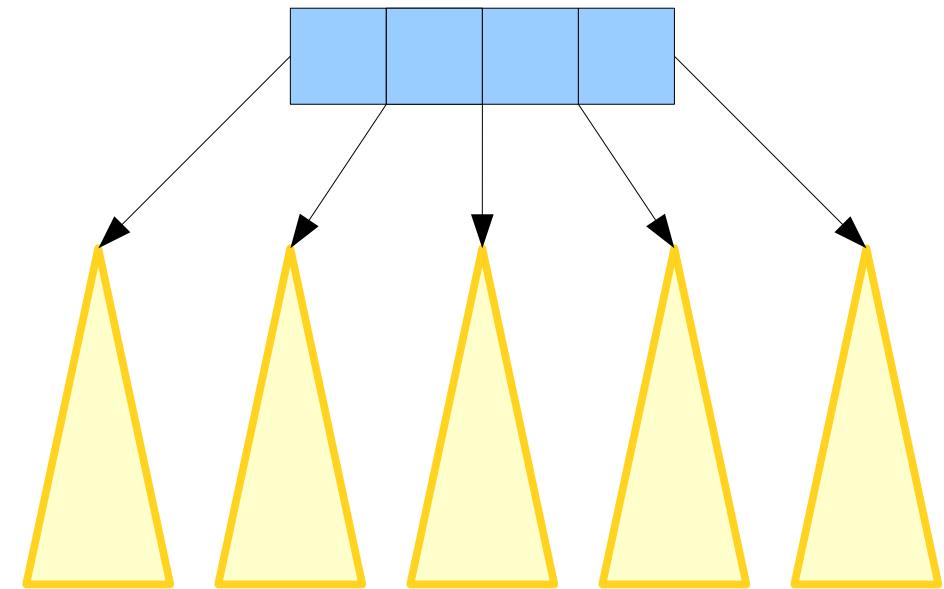
- To do a lookup in a B-tree, we need to determine which child tree to descend into.
- This means we need to compare our query key against the keys in the node.
- **Question:** How should we do this?



Analyzing Efficiency

- **Option 1:** Use a linear search!
- Cost per node: $O(b)$.
- Nodes visited: $O(\log_b n)$.
- Total cost:
$$O(b) \cdot O(\log_b n)$$

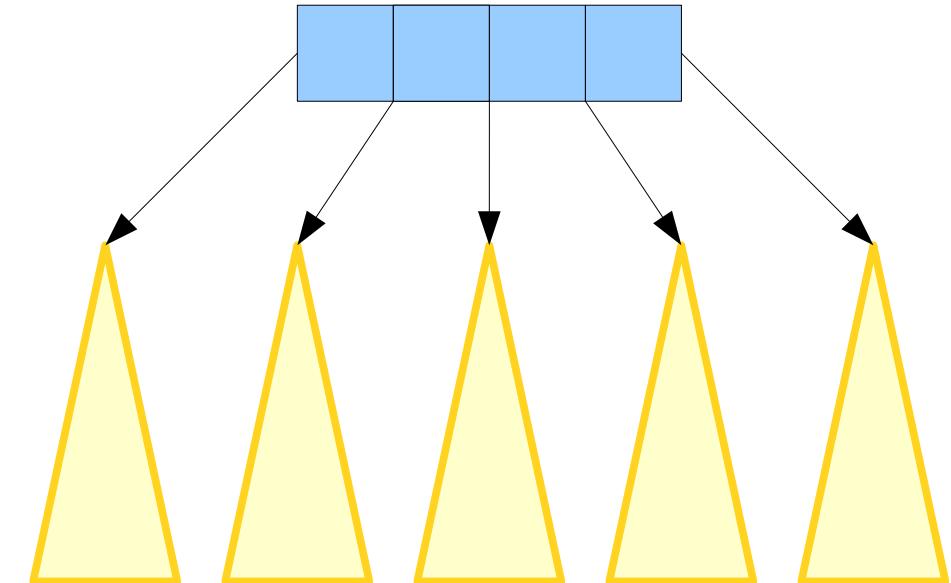
$$= \mathbf{O(b \log_b n)}$$



Analyzing Efficiency

- **Option 2:** Use a binary search!
- Cost per node: $O(\log b)$.
- Nodes visited: $O(\log_b n)$.
- Total cost:

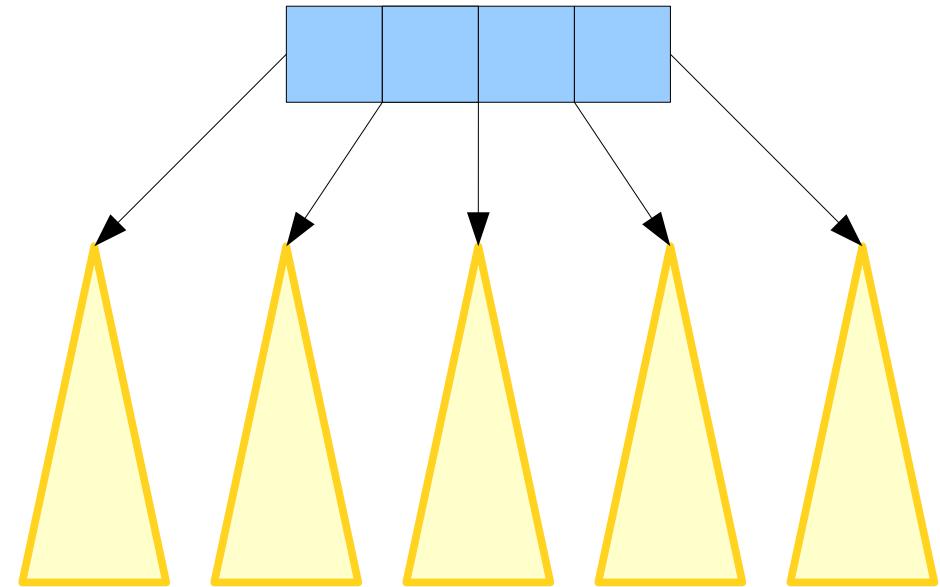
$$\begin{aligned} & O(\log b) \cdot O(\log_b n) \\ &= O(\log b \cdot \log_b n) \\ &= O(\log b \cdot (\log n) / (\log b)) \\ &= \mathbf{O(\log n)}. \end{aligned}$$



Intuition: We can't do better than $O(\log n)$ for arbitrary data, because it's the information-theoretic minimum number of comparisons needed to find something in a sorted collection!

Analyzing Efficiency

- Suppose we have a B-tree of order b .
- What is the worst-case runtime of inserting a key into the B-tree?
- Each insertion visits $O(\log_b n)$ nodes, and in the worst case we have to split every node we see.
- **Answer:** $O(b \log_b n)$.



Analyzing Efficiency

- The cost of an insertion in a B-tree of order b is $O(b \log_b n)$.
- What's the best choice of b to use here?
- Note that

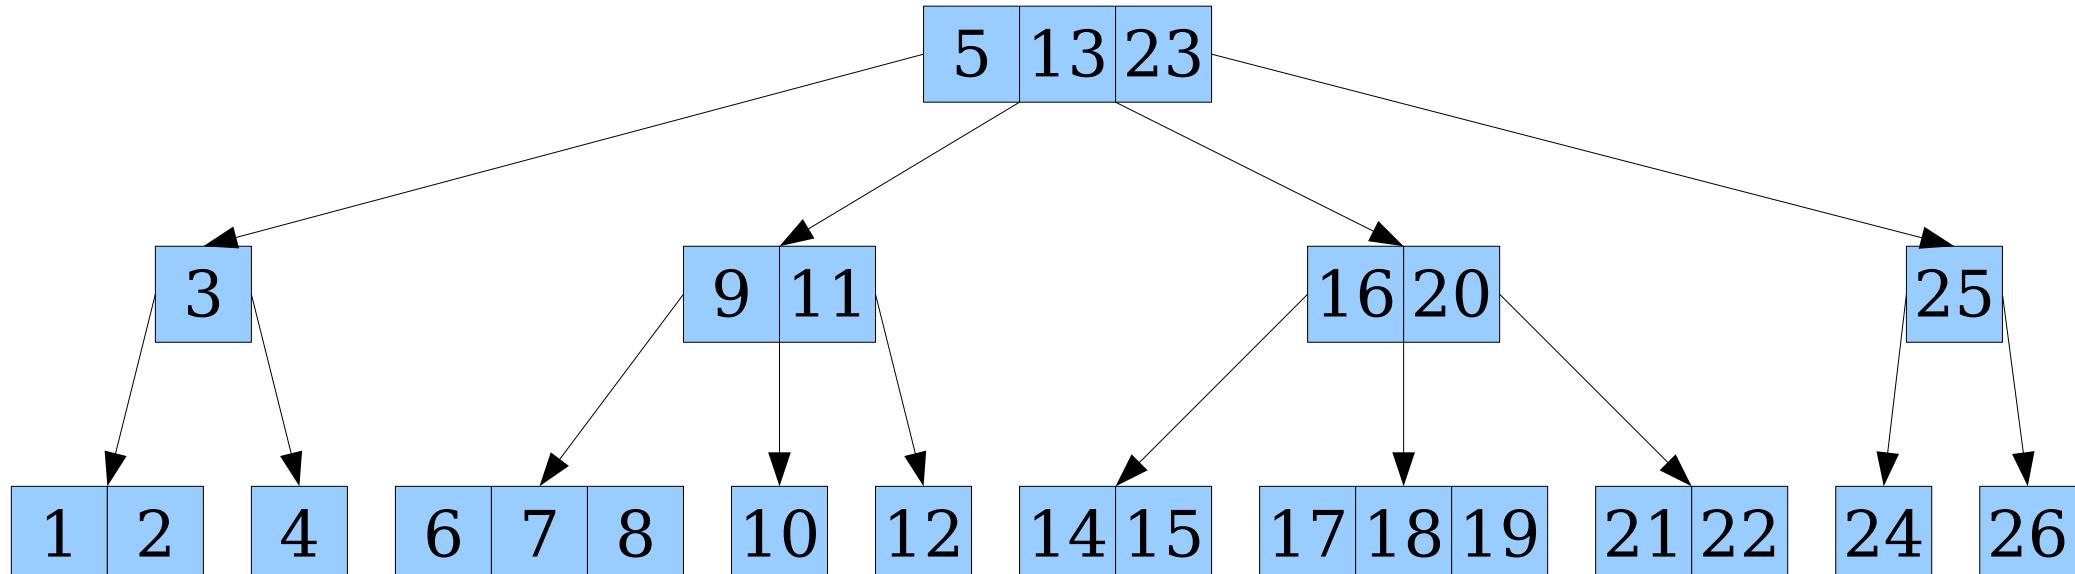
$$\begin{aligned} & b \log_b n \\ &= b (\log n / \log b) \\ &= (b / \log b) \log n. \end{aligned}$$

Fun fact: This is the same time bound you'd get if you used a b -ary heap instead of a binary heap for a priority queue.

- What choice of b minimizes $b / \log b$?
- **Answer:** Pick $b = e$.

2-3-4 Trees

- A **2-3-4 tree** is a B-tree of order 2. Specifically:
 - each node has between 1 and 3 keys;
 - each node is either a leaf or has one more child than key; and
 - all leaves are at the same depth.
- You actually saw this B-tree earlier! It's the type of tree from our insertion example.



The Story So Far

- A B-tree supports
 - lookups in time $O(\log n)$, and
 - insertions in time $O(b \log_b n)$.
- Picking b to be around 2 or 3 makes this optimal in Theoryland.
 - The 2-3-4 tree is great for that reason.
- ***Plot Twist:*** In practice, you most often see choices of b like 1,024 or 4,096.
- ***Question:*** Why would anyone do that?



Theoryland

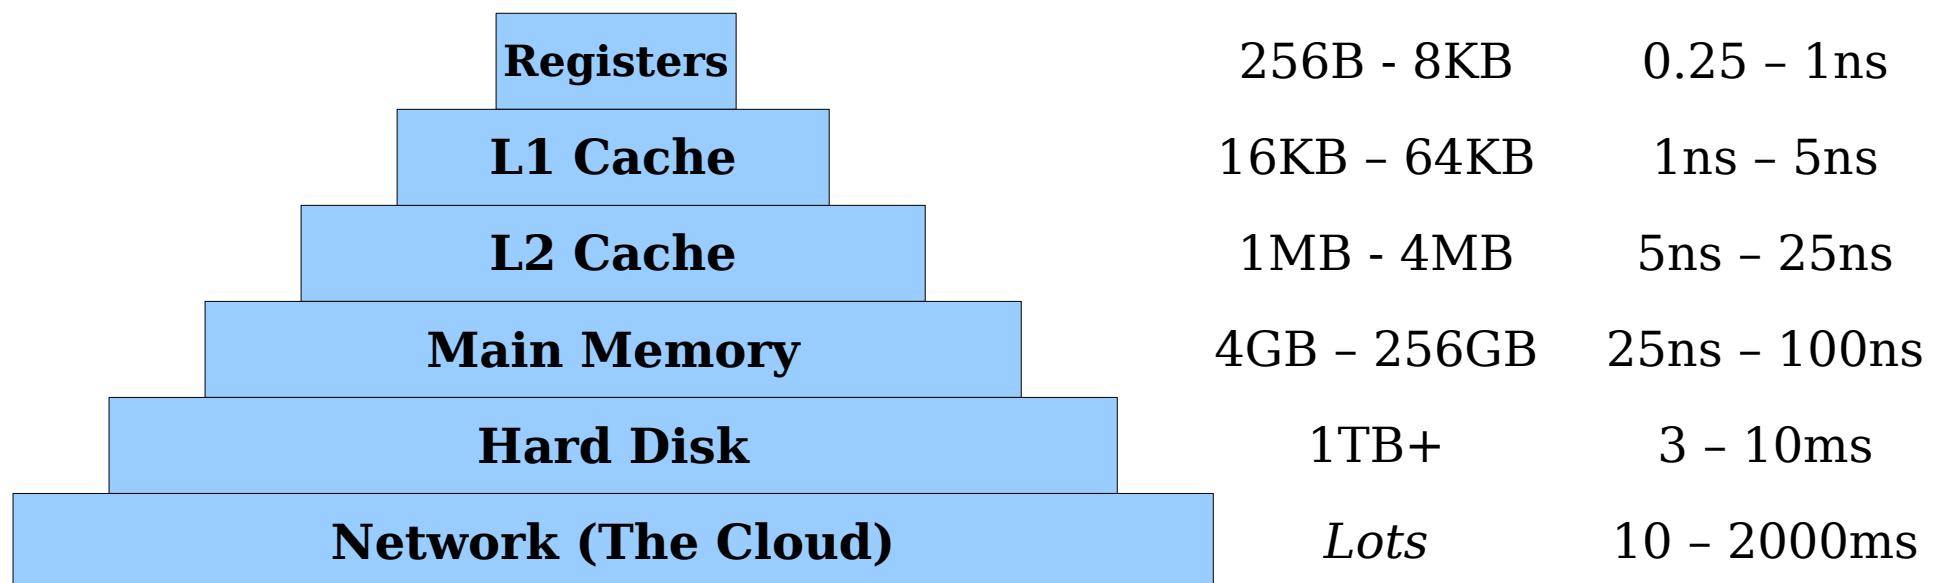


IRL

The Memory Hierarchy

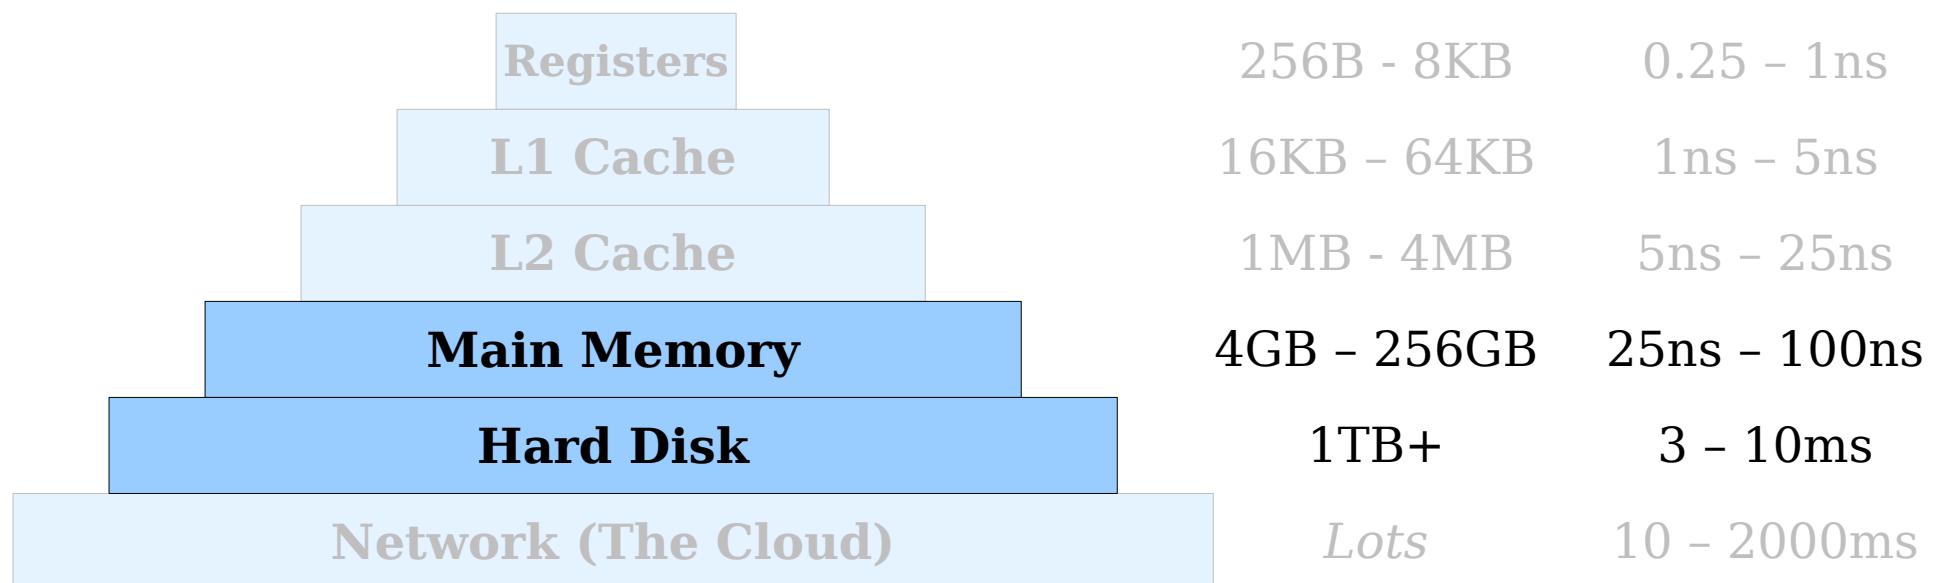
The Memory Hierarchy

- **Idea:** Try to get the best of all worlds by using multiple types of memory.



The Memory Hierarchy

- **Idea:** Try to get the best of all worlds by using multiple types of memory.



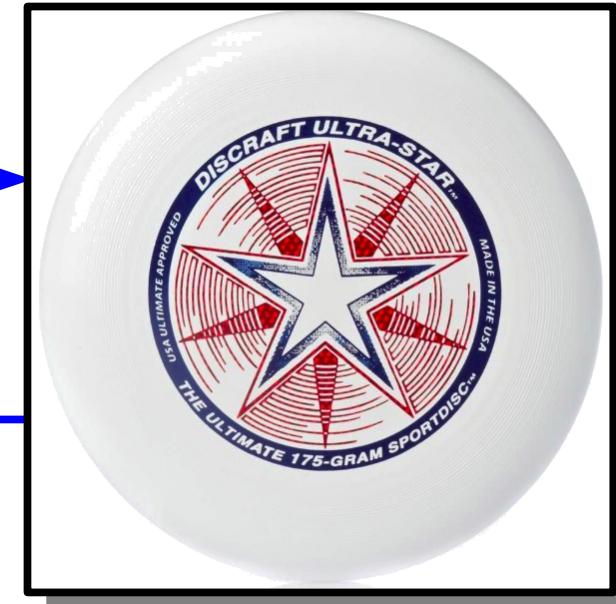
External Data Structures

- Suppose you have a data set that's *way* too big to fit in RAM.
- The data structure is on disk and read into RAM as needed.
- Data from disk doesn't come back one *byte* at a time, but rather one *page* at a time.
- **Goal:** Minimize the number of disk reads and writes, not the number of instructions executed.



"Please give me 4KB
starting at location *addr2*"

001101010001010001010001...



Analyzing B-Trees

- Suppose we tune b so that each node in the B-tree fits inside a single disk page.
- We *only* care about the number of disk pages read or written.
 - It's so much slower than RAM that it'll dominate the runtime.
- **Question:** What is the cost of a lookup in a B-tree in this model?
- **Question:** What is the cost of inserting into a B-tree in this model?

Analyzing B-Trees

- Suppose we tune b so that each node in the B-tree fits inside a single disk page.
- We *only* care about the number of disk pages read or written.
 - It's so much slower than RAM that it'll dominate the runtime.
- **Question:** What is the cost of a lookup in a B-tree in this model?
 - Answer: The height of the tree, $O(\log_b n)$.
- **Question:** What is the cost of inserting into a B-tree in this model?
 - Answer: The height of the tree, $O(\log_b n)$.

External Data Structures

- Because B-trees have a huge branching factor, they're great for on-disk storage.
 - Disk block reads/writes are slow compared to CPU operations.
 - The high branching factor minimizes the number of blocks to read during a lookup.
 - Extra work scanning inside a block offset by these savings.
- Major use cases for B-trees and their variants (B⁺-trees, H-trees, etc.) include
 - databases (huge amount of data stored on disk);
 - file systems (ext4, NTFS, ReFS); and, recently,
 - in-memory data structures (due to cache effects).

Analyzing B-Trees

- The cost model we use will change our overall analysis.
- Cost is number of operations:
O(log n) per lookup, **O(b log_b n)** per insertion.
- Cost is number of blocks accessed:
O(log_b n) per lookup, **O(log_b n)** per insertion.
- Going forward, we'll use operation counts as our cost model, though looking at caching effects of data structures would make for an awesome final project!

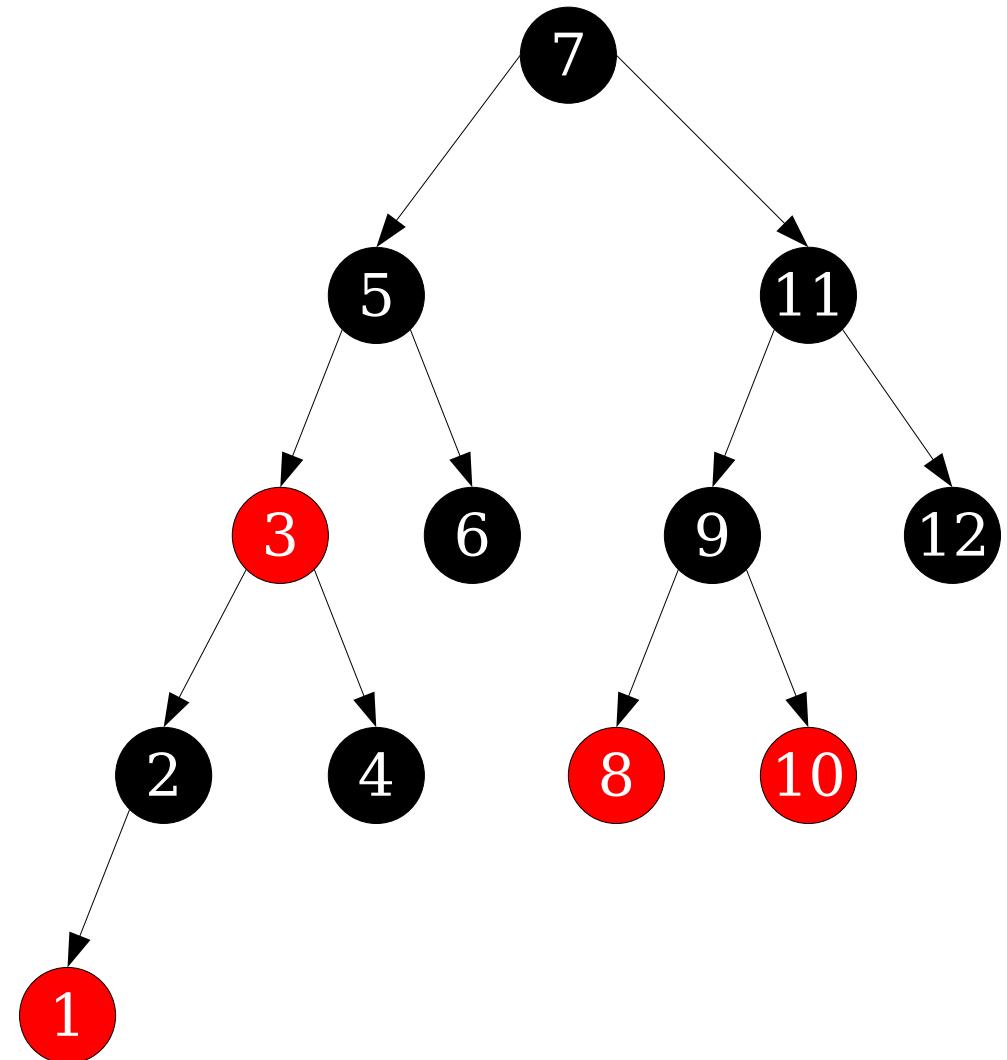
The Story So Far

- We've just built a simple, elegant, balanced multiway tree structure.
- We can use them as balanced trees in main memory (2-3-4 trees).
- We can use them to store huge quantities of information on disk (B-trees).
- We've seen that different cost models are appropriate in different situations.

So... red/black trees?

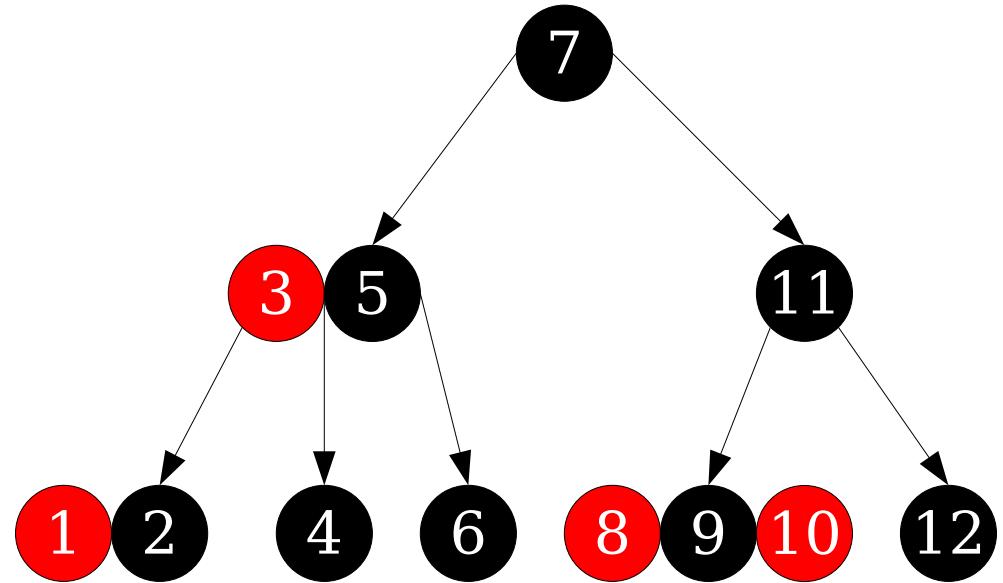
Red/Black Trees

- A **red/black tree** is a BST with the following properties:
 - Every node is either red or black.
 - The root is black.
 - No red node has a red child.
 - Every root-null path in the tree passes through the same number of black nodes.



Red/Black Trees

- A **red/black tree** is a BST with the following properties:
 - Every node is either red or black.
 - The root is black.
 - No red node has a red child.
 - Every root-null path in the tree passes through the same number of black nodes.
- After we hoist red nodes into their parents:
 - Each “meta node” has 1, 2, or 3 keys in it. (No red node has a red child.)
 - Each “meta node” is either a leaf or has one more child than key. (Root-null path property.)
 - Each “meta leaf” is at the same depth. (Root-null path property.)

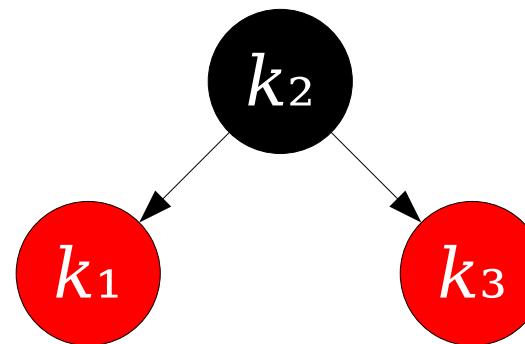
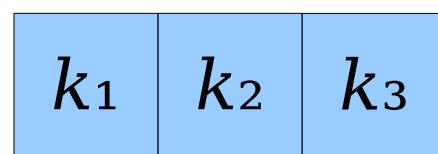
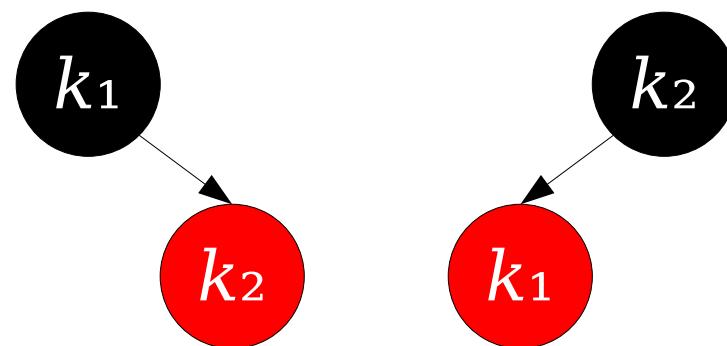
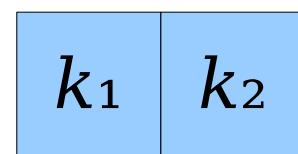
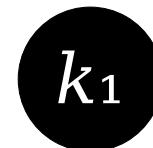
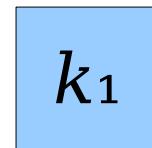


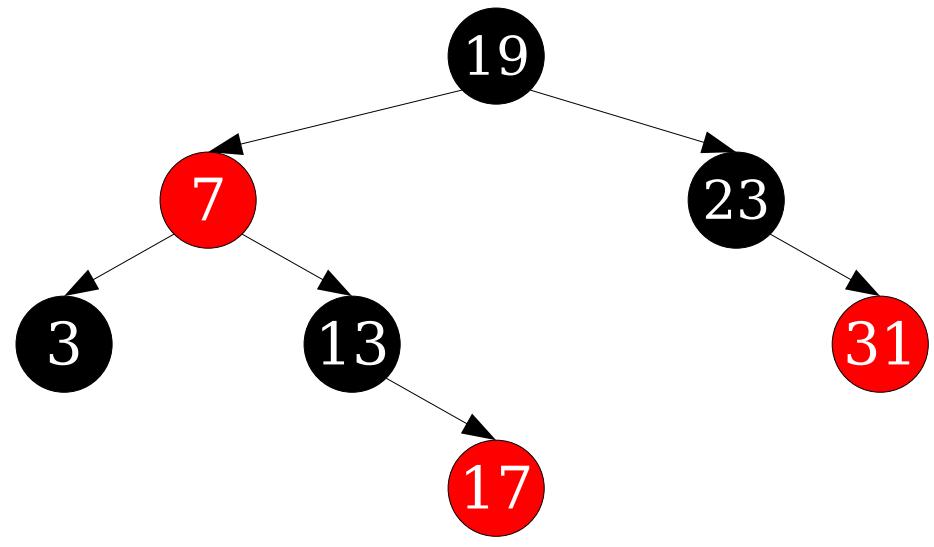
***This is a
2-3-4 tree!***

Data Structure Isometries

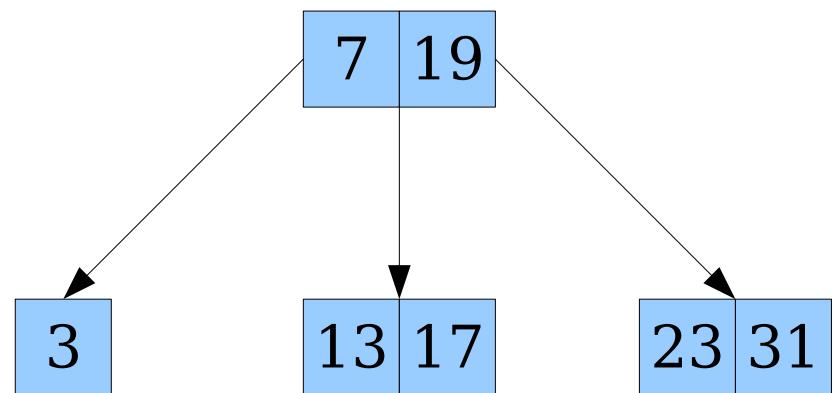
- Red/black trees are an *isometry* of 2-3-4 trees; they represent the structure of 2-3-4 trees in a different way.
- Many data structures can be designed and analyzed in the same way.
- **Huge advantage:** Rather than memorizing a complex list of red/black tree rules, just think about what the equivalent operation on the corresponding 2-3-4 tree would be and simulate it with BST operations.

Exploring the Isometry

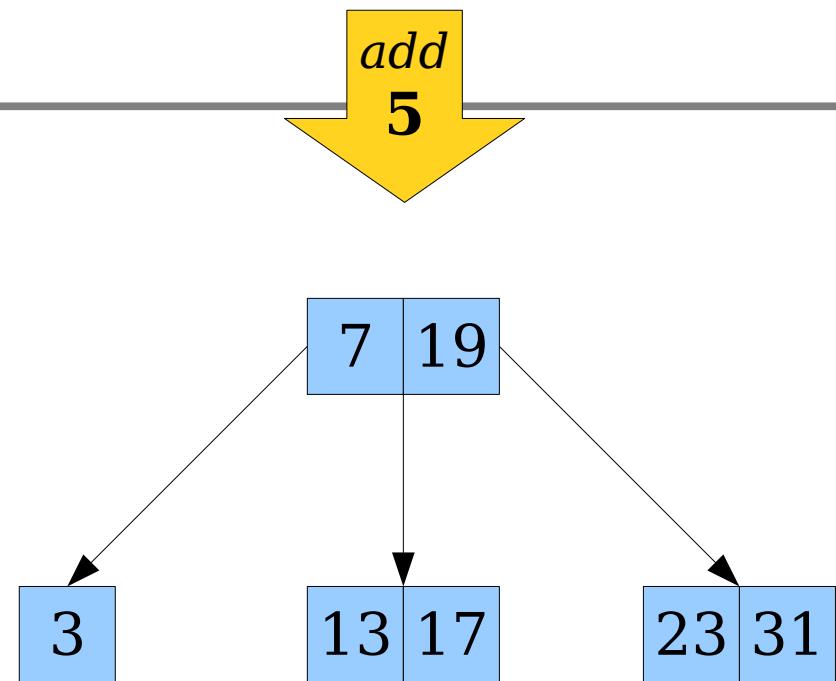
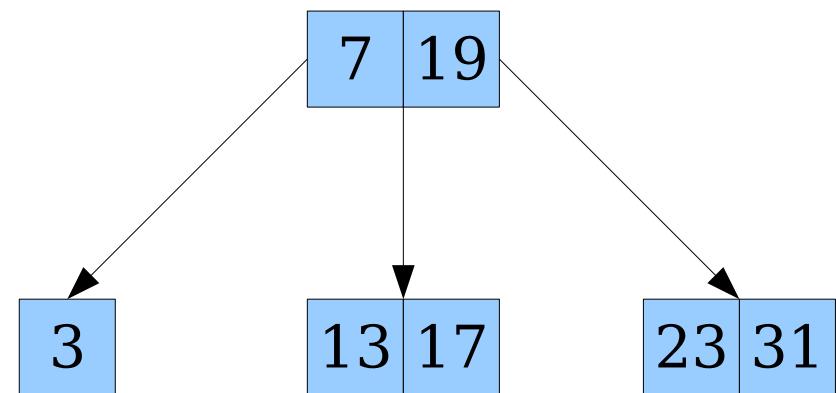
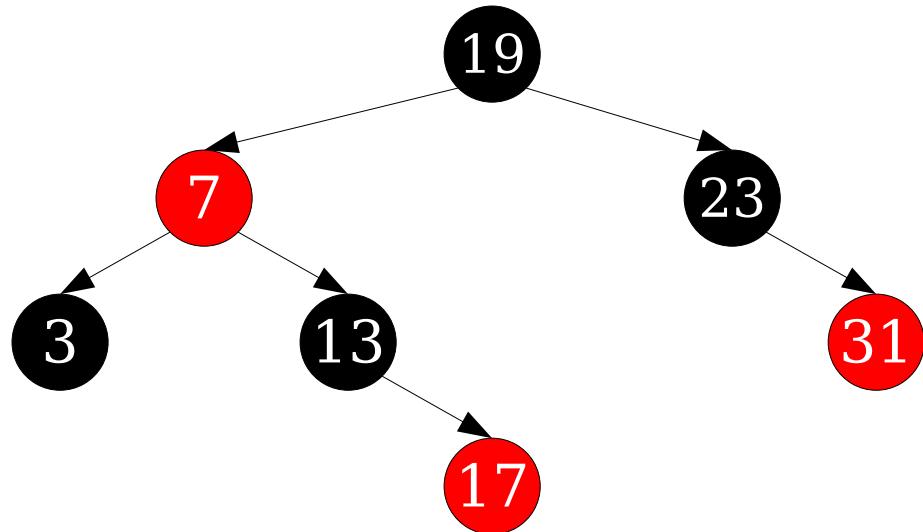


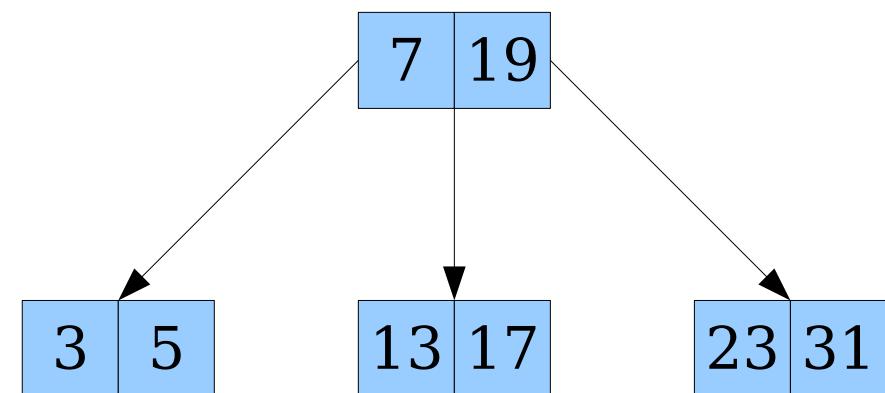
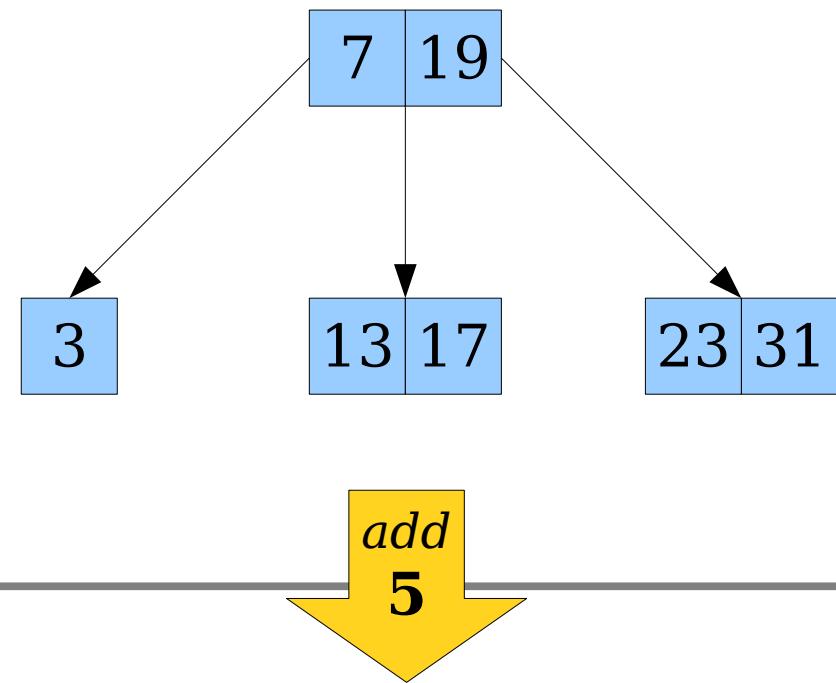
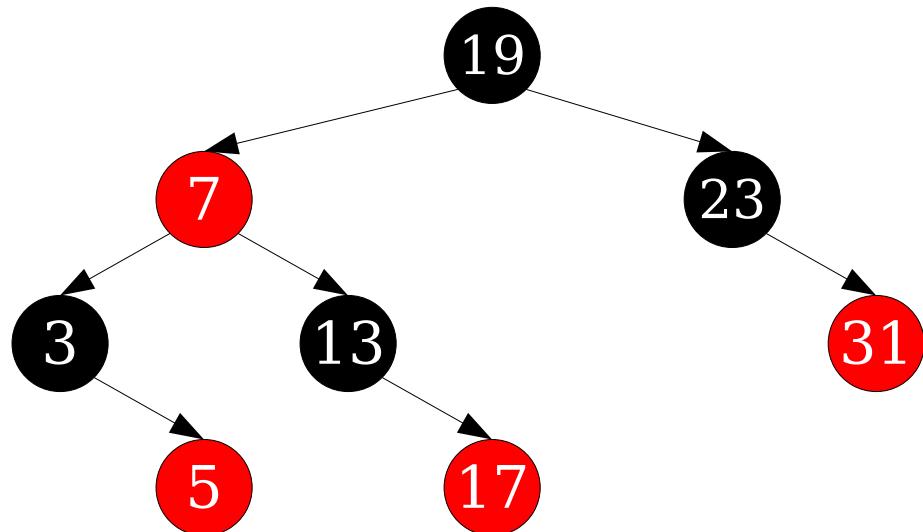
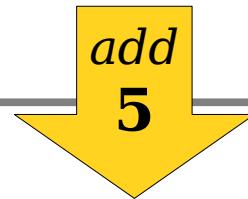
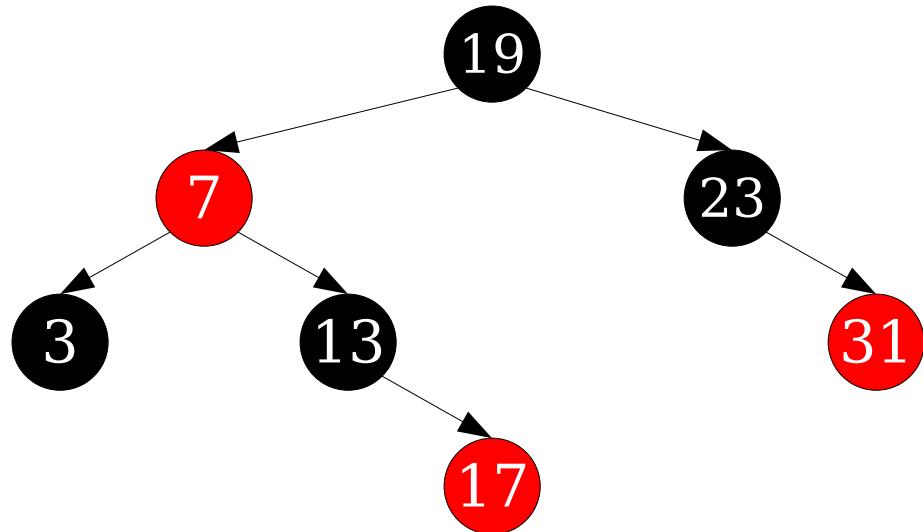


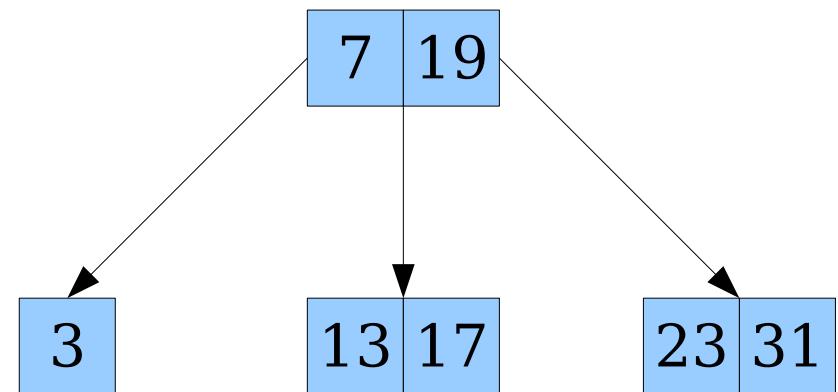
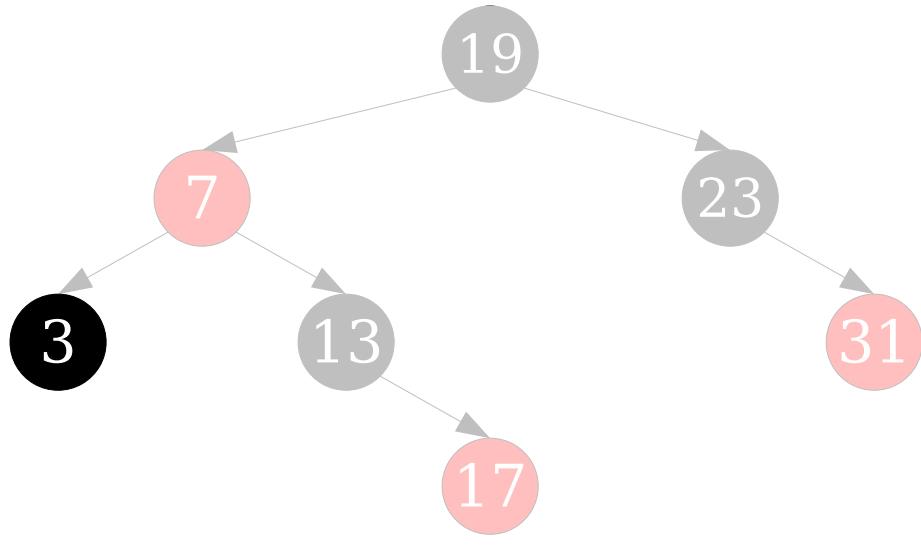
add
5



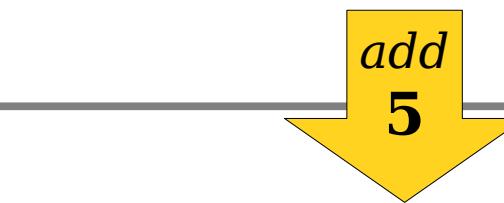
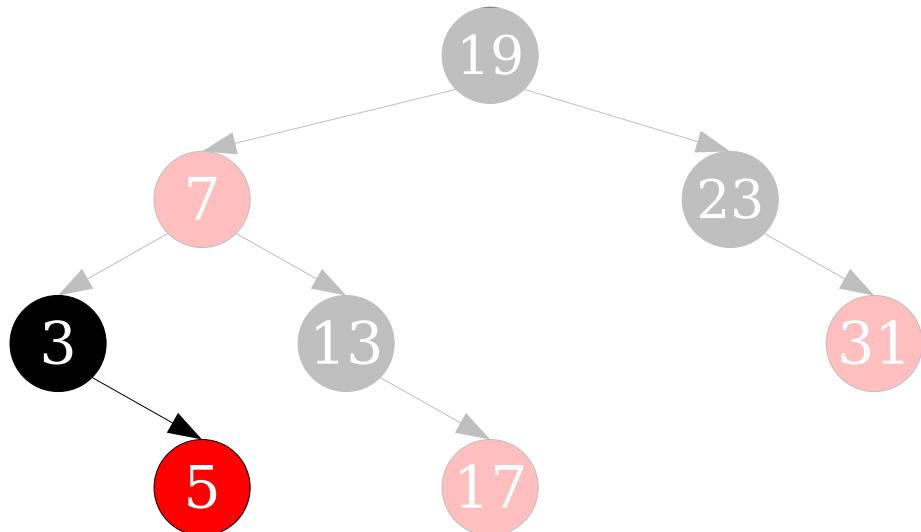
add
5



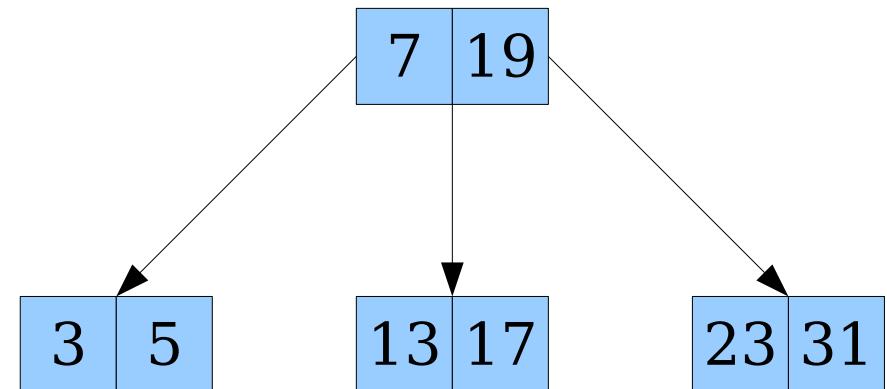


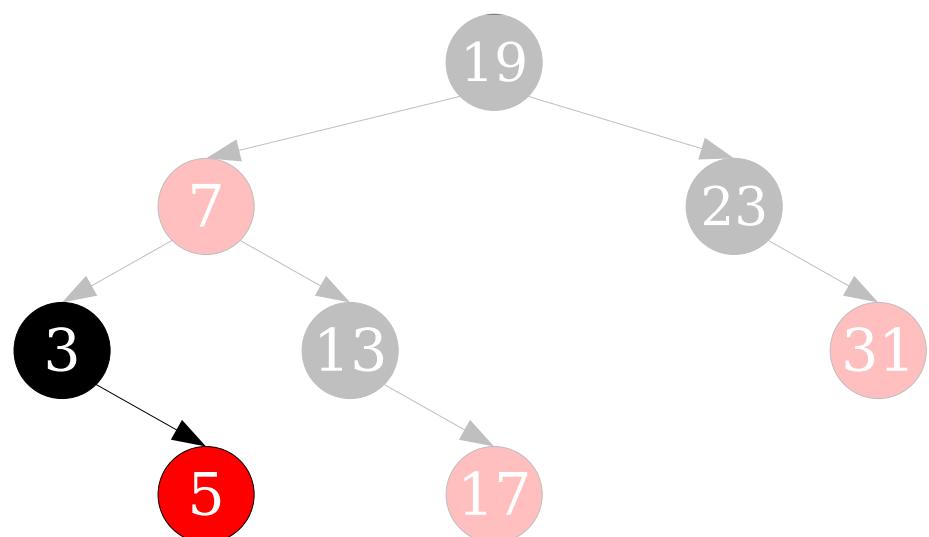
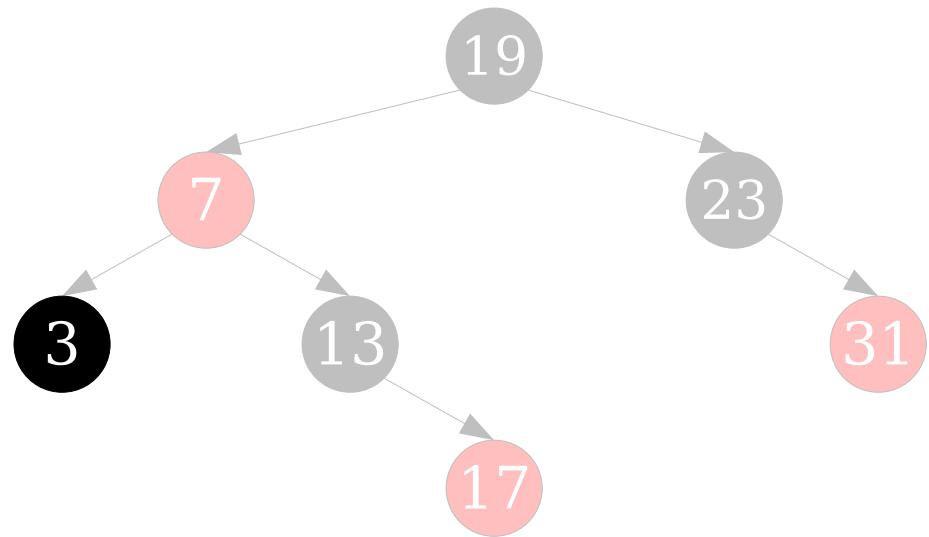


add
5

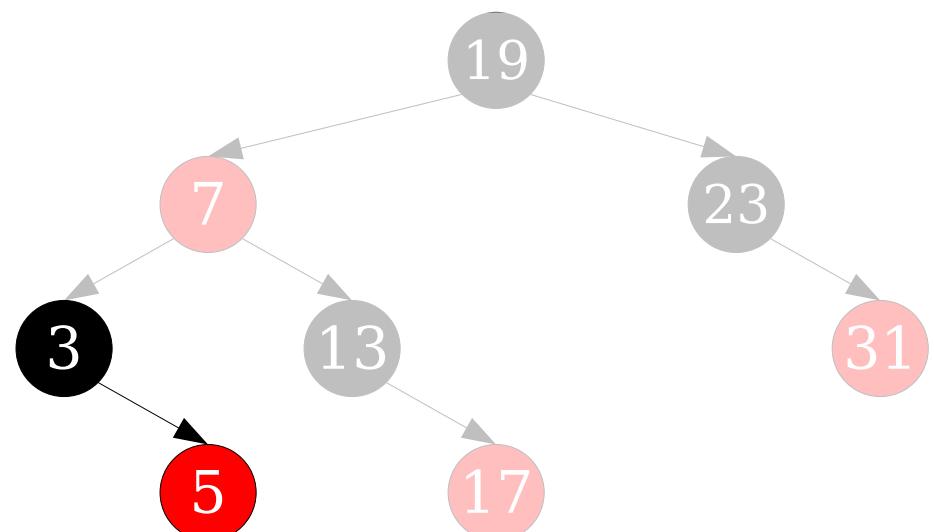
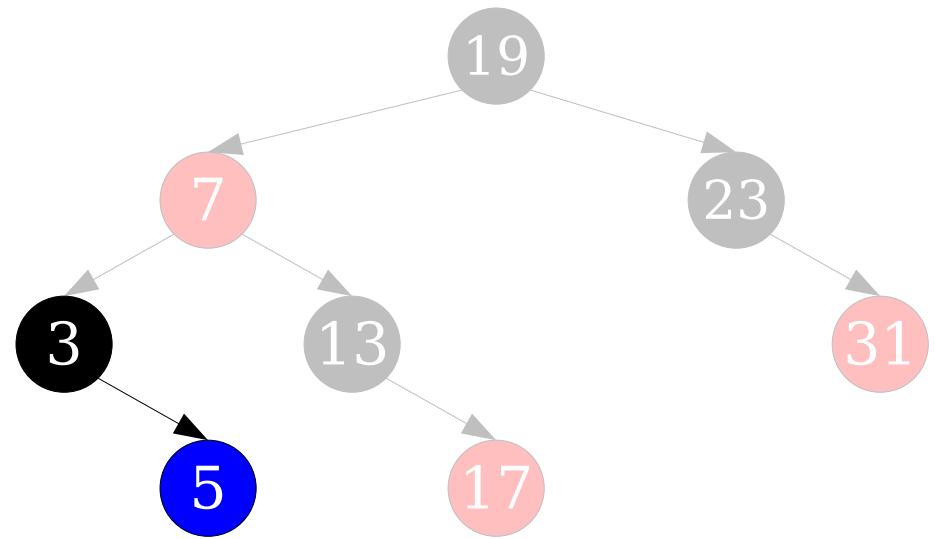


add
5

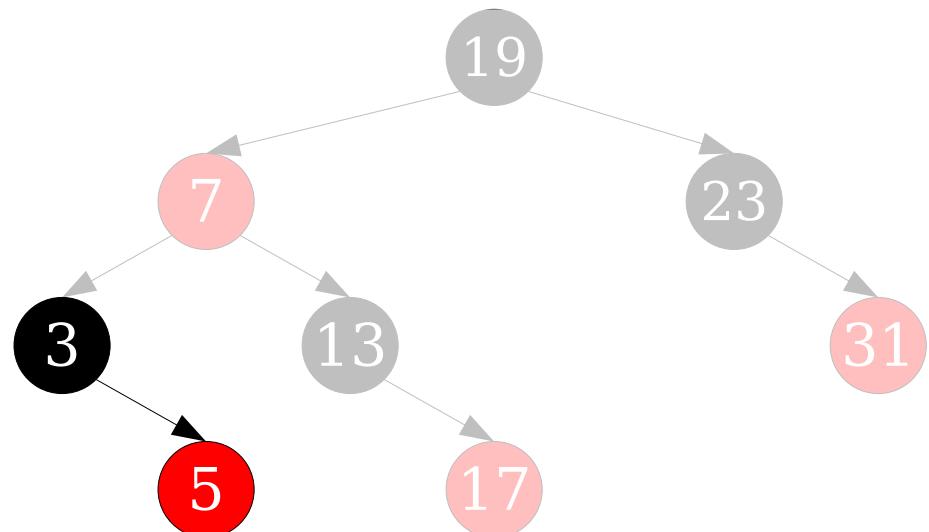
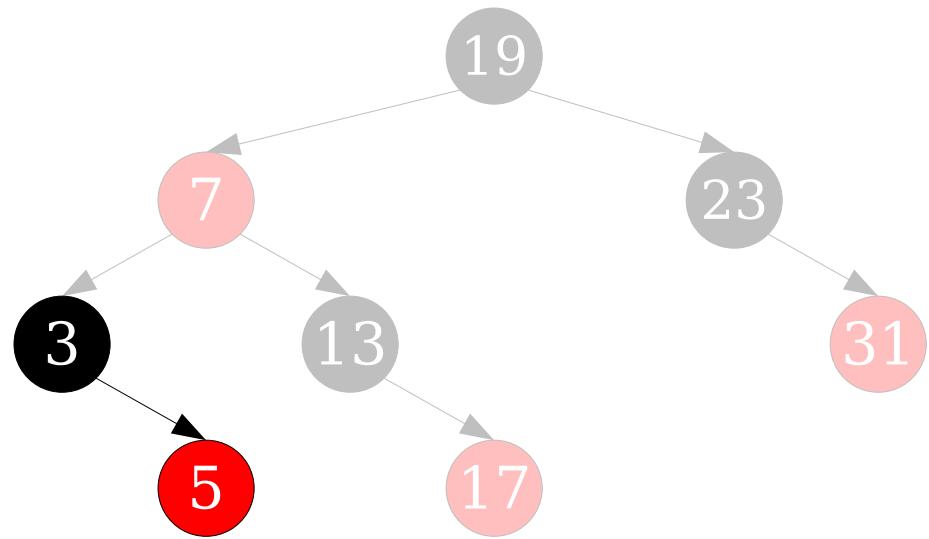




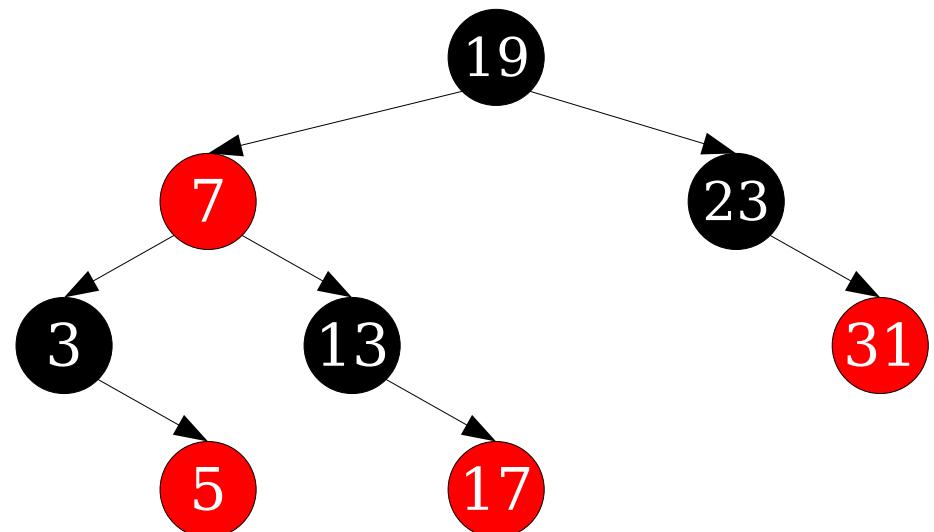
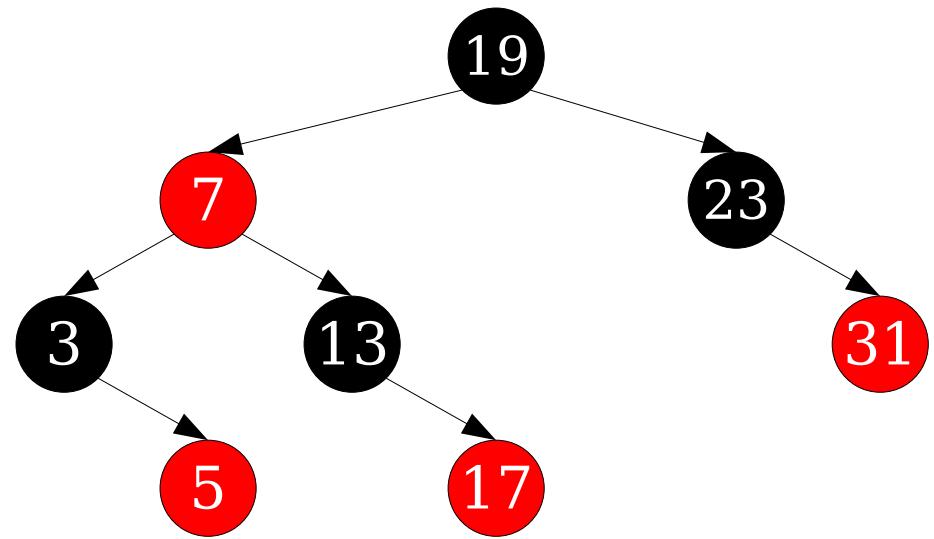
Goal



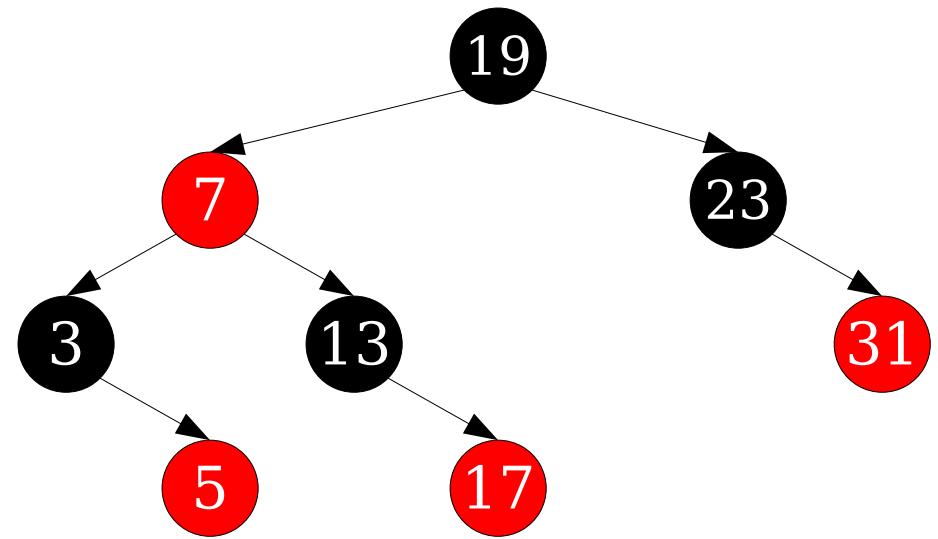
Goal

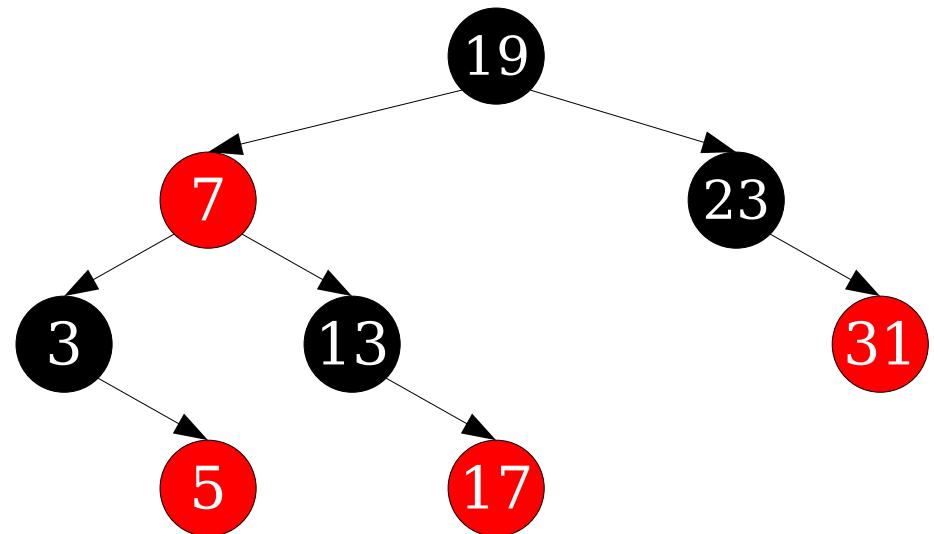


Goal

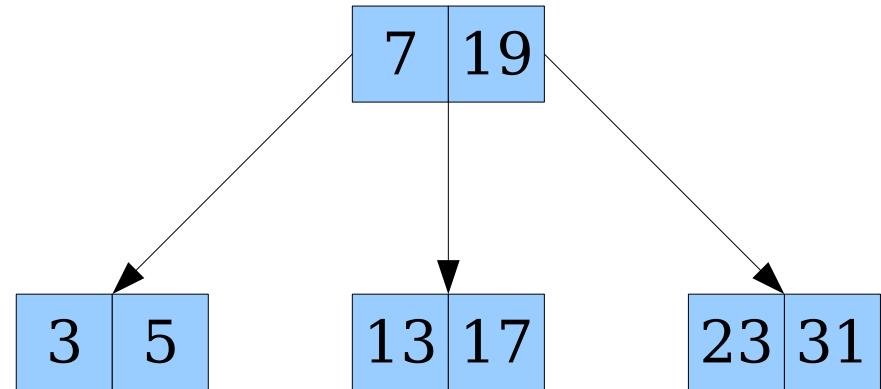


Goal

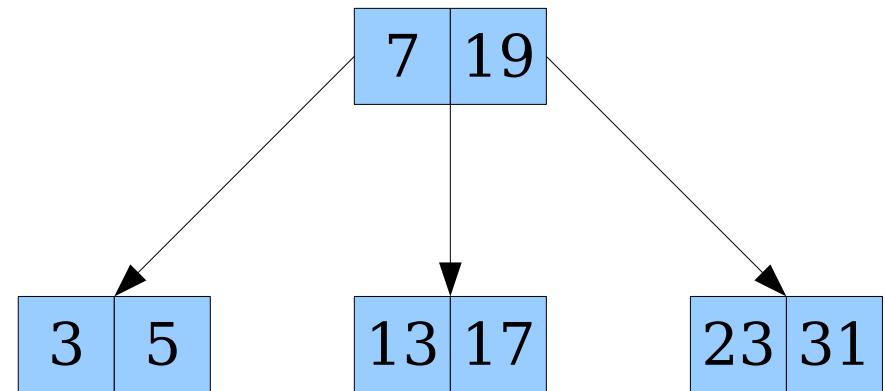
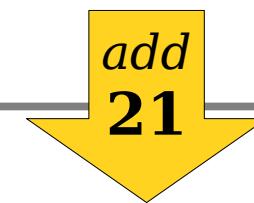
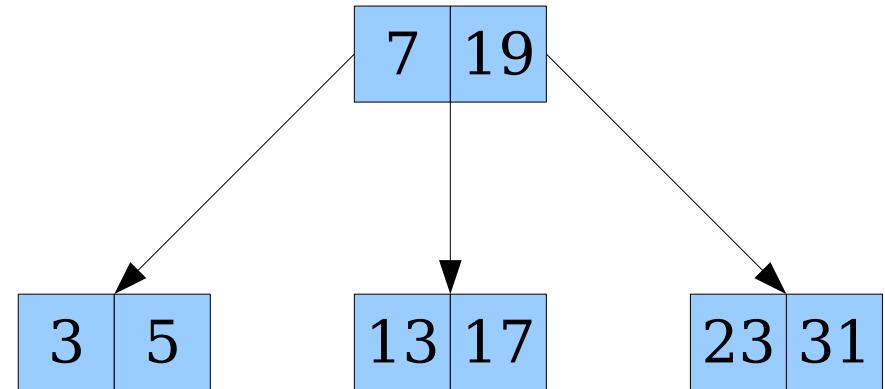
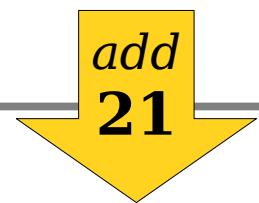
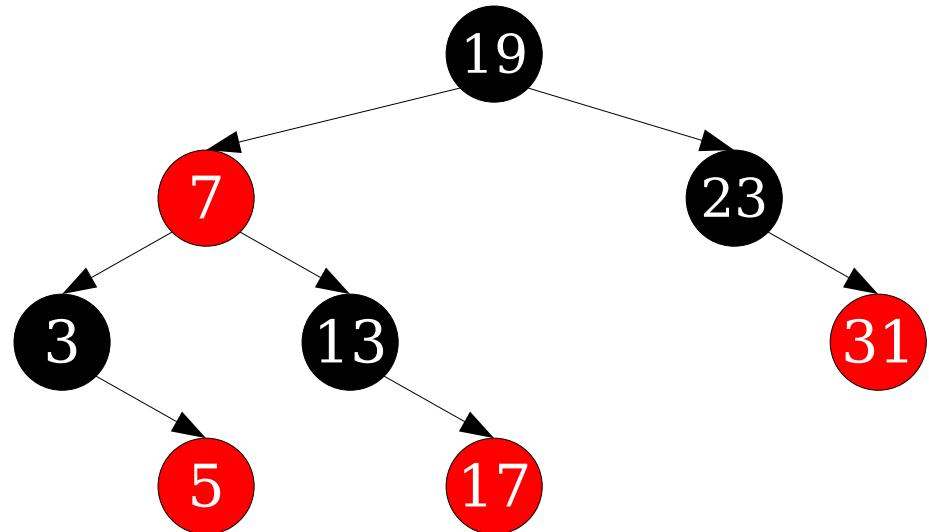


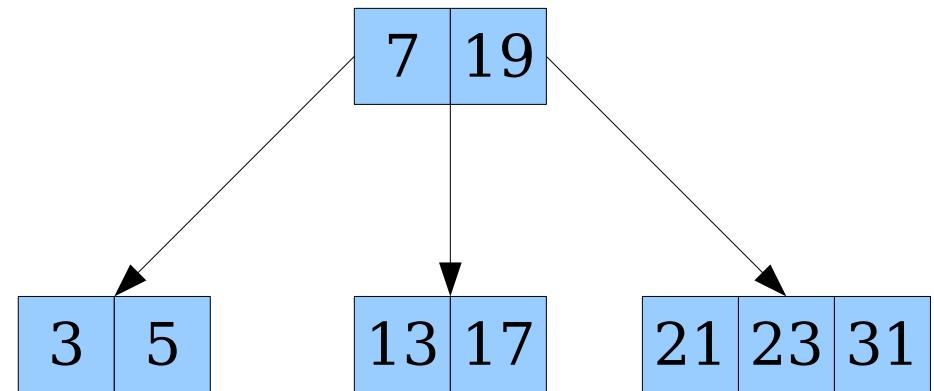
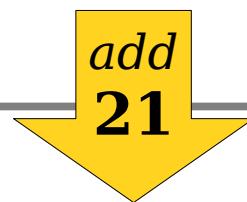
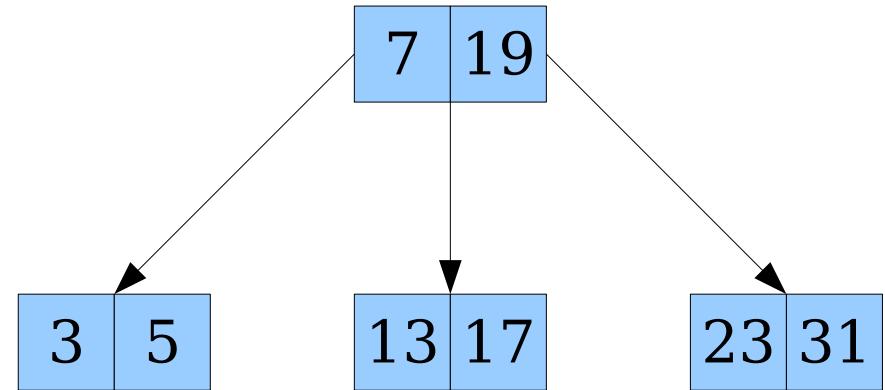
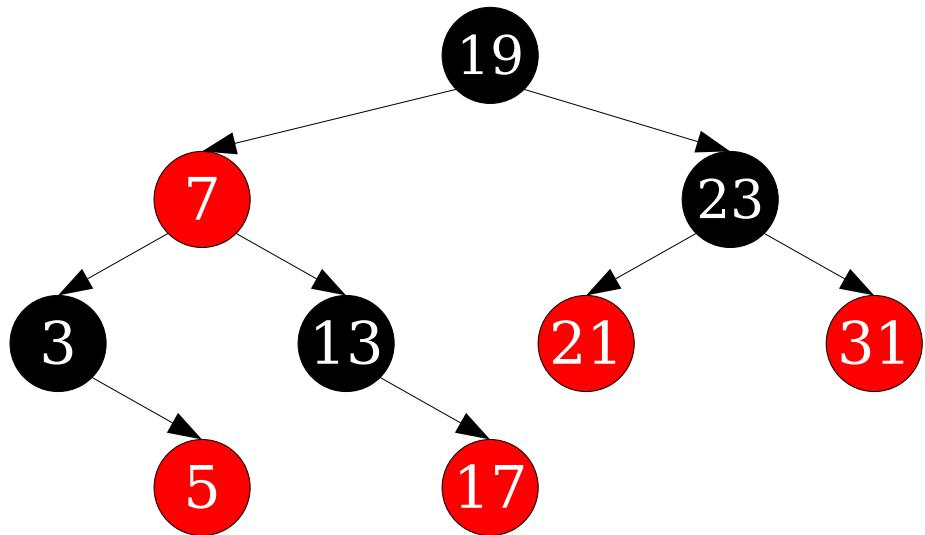
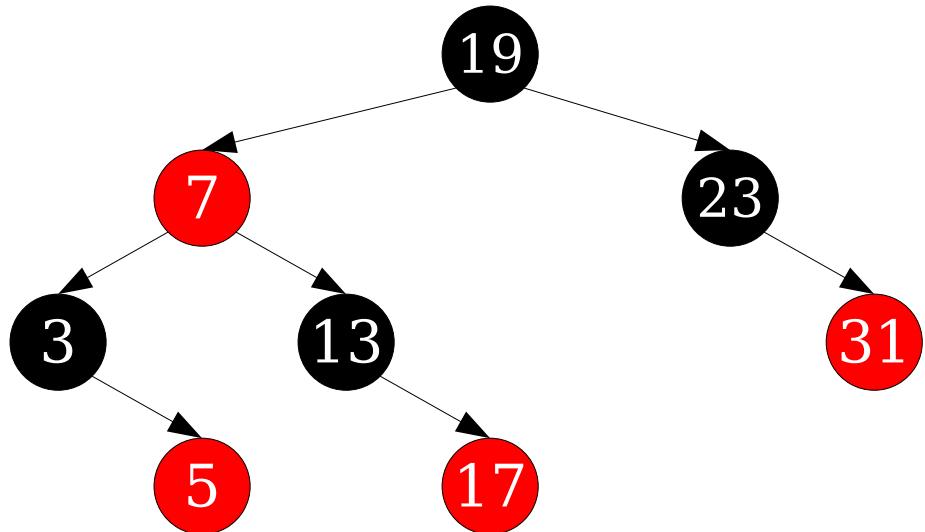


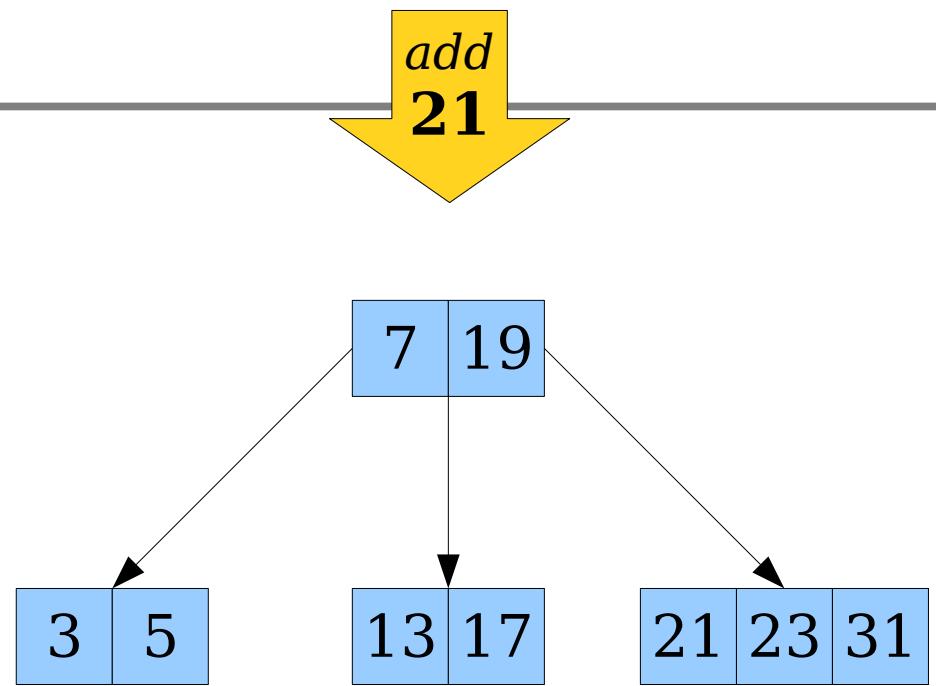
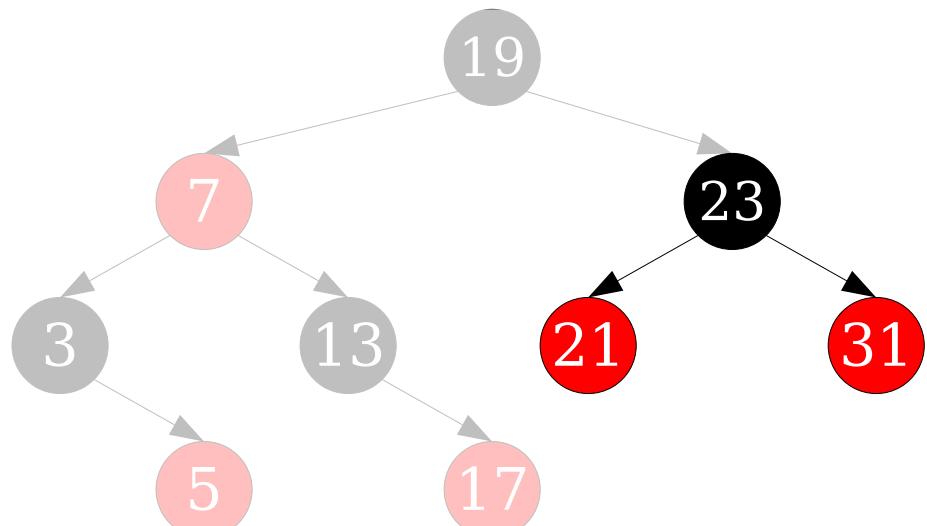
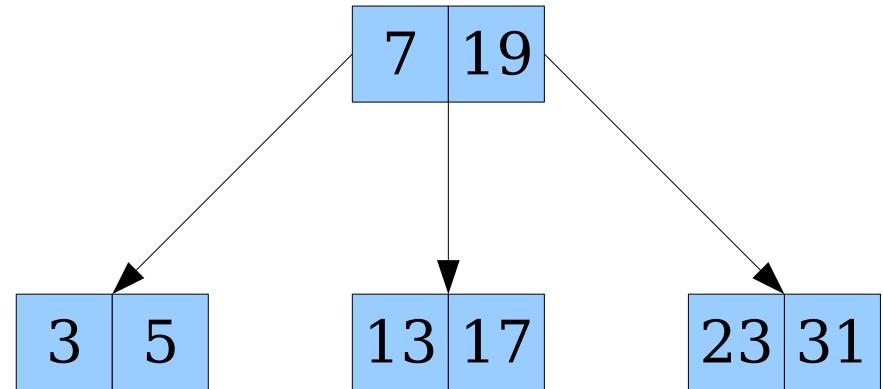
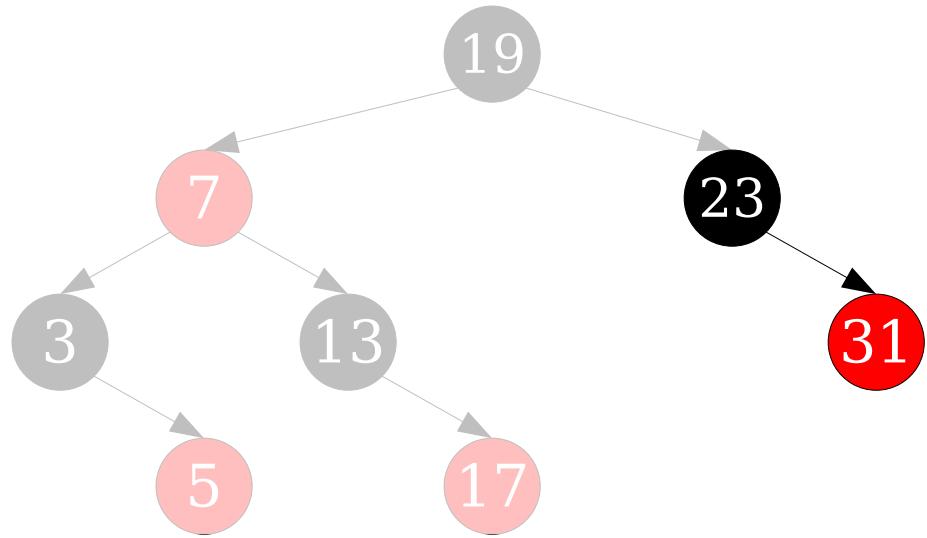
add
21

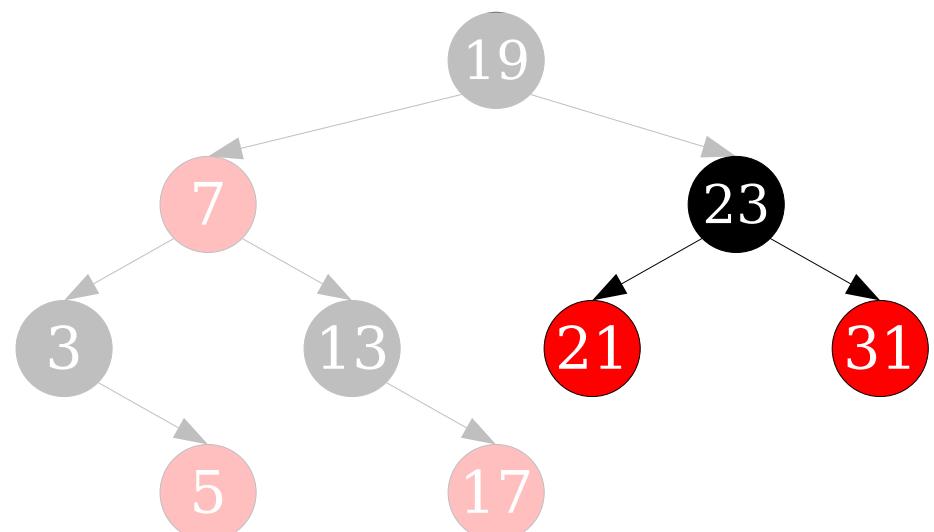
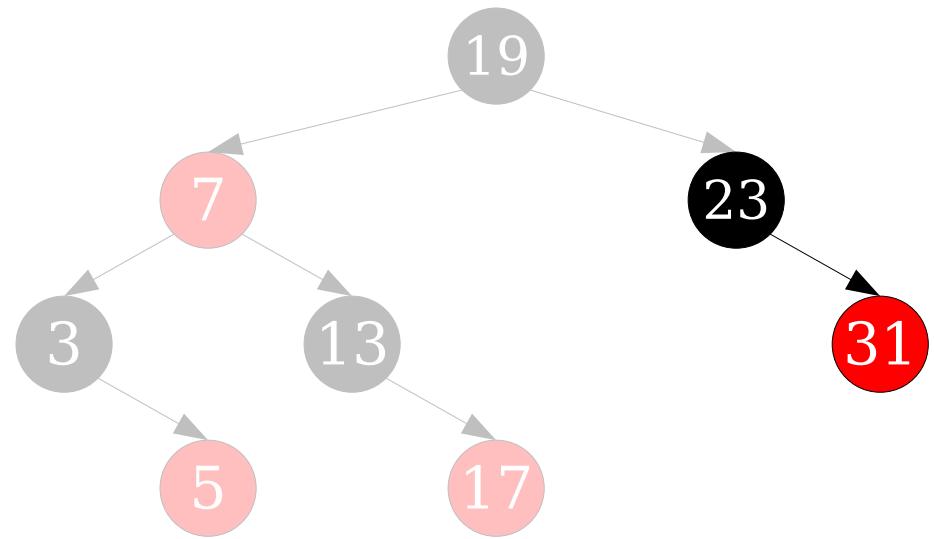


add
21

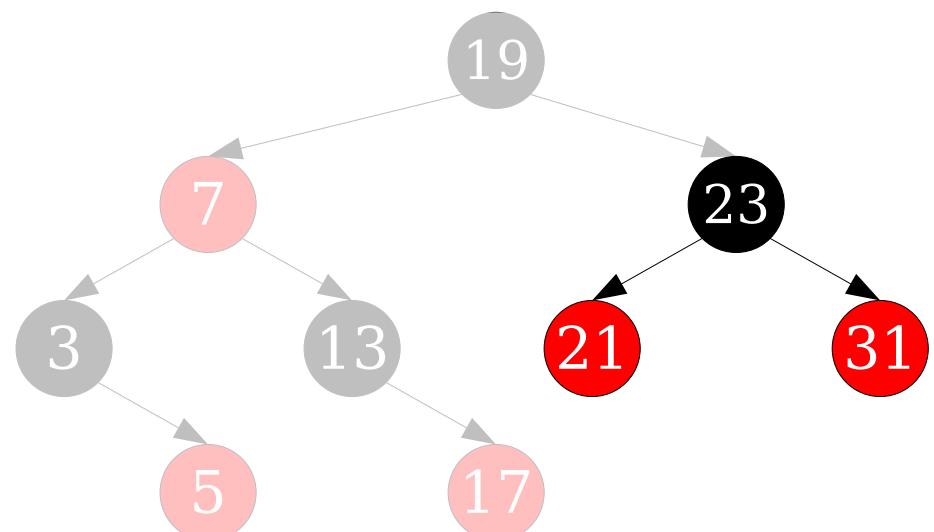
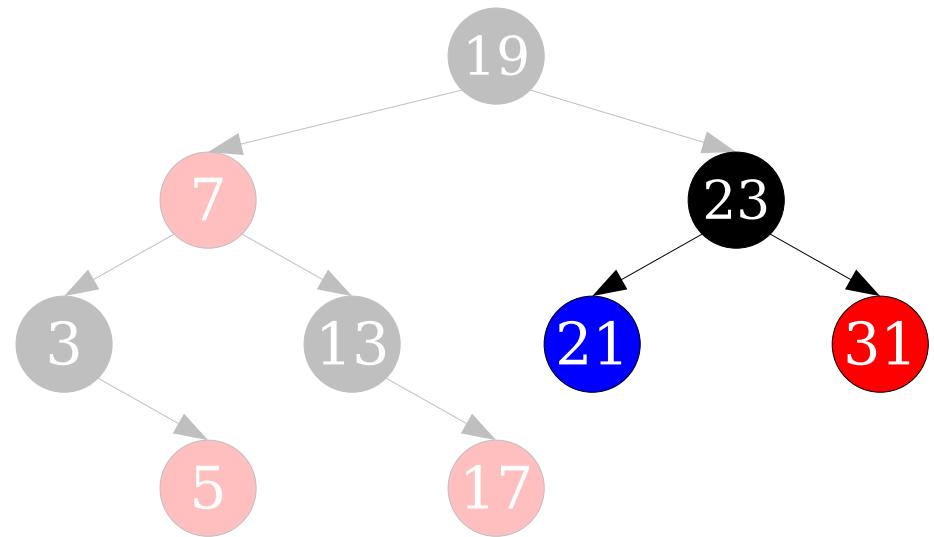




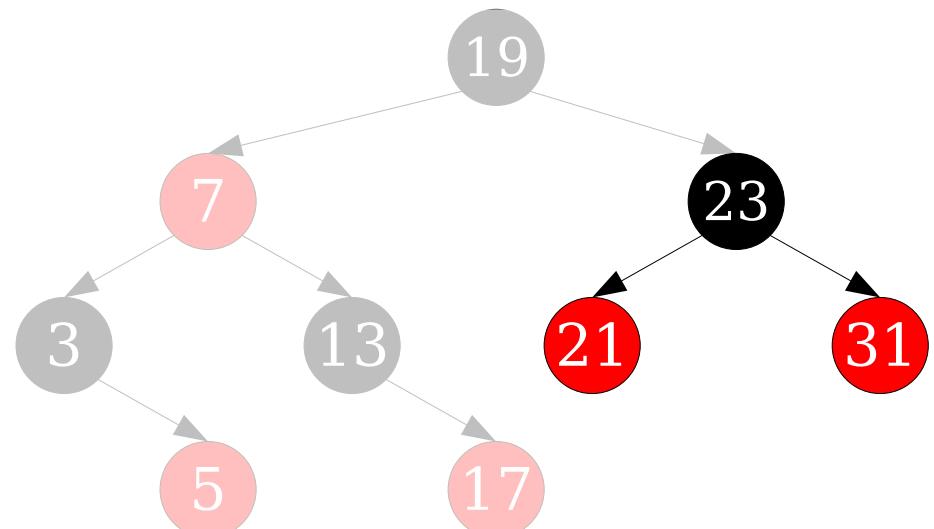
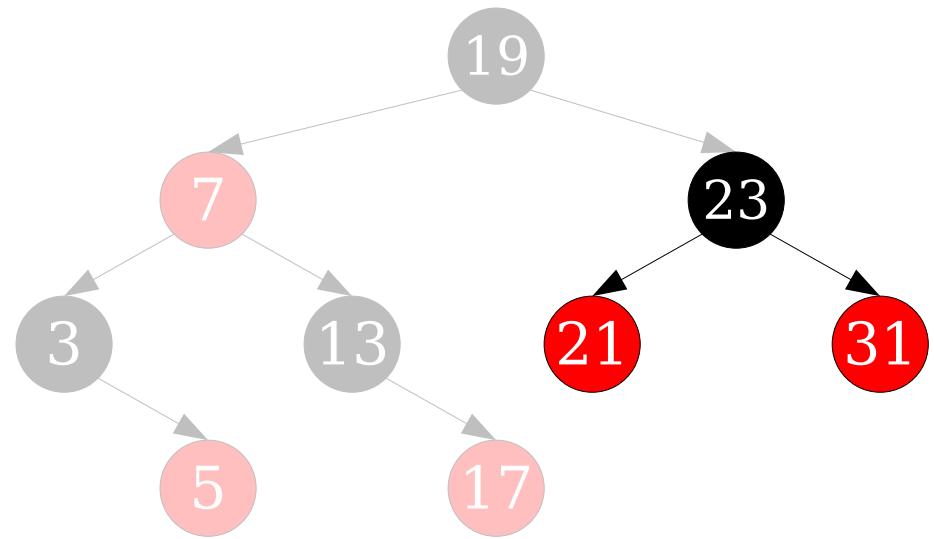




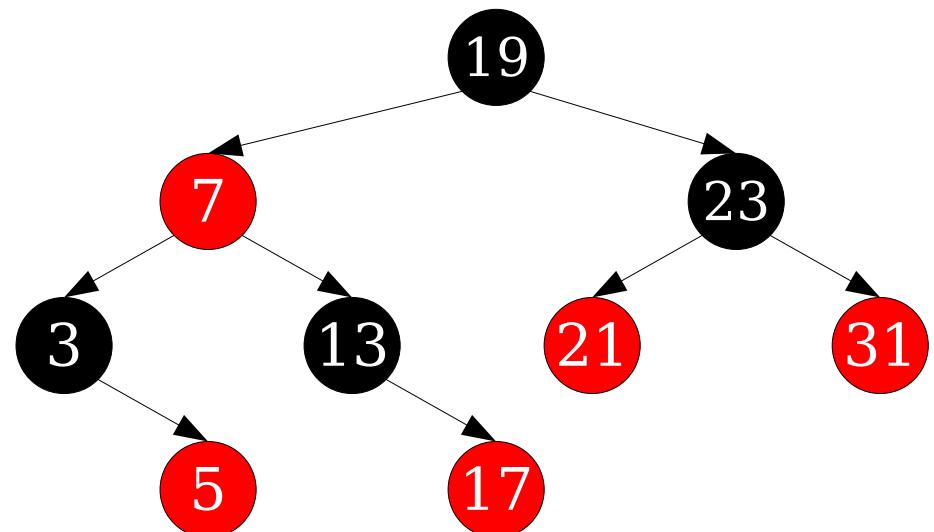
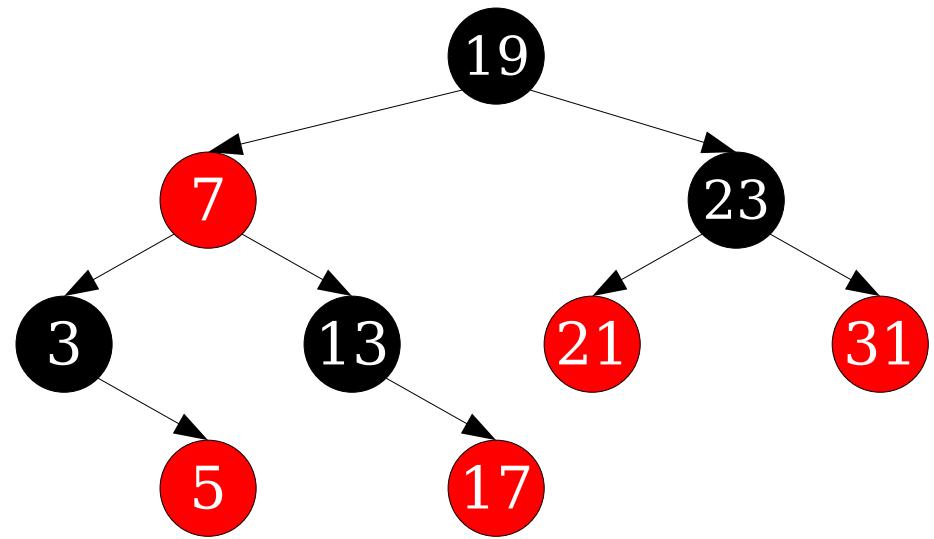
Goal



Goal



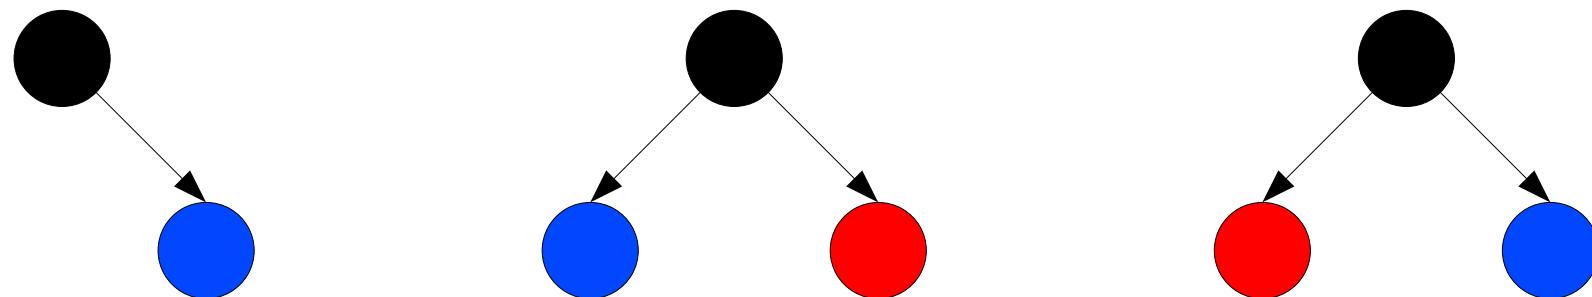
Goal

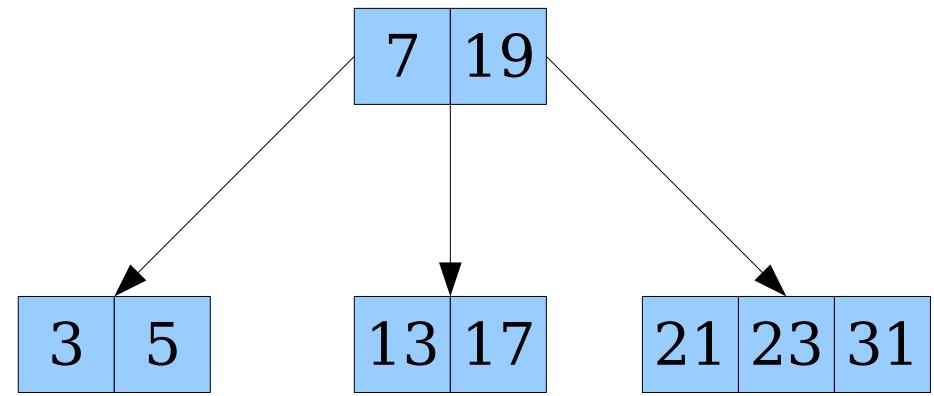
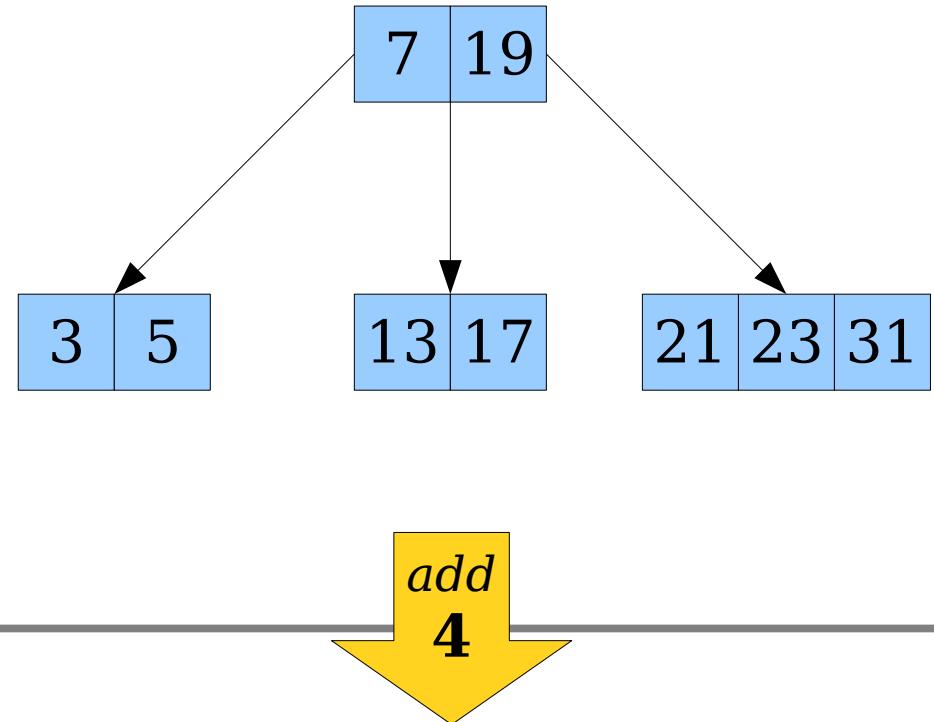
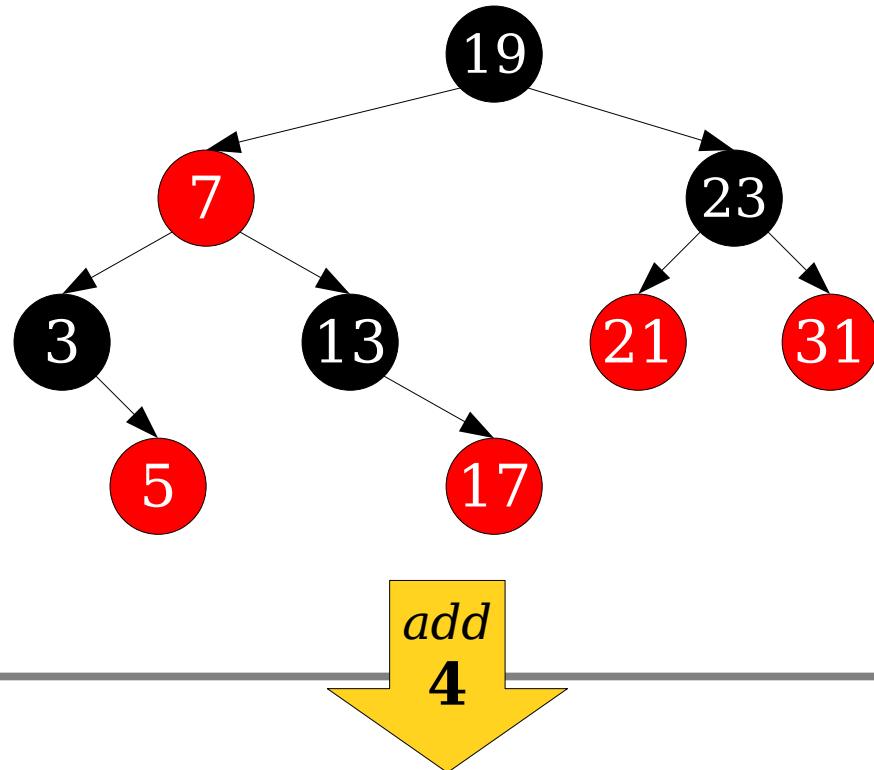


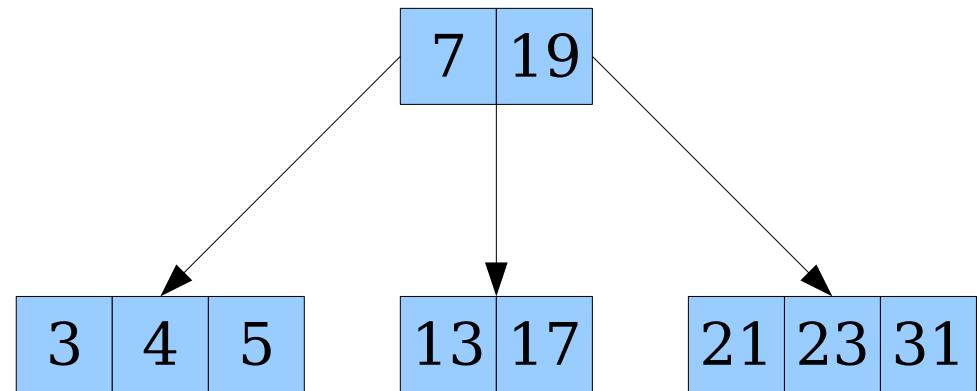
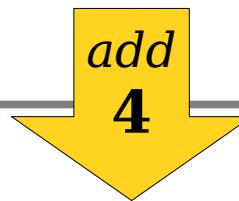
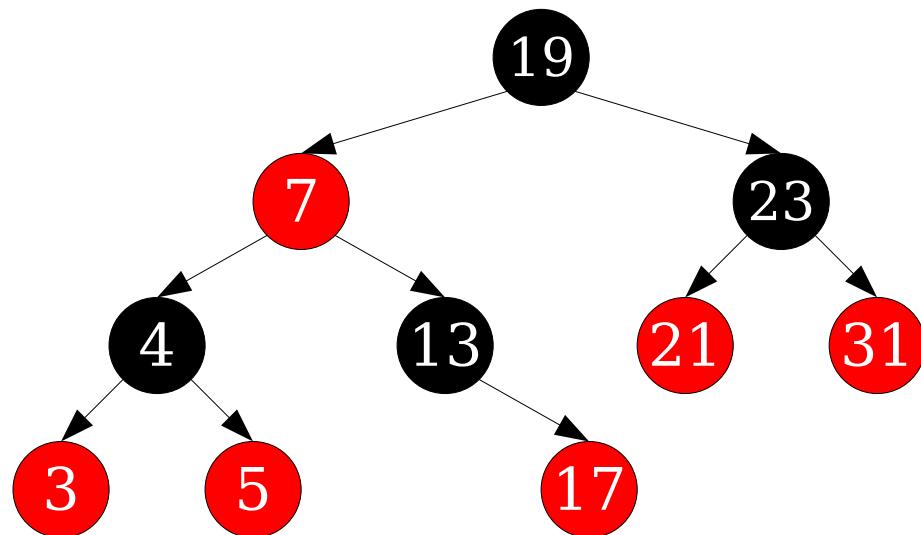
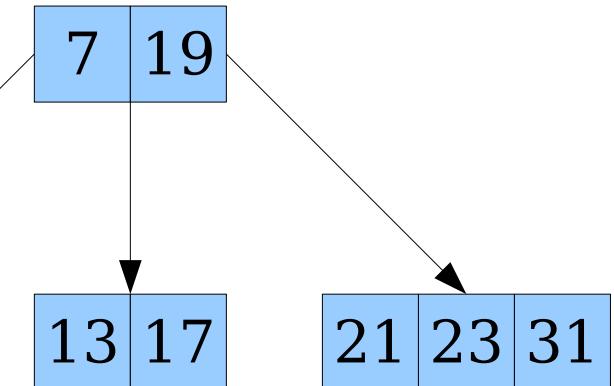
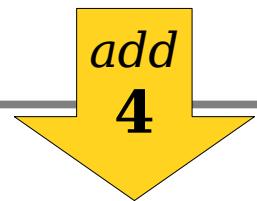
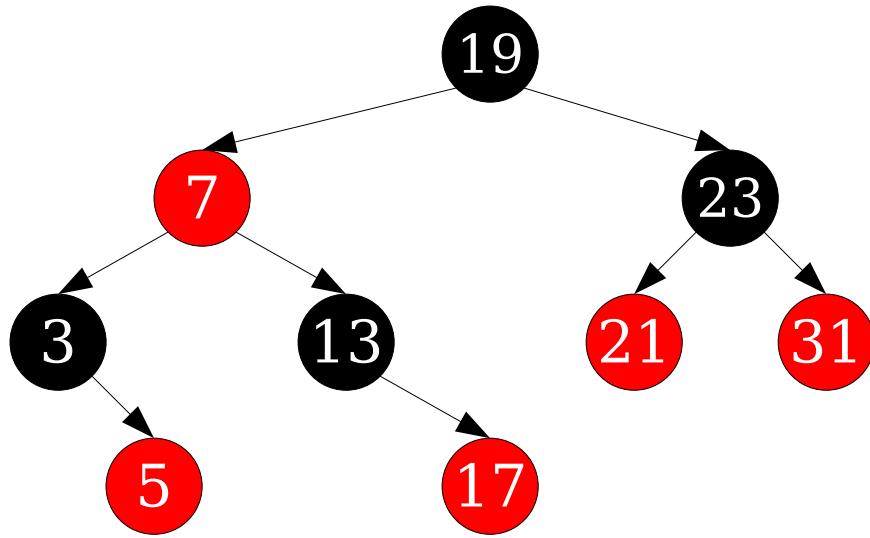
Goal

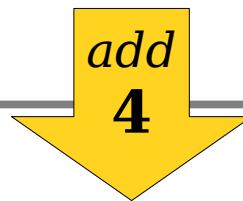
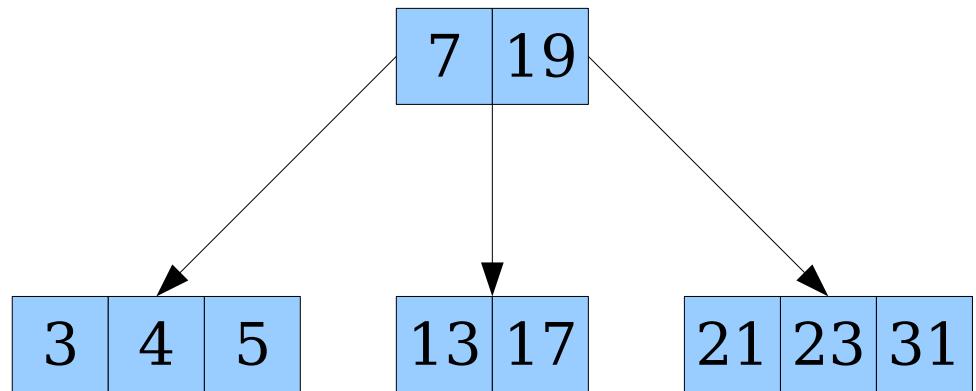
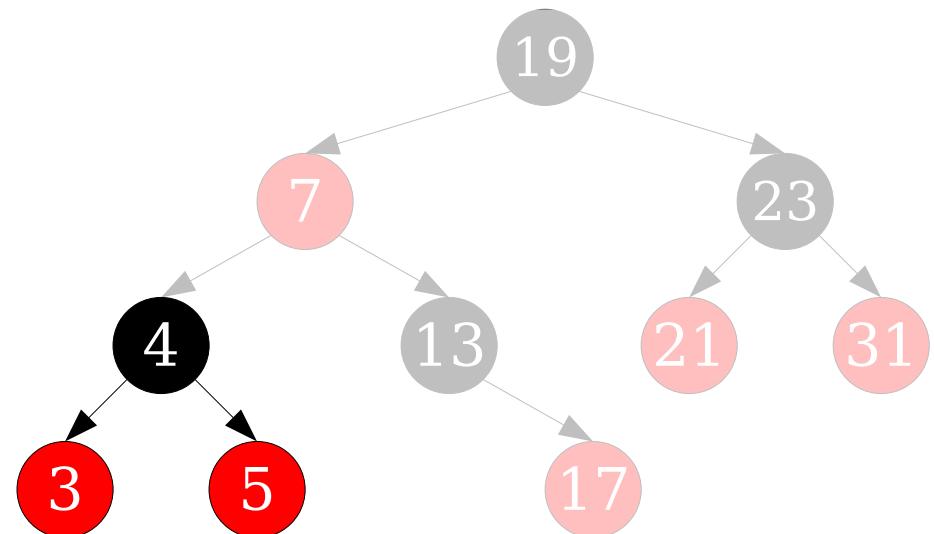
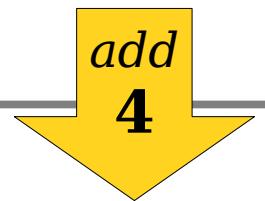
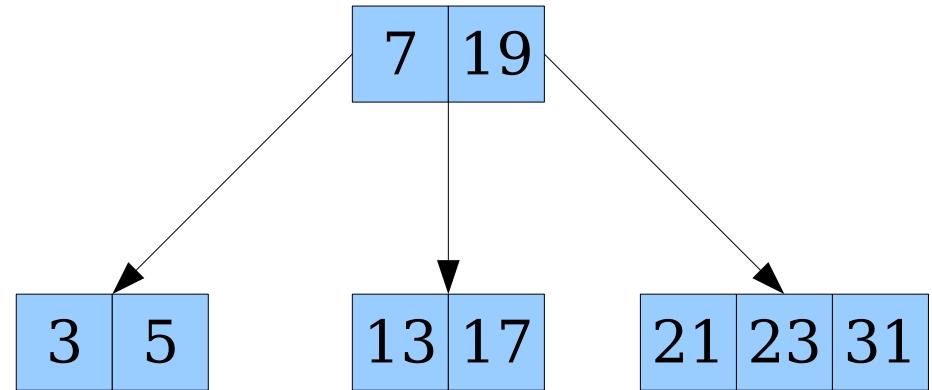
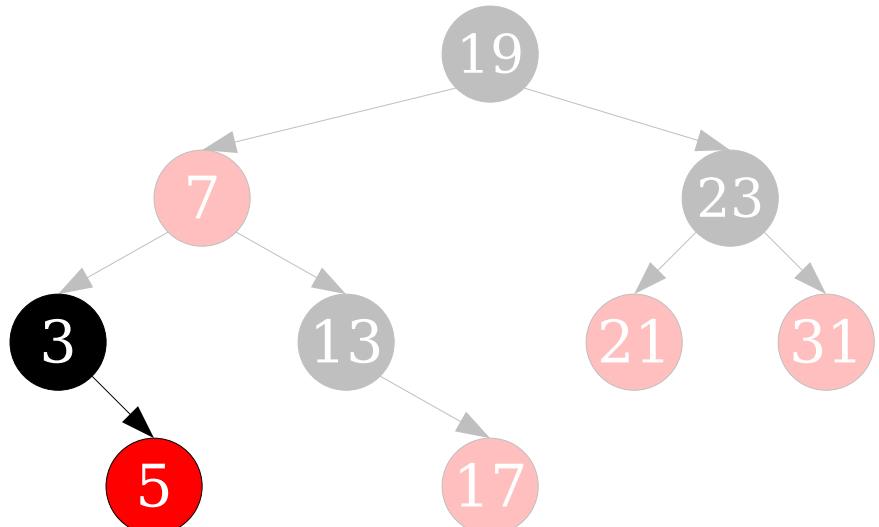
Red/Black Tree Insertion

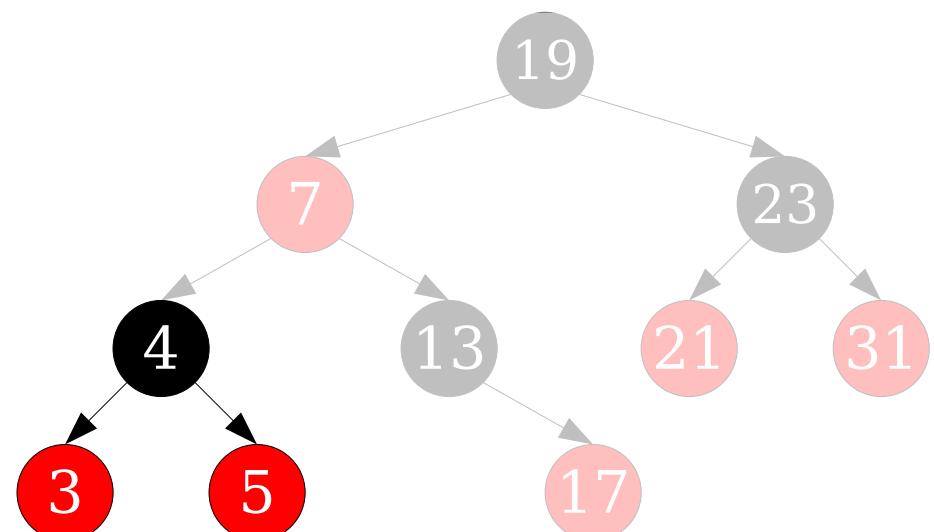
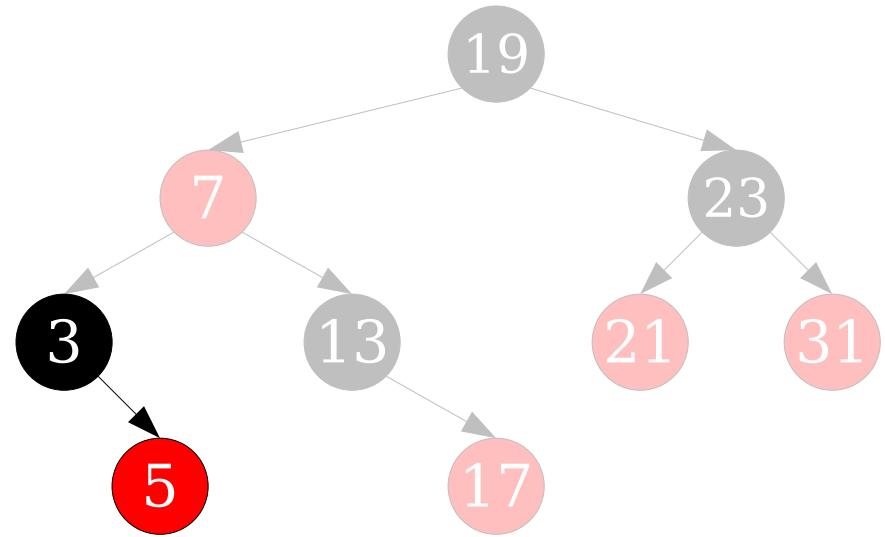
- **Rule #1:** When inserting a node, if its parent is black, make the node red and stop.
- **Justification:** This simulates inserting a key into an existing 2-node or 3-node.



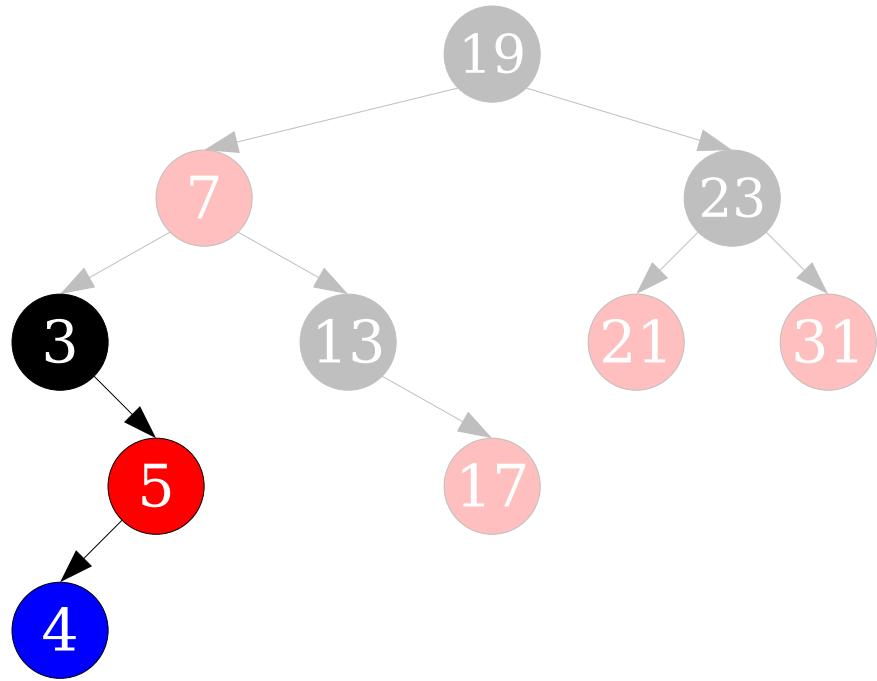




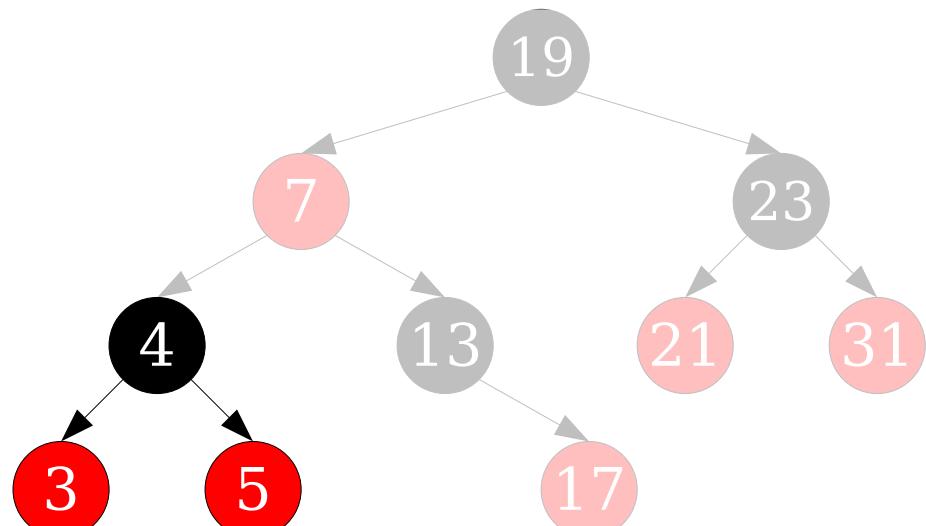




Goal

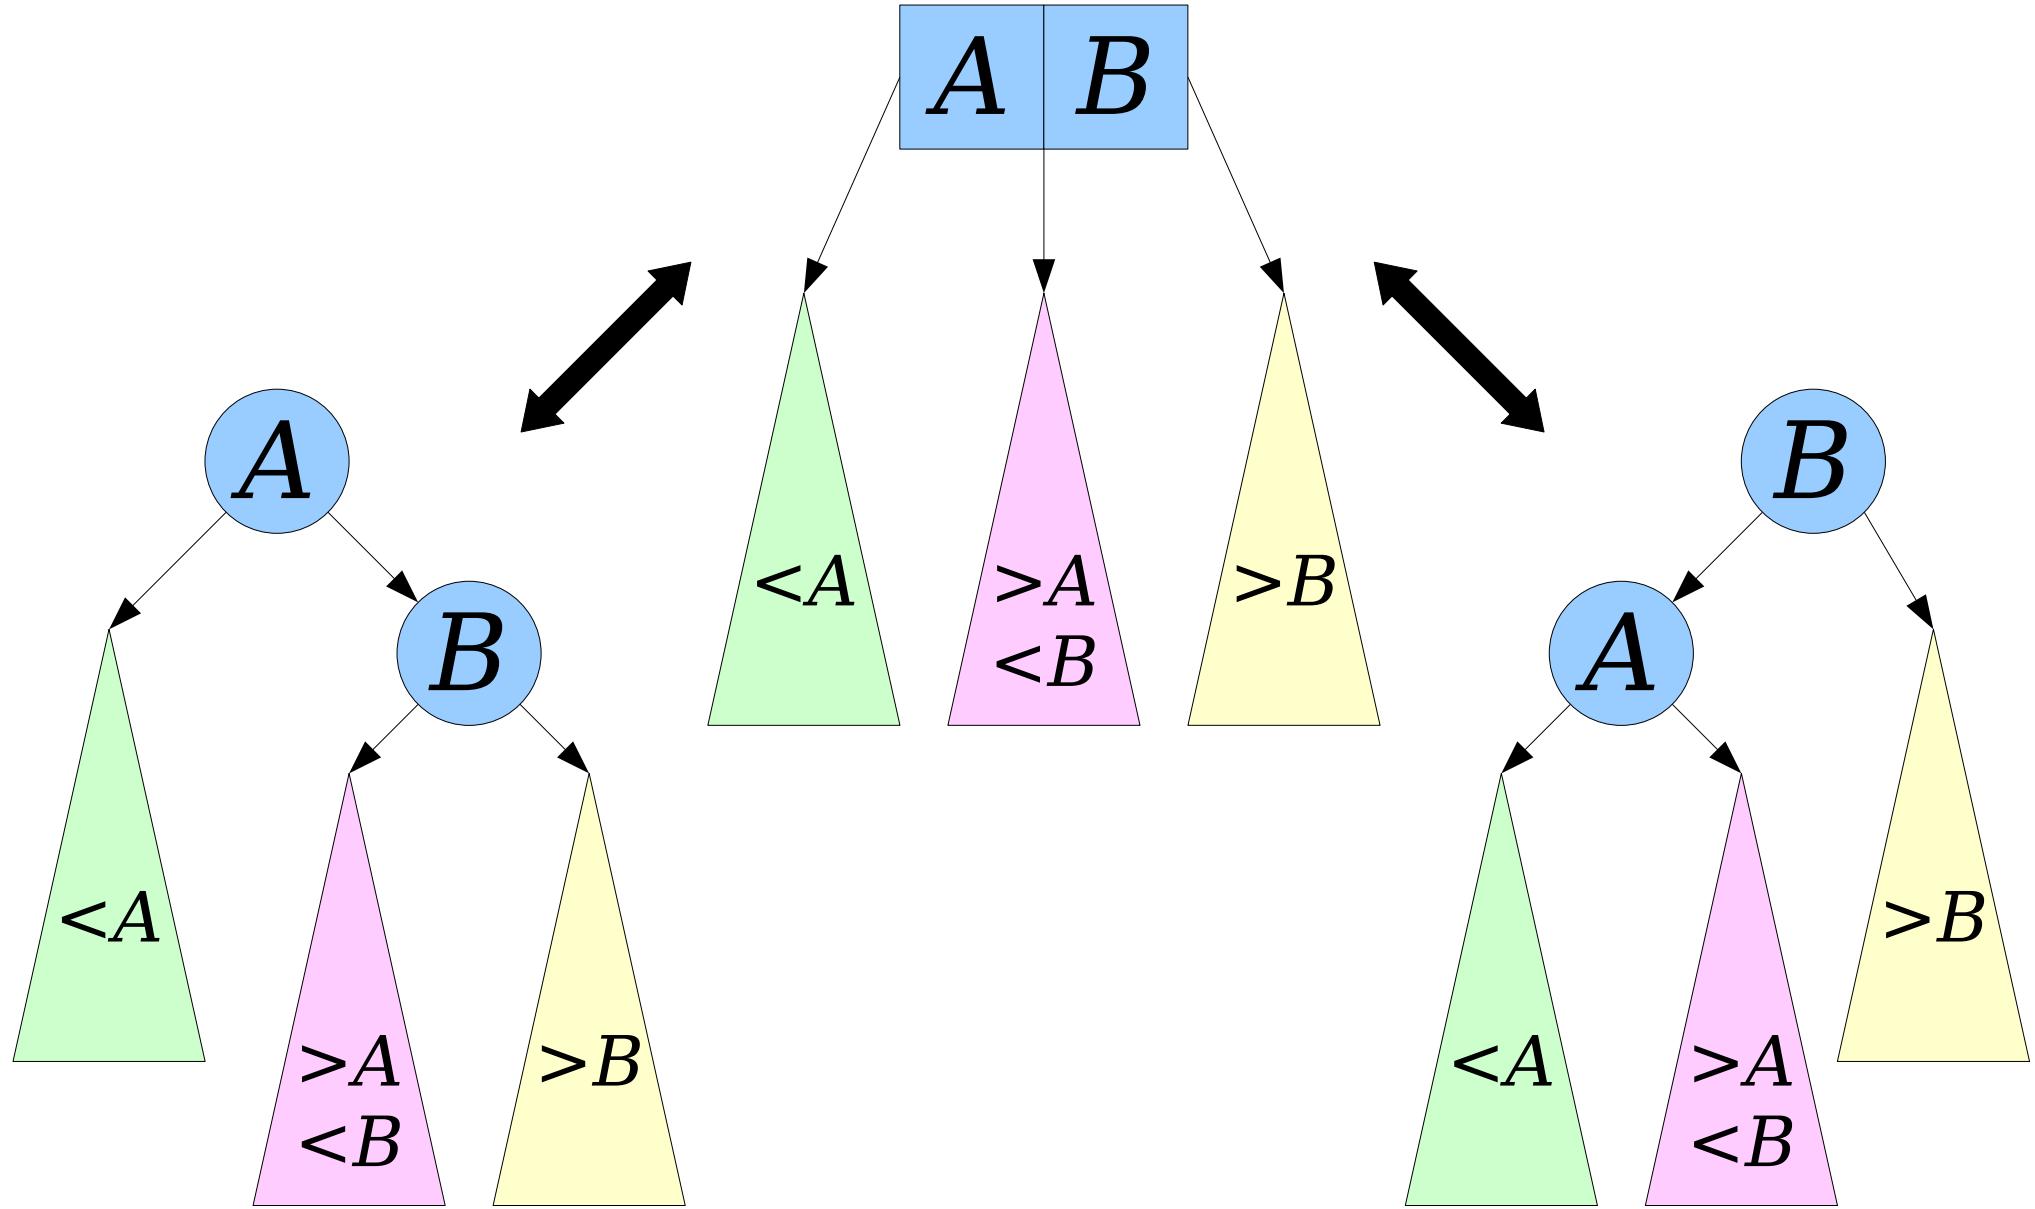


We need to move nodes around in a binary search tree. How do we do this?

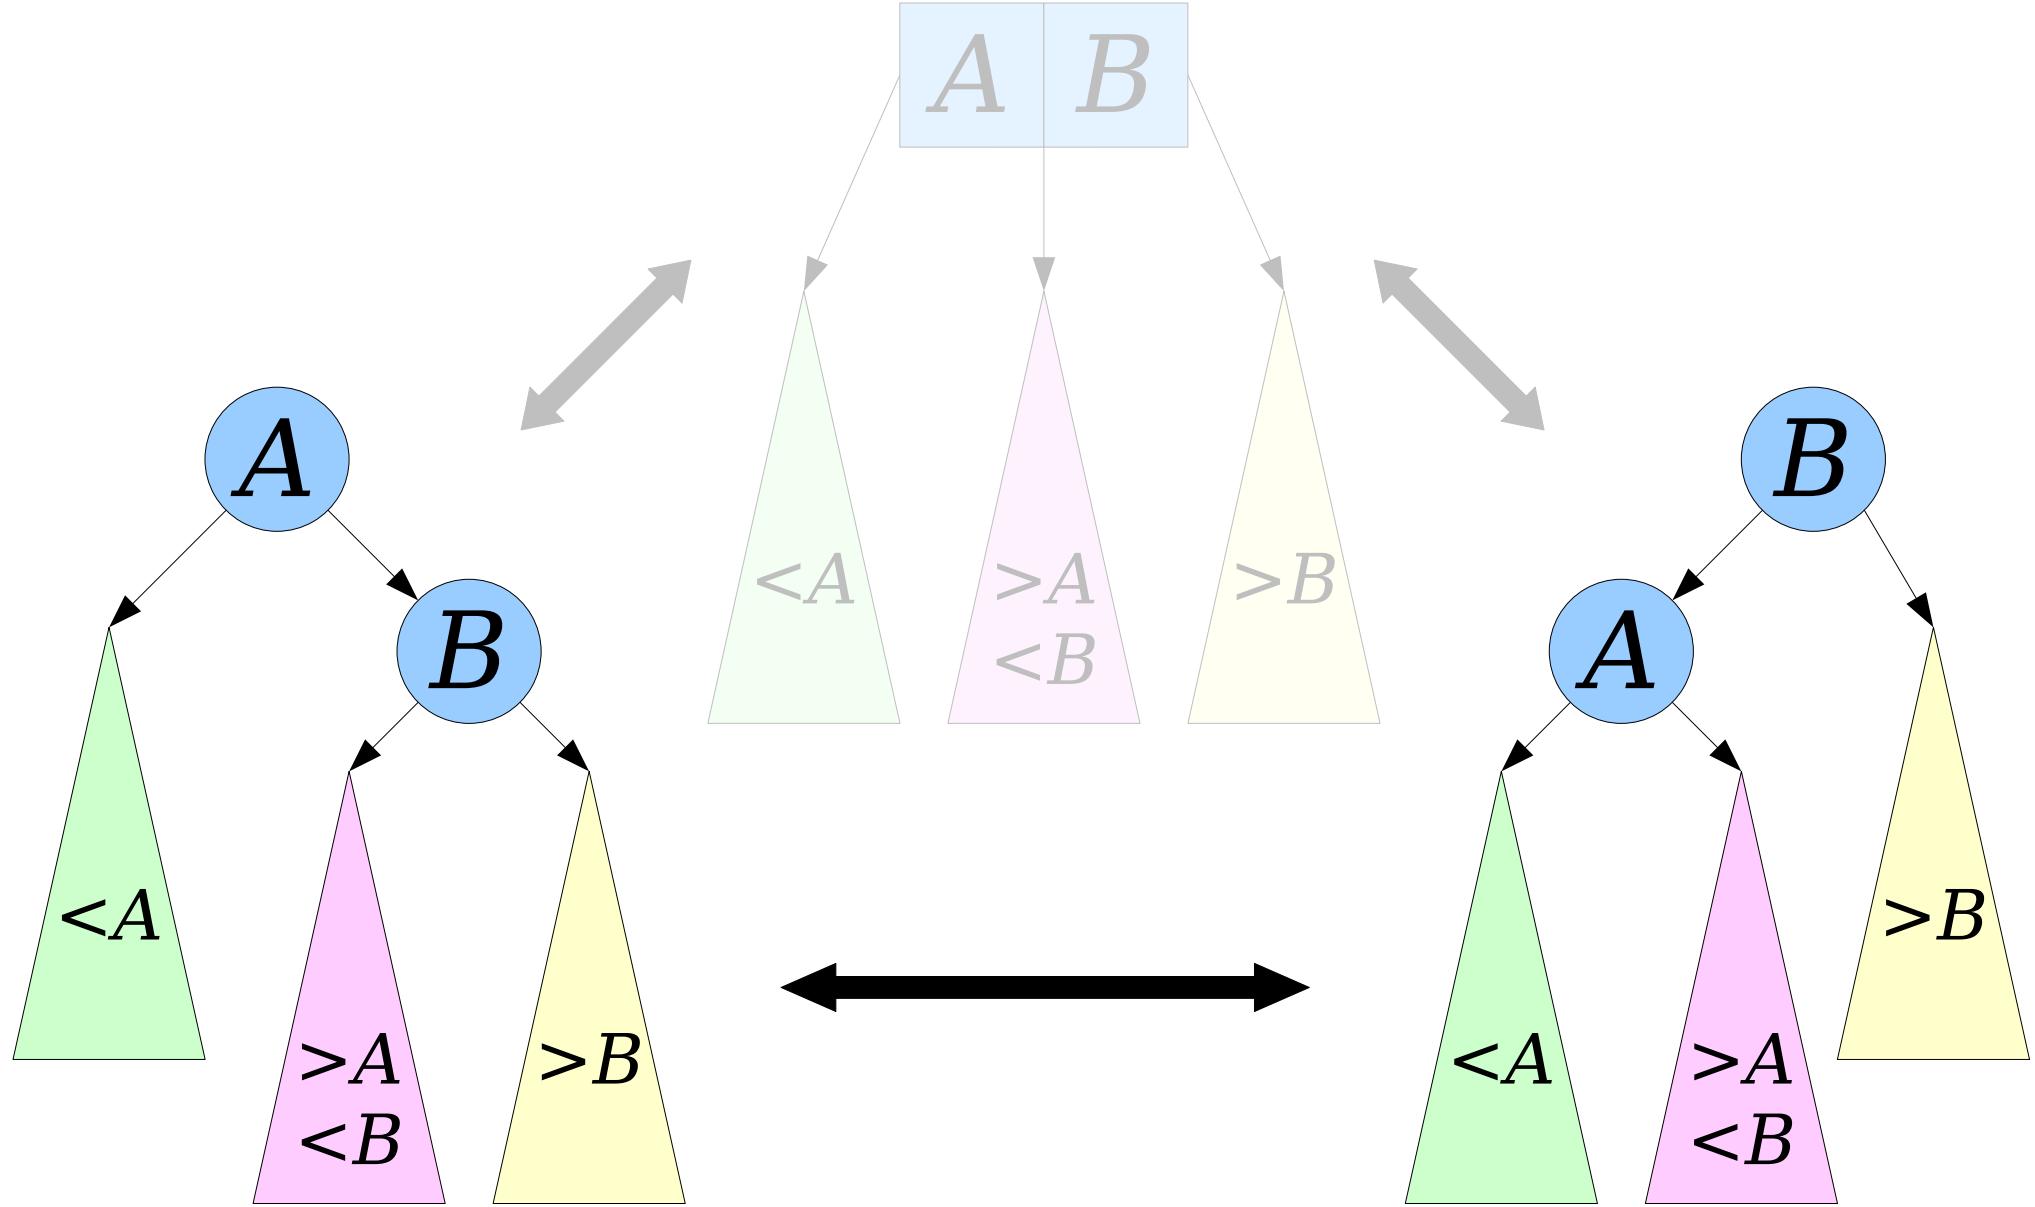


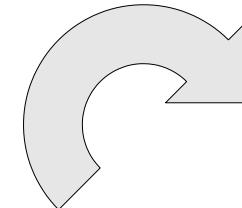
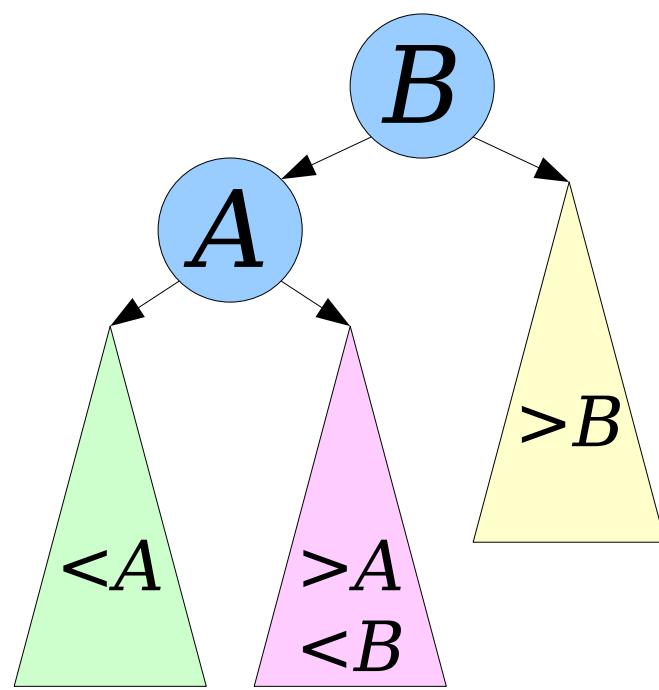
Goal

Tree Rotations

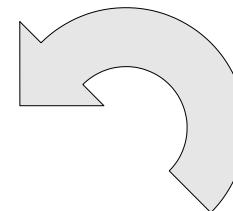
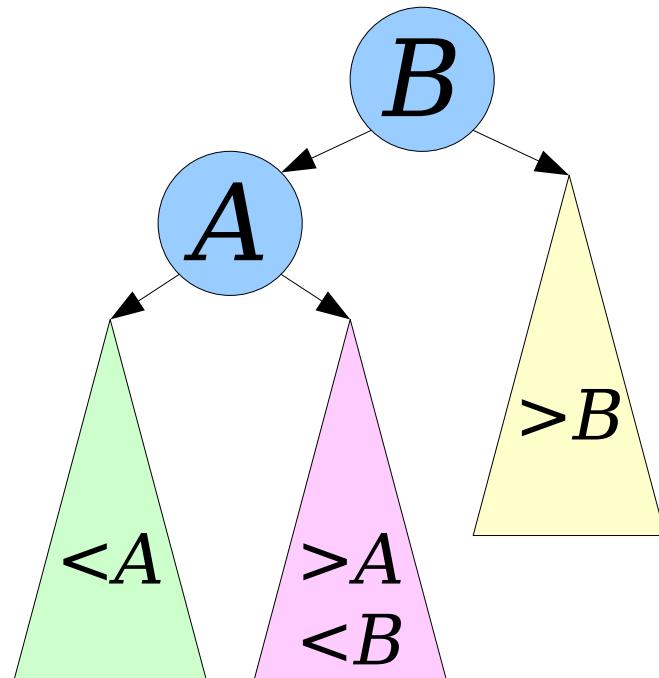
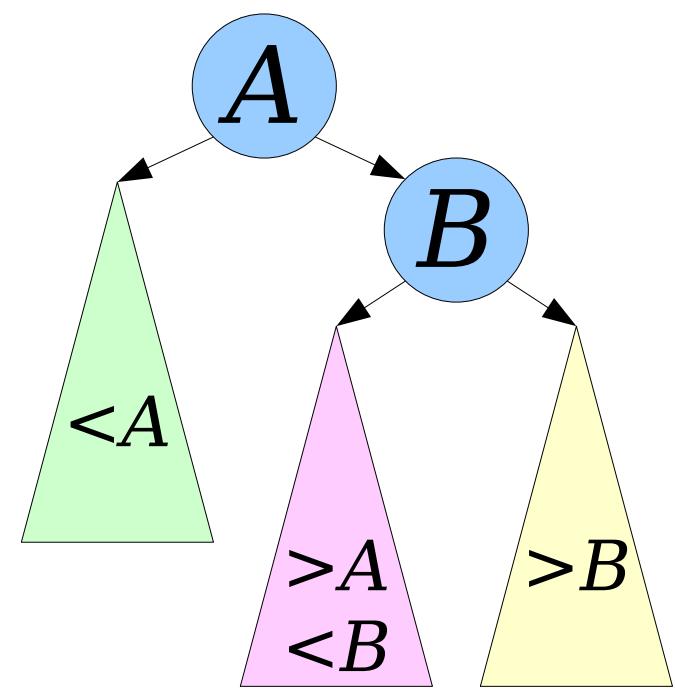


Tree Rotations

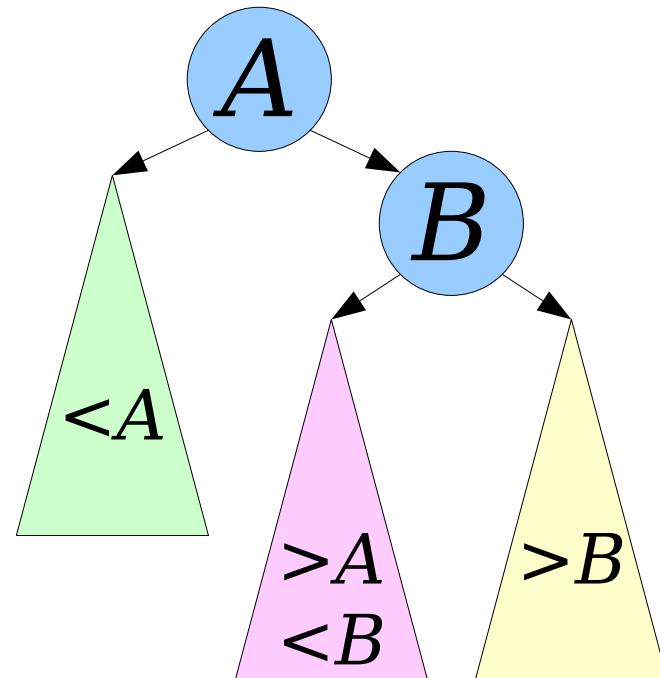


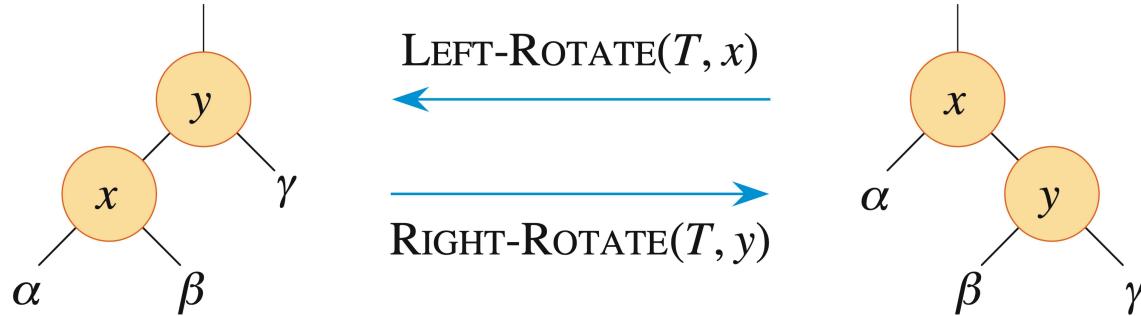


*Rotate
Right*



*Rotate
Left*





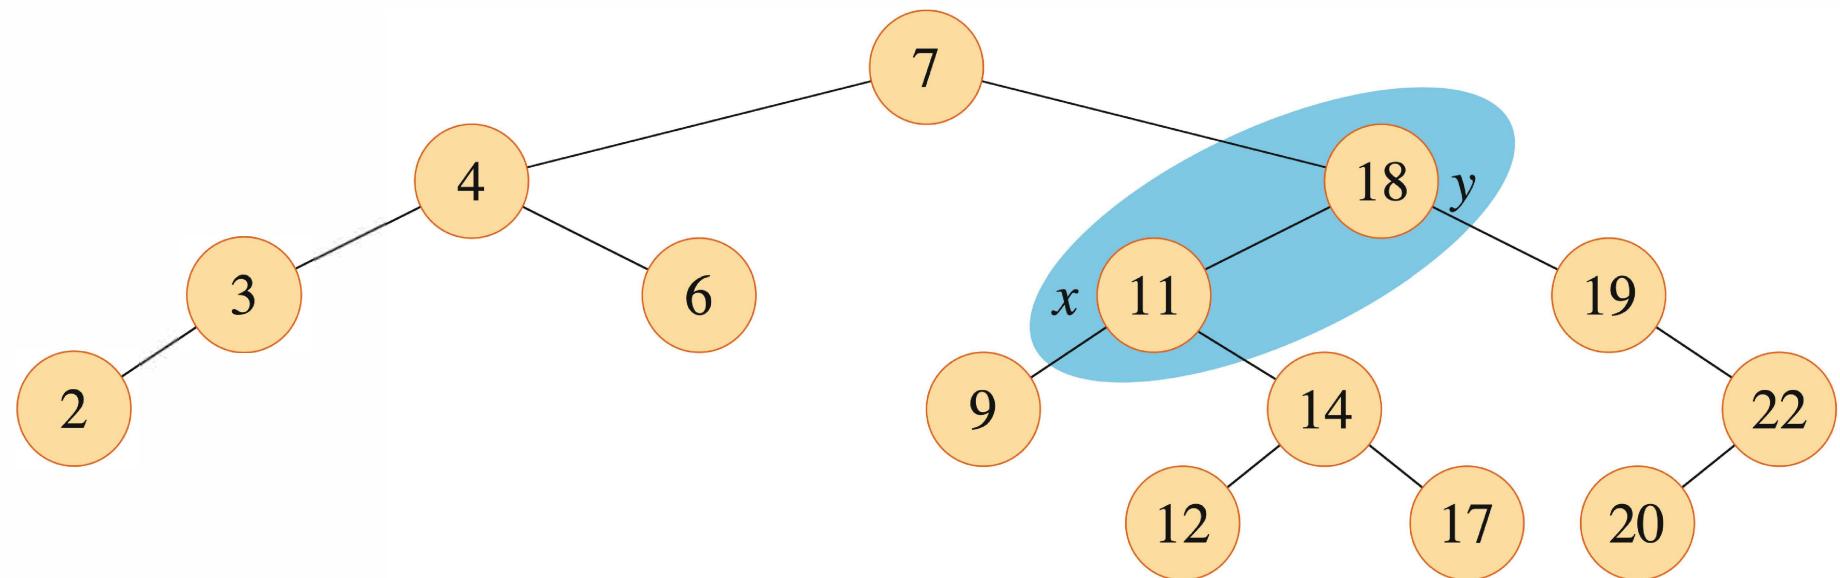
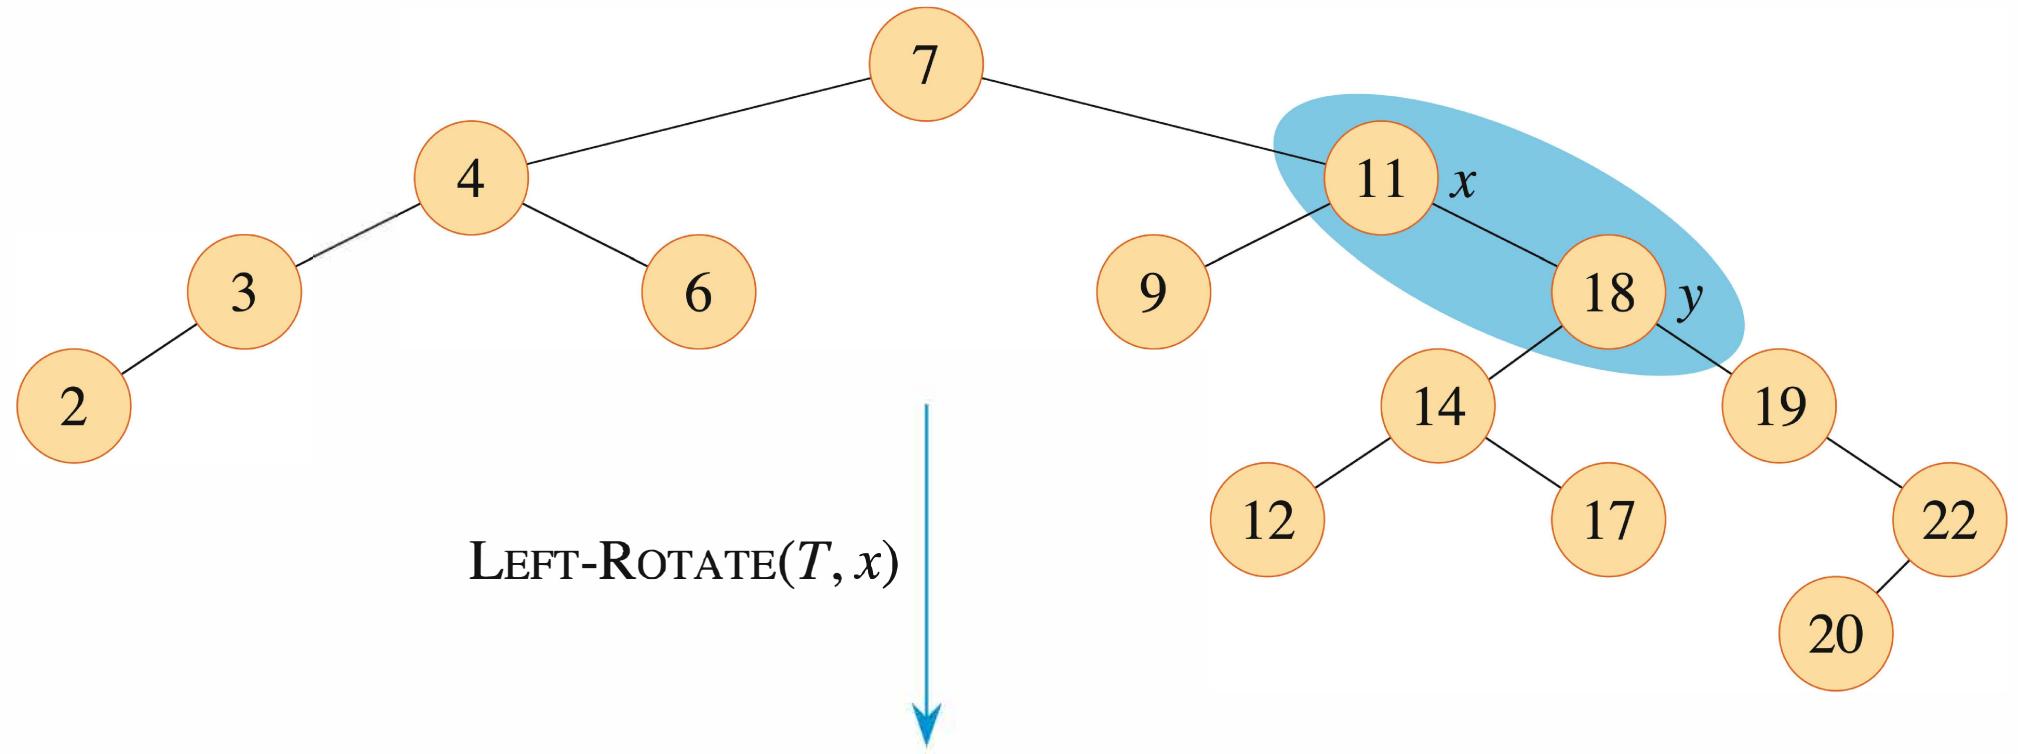
LEFT-ROTATE(T, x)

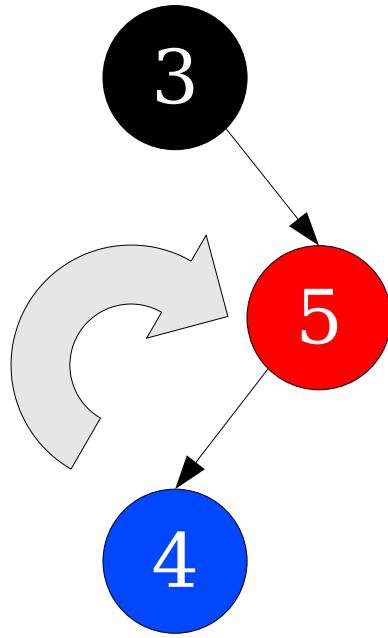
```

1   $y = x.right$                                 // turn  $y$ 's left subtree into  $x$ 's right subtree
2   $x.right = y.left$                            // if  $y$ 's left subtree is not empty ...
3  if  $y.left \neq T.nil$                       // ... then  $x$  becomes the parent of the subtree's root
4       $y.left.p = x$                             //  $x$ 's parent becomes  $y$ 's parent
5   $y.p = x.p$                                  // if  $x$  was the root ...
6  if  $x.p == T.nil$                           // ... then  $y$  becomes the root
7       $T.root = y$                             // otherwise, if  $x$  was a left child ...
8  elseif  $x == x.p.left$                      // ... then  $y$  becomes a left child
9       $x.p.left = y$                            // otherwise,  $x$  was a right child, and now  $y$  is
10     else  $x.p.right = y$                    // make  $x$  become  $y$ 's left child
11      $y.left = x$ 
12      $x.p = y$ 

```

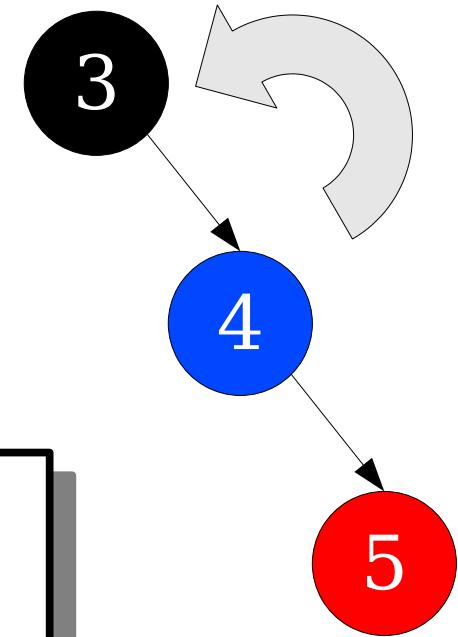
// turn y 's left subtree into x 's right subtree
 // if y 's left subtree is not empty ...
 // ... then x becomes the parent of the subtree's root
 // x 's parent becomes y 's parent
 // if x was the root ...
 // ... then y becomes the root
 // otherwise, if x was a left child ...
 // ... then y becomes a left child
 // otherwise, x was a right child, and now y is
 // make x become y 's left child





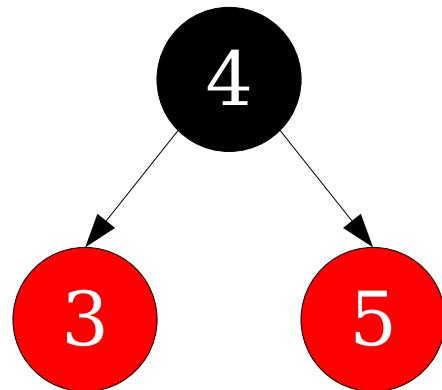
apply
rotation

A large blue arrow pointing to the right, positioned between the initial tree and the first step of the rotation.



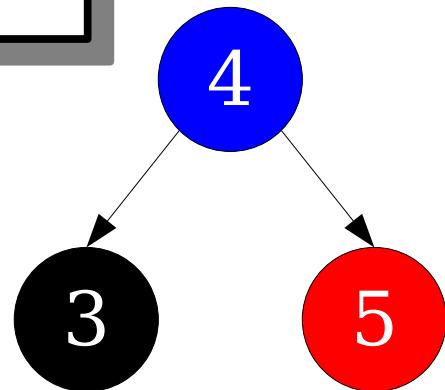
This applies any time we're inserting a new node into the middle of a “3-node” in this pattern.

By making observations like these, we can determine how to update a red/black tree after an insertion.



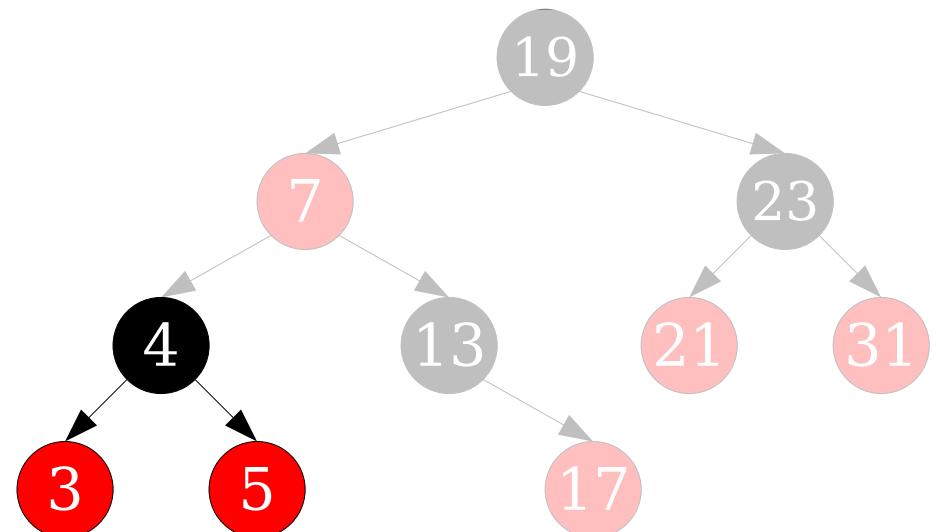
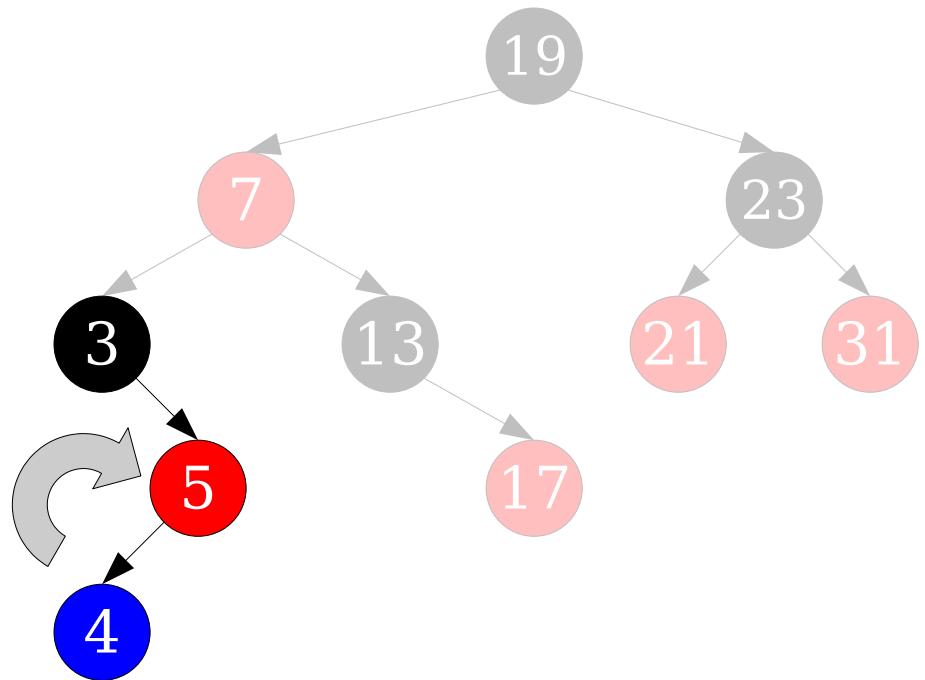
change
colors

A large blue arrow pointing to the left, positioned between the first rotated tree and the final colored tree.

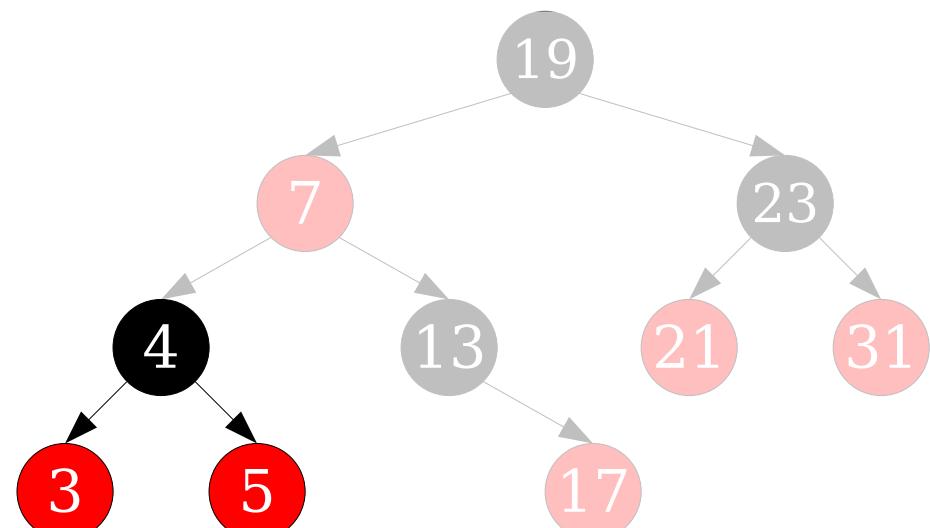
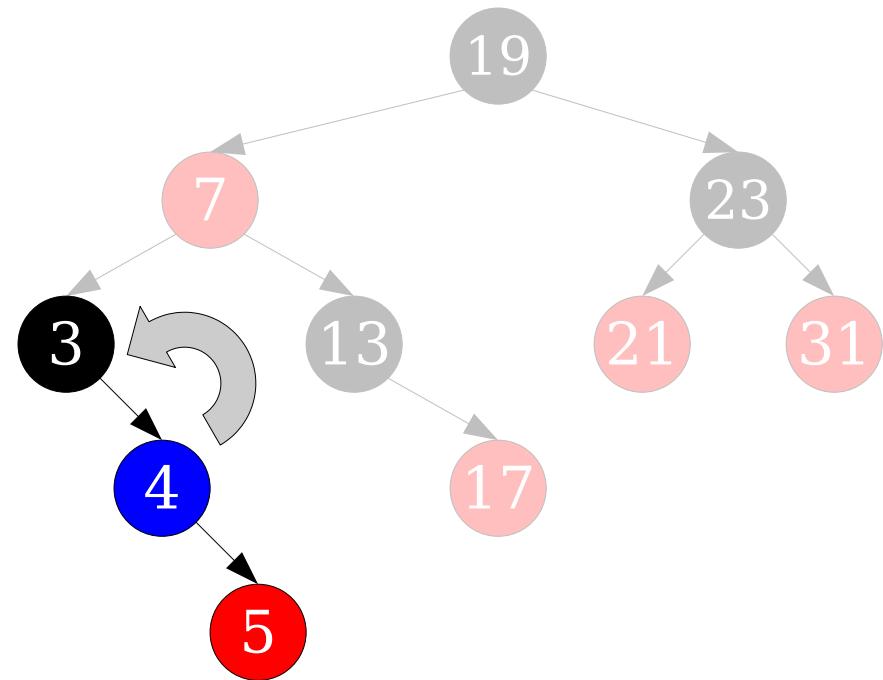


apply
rotation

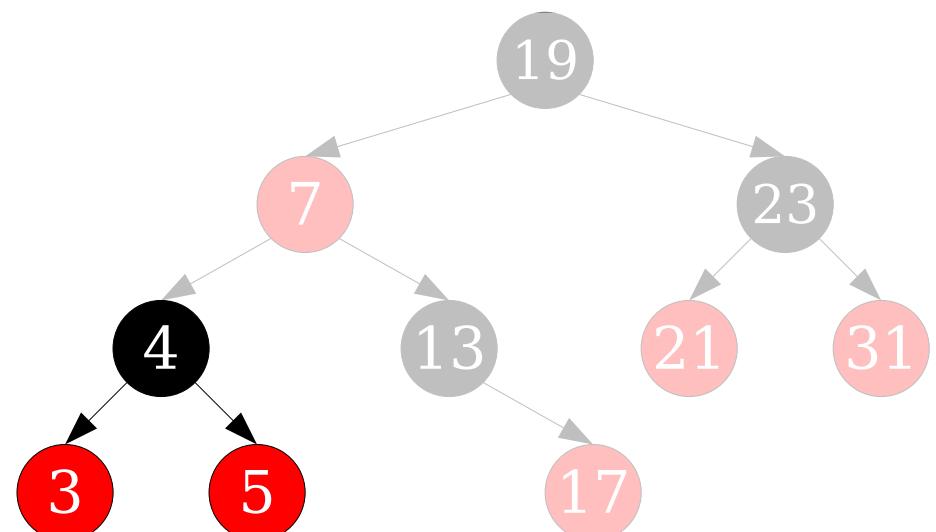
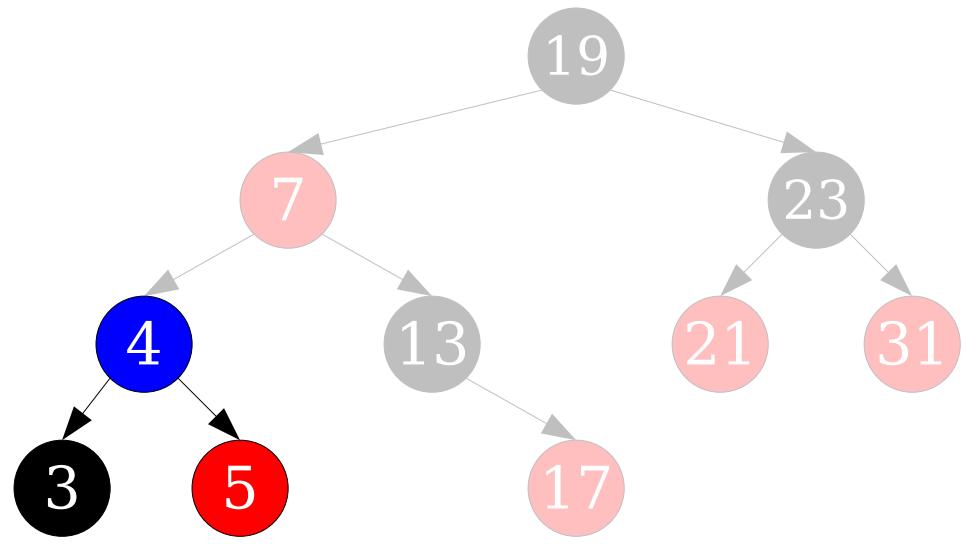
A large blue arrow pointing downwards, positioned between the final colored tree and the final state of the tree.



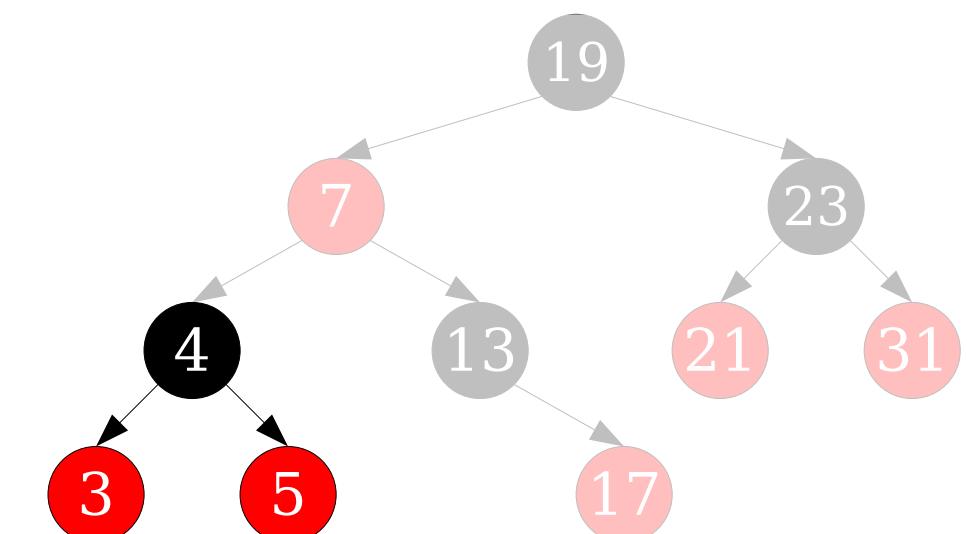
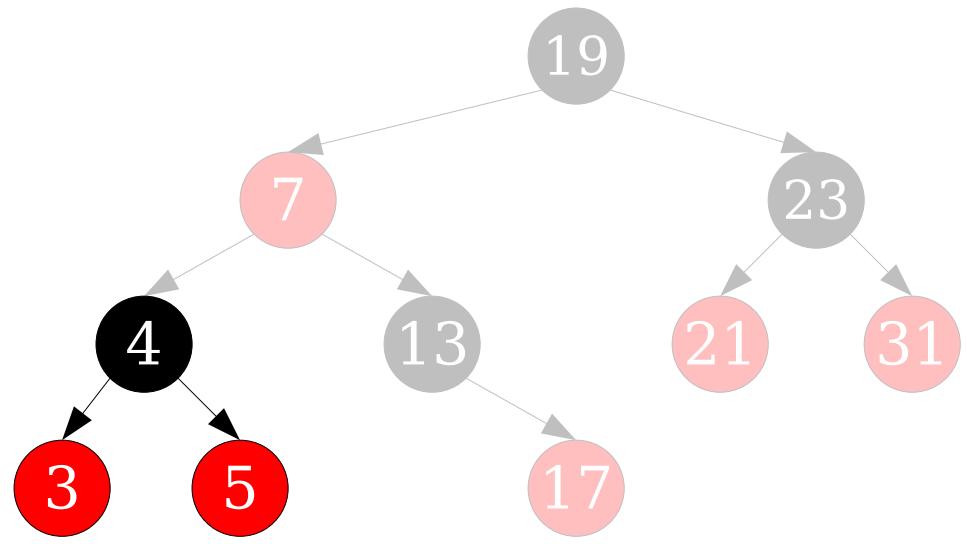
Goal



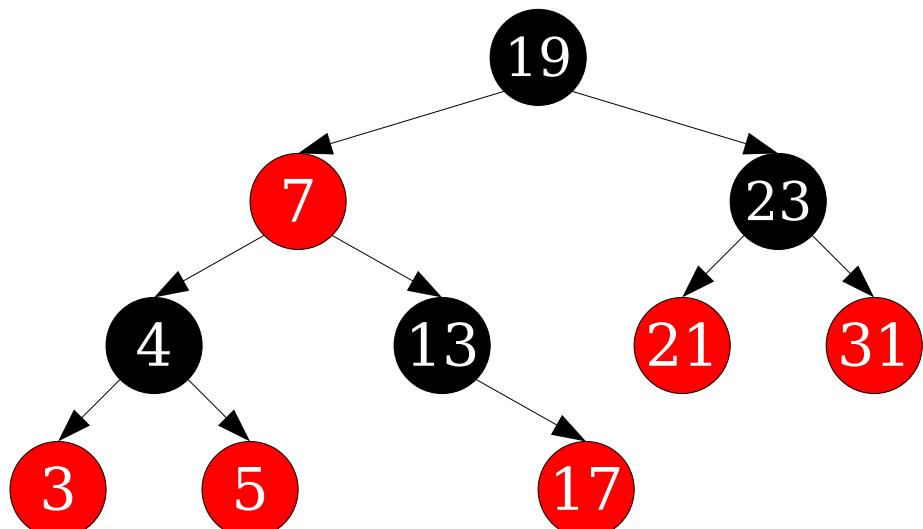
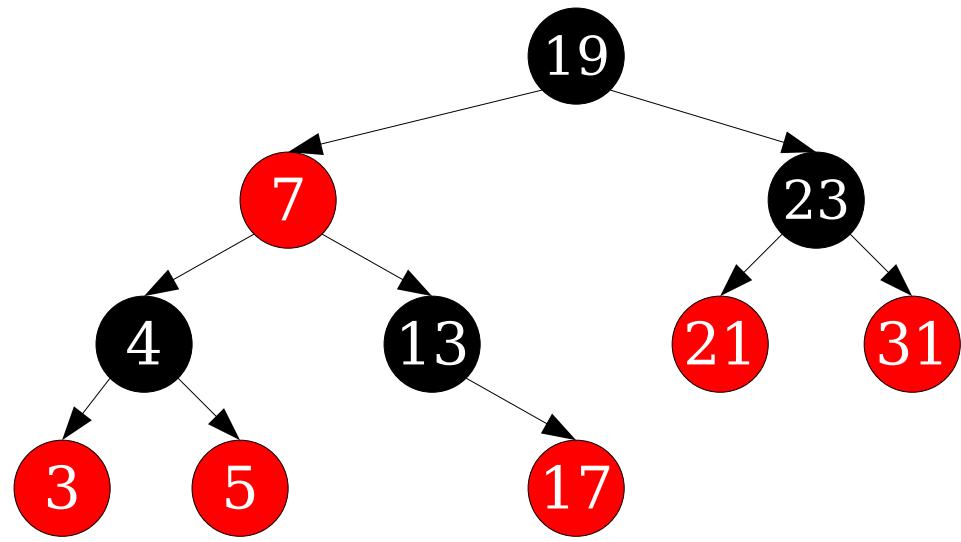
Goal



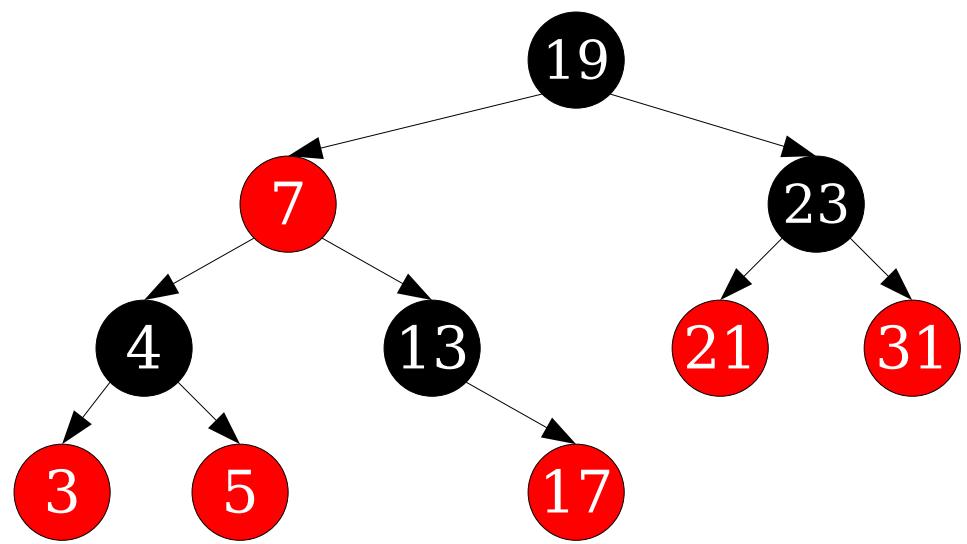
Goal

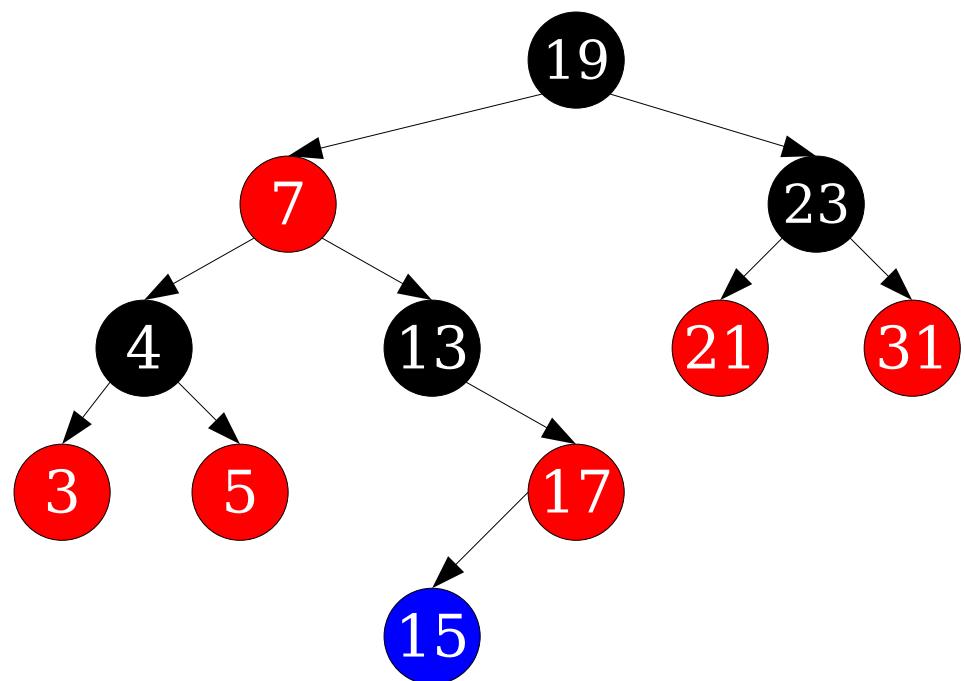


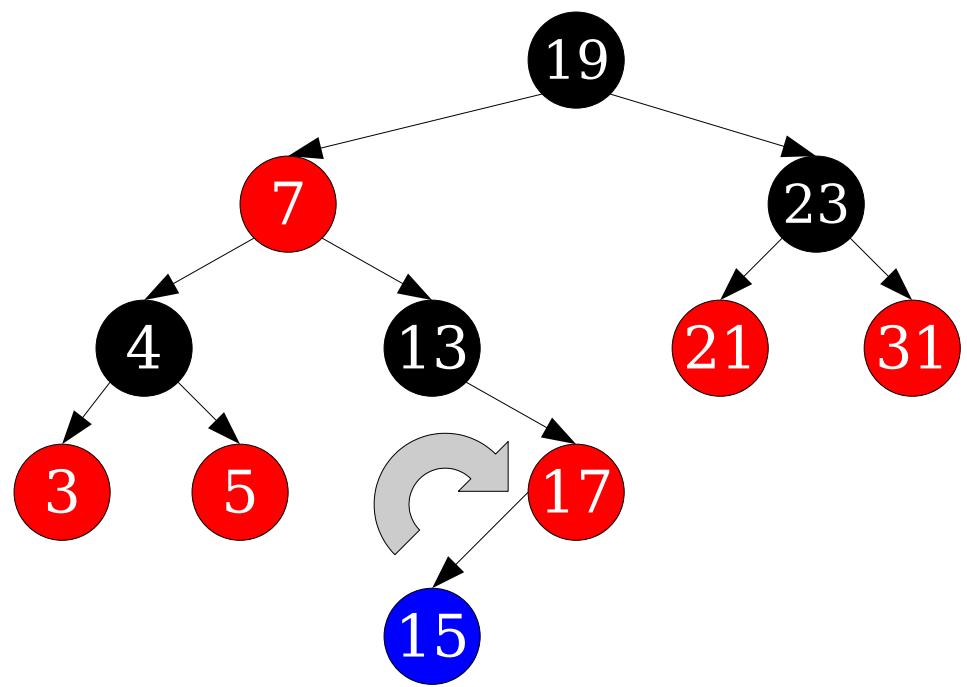
Goal

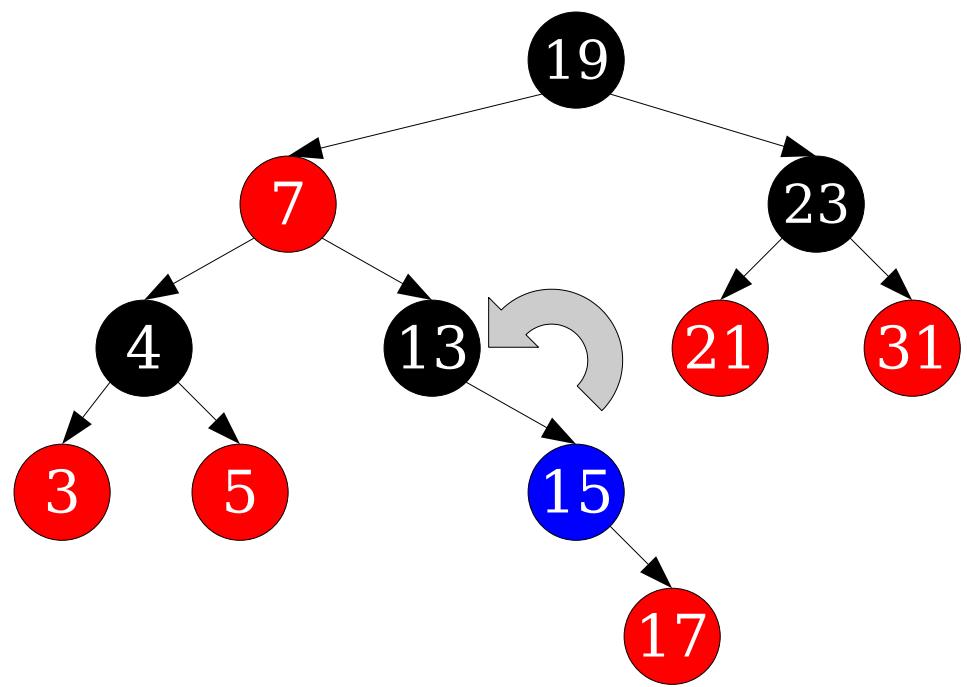


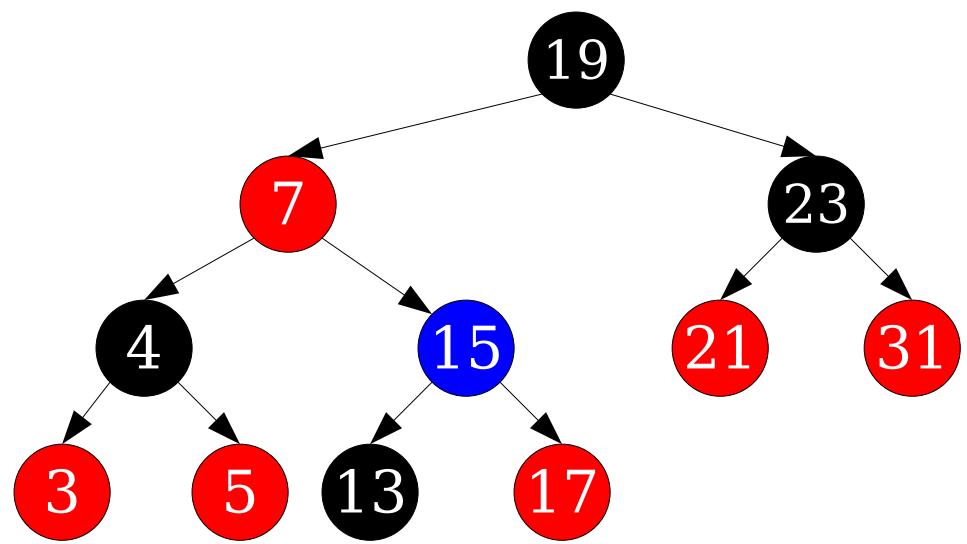
Goal

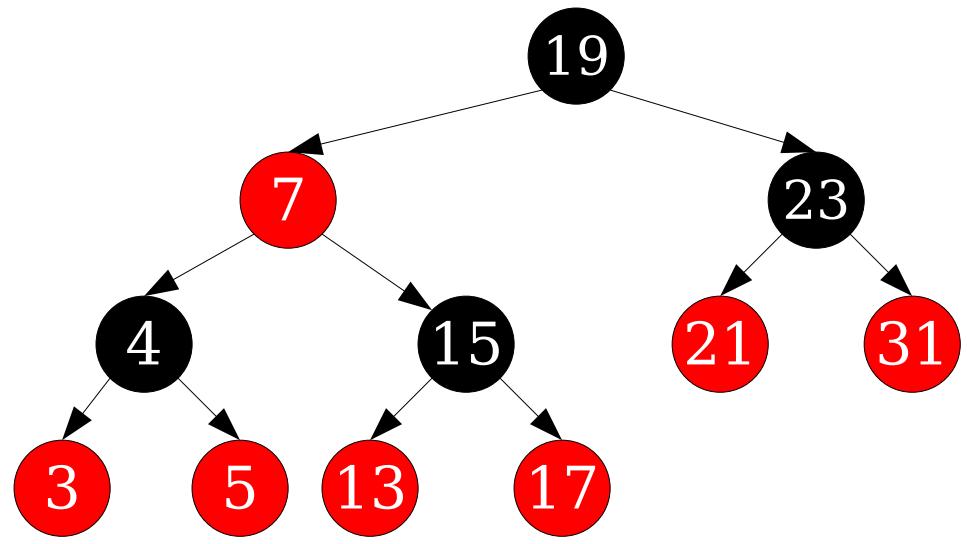


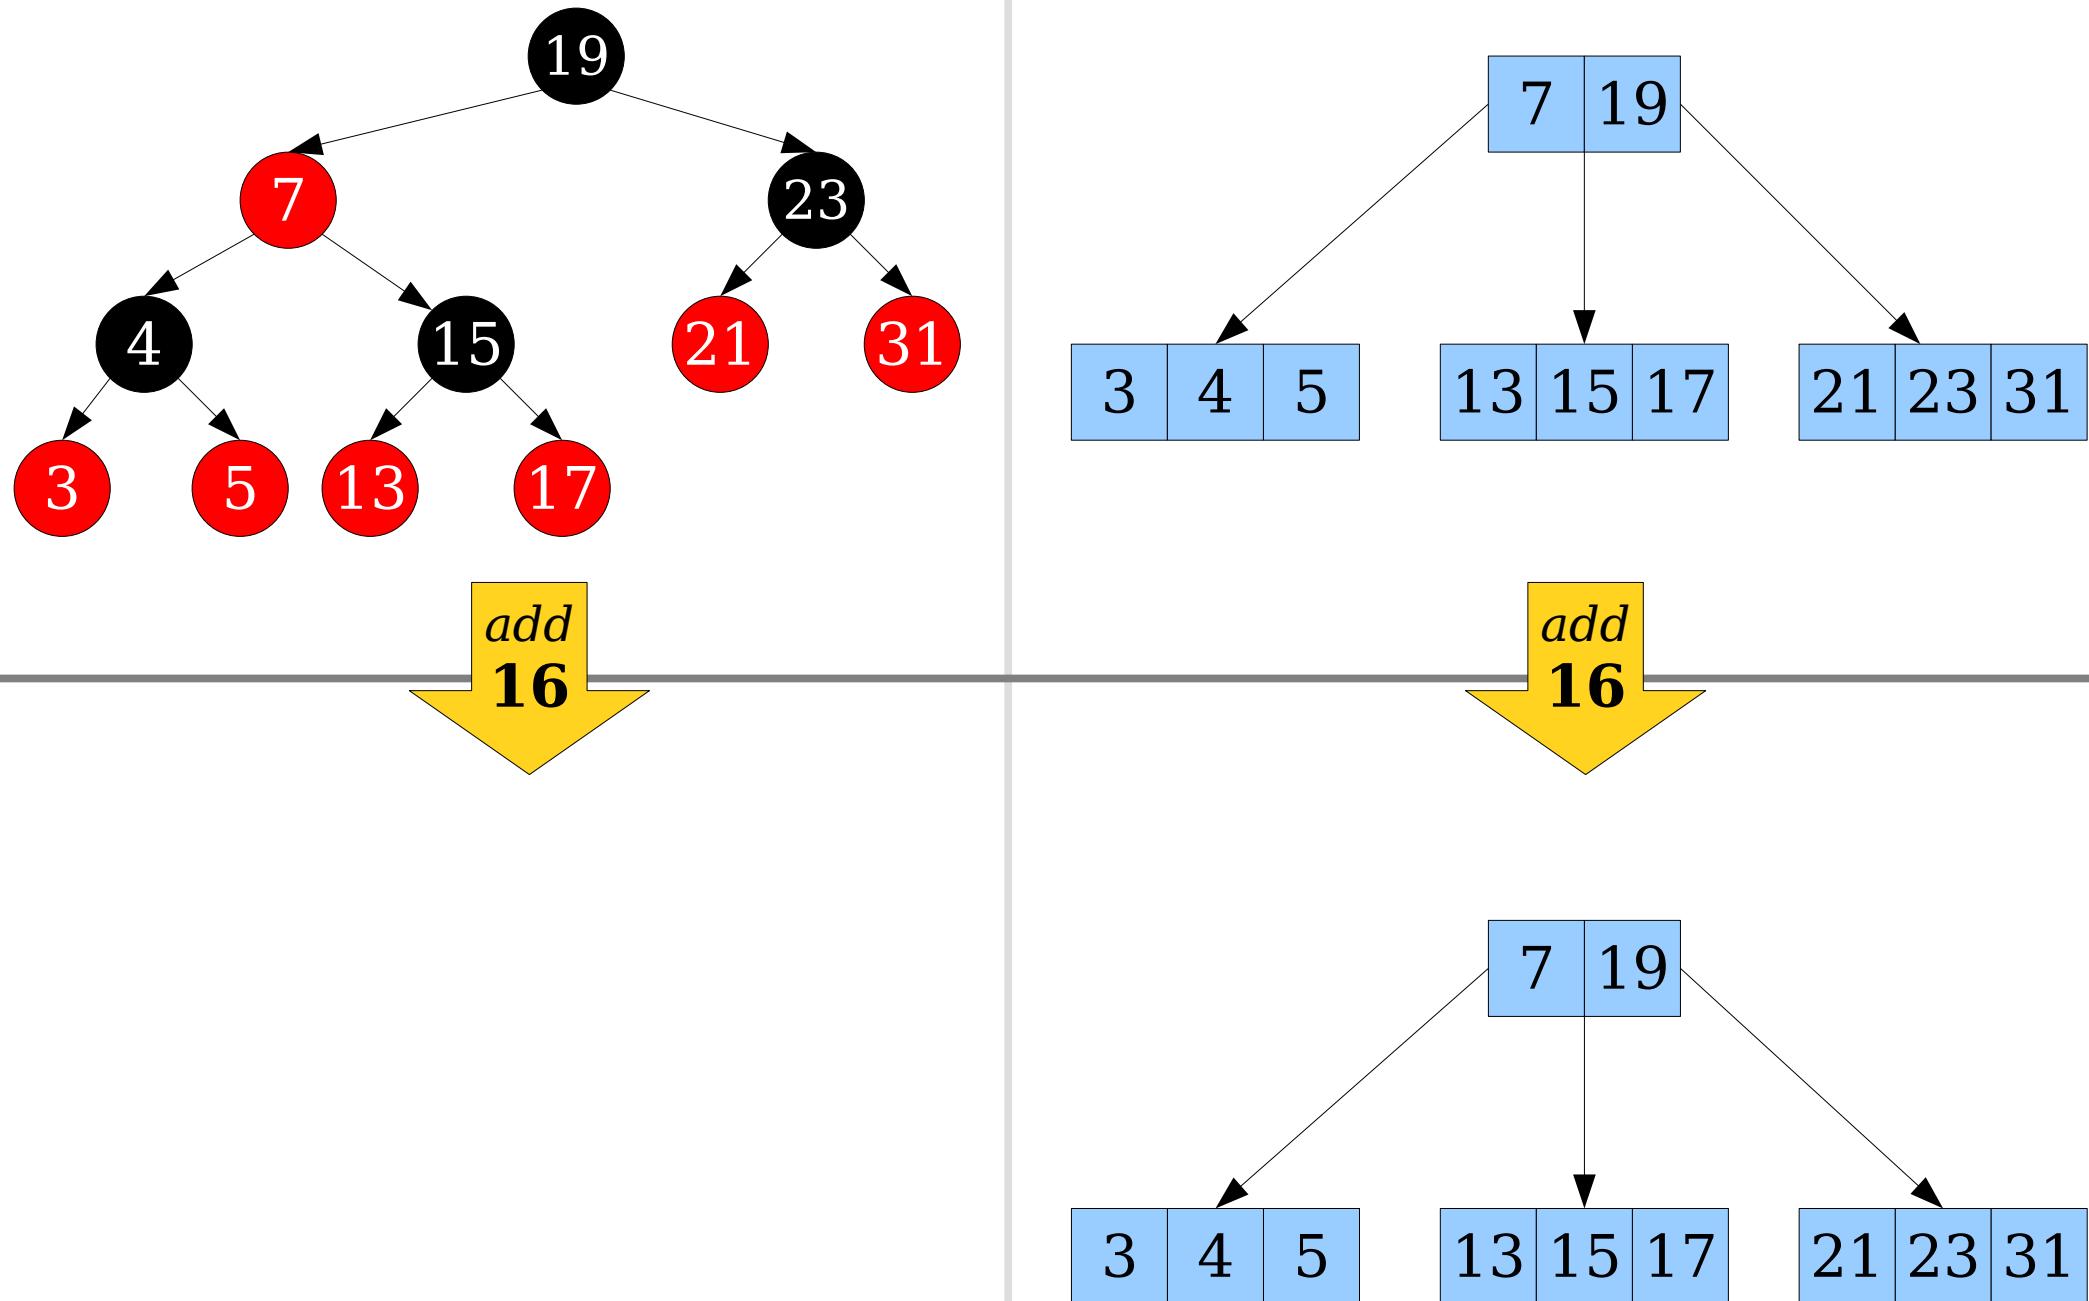


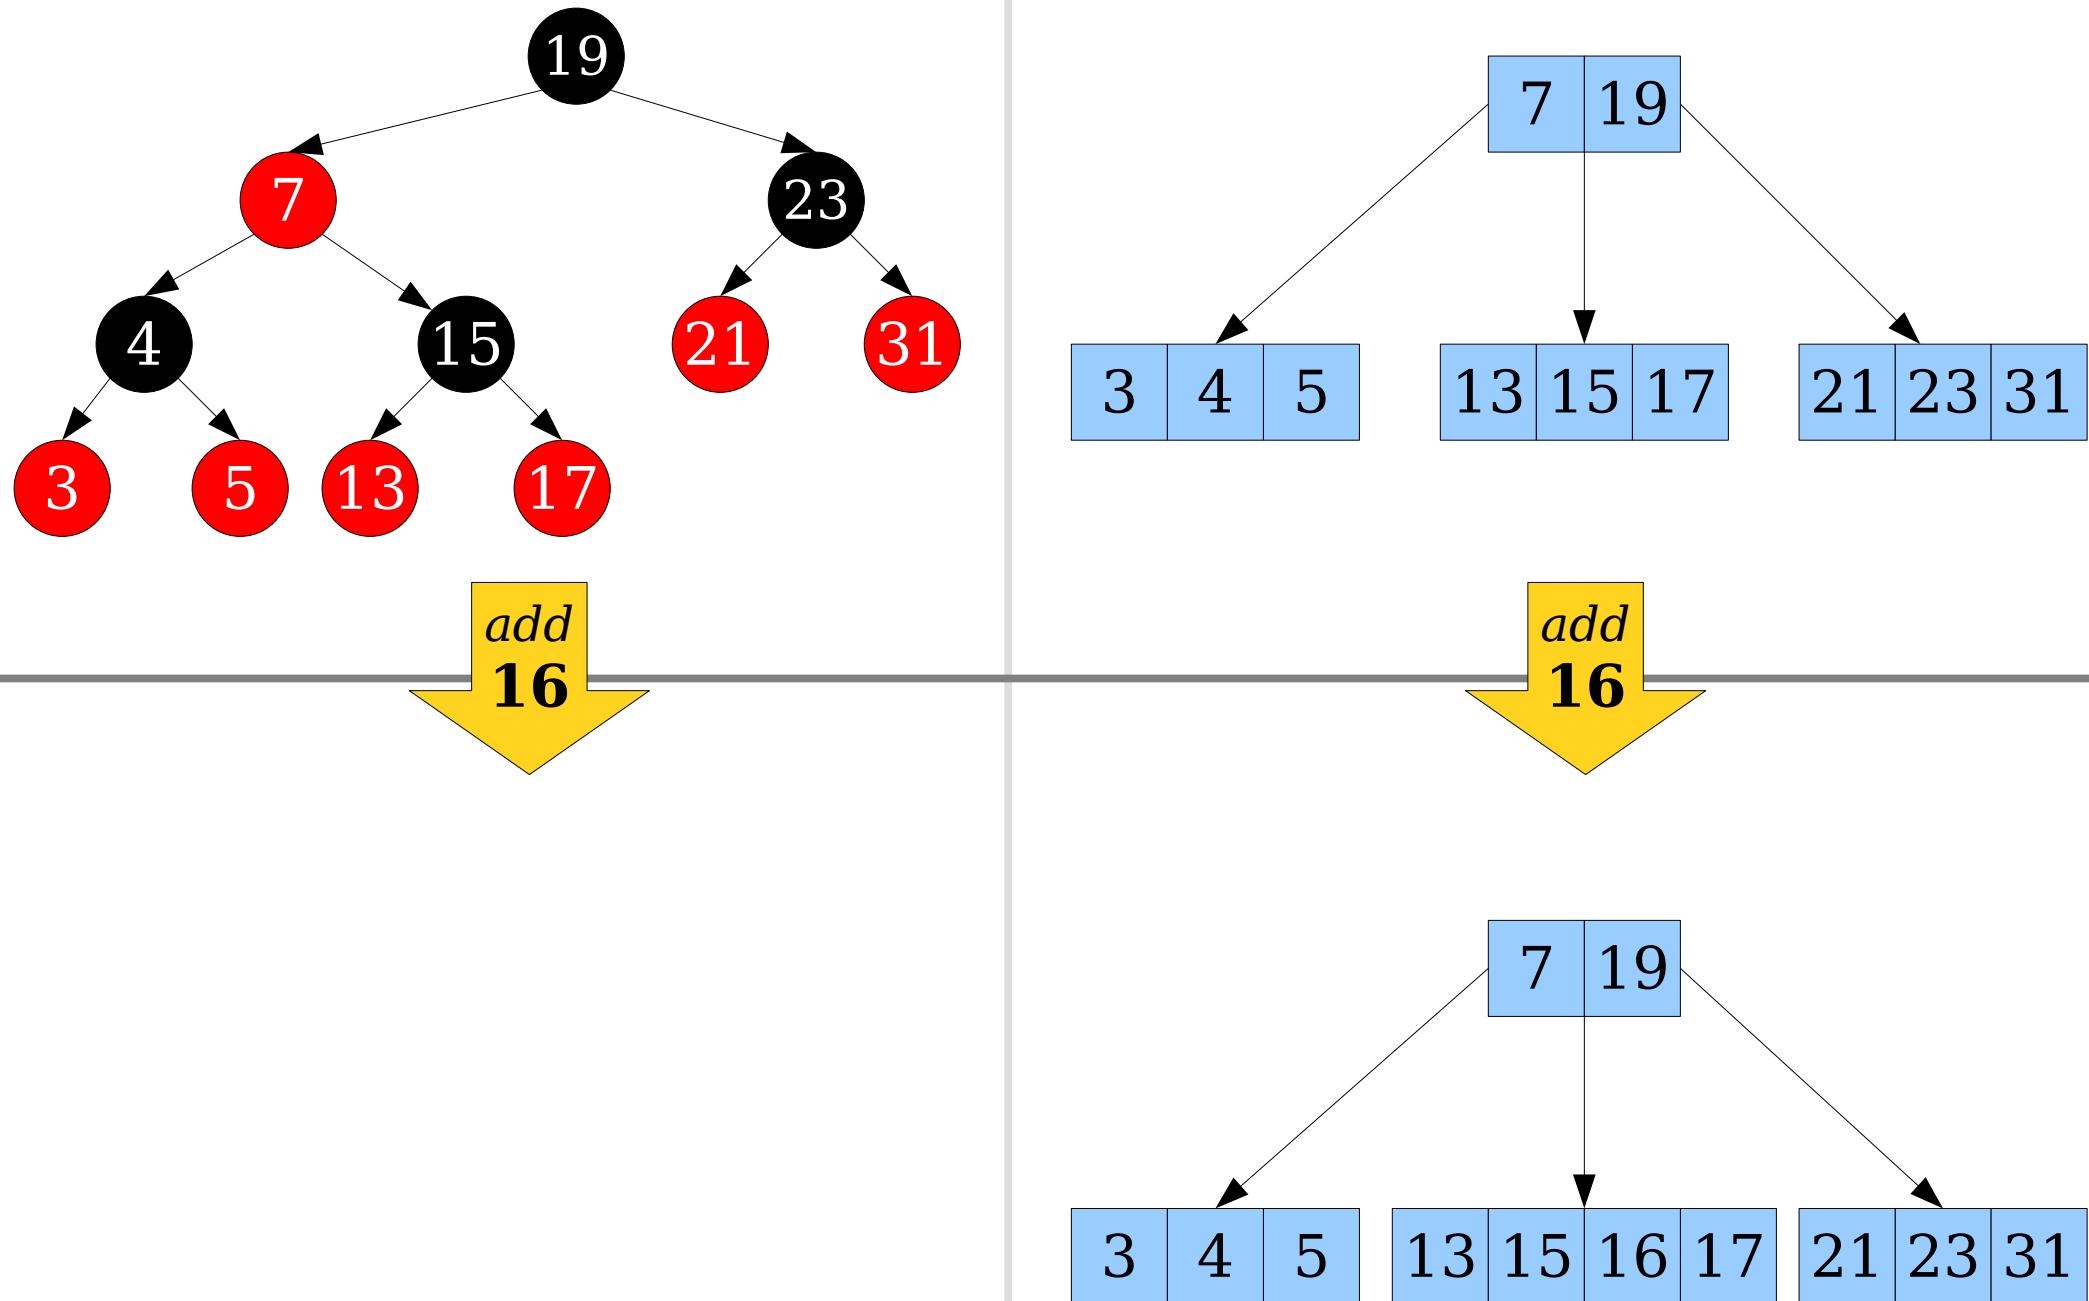


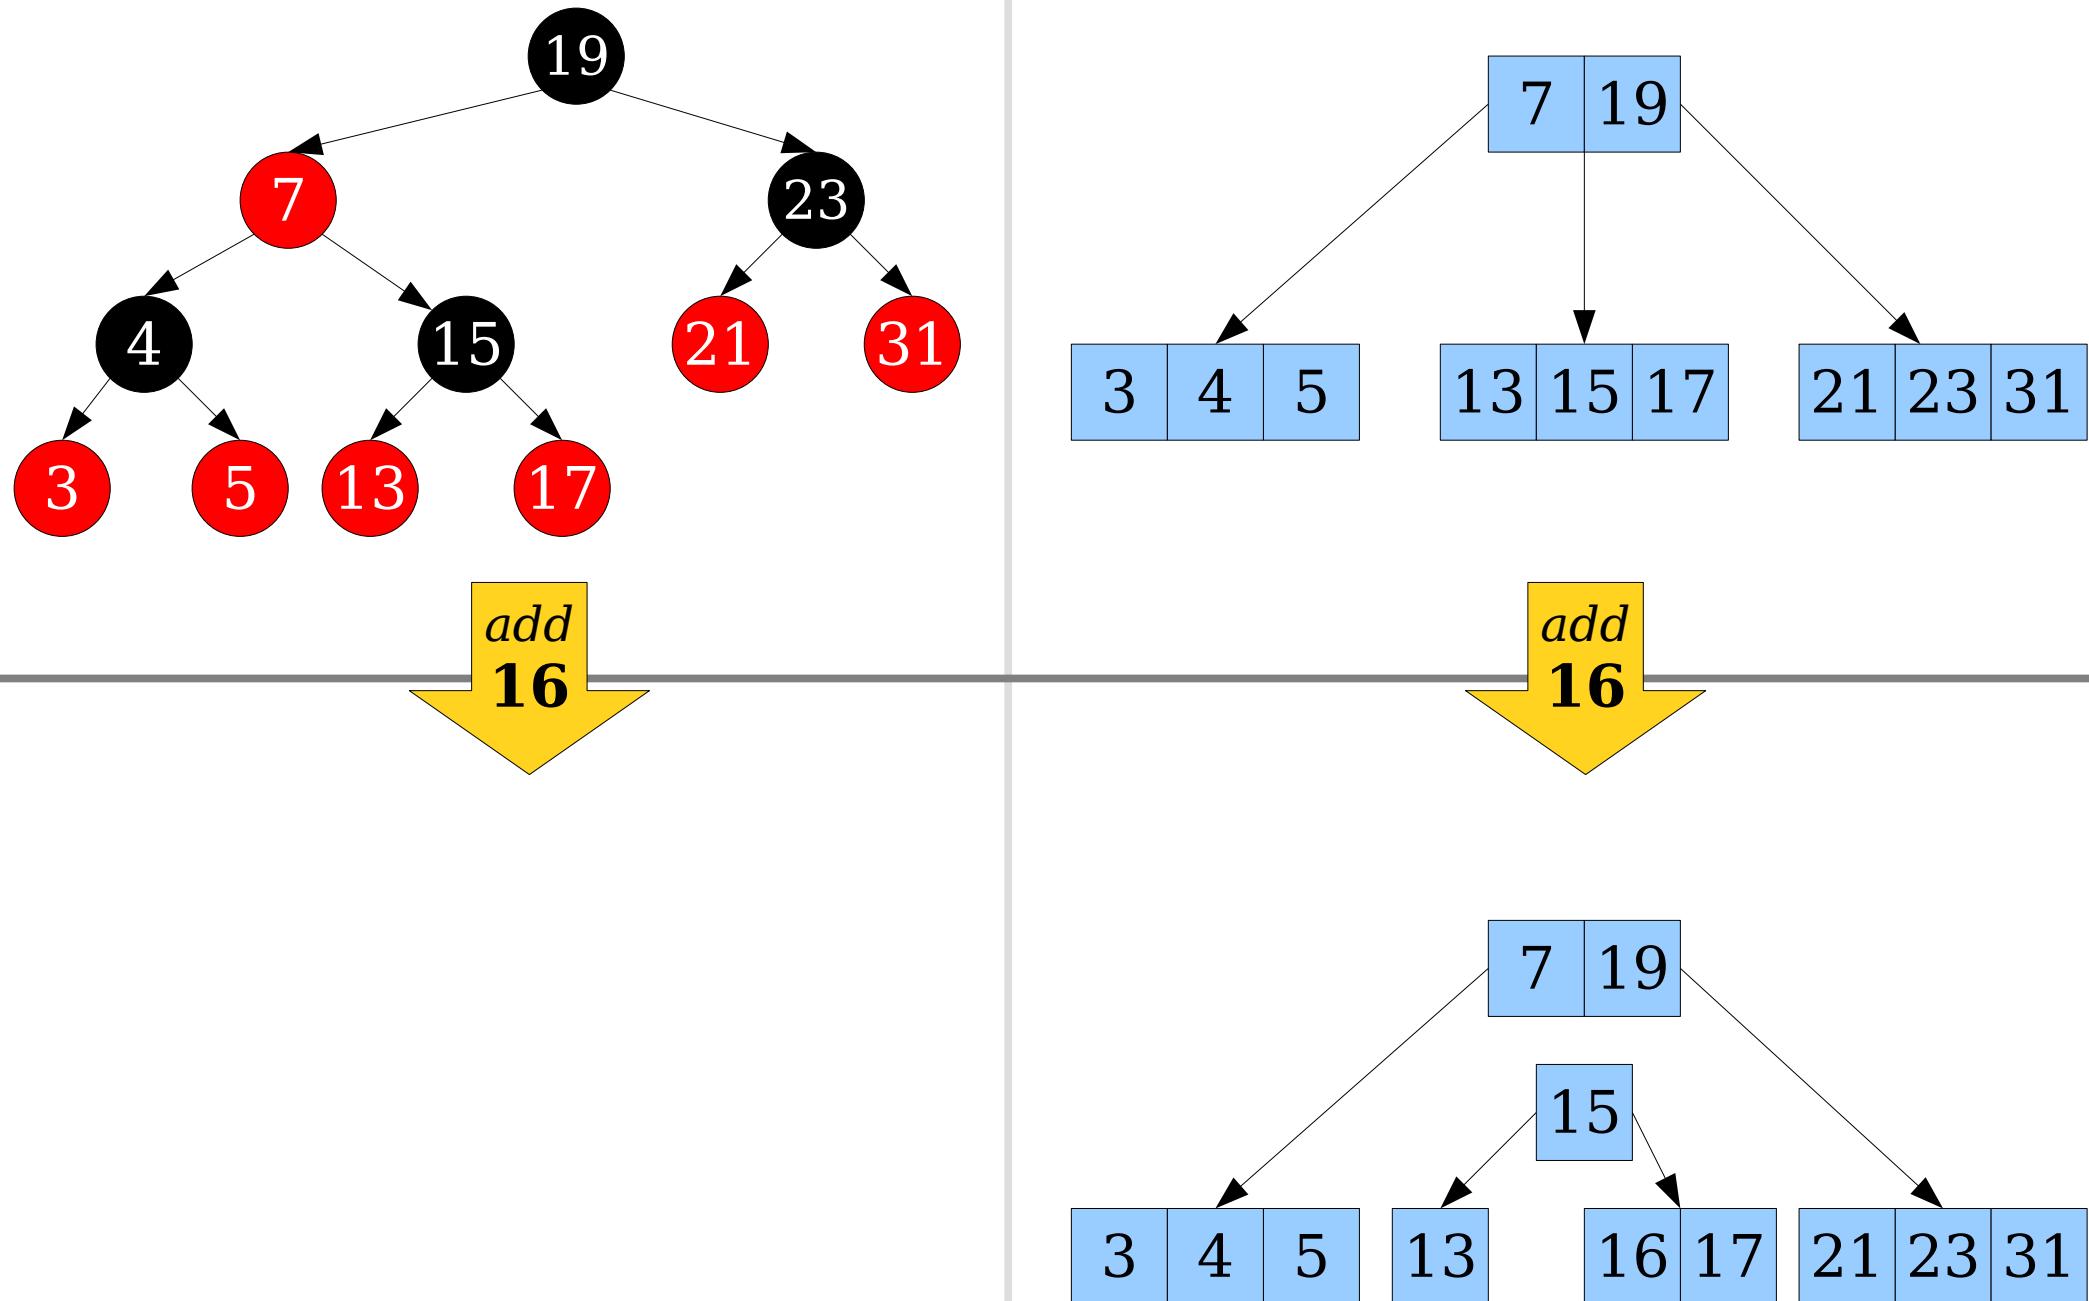


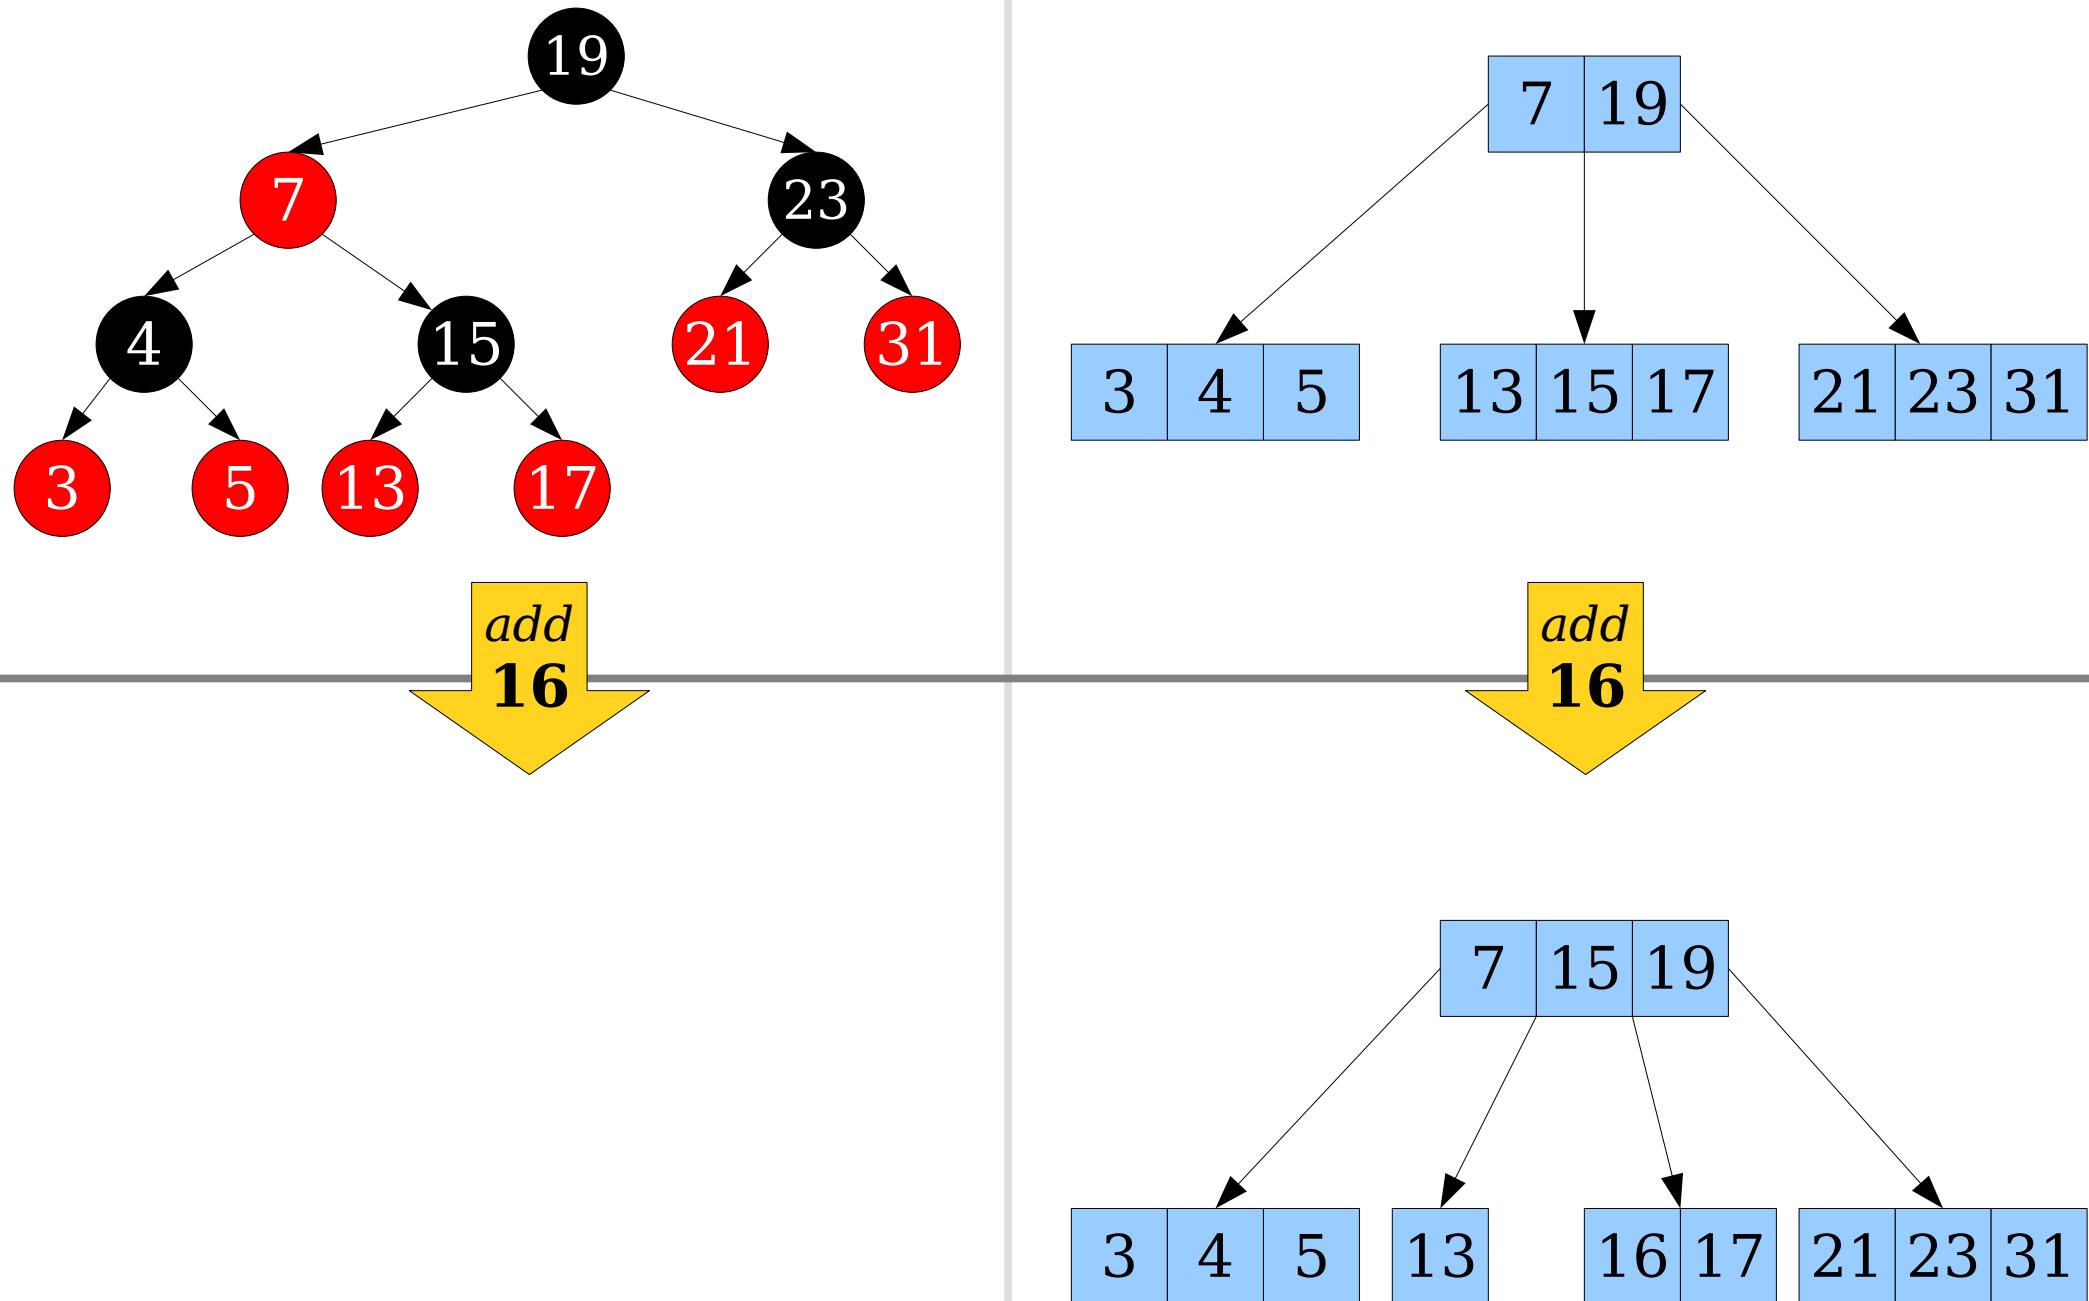


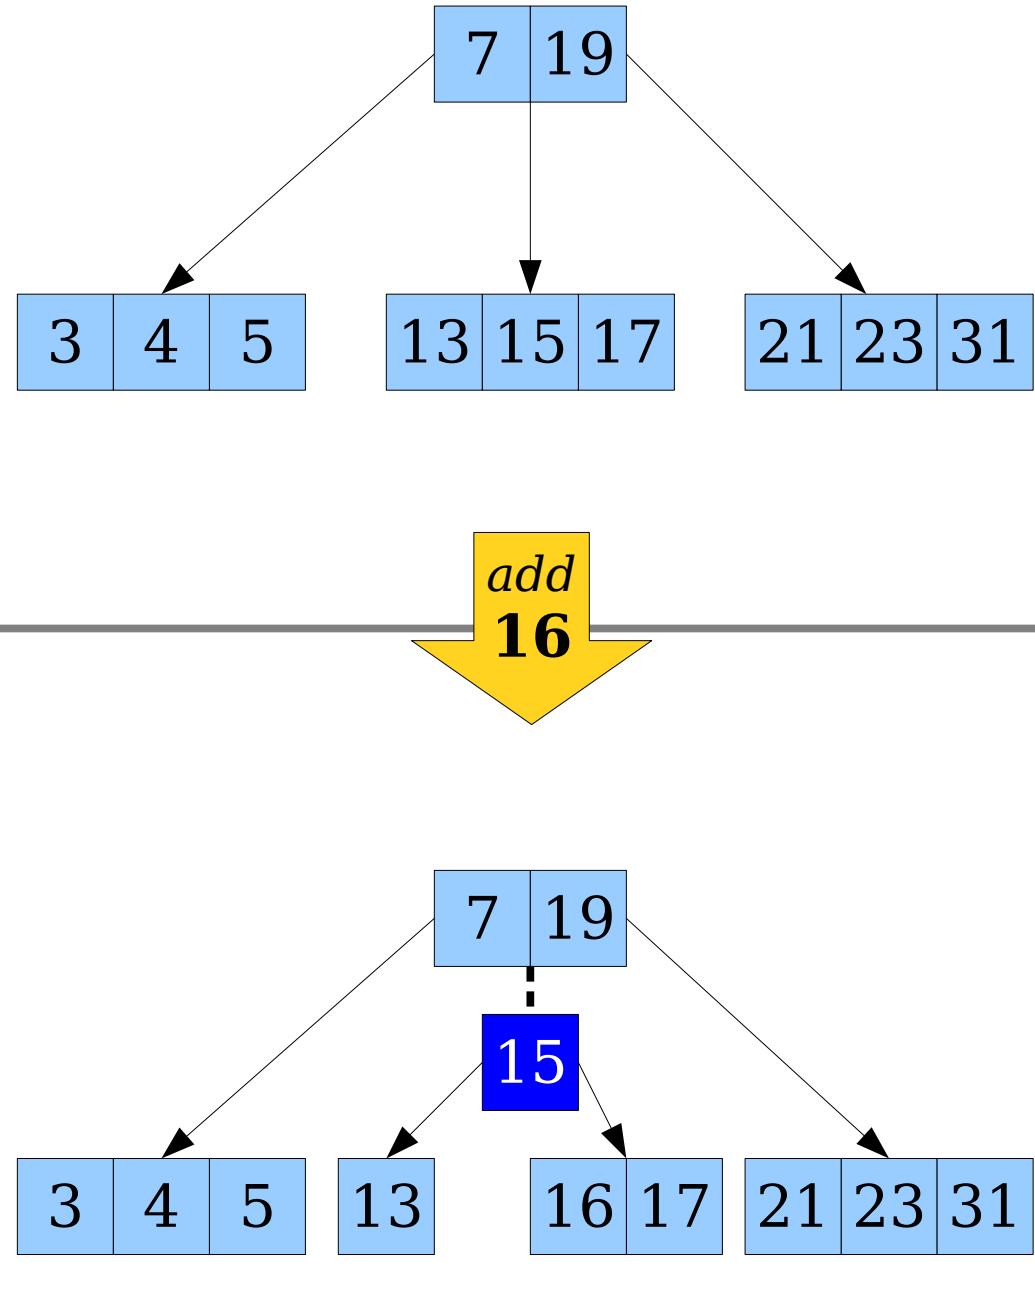
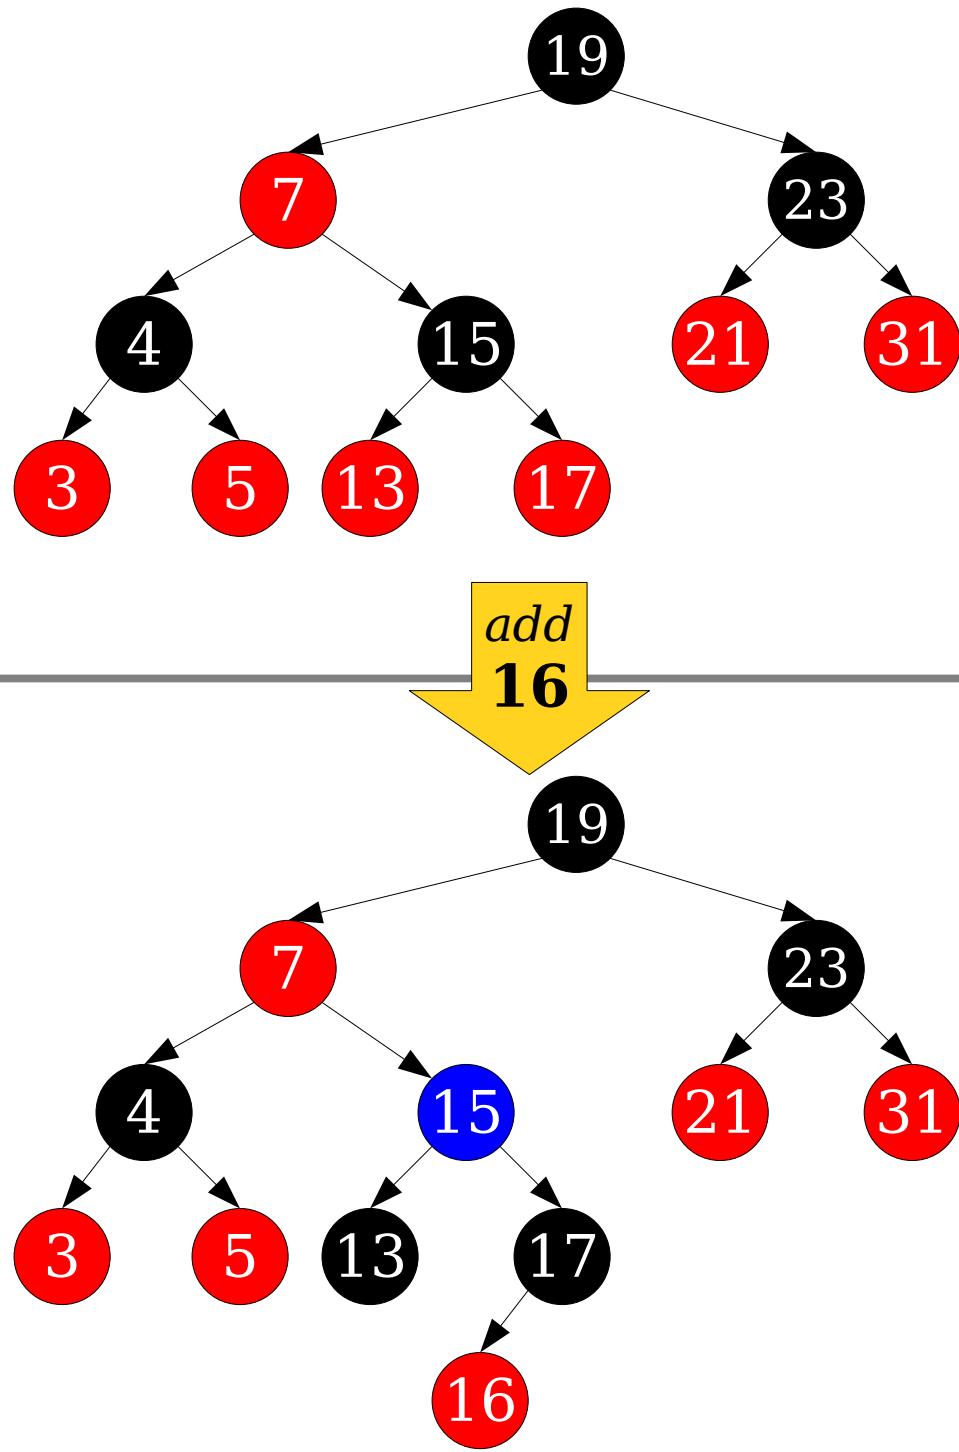


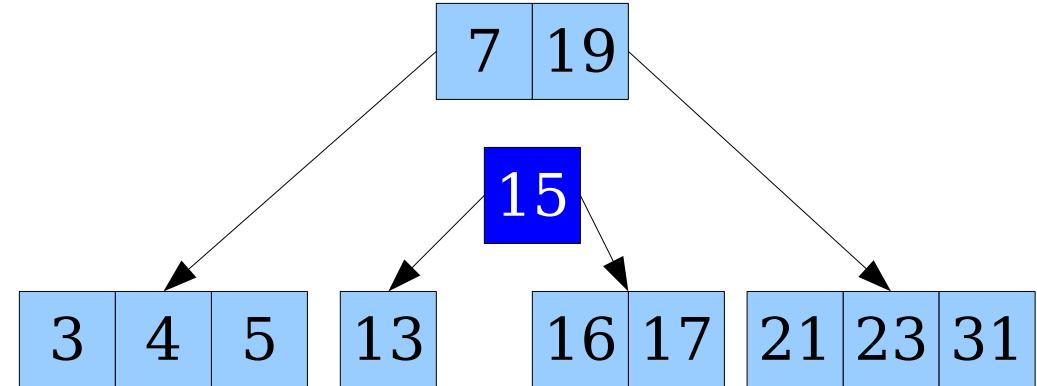
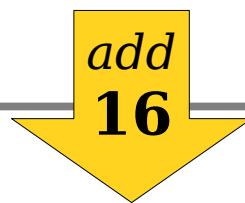
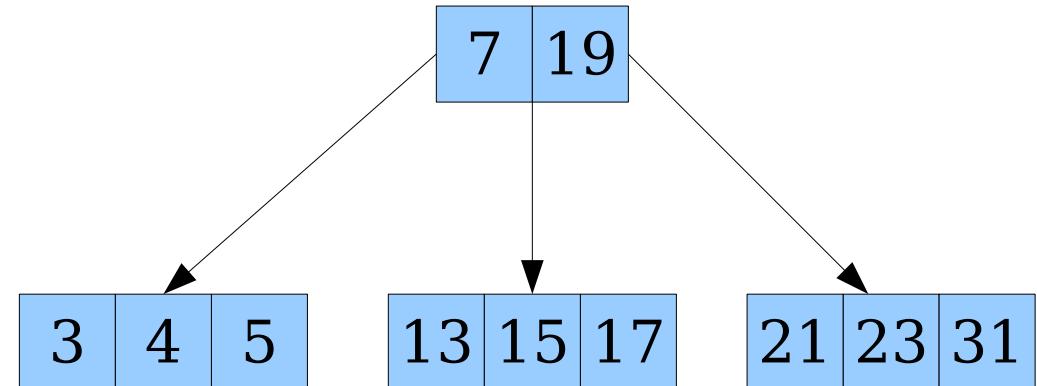
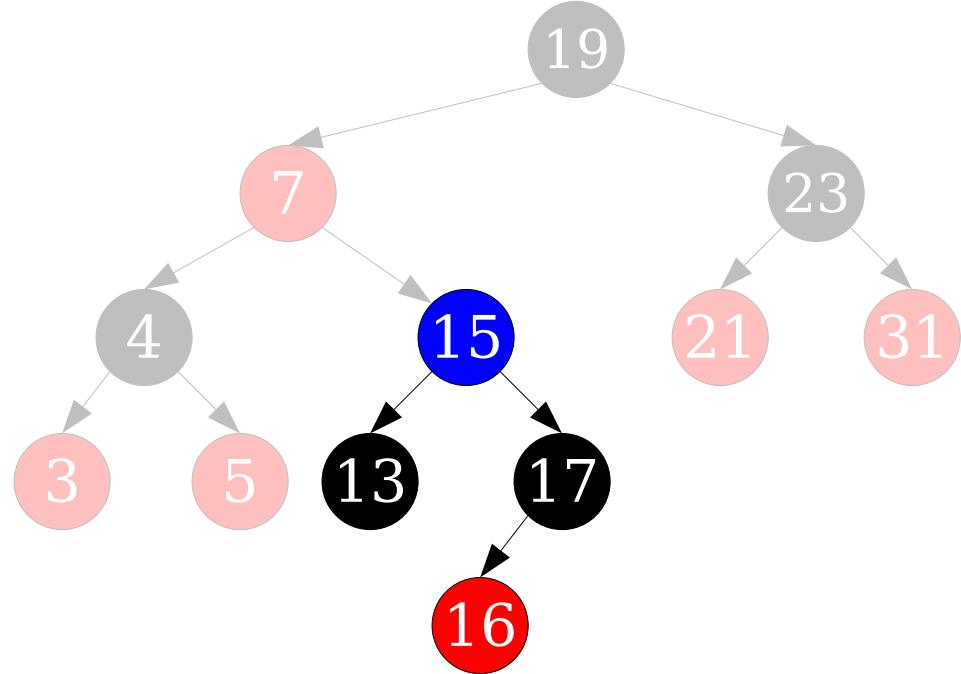
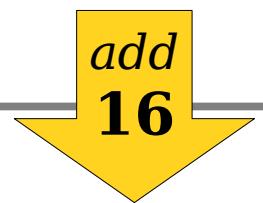
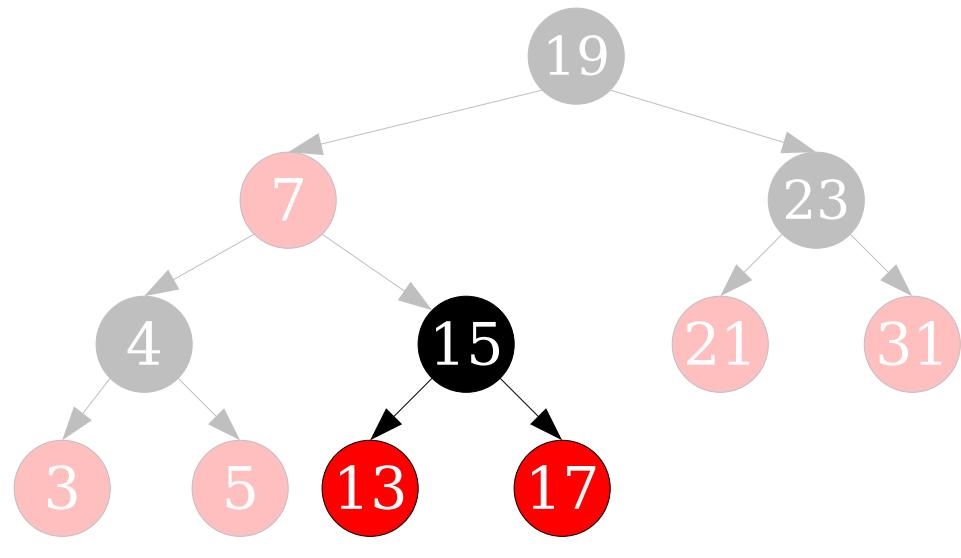


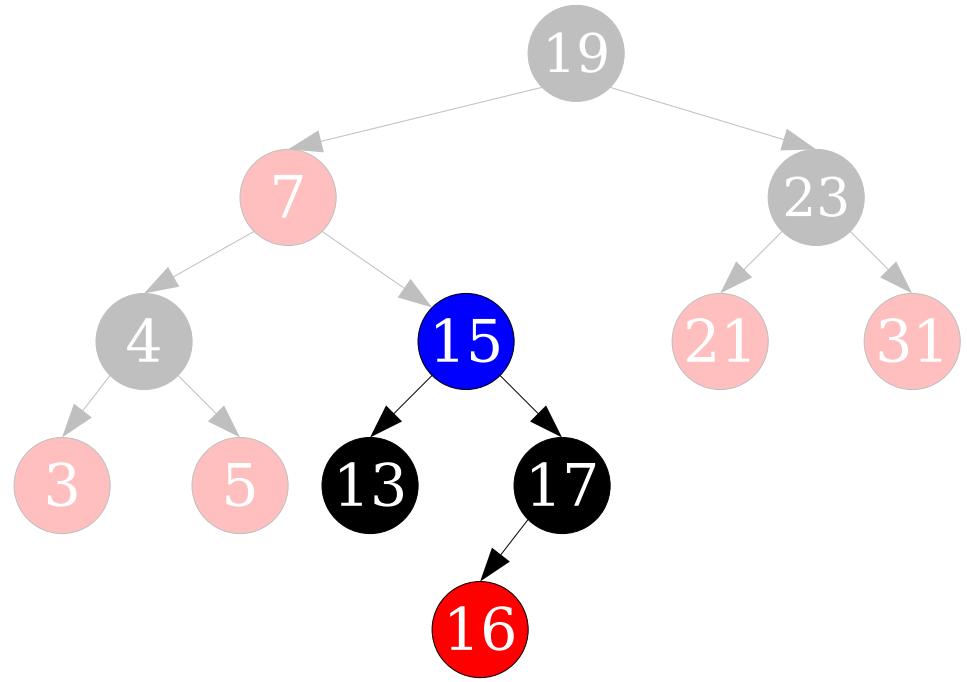
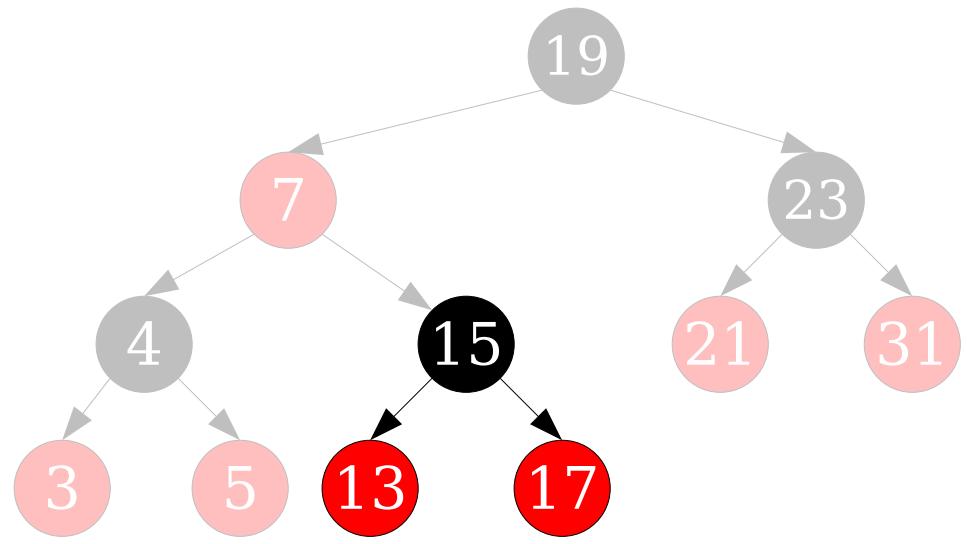


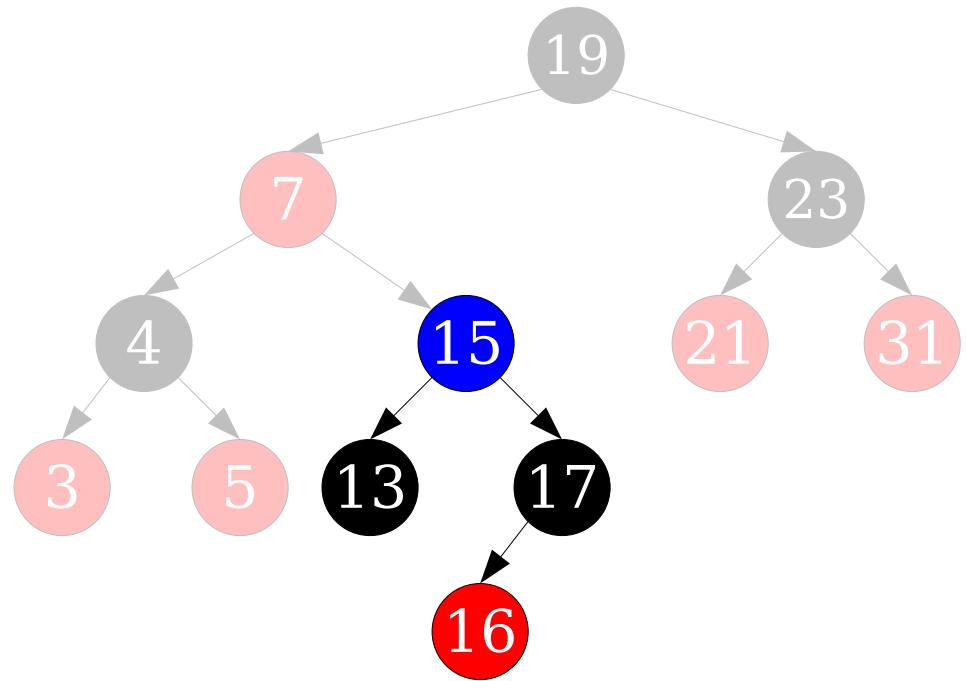
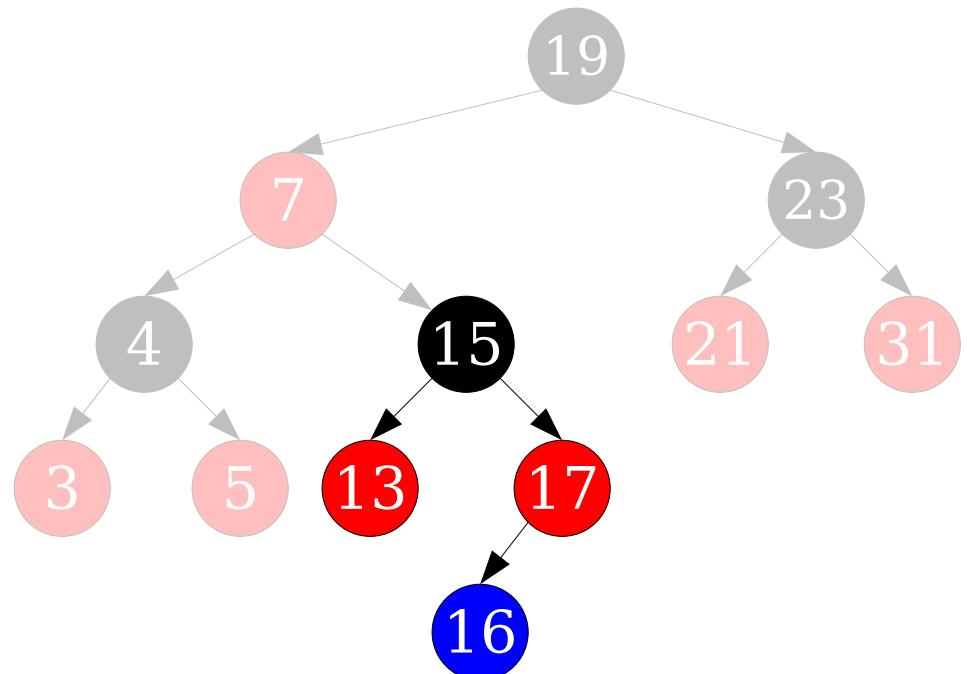


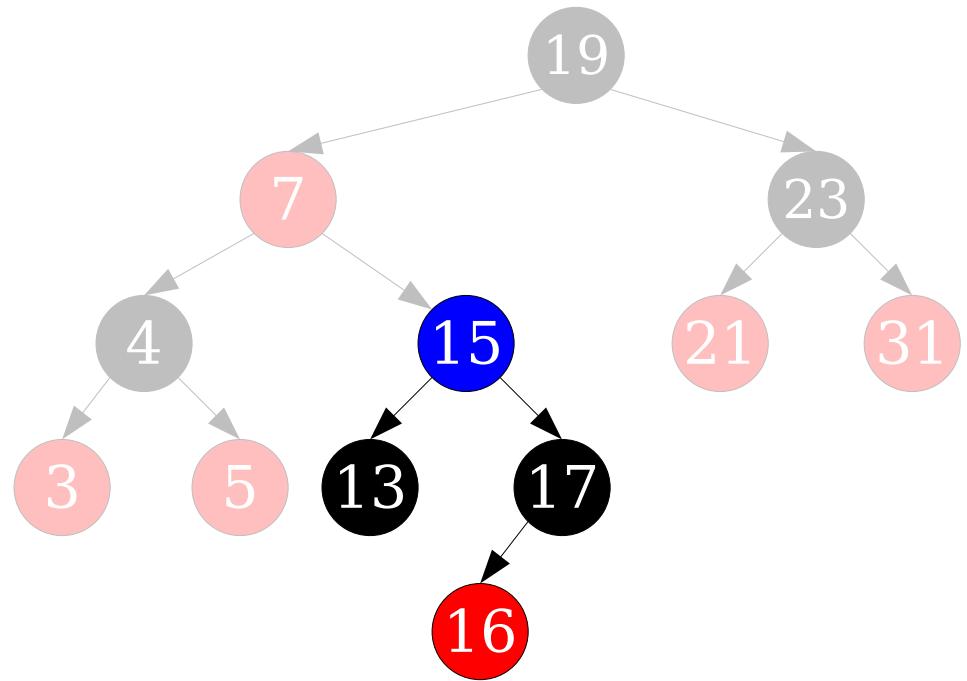
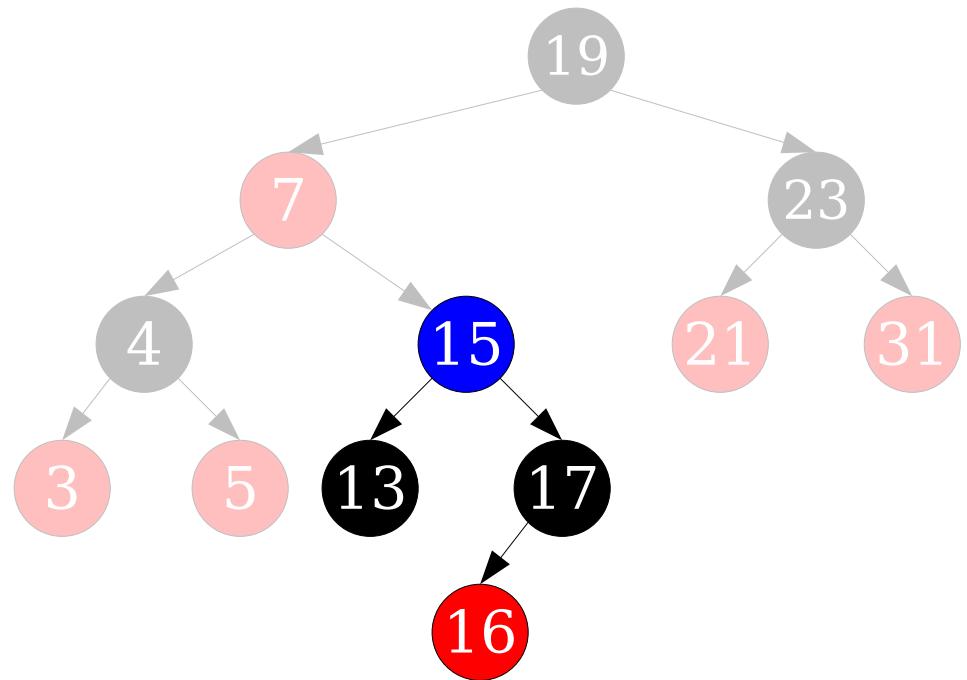


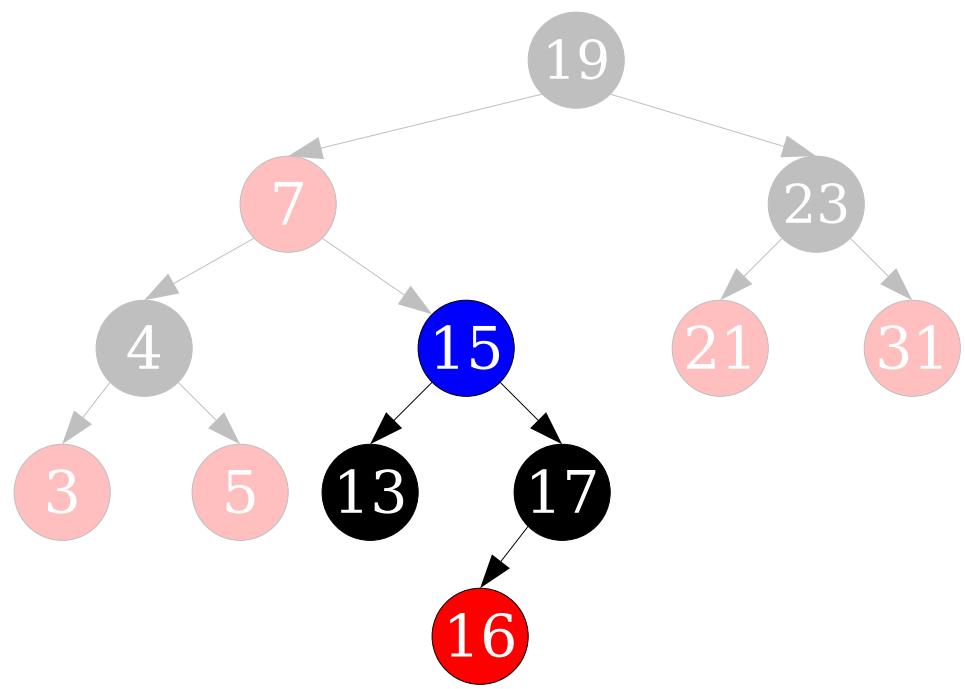


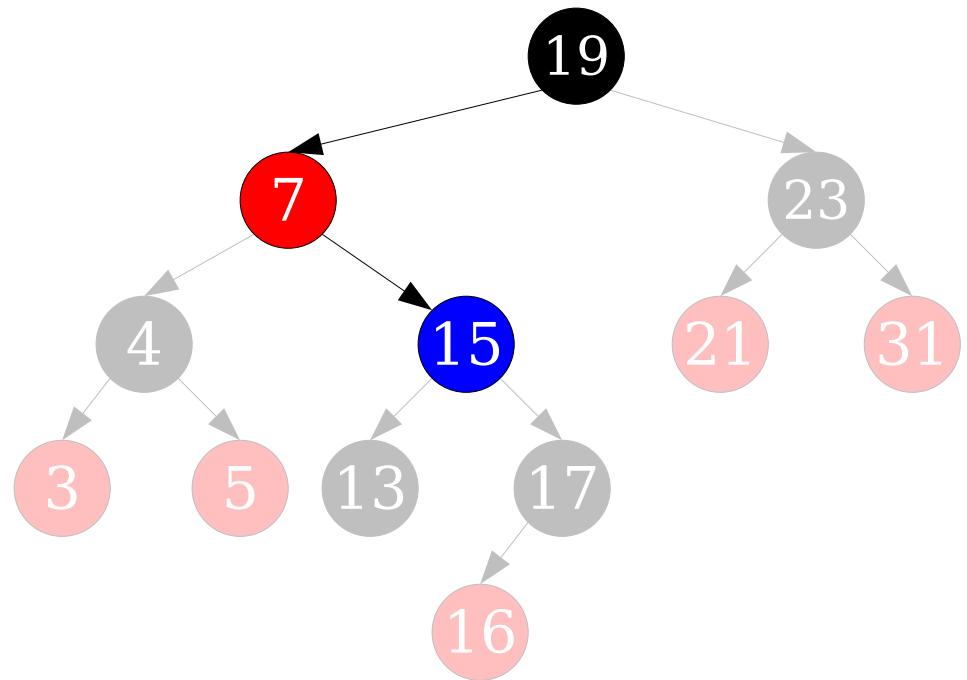


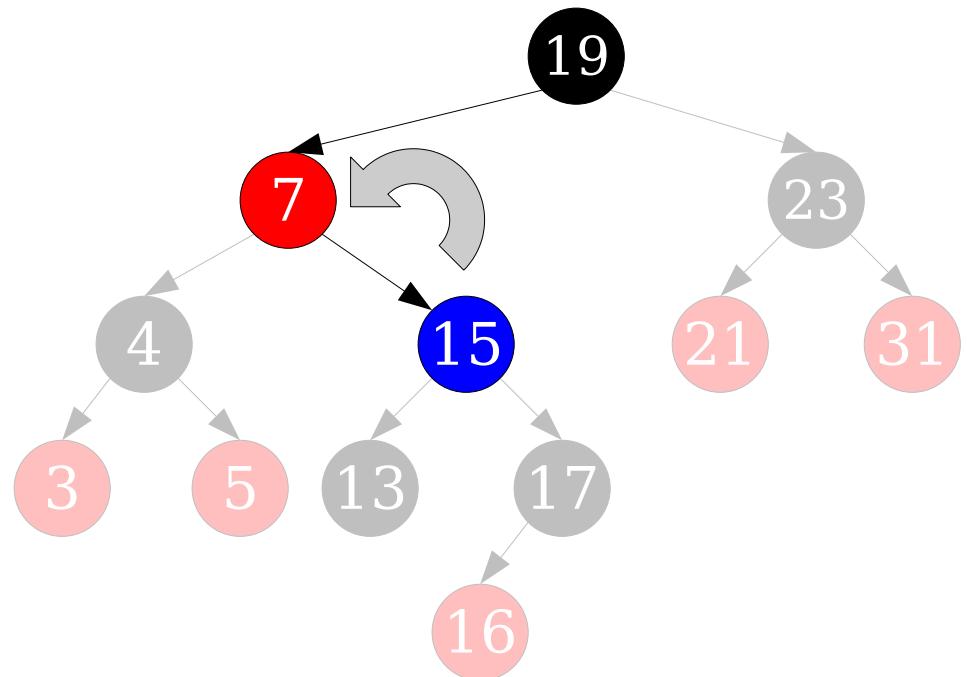


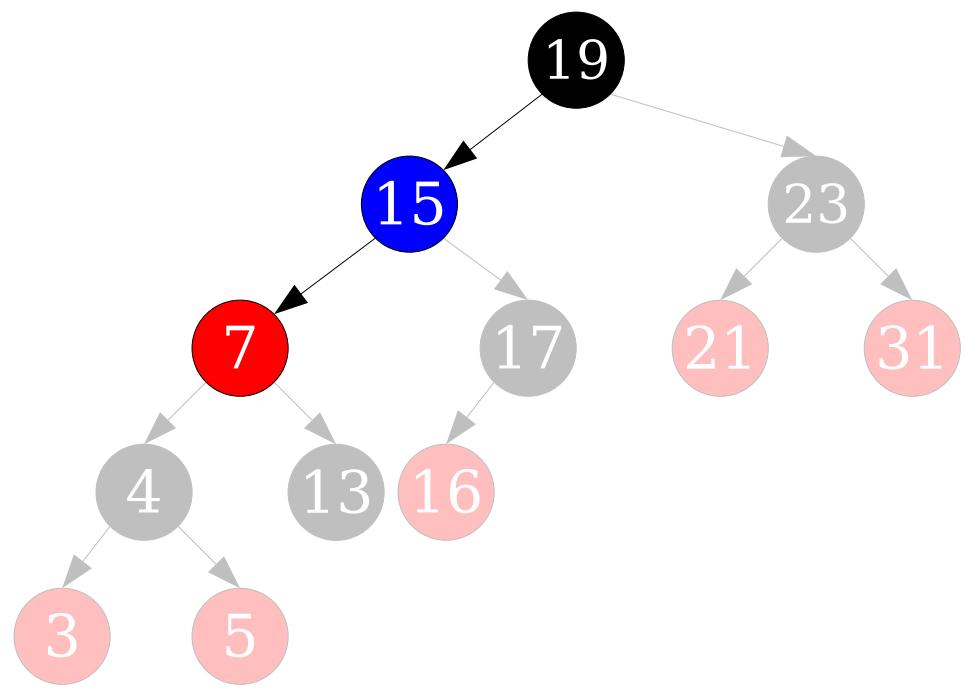


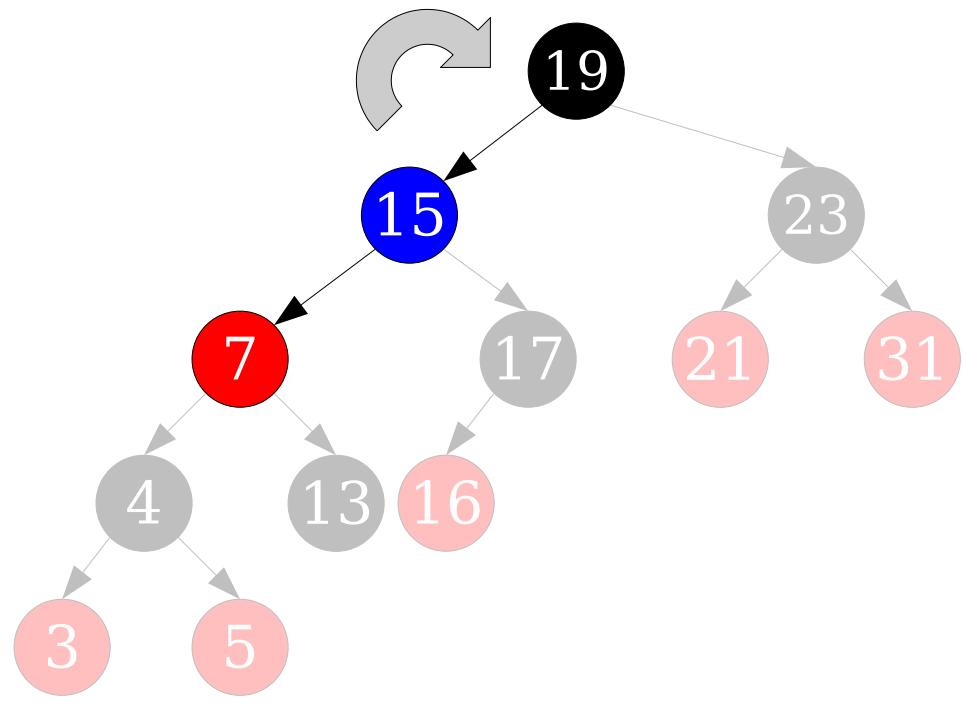


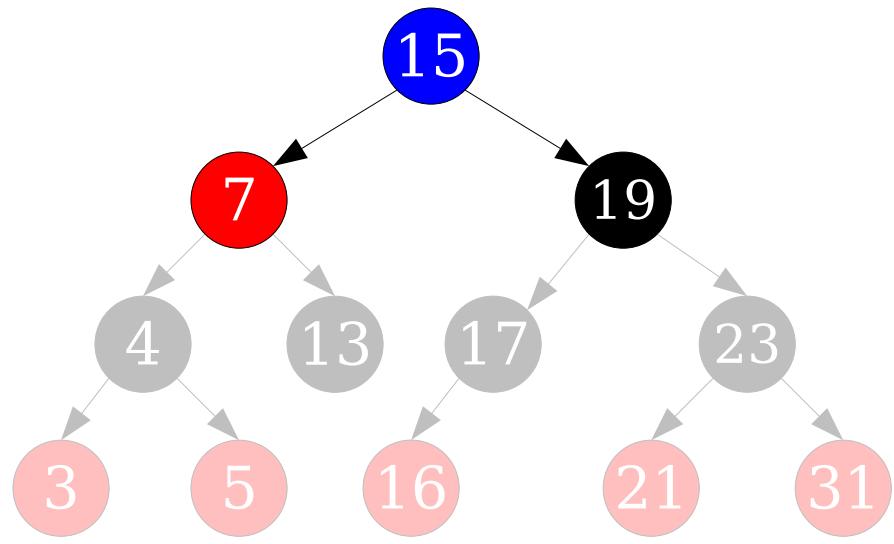


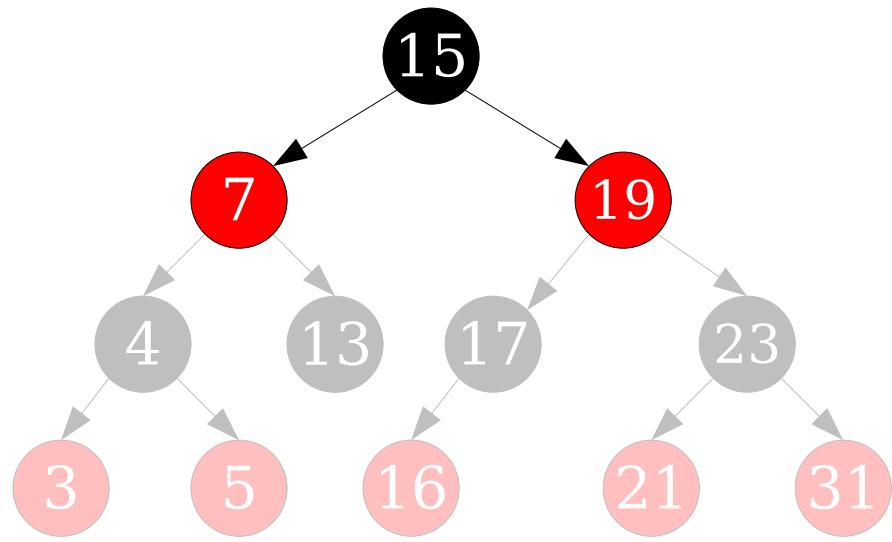


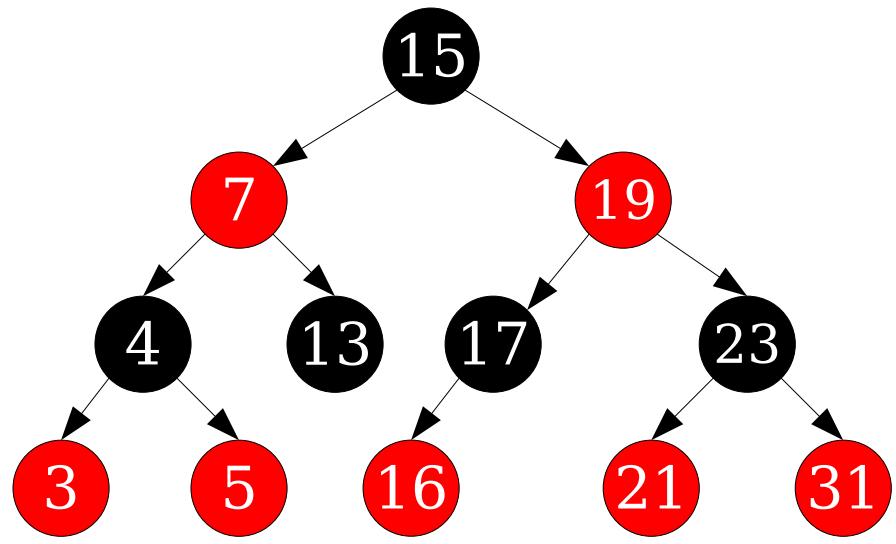








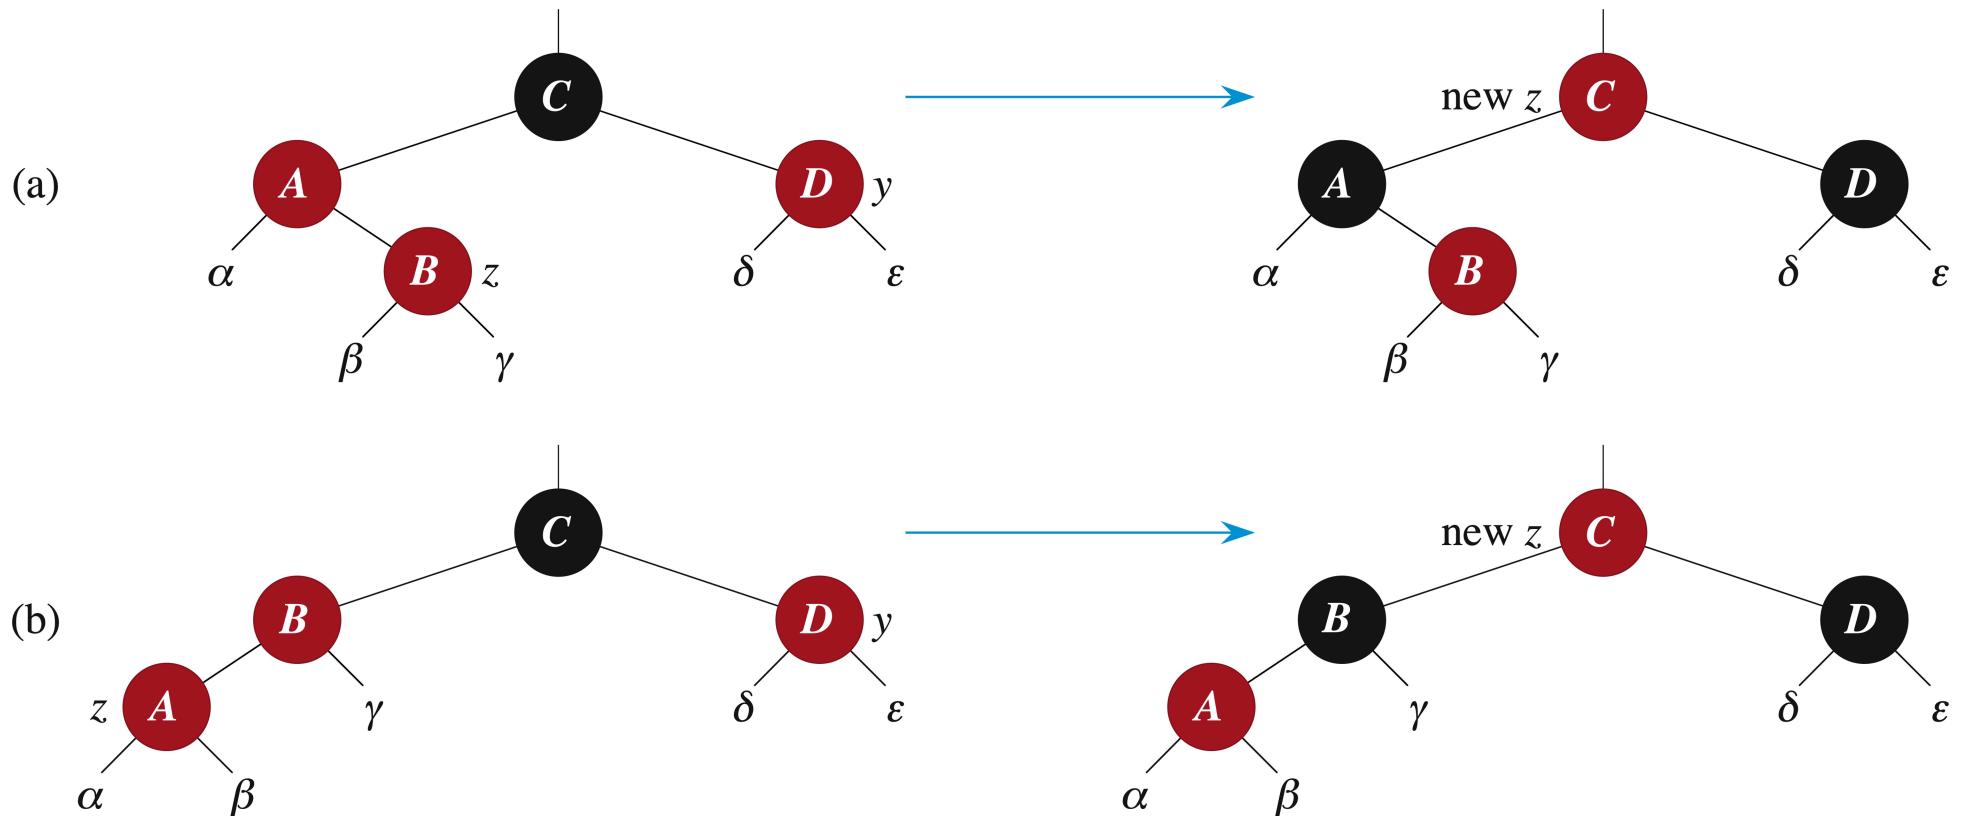




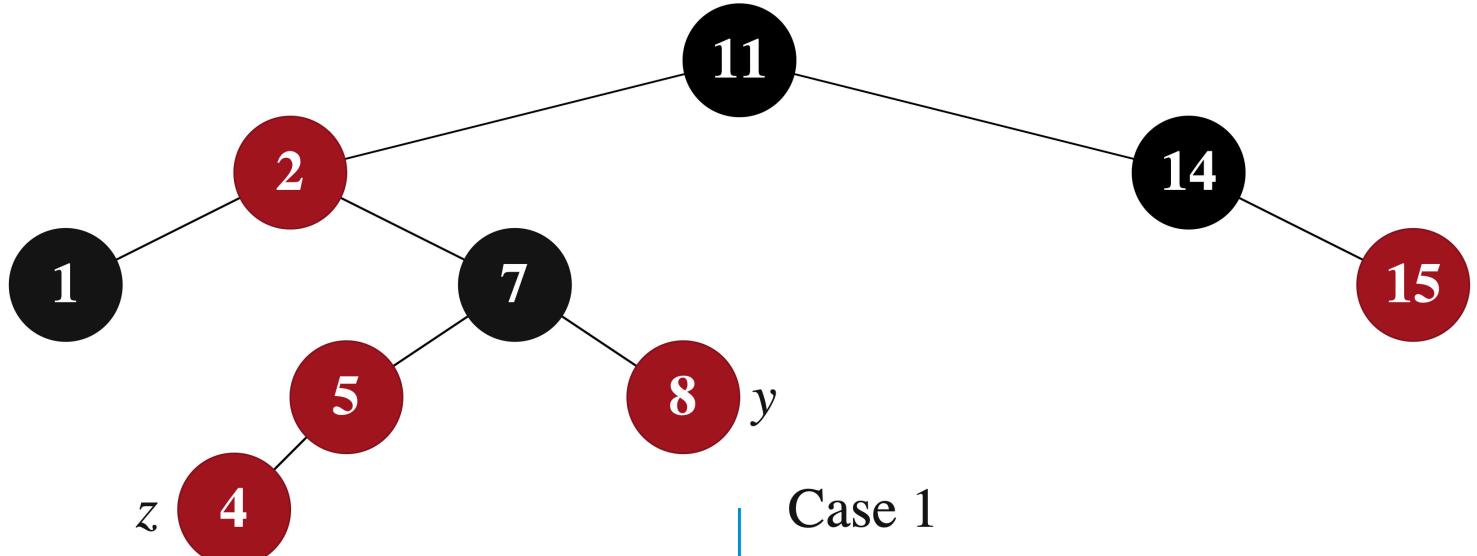
Insert Into Red Black Tree

RB-INSERT(T, z)

```
1   $x = T.root$                                 // node being compared with  $z$ 
2   $y = T.nil$                                  //  $y$  will be parent of  $z$ 
3  while  $x \neq T.nil$                       // descend until reaching the sentinel
4       $y = x$ 
5      if  $z.key < x.key$ 
6           $x = x.left$ 
7      else  $x = x.right$ 
8   $z.p = y$                                     // found the location—insert  $z$  with parent  $y$ 
9  if  $y == T.nil$ 
10      $T.root = z$                             // tree  $T$  was empty
11  elseif  $z.key < y.key$ 
12       $y.left = z$ 
13  else  $y.right = z$ 
14   $z.left = T.nil$                            // both of  $z$ 's children are the sentinel
15   $z.right = T.nil$ 
16   $z.color = \text{RED}$                          // the new node starts out red
17  RB-INSERT-FIXUP( $T, z$ )                    // correct any violations of red-black properties
```

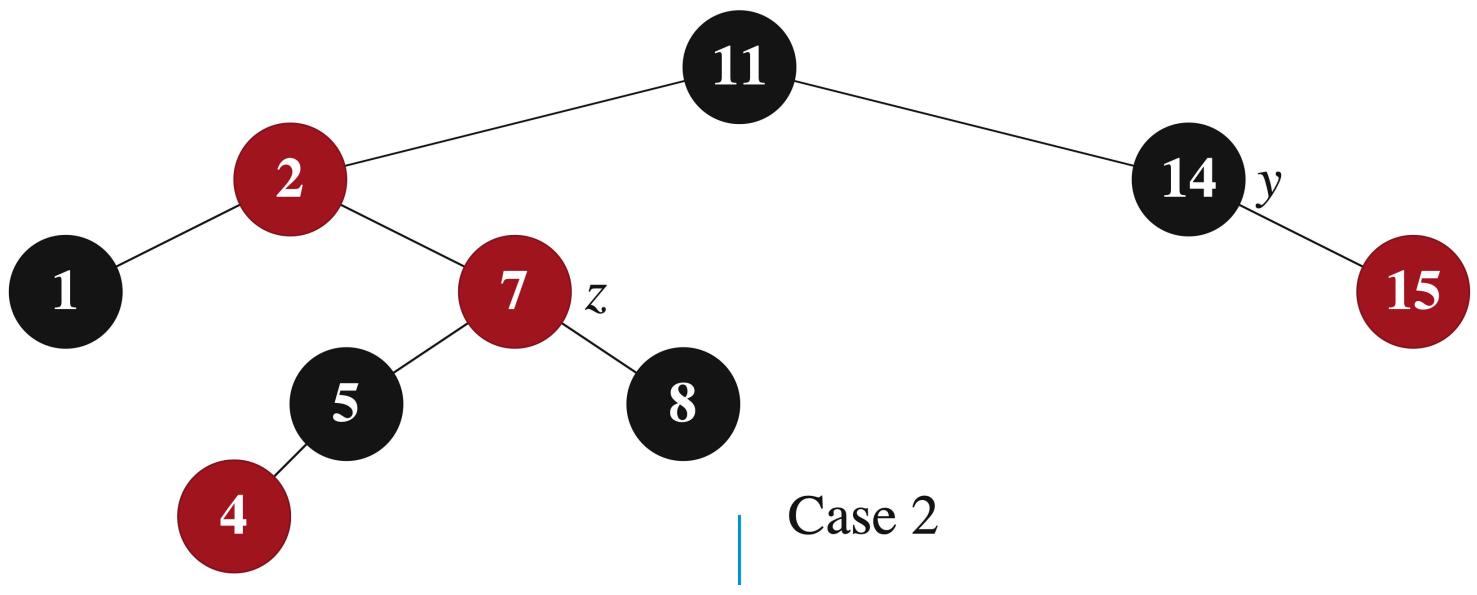


(a)

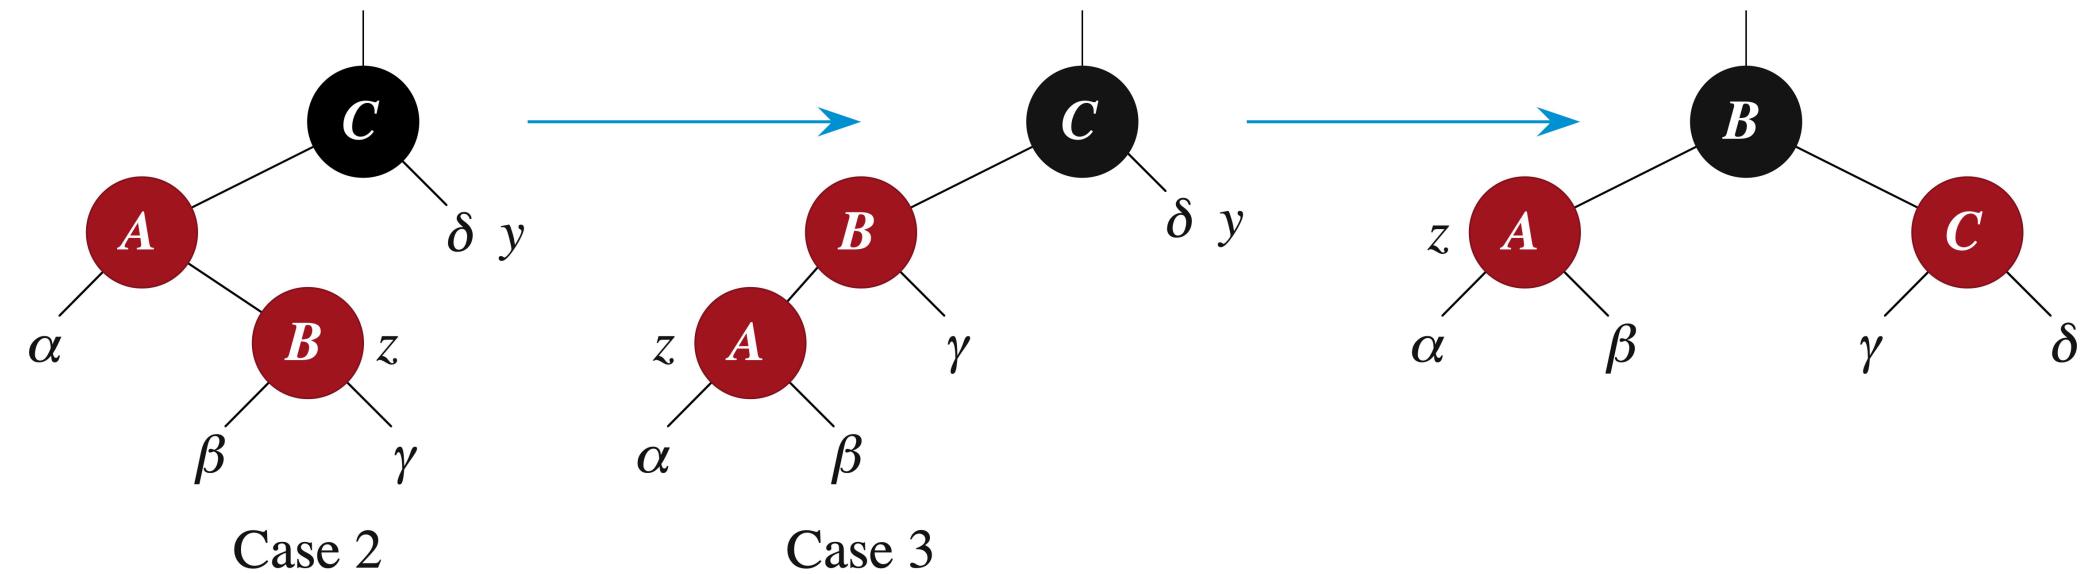


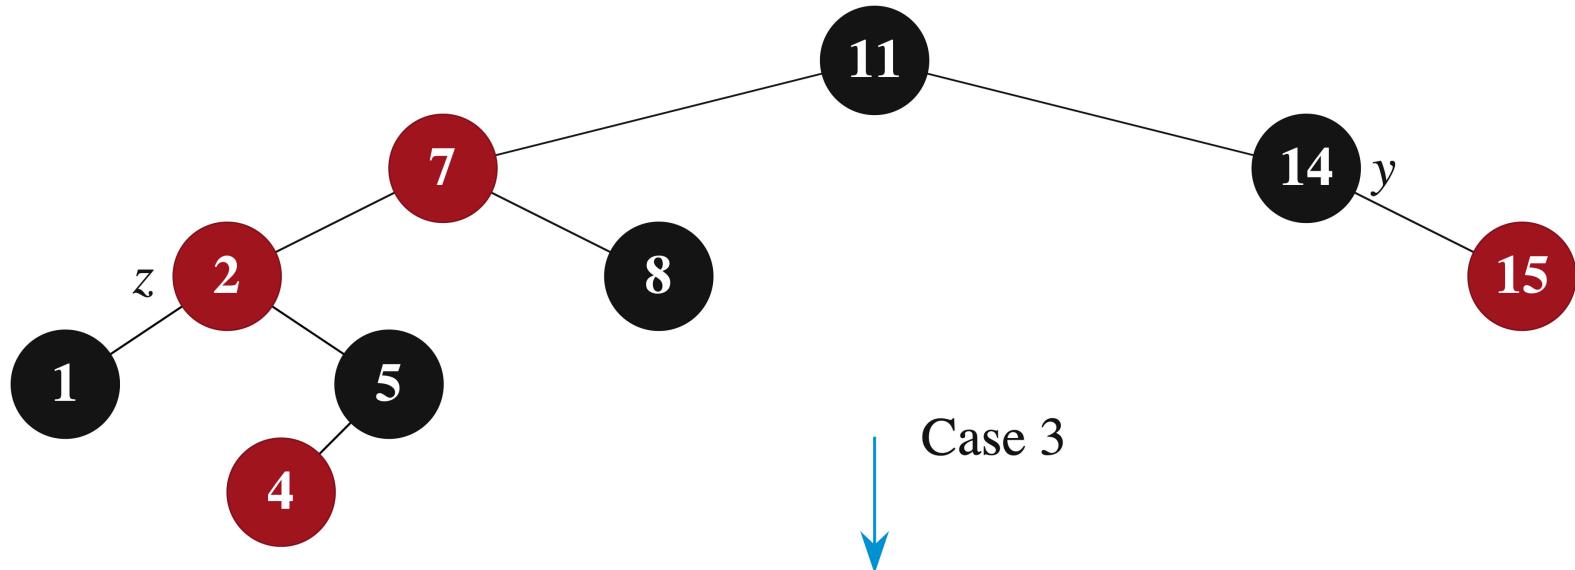
Case 1

(b)

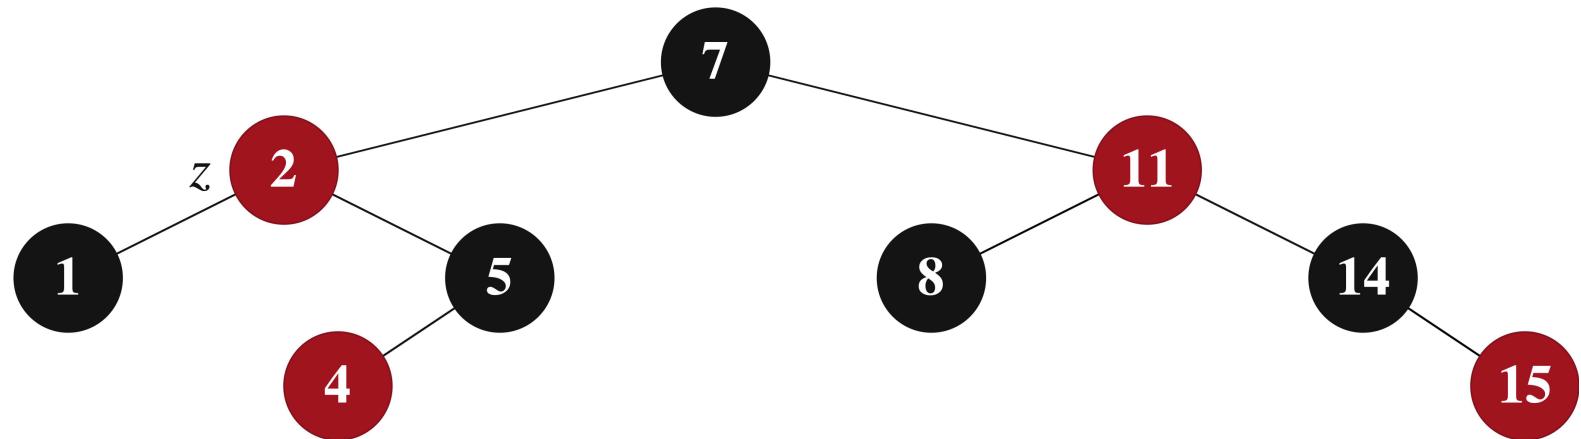


Case 2





Case 3



RB-INSERT-FIXUP(T, z)

```
1  while  $z.p.color == \text{RED}$ 
2      if  $z.p == z.p.p.left$                                 // is  $z$ 's parent a left child?
3           $y = z.p.p.right$                                //  $y$  is  $z$ 's uncle
4          if  $y.color == \text{RED}$                           // are  $z$ 's parent and uncle both red?
5               $z.p.color = \text{BLACK}$ 
6               $y.color = \text{BLACK}$ 
7               $z.p.p.color = \text{RED}$ 
8               $z = z.p.p$ 
9      else
10         if  $z == z.p.right$ 
11              $z = z.p$ 
12             LEFT-ROTATE( $T, z$ )                         } case 2
13              $z.p.color = \text{BLACK}$ 
14              $z.p.p.color = \text{RED}$ 
15             RIGHT-ROTATE( $T, z.p.p$ )                  } case 3
```