

Single Source Shortest Path Problem

Outline for Today

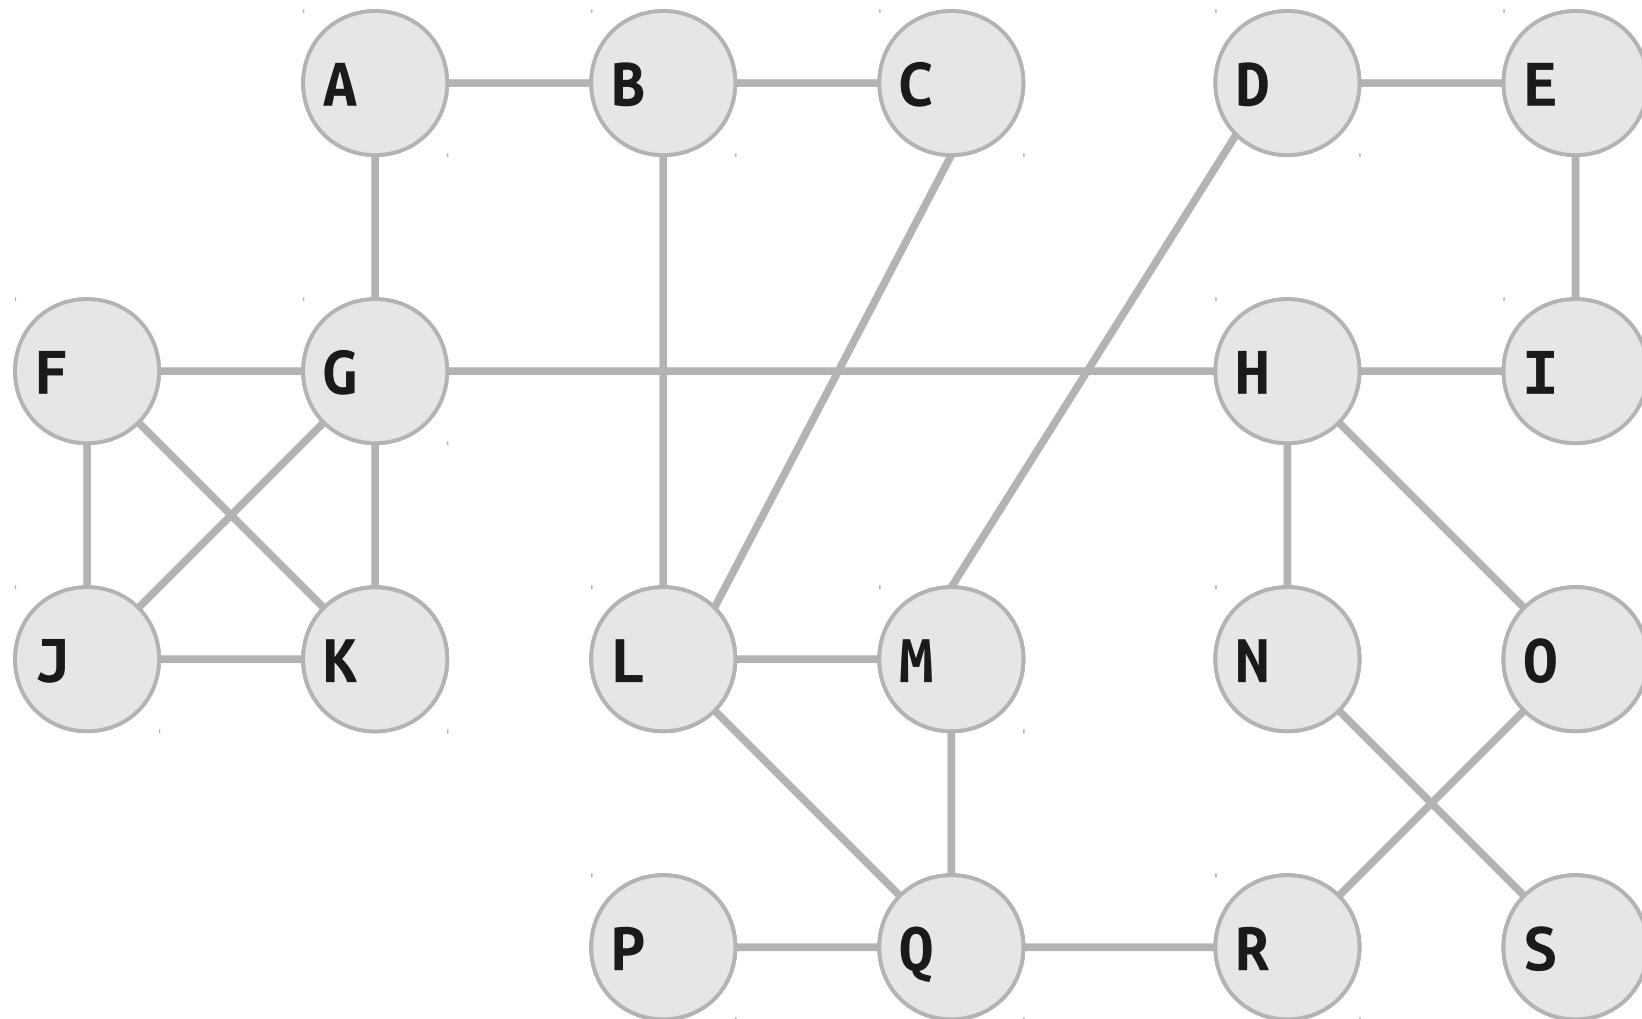
- **BFS/DFS**
- **Dijkstra's Algorithm**
 - An algorithm for finding shortest paths in more realistic settings
- **Bellman-Ford**
- **Practical Concerns**

Breadth-First Search

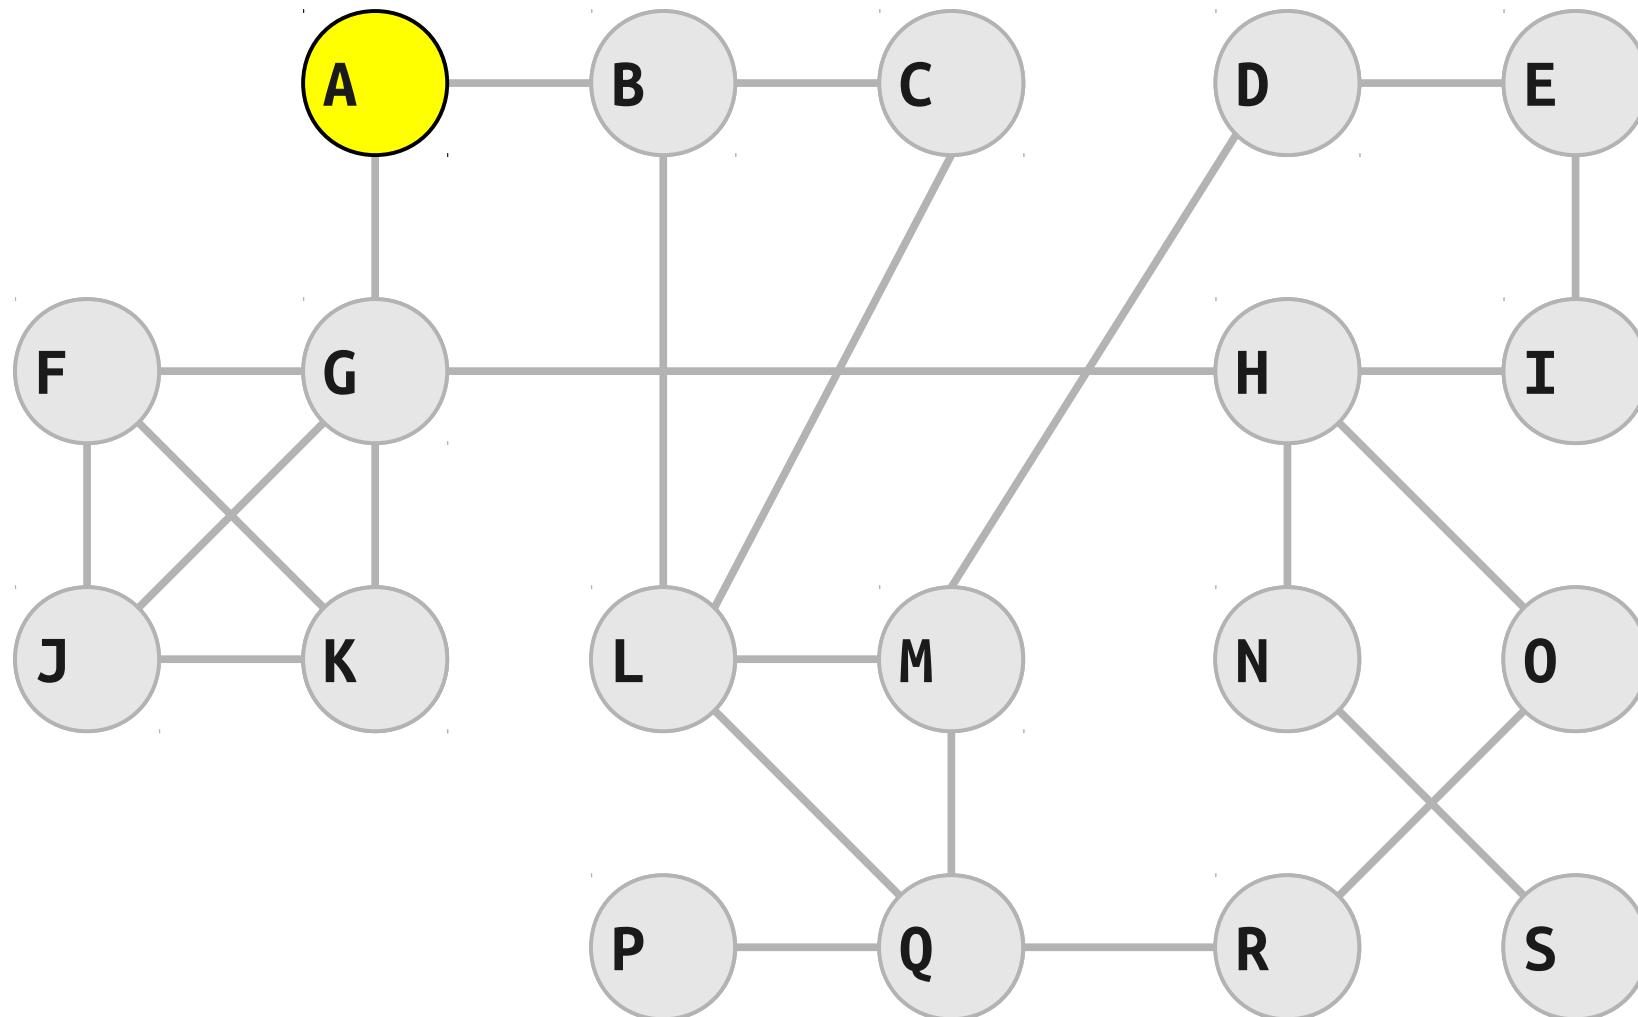
- Given an arbitrary graph $G = (V, E)$ and a starting node $s \in V$, **breadth-first search** finds shortest paths from s to each reachable node v .
- When implemented using an adjacency list, runs in $O(m + n)$ time, which we defined to be linear time on a graph.
- One correctness proof worked in terms of “layers:” the algorithm finds all nodes at distance 0, 1, 2, ... in order.

A Second Intuition for BFS

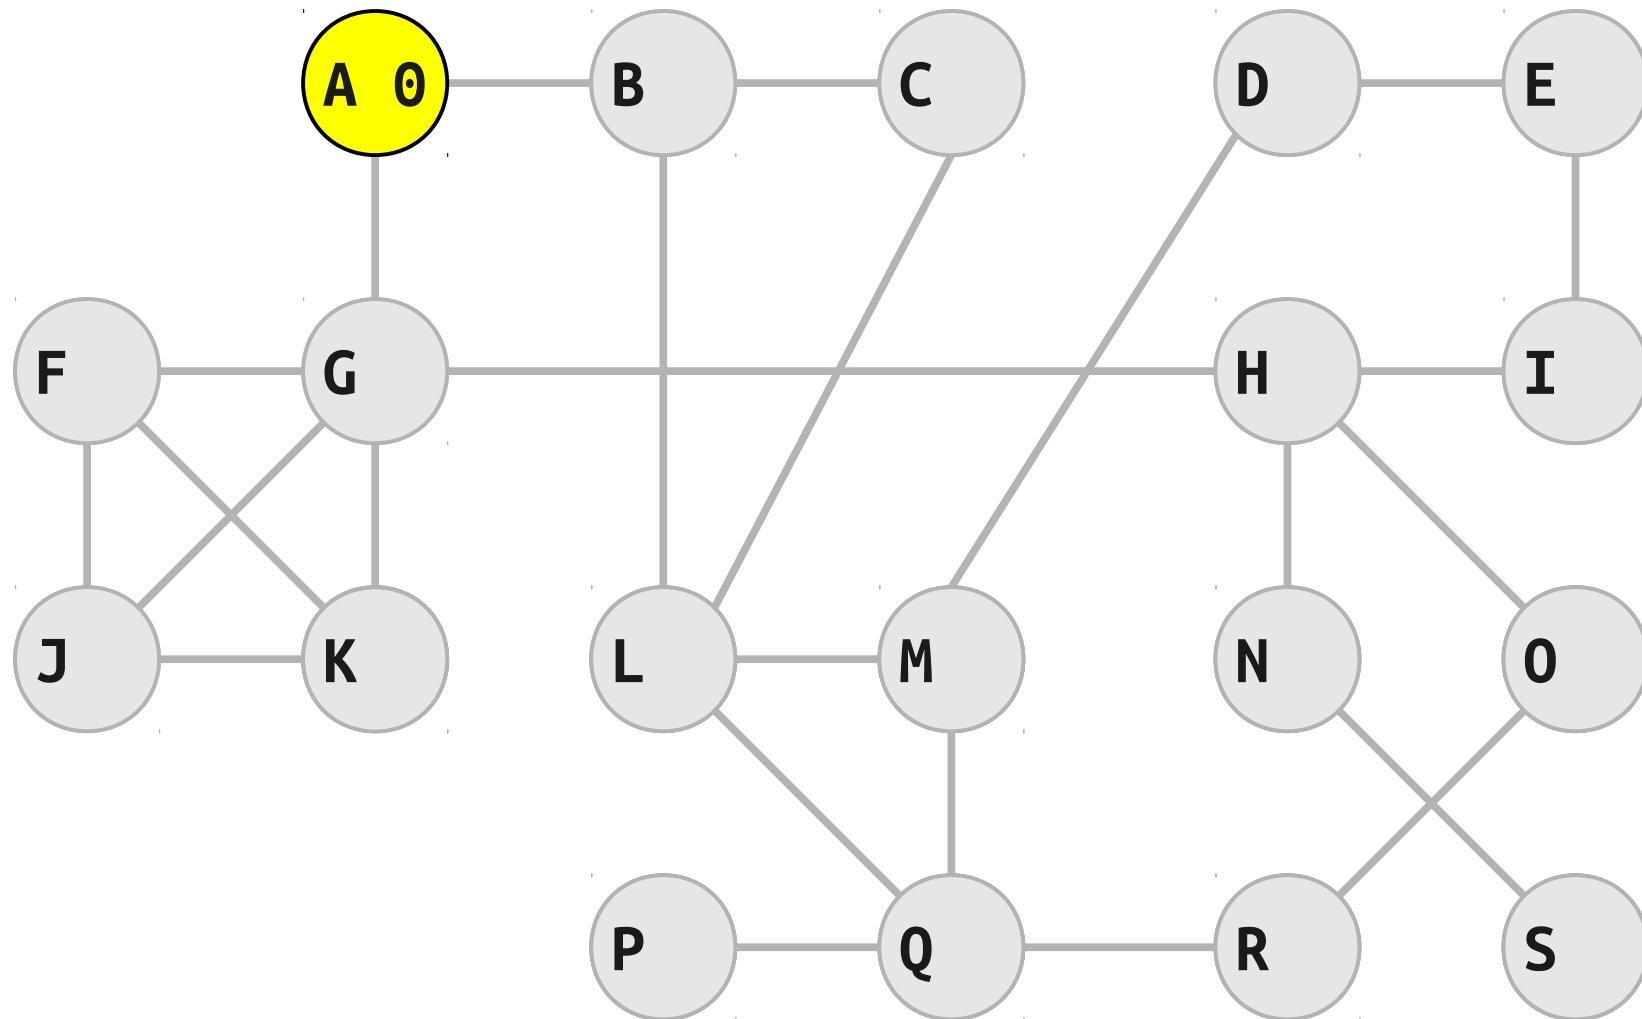
Breadth-First Search



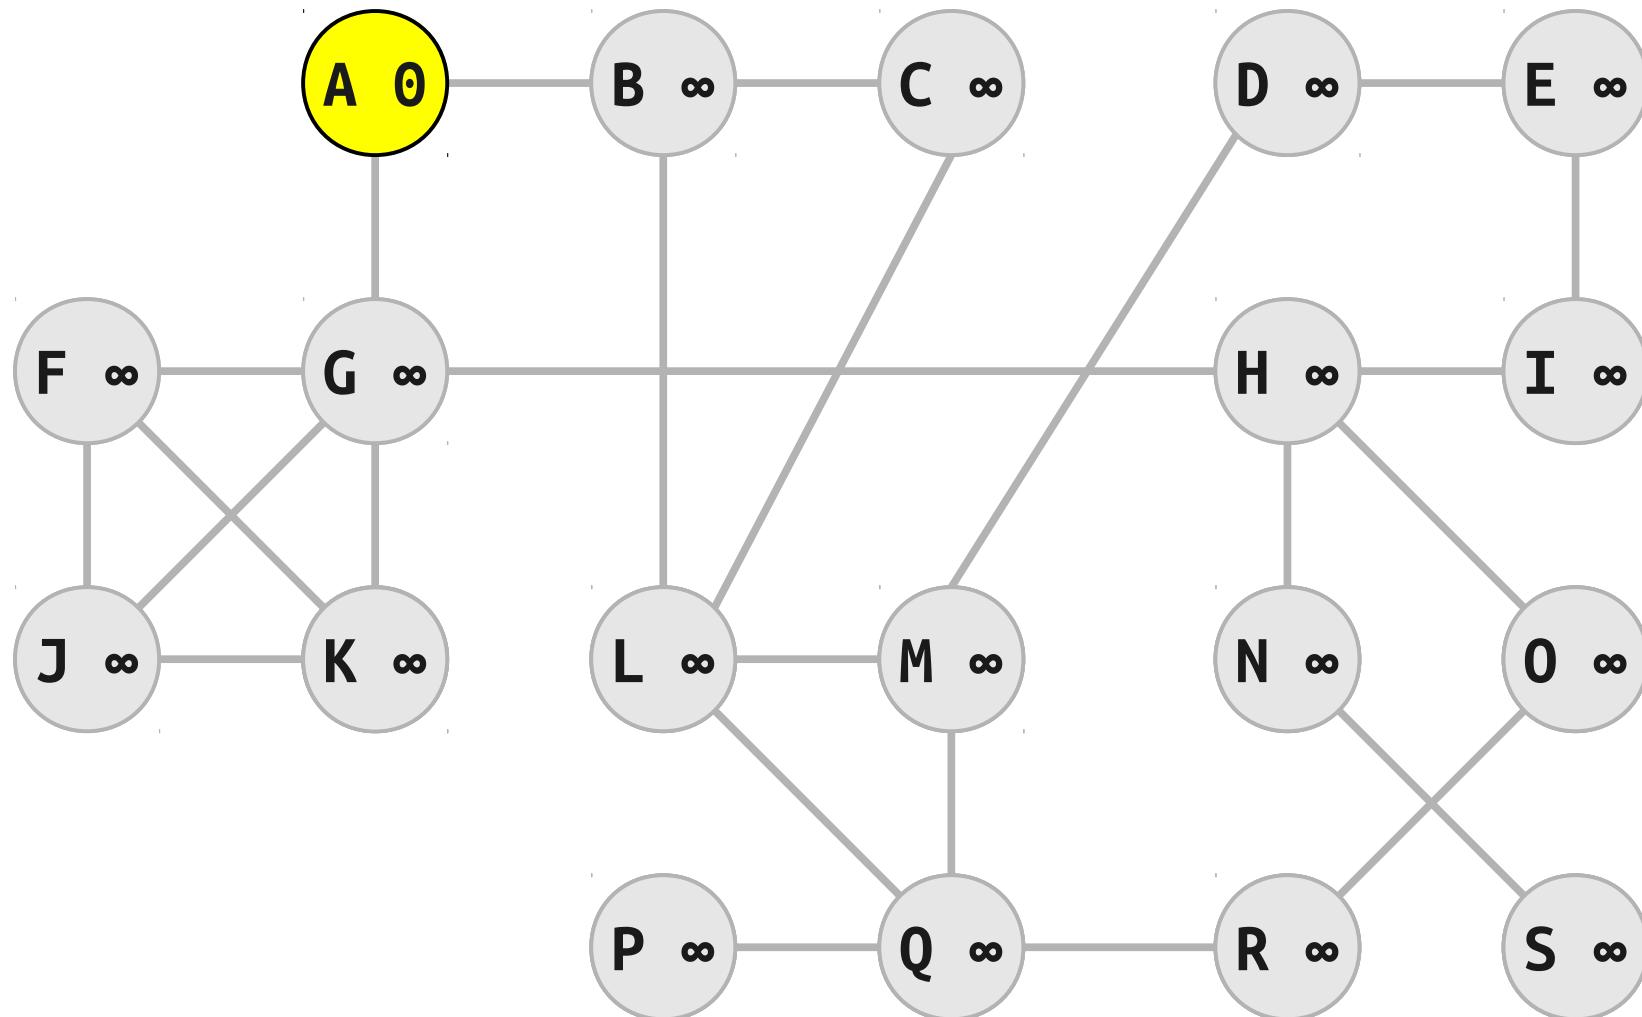
Breadth-First Search



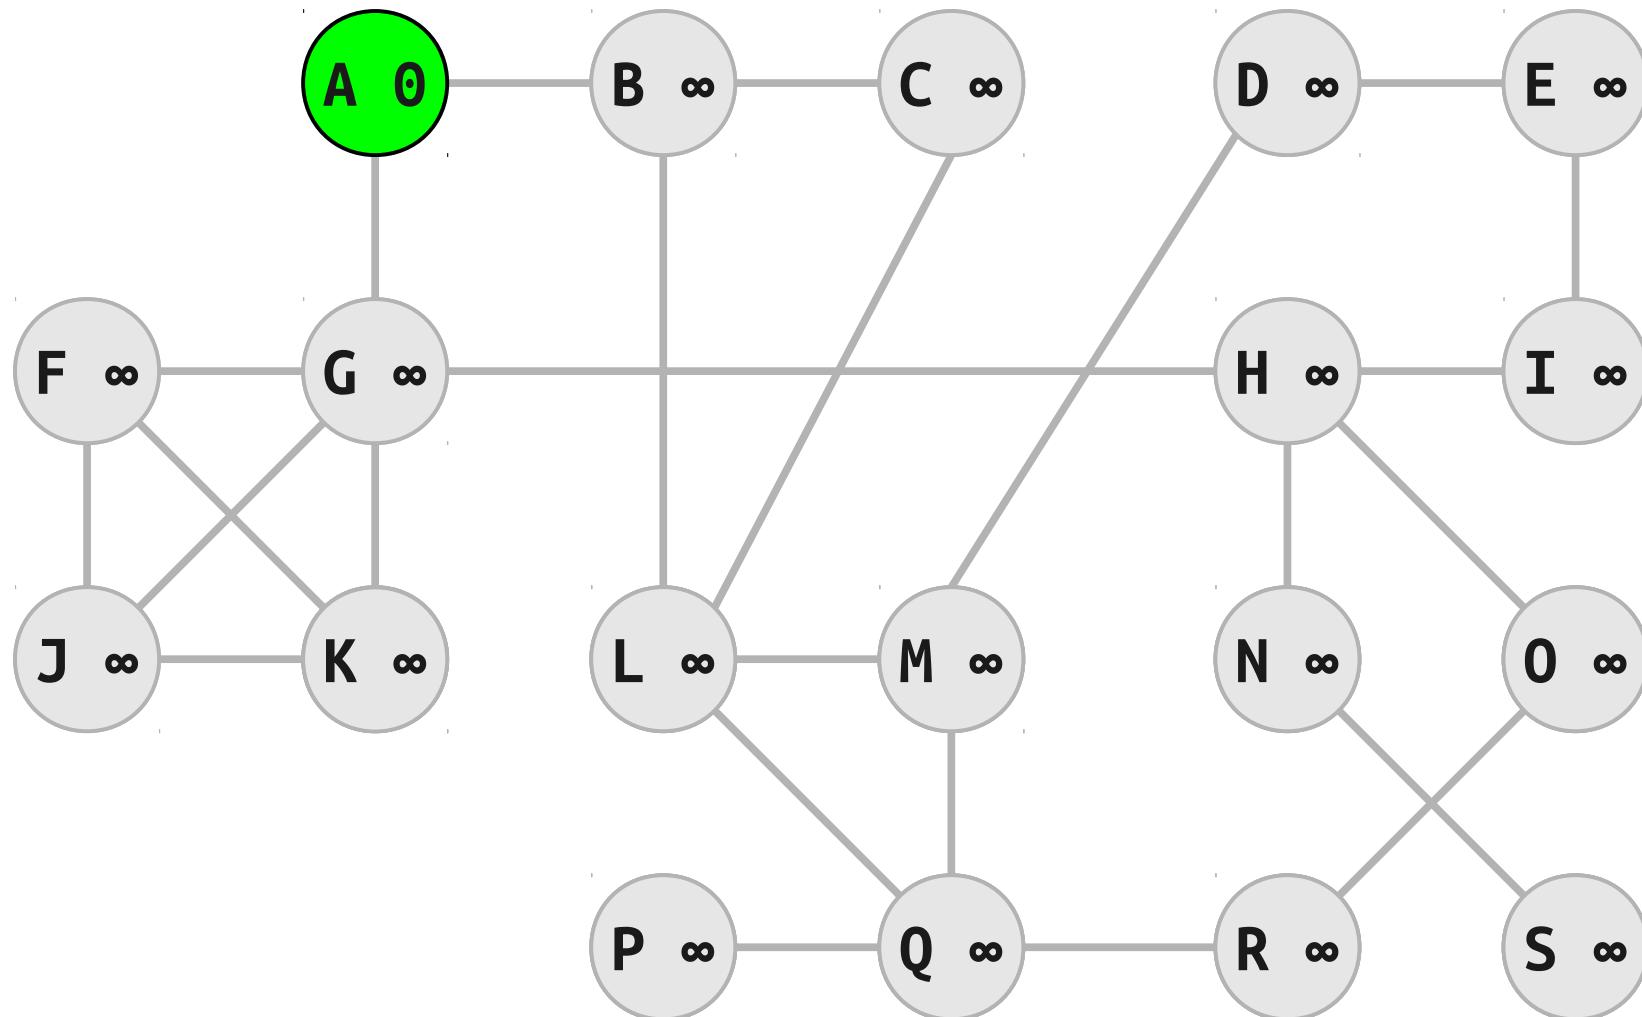
Breadth-First Search



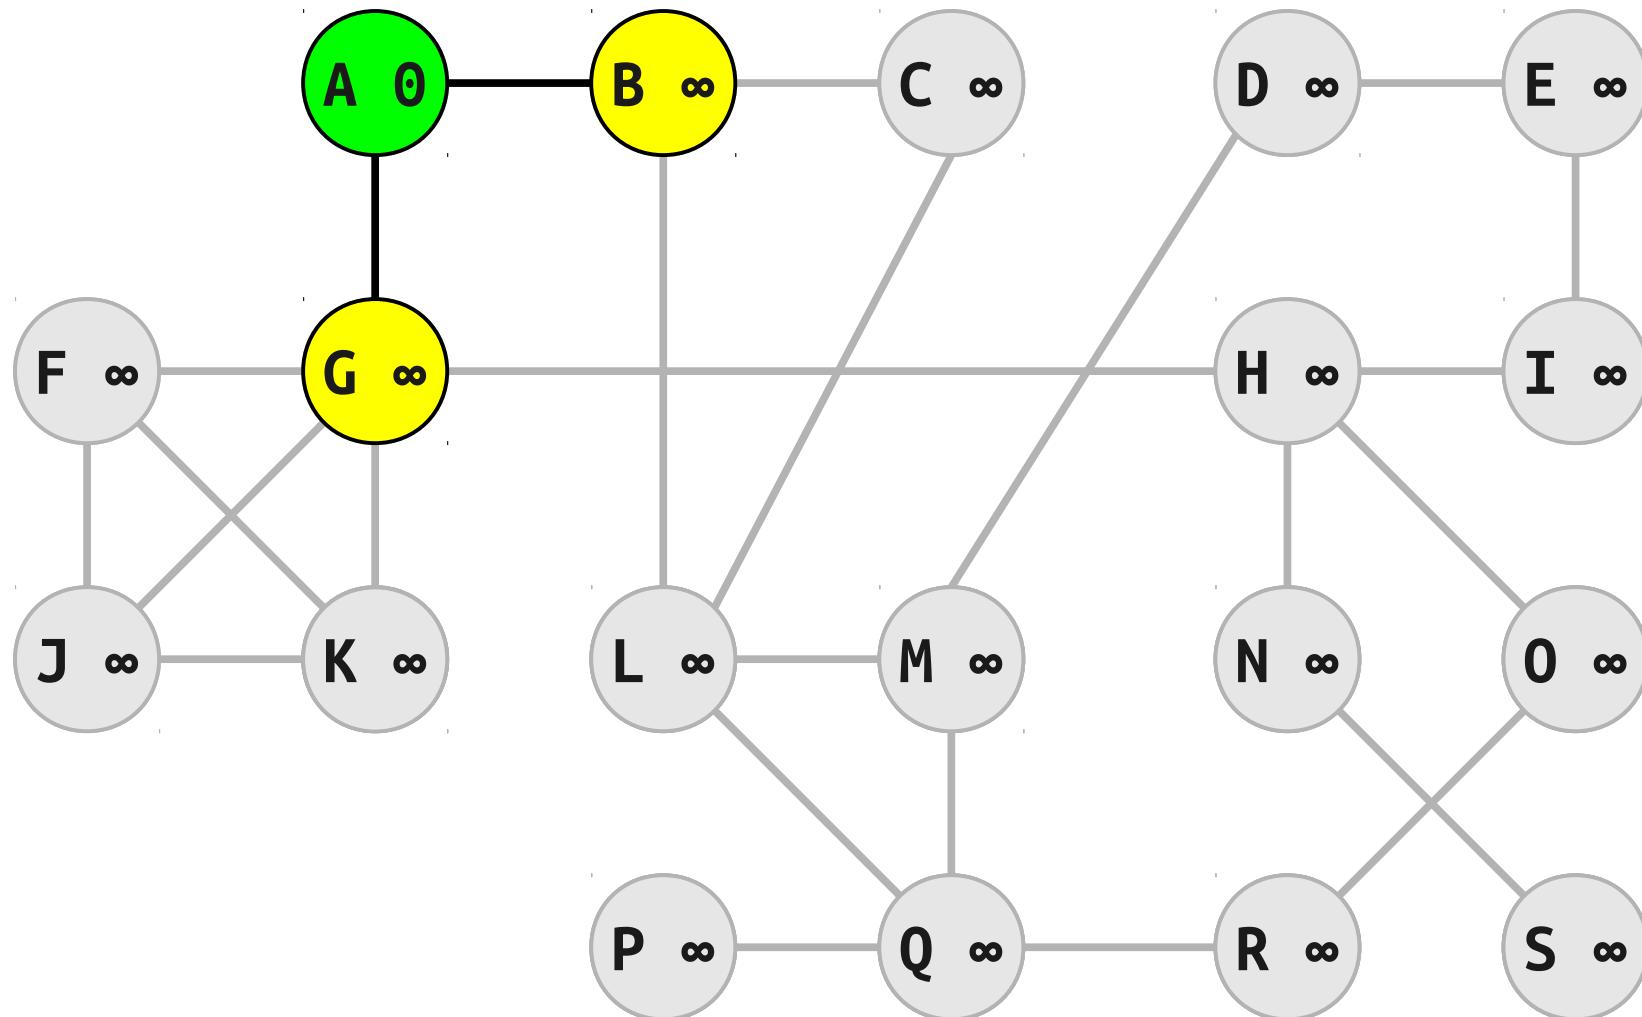
Breadth-First Search



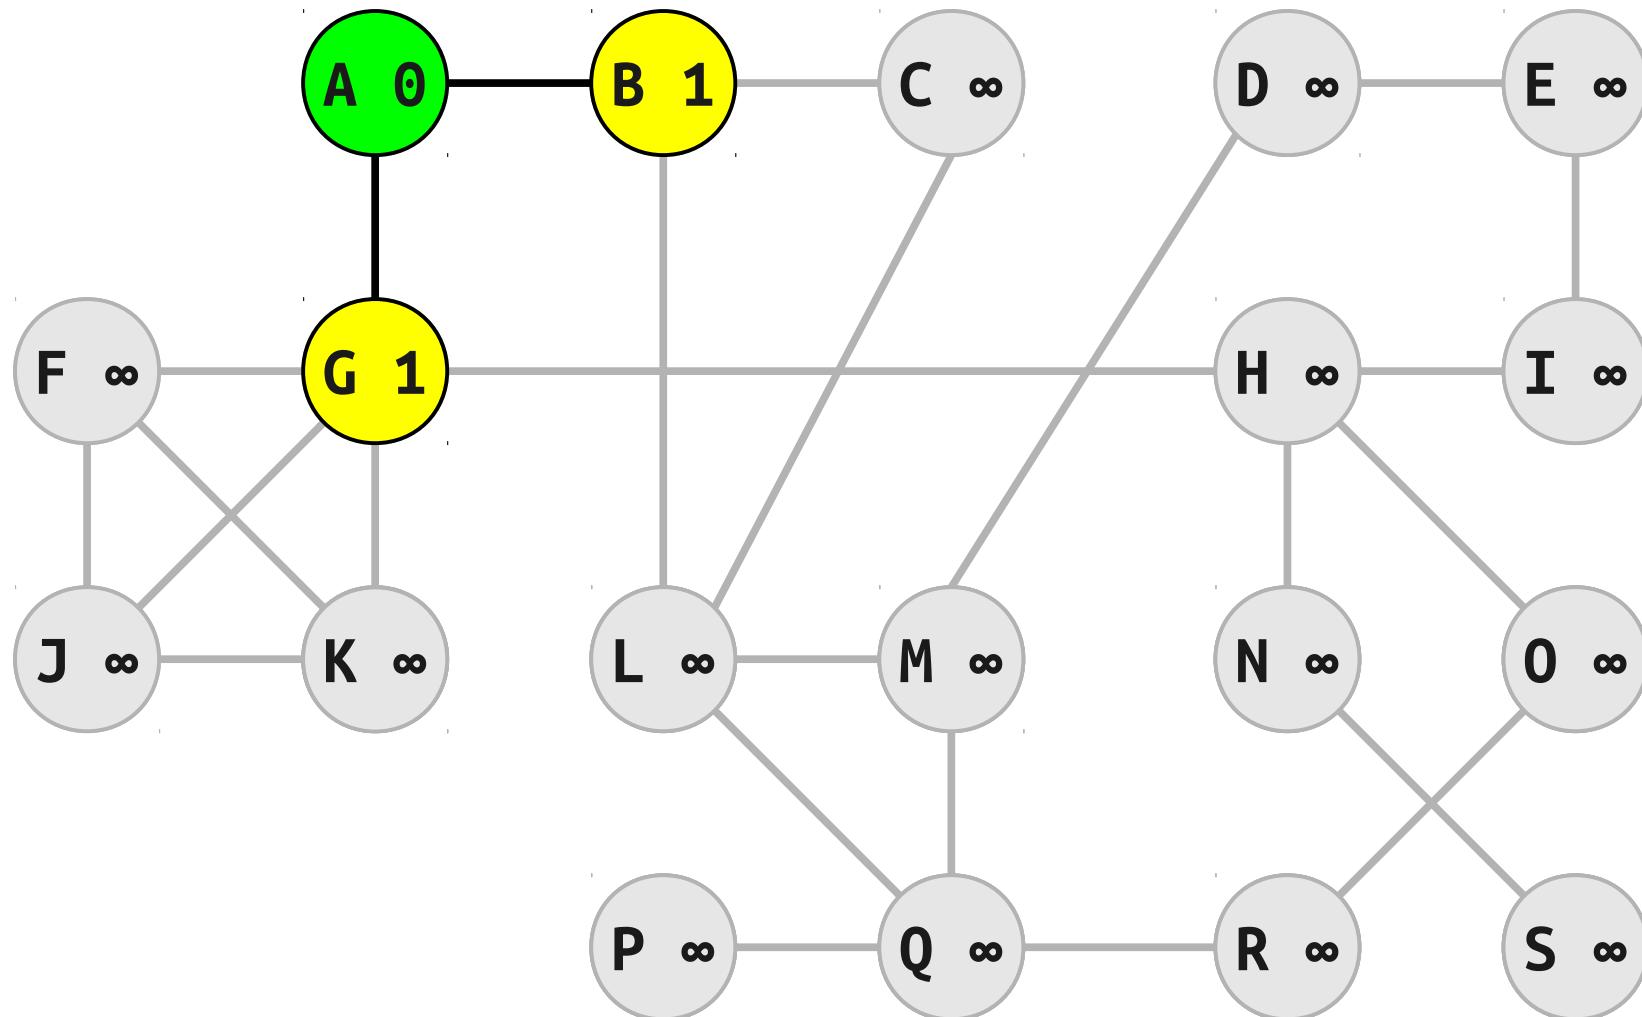
Breadth-First Search



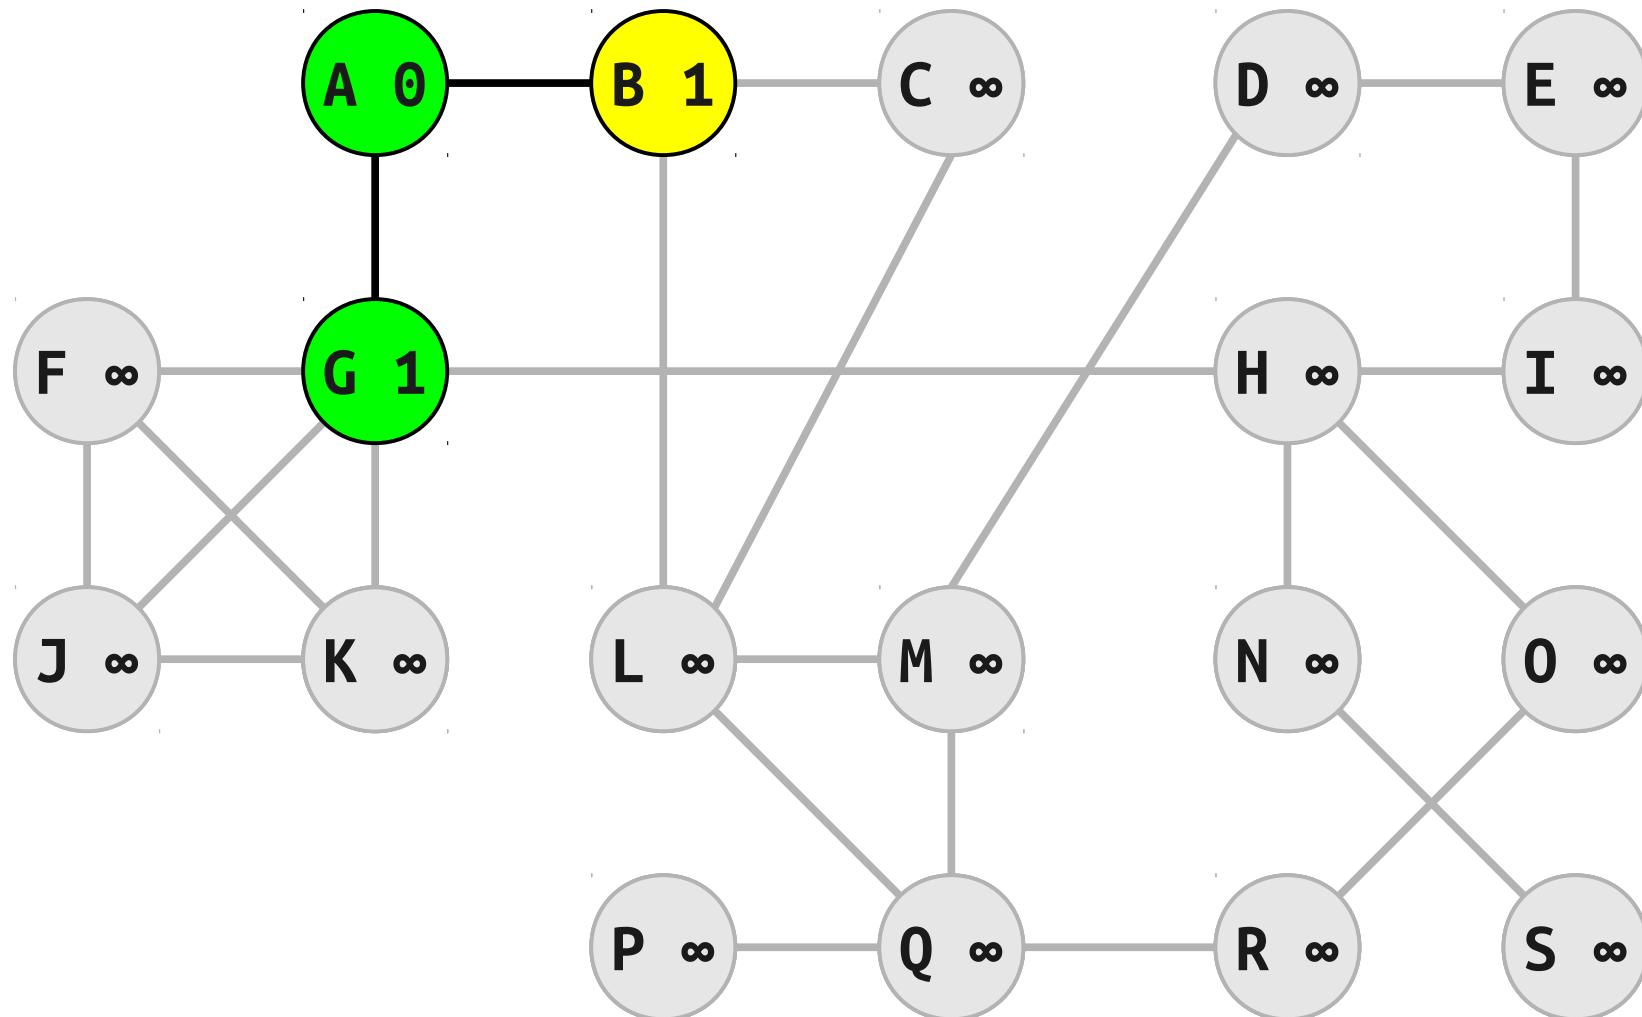
Breadth-First Search



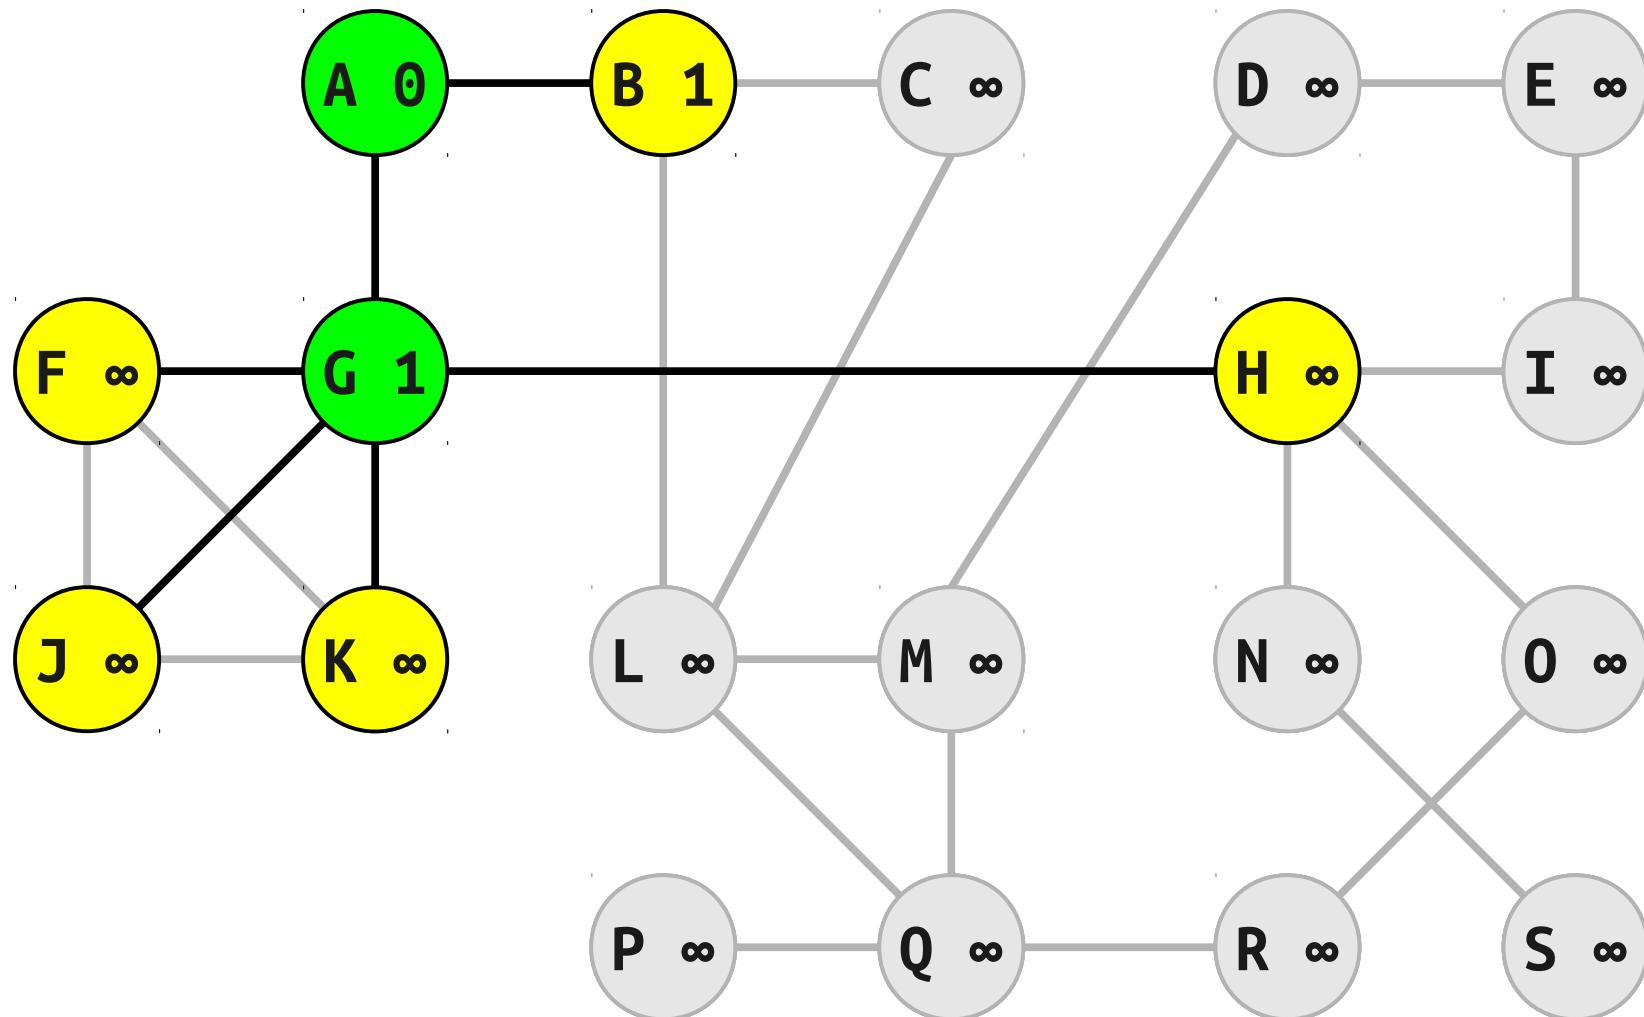
Breadth-First Search



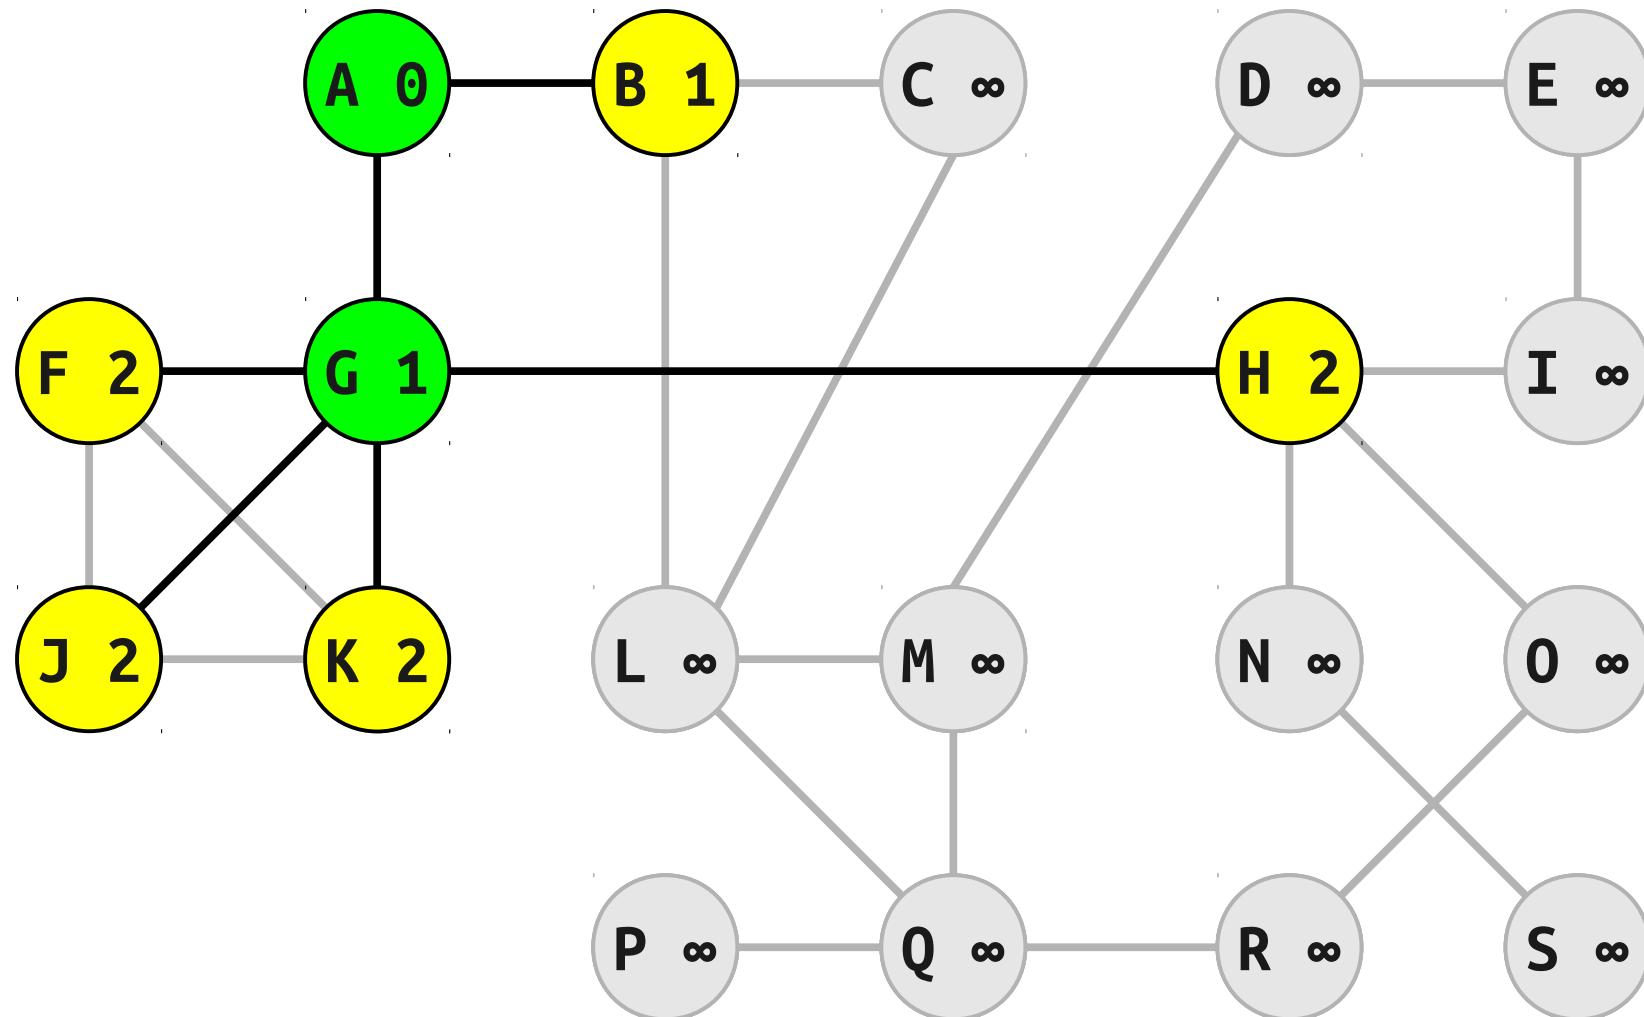
Breadth-First Search



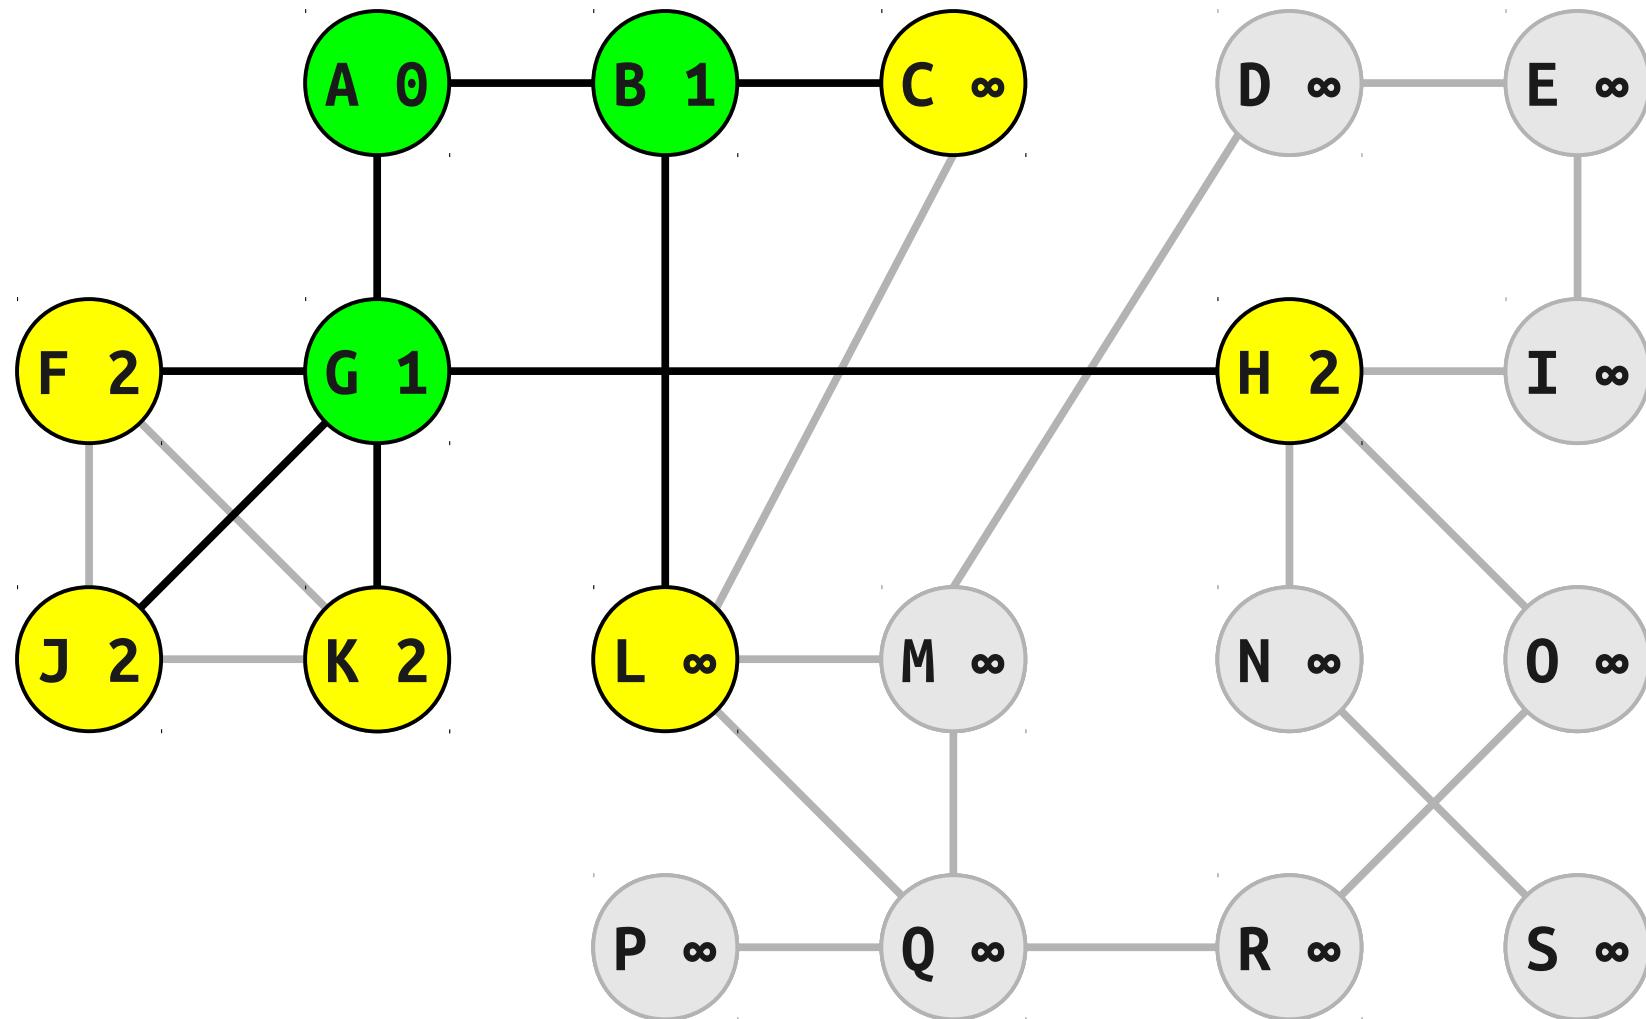
Breadth-First Search



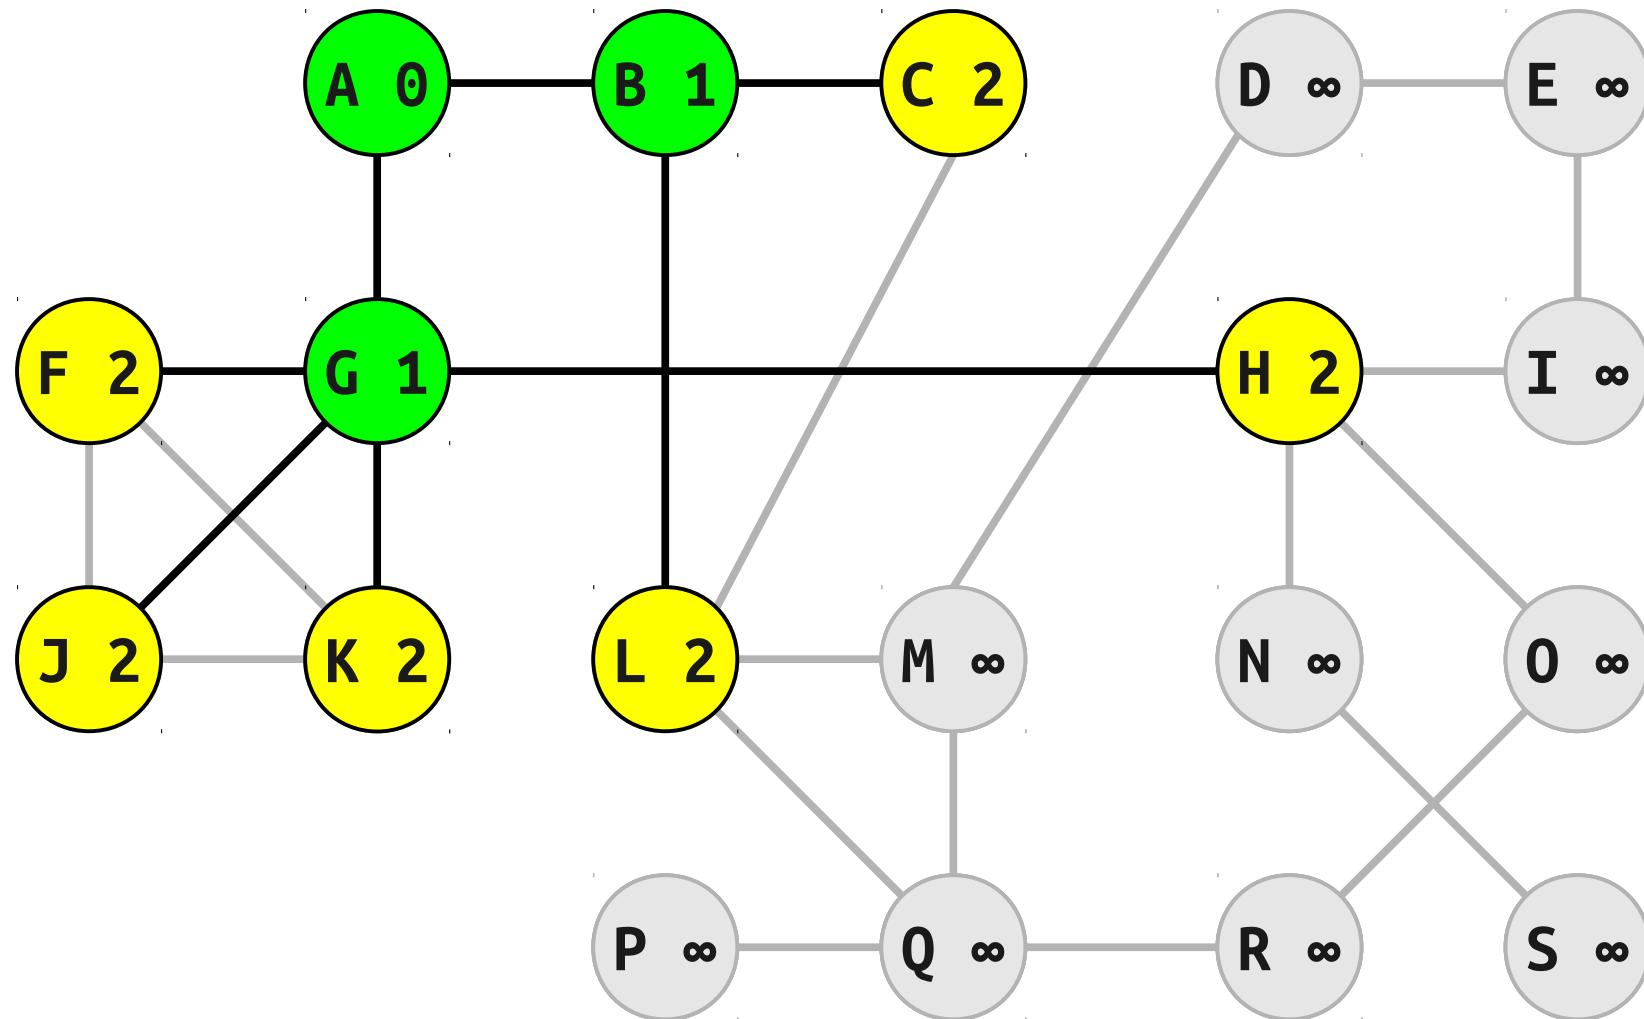
Breadth-First Search



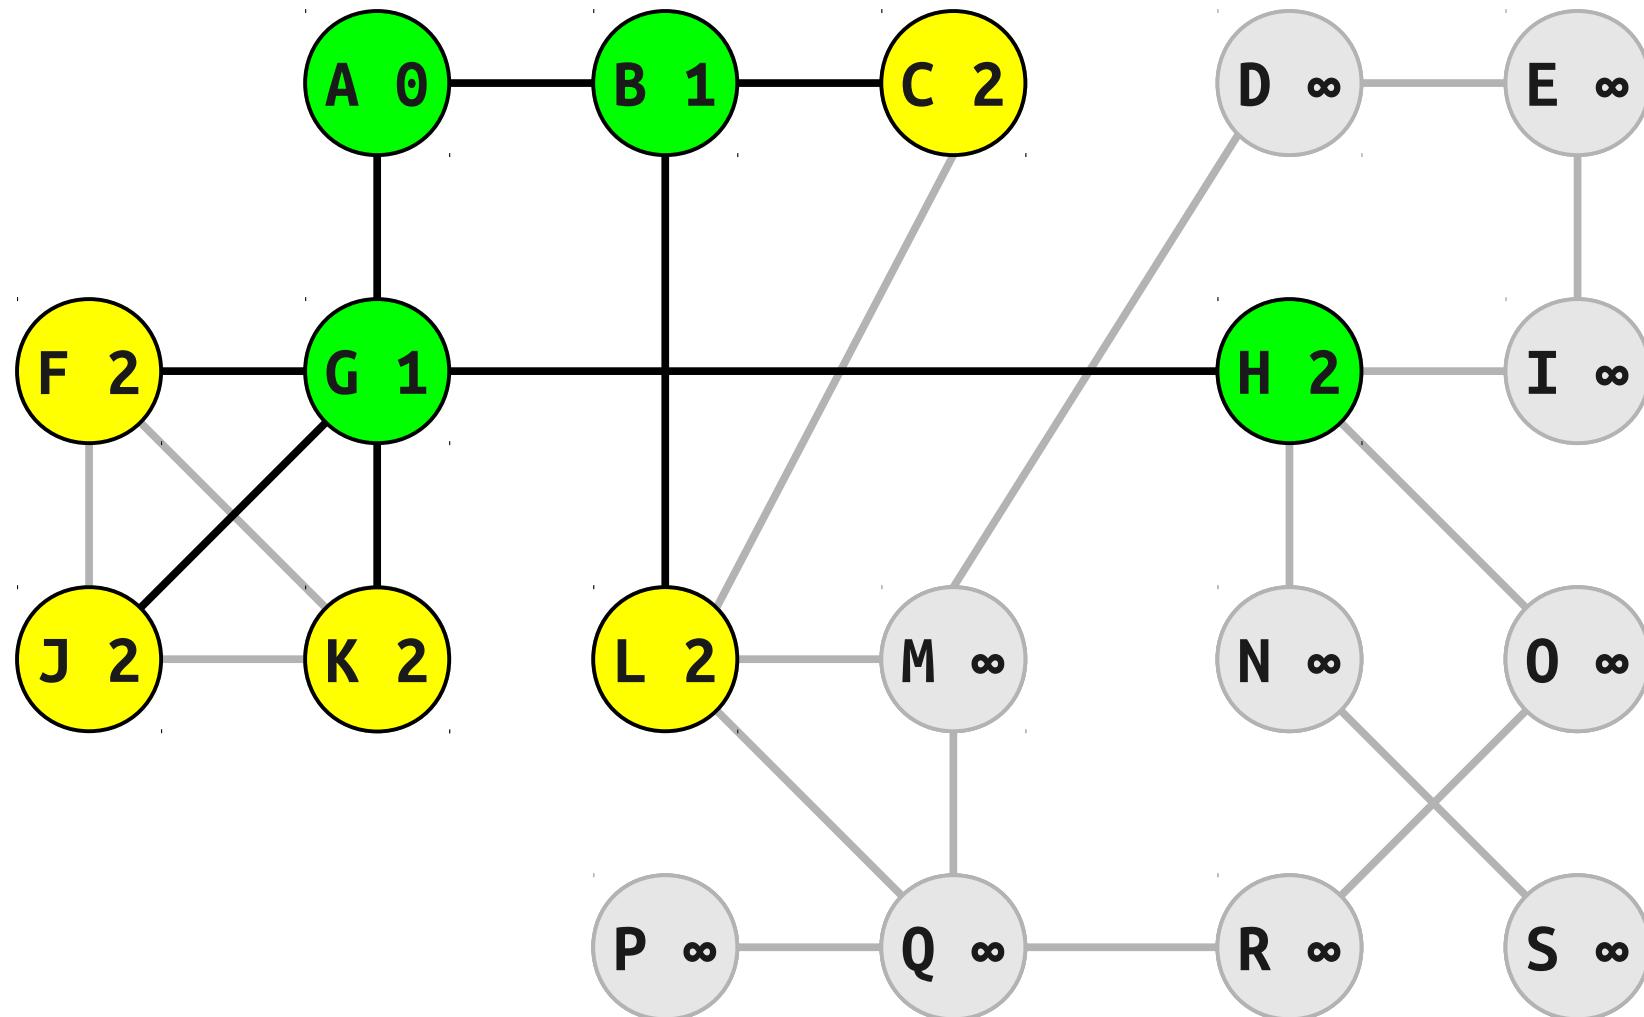
Breadth-First Search



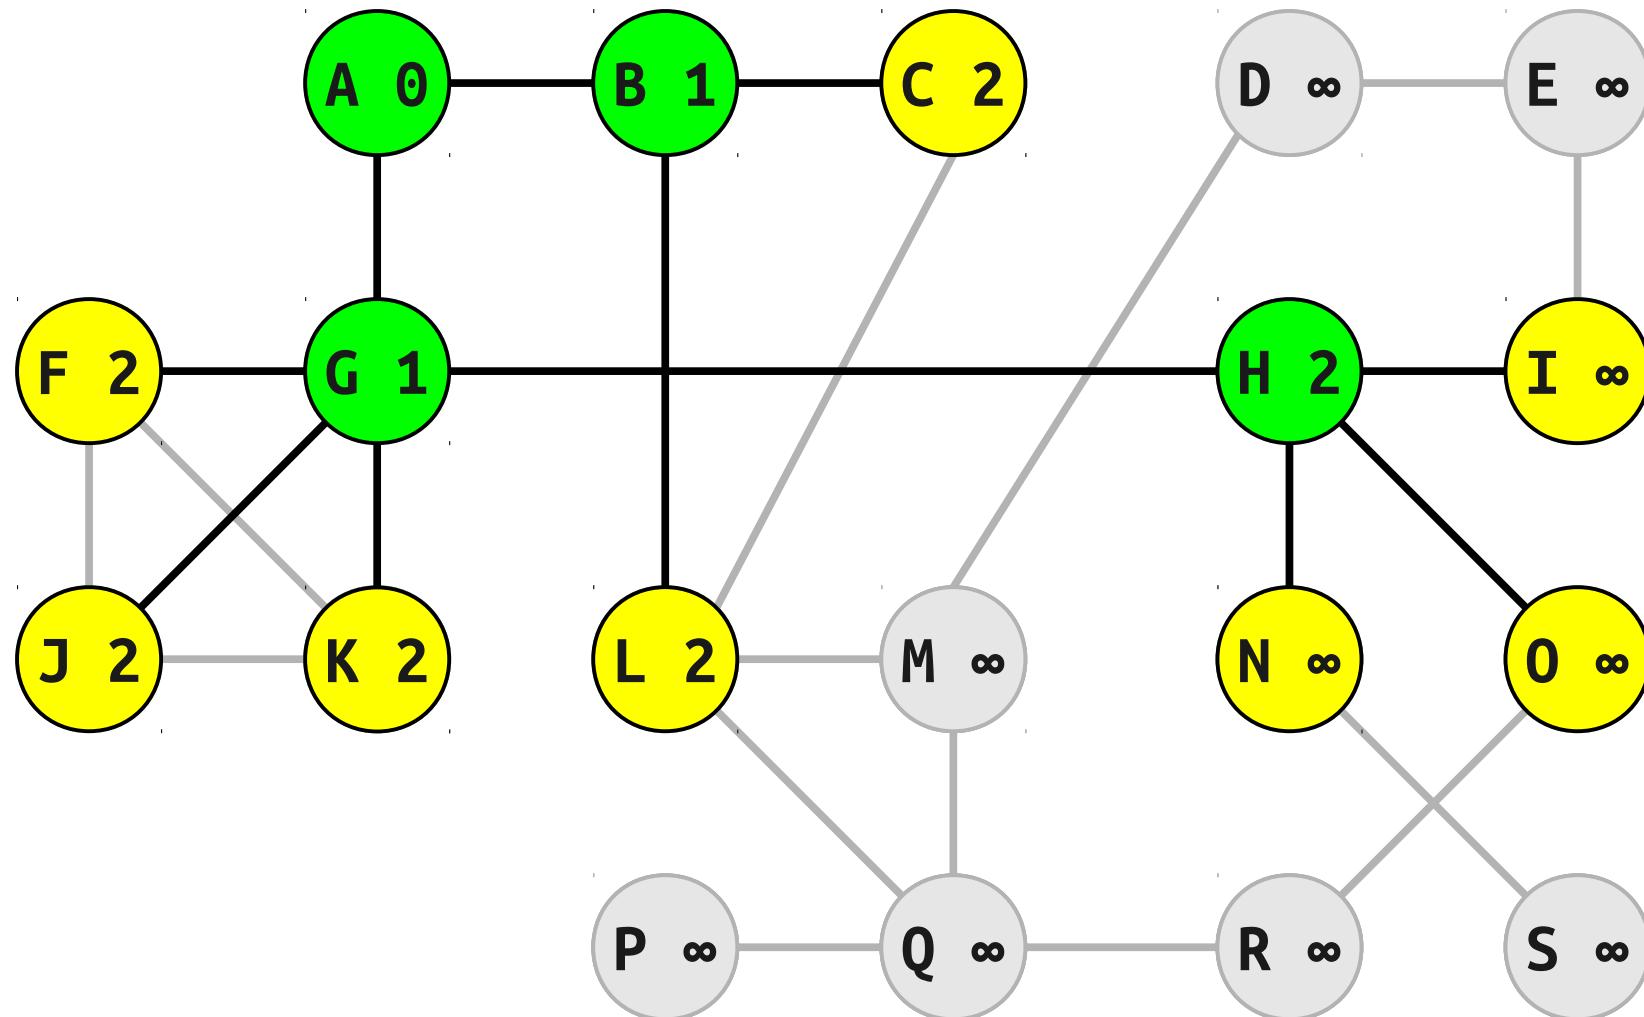
Breadth-First Search



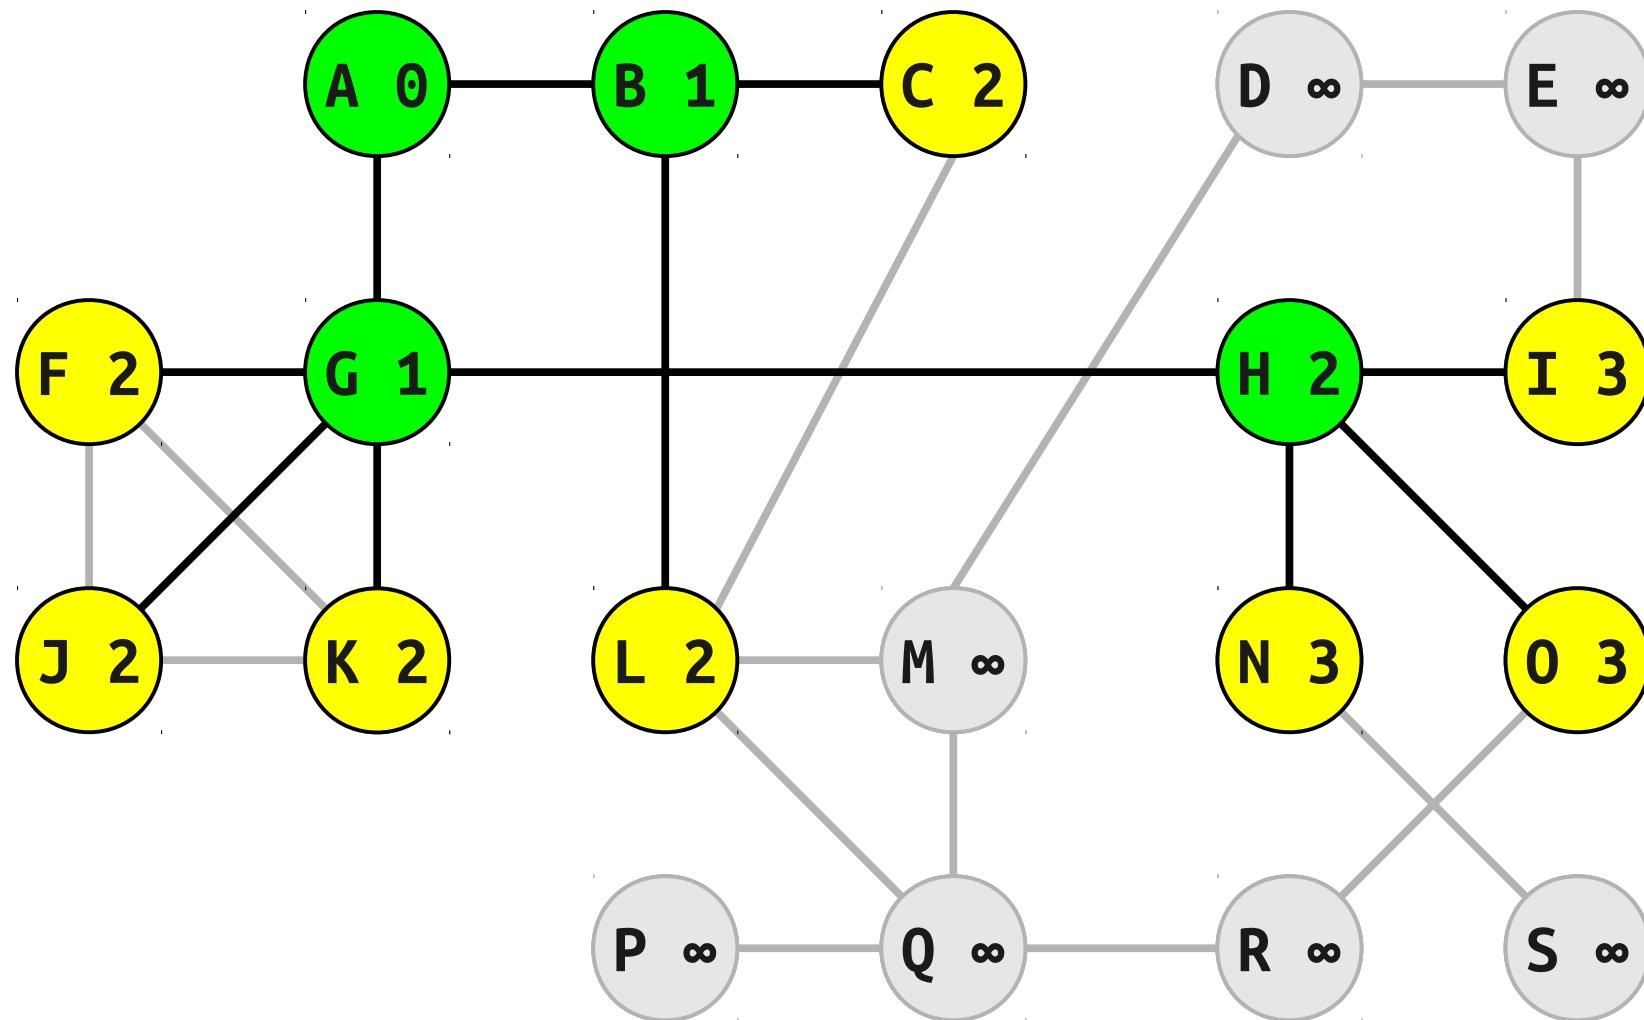
Breadth-First Search



Breadth-First Search



Breadth-First Search



$O(1)$

$O(n)$

$O(1)$

$O(n)$

$O(m + n)$

```
procedure breadthFirstSearch(s, G):
    let q be a new queue.
    for each node v in G:
        dist[v] = ∞

    dist[s] = 0
    enqueue(s, q)

    while q is not empty:
        let v = dequeue(q)
        for each neighbor u of v:
            if dist[u] = ∞:
                dist[u] = dist[v] + 1
                enqueue(u, q)
```

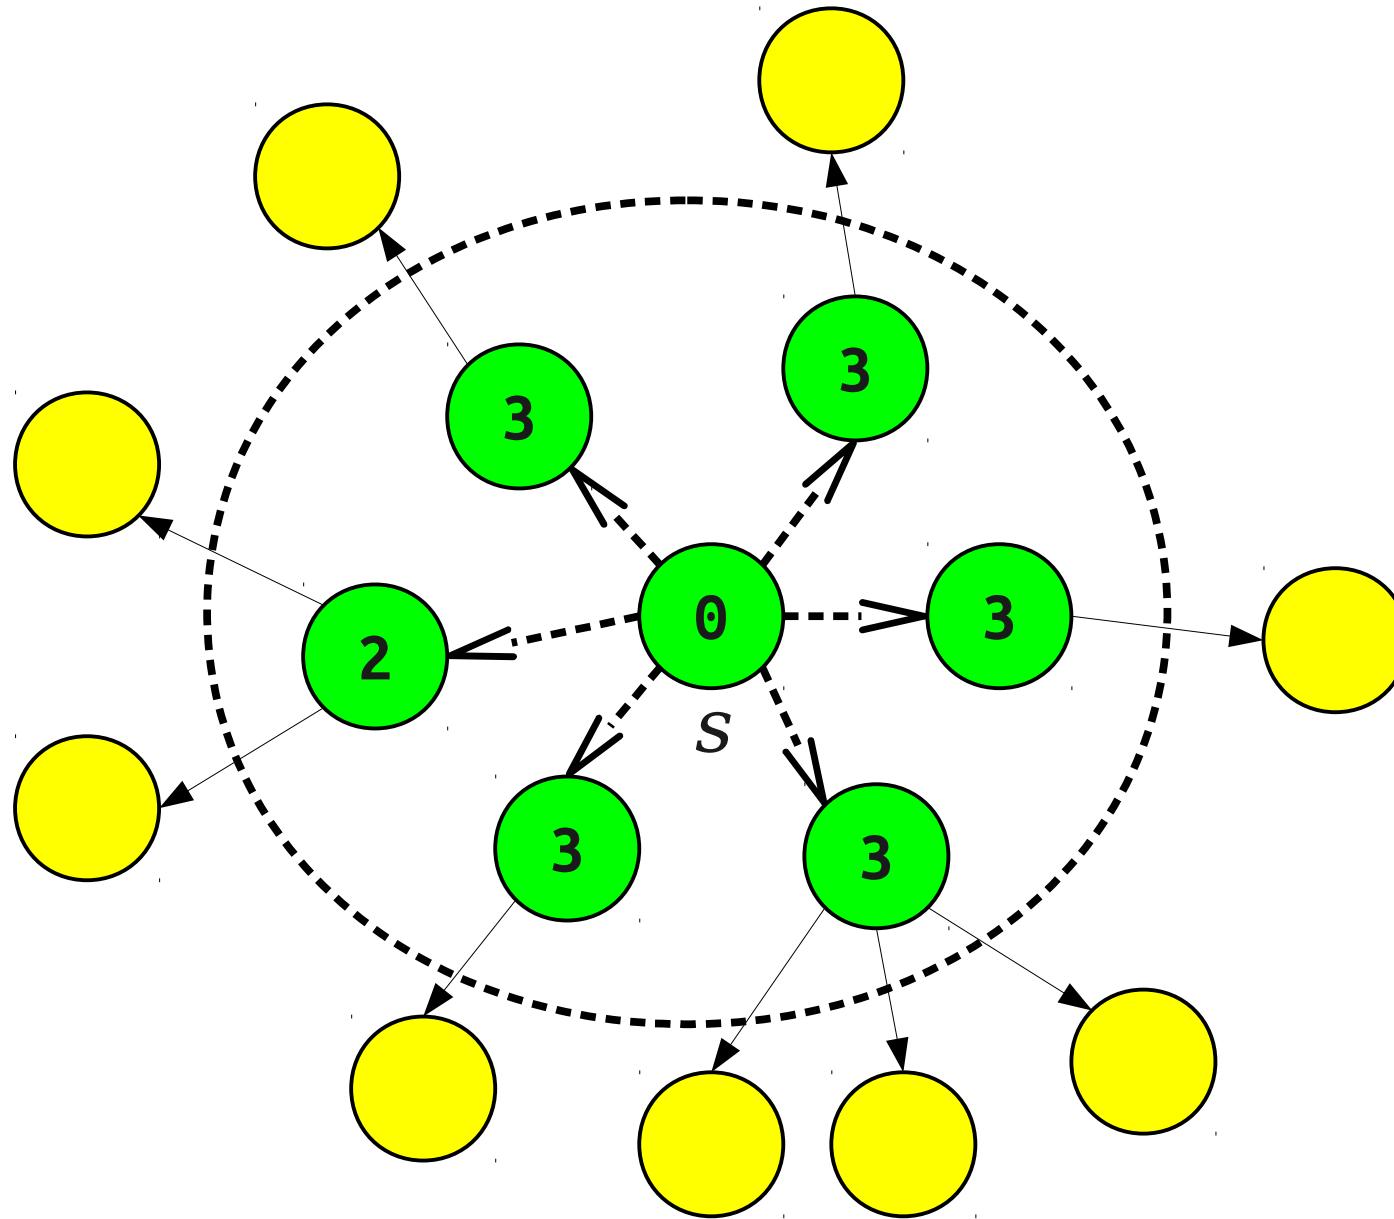
A Second Intuition

- At each point in the execution of BFS, a node v is either
 - green, and we have the shortest path to v ;
 - yellow, and it is connected to some green node; or
 - gray, and v is undiscovered.
- Each iteration, we pick a yellow node with minimal distance from the start node and color it green. So what is the cost of the lowest-cost yellow node?
- If v is yellow, it is connected to a green node u by an edge.
- The cost of getting from s to v is then $d(s, u) + 1$.
- BFS works by picking the yellow node v minimizing

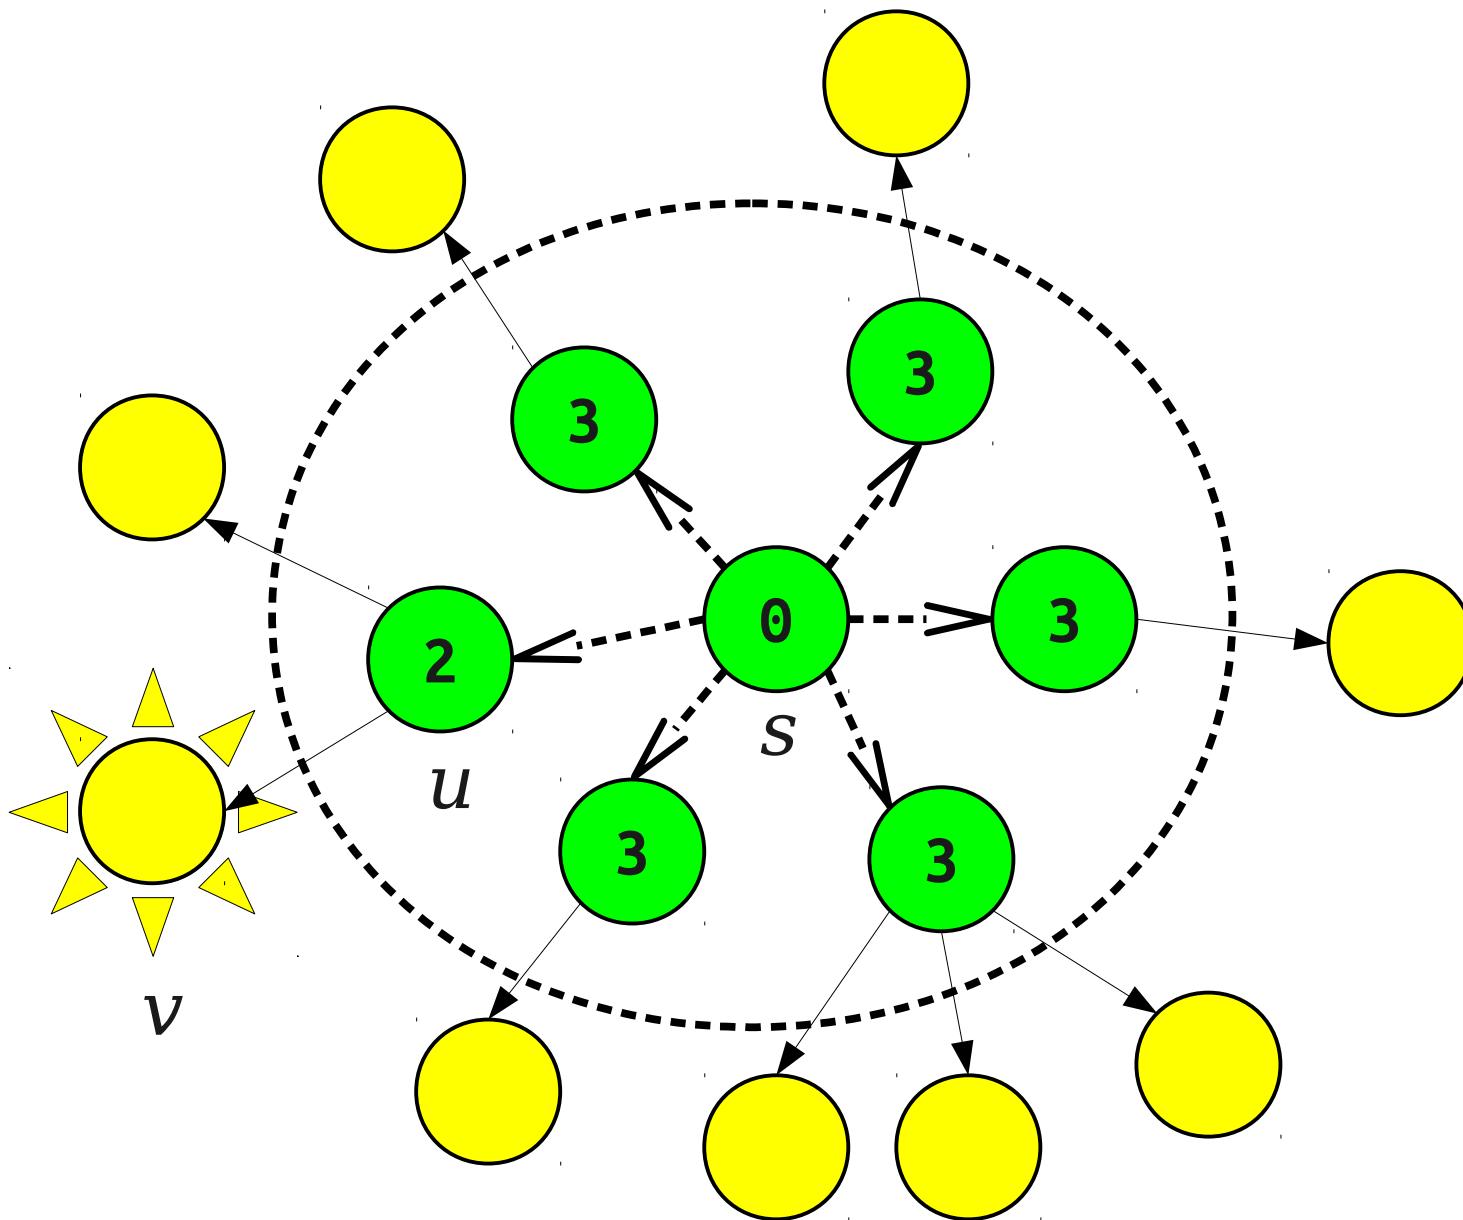
$$\mathbf{d(s, u) + 1}$$

where (u, v) is an edge and u is green.

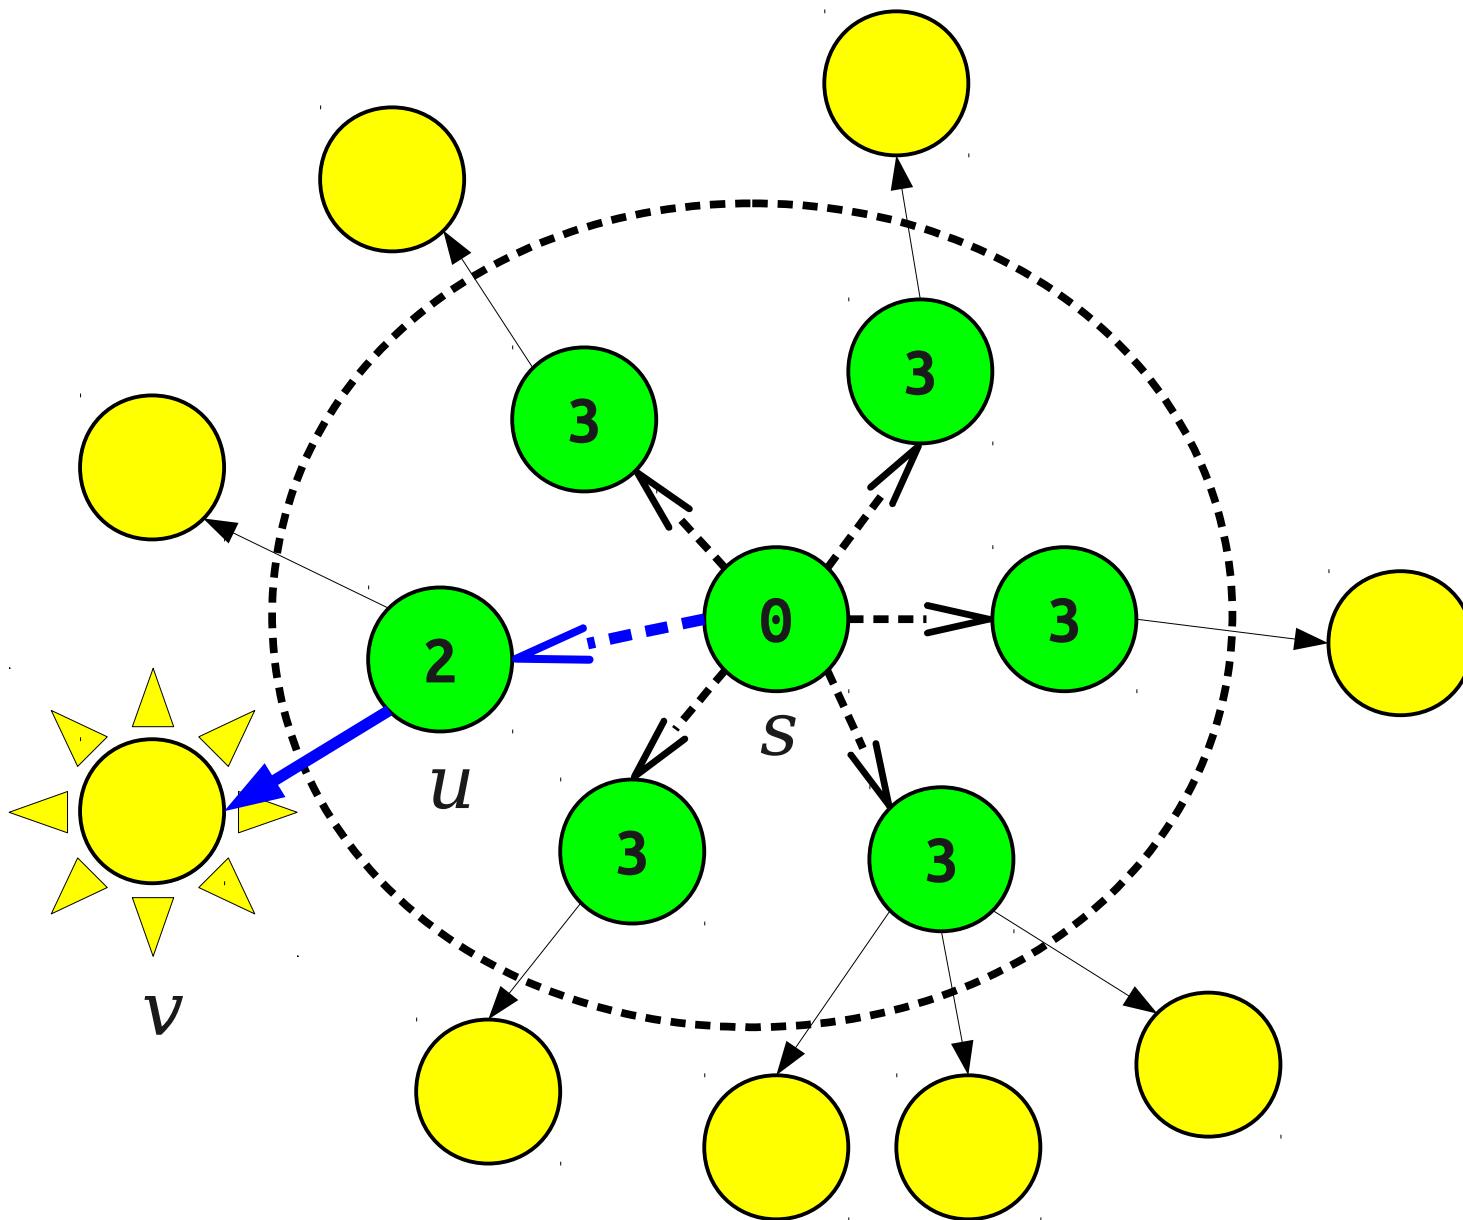
Pick yellow node v minimizing $d(s, u) + 1$,
where (u, v) is an edge and u is green.



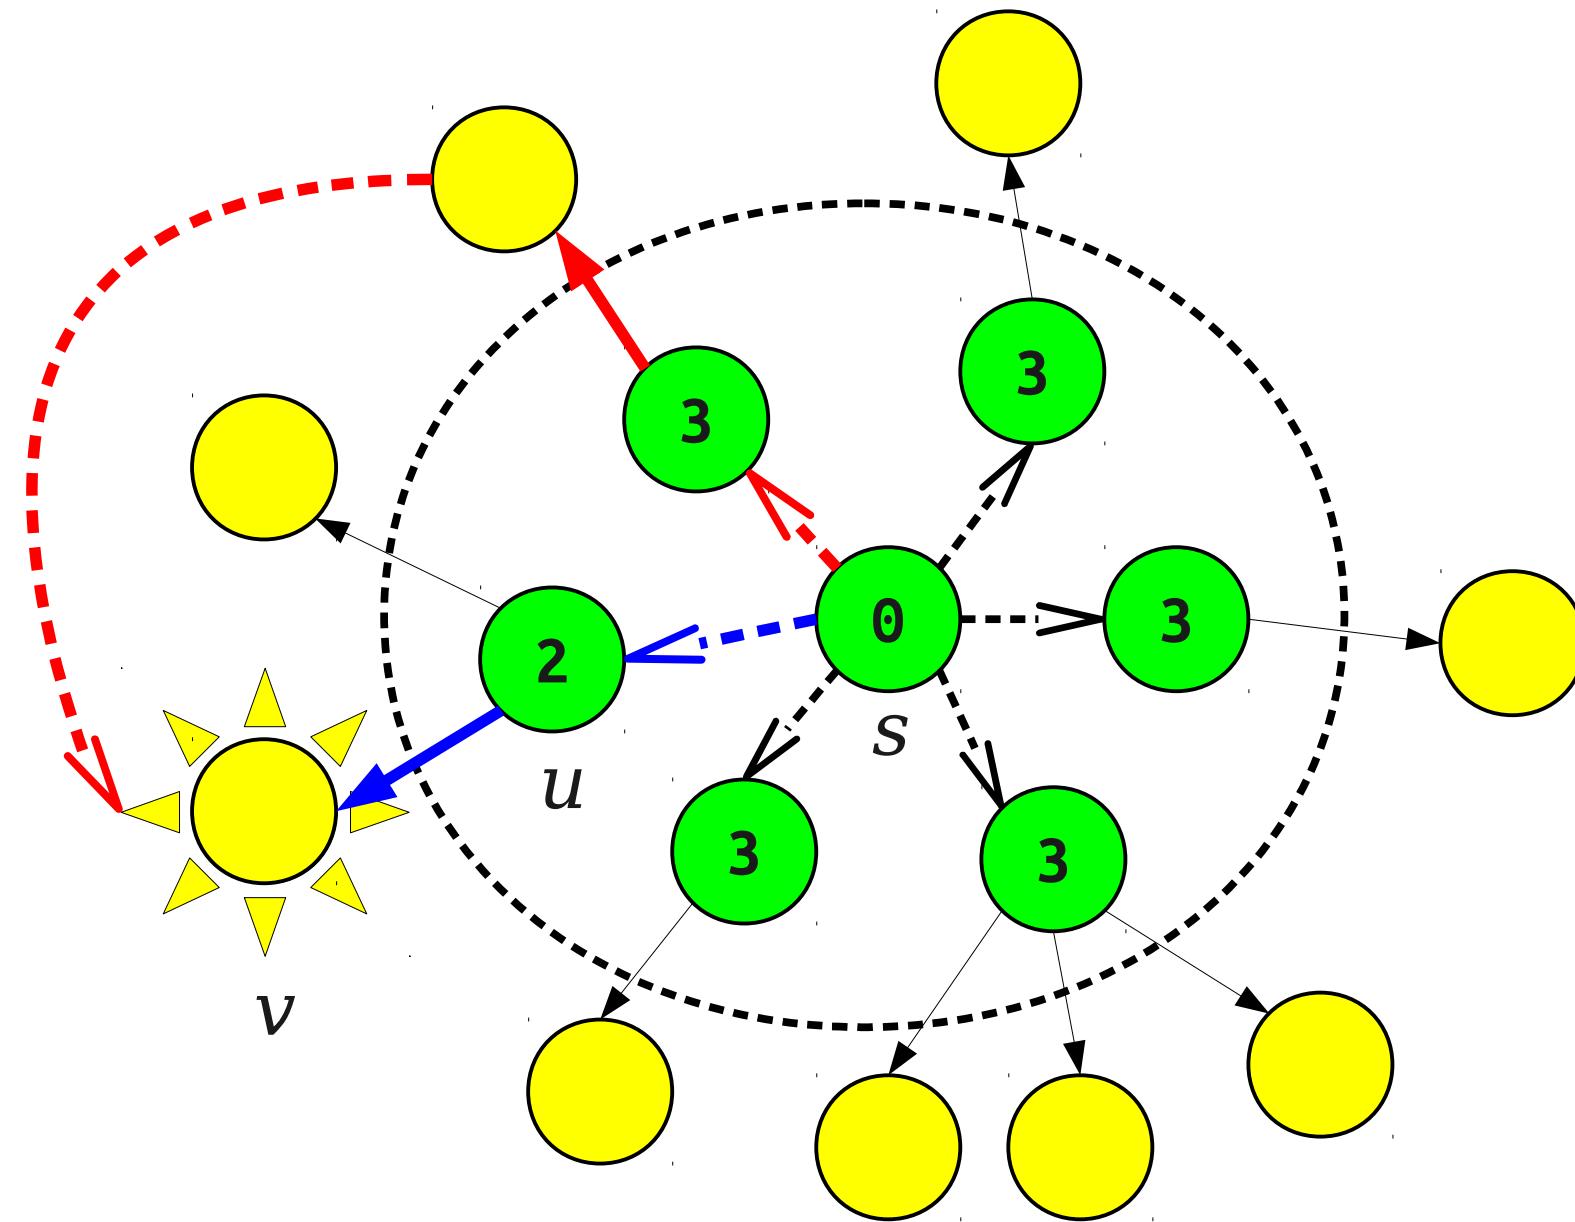
Pick yellow node v minimizing $d(s, u) + 1$,
where (u, v) is an edge and u is green.



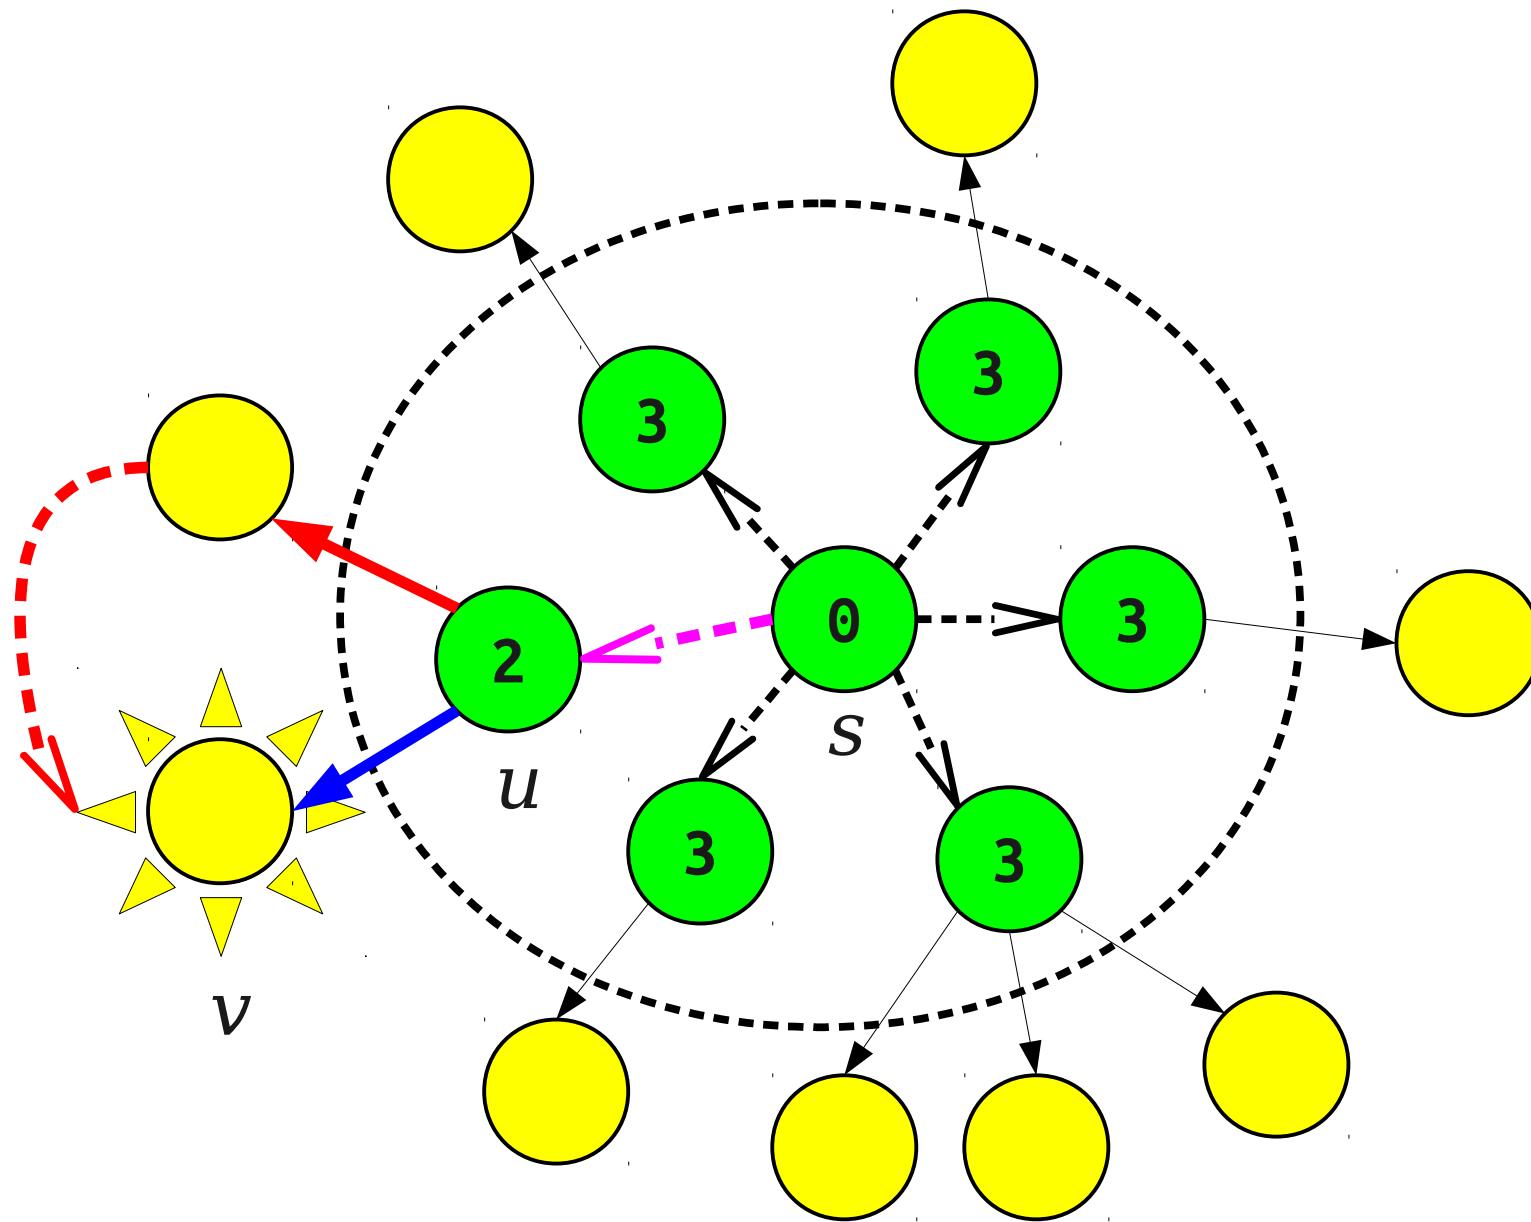
Pick yellow node v minimizing $d(s, u) + 1$,
where (u, v) is an edge and u is green.



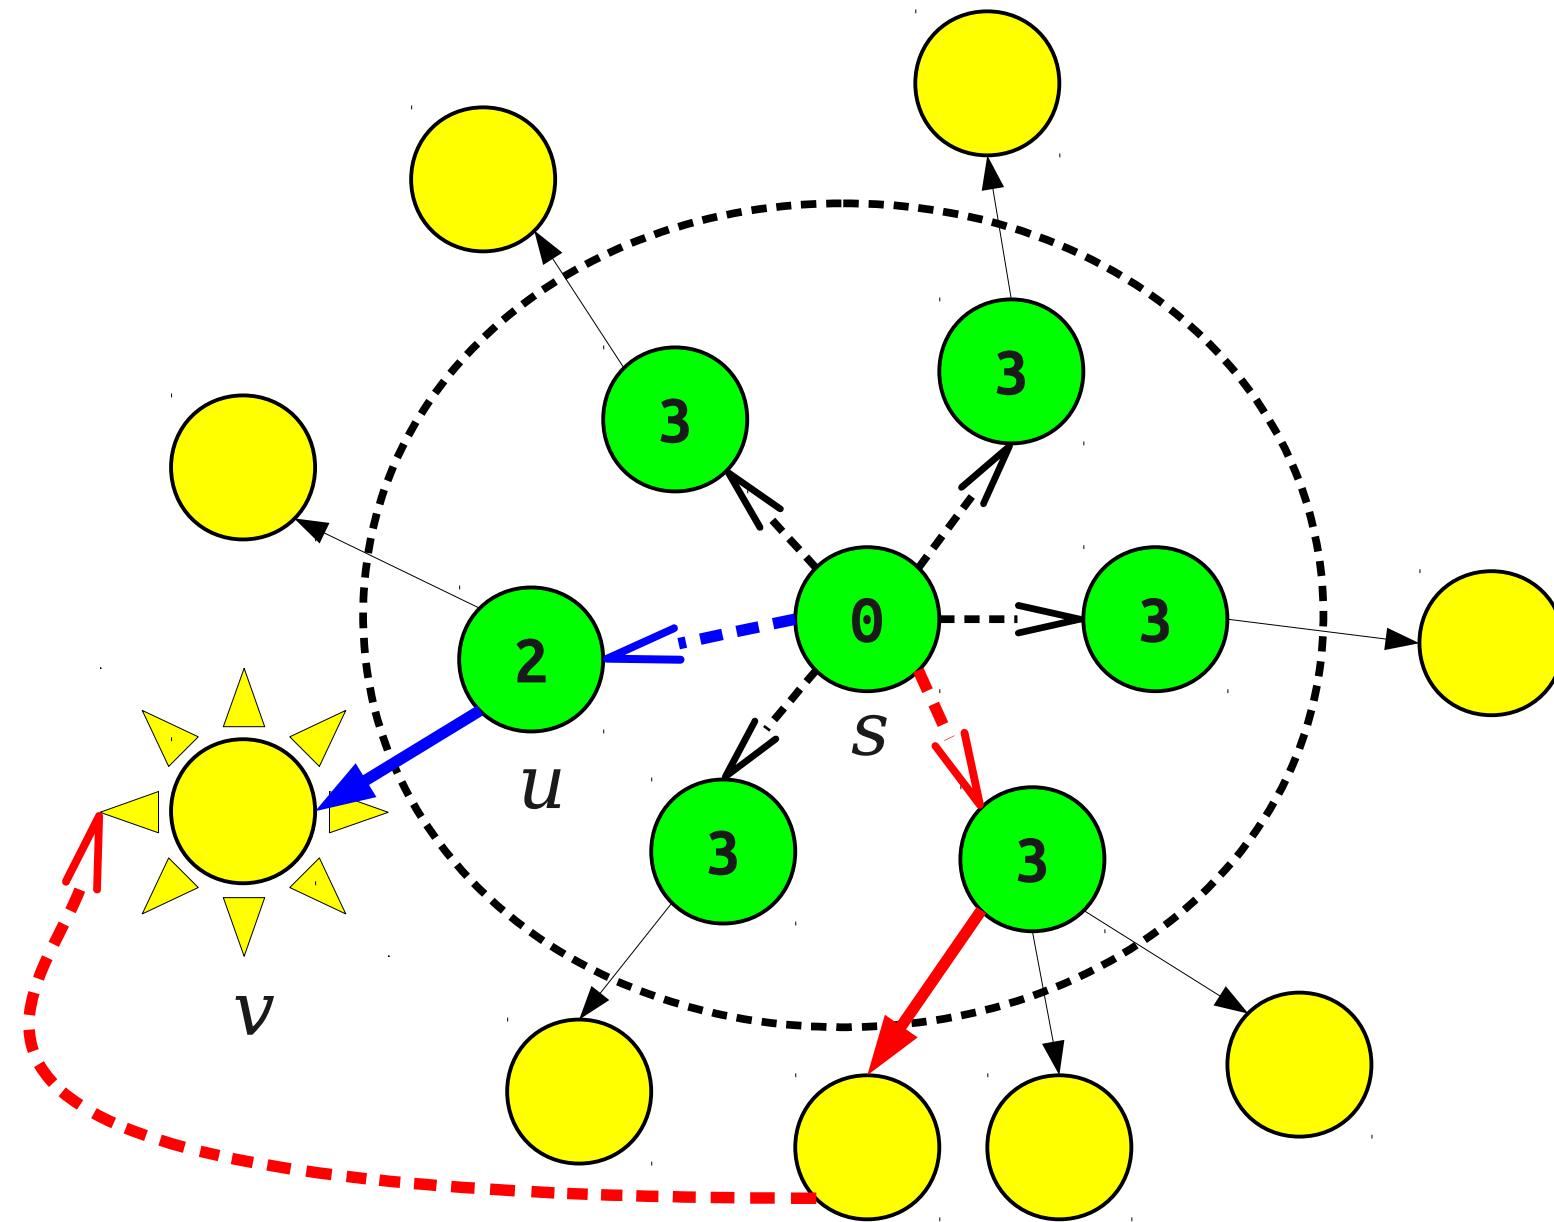
Pick yellow node v minimizing $d(s, u) + 1$,
where (u, v) is an edge and u is green.



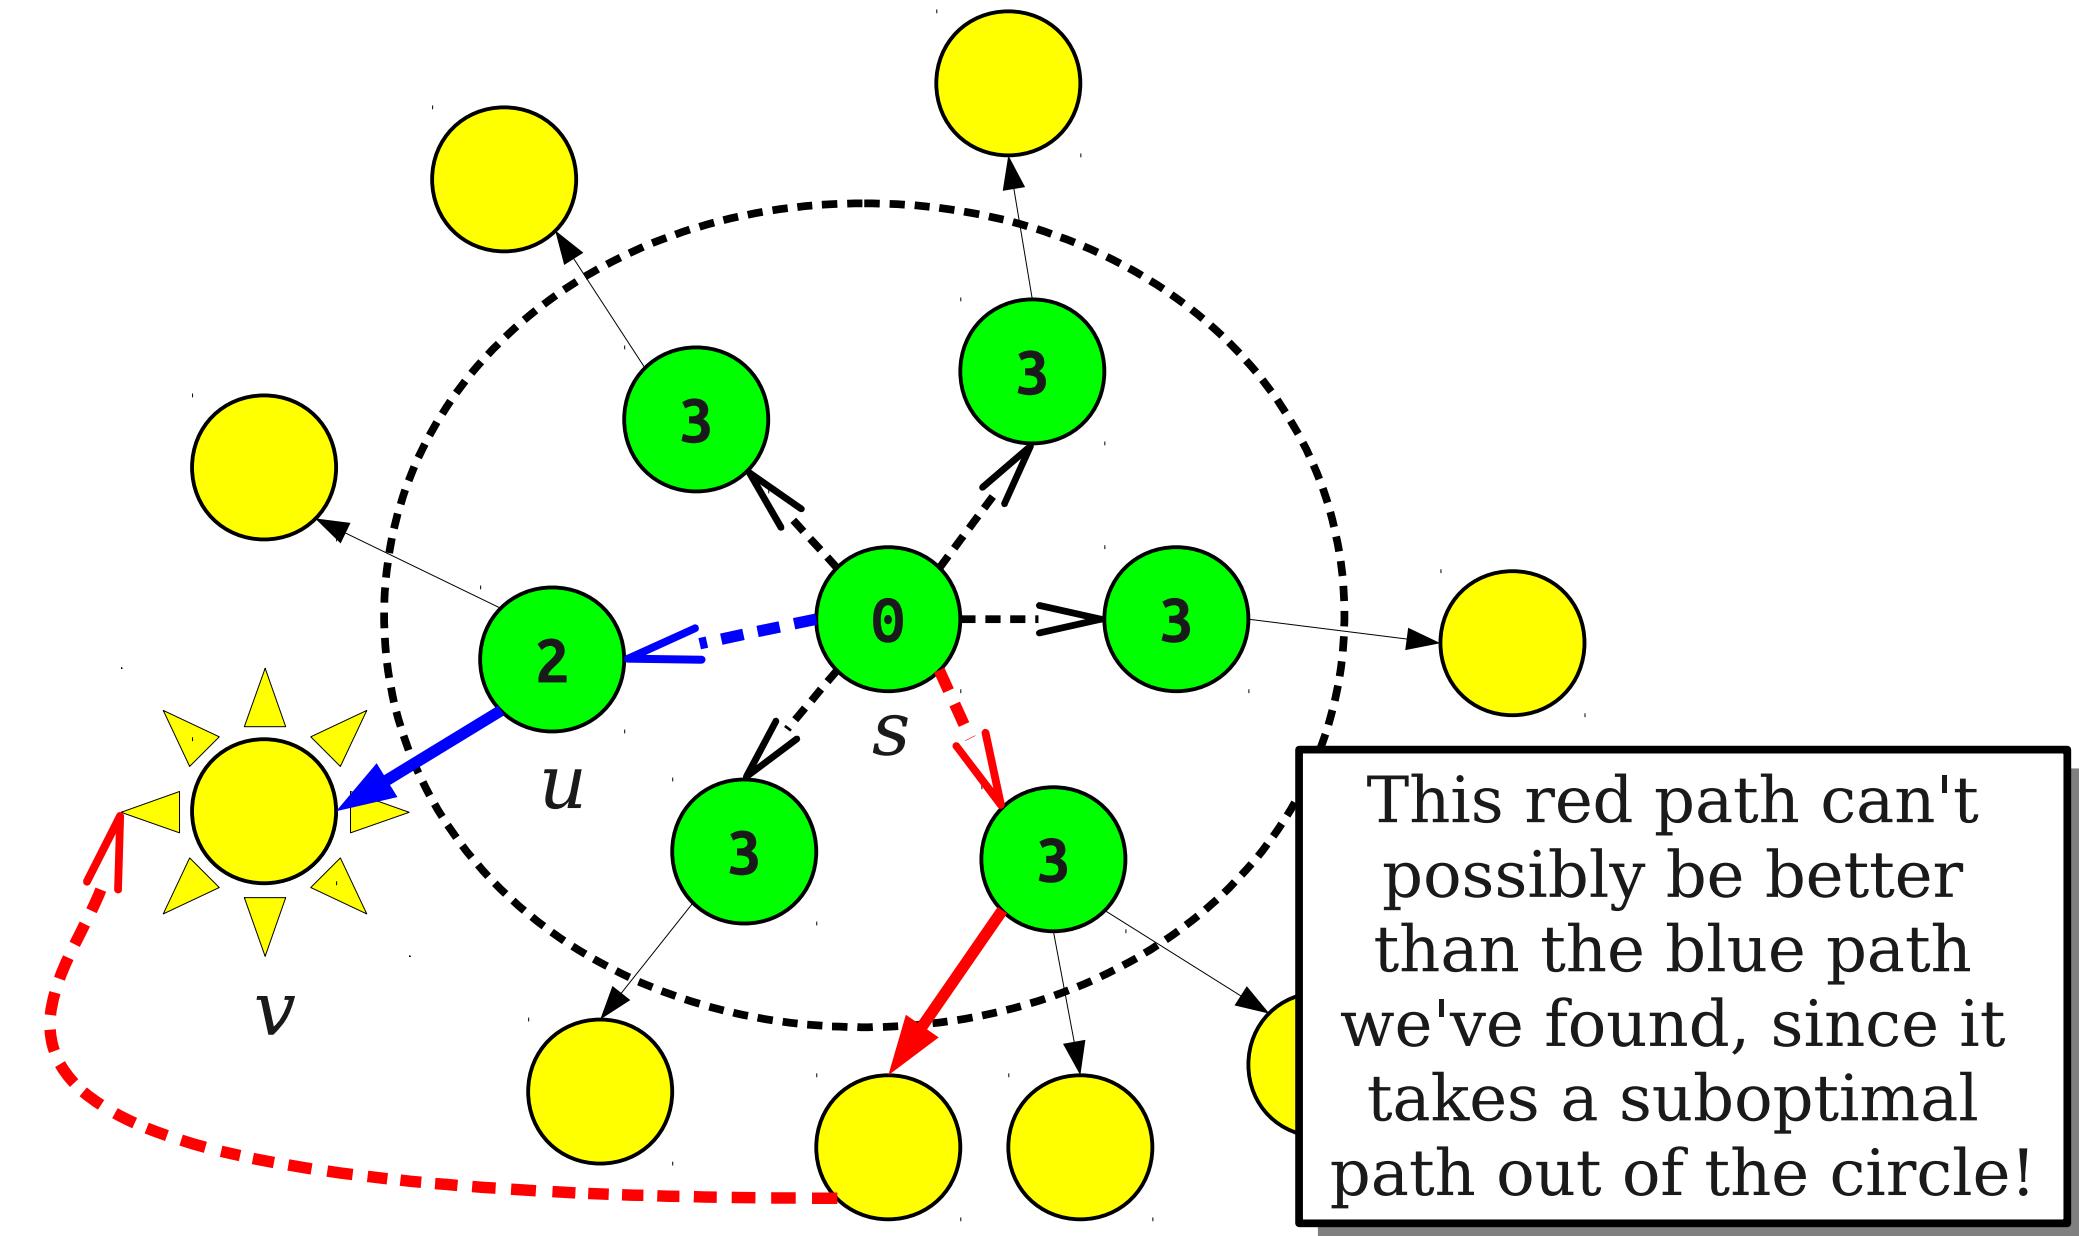
Pick yellow node v minimizing $d(s, u) + 1$,
where (u, v) is an edge and u is green.



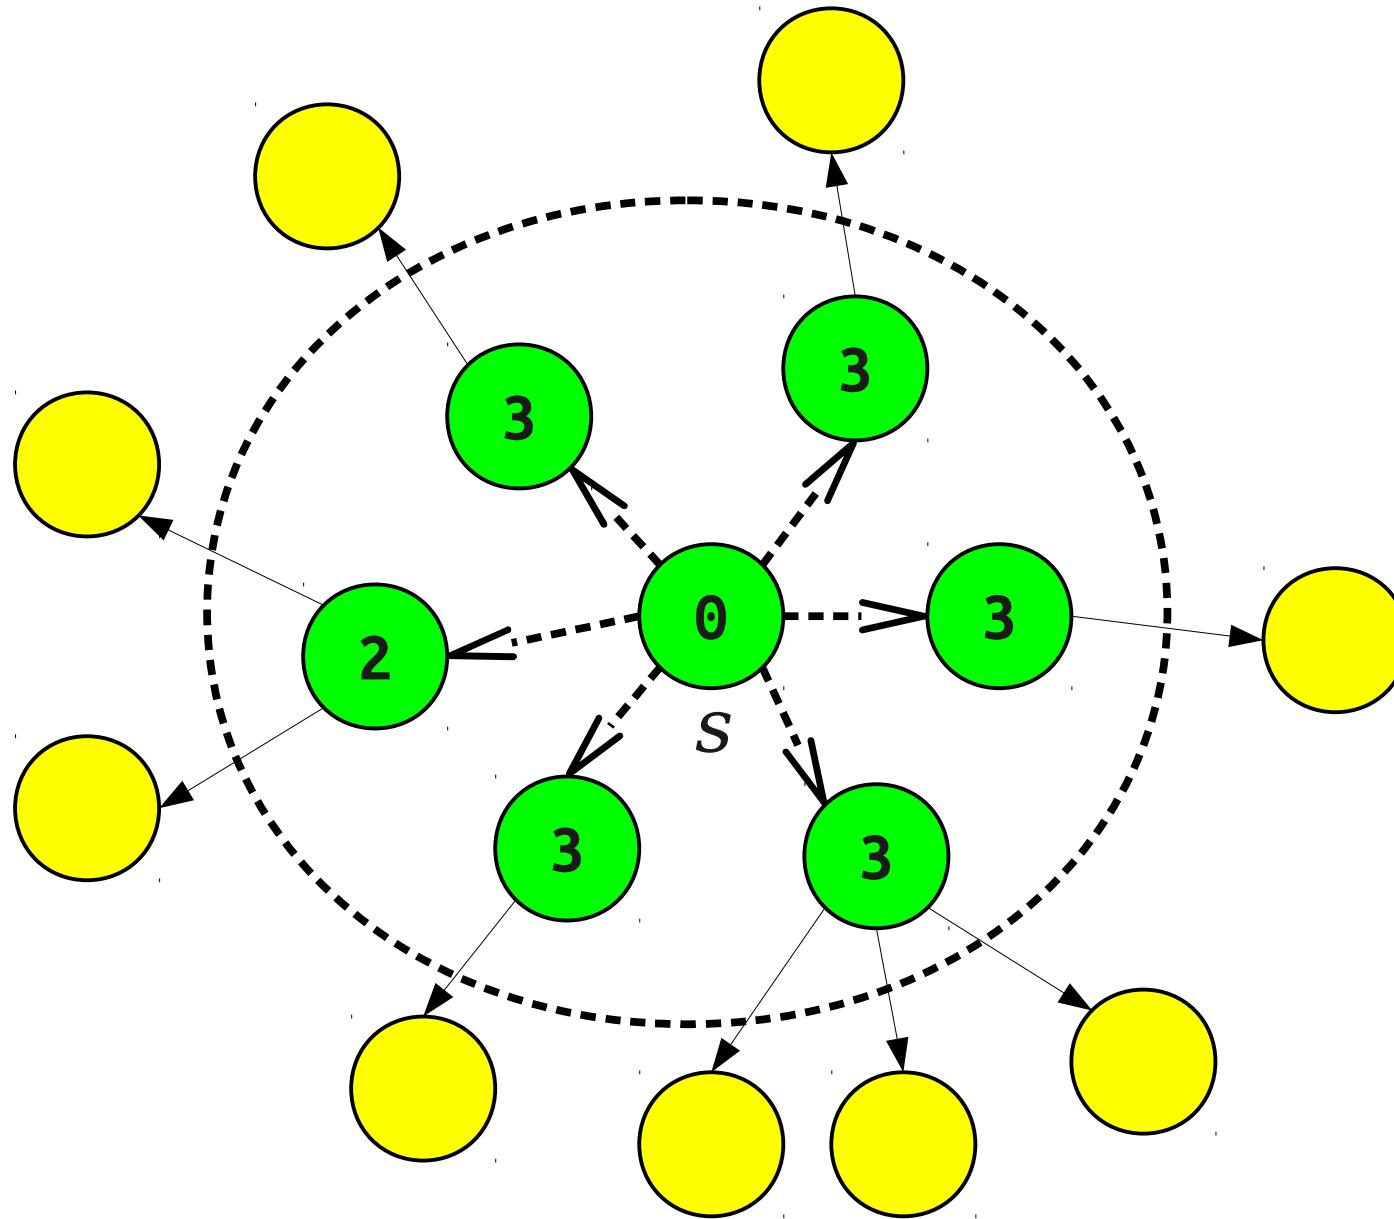
Pick yellow node v minimizing $d(s, u) + 1$,
where (u, v) is an edge and u is green.



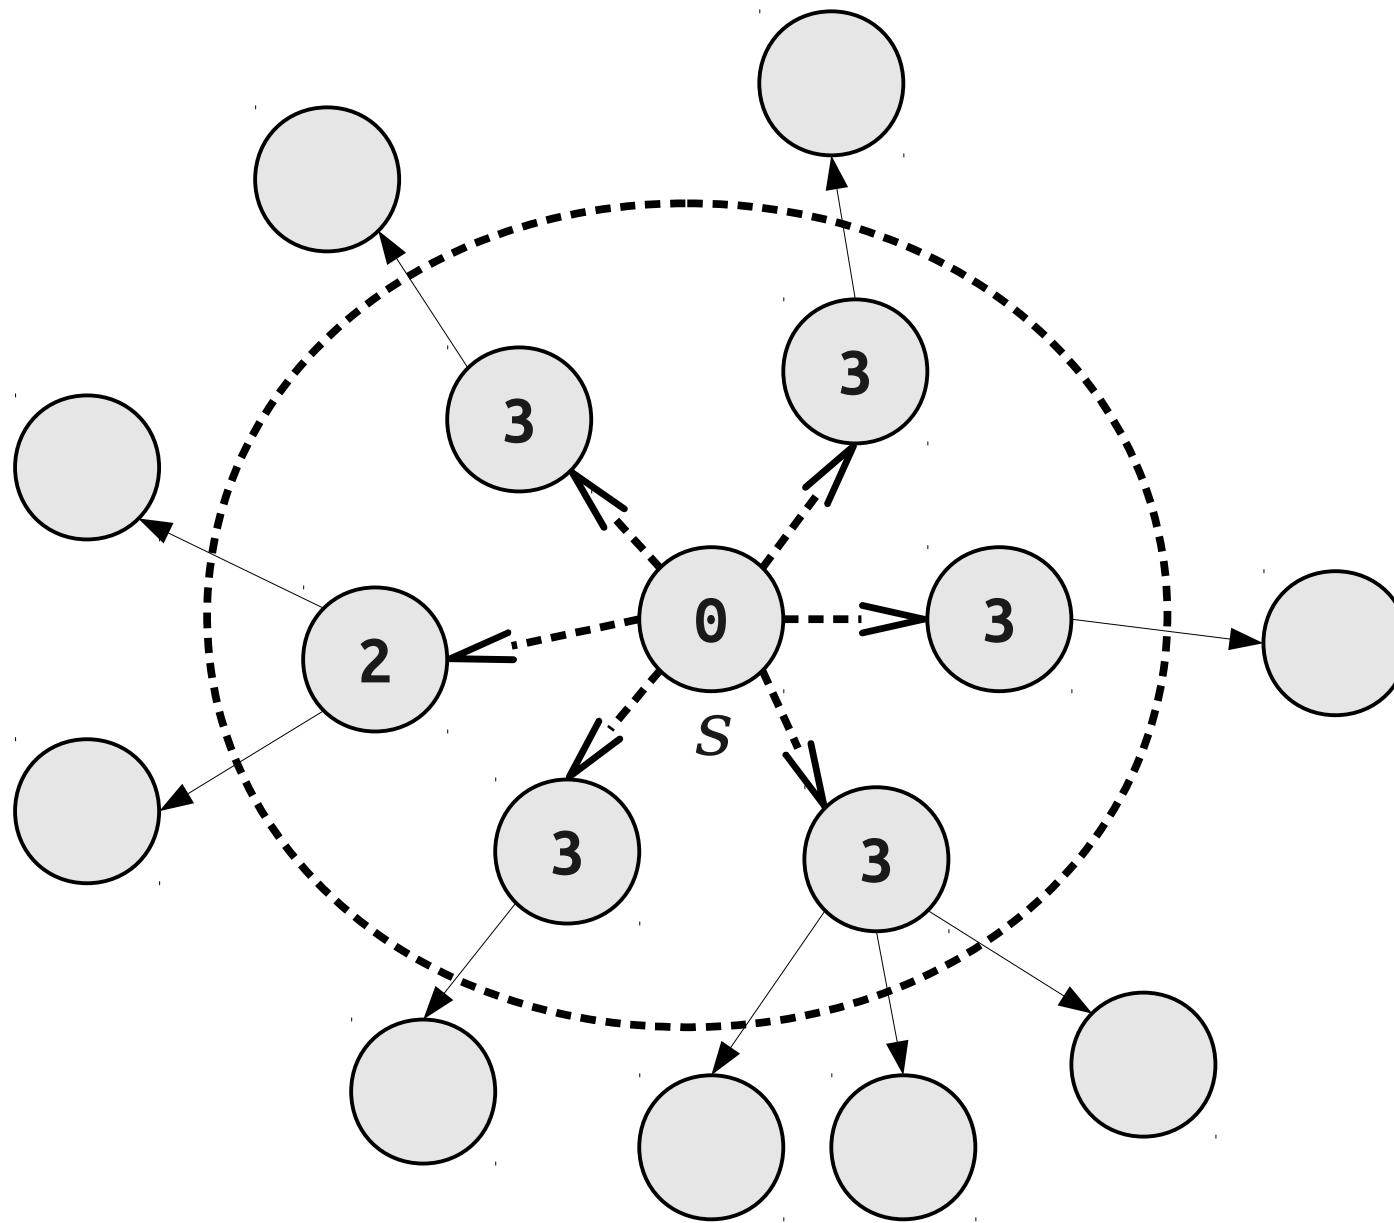
Pick yellow node v minimizing $d(s, u) + 1$,
where (u, v) is an edge and u is green.



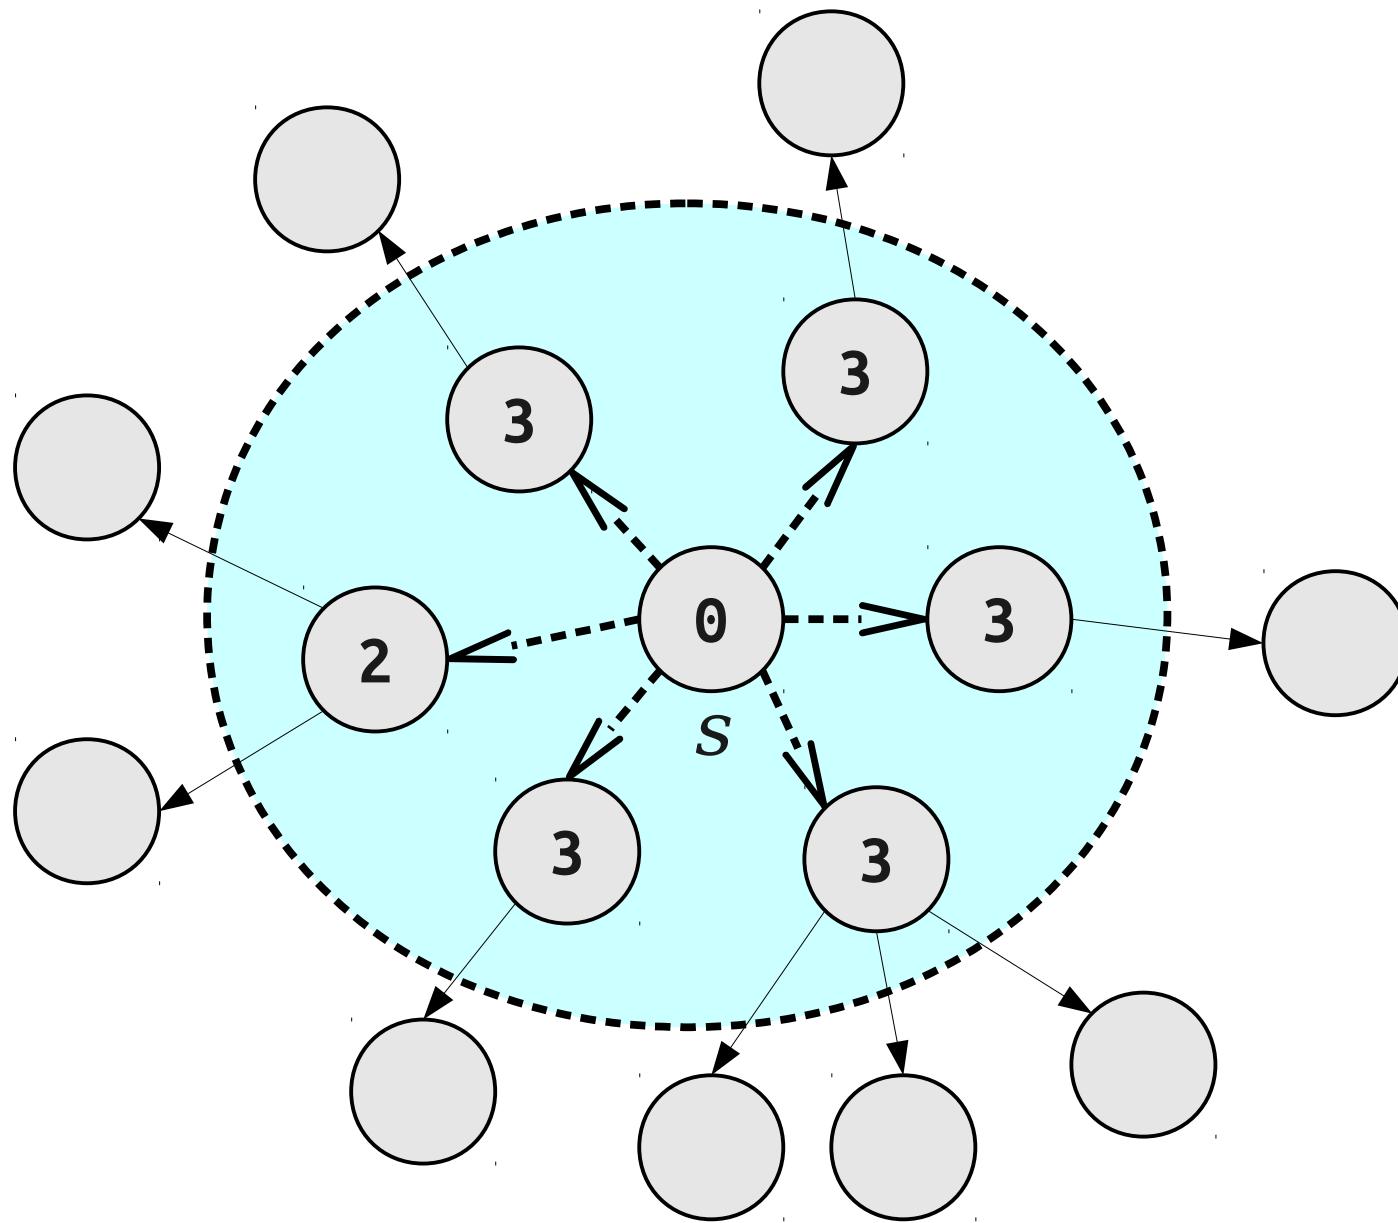
Pick yellow node v minimizing $d(s, u) + 1$,
where (u, v) is an edge and u is green.



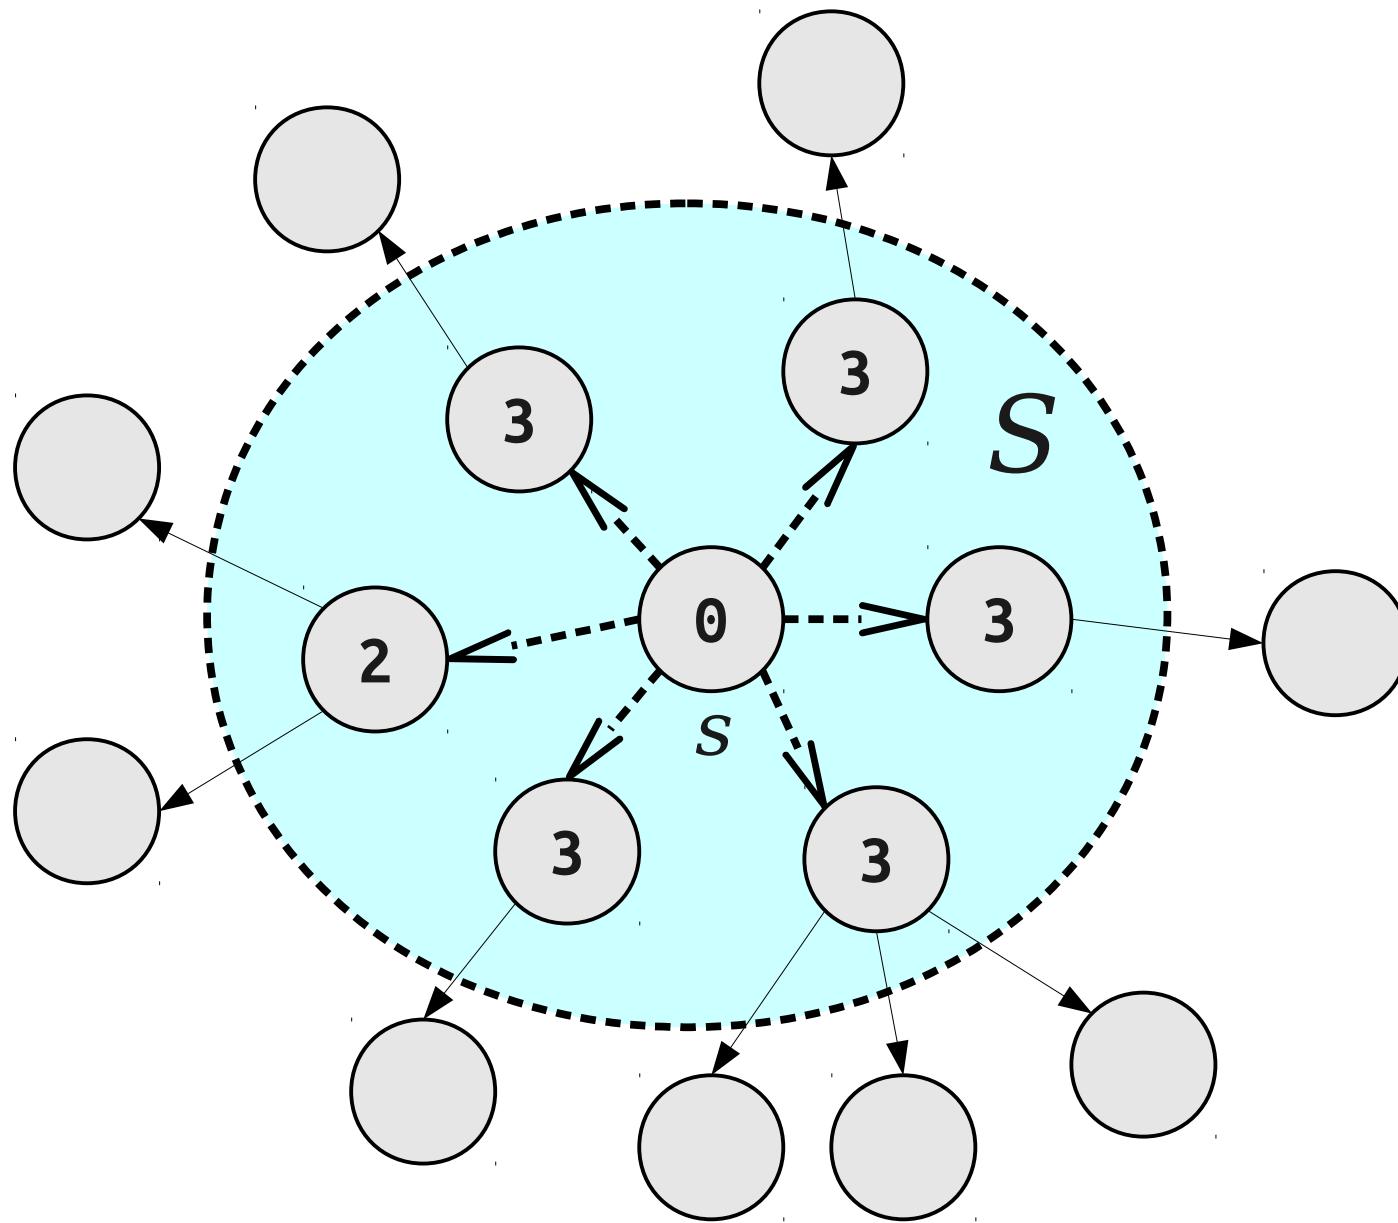
Pick yellow node v minimizing $d(s, u) + 1$,
where (u, v) is an edge and u is green.



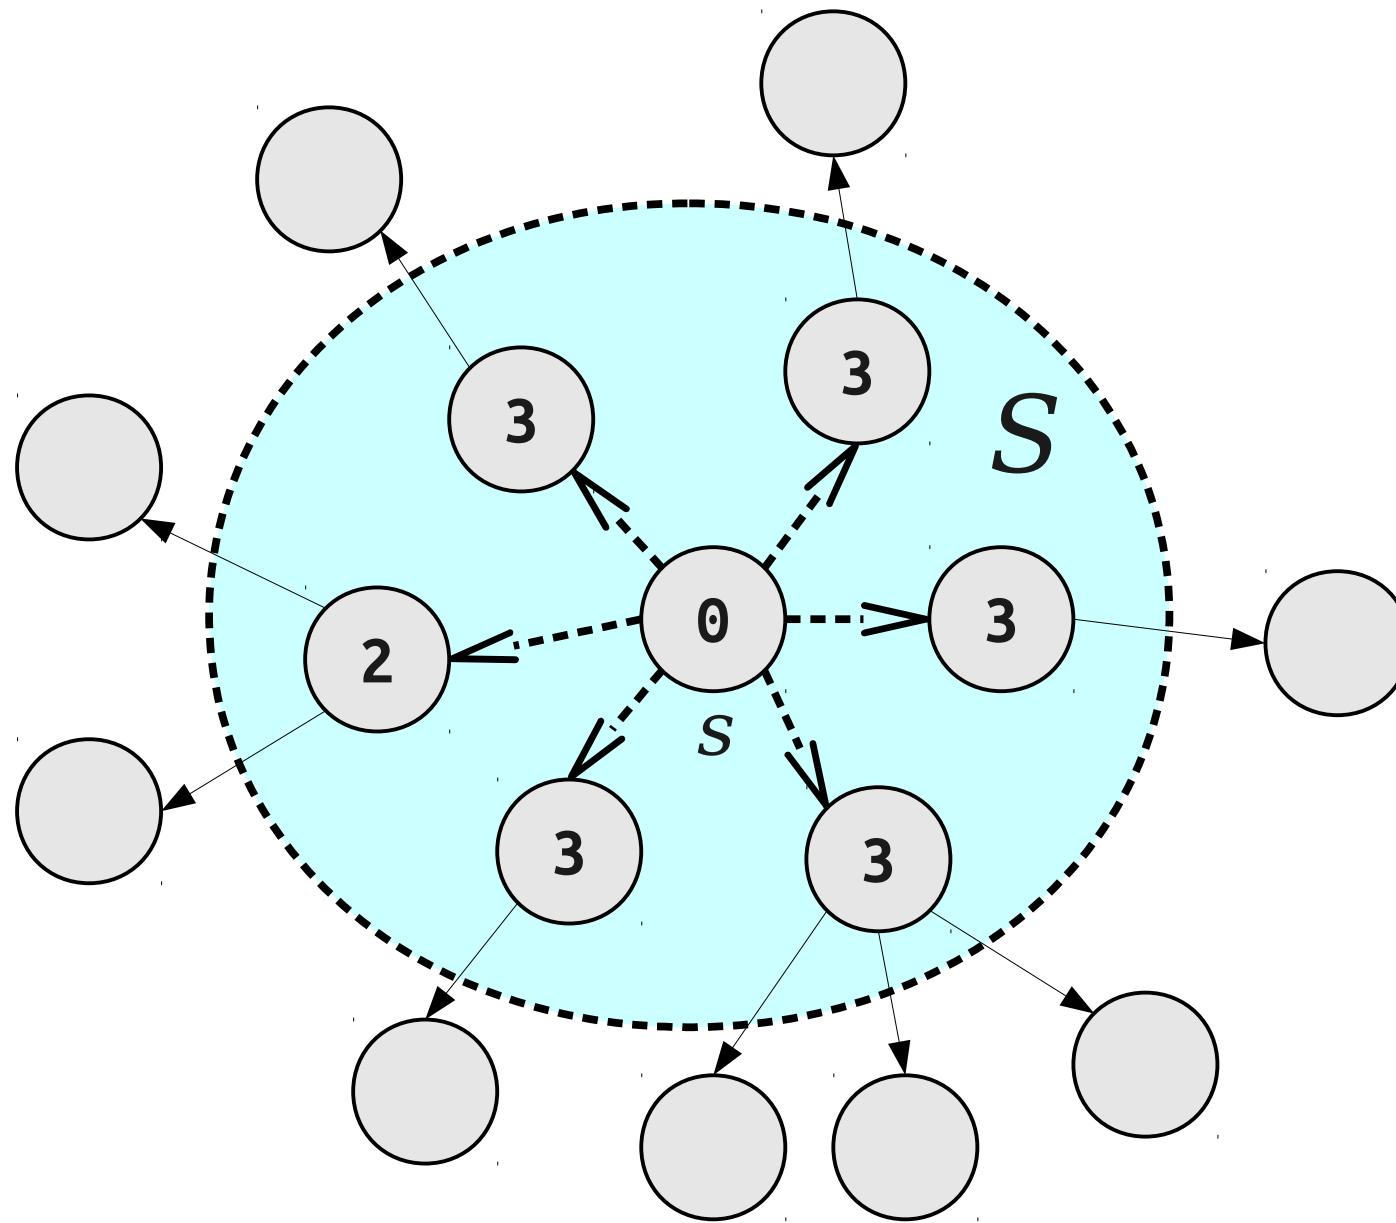
Pick yellow node v minimizing $d(s, u) + 1$,
where (u, v) is an edge and u is green.



Pick yellow node v minimizing $d(s, u) + 1$,
where (u, v) is an edge and u is green.



Pick node $v \notin S$ minimizing $d(s, u) + 1$,
where (u, v) is an edge and $u \in S$



Lemma: Suppose we have shortest paths computed for nodes $S \subseteq V$, where $s \in S$. Consider a node v where $(u, v) \in E$, $u \in S$, and the quantity $d(s, u) + 1$ is minimized. Then $d(s, v) = d(s, u) + 1$.

Proof: There is a path to v of cost $d(s, u) + 1$: follow the shortest path to u (which has cost $d(s, u)$), then follow one more edge to v for total cost $d(s, u) + 1$.

Now suppose for the sake of contradiction that there is a shorter path P to v . This path must start in S (since $s \in S$) and leave S (since $v \notin S$). So consider when P leaves S . When this happens, P must go from s to some node $x \in S$, cross an edge (x, y) to some node y , then continue from y to v . This means that $|P|$ is at least $d(s, x) + 1$, since the path goes from s to x and then follows at least one more edge.

Since v was picked to minimize $d(s, u) + 1$ for any choice of $u \in S$ adjacent to an edge (u, v) , we know

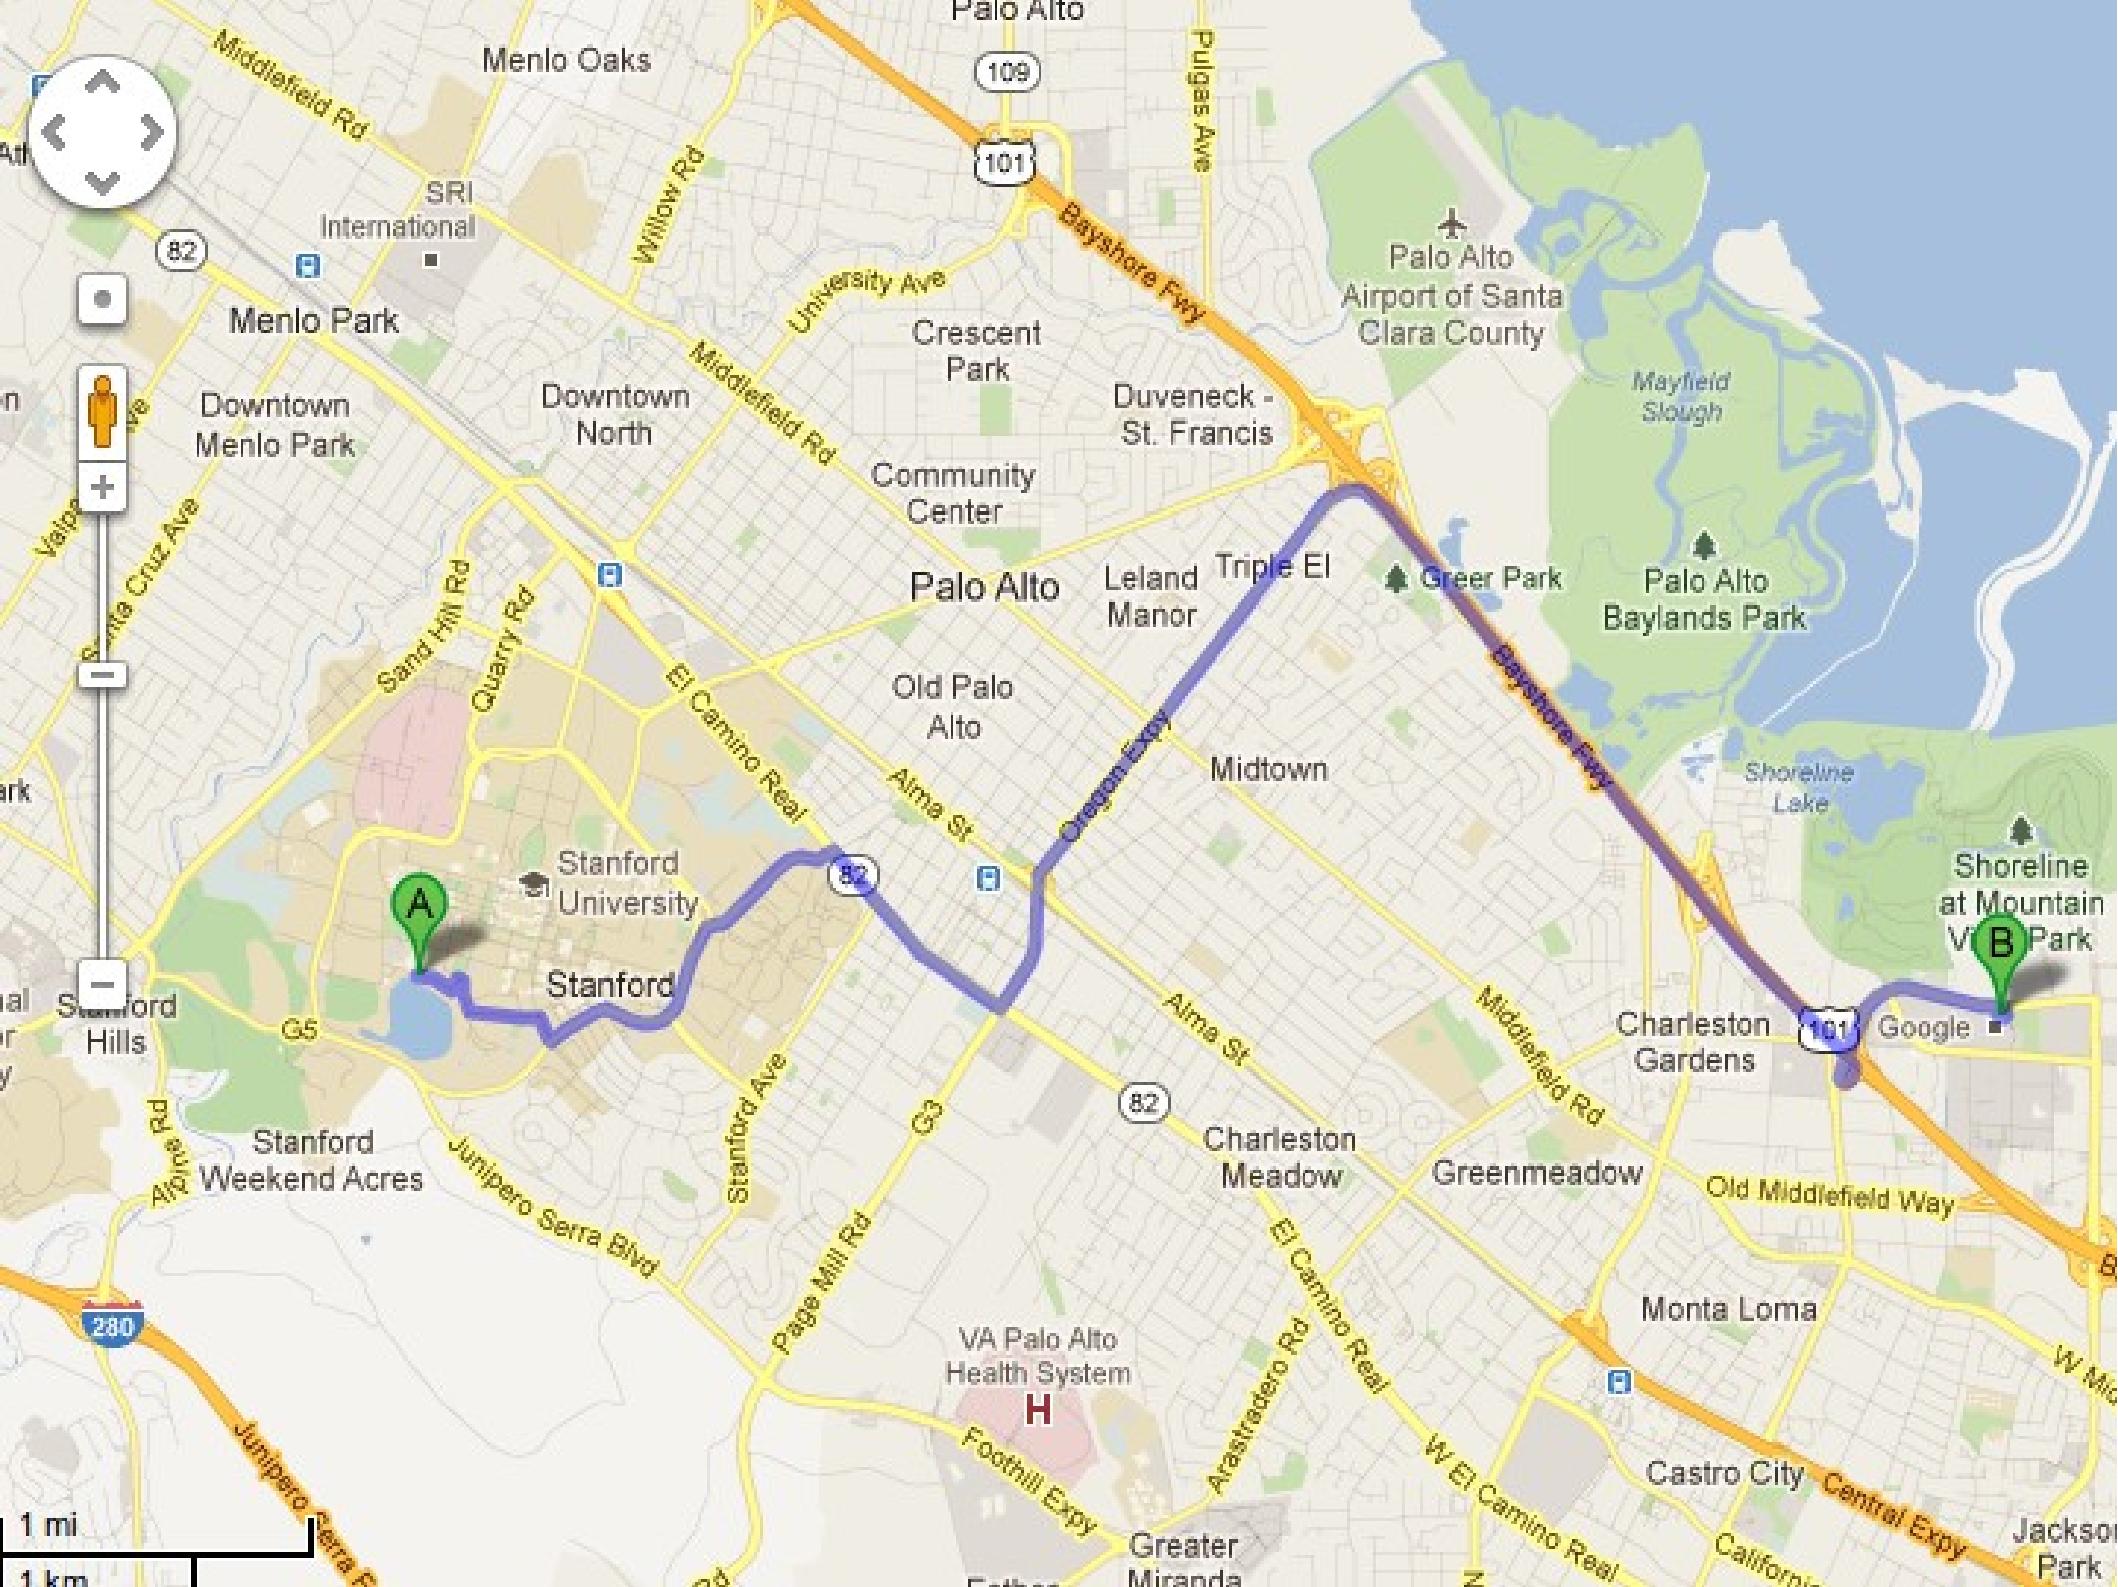
$$d(s, u) + 1 \leq d(s, x) + 1 \leq |P|$$

contradicting the fact that $|P| < d(s, u) + 1$. We have reached a contradiction, so our assumption was wrong and no shorter path exists.

Since there is a path of length $d(s, u) + 1$ from s to v and no shorter path, this means that $d(s, v) = d(s, u) + 1$. ■

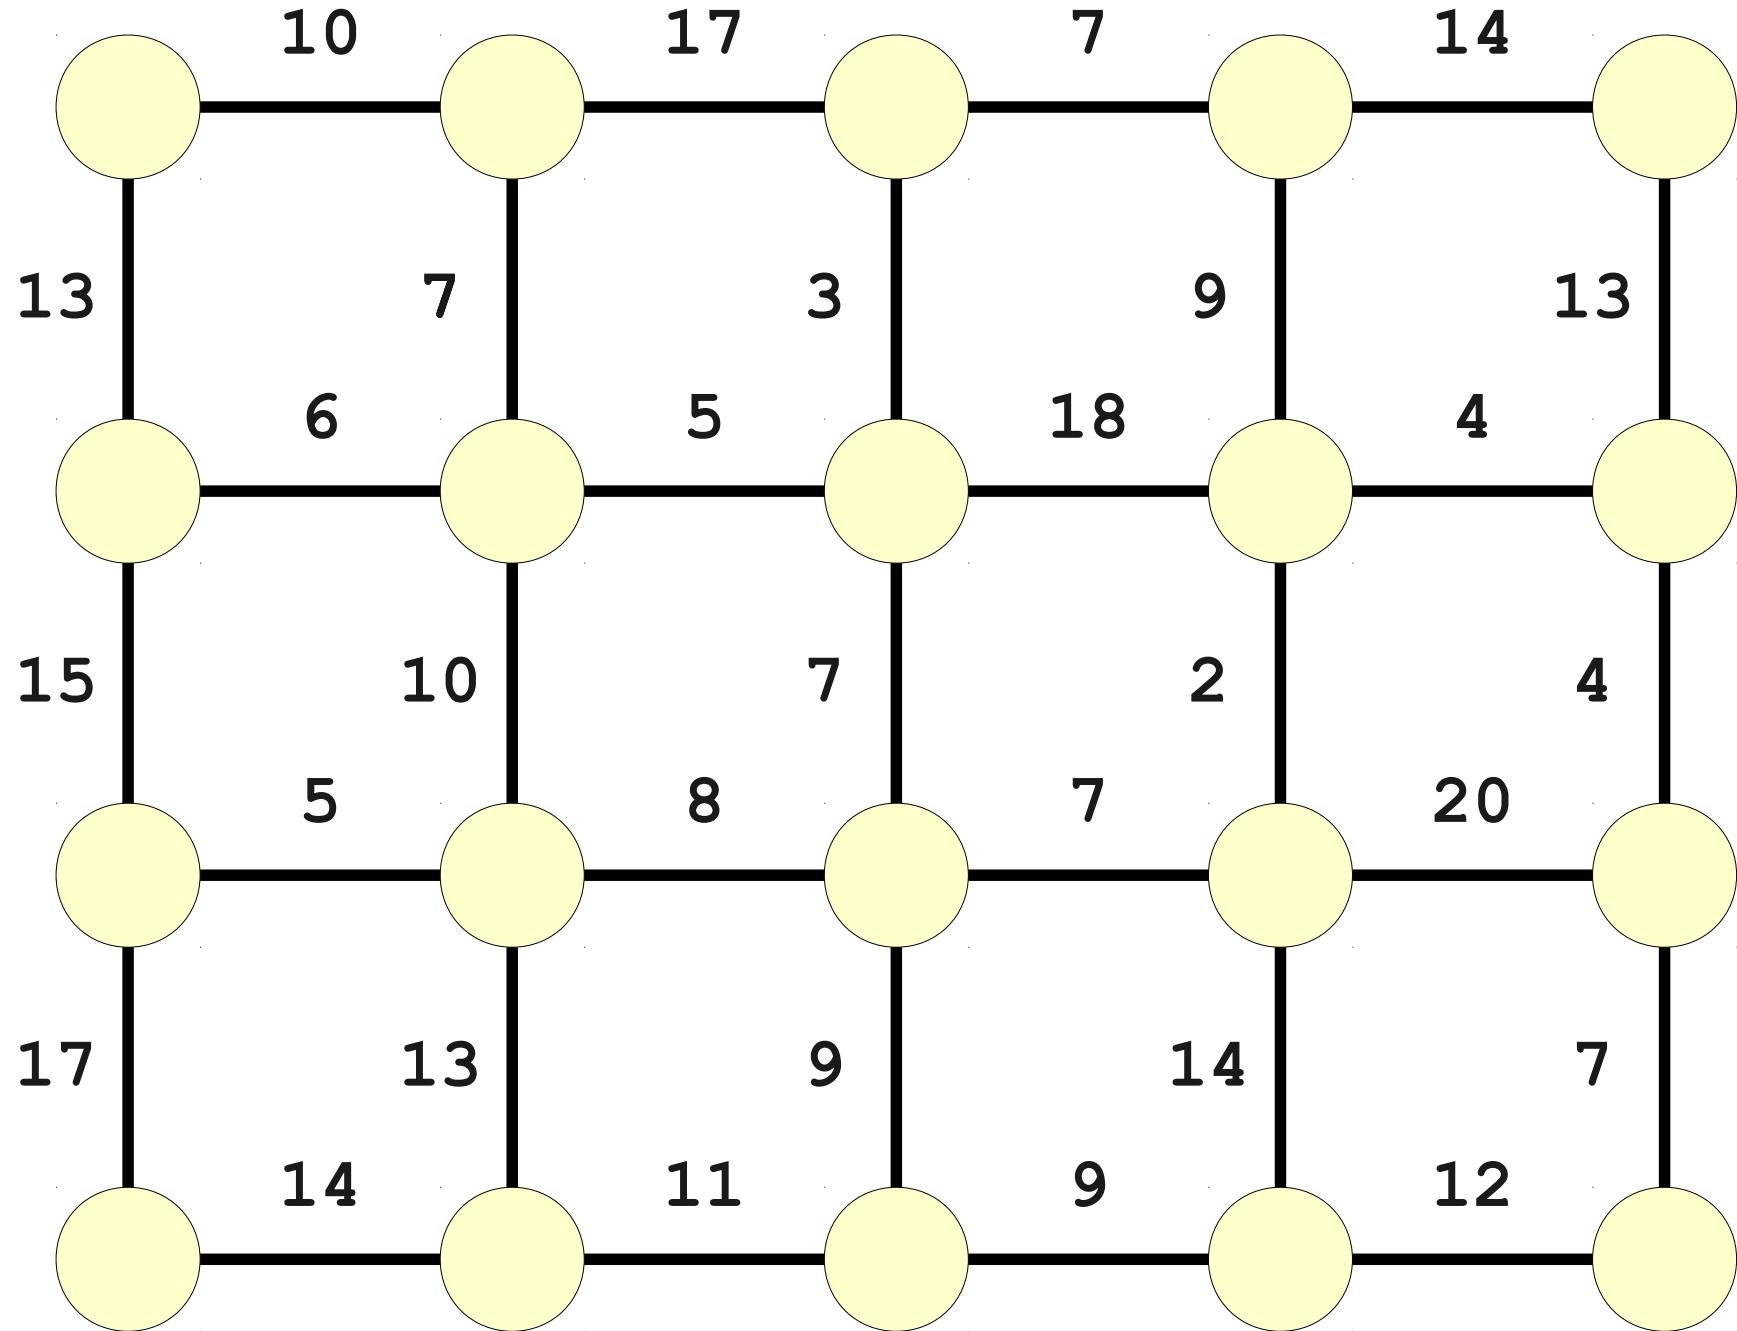
Why These Two Proofs Matter

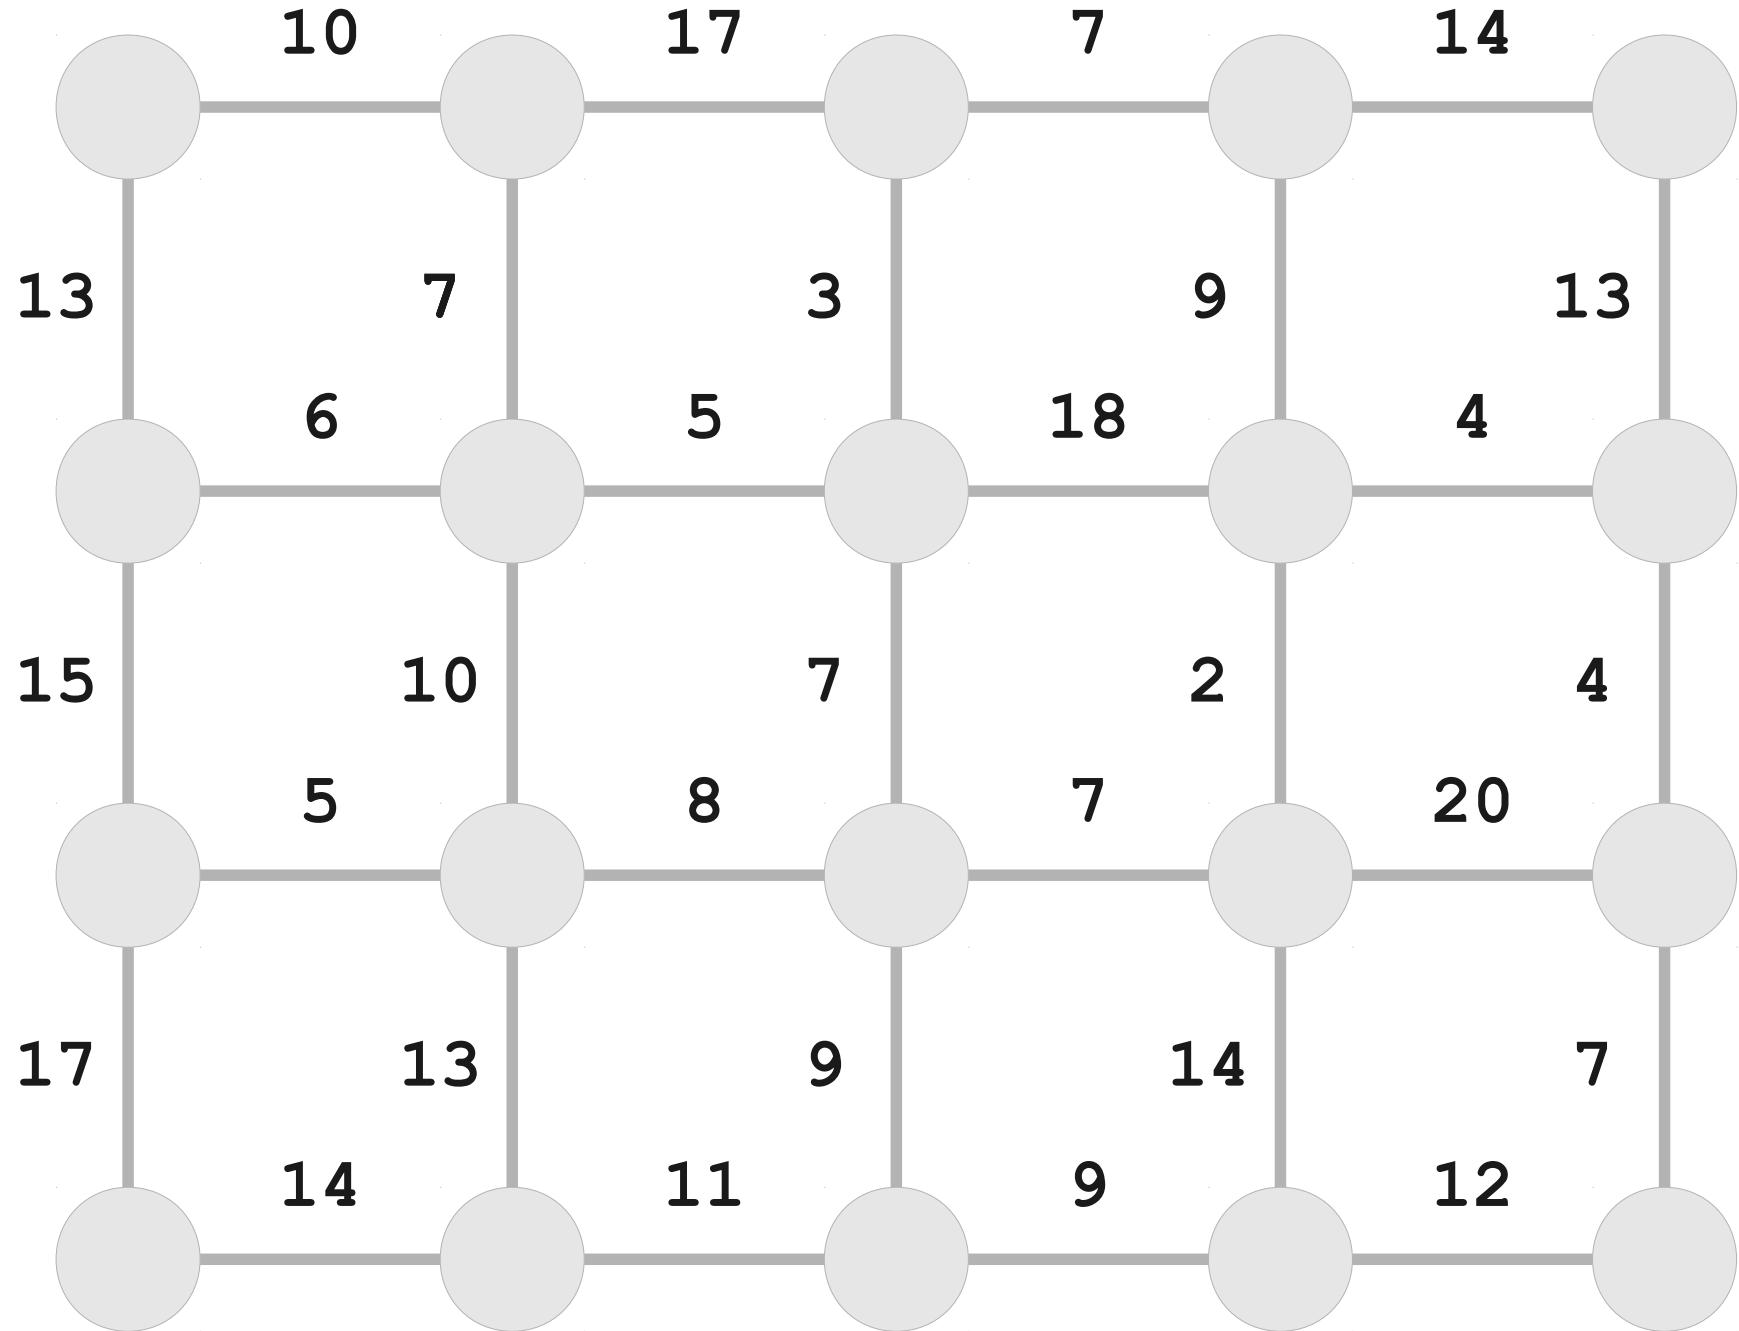
- The first proof of correctness (based on layers) is based on our first observation: the nodes visited in BFS radiate outward from the start node in ascending order of distance.
- The second proof of correctness (based on picking the lowest yellow node) is based on our second observation: picking the lowest-cost yellow node correctly computes a shortest path.
- Interestingly, this second correctness proof can be generalized to a larger setting...

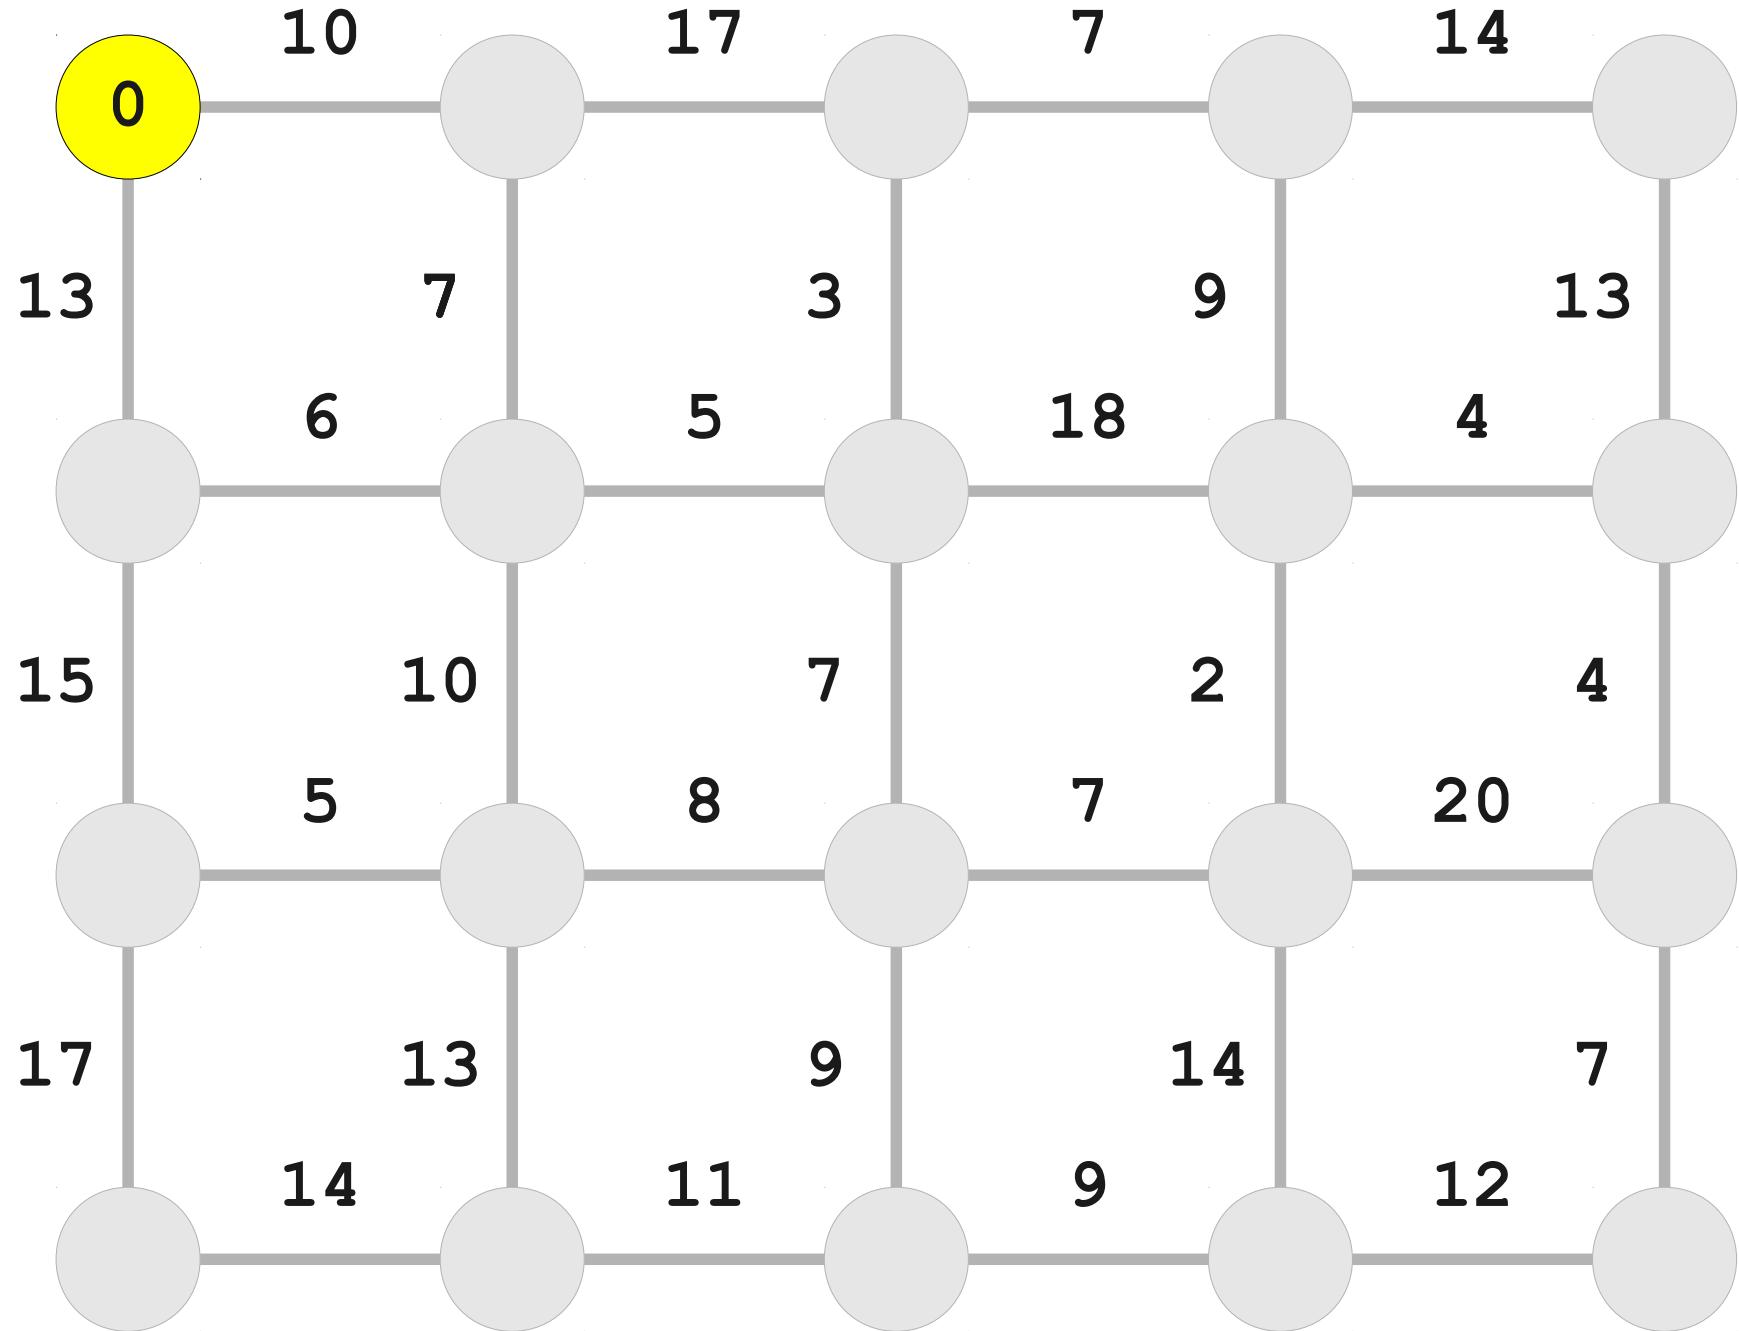


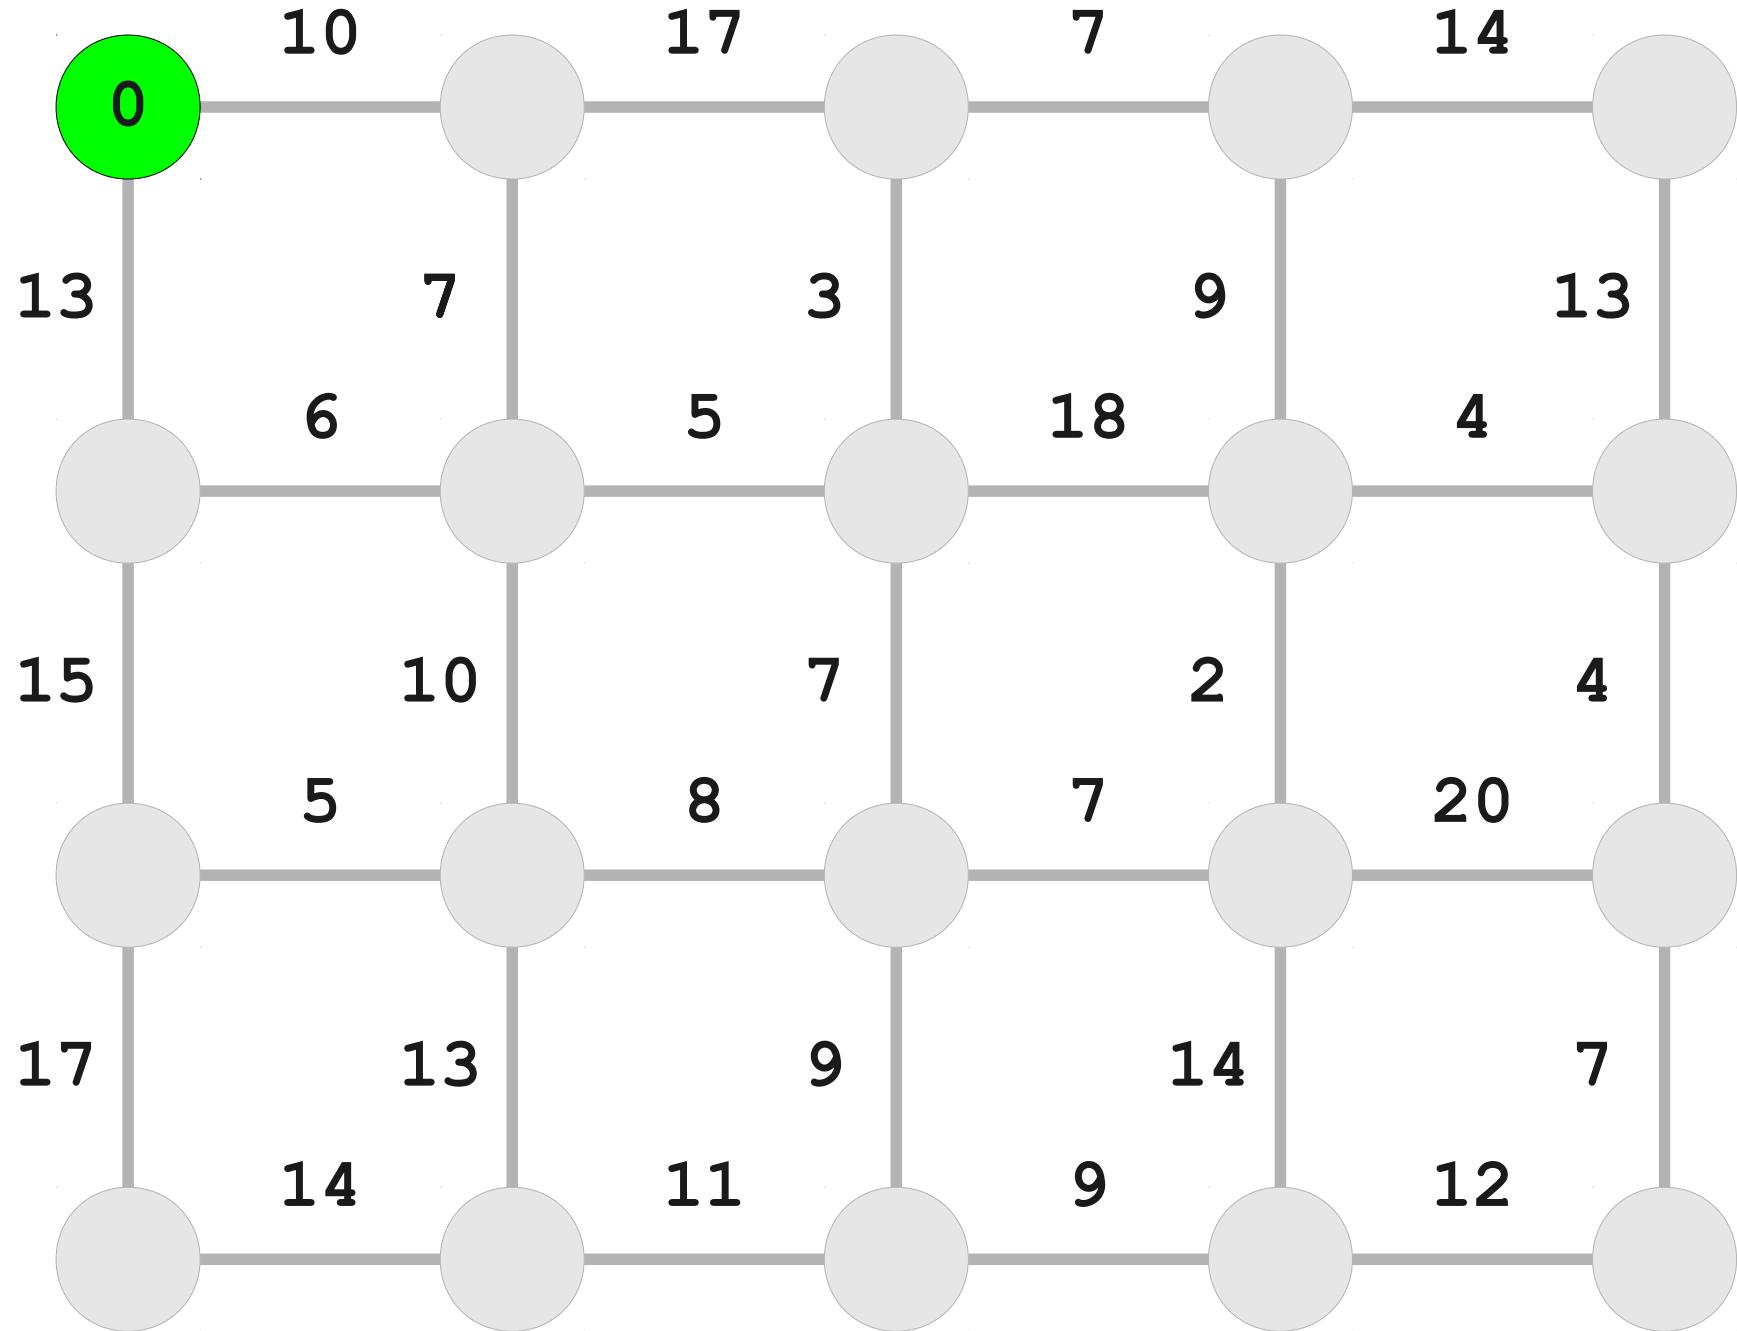
Edges with Costs

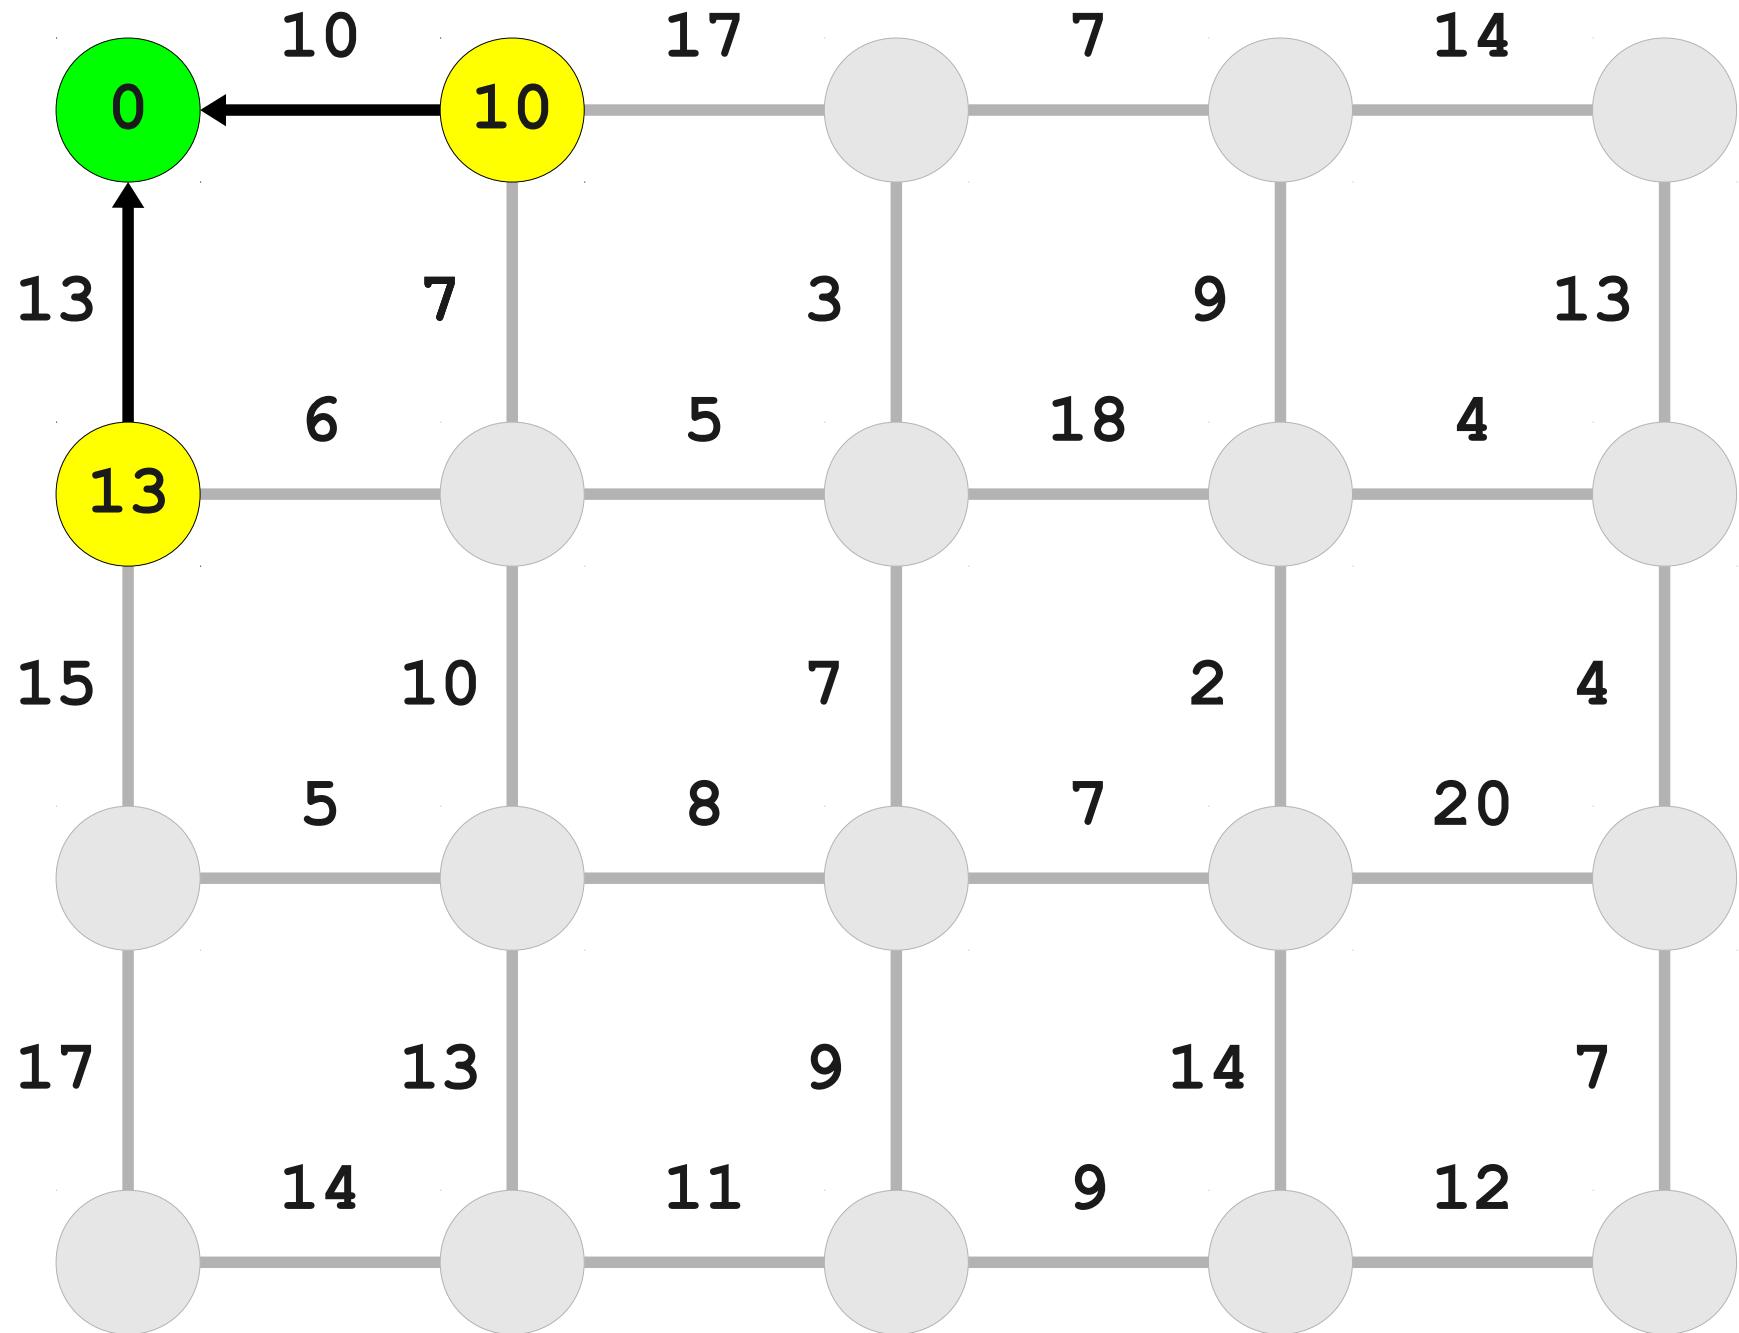
- In many applications, edges have an associated length (or cost, weight, etc.), denoted $\mathbf{l}(u, v)$.
- **Assumption:** Lengths are nonnegative. (We'll revisit this later in the quarter.)
- Let's say that the length of a path P (denoted $\mathbf{l}(P)$) is the sum of all the edge lengths in the path P .
- Goal: find the shortest path from s to every node in V , taking costs into account.

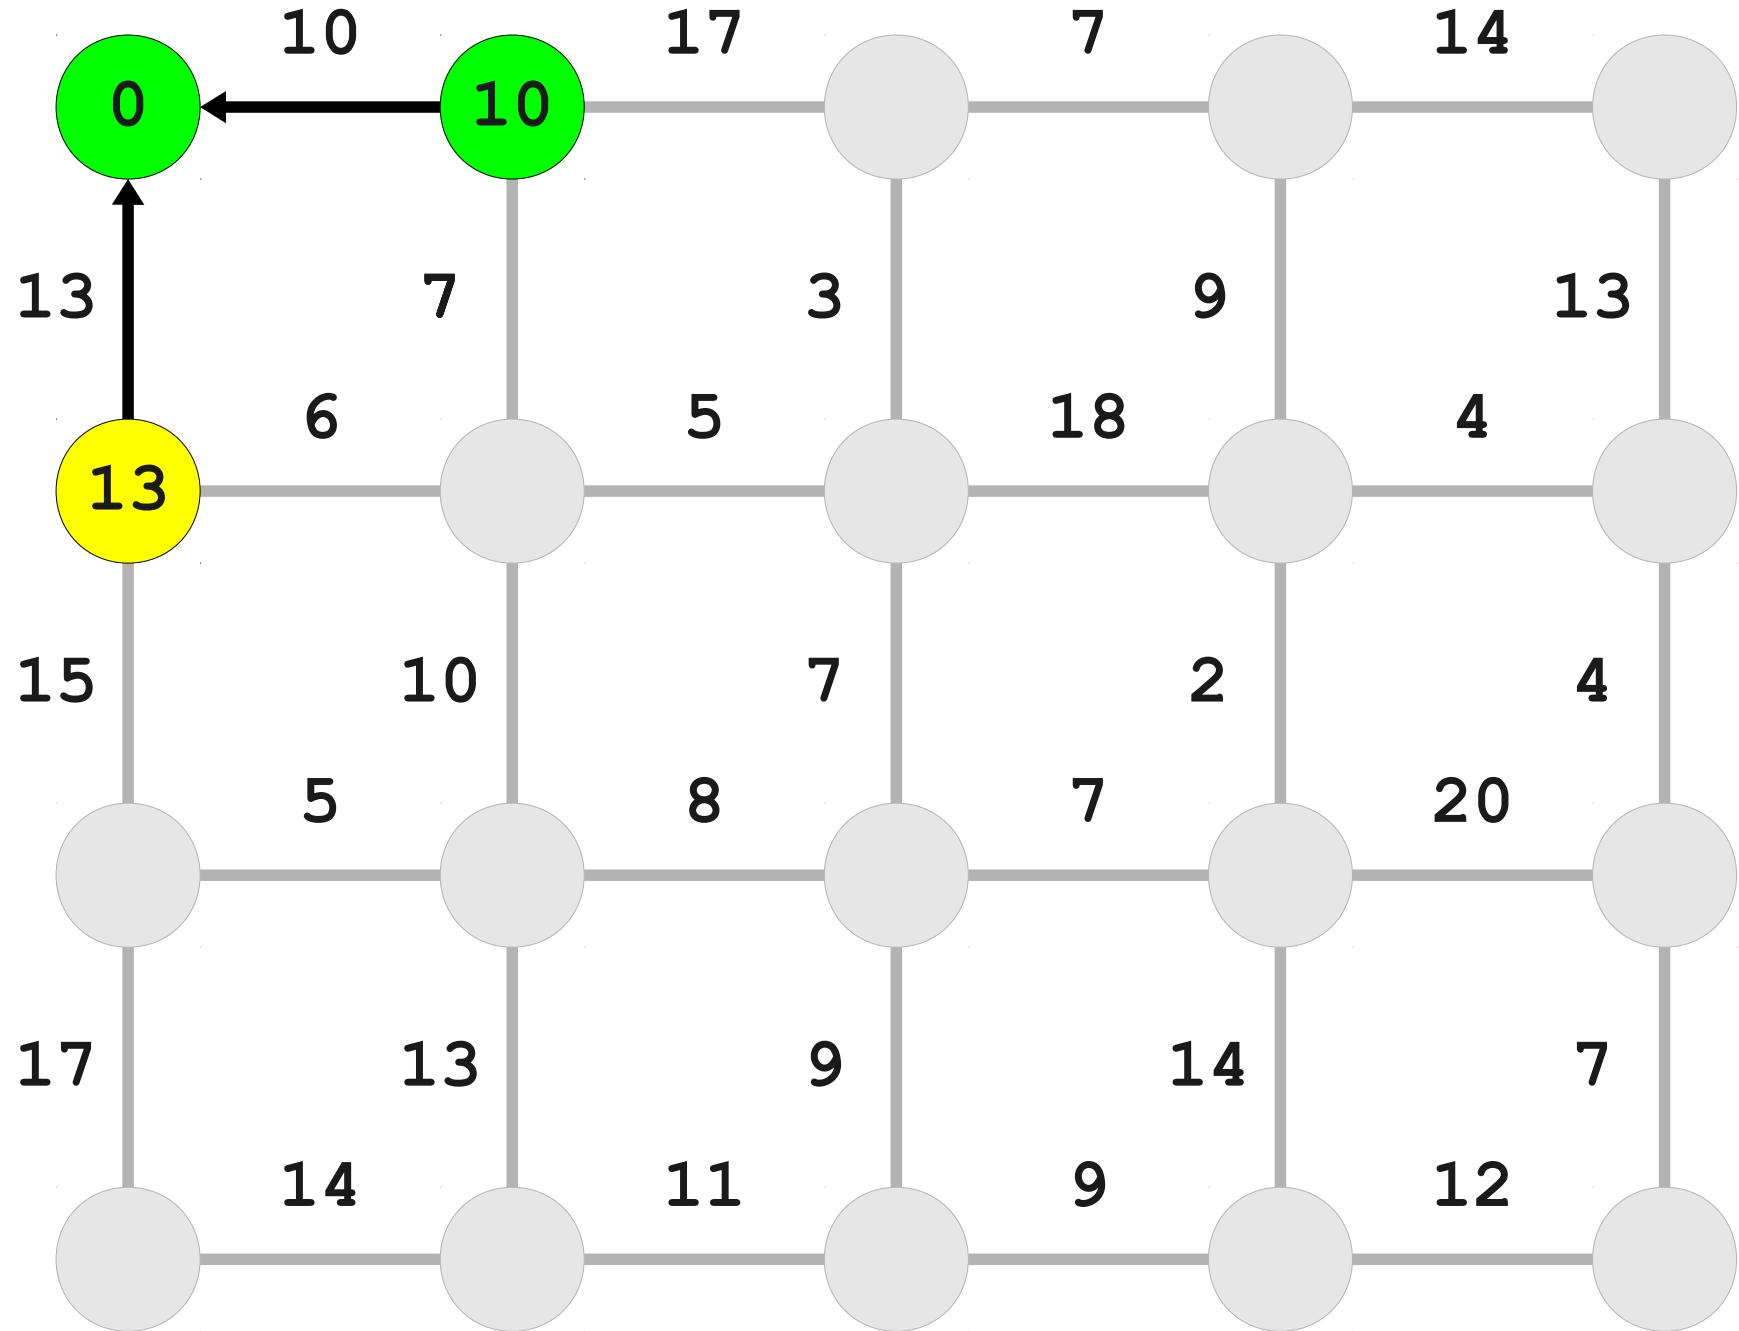


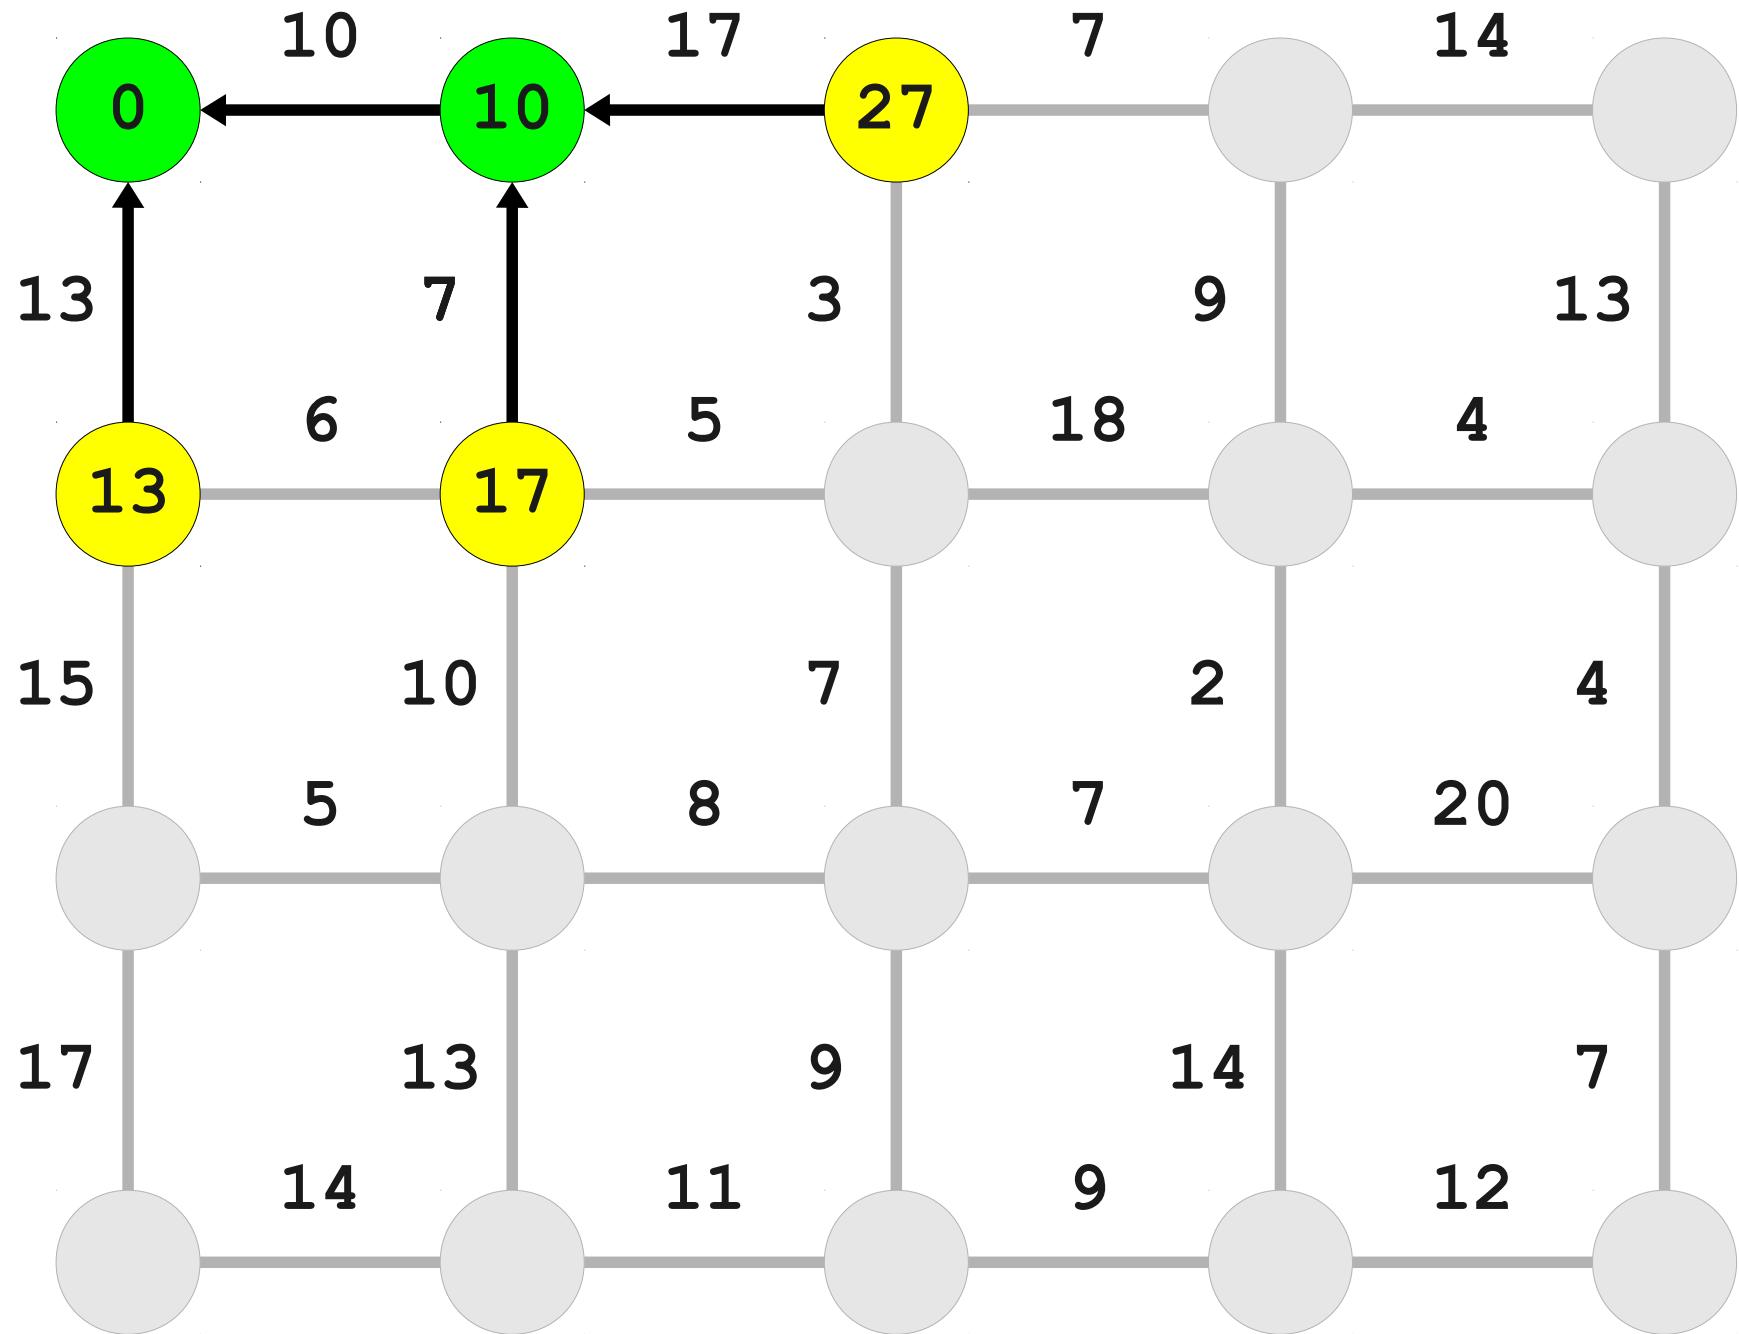


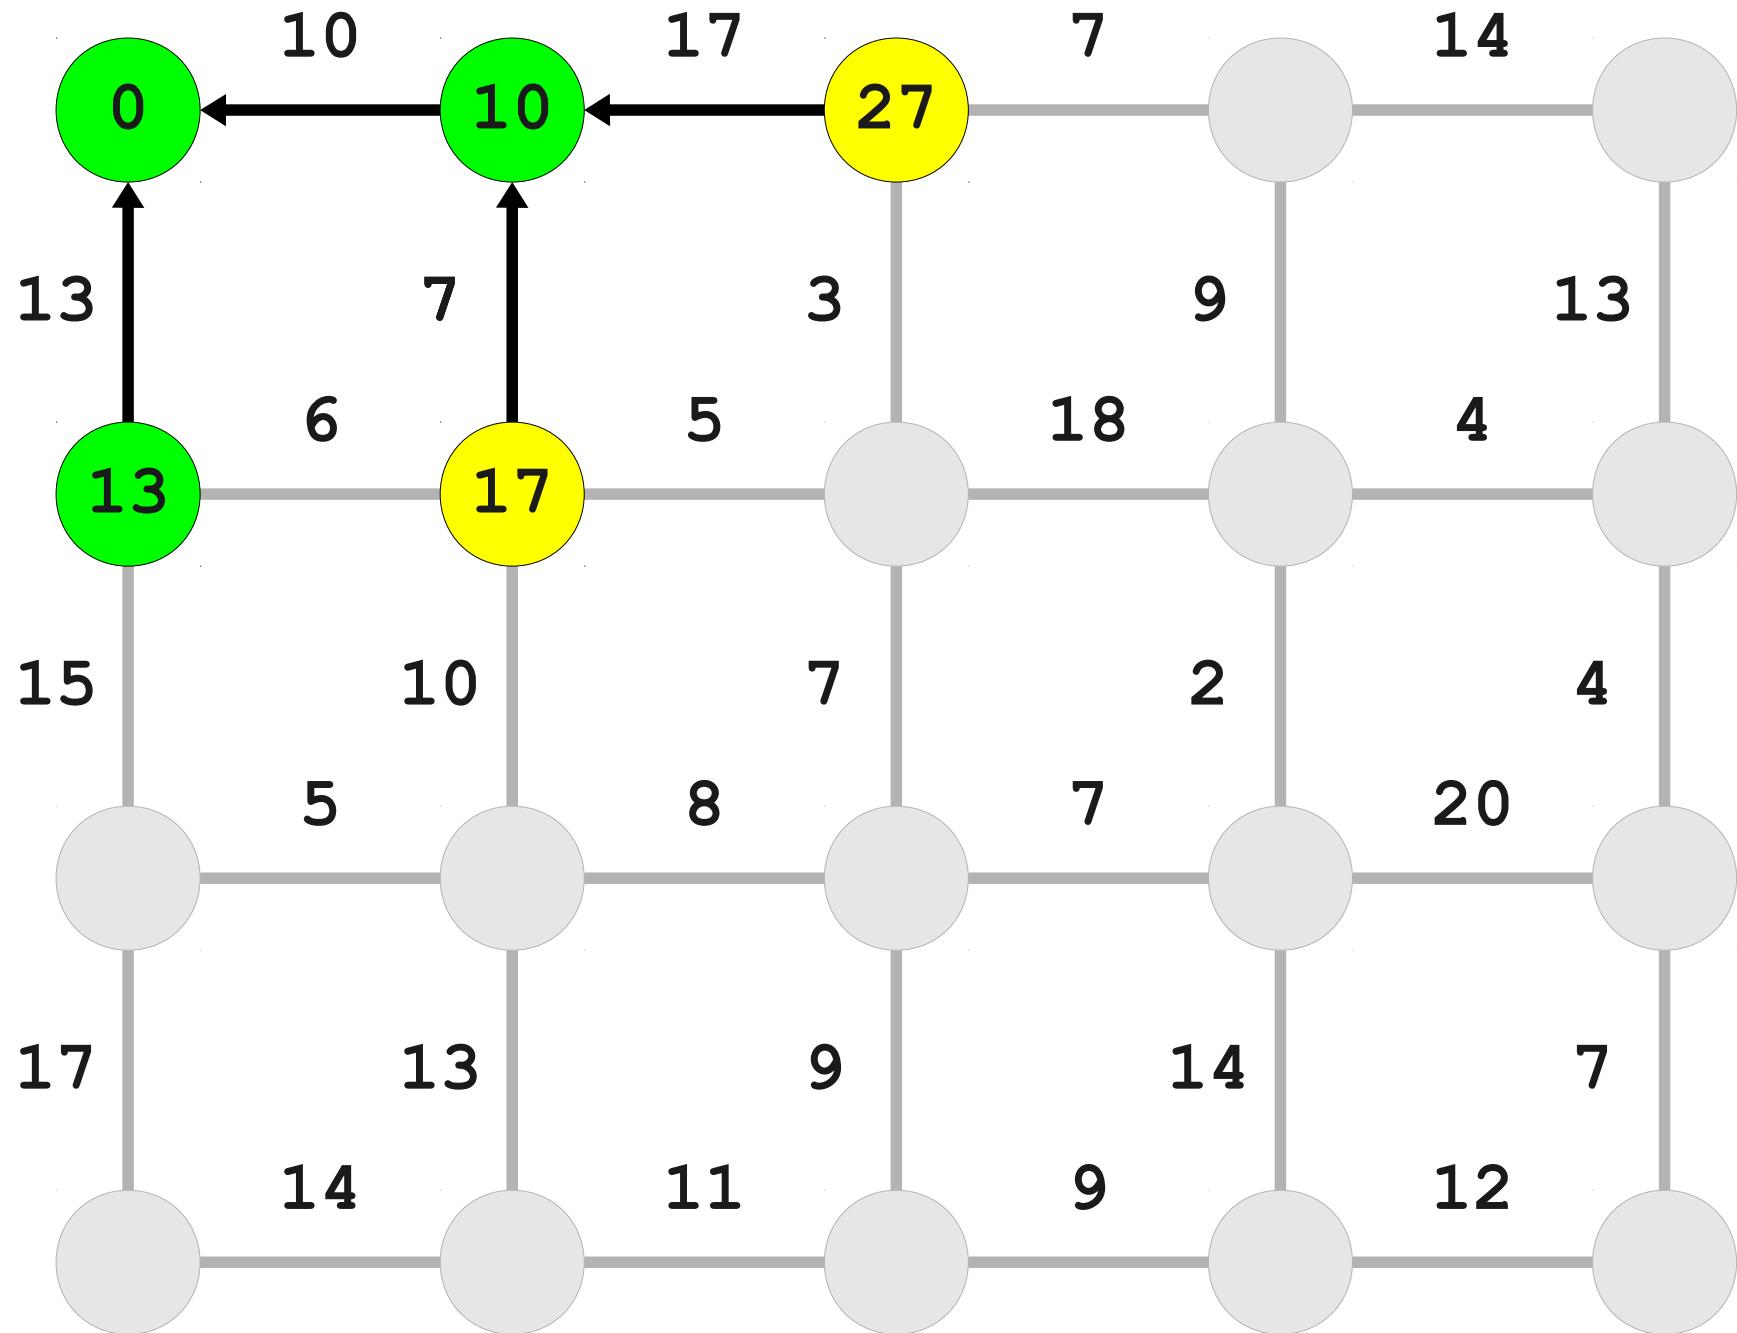


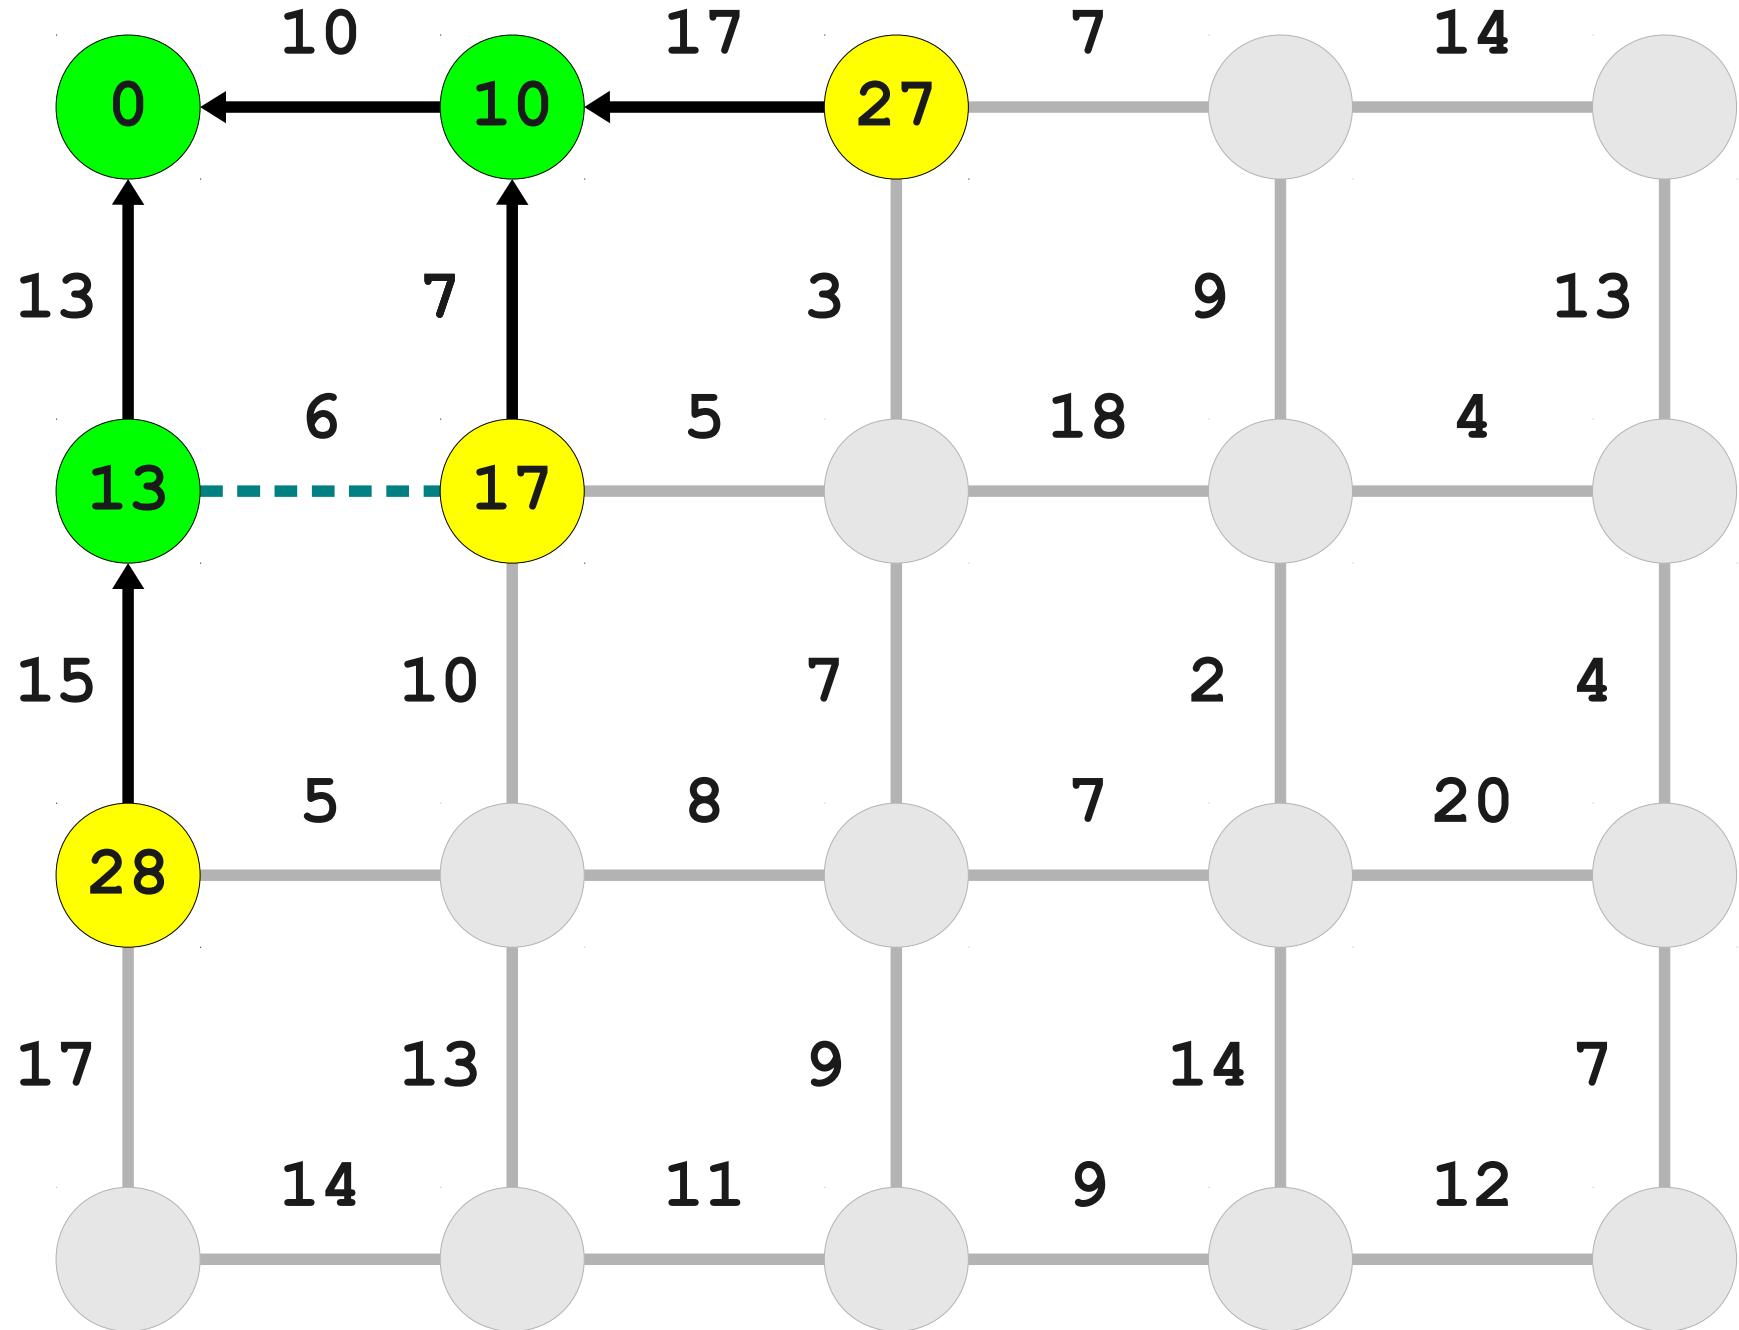


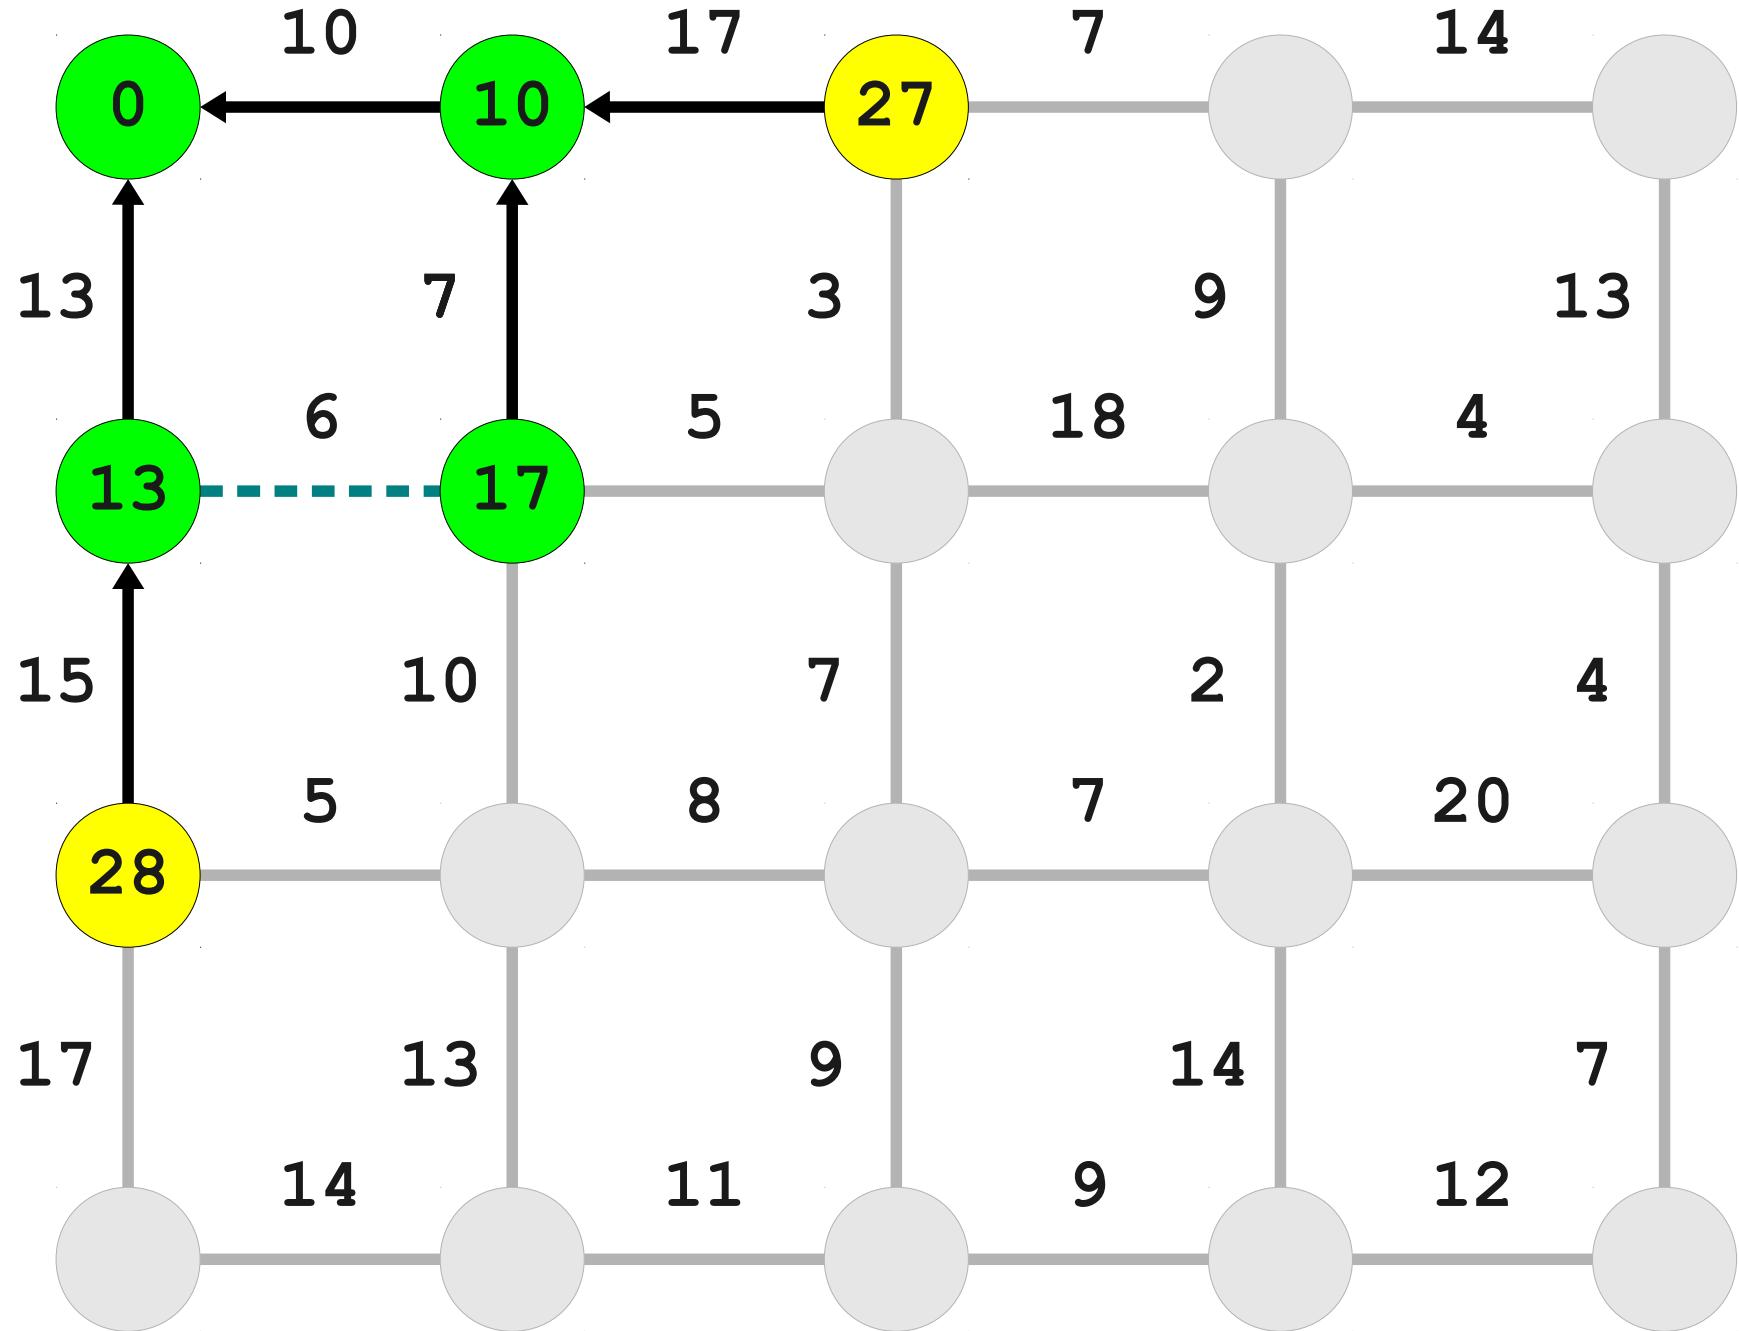


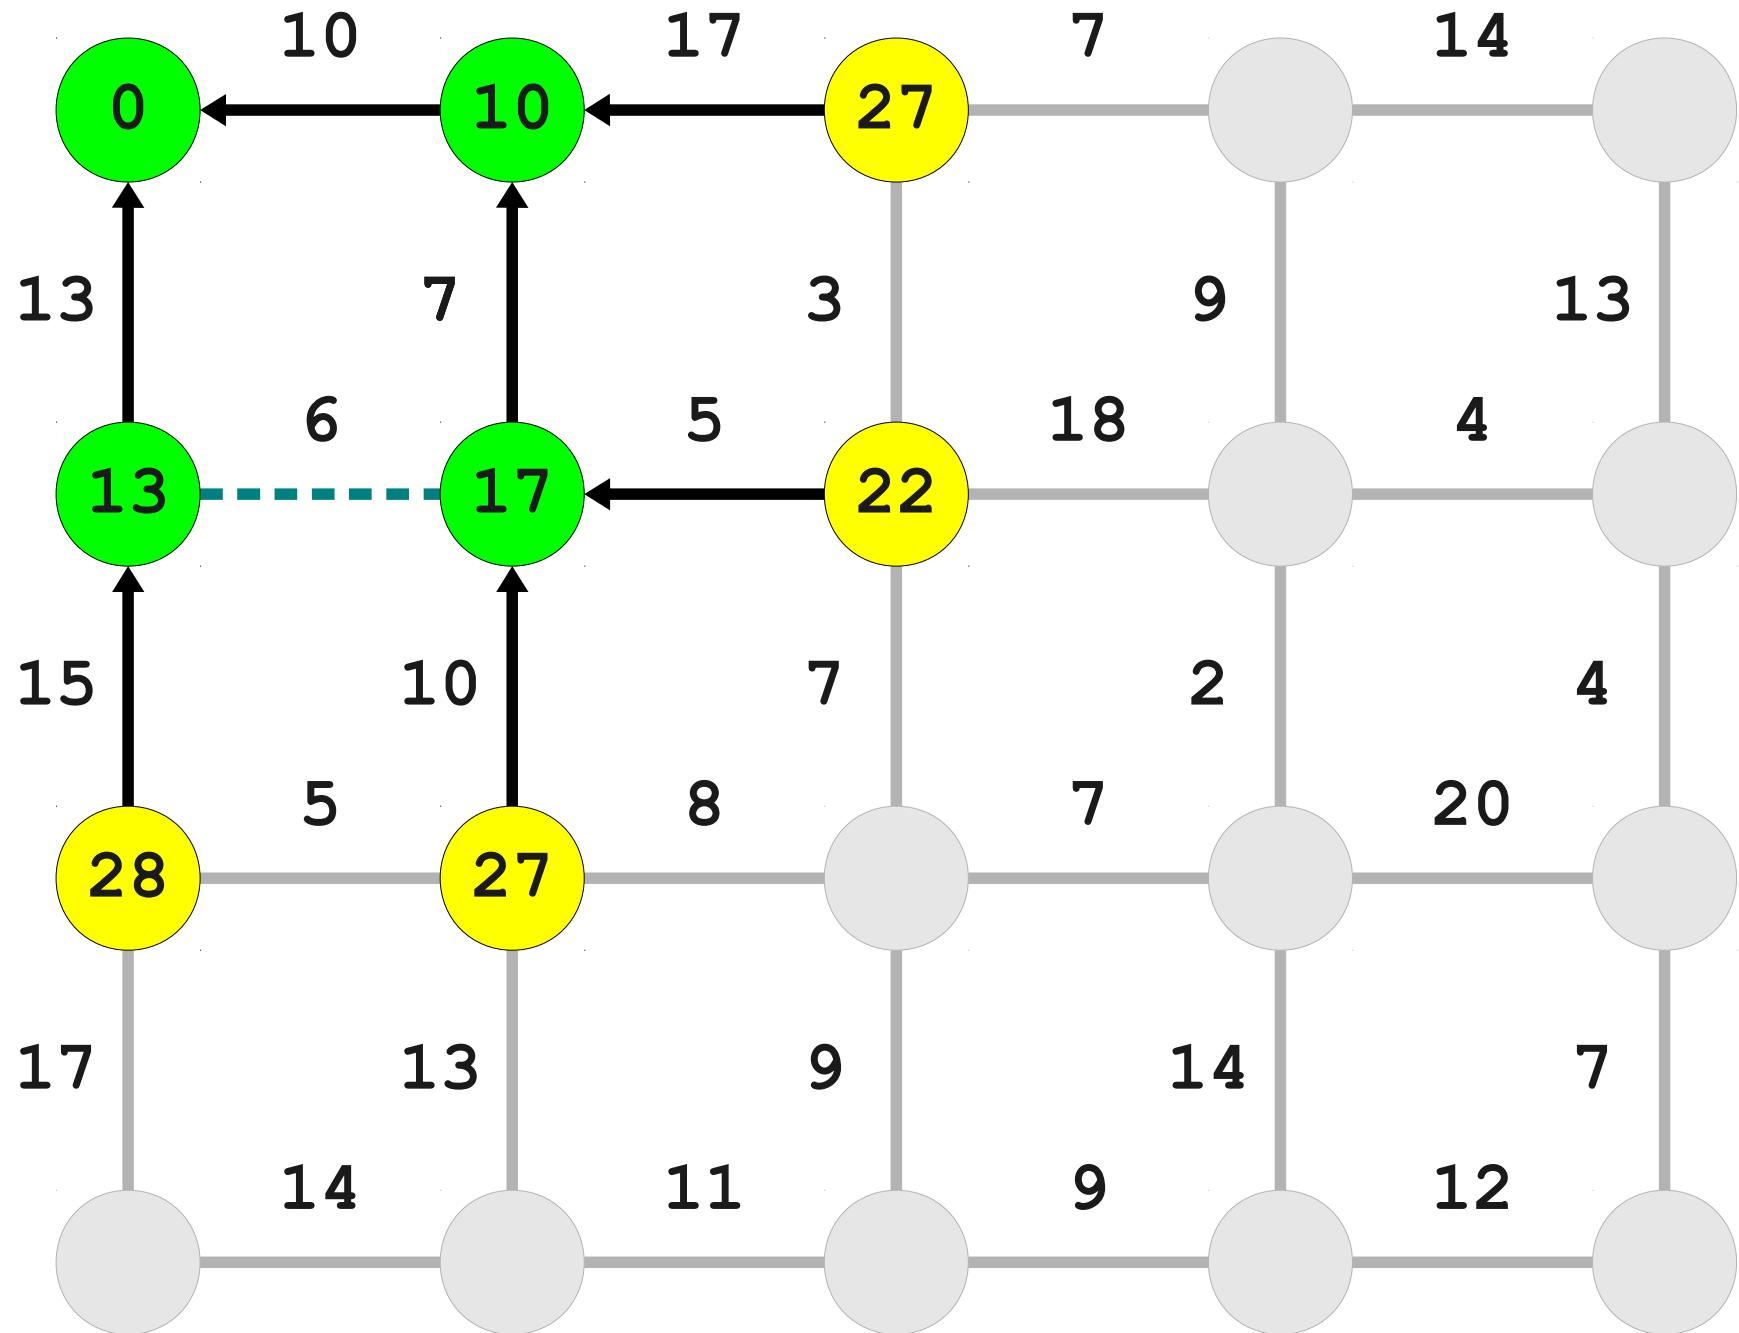


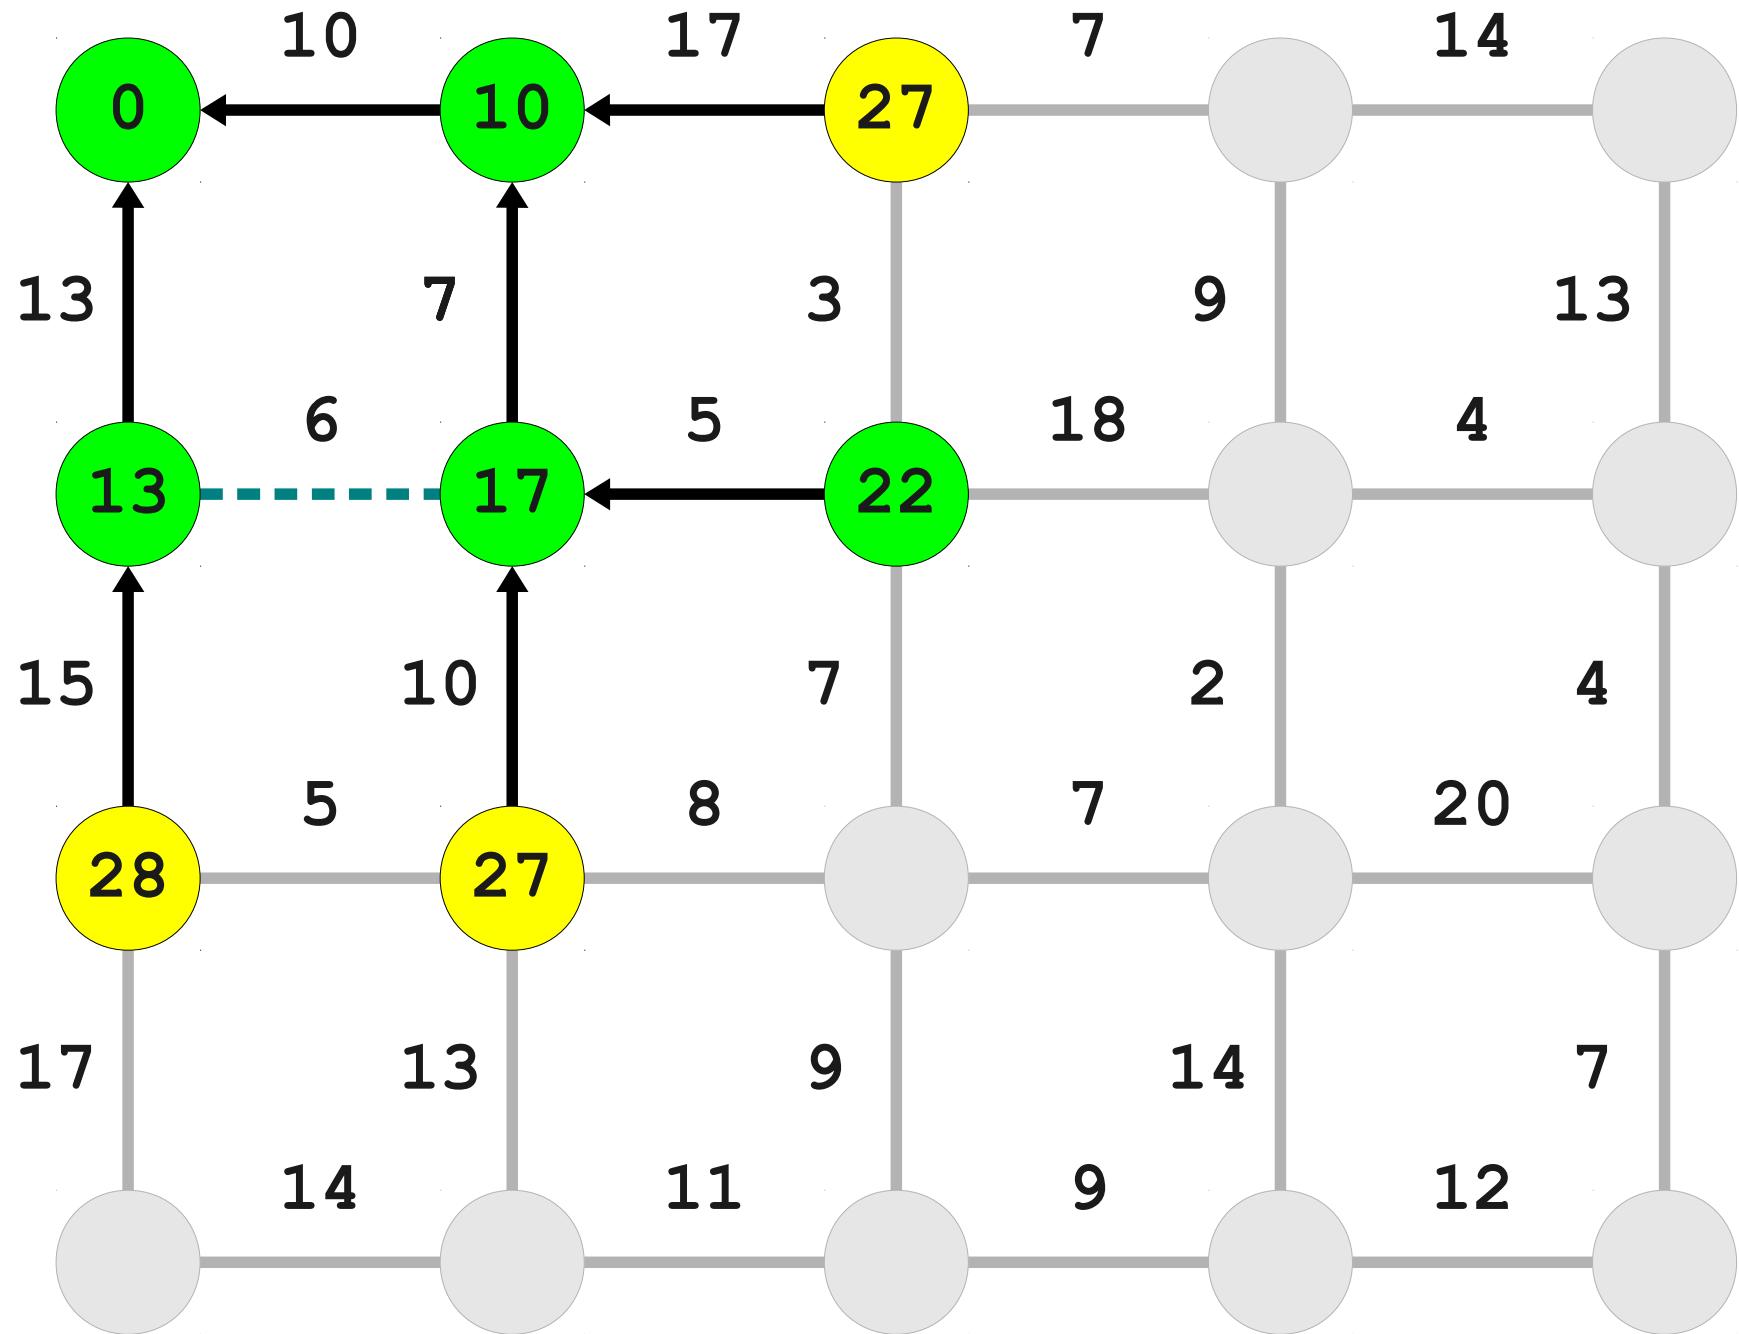


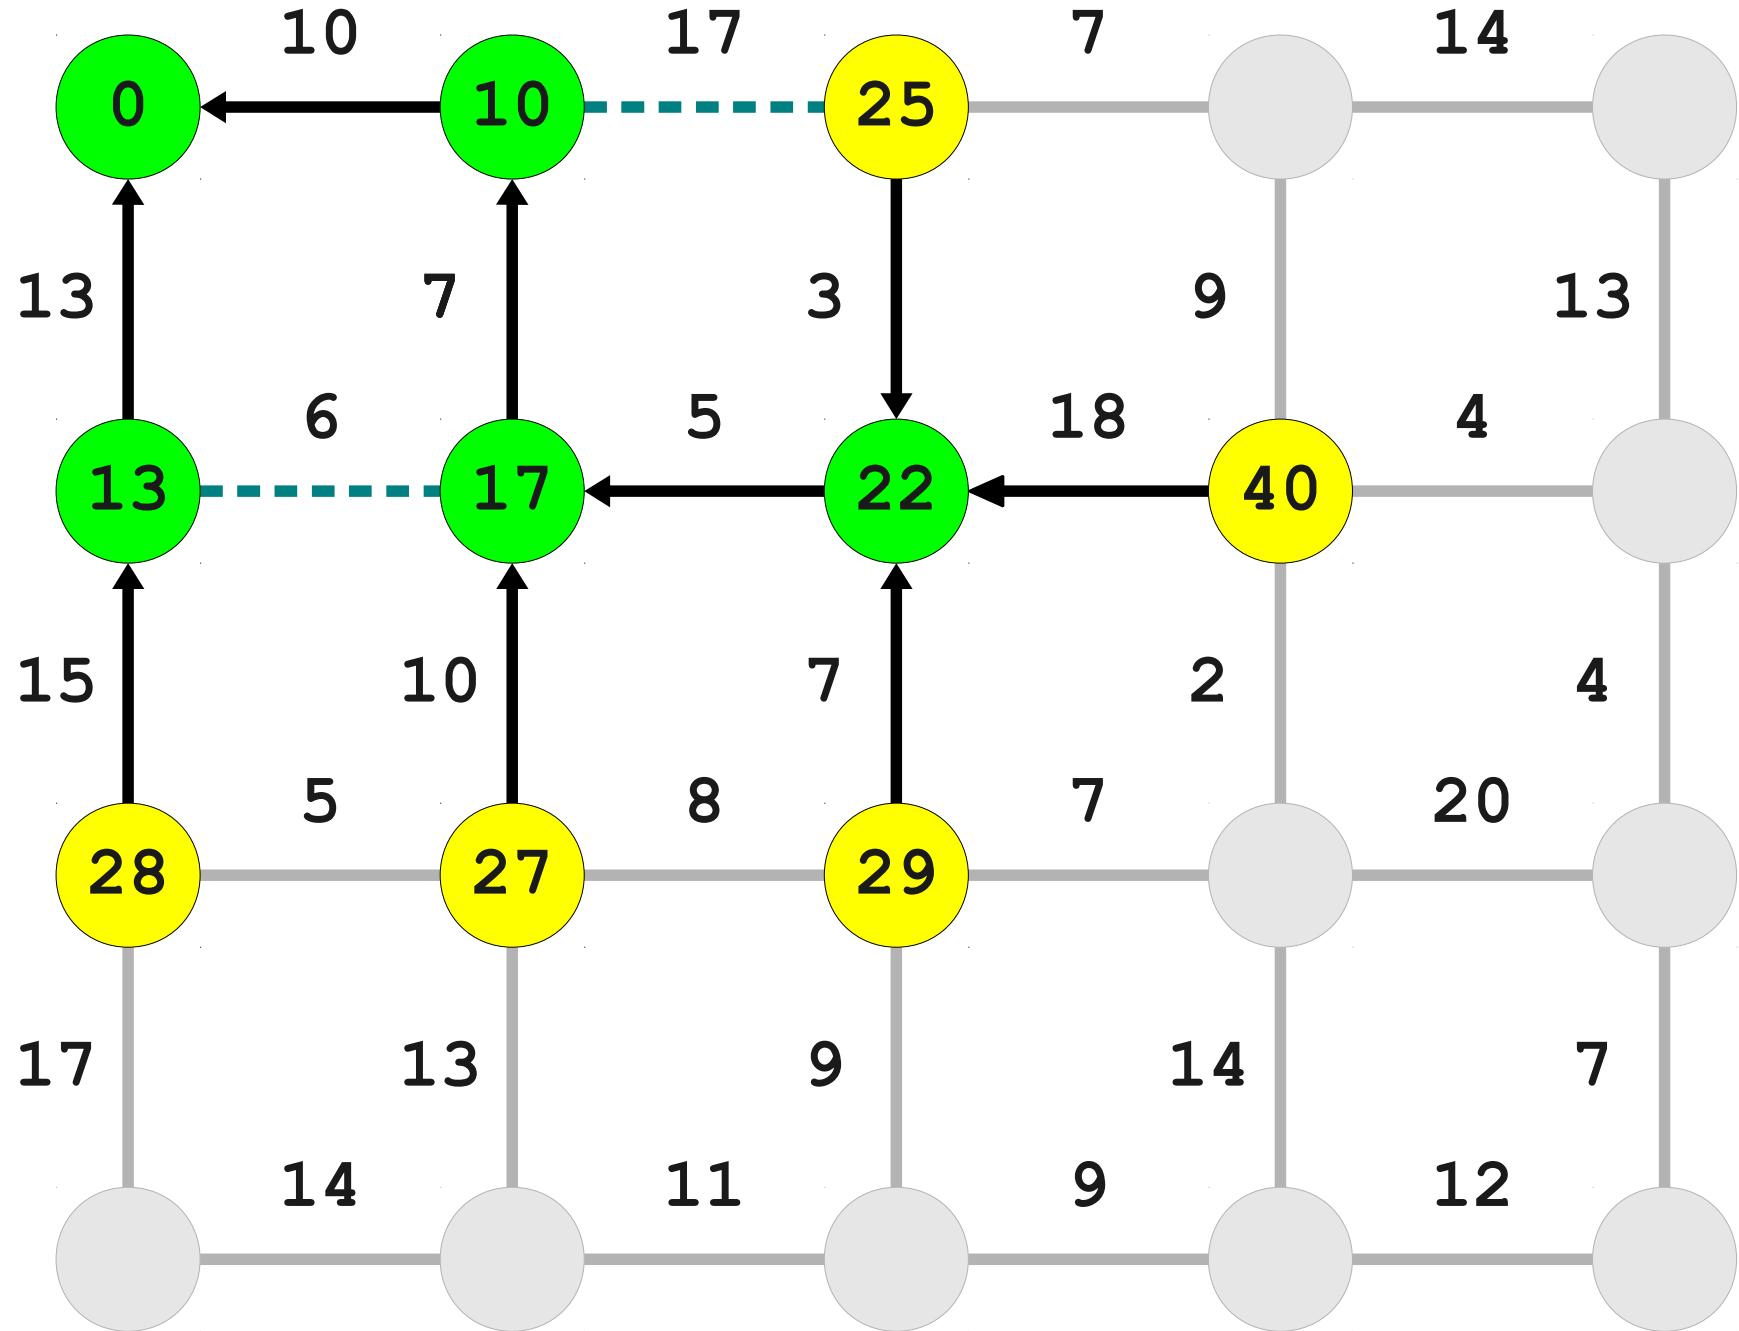


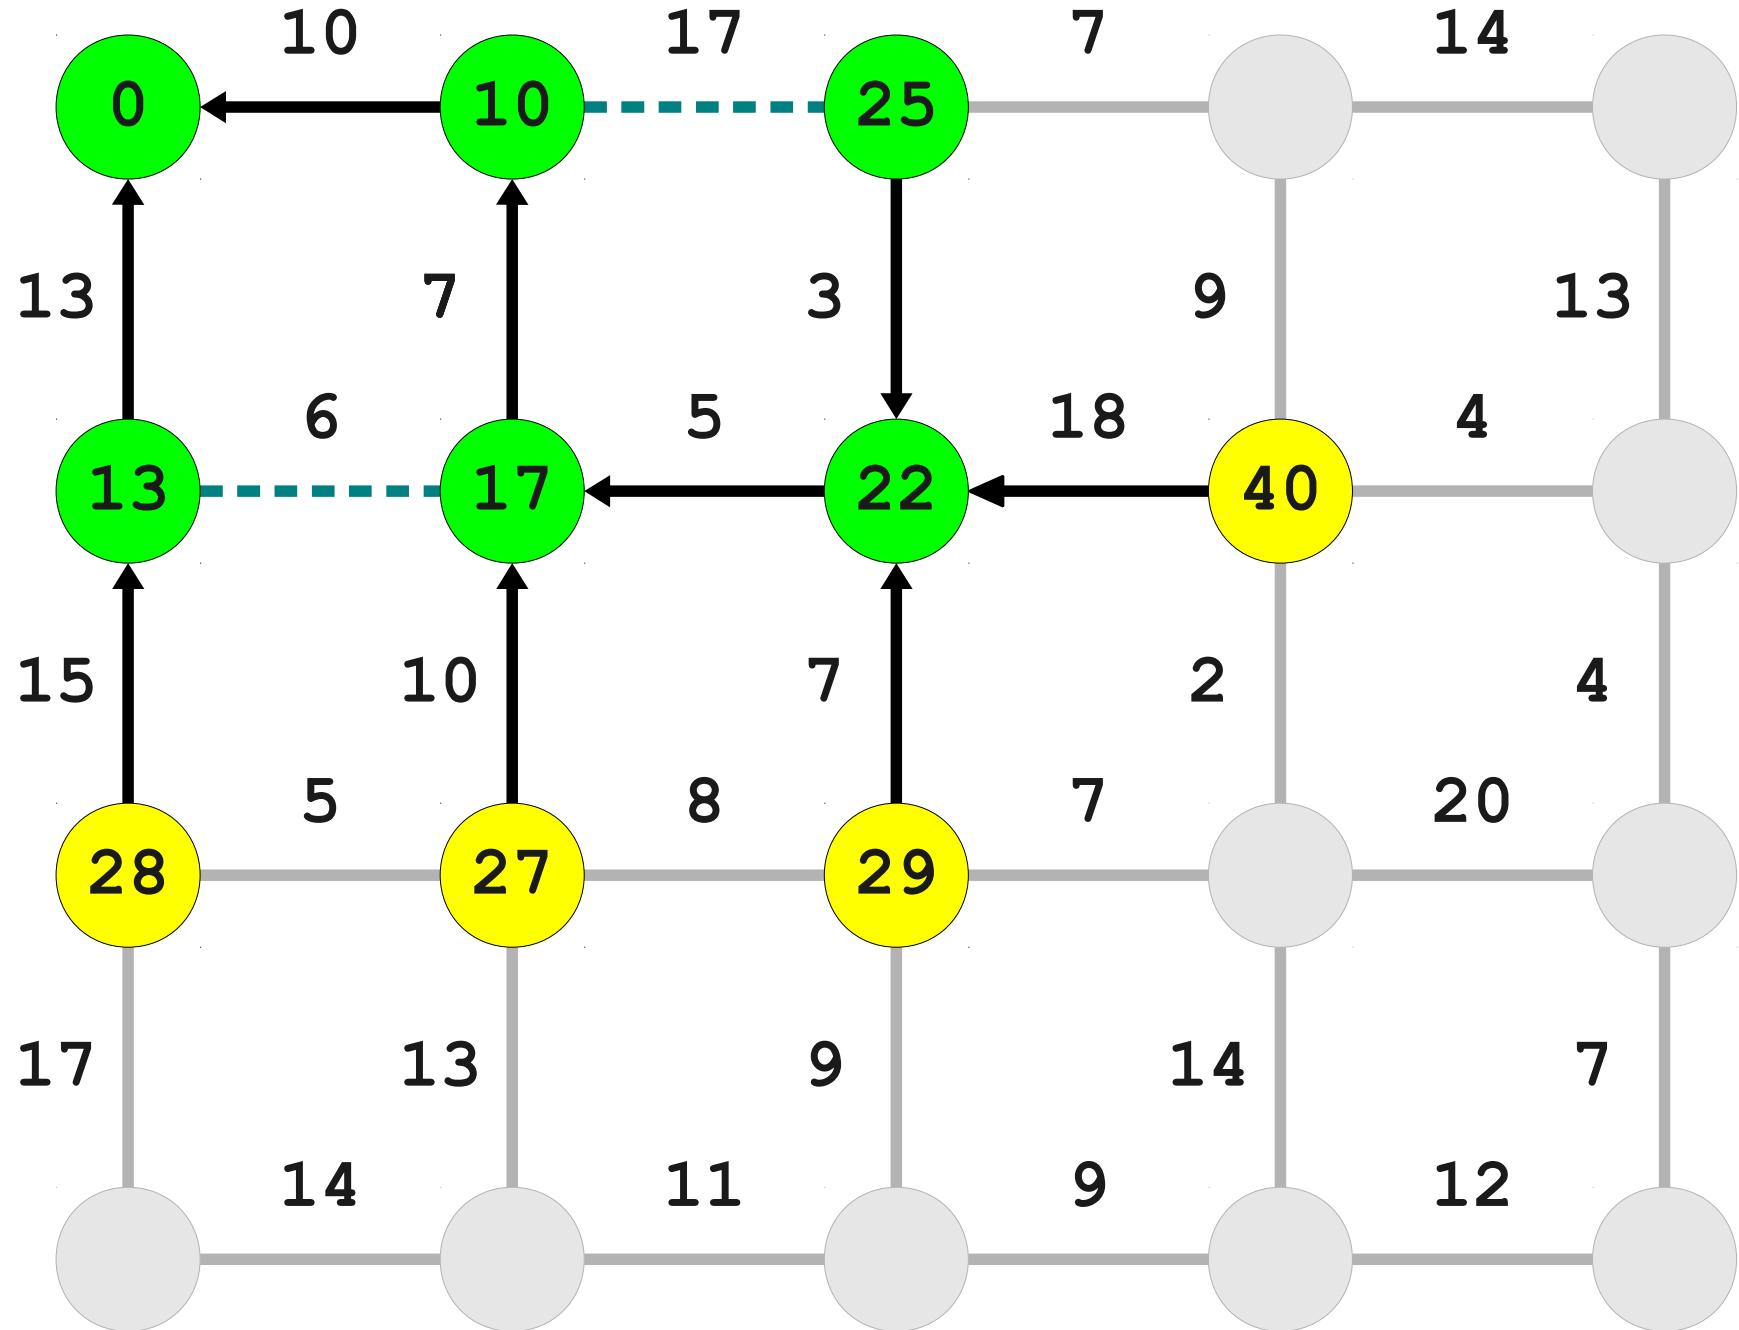


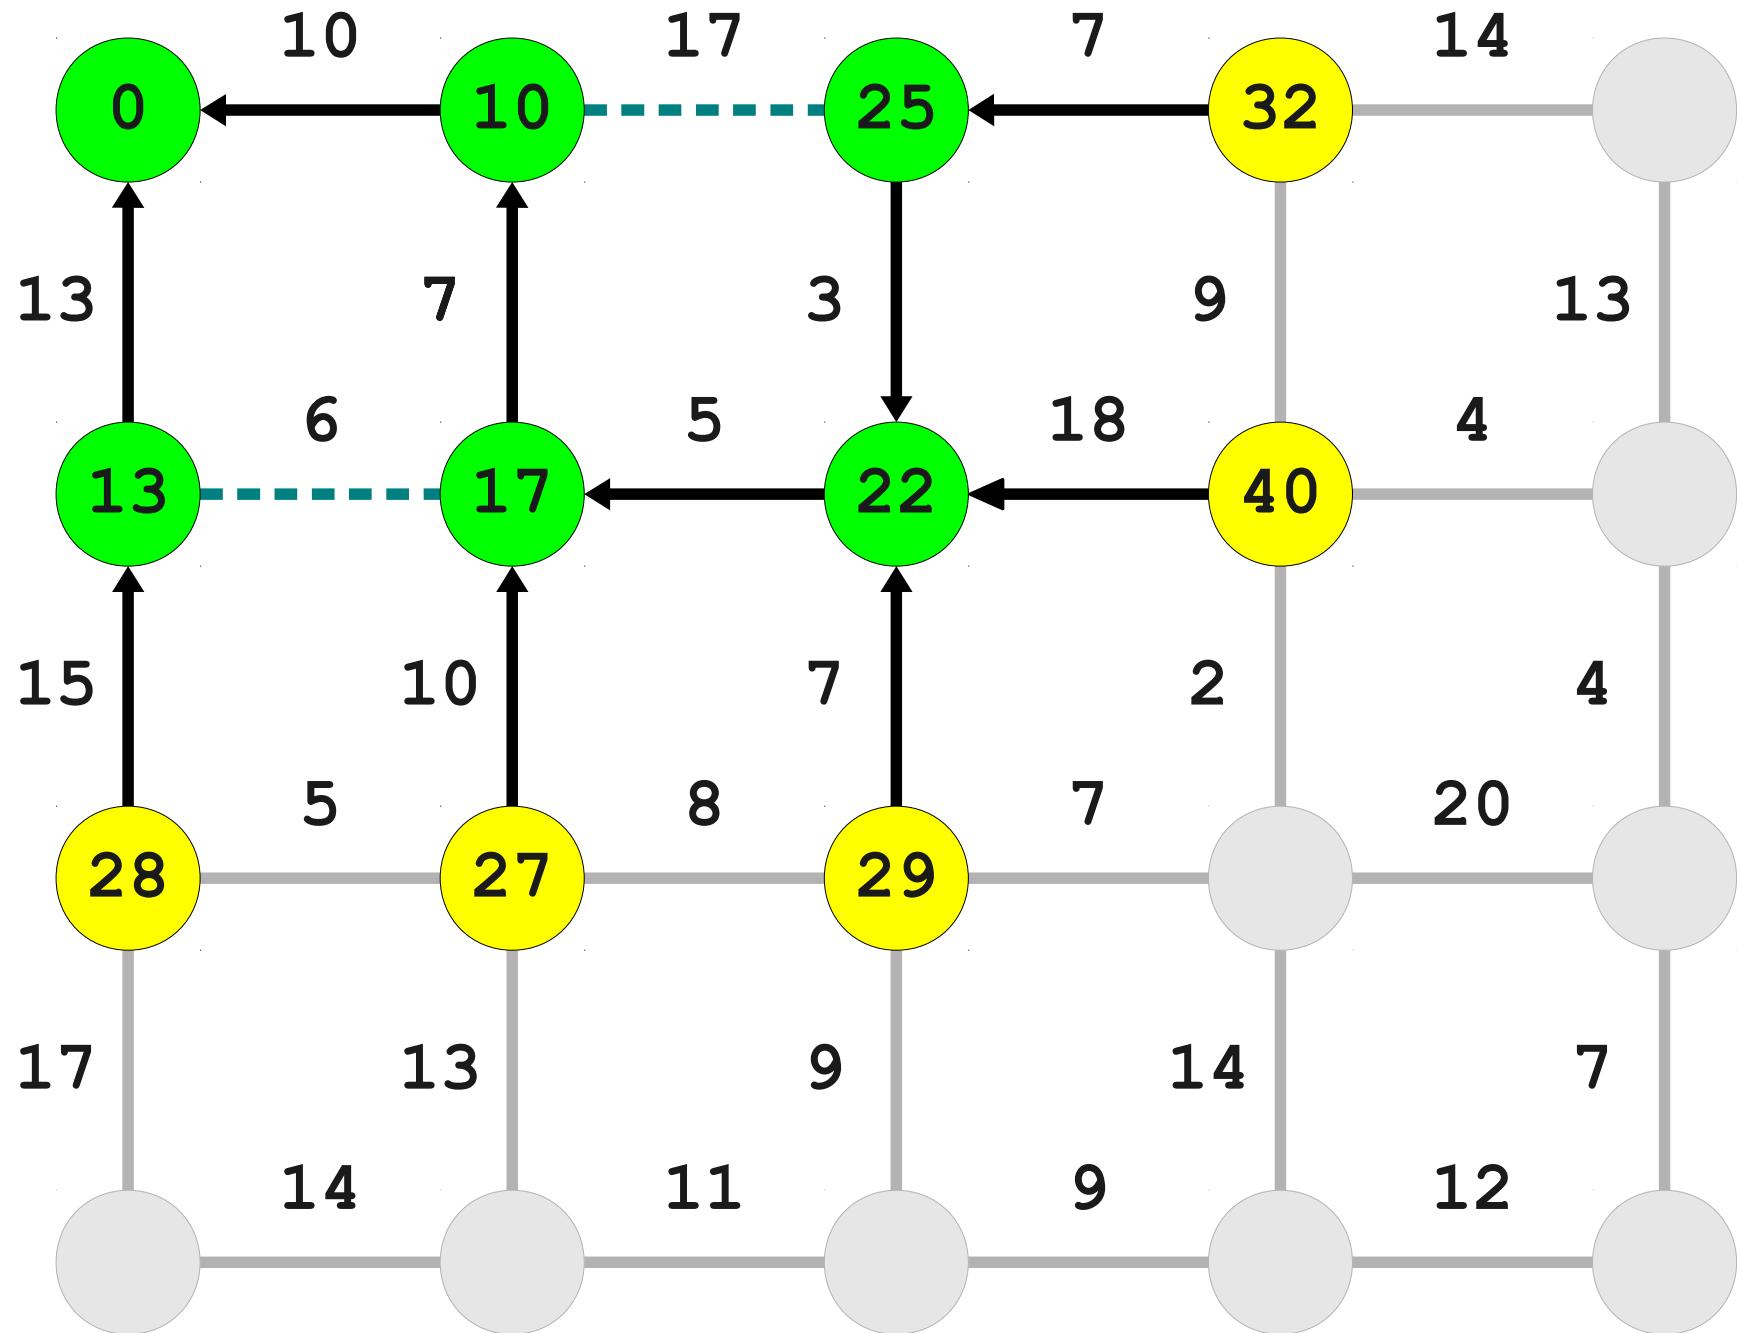


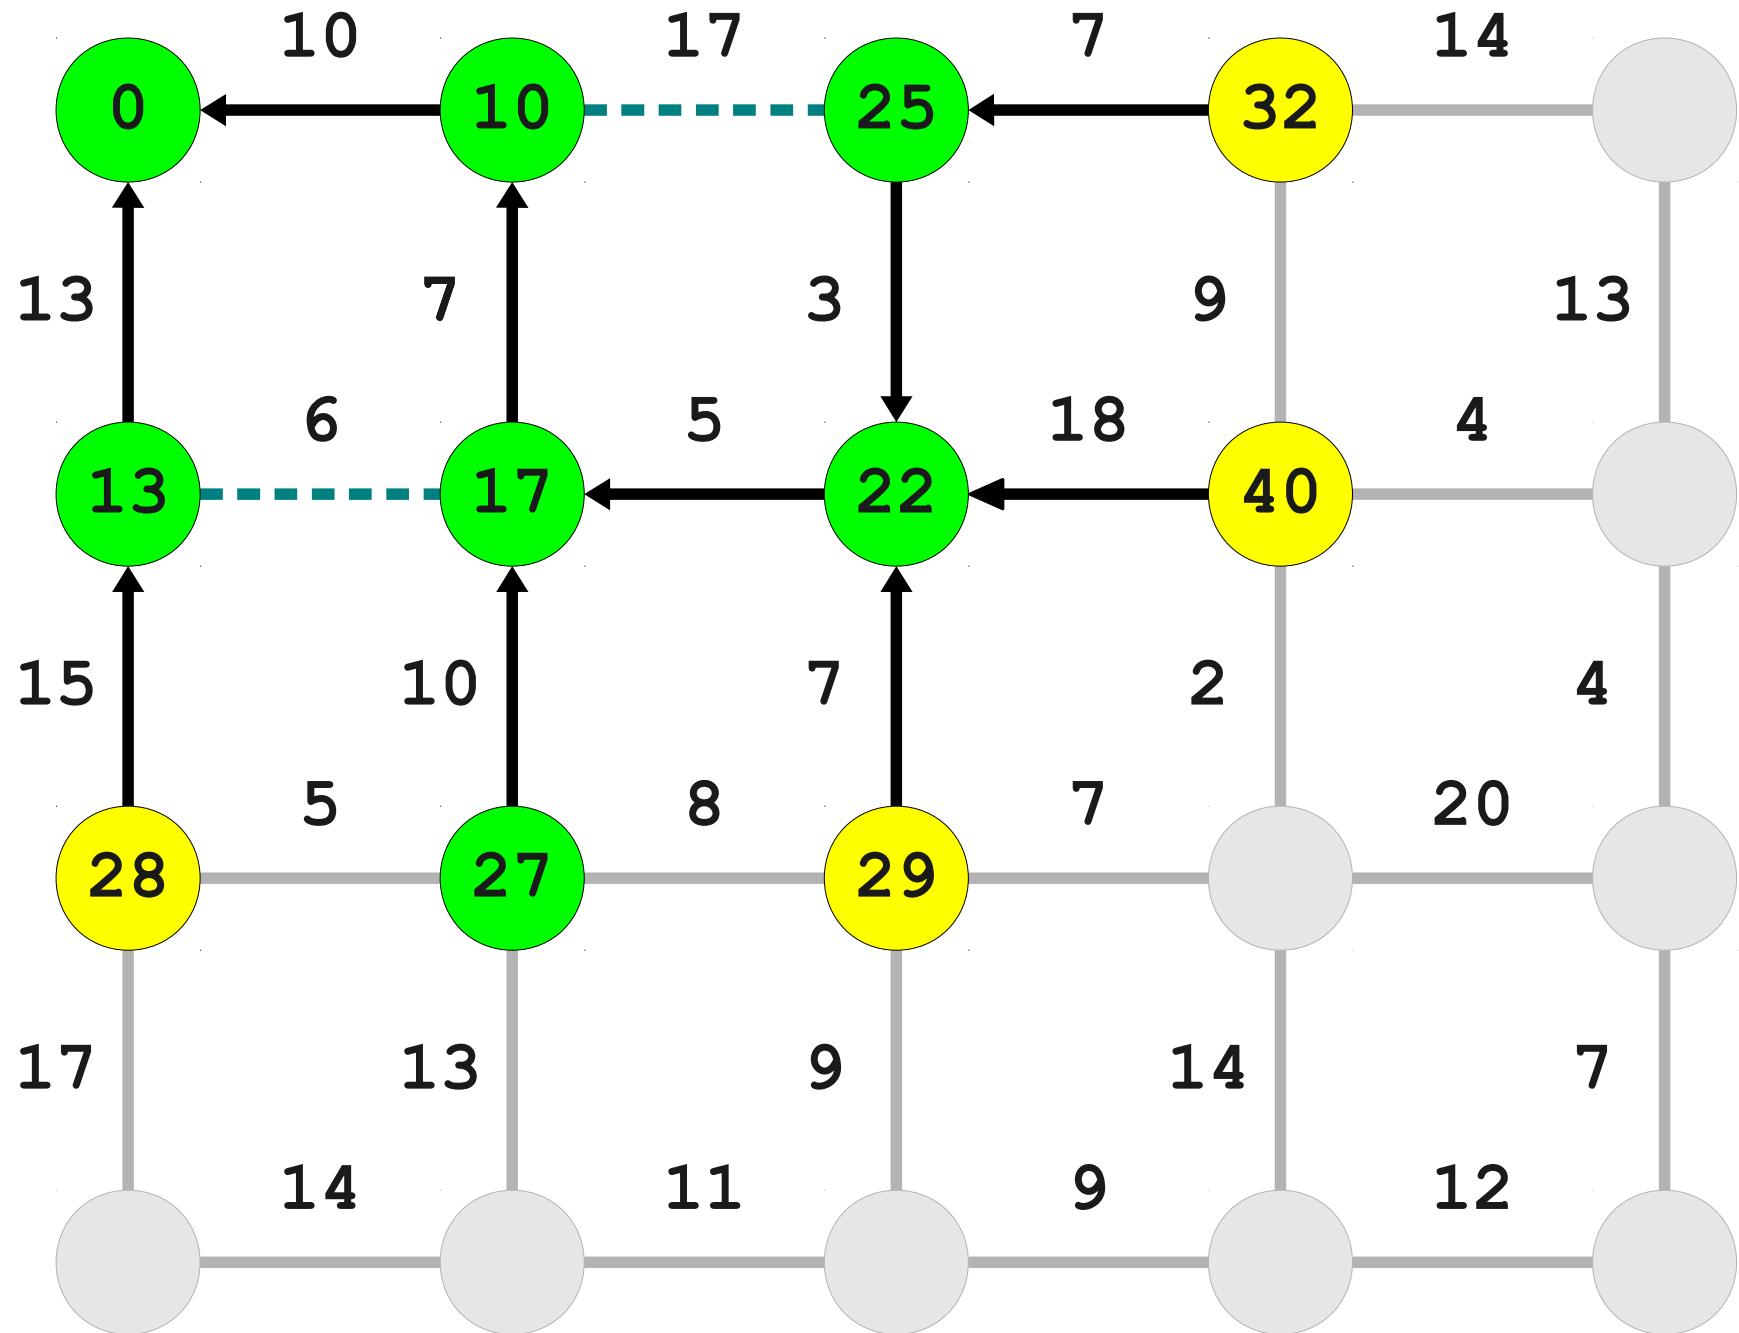


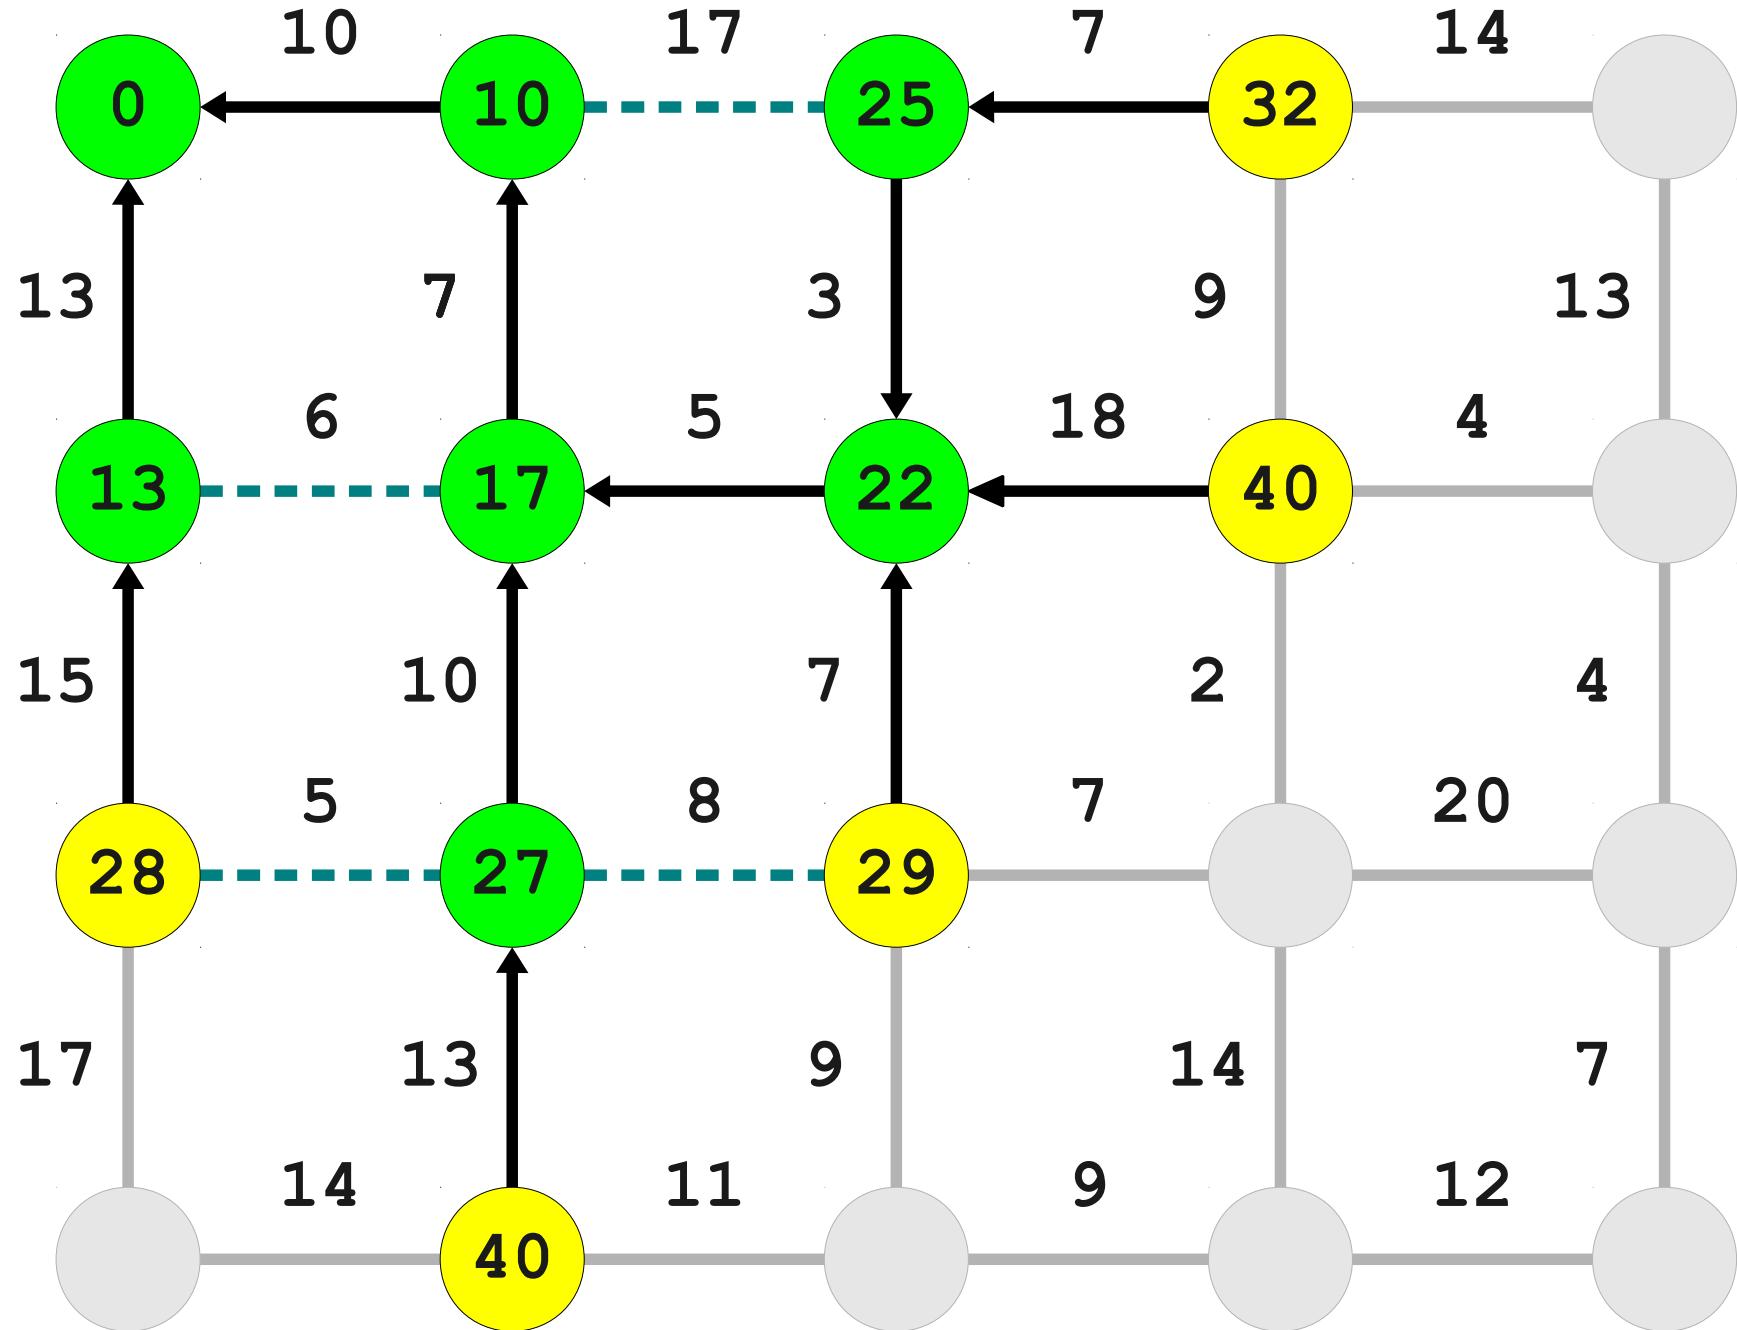


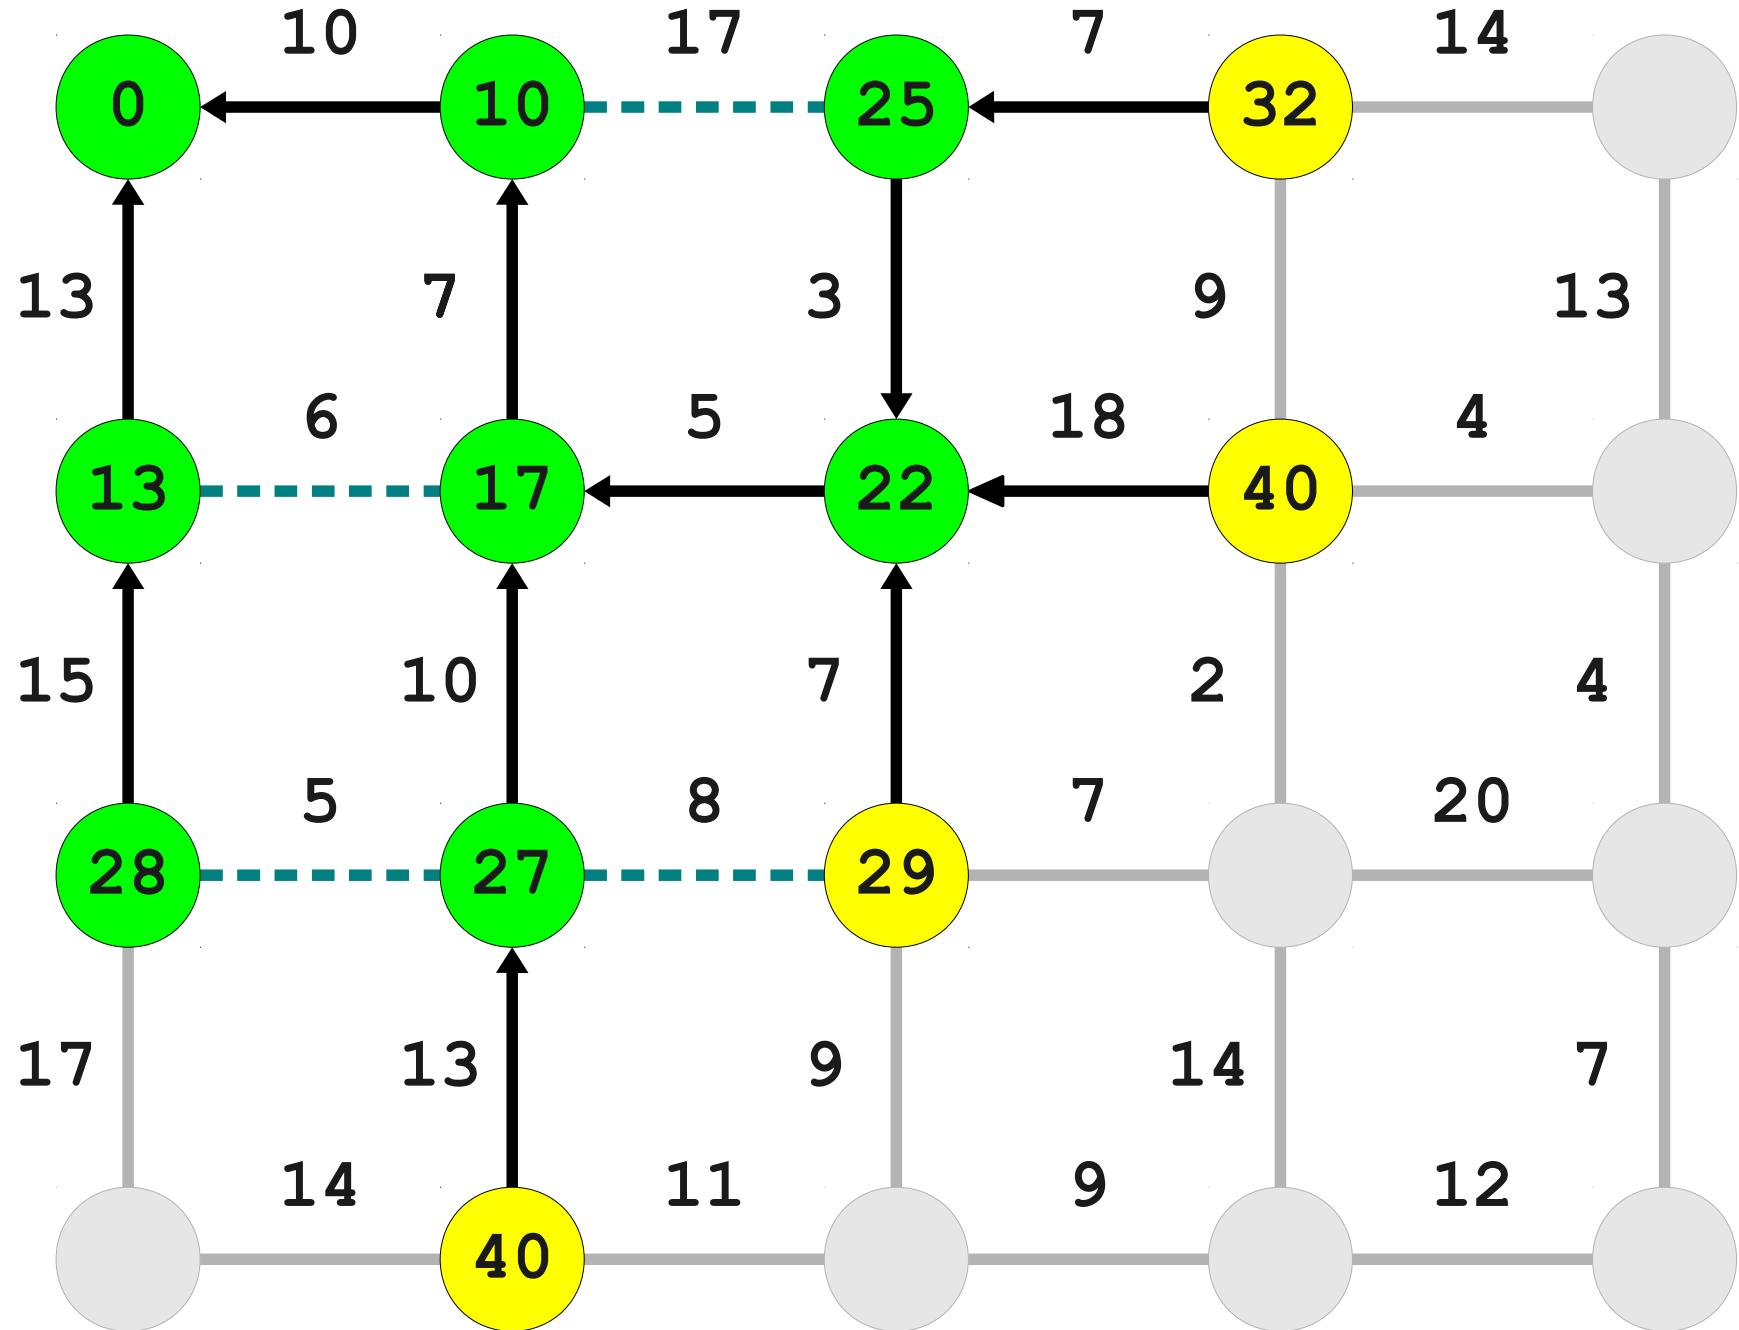


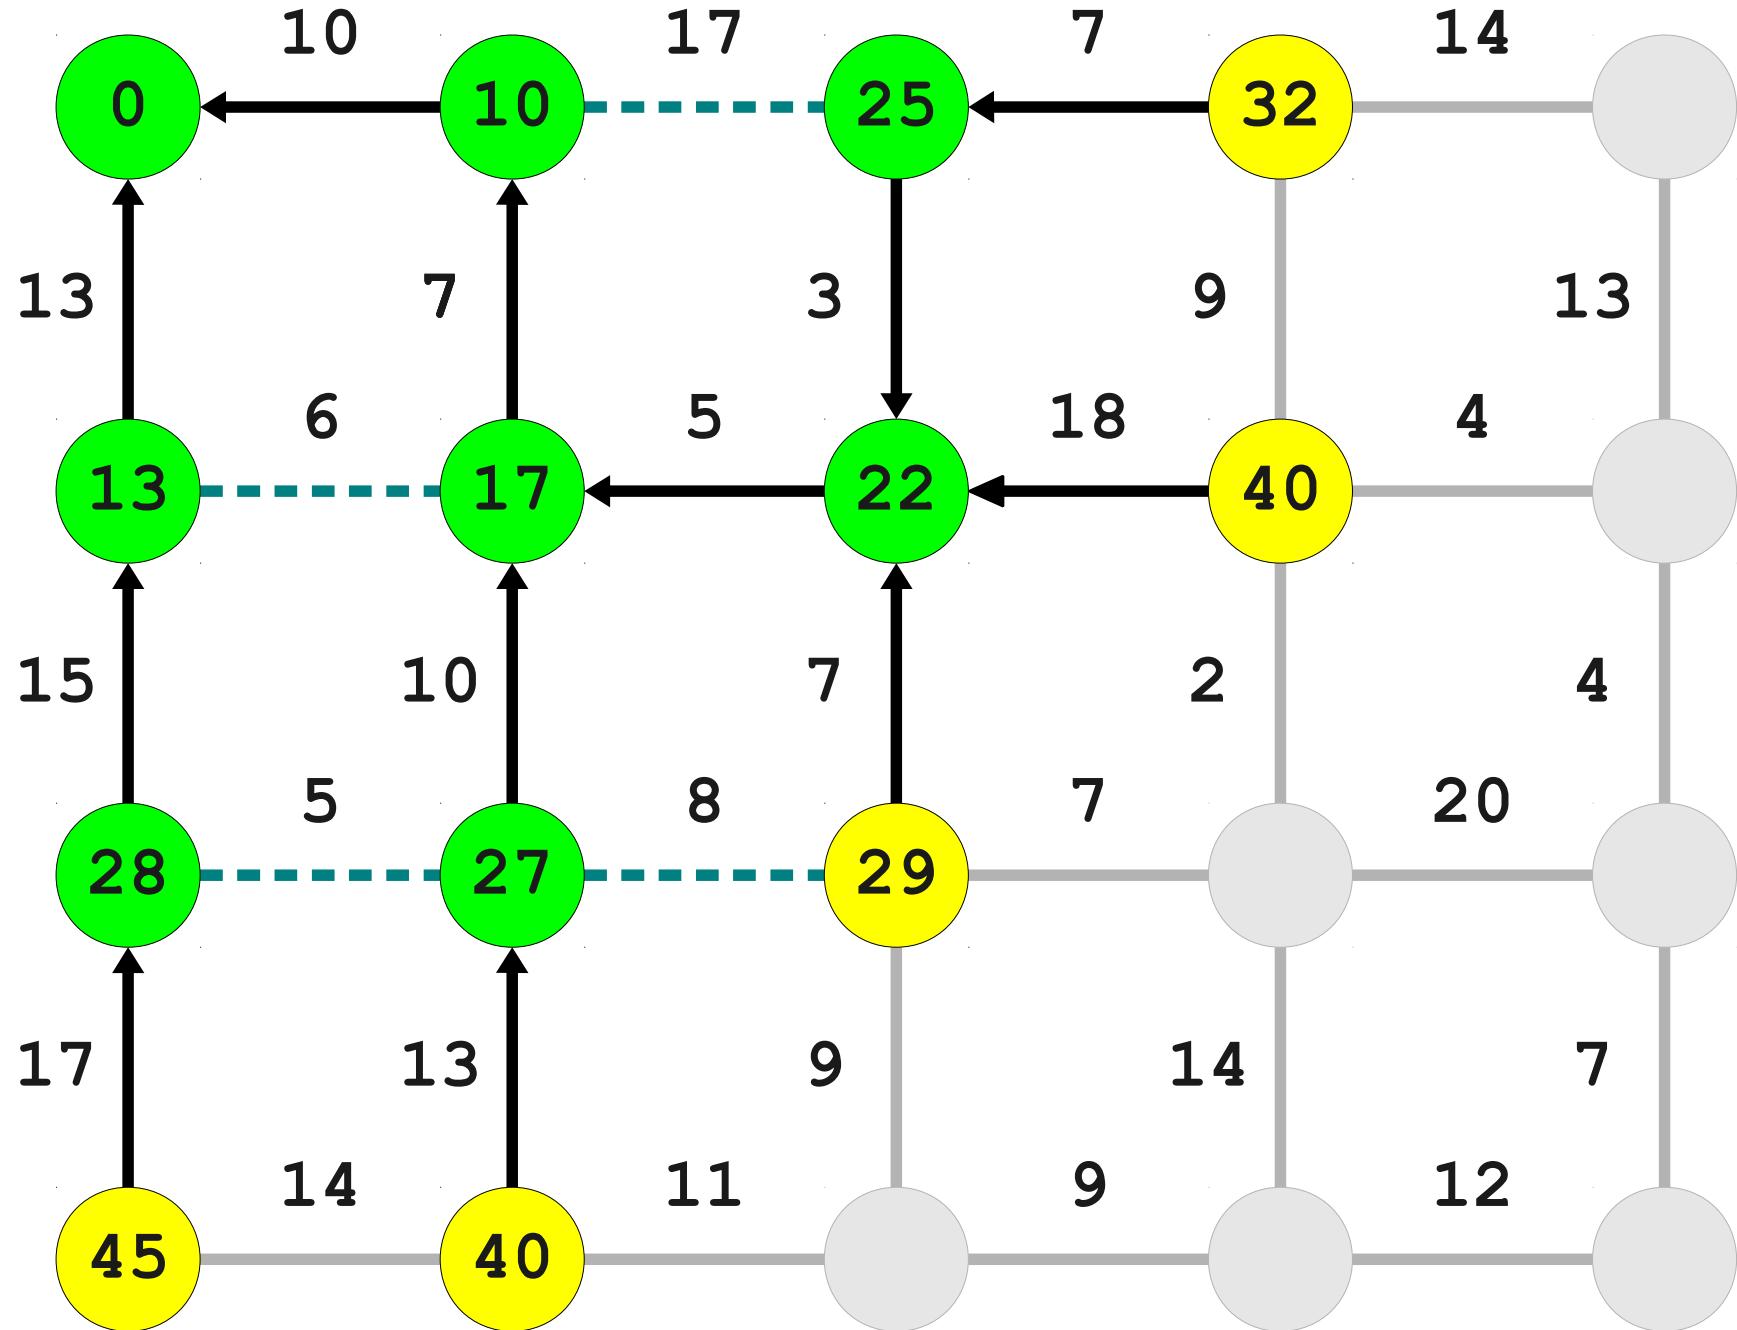


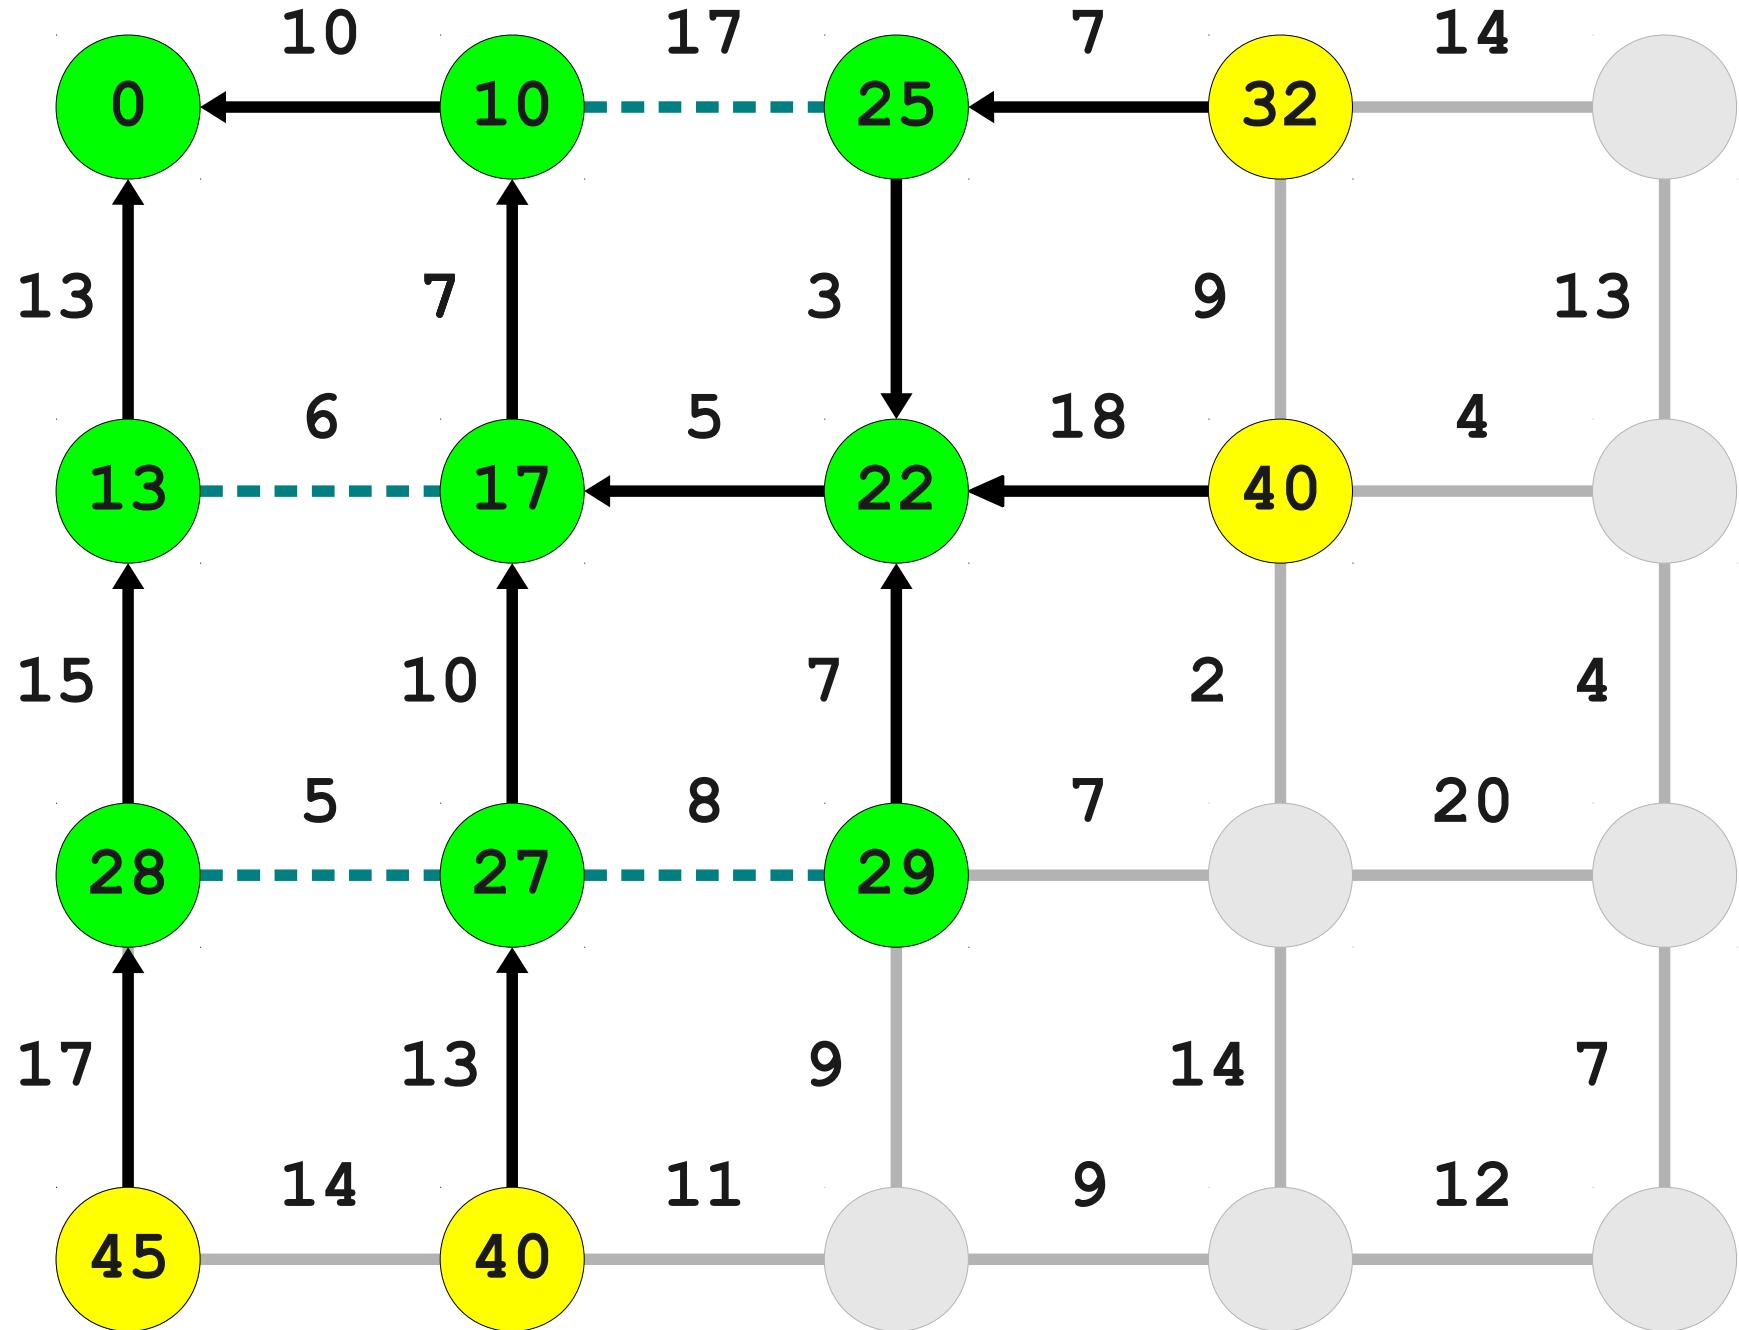


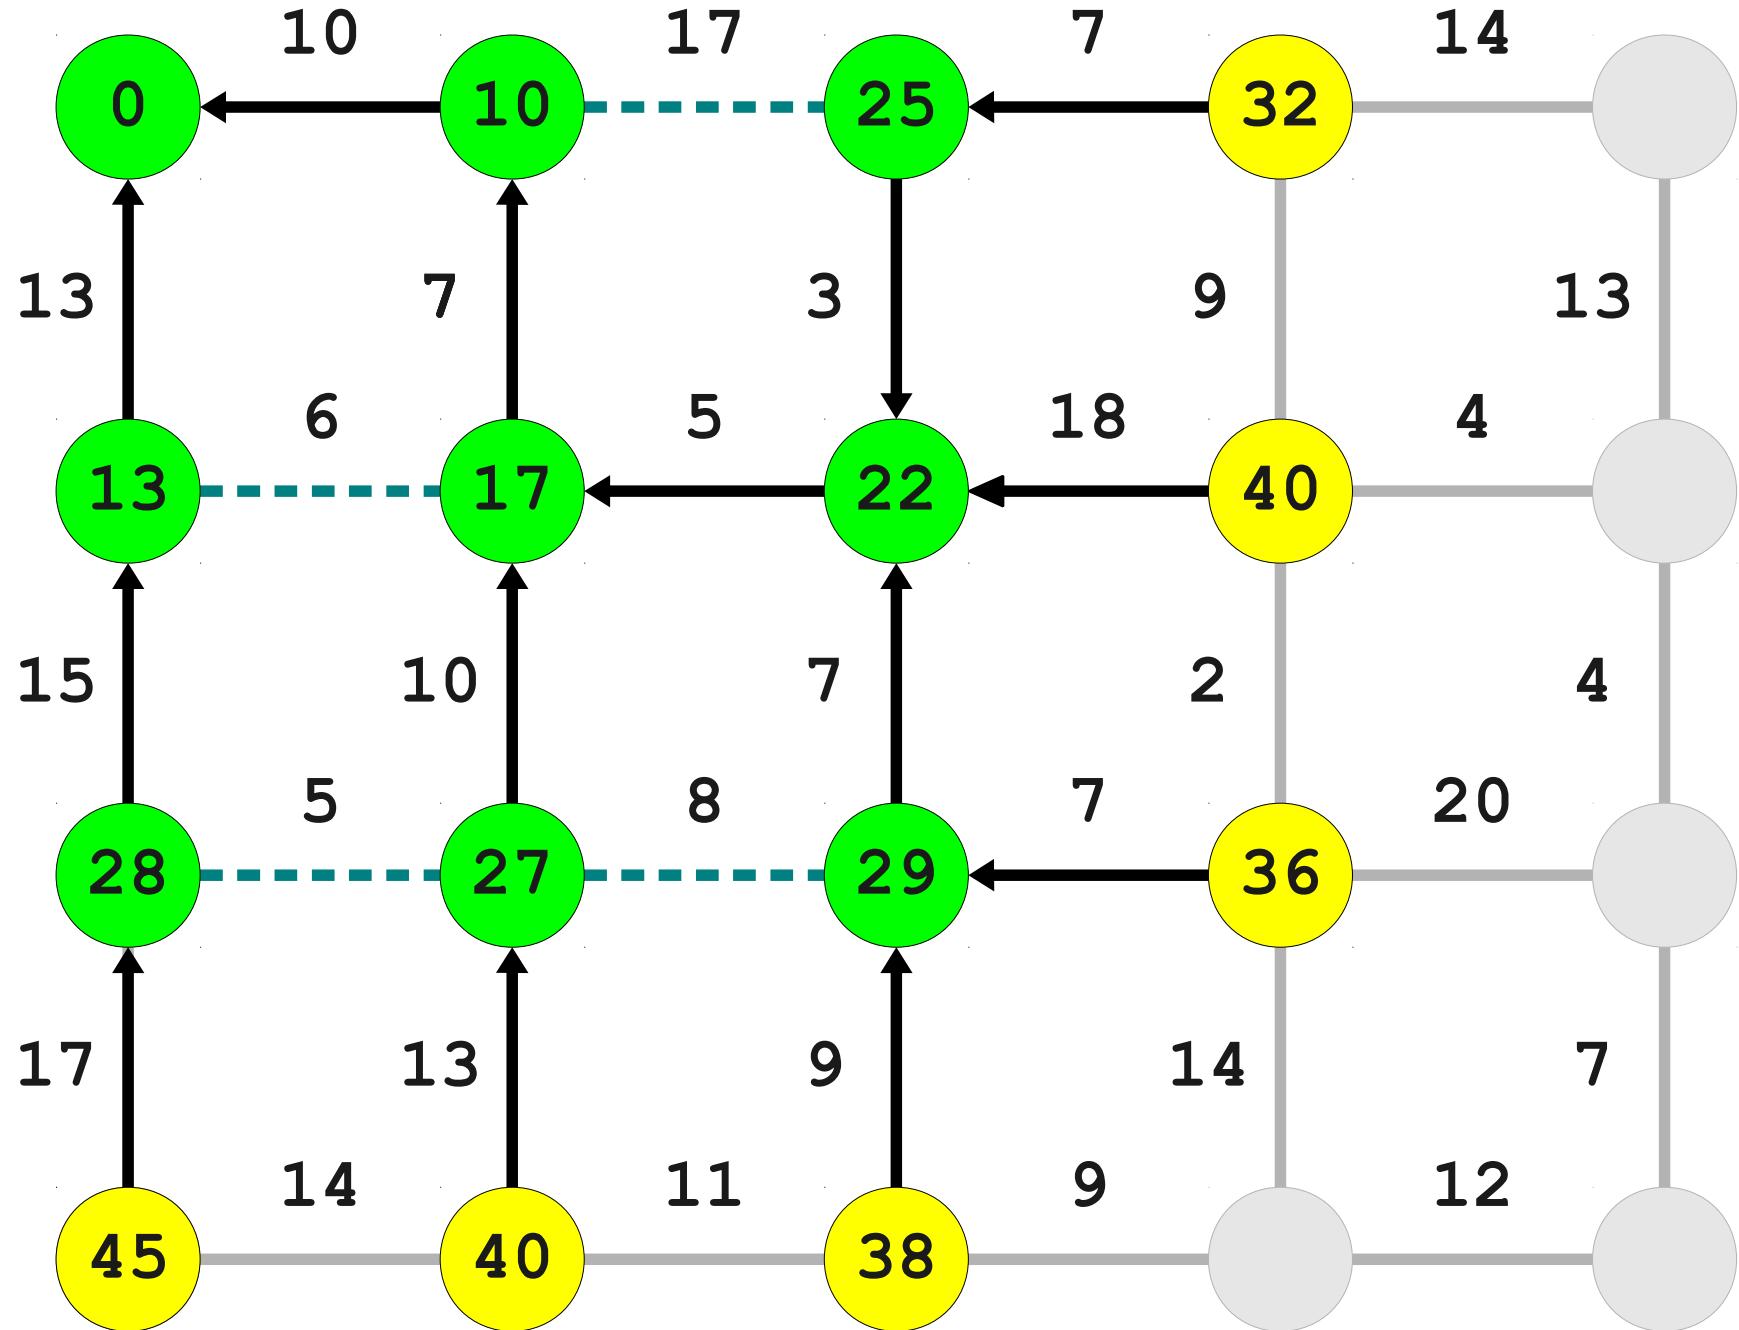


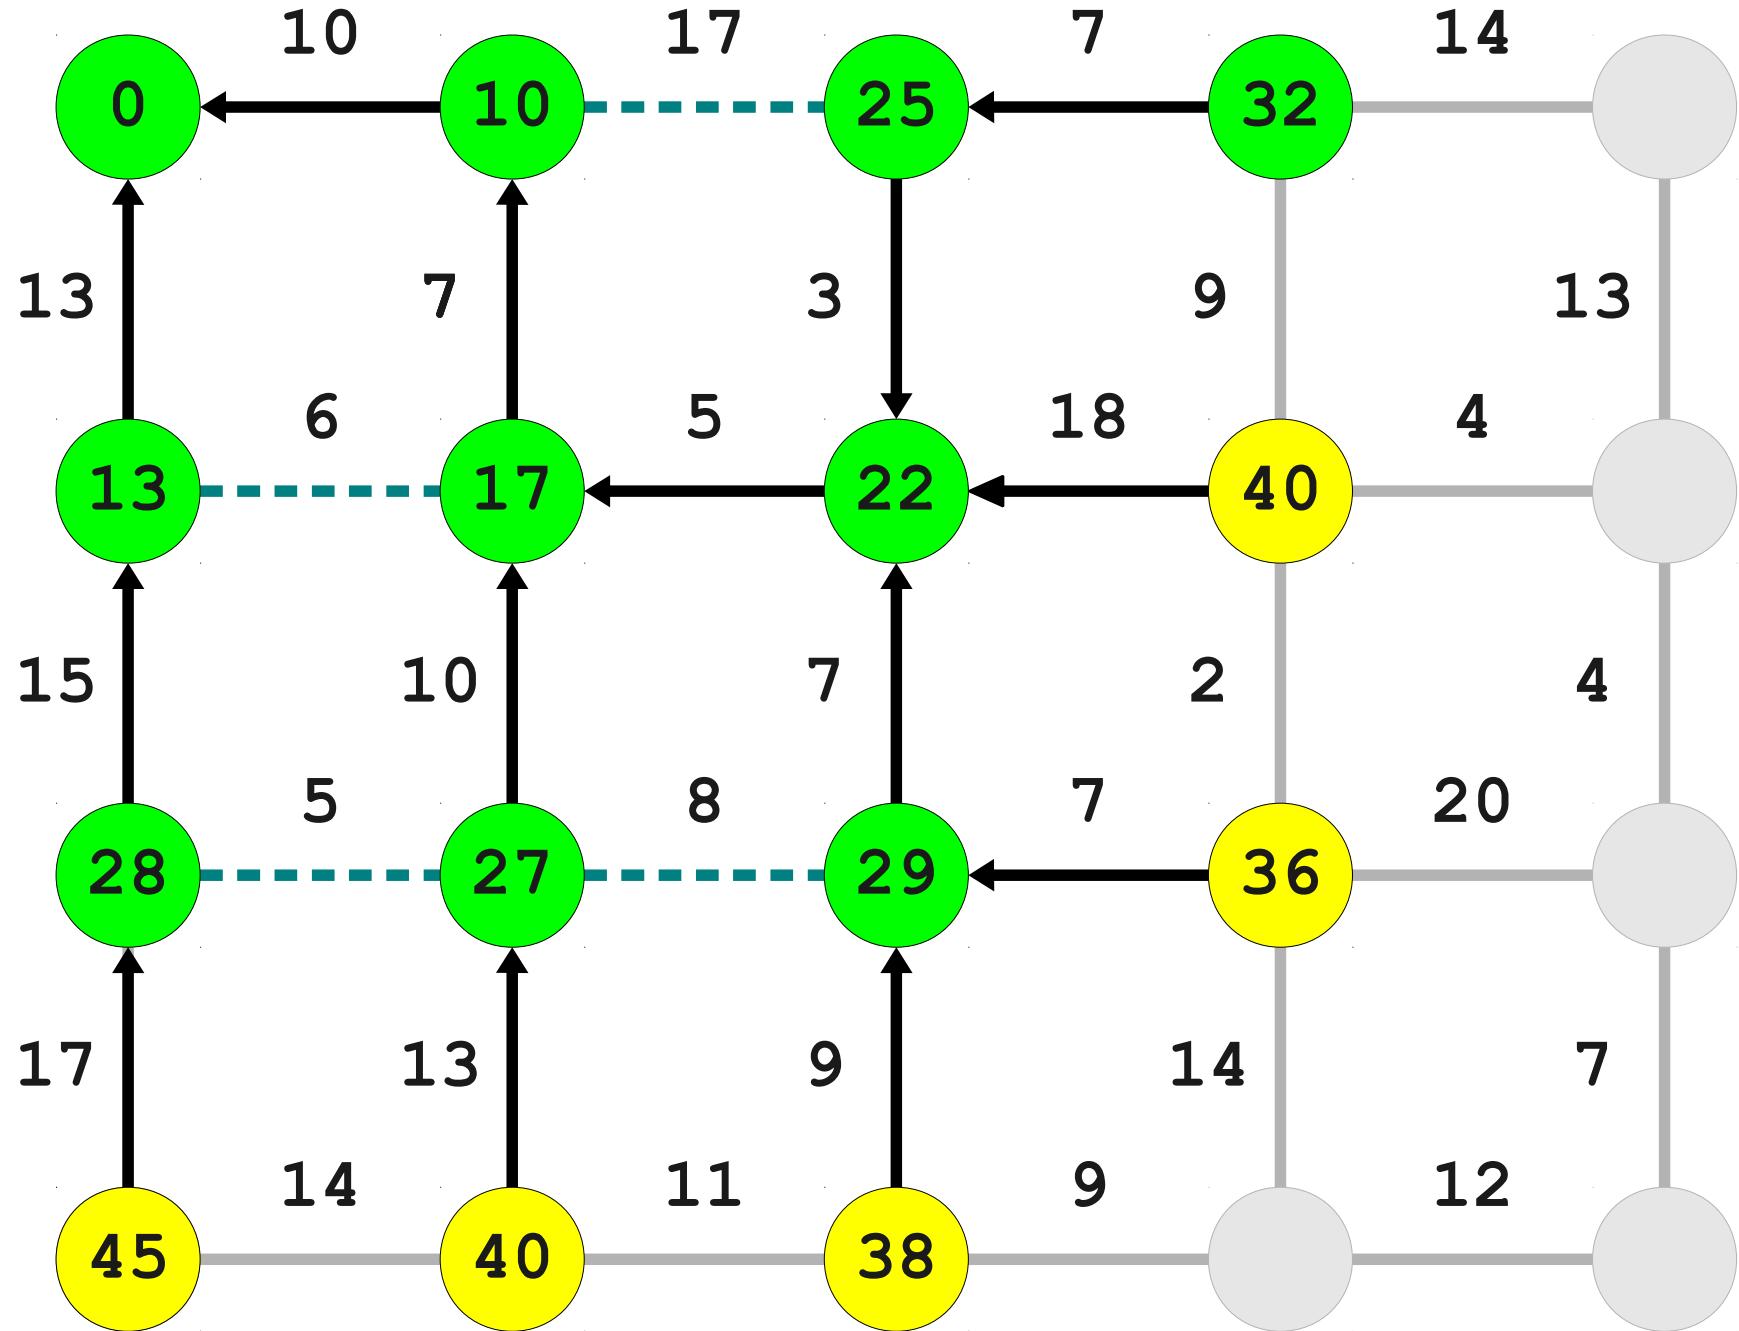


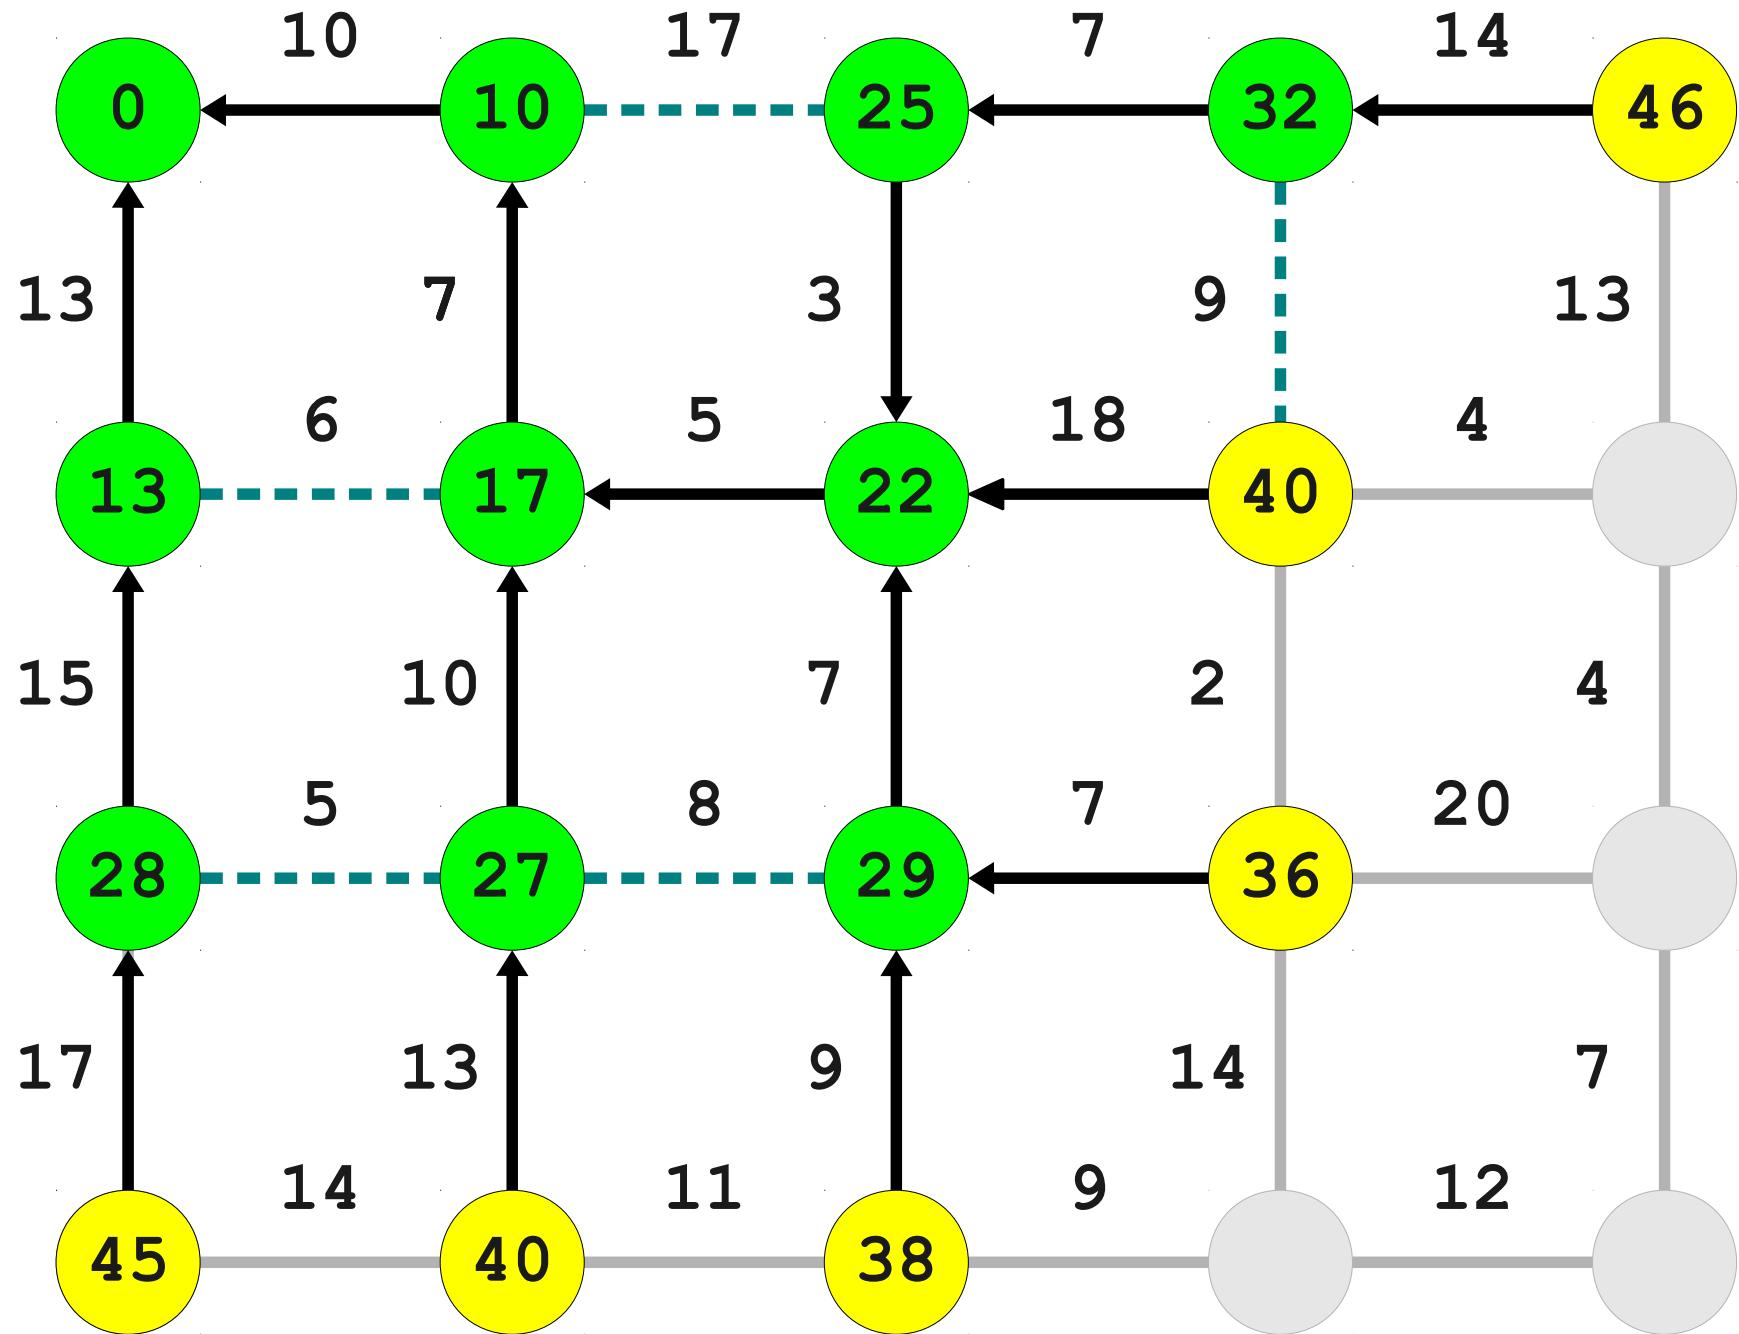


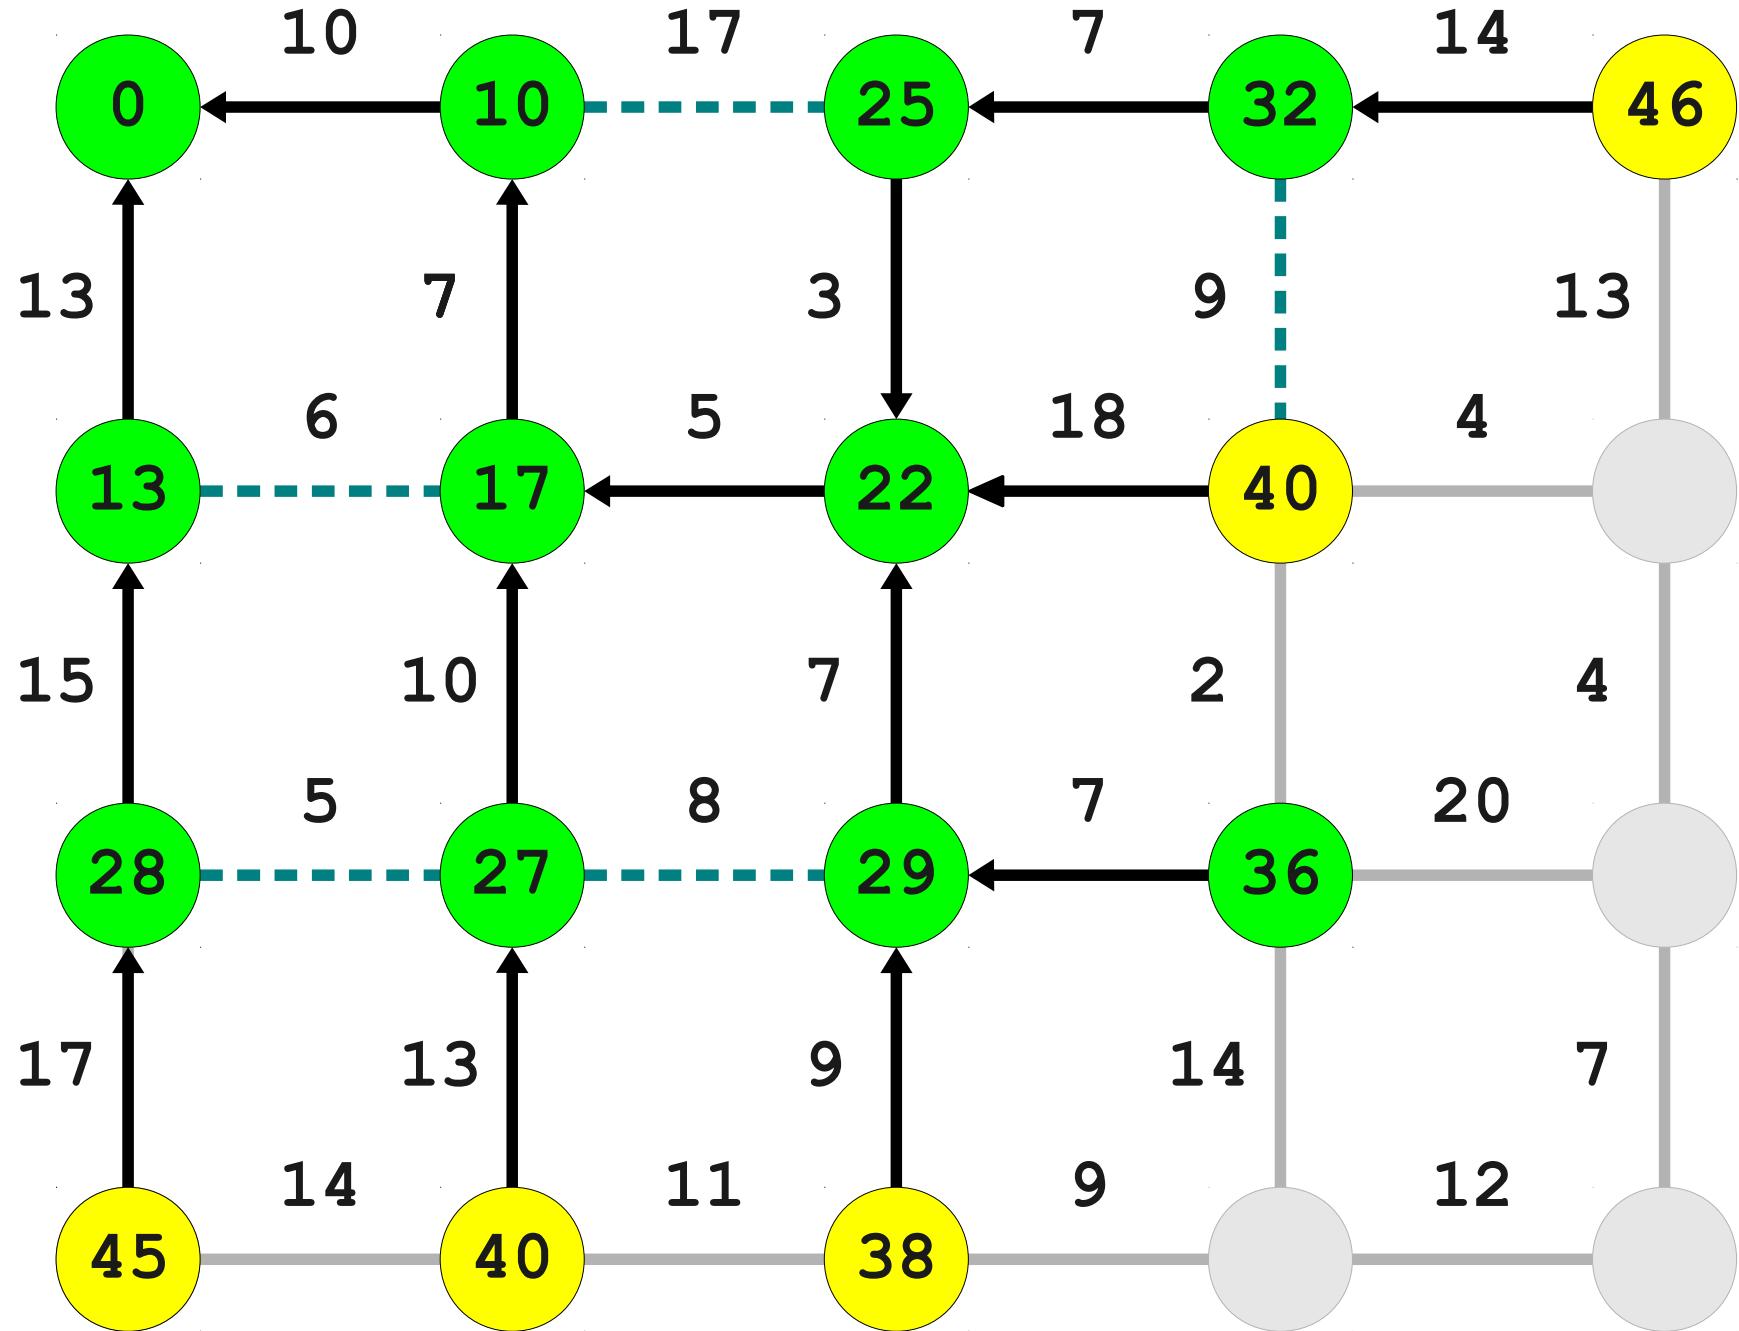


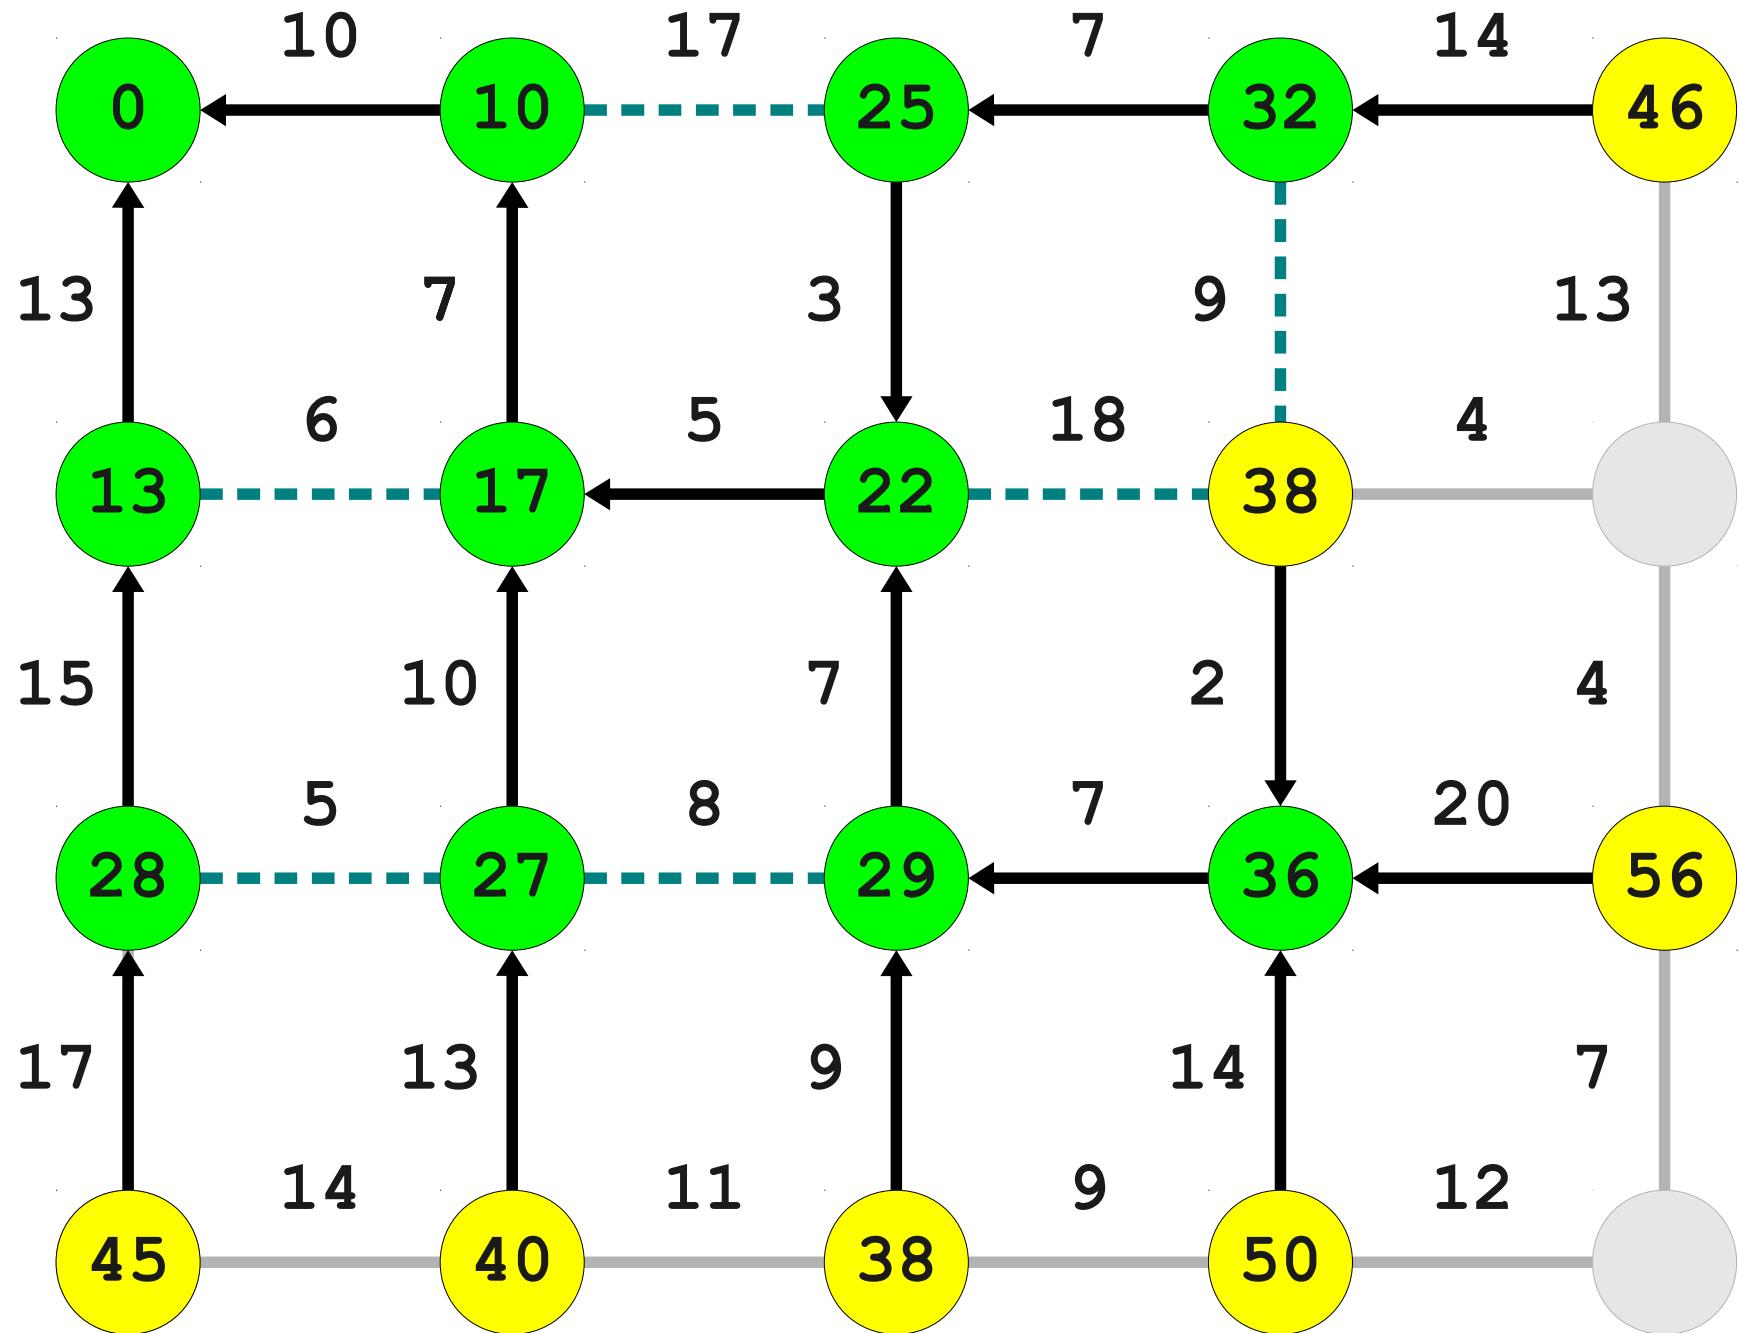


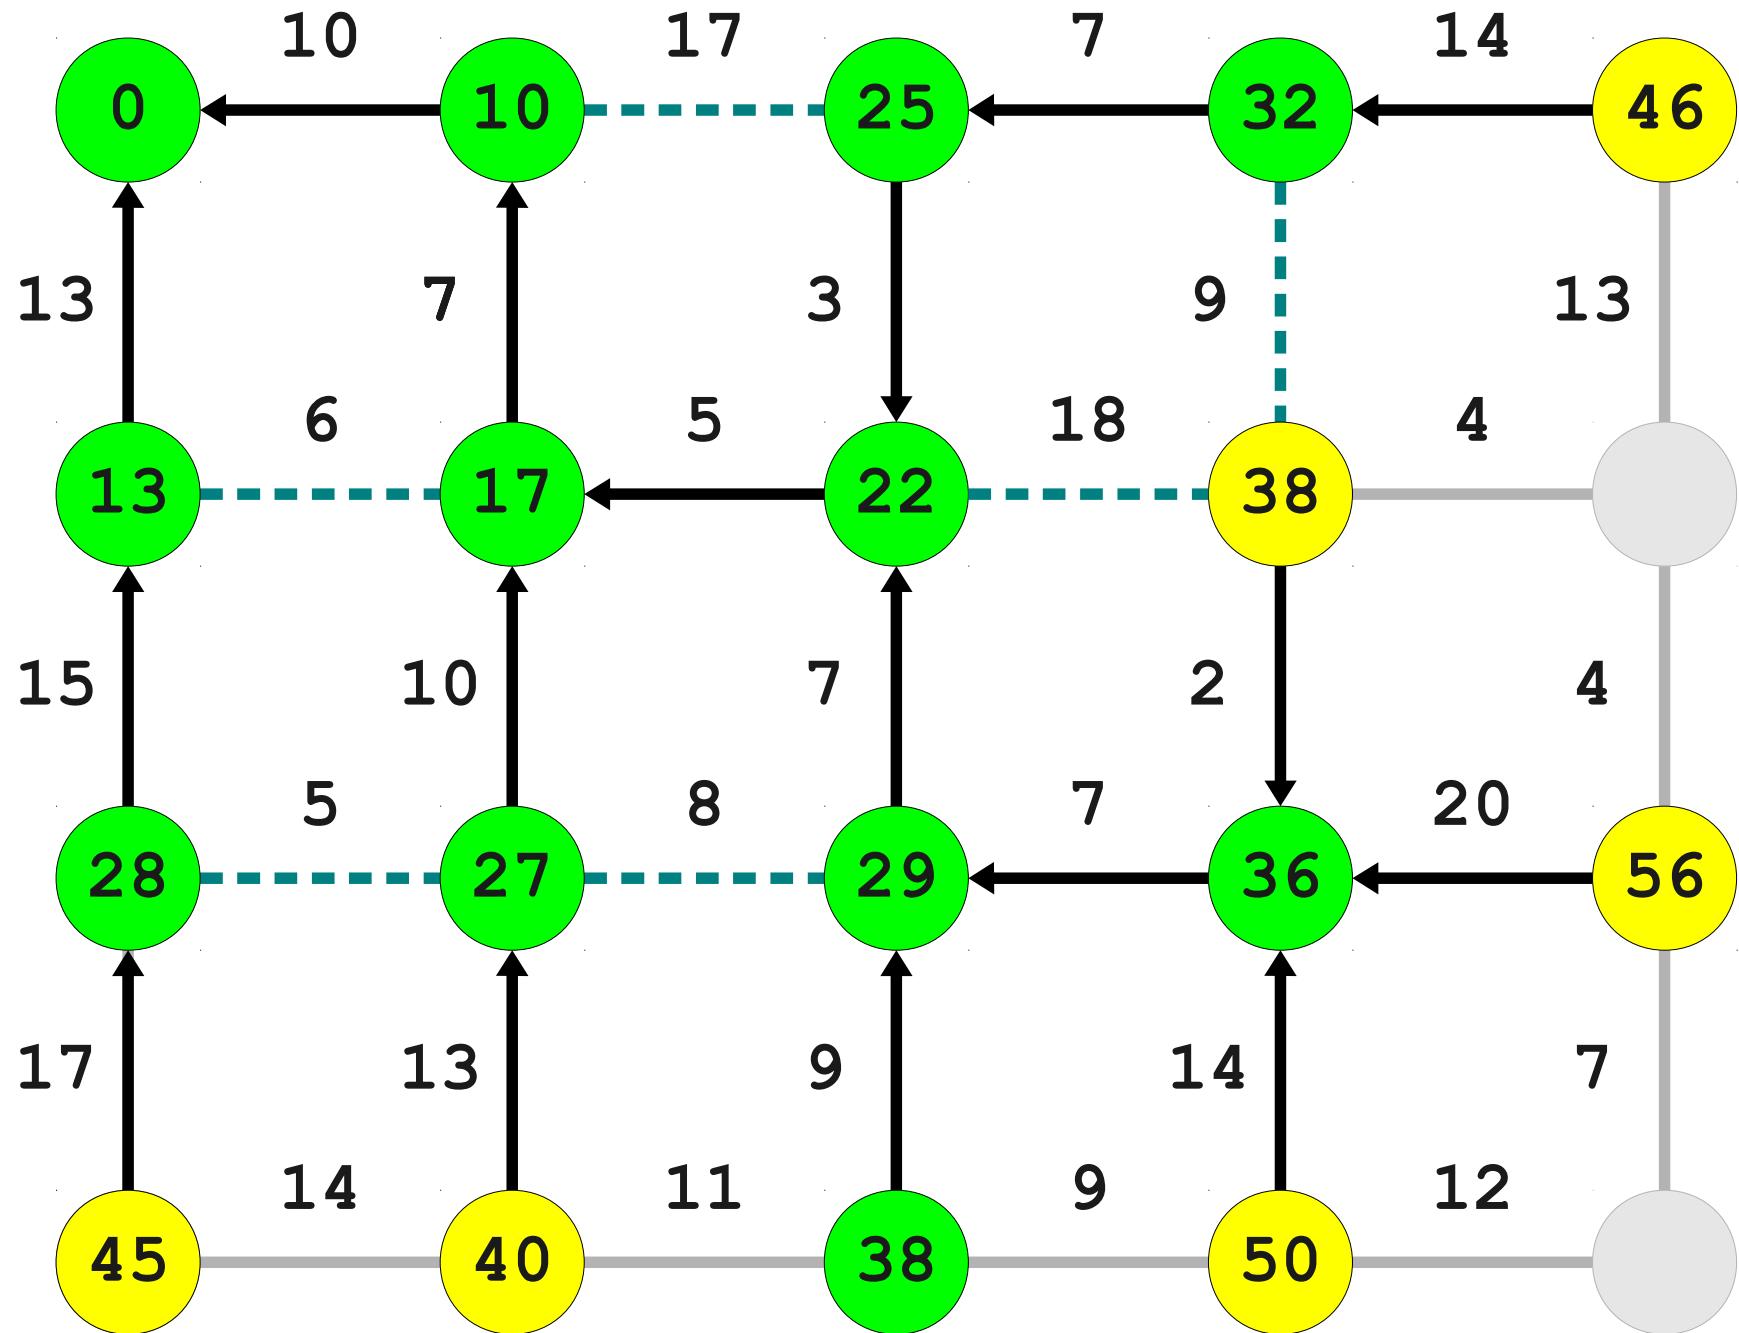


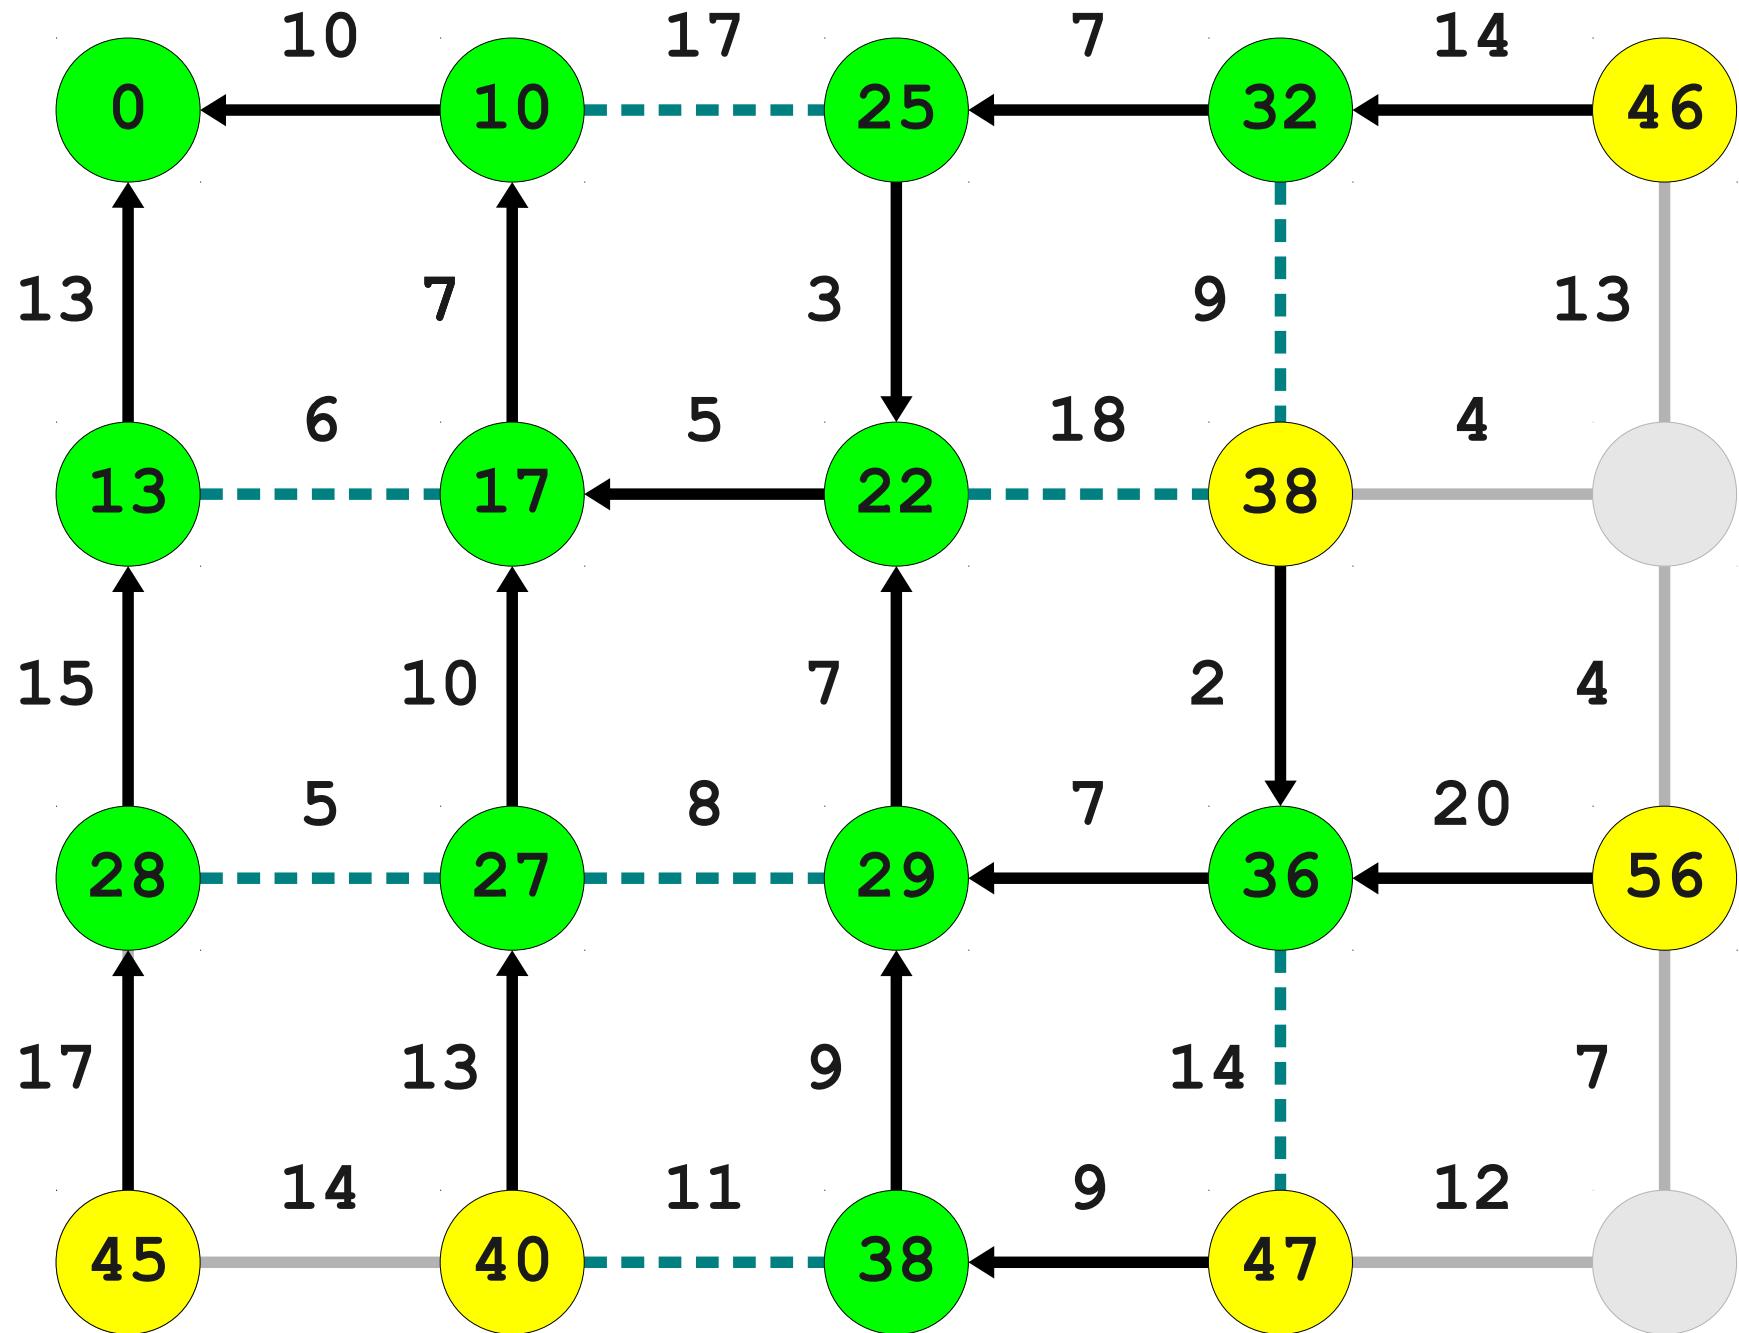


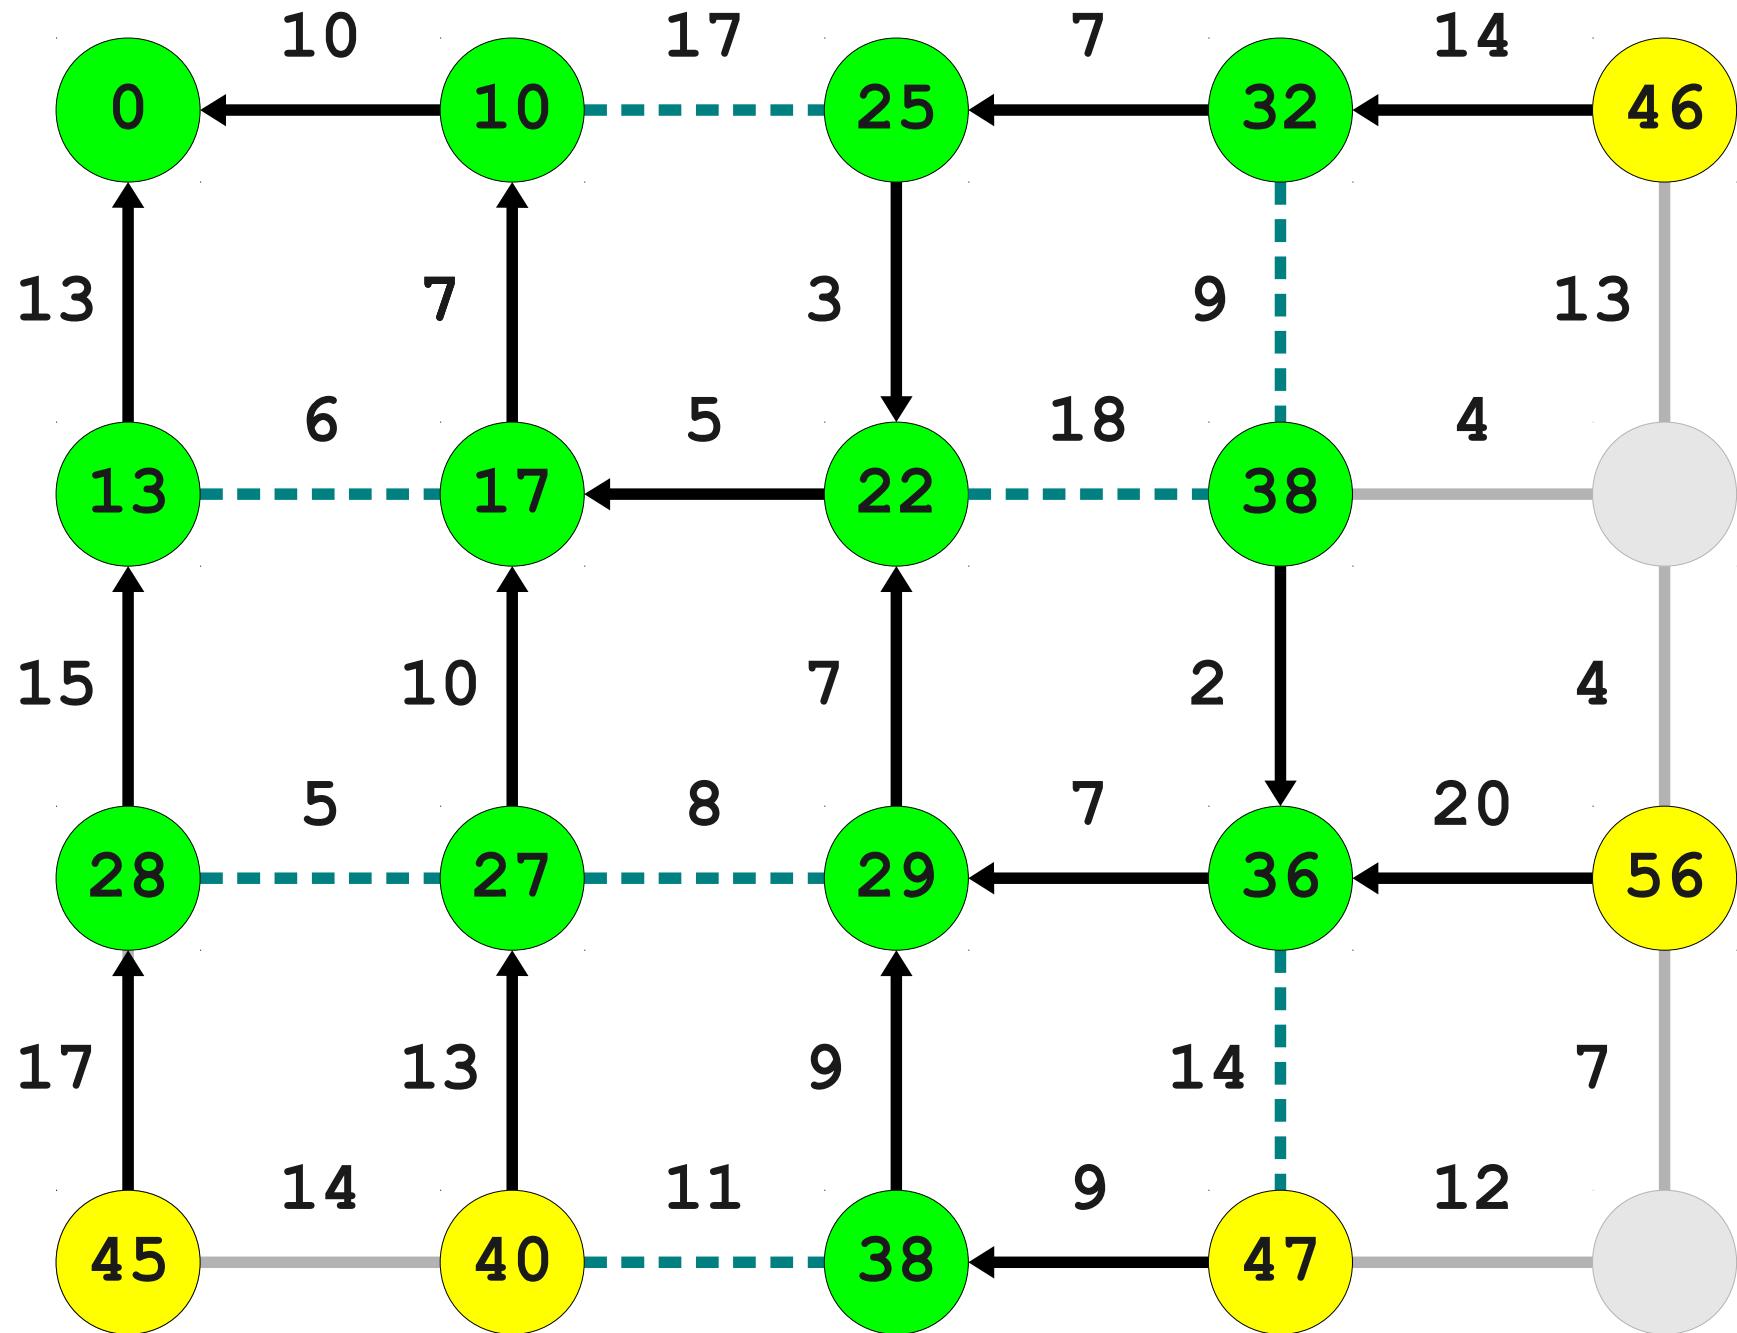


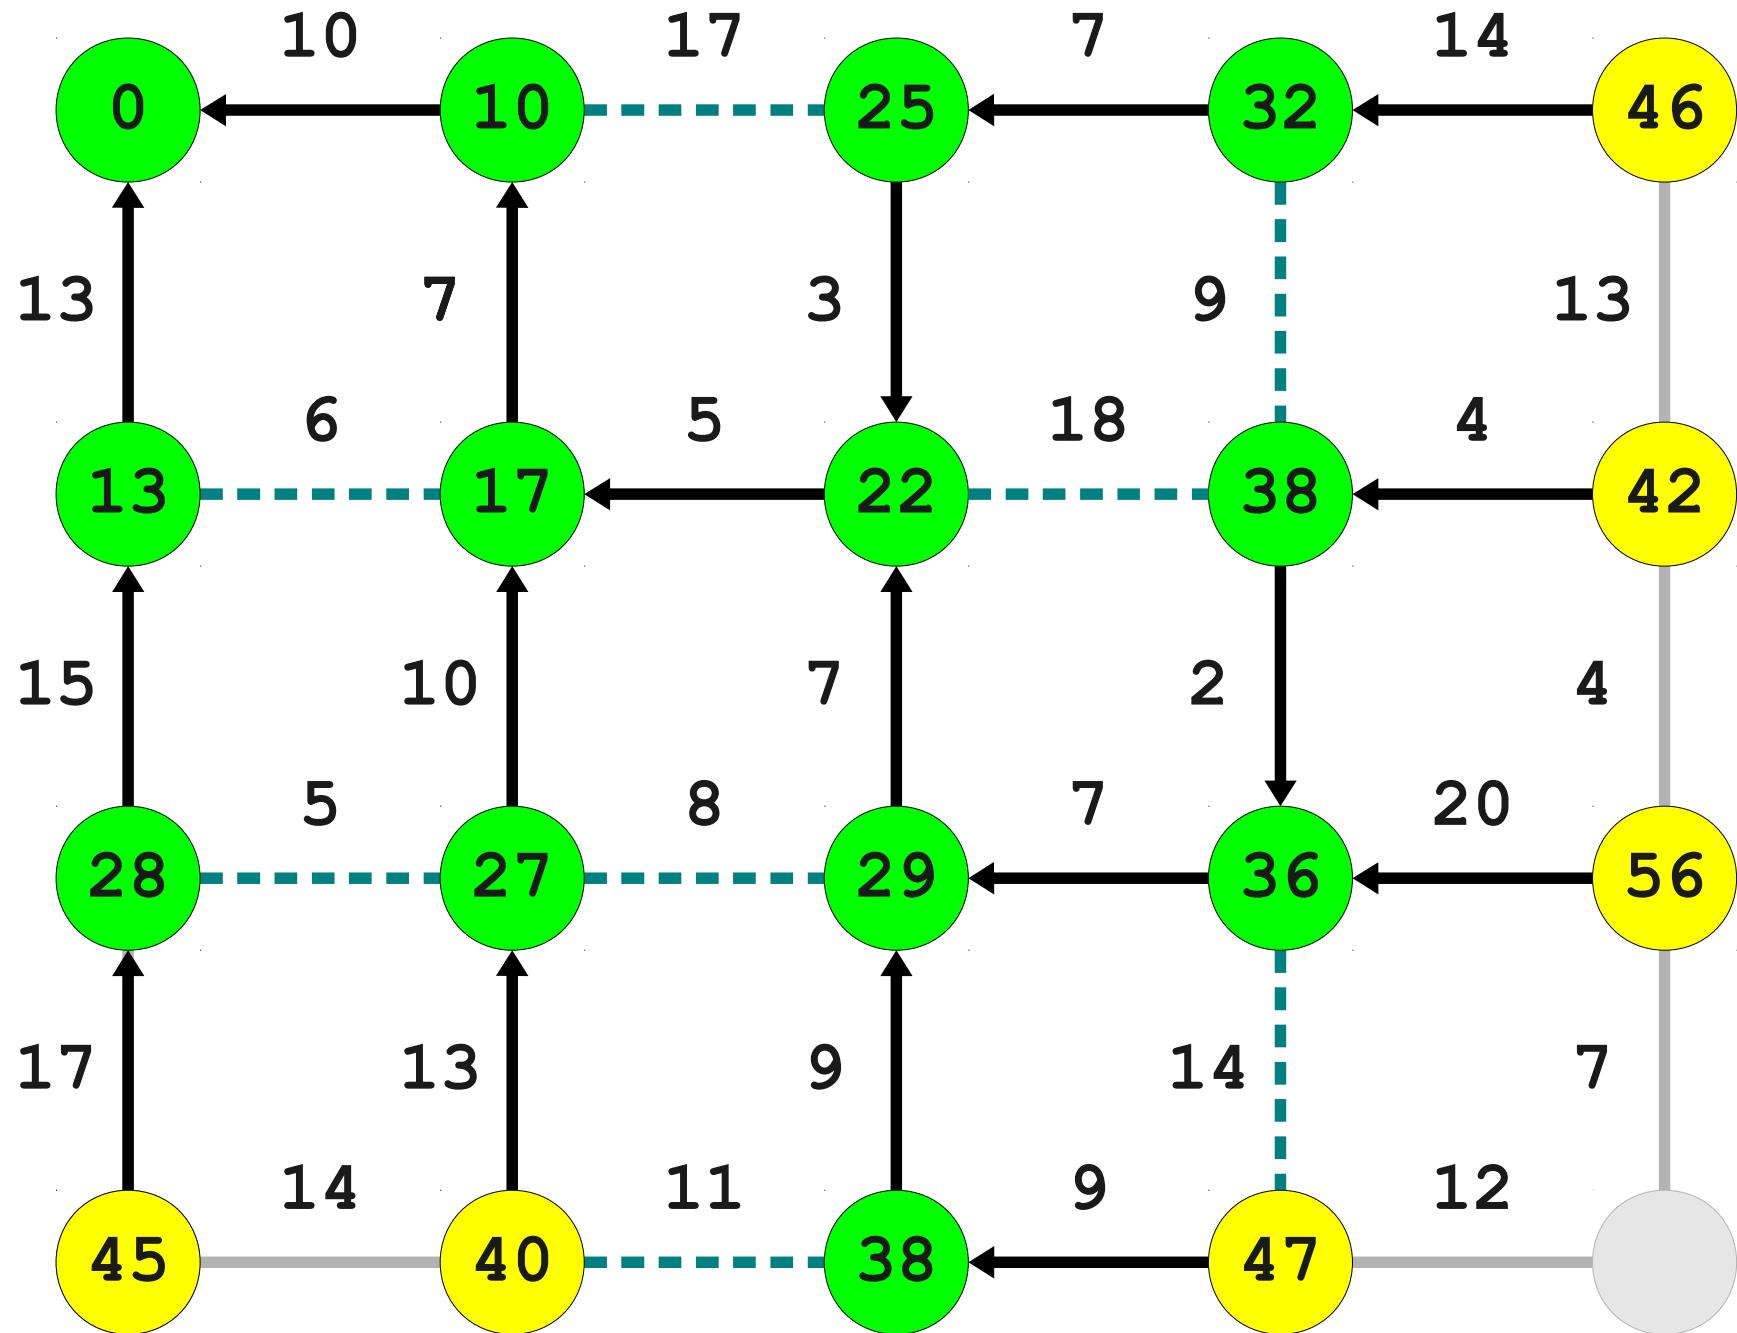


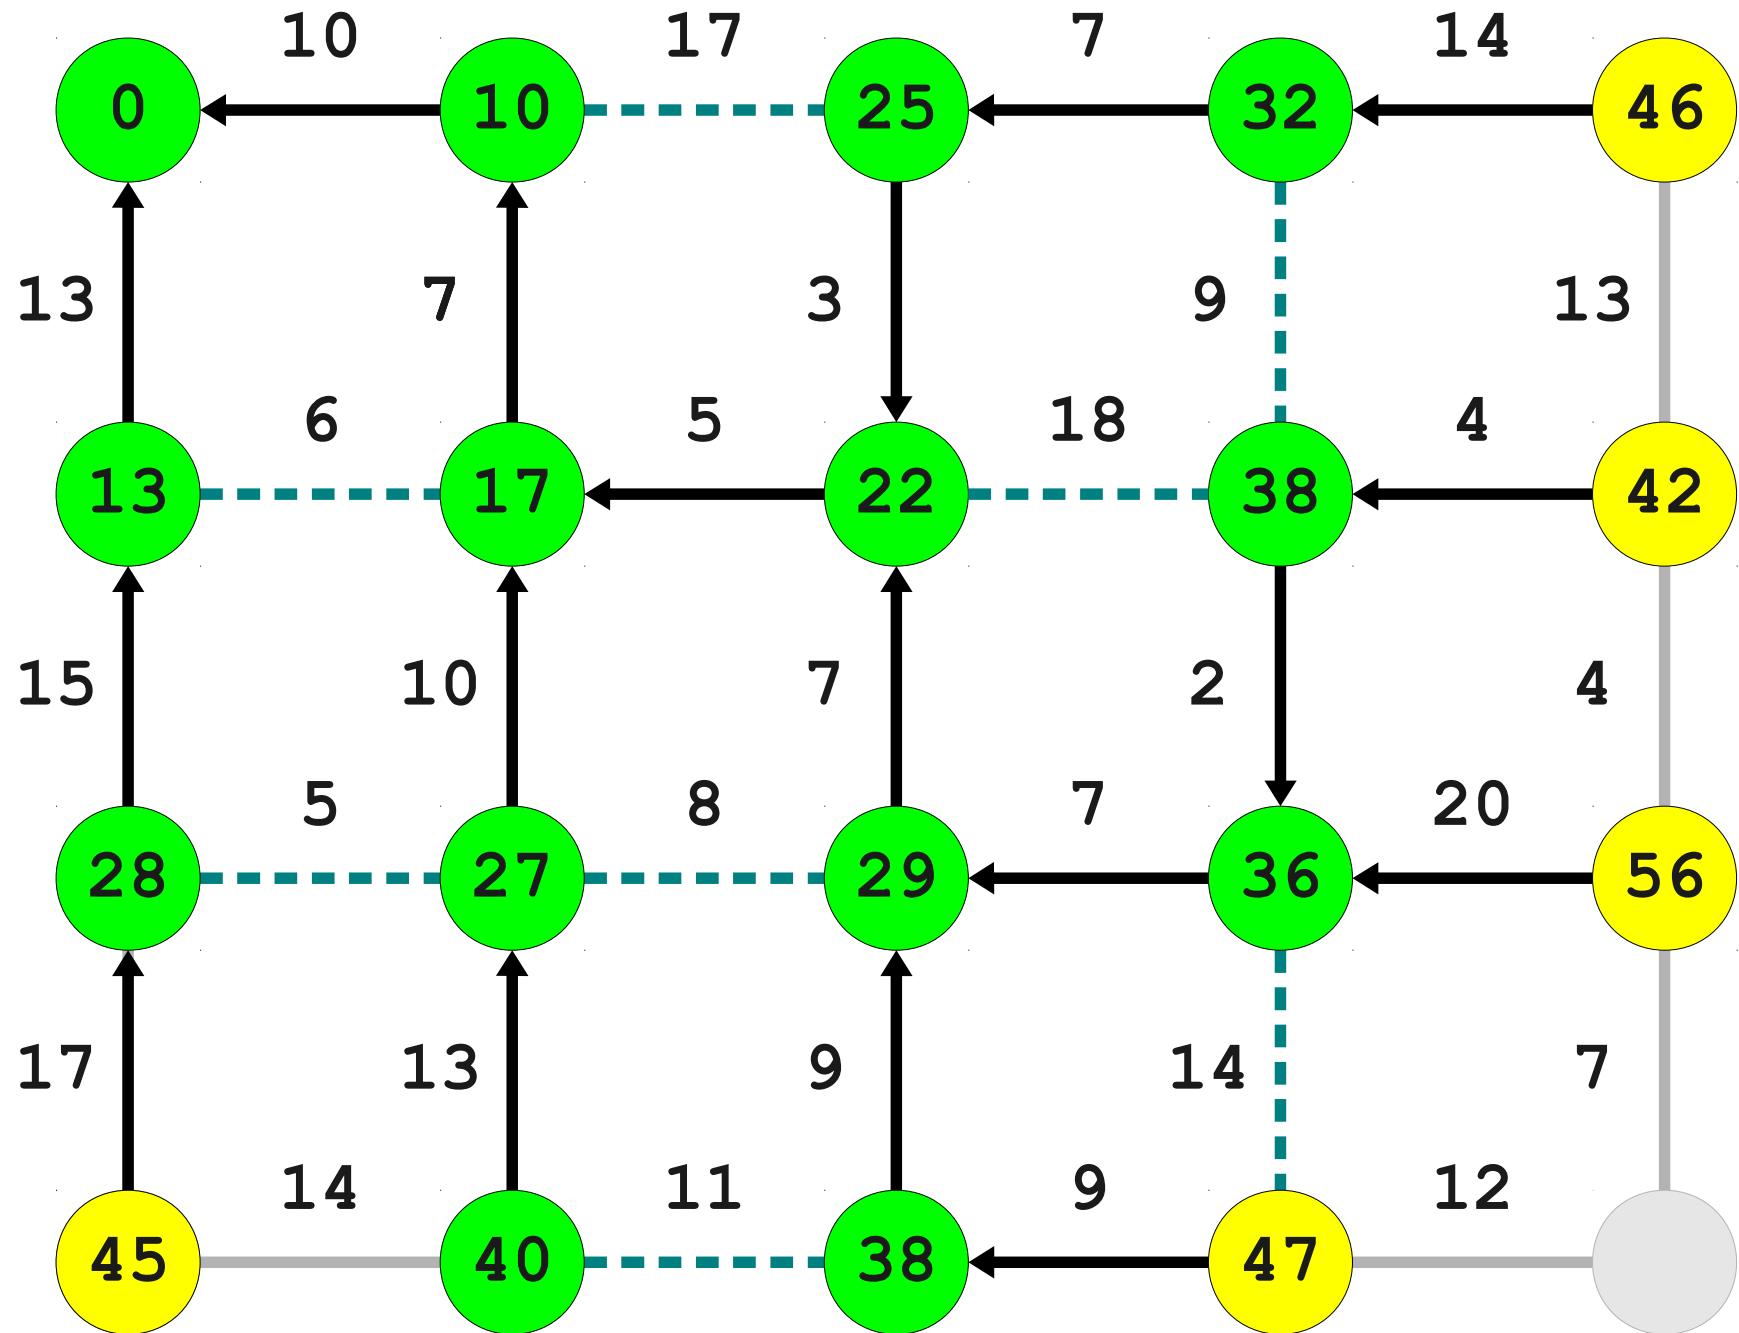


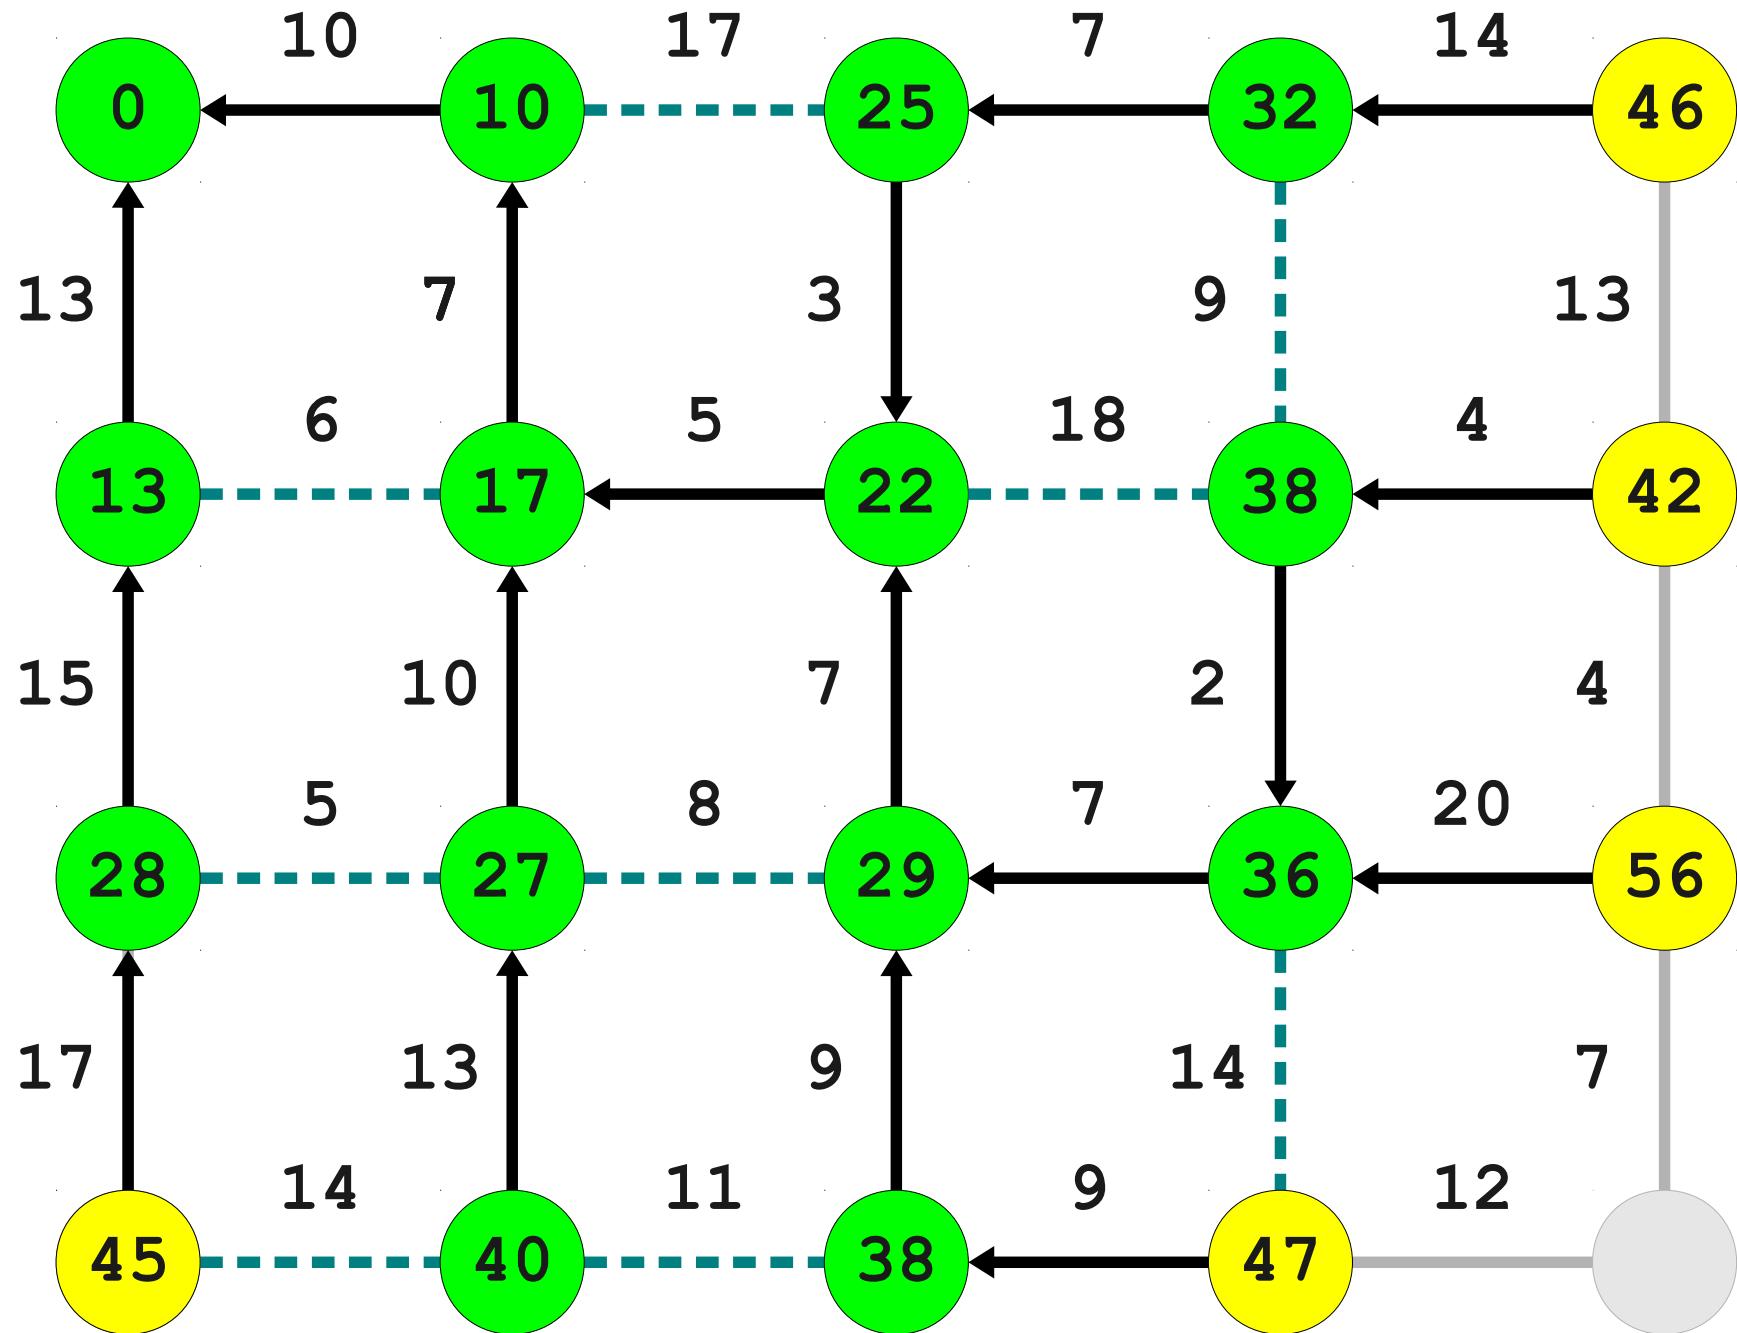


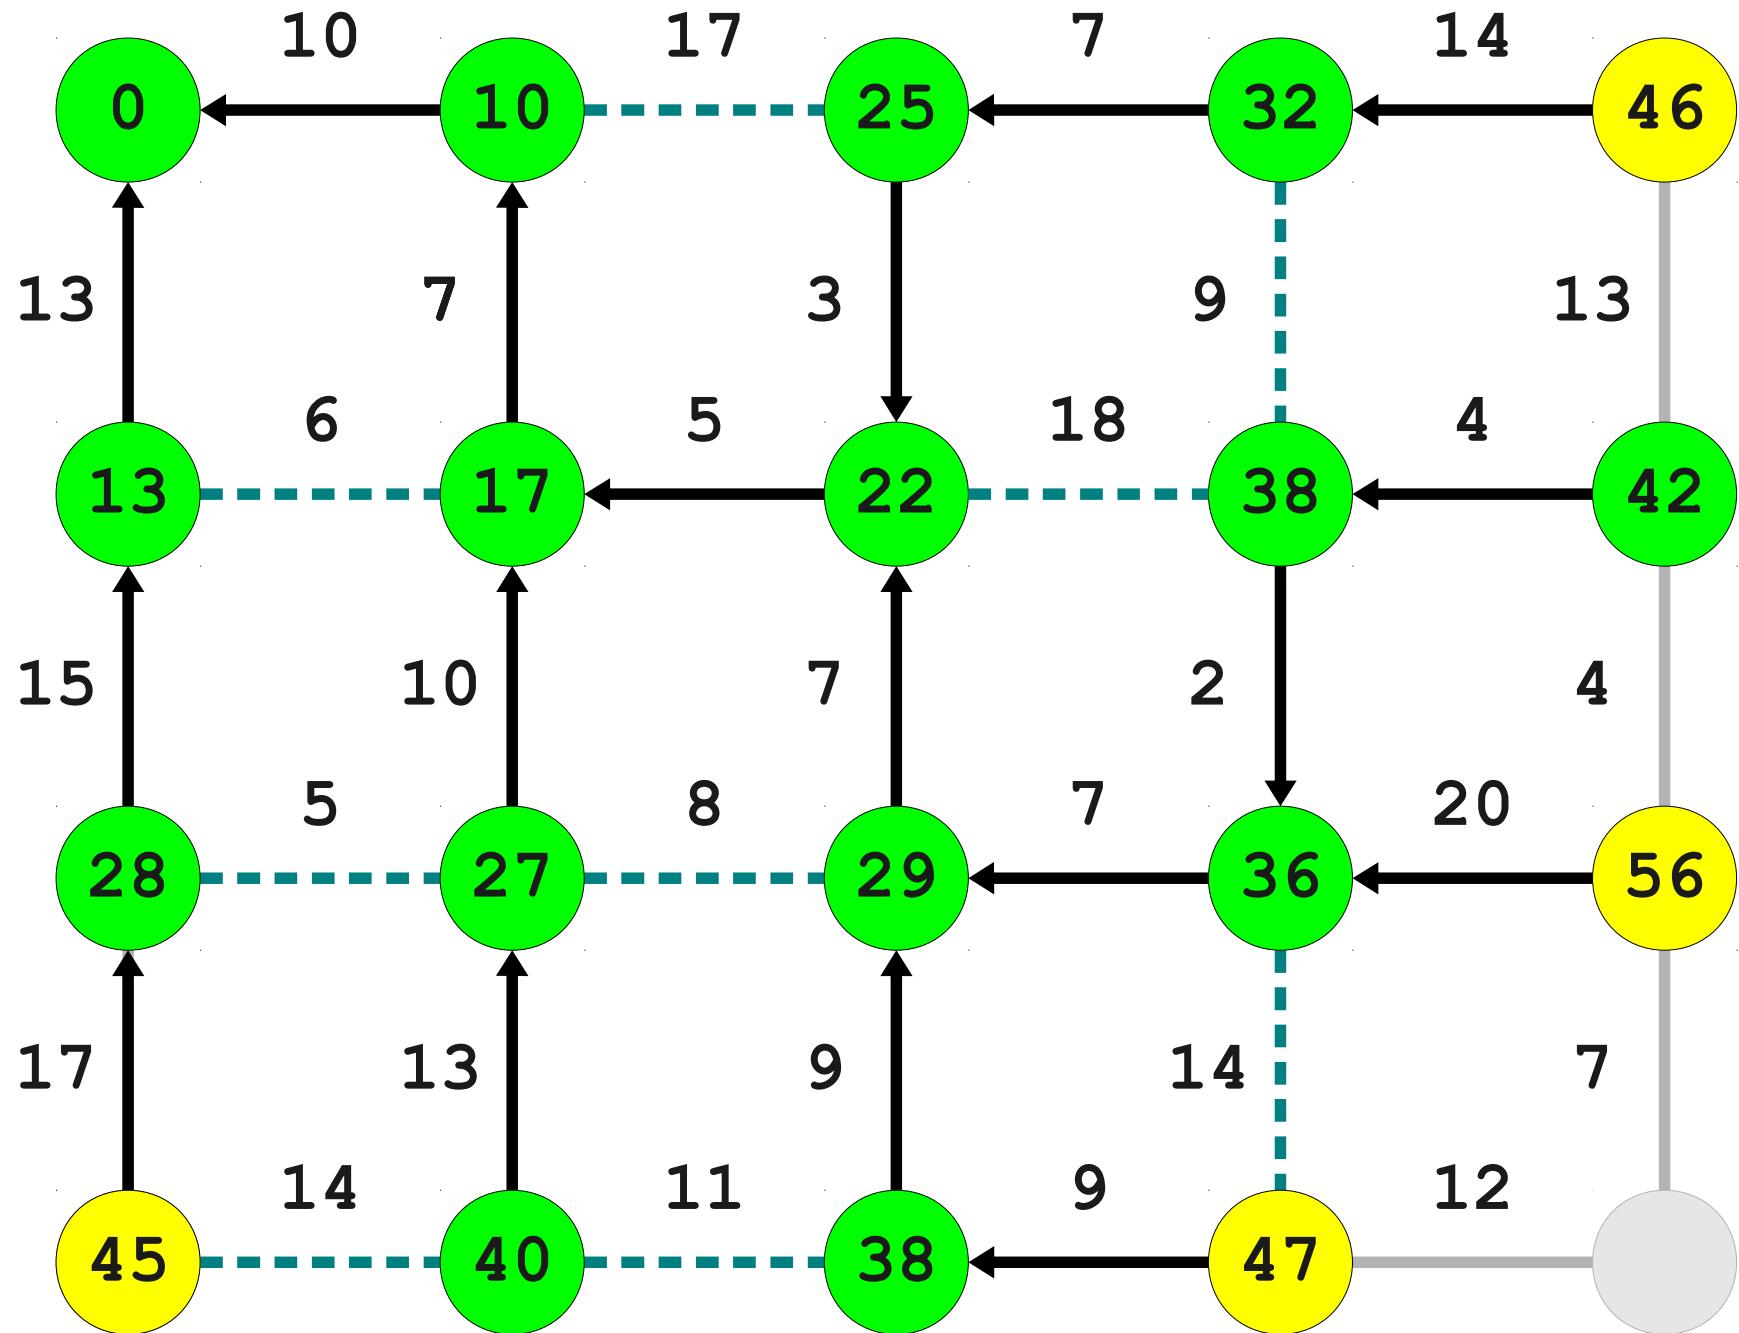


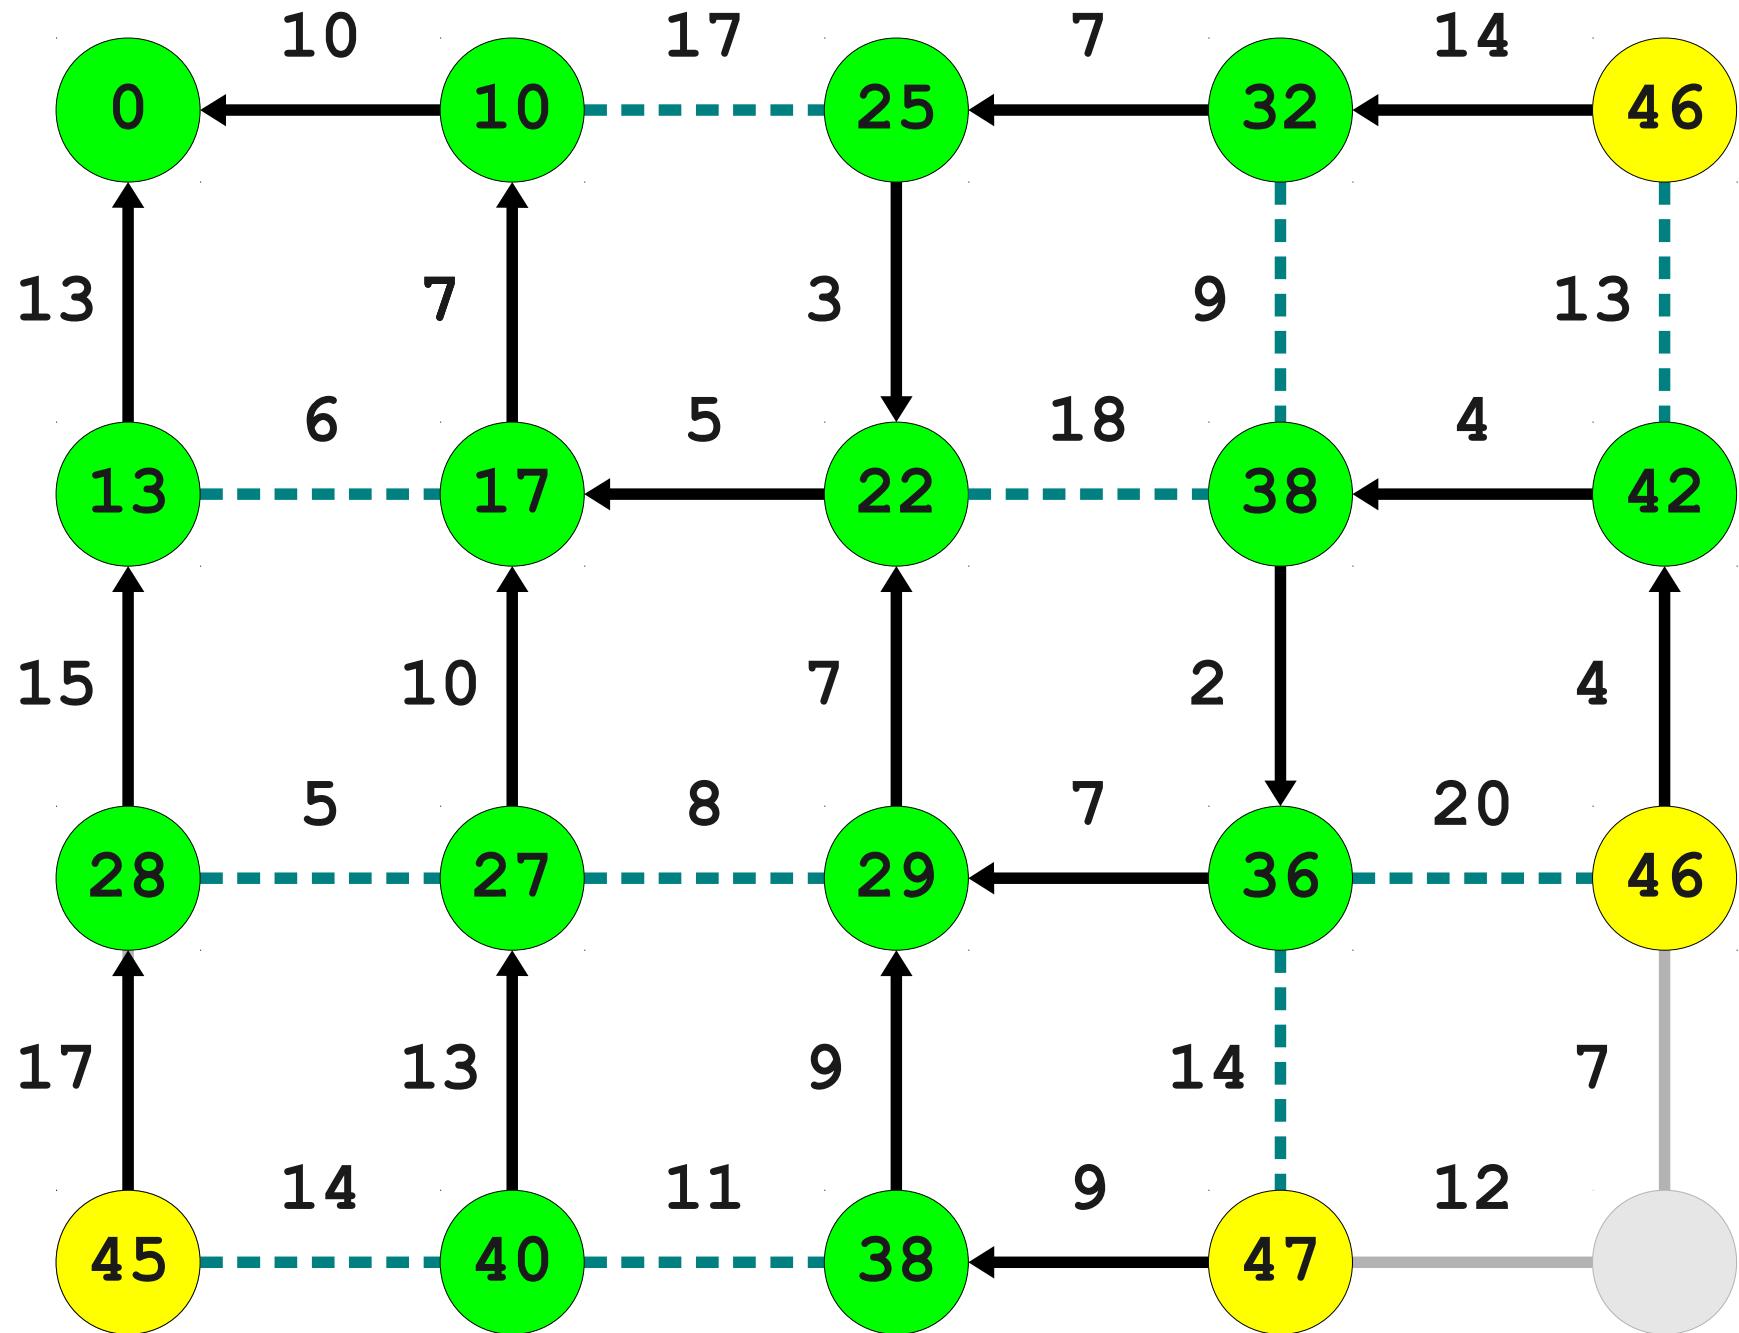


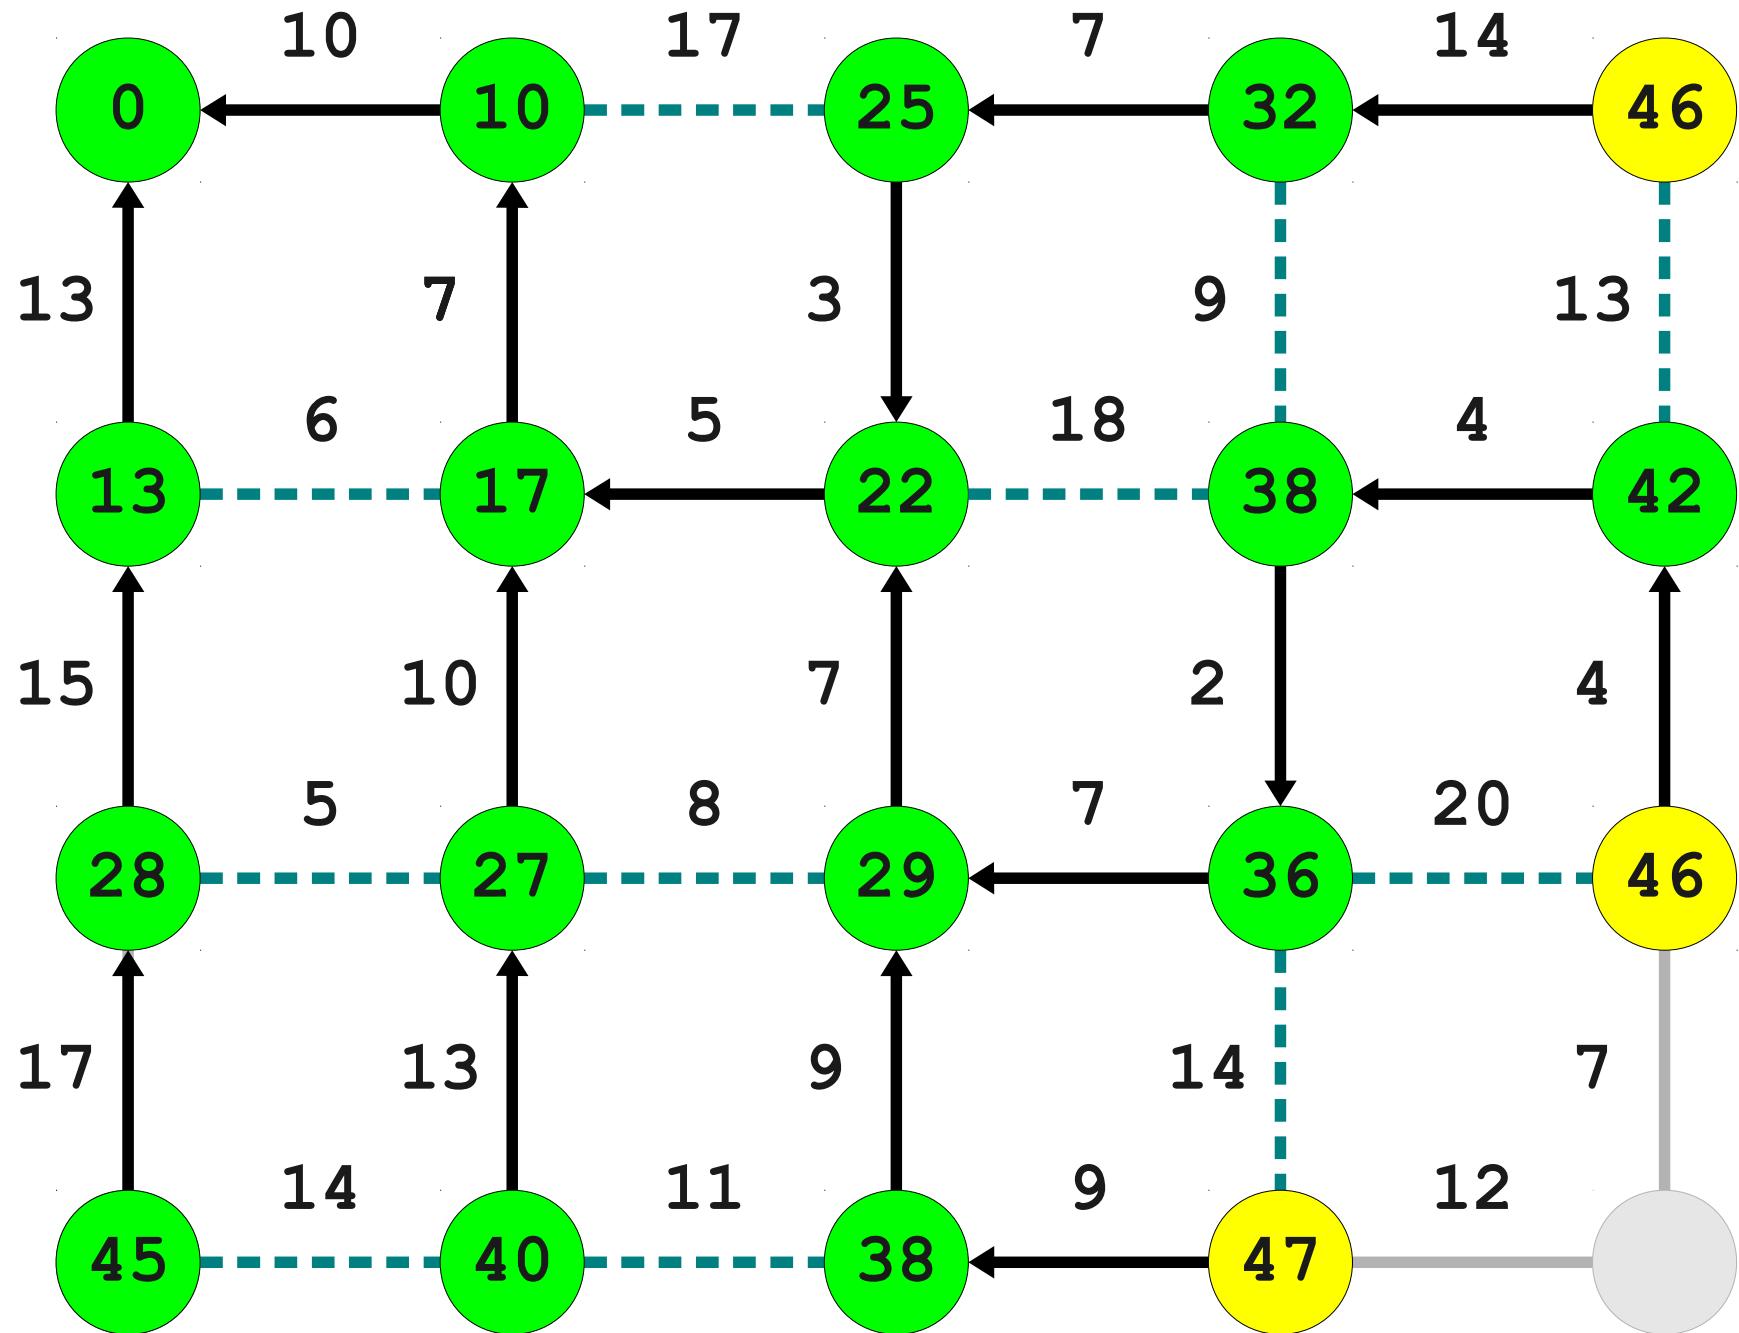


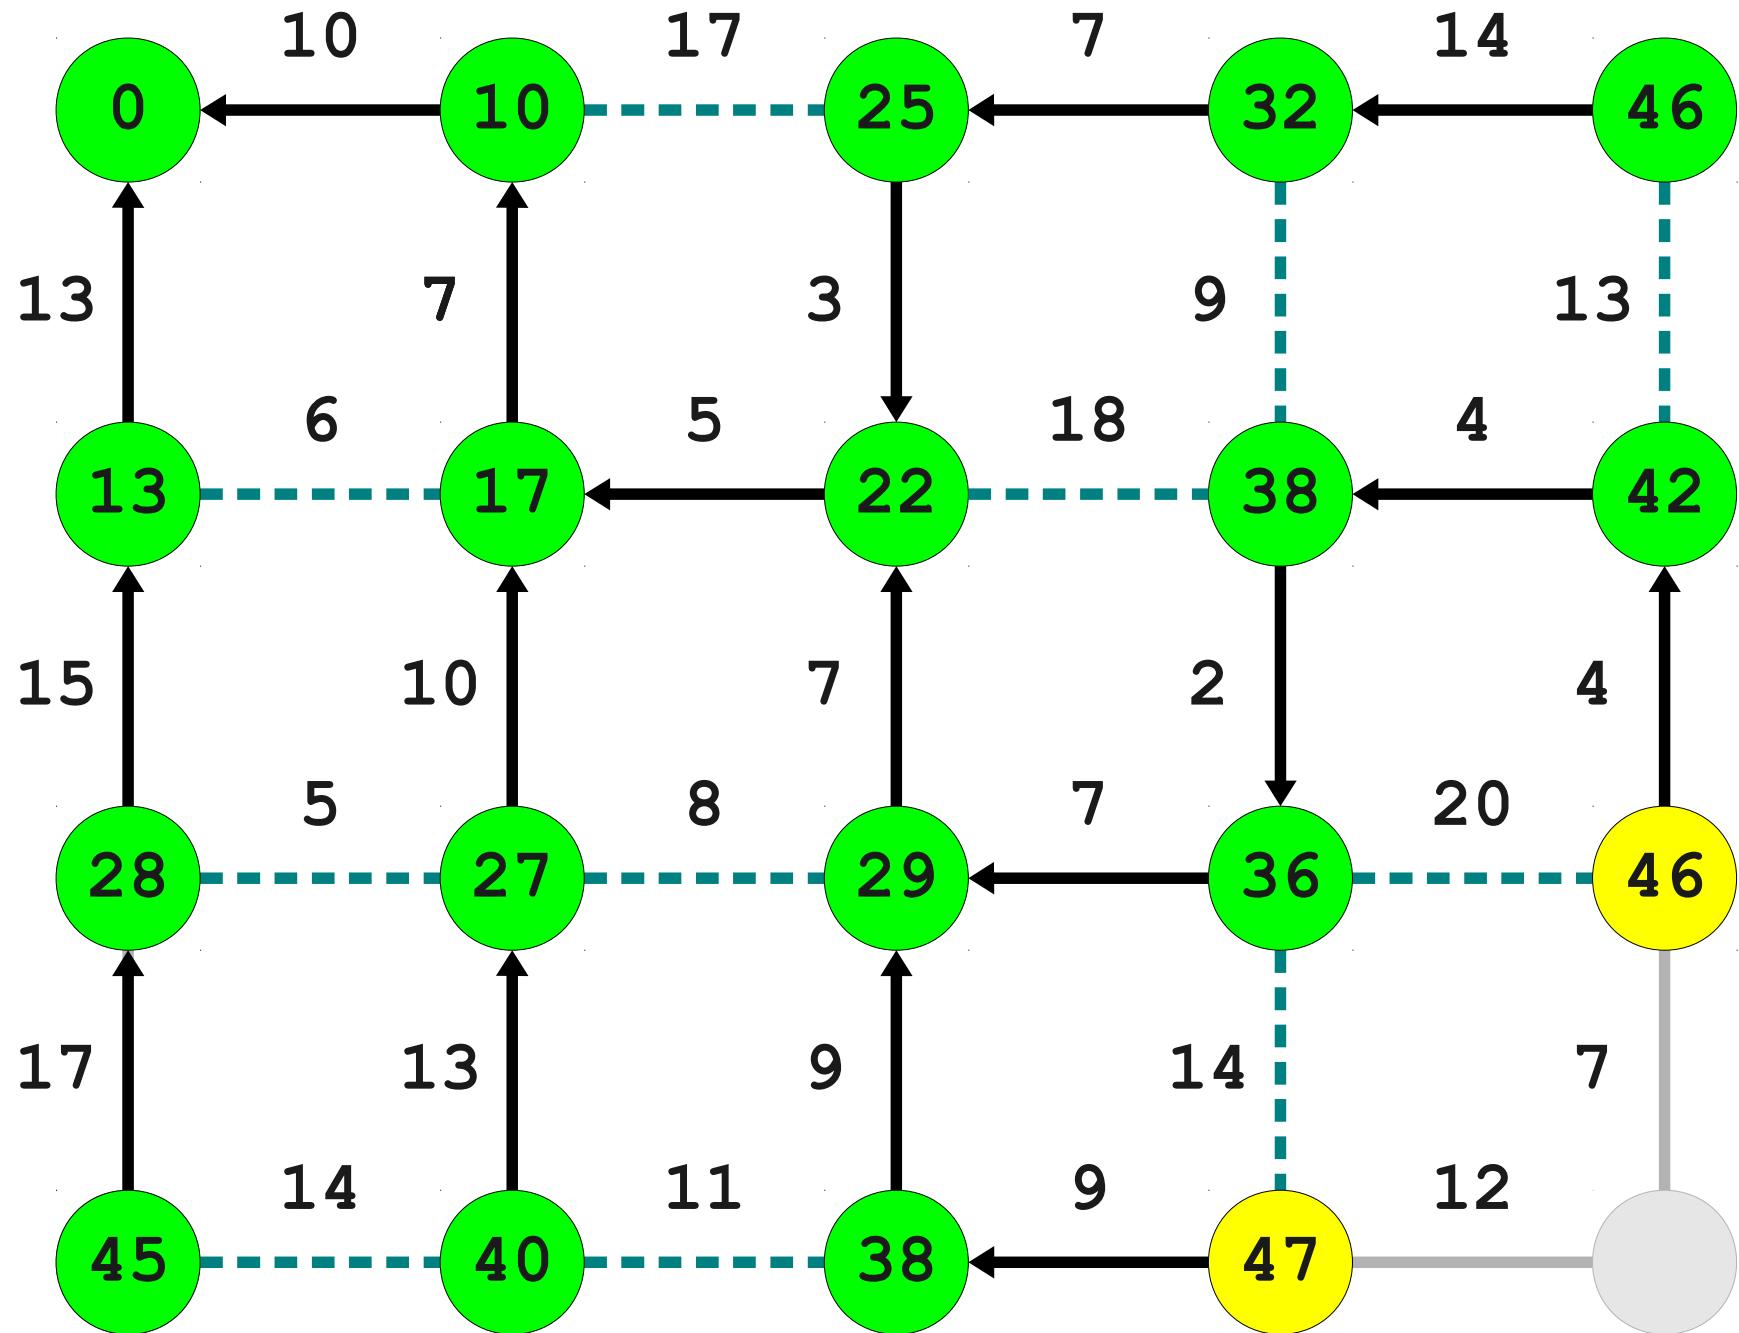


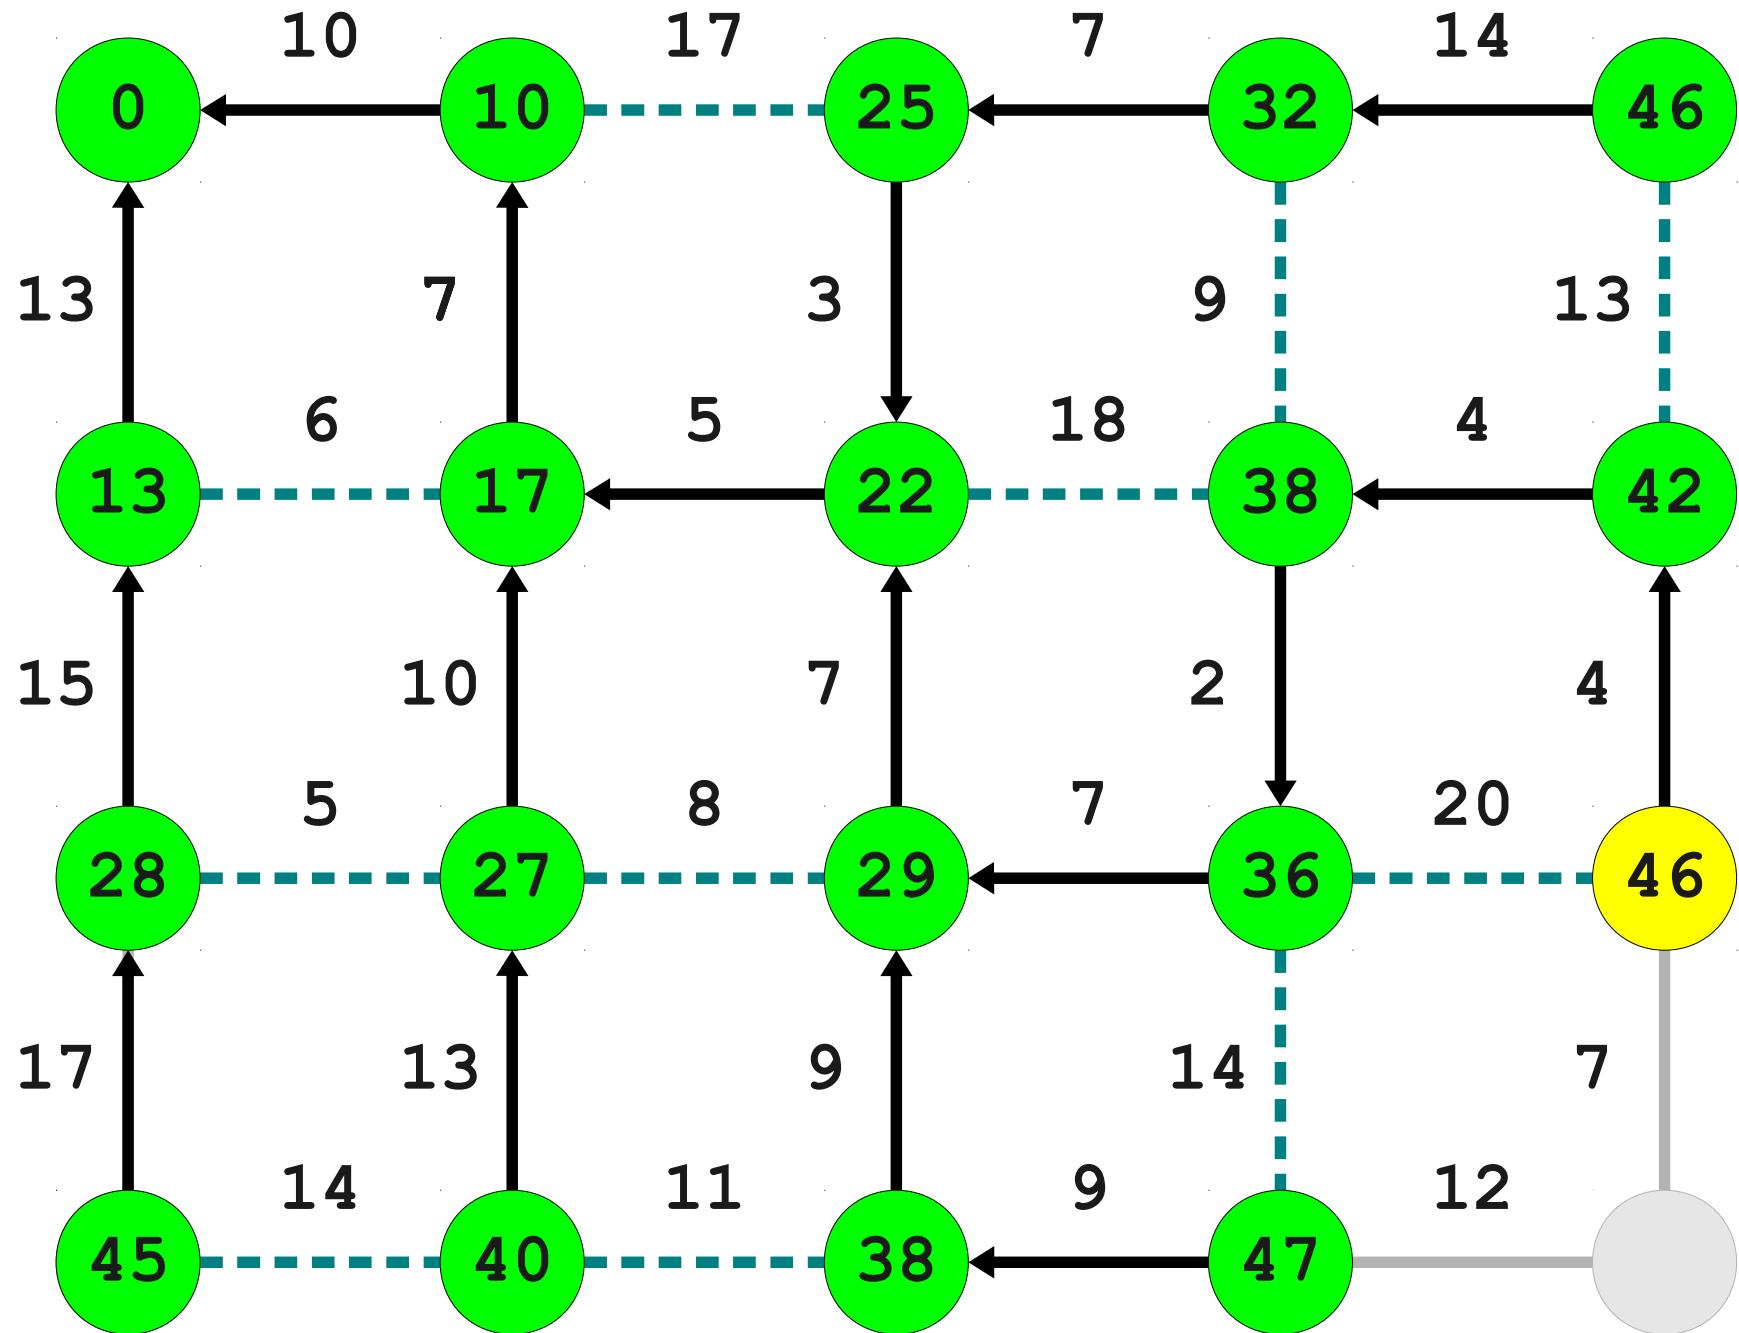


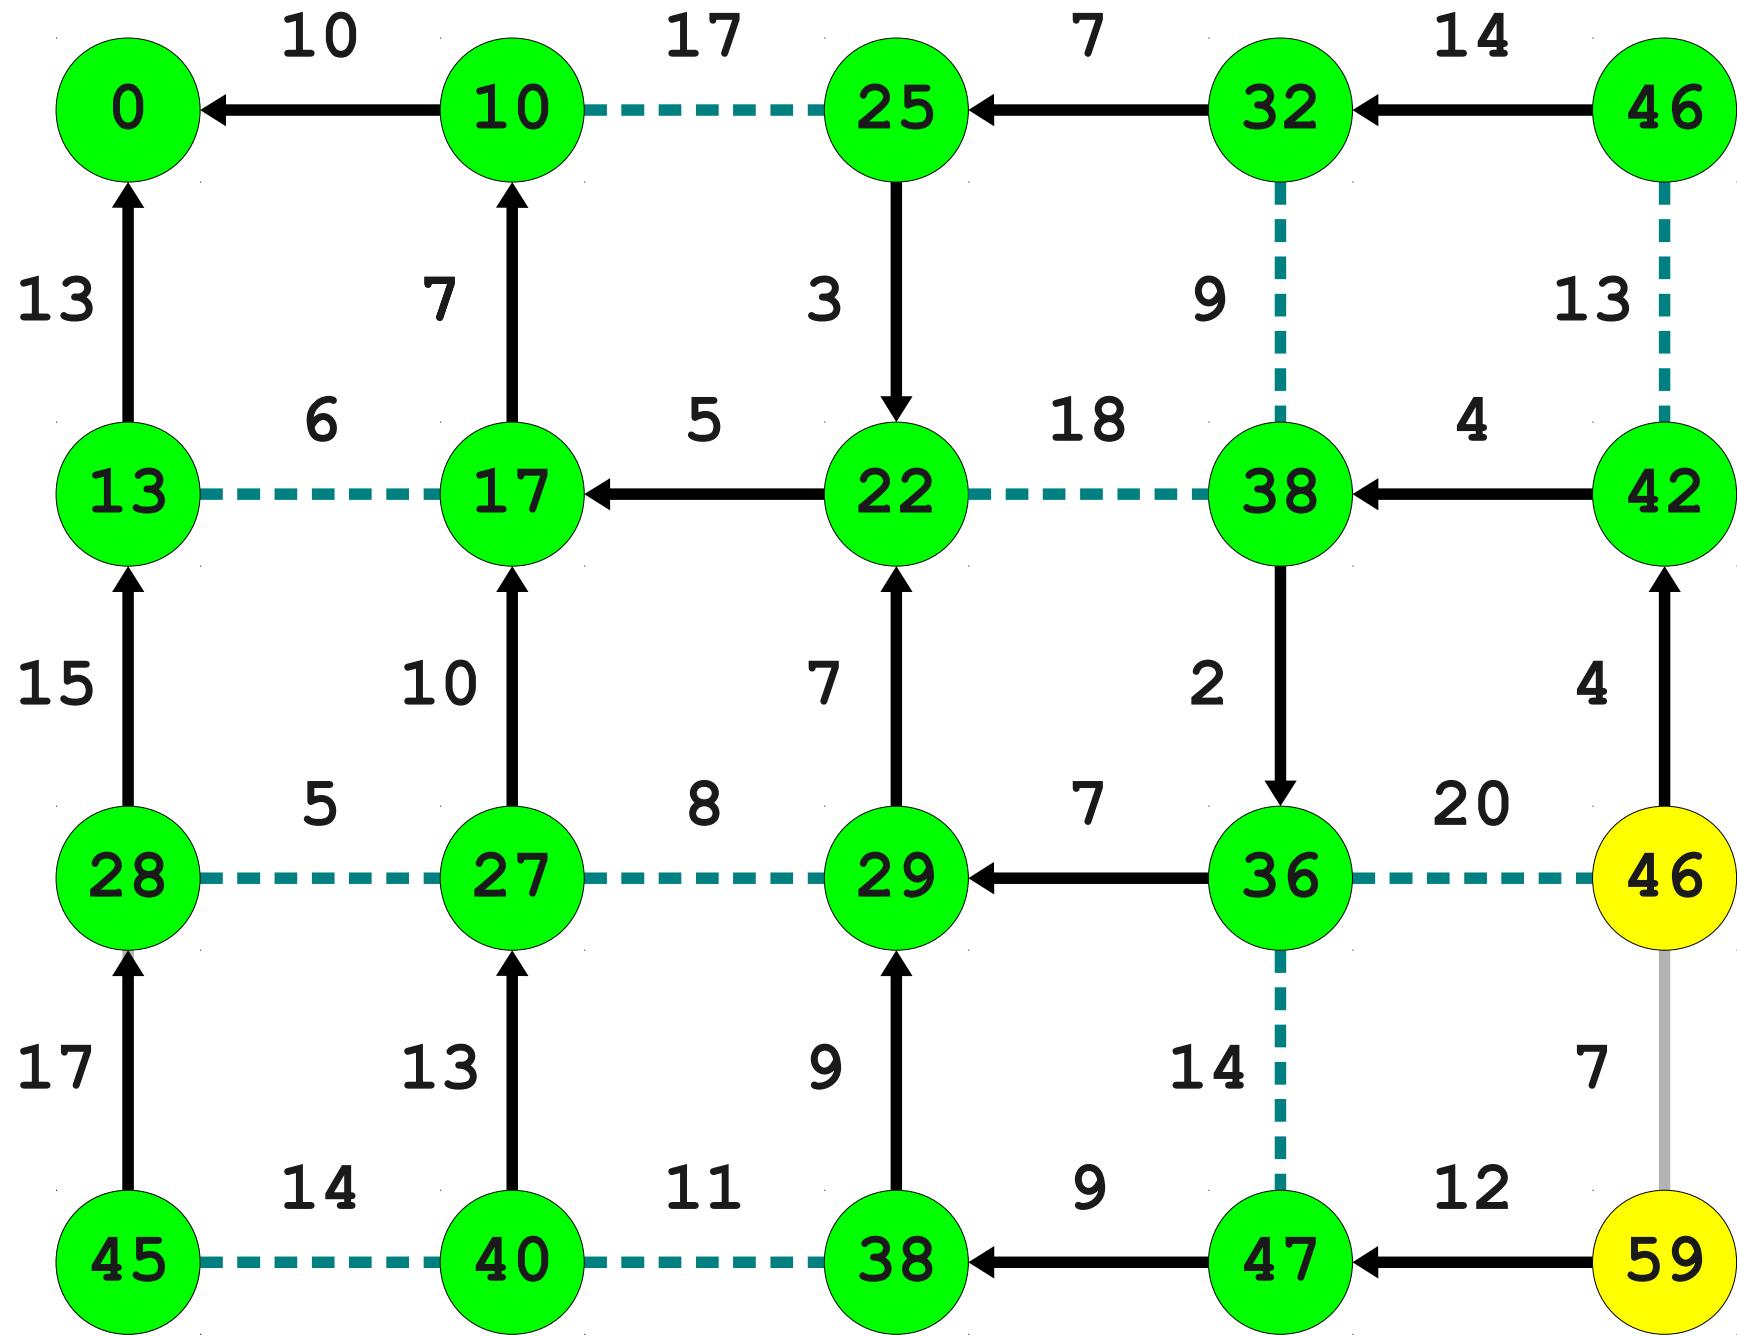


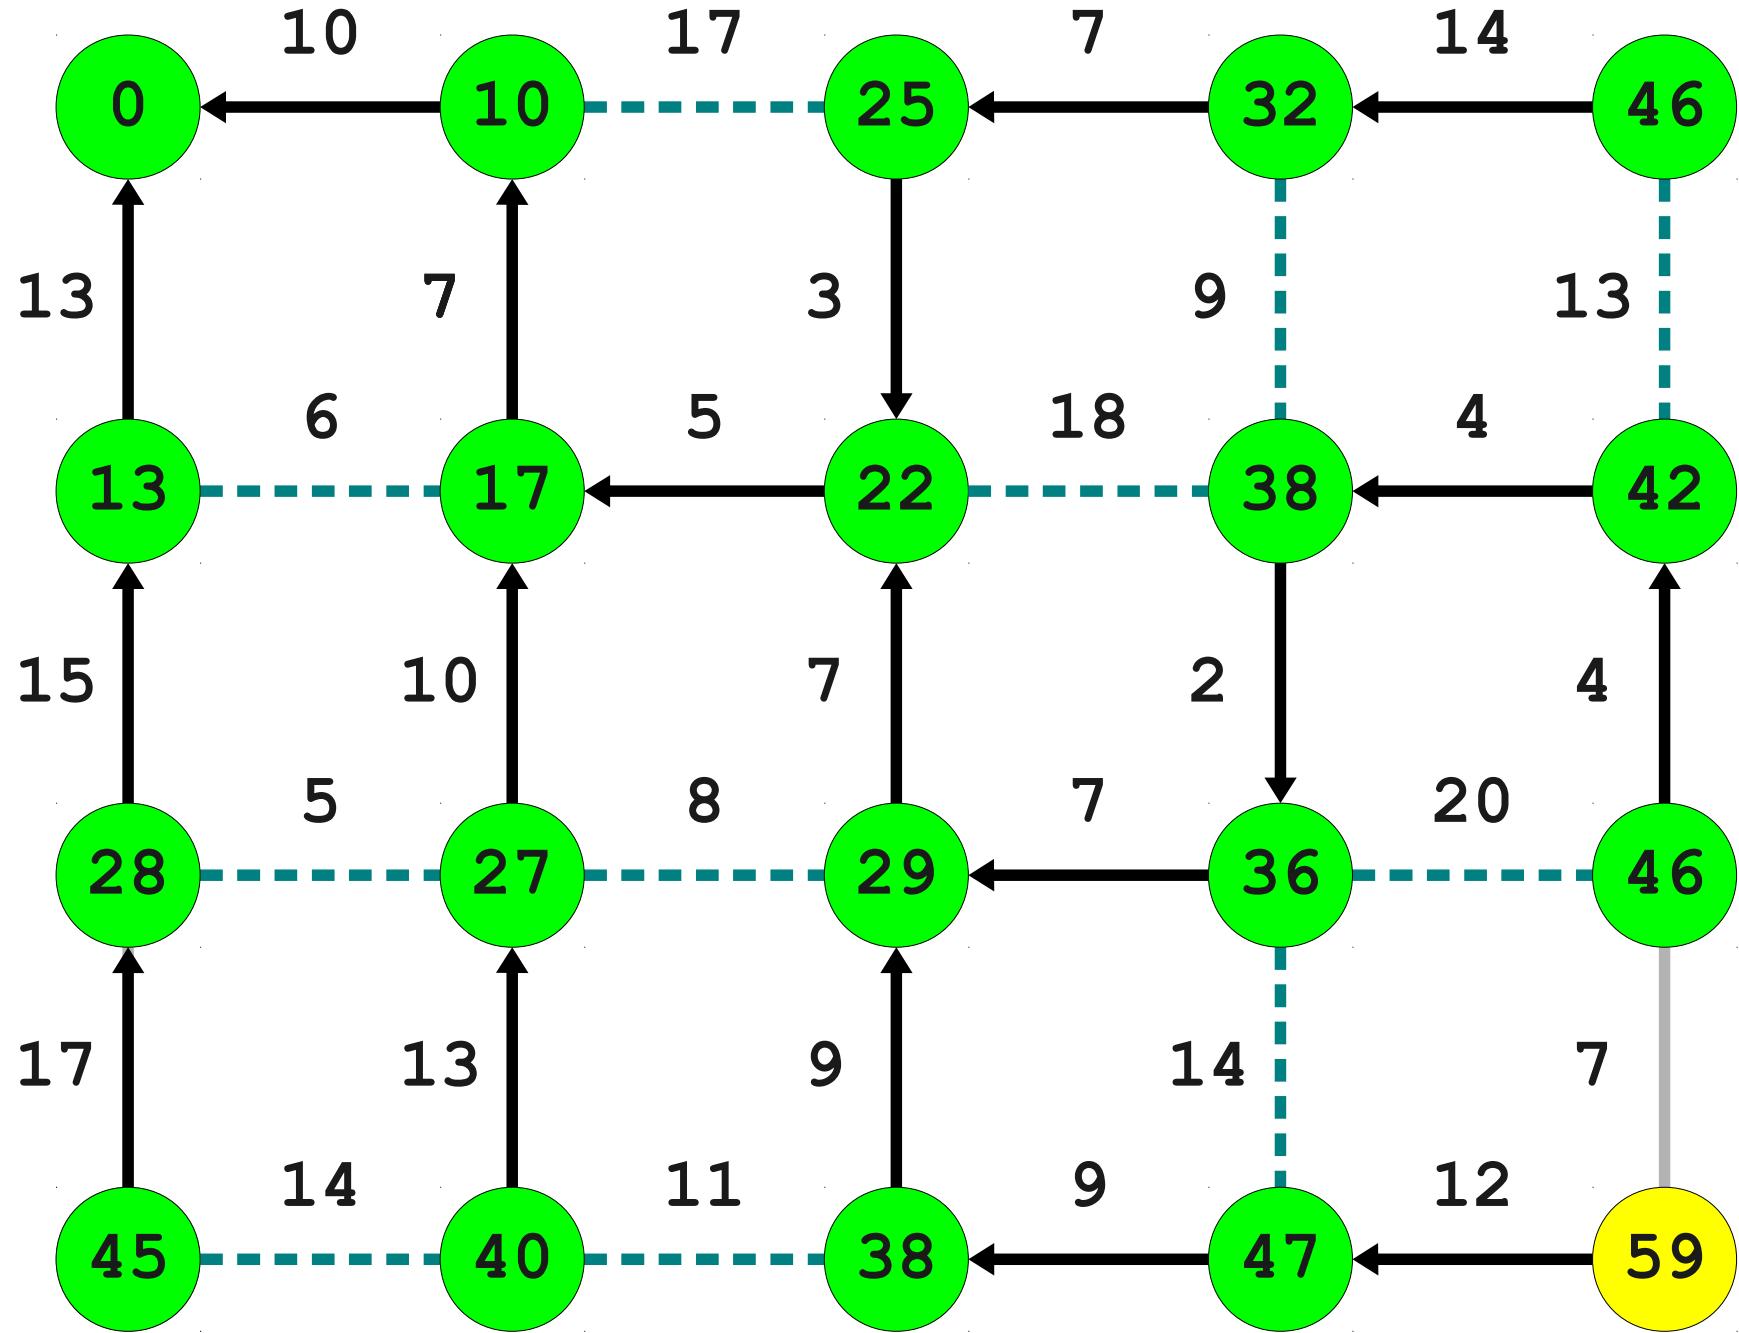


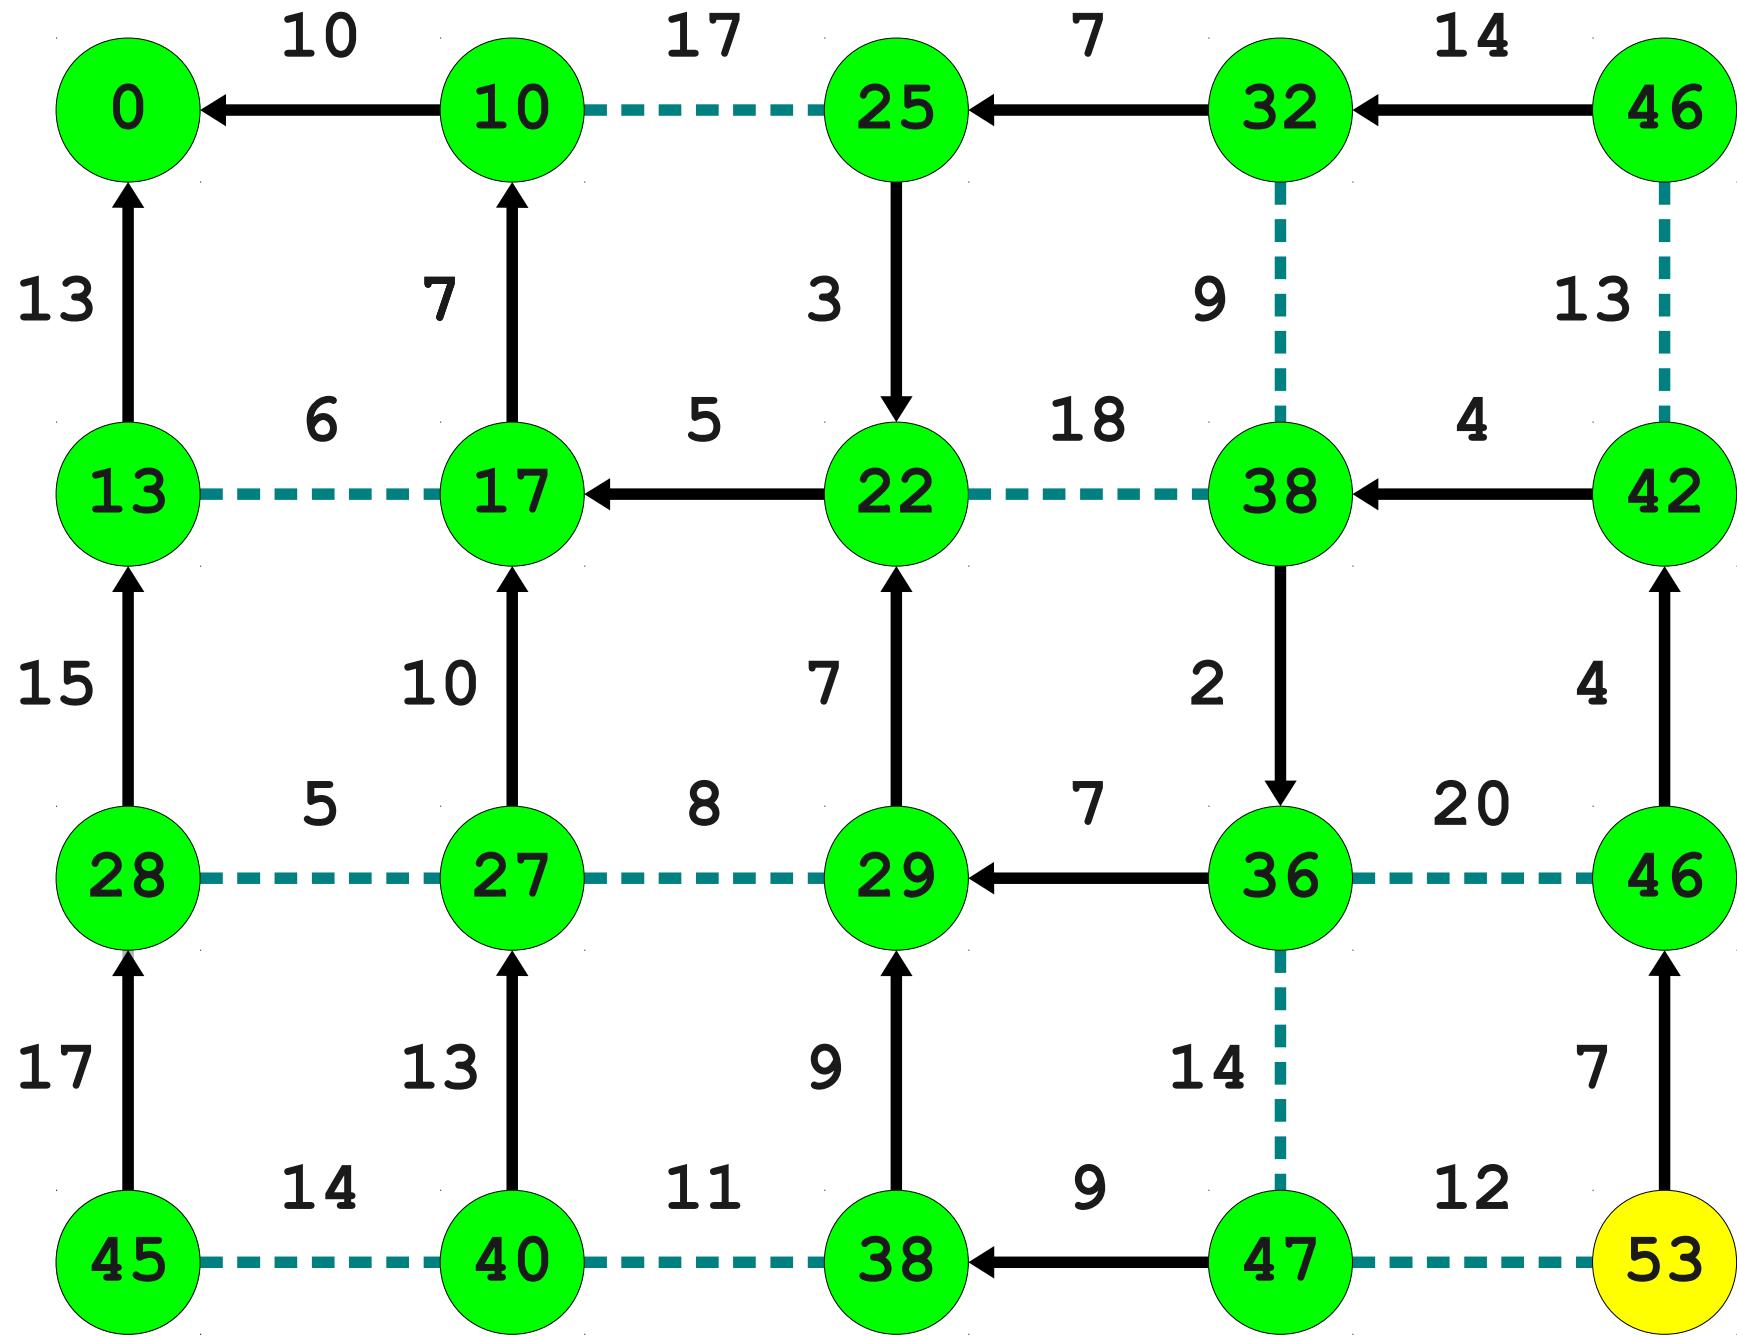


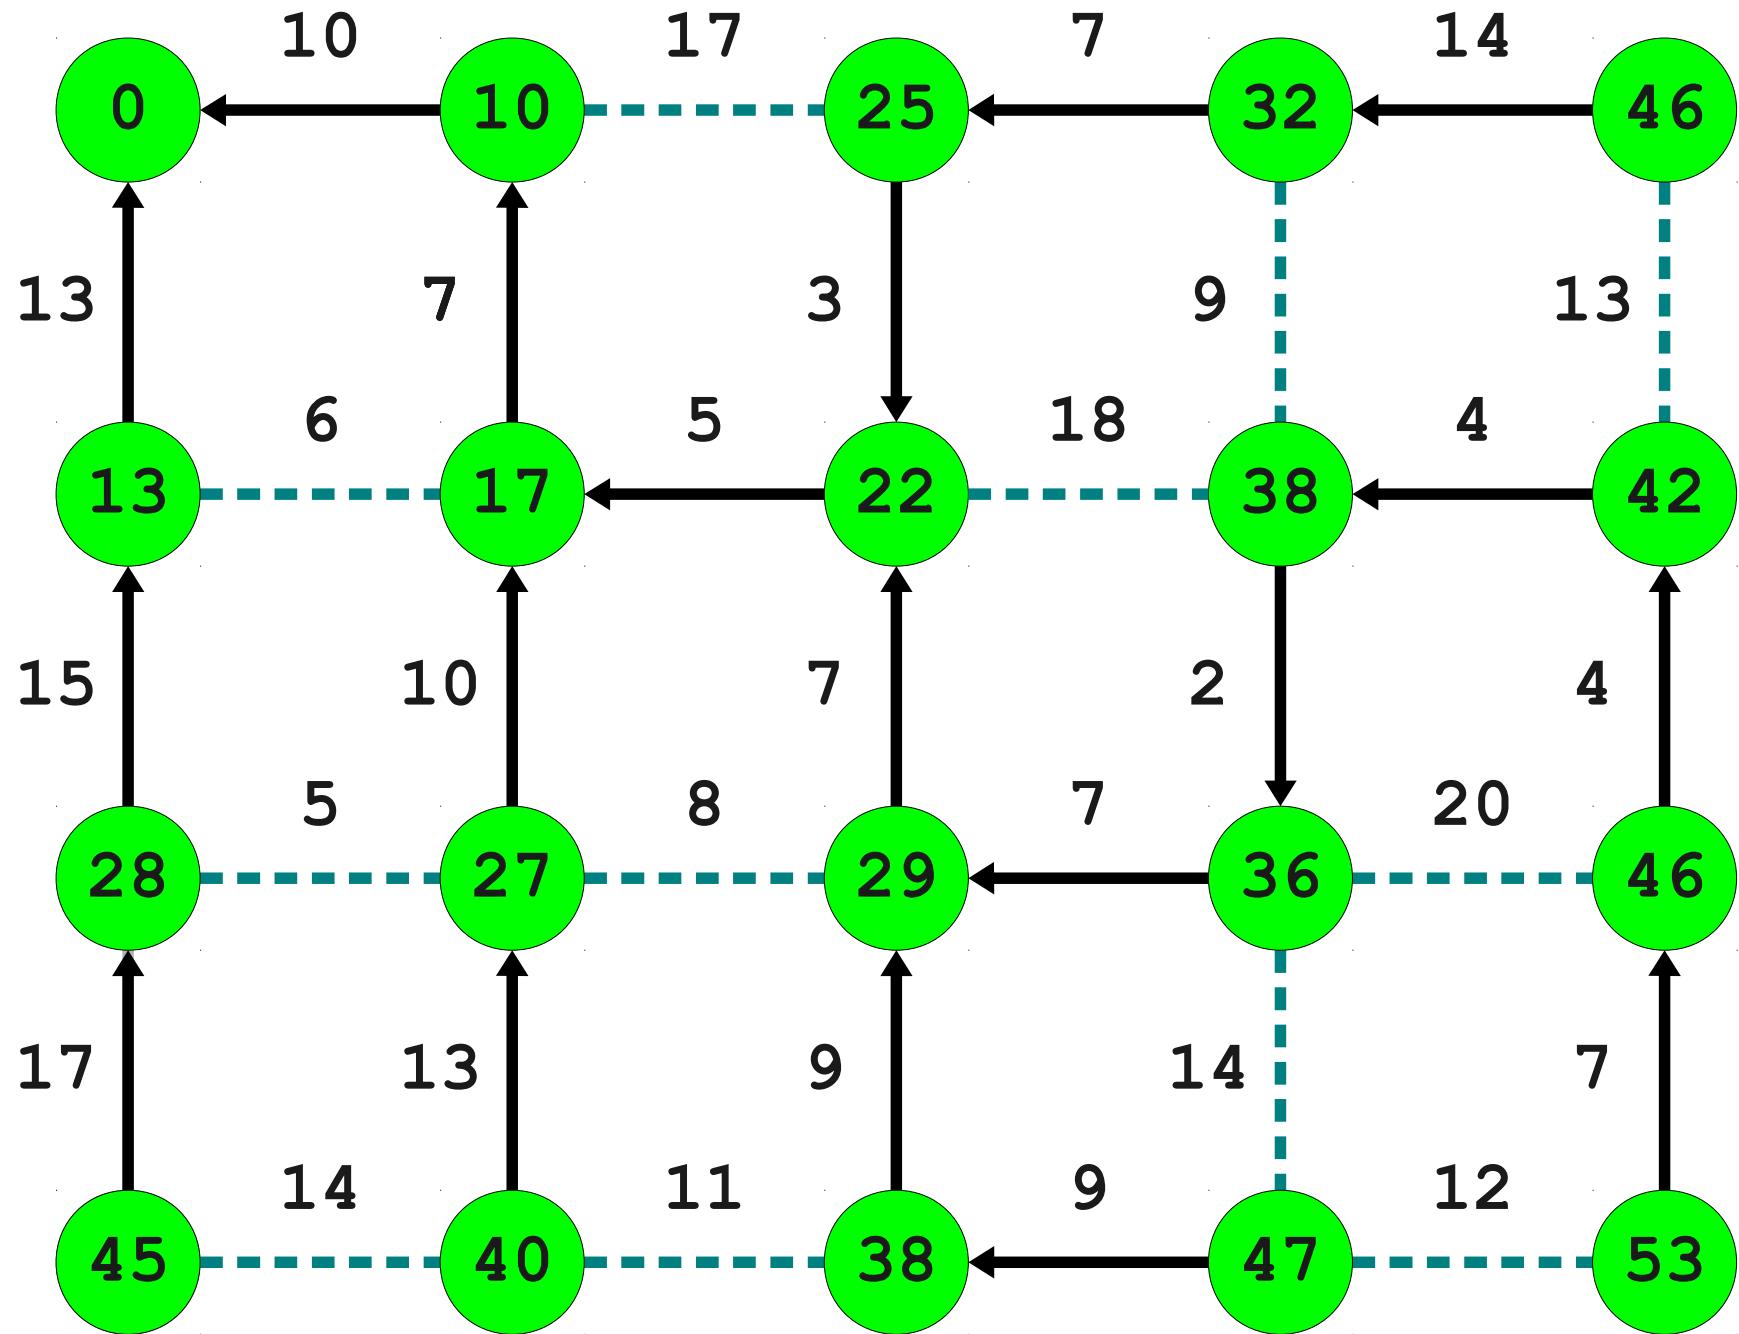


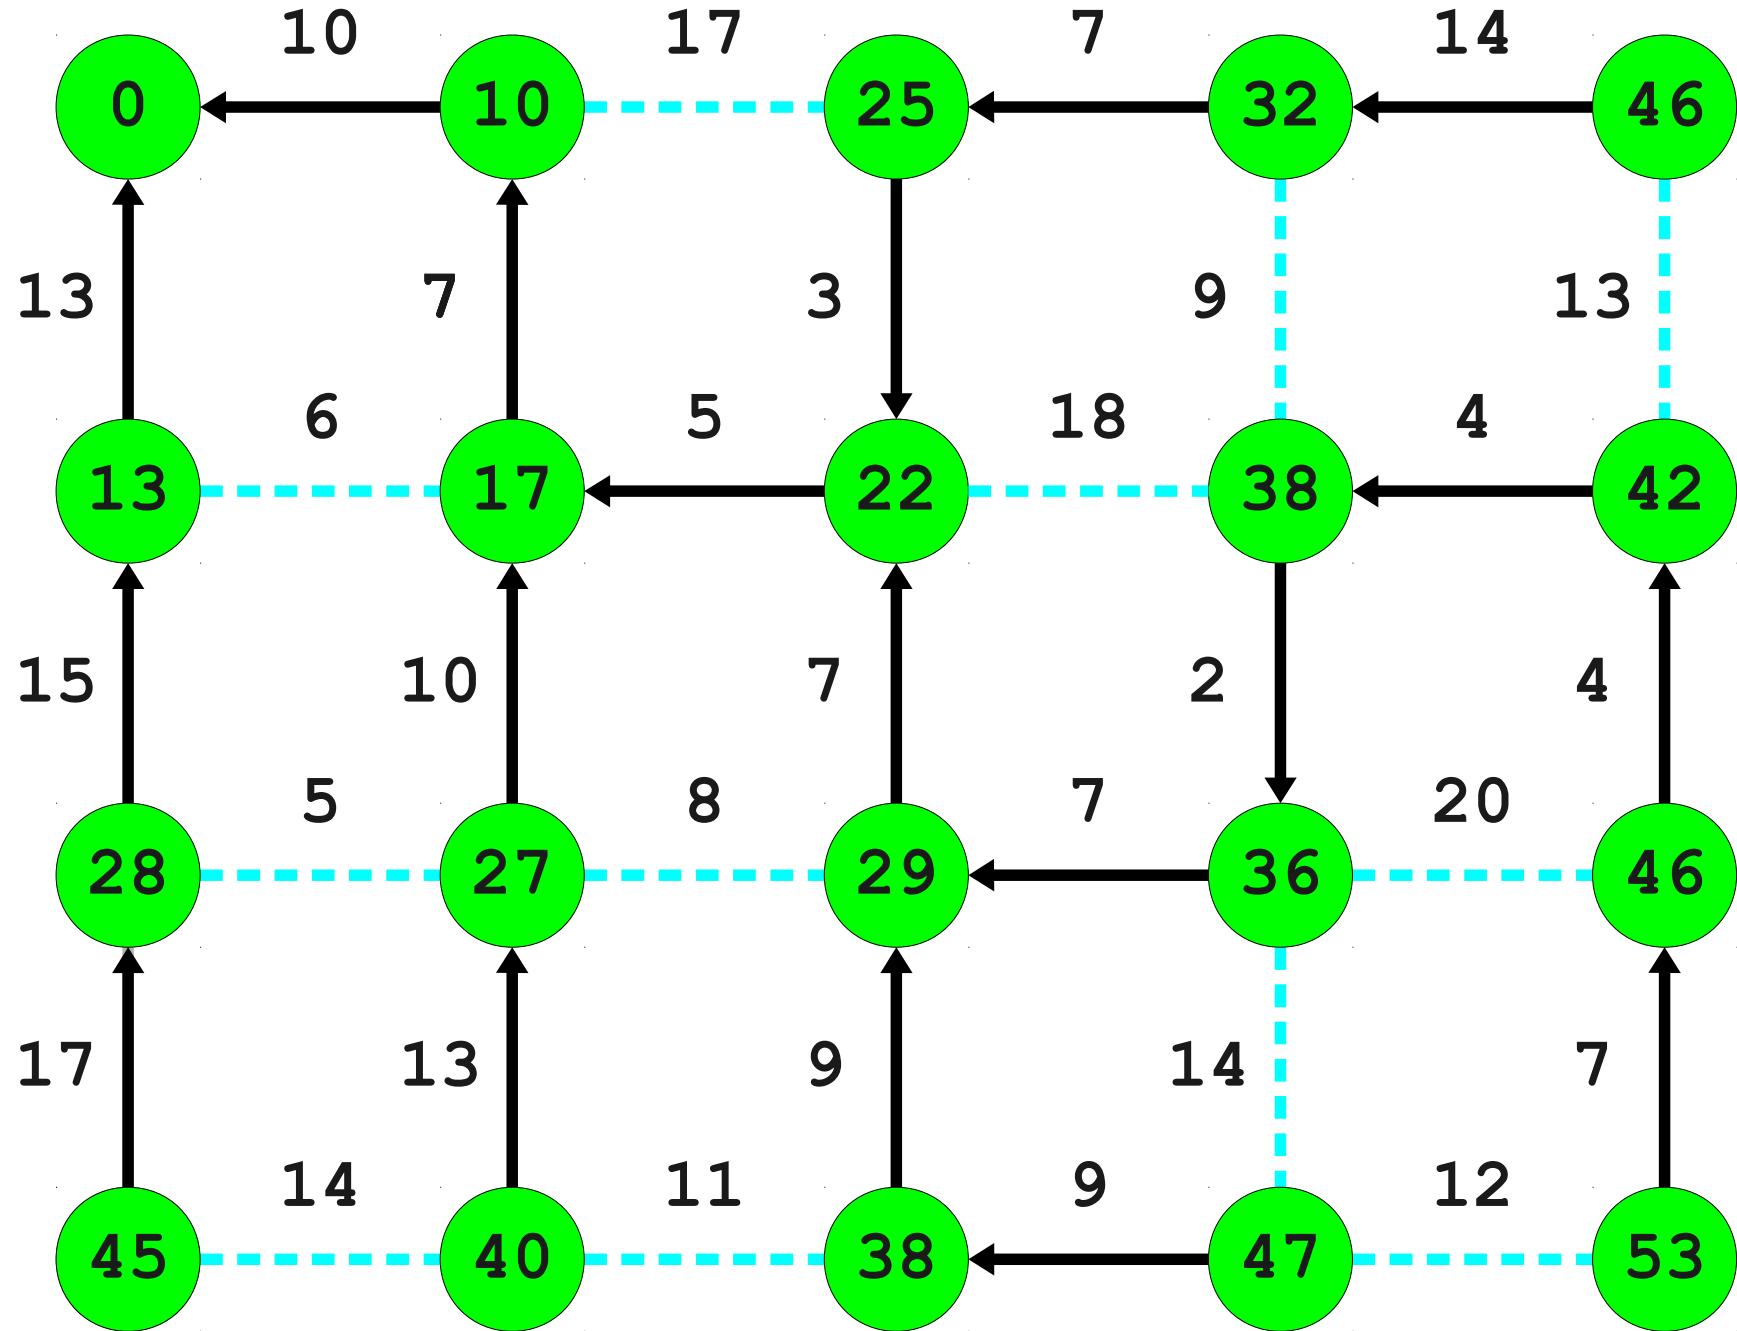


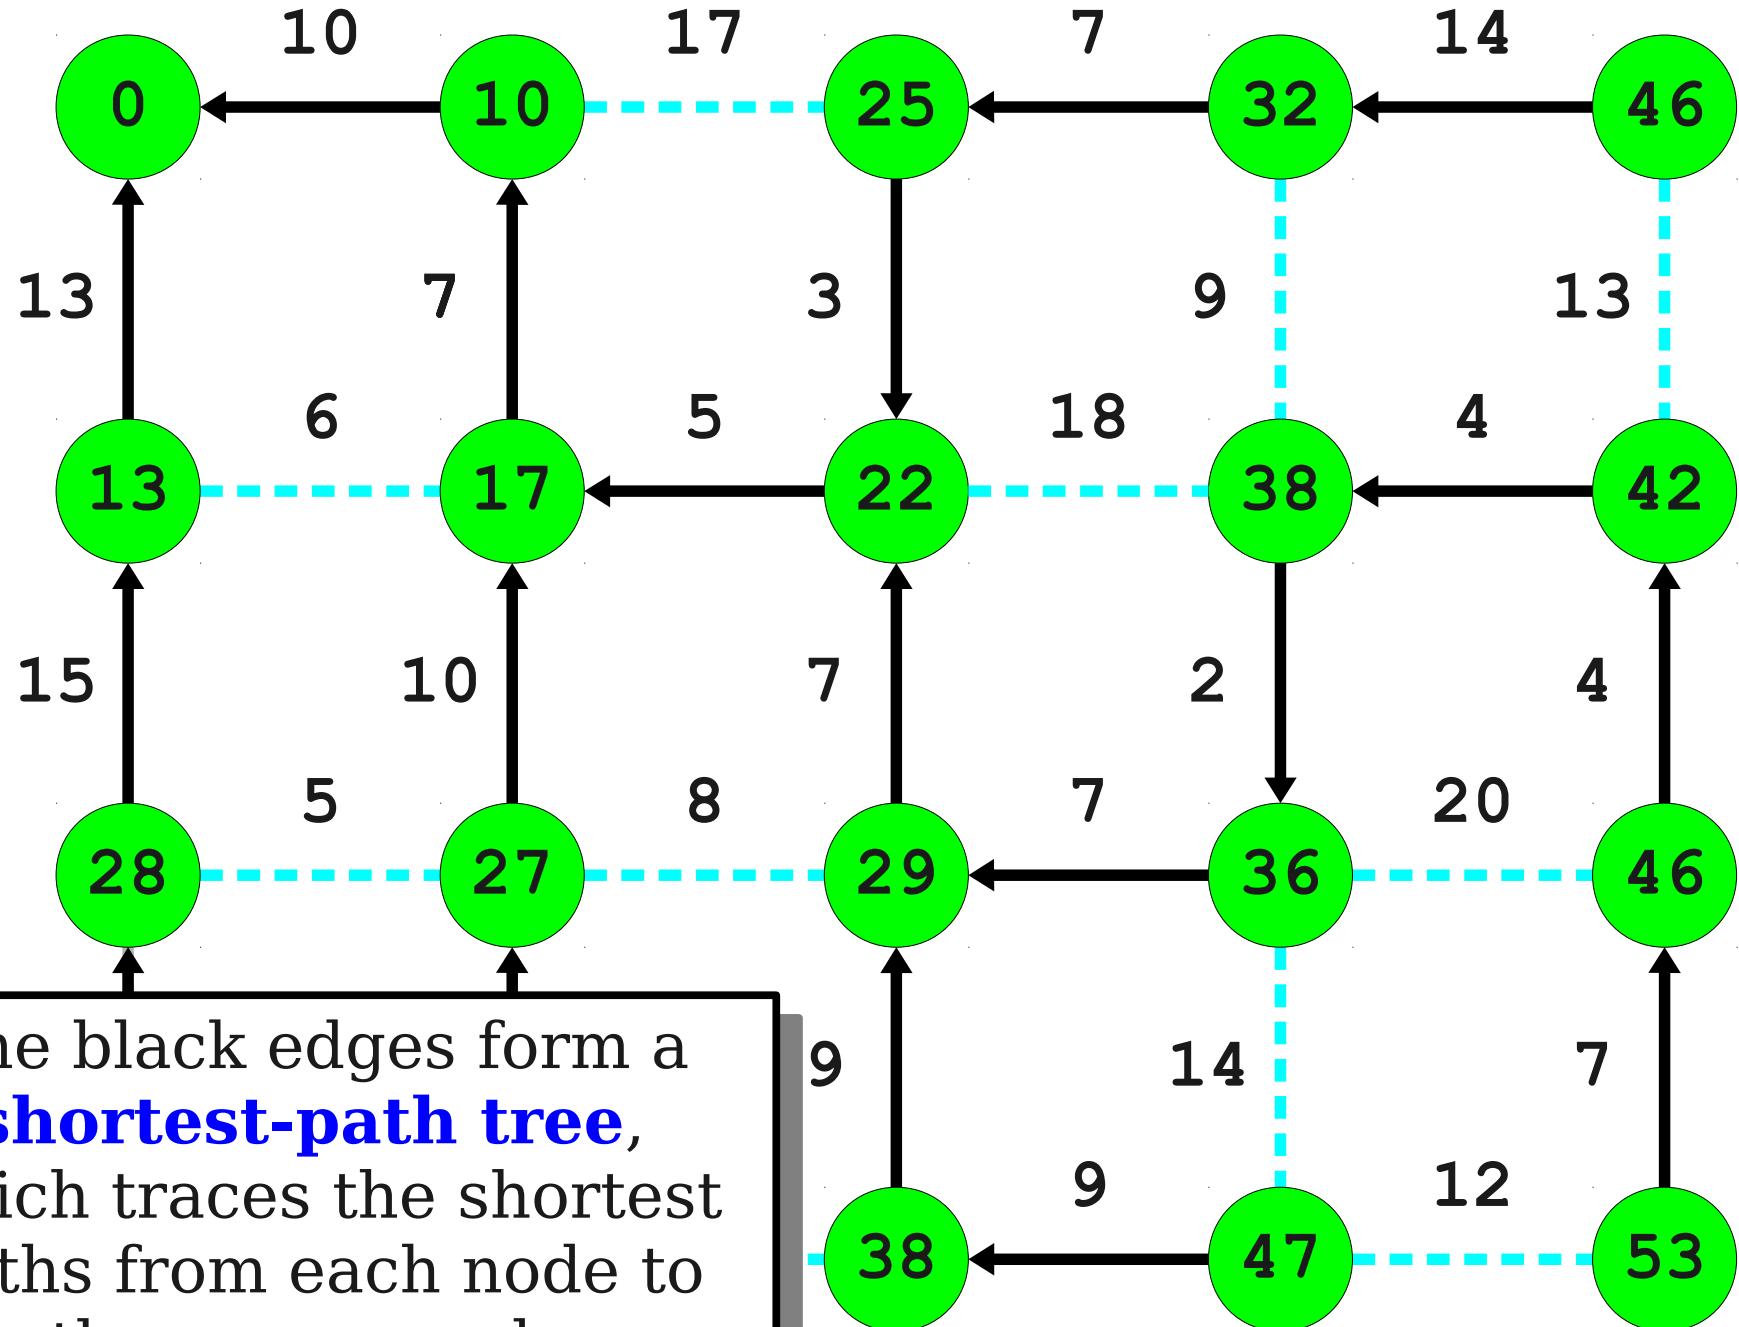












```
procedure dijkstrasAlgorithm(s, G):
    let q be a new queue
    for each v in V:
        dist[v] =  $\infty$ 

    dist[s] = 0
    enqueue(s, q)

    while q is not empty:
        let v be a node in q minimizing dist[v]
        remove(v, q)

        for each node u connected to v:
            if dist[u] > dist[v] + l(u, v):
                dist[u] = dist[v] + l(u, v)
                if u is not enqueued into q:
                    enqueue(u, q)
```

Dijkstra's Algorithm

- Assuming nonnegative edge lengths, finds the shortest path from s to each node in G .
- Correctness proof sketch is based on the second argument for breadth-first search:
 - Always picks the node v minimizing $d(s, u) + l(u, v)$ for yellow v and green u .
 - If a shorter path P exists to v , it must leave the set of green nodes through some edge (x, y) .
 - But then $l(P)$ is at least $d(s, x) + l(x, y)$, which is at least $d(s, u) + l(u, v)$.
 - So the “shorter” path costs at least as much as the path we found.

```
procedure dijkstrasAlgorithm(s, G):
    let q be a new queue
    for each v in V:
        dist[v] =  $\infty$ 

    dist[s] = 0
    enqueue(s, q)

    while q is not empty:
        let v be a node in q minimizing dist[v]
        remove(v, q)

        for each node u connected to v:
            if dist[u] > dist[v] + l(u, v):
                dist[u] = dist[v] + l(u, v)
                if u is not enqueued into q:
                    enqueue(u, q)
```

```
procedure dijkstrasAlgorithm(s, G):
    let q be a new queue
    for each v in V:
        dist[v] =  $\infty$ 

    dist[s] = 0
    enqueue(s, q)

    while q is not empty:
        let v be a node in q minimizing dist[v]
        remove(v, q)

        for each node u connected to v:
            if dist[u] > dist[v] + l(u, v):
                dist[u] = dist[v] + l(u, v)
                if u is not enqueued into q:
                    enqueue(u, q)
```

```
procedure dijkstrasAlgorithm(s, G):
```

```
    let q be a new queue
```

```
    for each v in V:
```

```
        dist[v] =  $\infty$ 
```

```
        dist[s] = 0
```

```
        enqueue(s, q)
```

O($m + n$)

```
    while q is not empty:
```

```
        let v be a node in q minimizing dist[v]
```

```
        remove(v, q)
```

```
        for each node u connected to v:
```

```
            if dist[u] > dist[v] + l(u, v):
```

```
                dist[u] = dist[v] + l(u, v)
```

```
                if u is not enqueued into q:
```

```
                    enqueue(u, q)
```

```
procedure dijkstrasAlgorithm(s, G):
    let q be a new queue
    for each v in V:
        dist[v] = ∞

    dist[s] = 0
    enqueue(s, q)

    while q is not empty:
        let v be a node in q minimizing dist[v]
        remove(v, q)

        for each node u connected to v:
            if dist[u] > dist[v] + l(u, v):
                dist[u] = dist[v] + l(u, v)
                if u is not enqueued into q:
                    enqueue(u, q)
```

```
procedure dijkstrasAlgorithm(s, G):
    let q be a new queue
    for each v in V:
        dist[v] = ∞

    dist[s] = 0
    enqueue(s, q)
```

O(n^2)

```
while q is not empty:
    let v be a node in q minimizing dist[v]
    remove(v, q)

    for each node u connected to v:
        if dist[u] > dist[v] + l(u, v):
            dist[u] = dist[v] + l(u, v)
            if u is not enqueued into q:
                enqueue(u, q)
```

Dijkstra Runtime

- Using a standard implementation of a queue, Dijkstra's algorithm runs in time **$O(n^2)$** .
 - $O(n + m)$ time processing nodes and edges, plus $O(n^2)$ time finding the lowest-cost node.
 - Since $m = O(n^2)$, $O(n + m + n^2) = O(n^2)$.
- Using a slightly fancier data structure (a binary heap), can be made to run in time **$O(m \log n)$** .
 - Is this necessarily more efficient?
 - More on how to do this later this quarter.
- Using a *much* fancier data structure (the *Fibonacci heap*), can be made to run in time **$O(m + n \log n)$** .
 - Take CE354 for details!

Shortest Path Algorithms

- If all edges have the same weight, can use breadth-first search to find shortest paths.
 - Takes time $O(m + n)$.
- If edges have nonnegative weight, can use Dijkstra's algorithm.
 - Takes time $O(n^2)$, or less using more complex data structures.
- What about the case where edges can have negative weight?
 - More on that later in the quarter...

INITIALIZE-SINGLE-SOURCE(G, s)

- 1 **for** each vertex $v \in G.V$
- 2 $v.d = \infty$
- 3 $v.\pi = \text{NIL}$
- 4 $s.d = 0$

DIJKSTRA(G, w, s)

```
1  INITIALIZE-SINGLE-SOURCE( $G, s$ )
2   $S = \emptyset$ 
3   $Q = \emptyset$ 
4  for each vertex  $u \in G.V$ 
5      INSERT( $Q, u$ )
6  while  $Q \neq \emptyset$ 
7       $u = \text{EXTRACT-MIN}(Q)$ 
8       $S = S \cup \{u\}$ 
9      for each vertex  $v$  in  $G.Adj[u]$ 
10         RELAX( $u, v, w$ )
11         if the call of RELAX decreased  $v.d$ 
12             DECREASE-KEY( $Q, v, v.d$ )
```

RELAX(u, v, w)

1 **if** $v.d > u.d + w(u, v)$

2 $v.d = u.d + w(u, v)$

3 $v.\pi = u$

Dijkstra's algorithm: which priority queue?

Performance. Depends on PQ: n INSERT, n DELETE-MIN, $\leq m$ DECREASE-KEY.

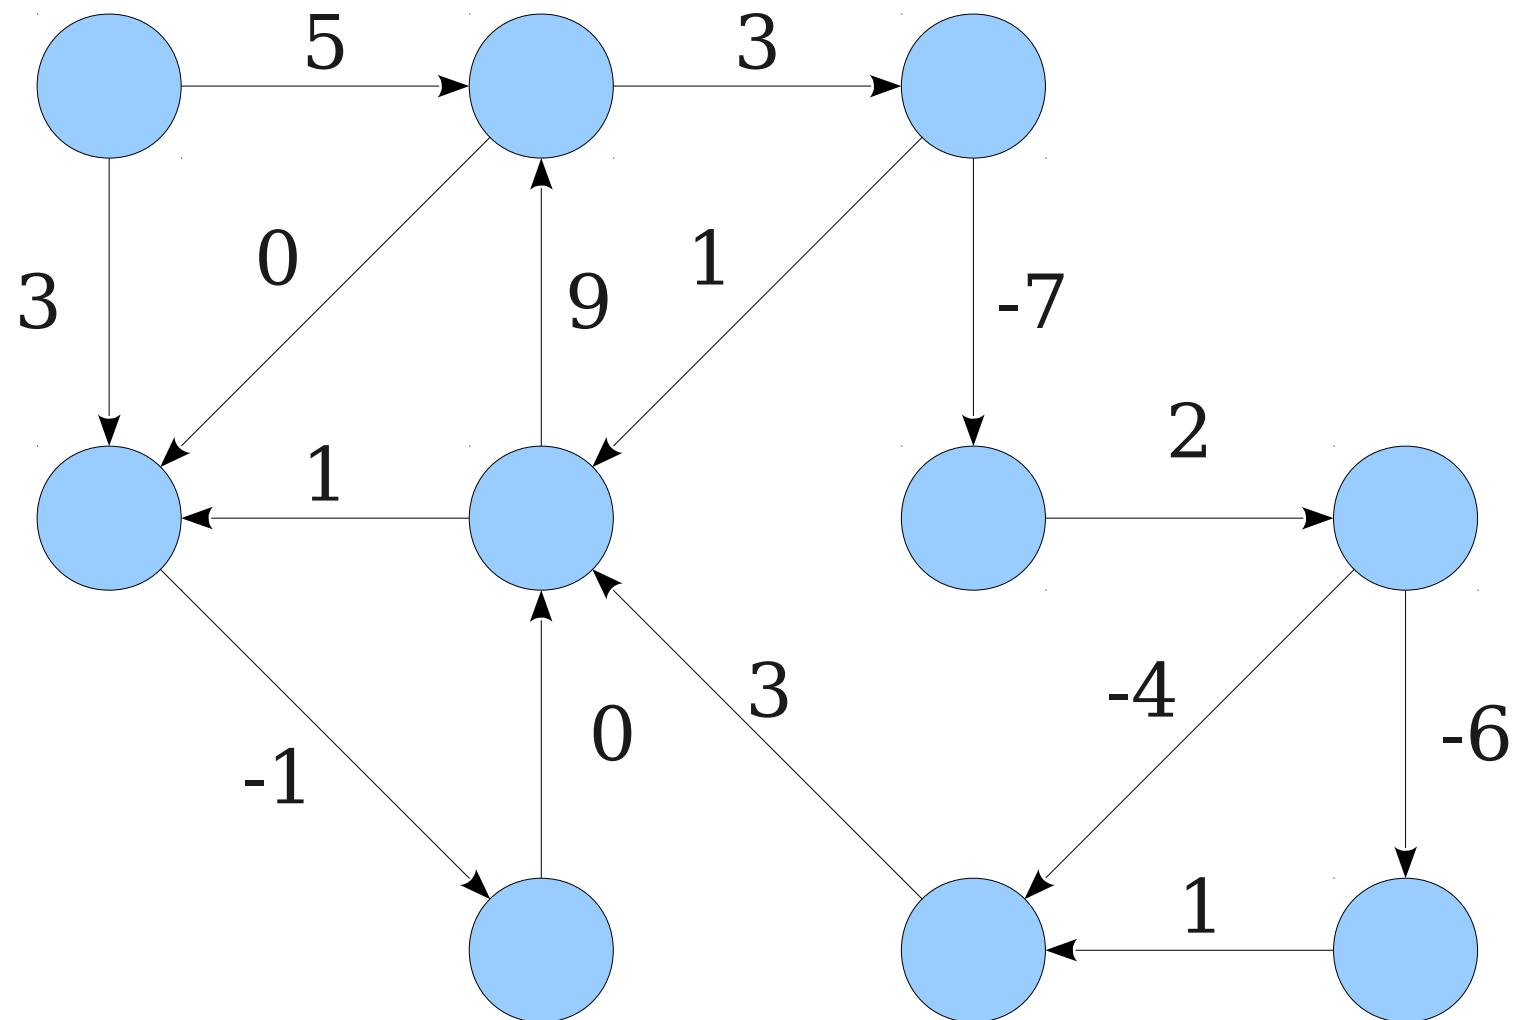
- Array implementation optimal for dense graphs. $\leftarrow \Theta(n^2)$ edges
- Binary heap much faster for sparse graphs. $\leftarrow \Theta(n)$ edges
- 4-way heap worth the trouble in performance-critical situations.

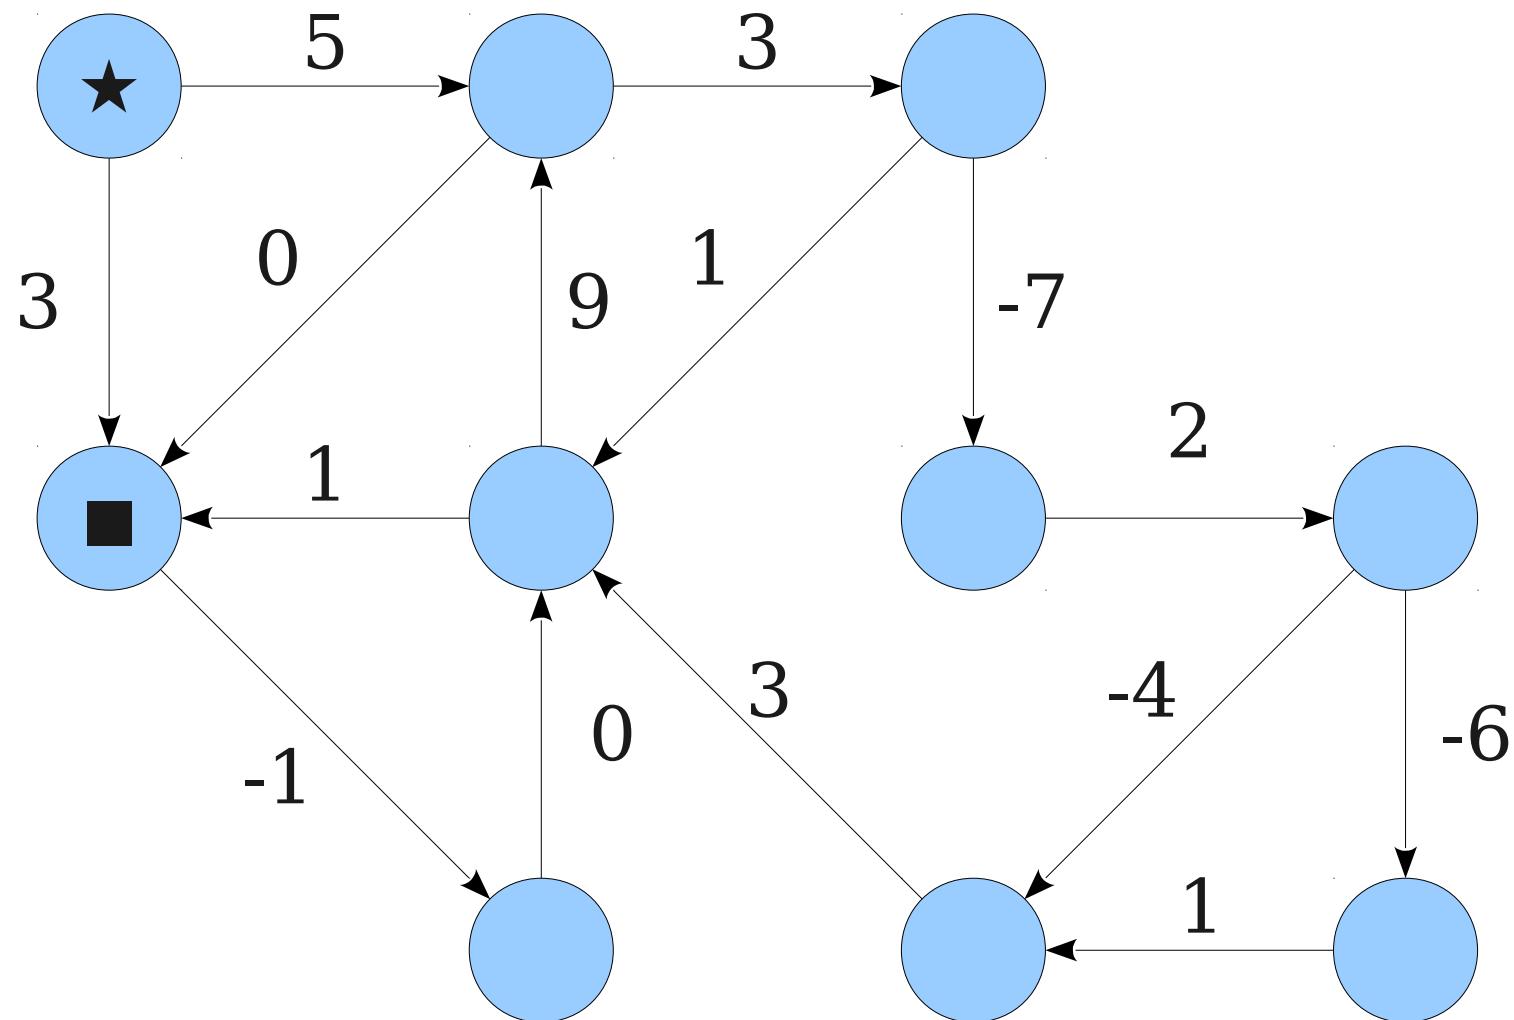
priority queue	INSERT	DELETE-MIN	DECREASE-KEY	total
node-indexed array ($A[i] = \text{priority of } i$)	$O(1)$	$O(n)$	$O(1)$	$O(n^2)$
binary heap	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(m \log n)$
d-way heap (Johnson 1975)	$O(d \log_d n)$	$O(d \log_d n)$	$O(\log_d n)$	$O(m \log_{m/n} n)$
Fibonacci heap (Fredman-Tarjan 1984)	$O(1)$	$O(\log n)^\dagger$	$O(1)^\dagger$	$O(m + n \log n)$
integer priority queue (Thorup 2004)	$O(1)$	$O(\log \log n)$	$O(1)$	$O(m + n \log \log n)$

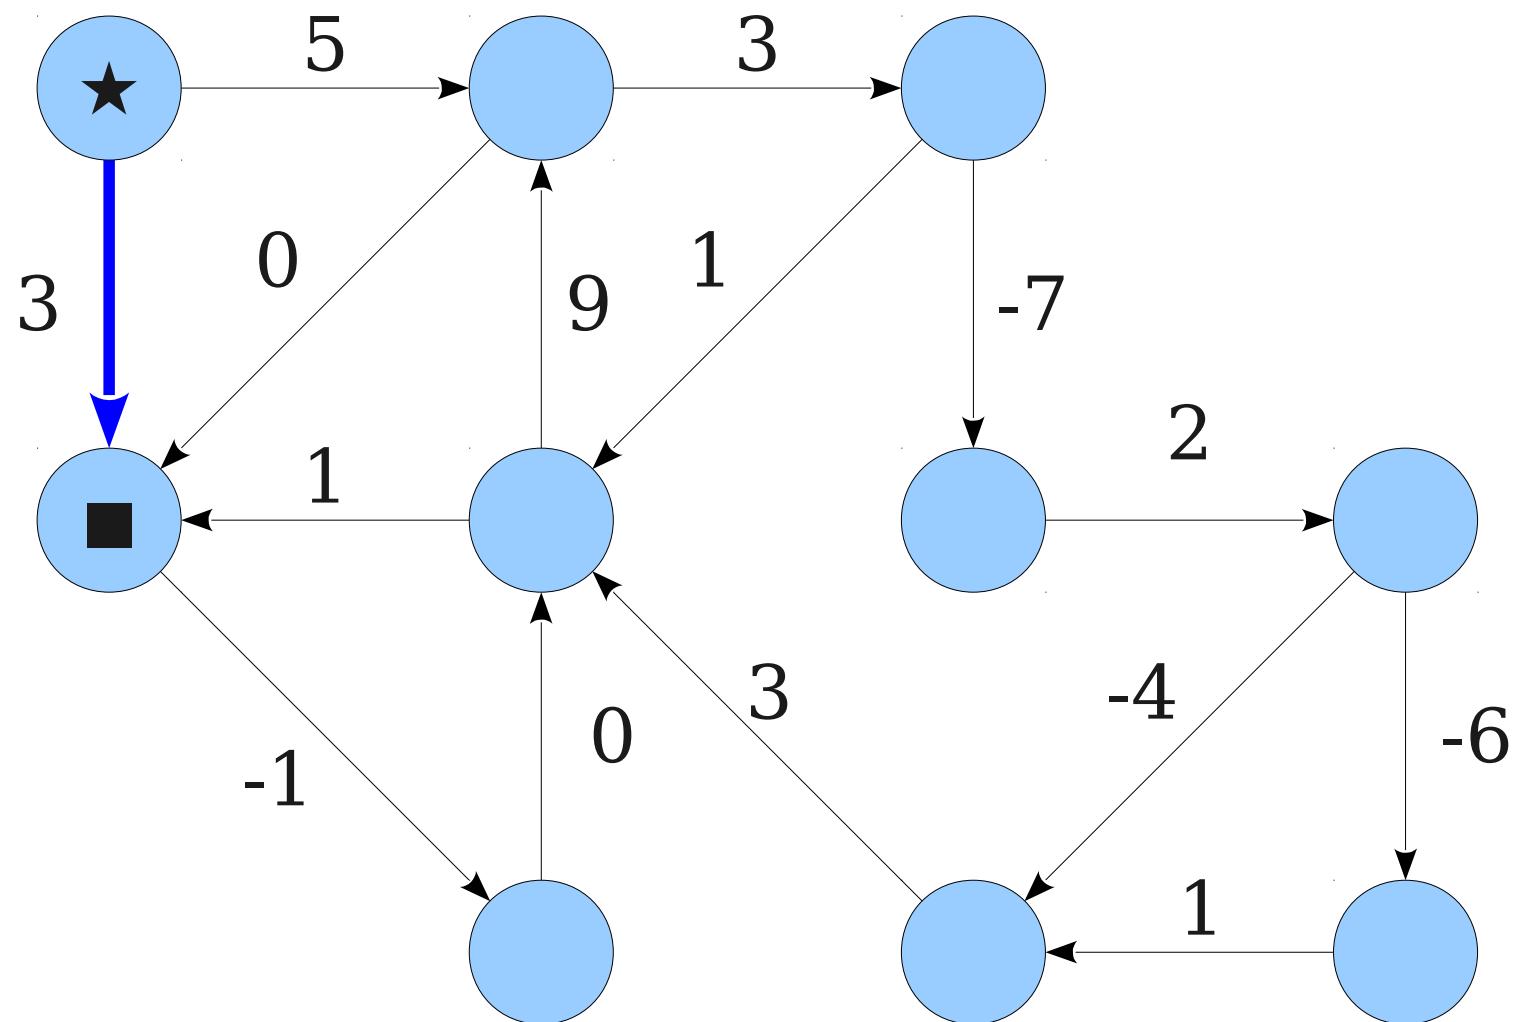
assumes $m \geq n$

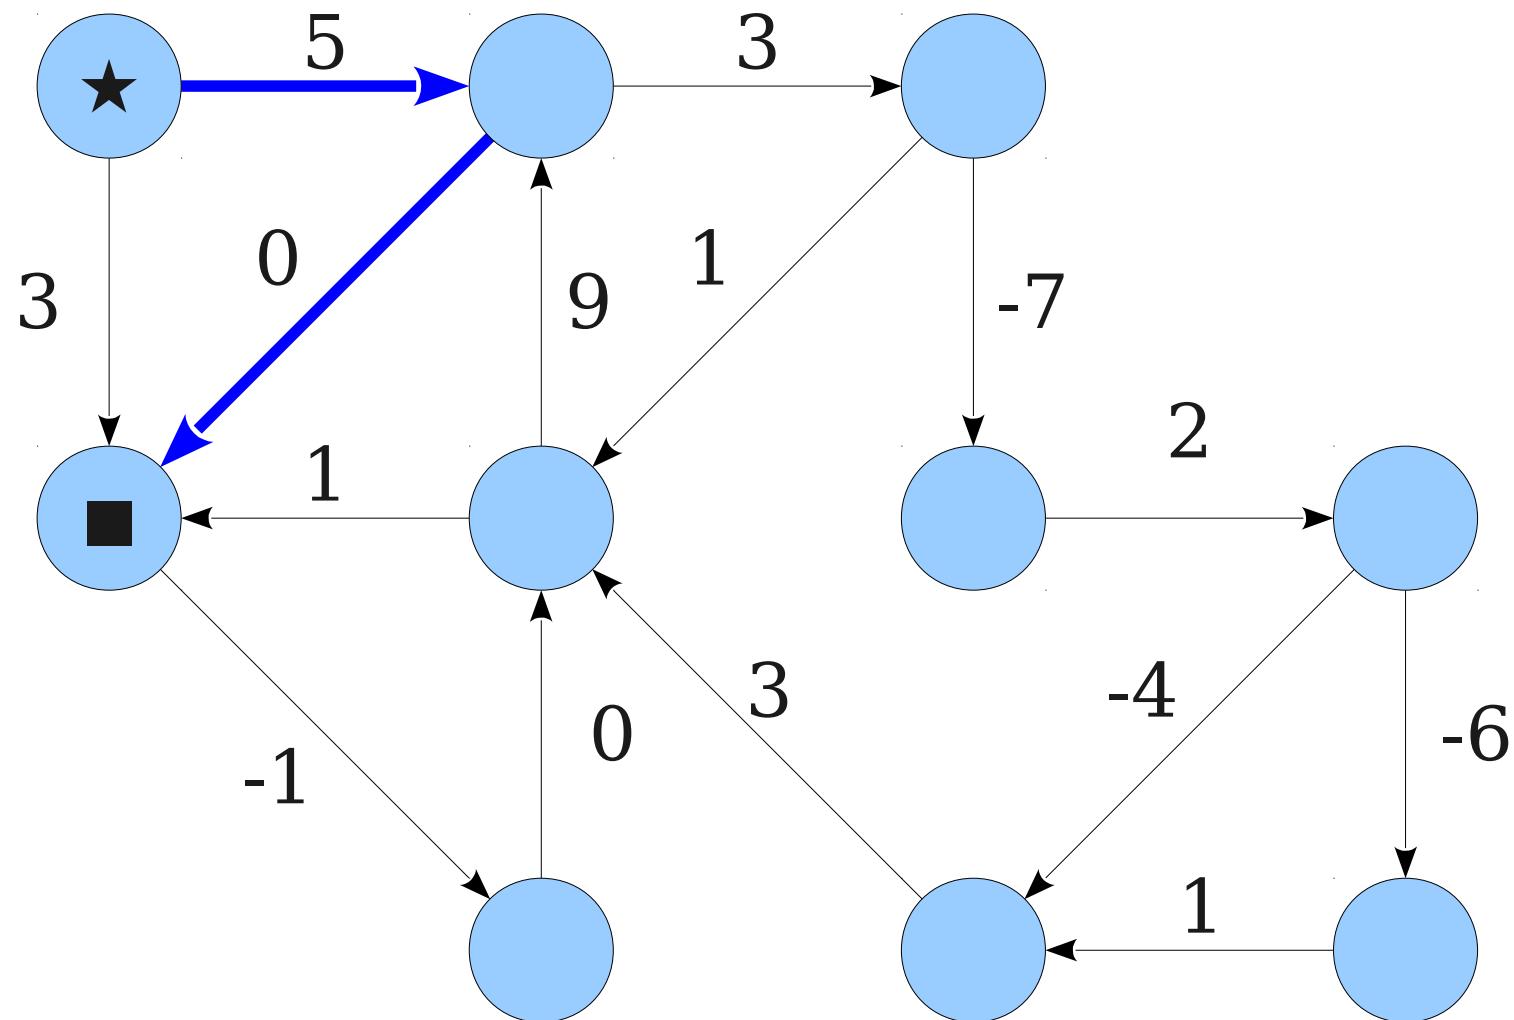
\dagger amortized

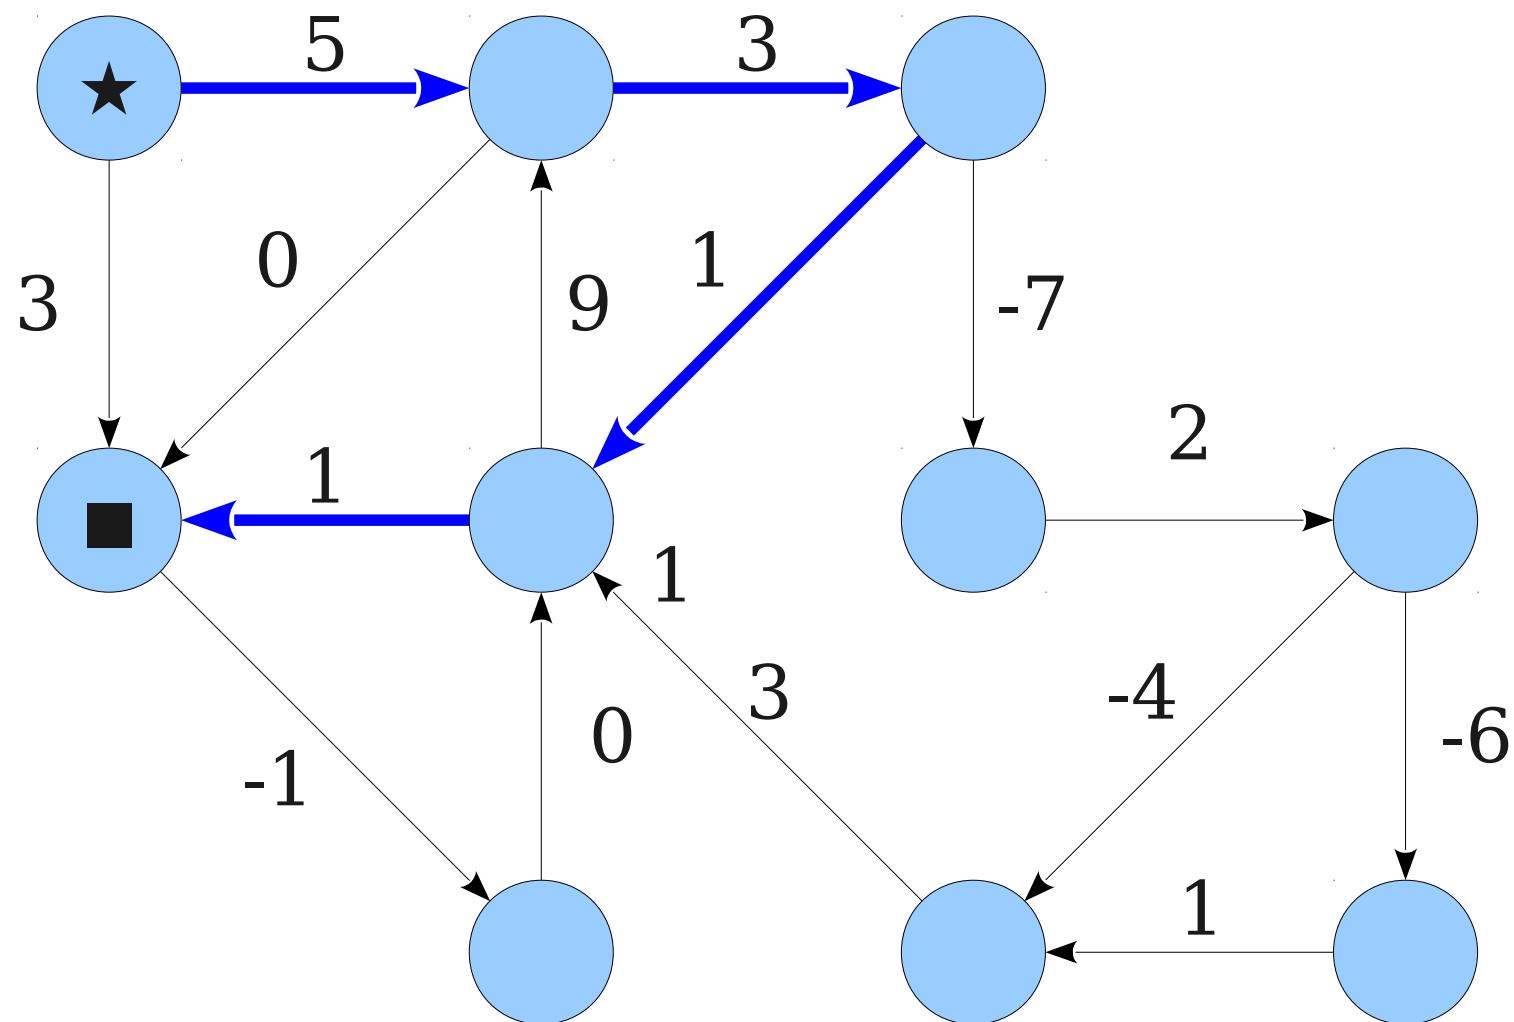
Negative Edge Weights

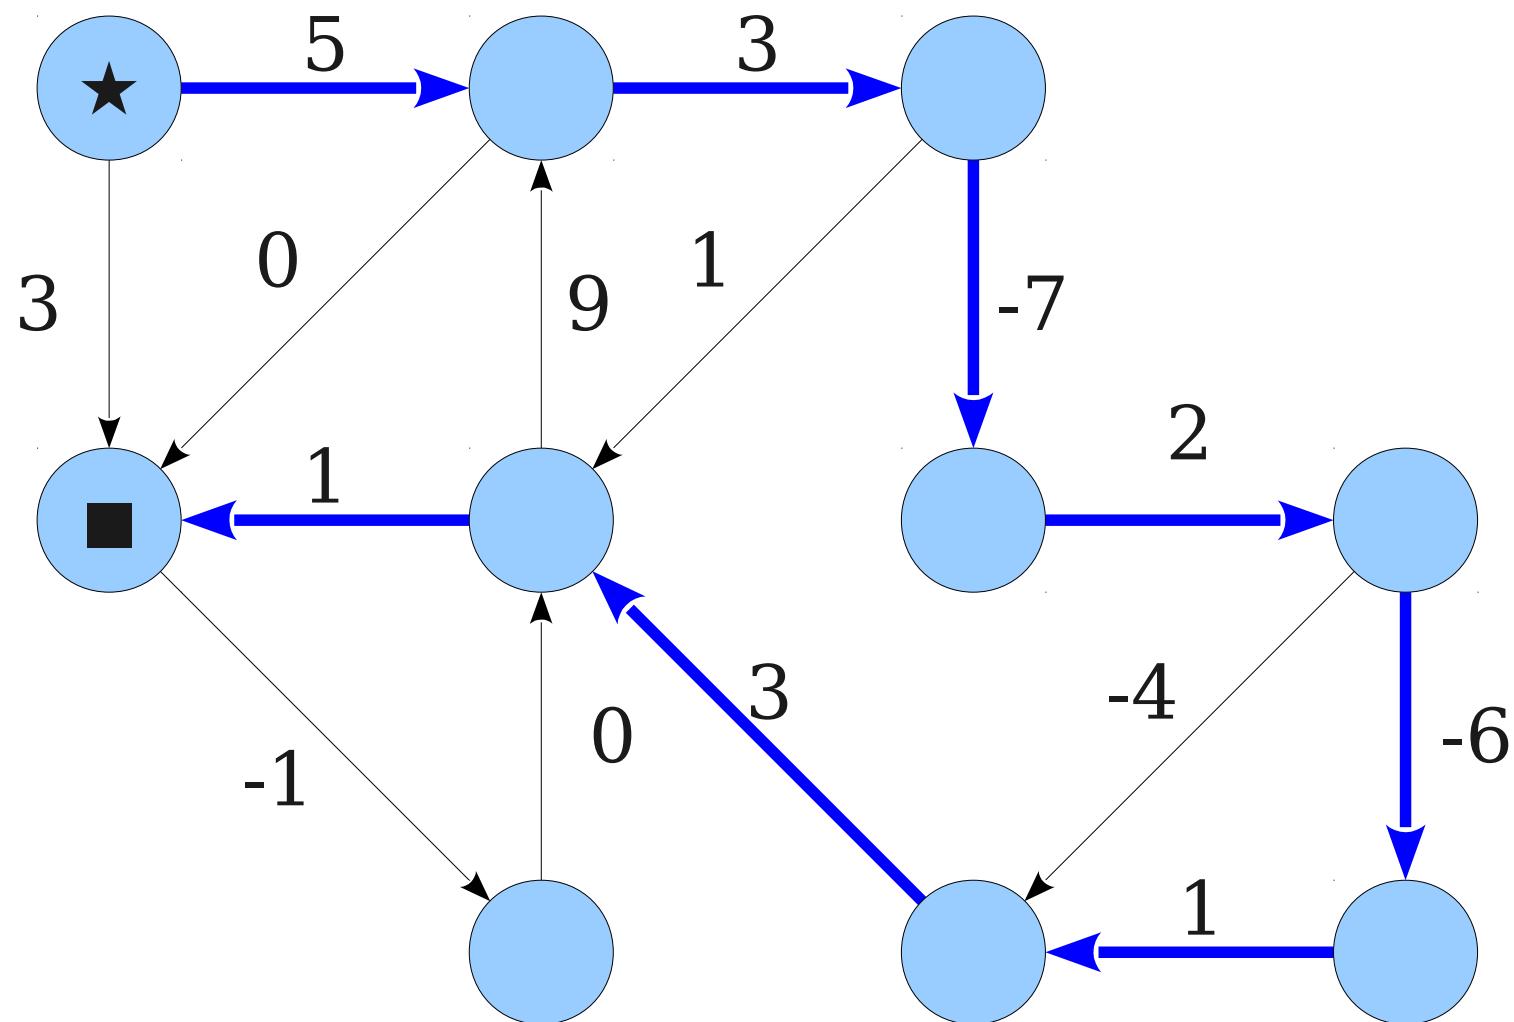






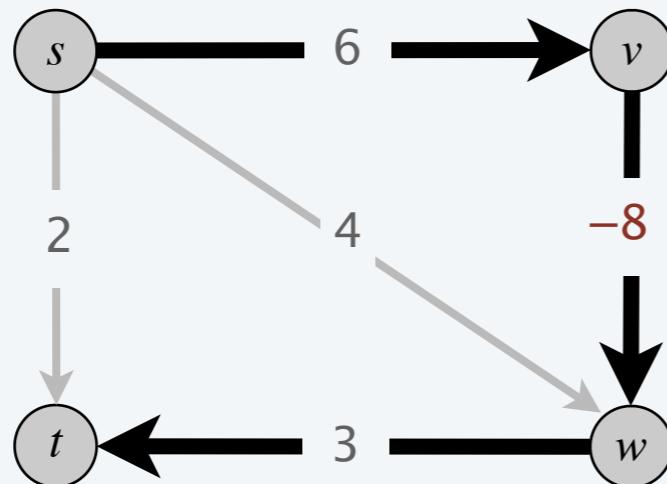






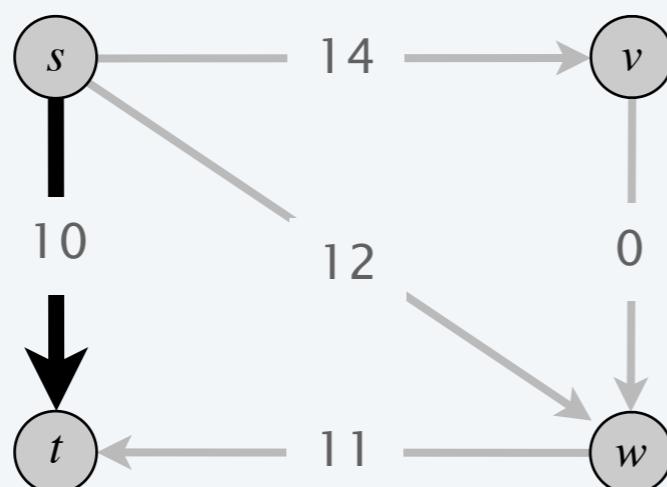
Shortest paths with negative weights: failed attempts

Dijkstra. May not produce shortest paths when edge lengths are negative.



Dijkstra selects the vertices in the order s, t, w, v
But shortest path from s to t is $s \rightarrow v \rightarrow w \rightarrow t$.

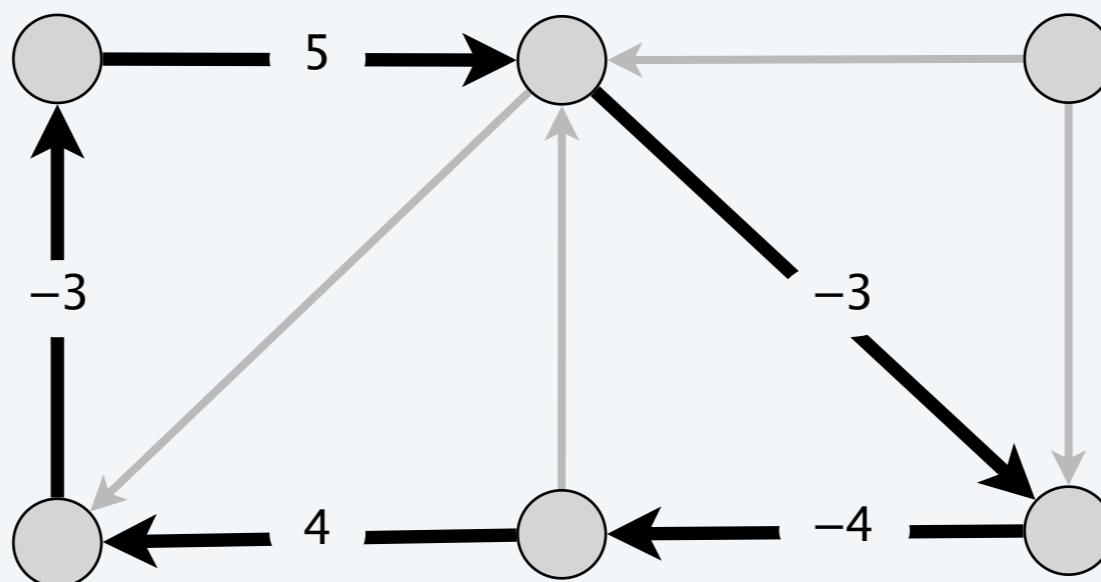
Reweighting. Adding a constant to every edge length does not necessarily make Dijkstra's algorithm produce shortest paths.



Adding 8 to each edge weight changes the shortest path from $s \rightarrow v \rightarrow w \rightarrow t$ to $s \rightarrow t$.

Negative cycles

Def. A **negative cycle** is a directed cycle for which the sum of its edge lengths is negative.

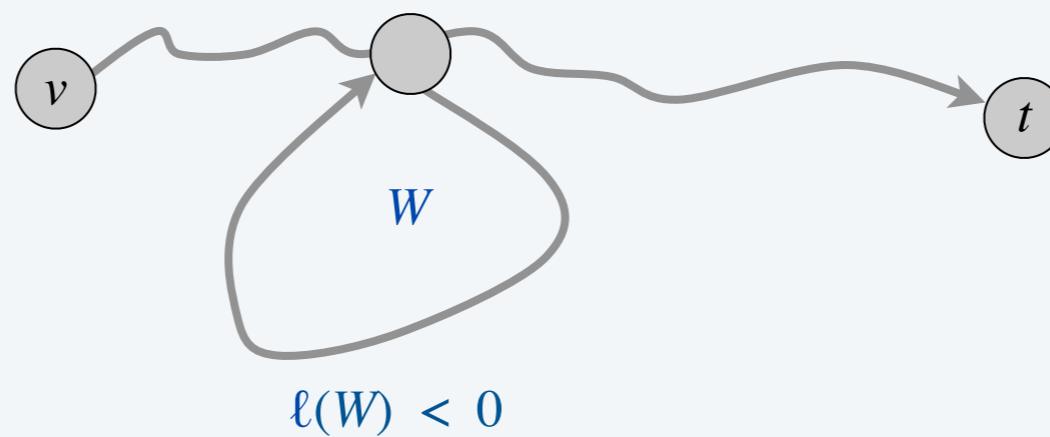


a **negative cycle** W : $\ell(W) = \sum_{e \in W} \ell_e < 0$

Shortest paths and negative cycles

Lemma 1. If some $v \rightsquigarrow t$ path contains a negative cycle, then there does not exist a shortest $v \rightsquigarrow t$ path.

Pf. If there exists such a cycle W , then can build a $v \rightsquigarrow t$ path of arbitrarily negative length by detouring around W as many times as desired. ▀

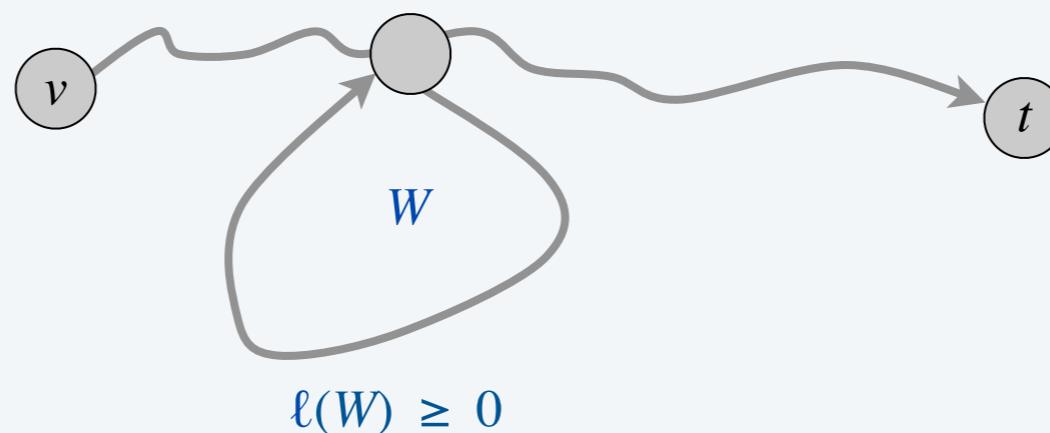


Shortest paths and negative cycles

Lemma 2. If G has no negative cycles, then there exists a shortest $v \rightsquigarrow t$ path that is simple (and has $\leq n - 1$ edges).

Pf.

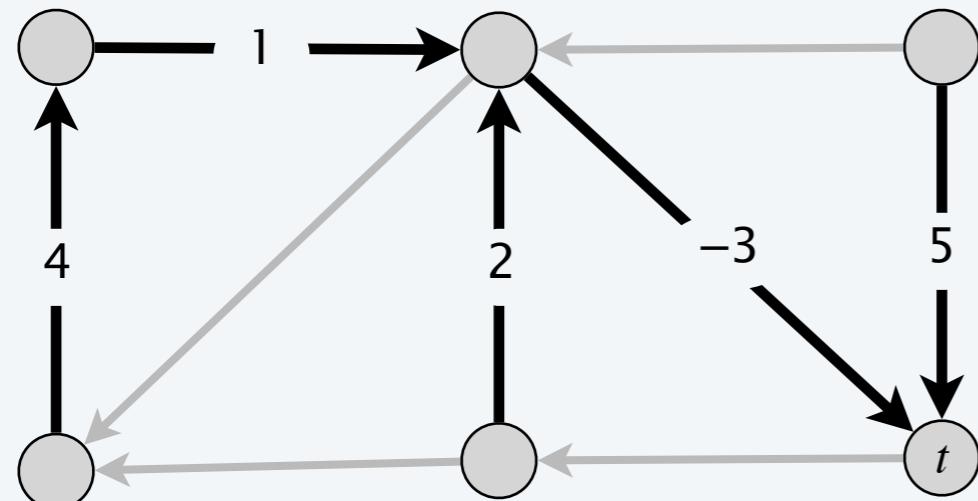
- Among all shortest $v \rightsquigarrow t$ paths, consider one that uses the fewest edges.
- If that path P contains a directed cycle W , can remove the portion of P corresponding to W without increasing its length. ▀



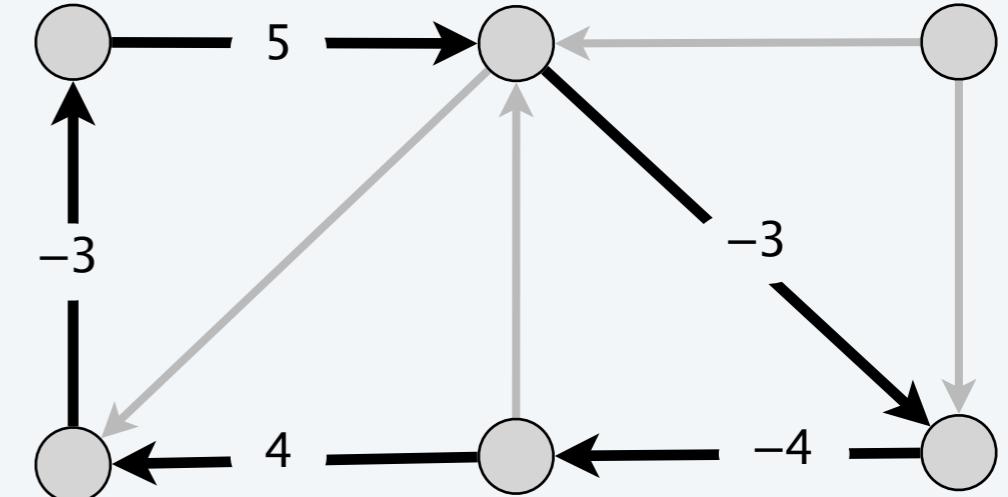
Shortest-paths and negative-cycle problems

Single-destination shortest-paths problem. Given a digraph $G = (V, E)$ with edge lengths ℓ_{vw} (but no negative cycles) and a distinguished node t , find a shortest $v \rightsquigarrow t$ path for every node v .

Negative-cycle problem. Given a digraph $G = (V, E)$ with edge lengths ℓ_{vw} , find a negative cycle (if one exists).

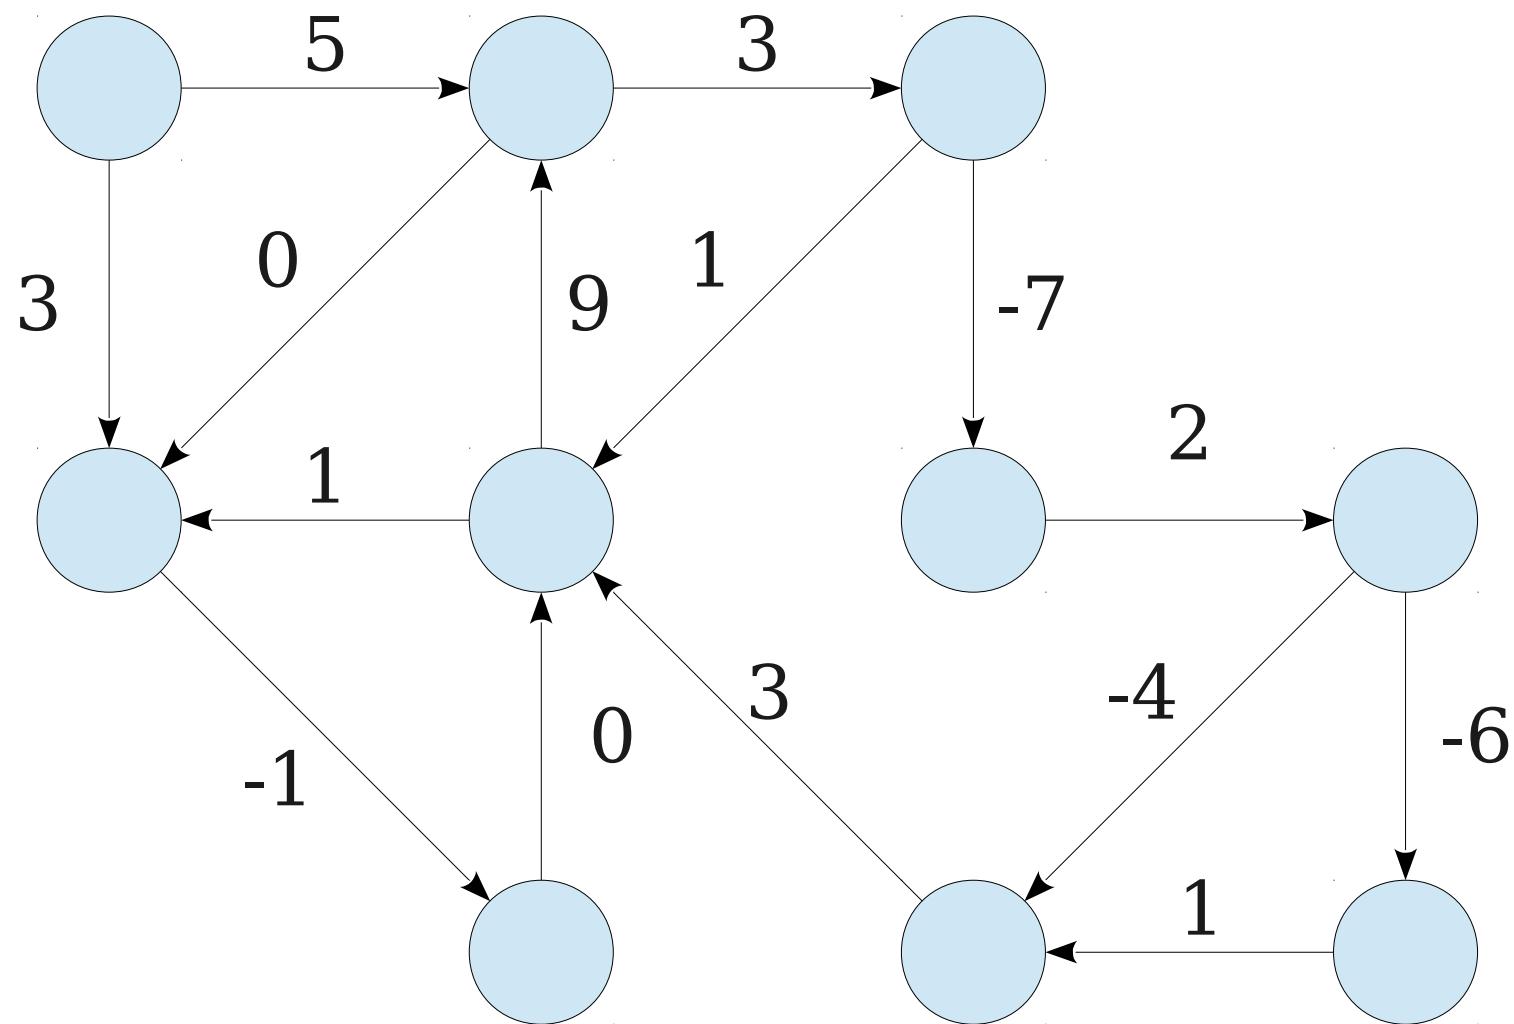


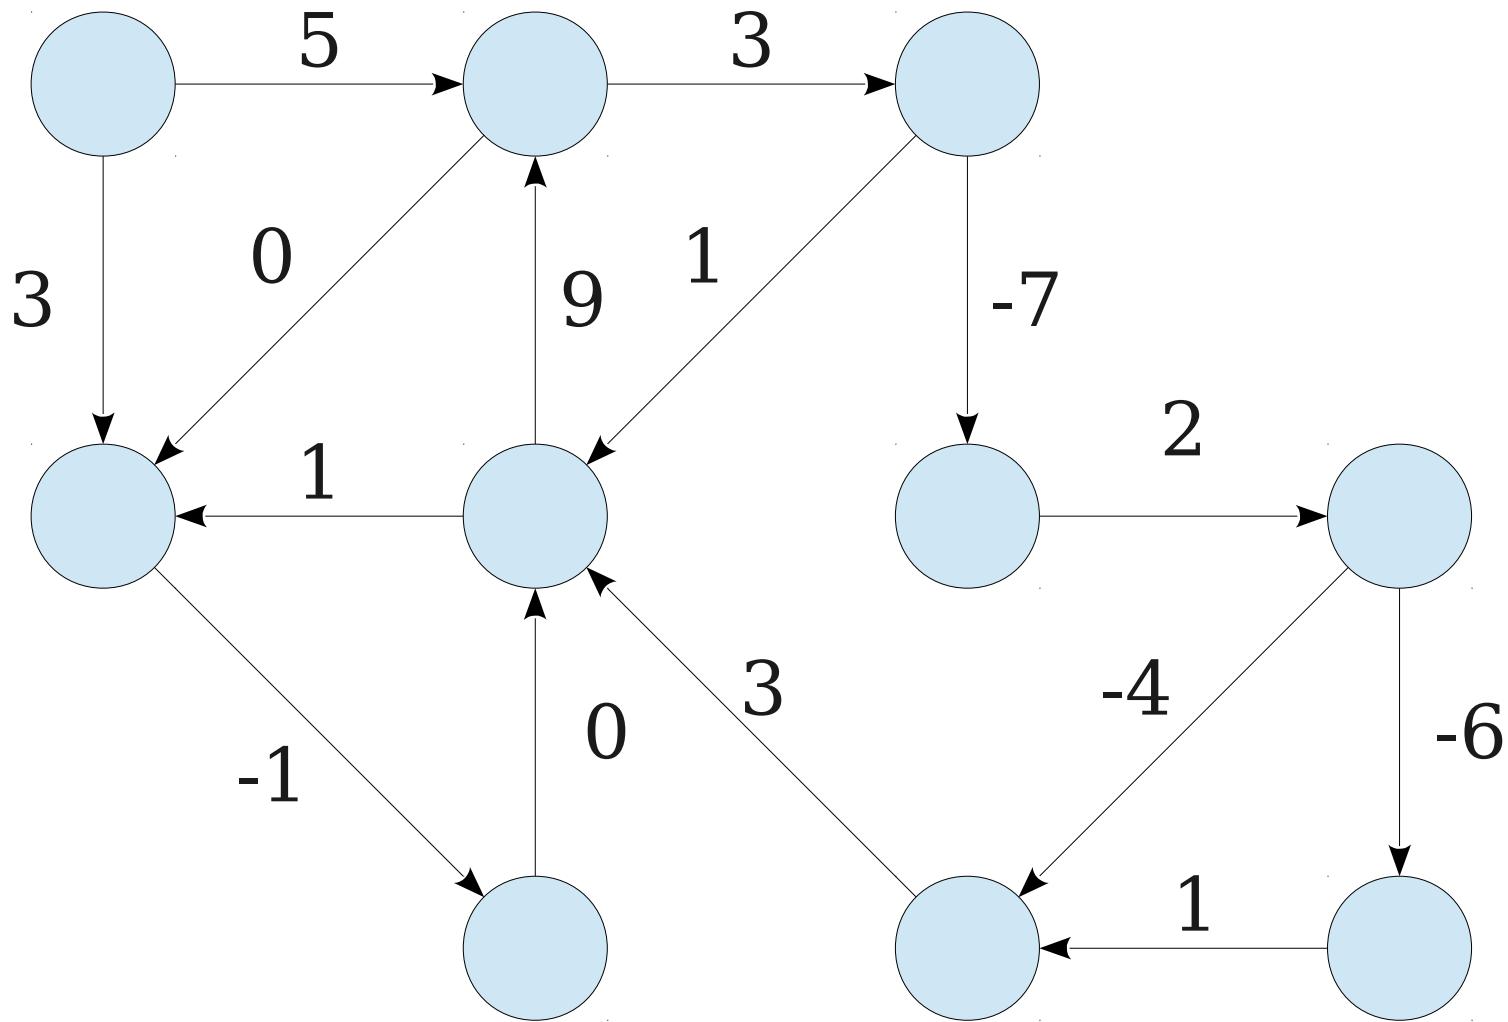
shortest-paths tree



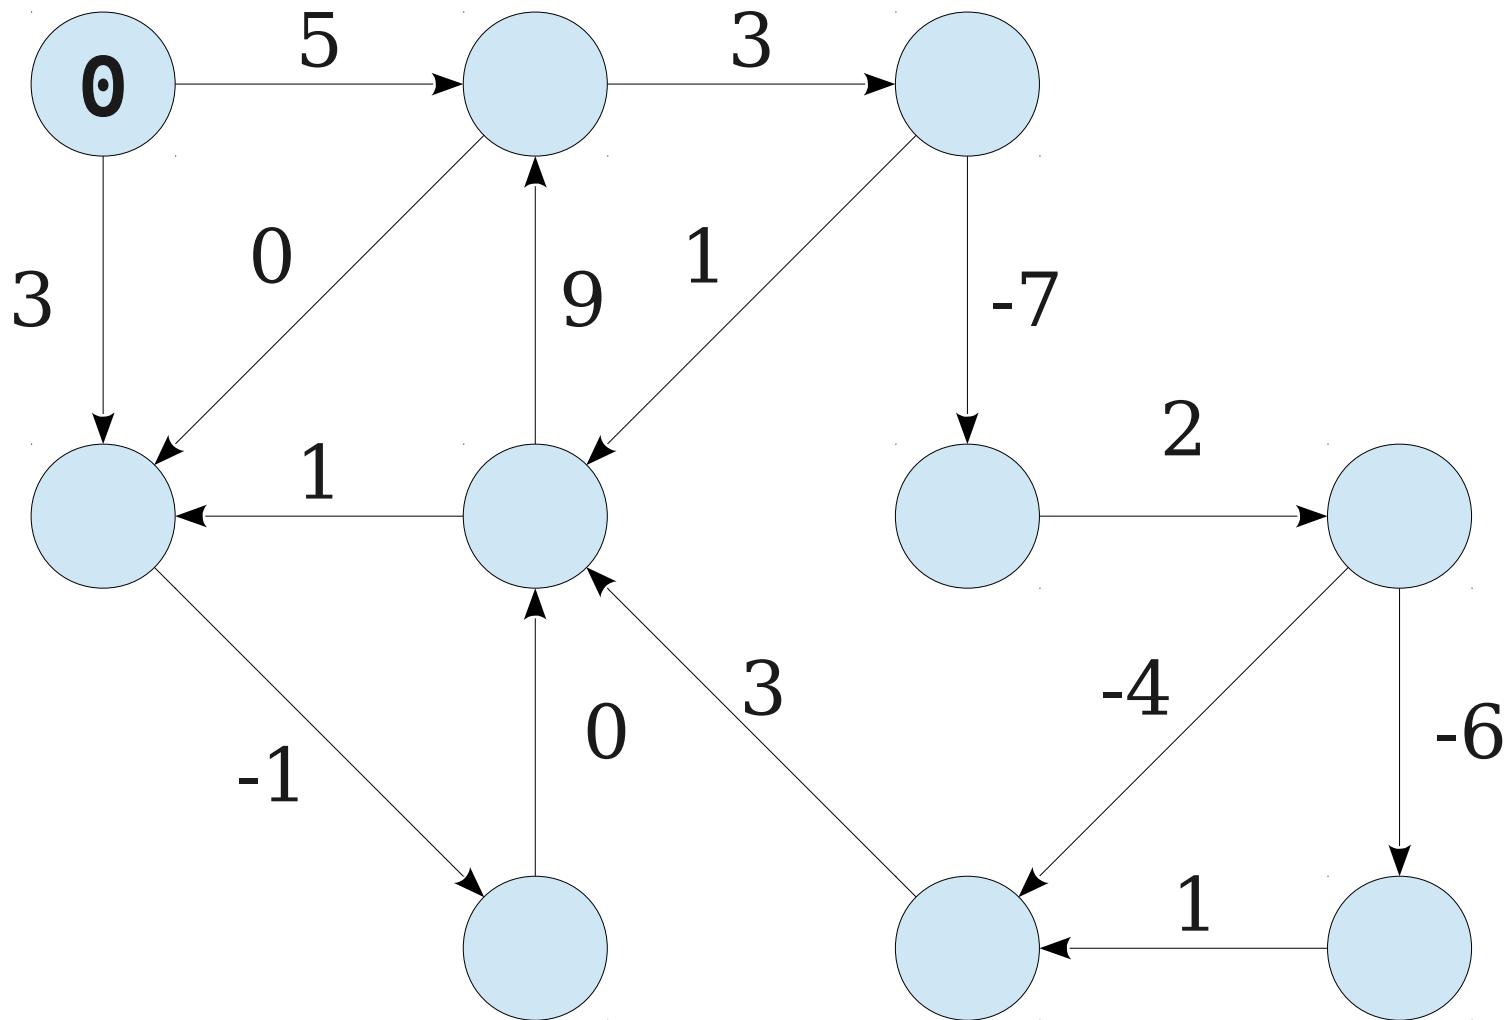
negative cycle

A Different Intuition

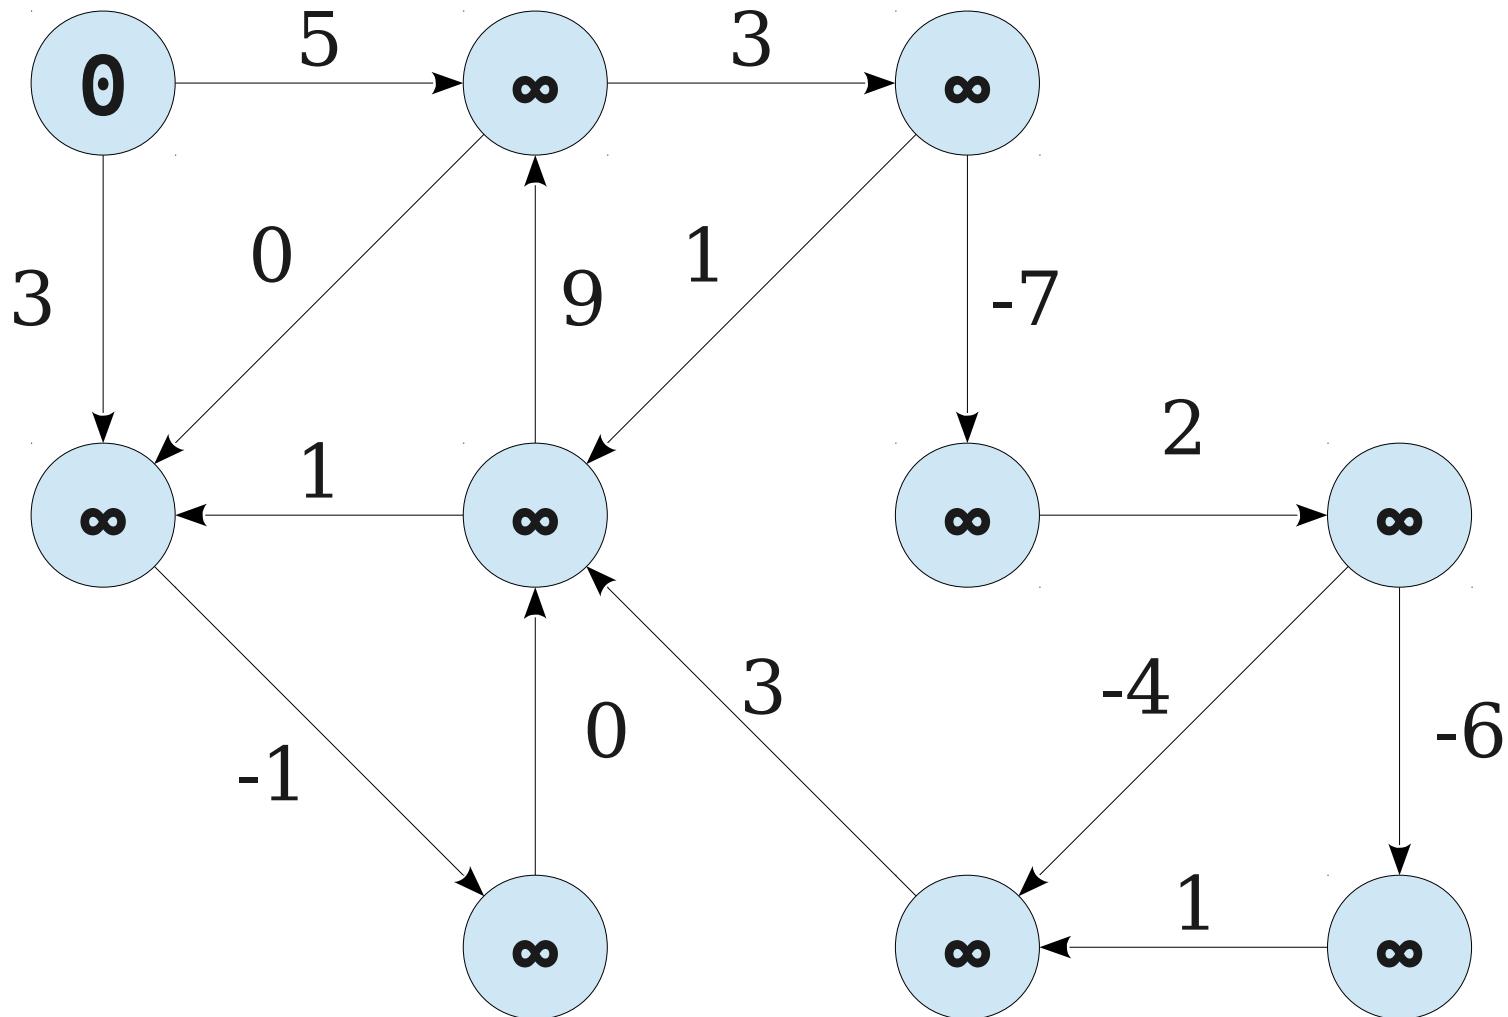




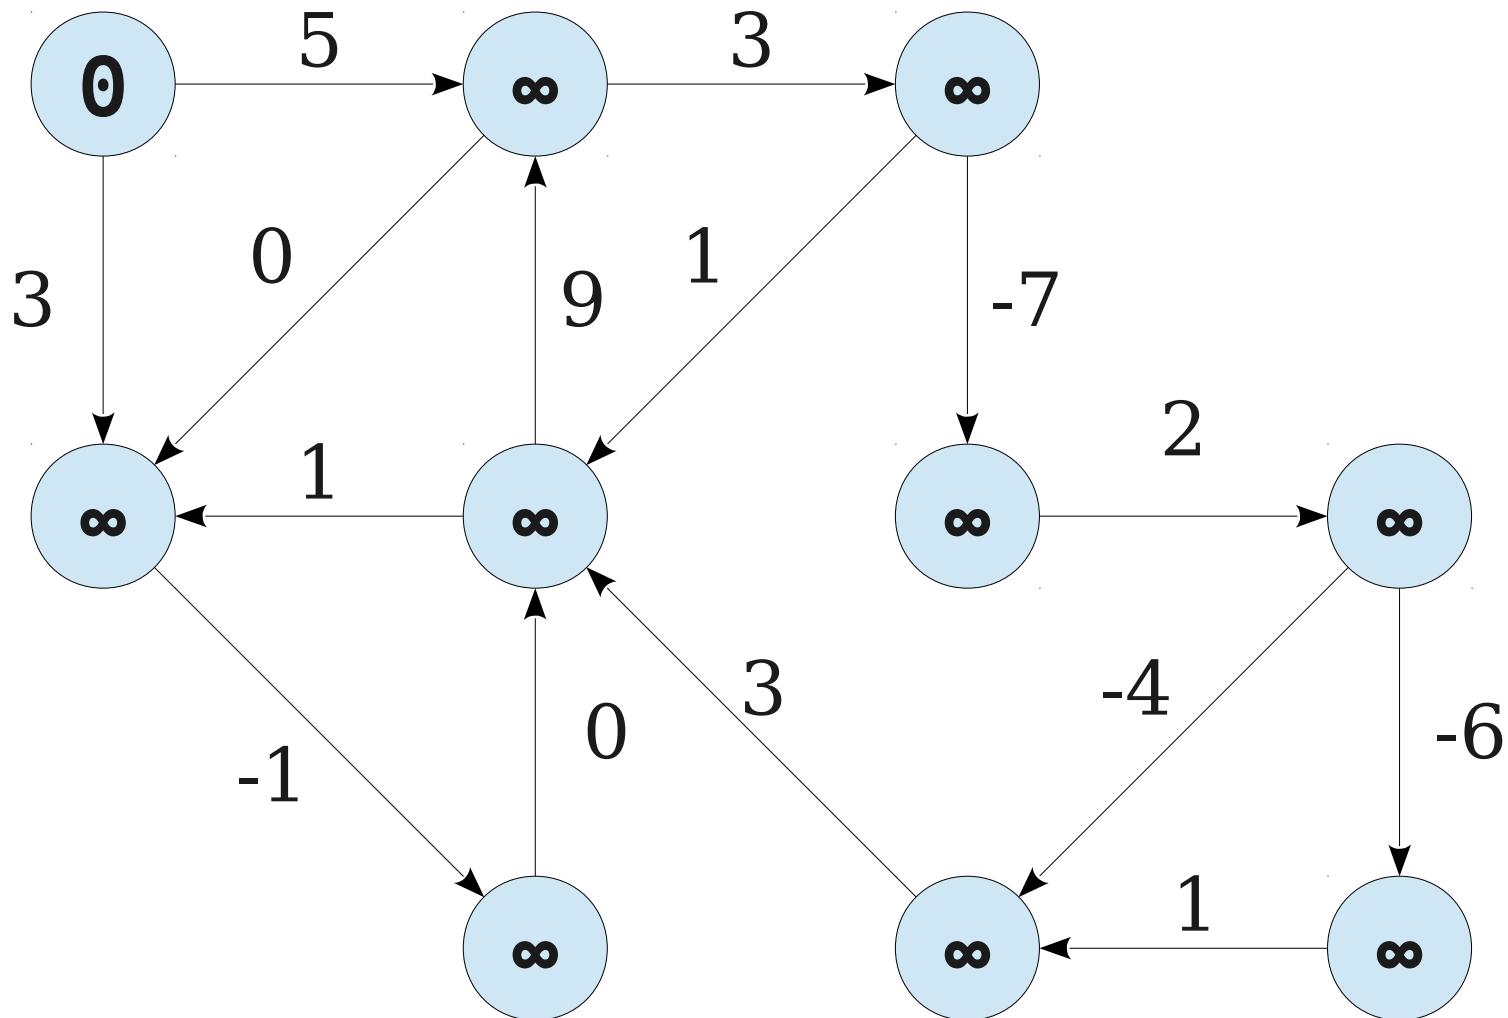
Consider paths of length **0**.



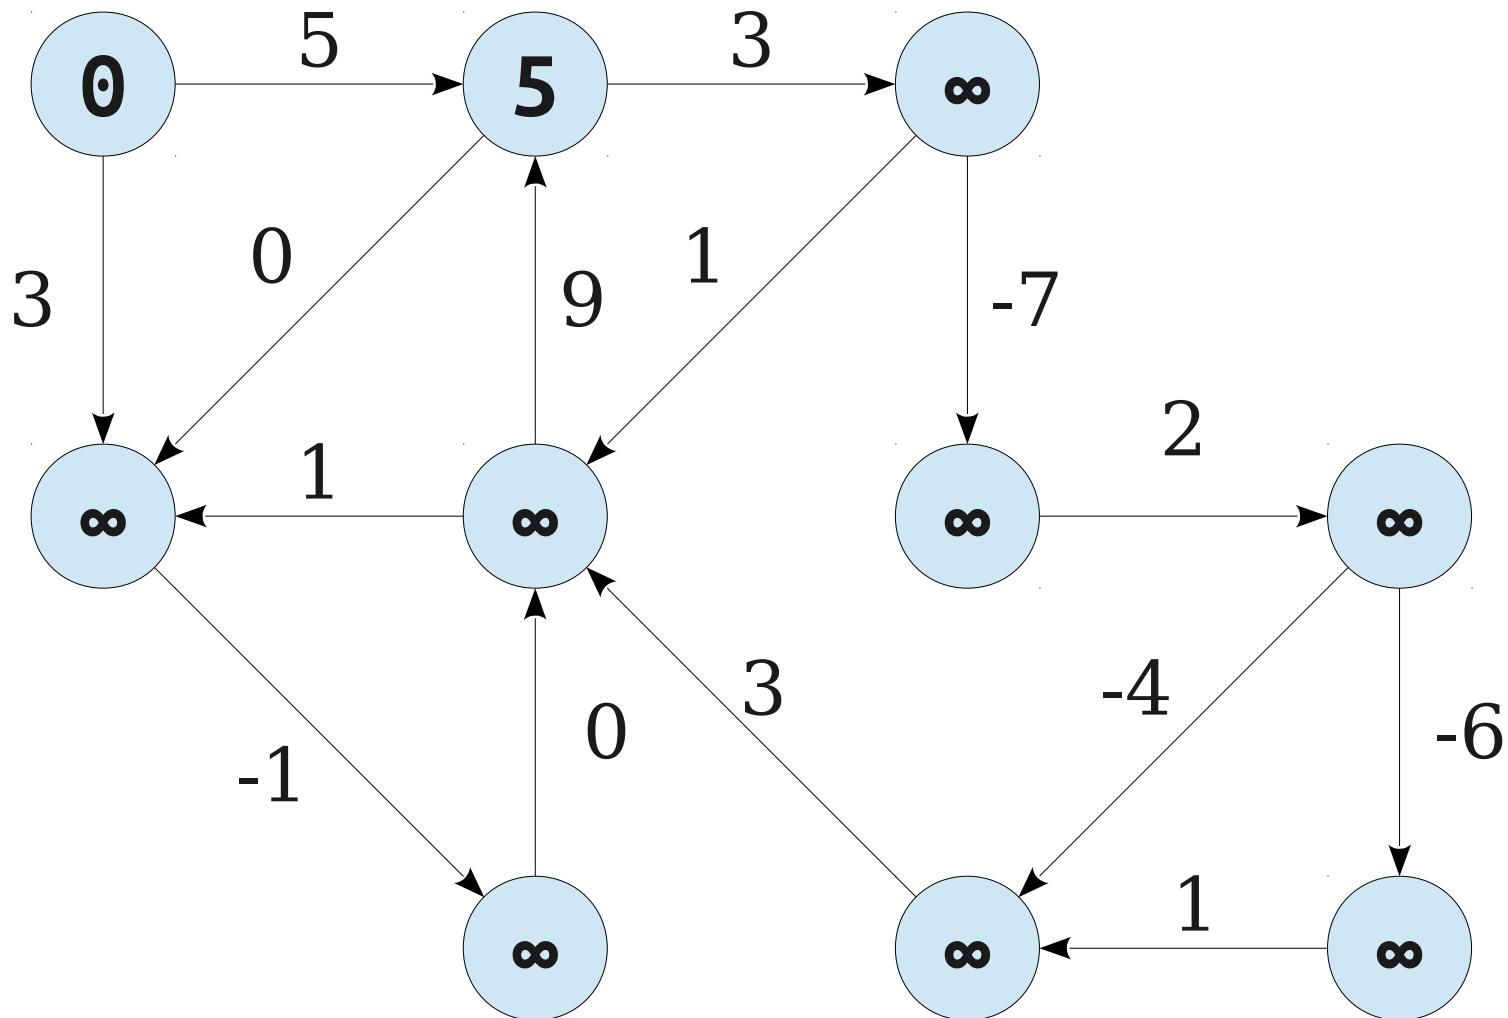
Consider paths of length **0**.



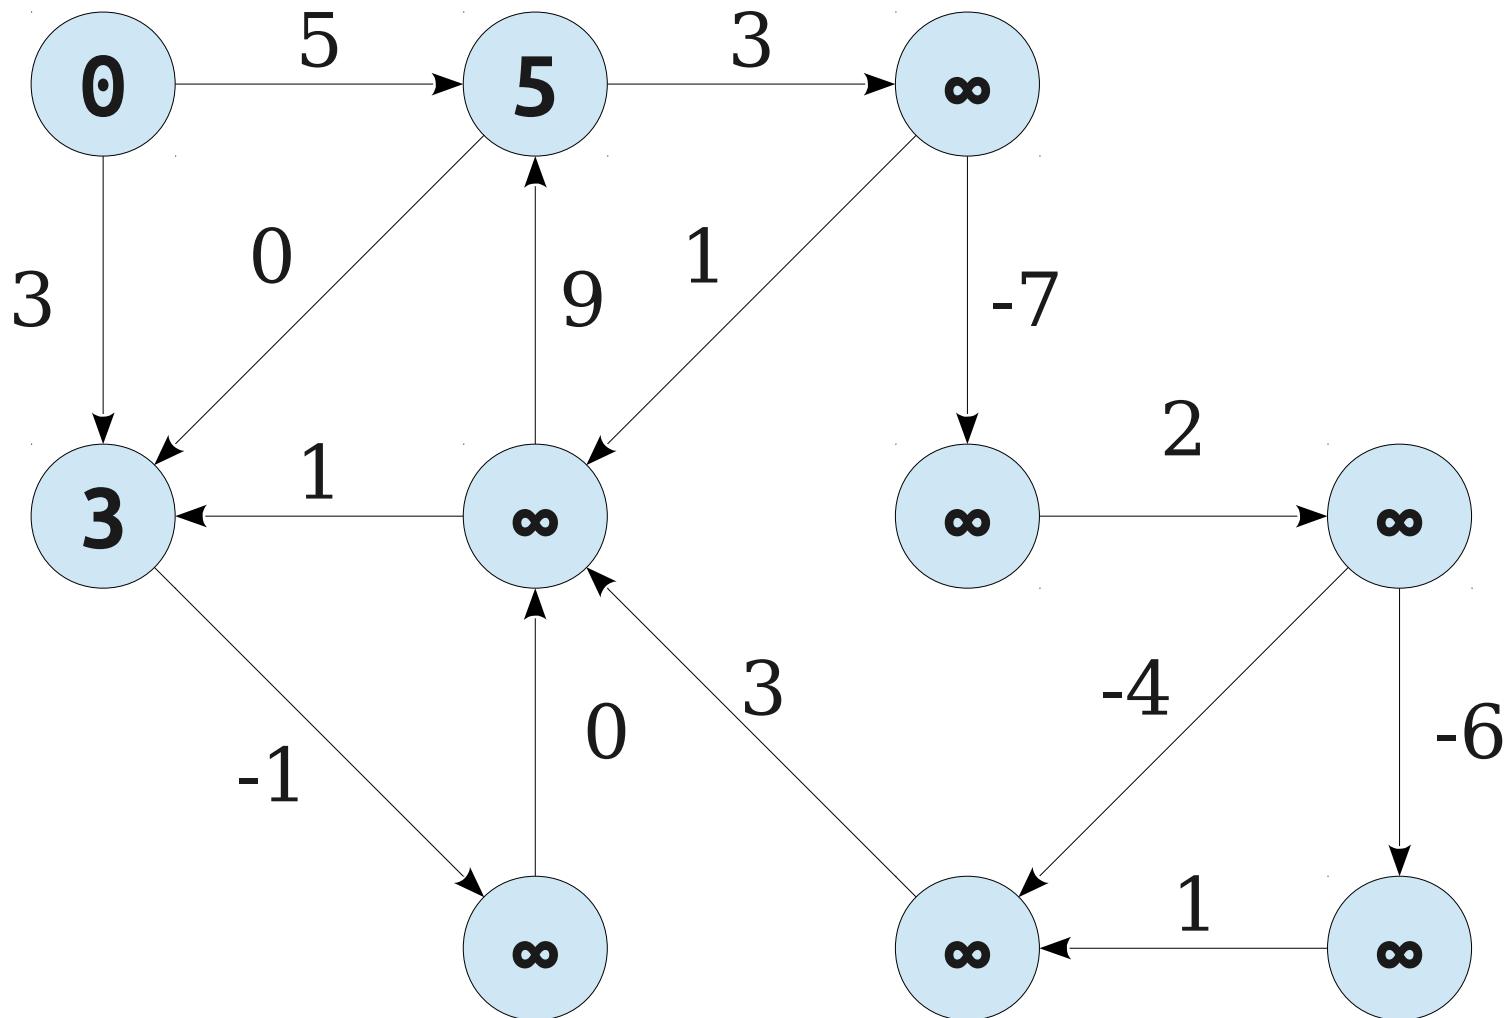
Consider paths of length **0**.



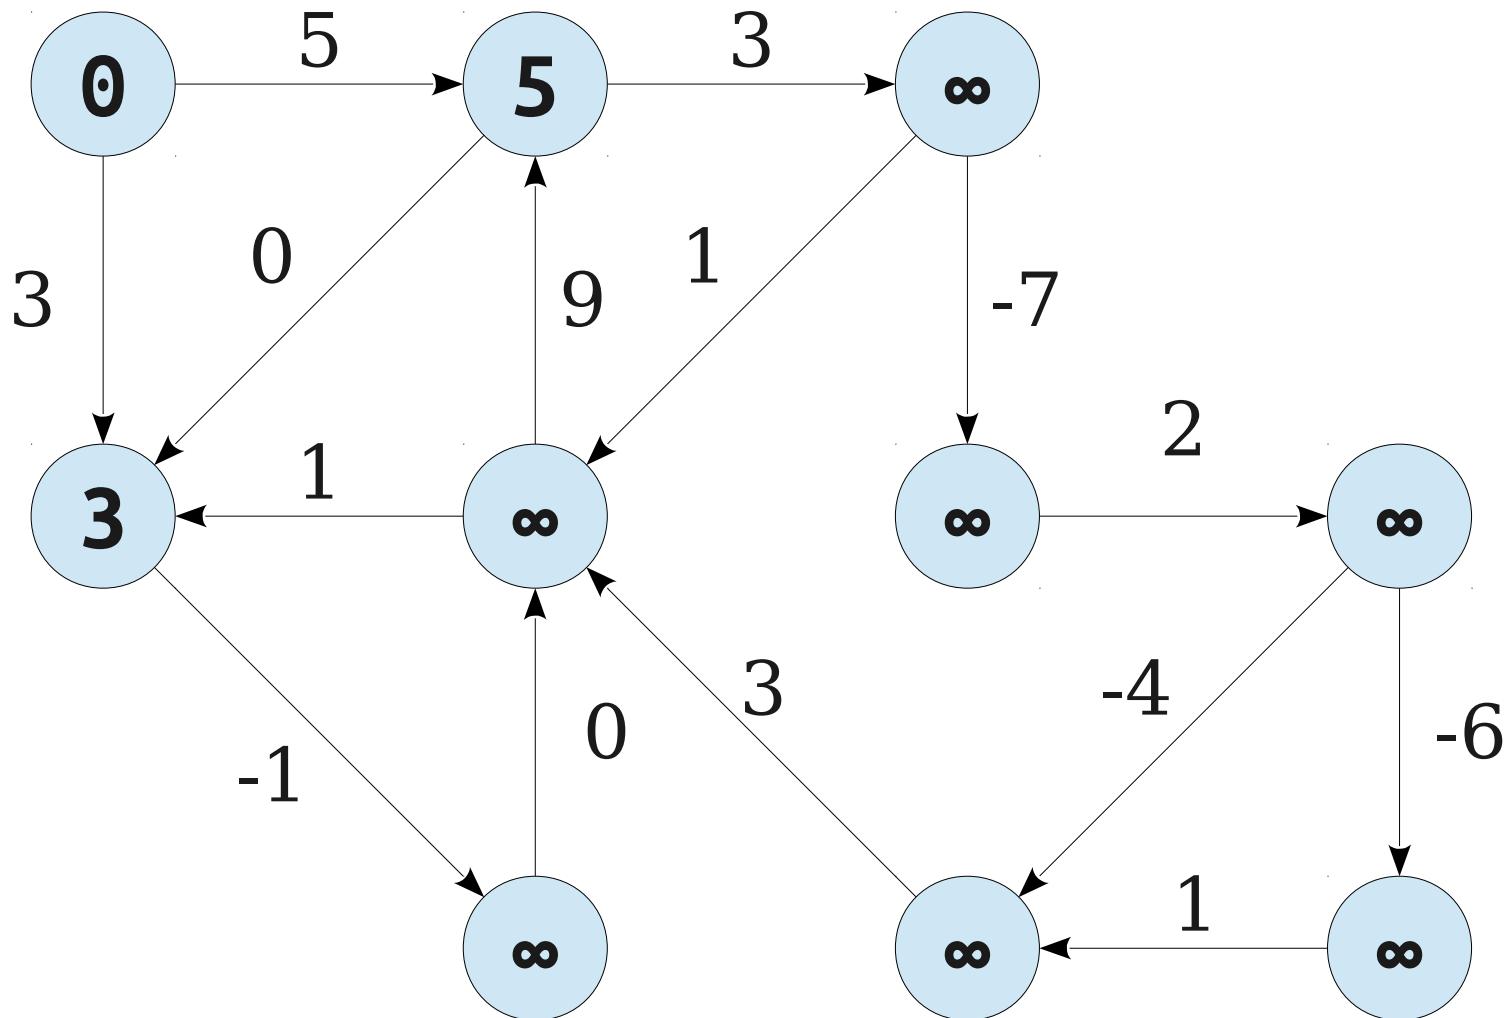
Consider paths of length **1**.



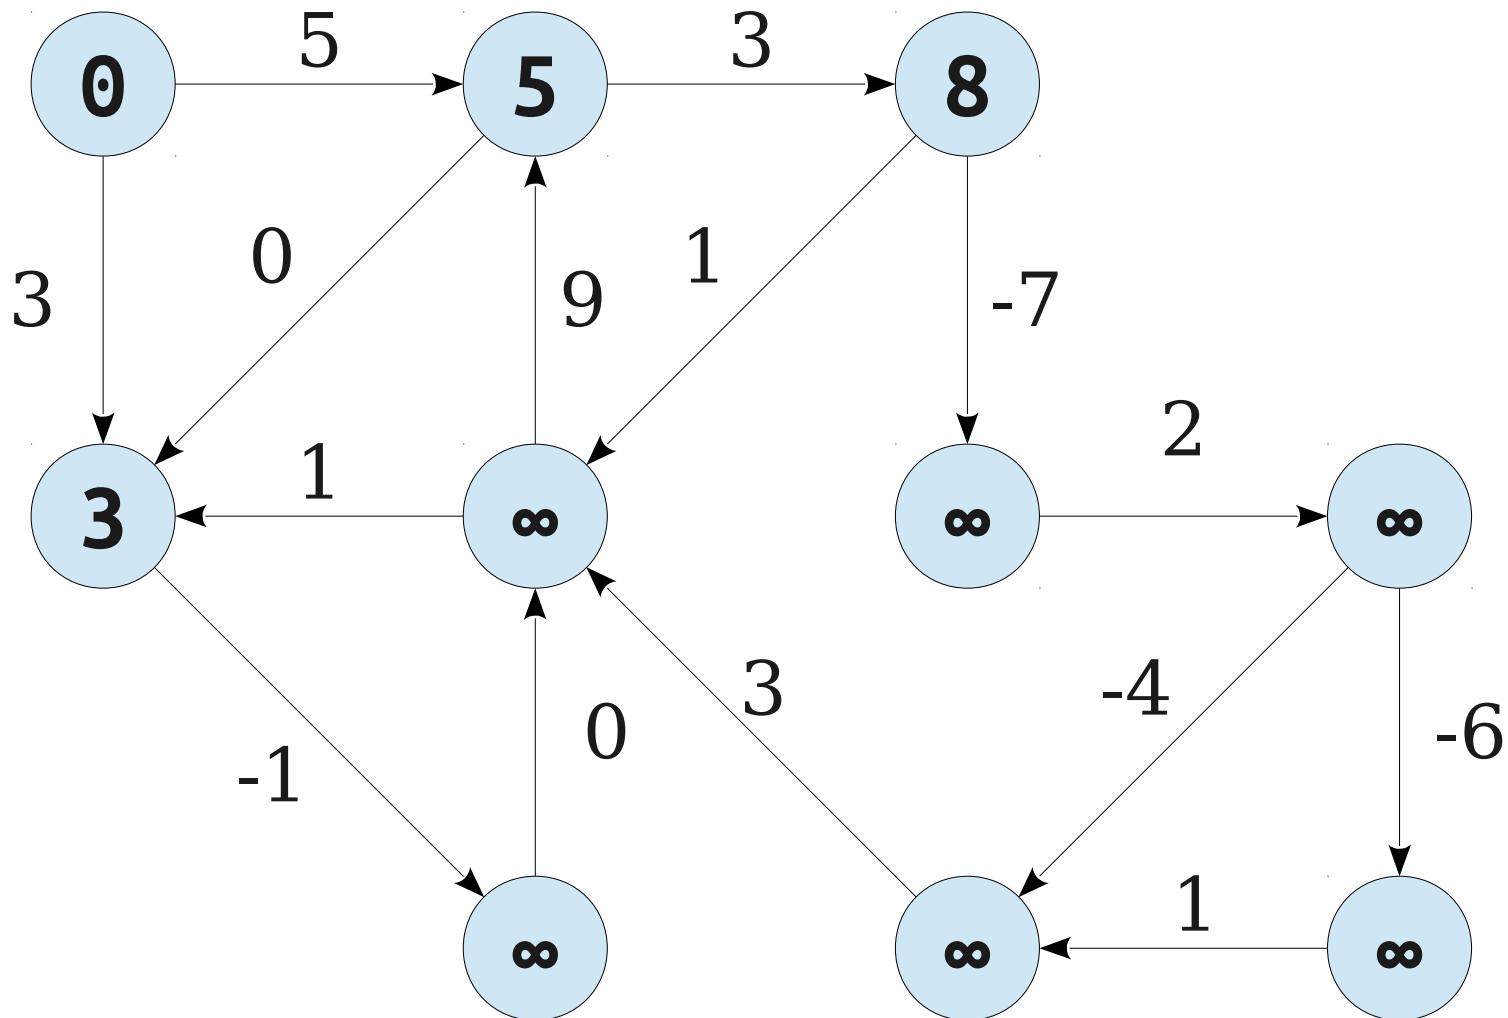
Consider paths of length **1**.



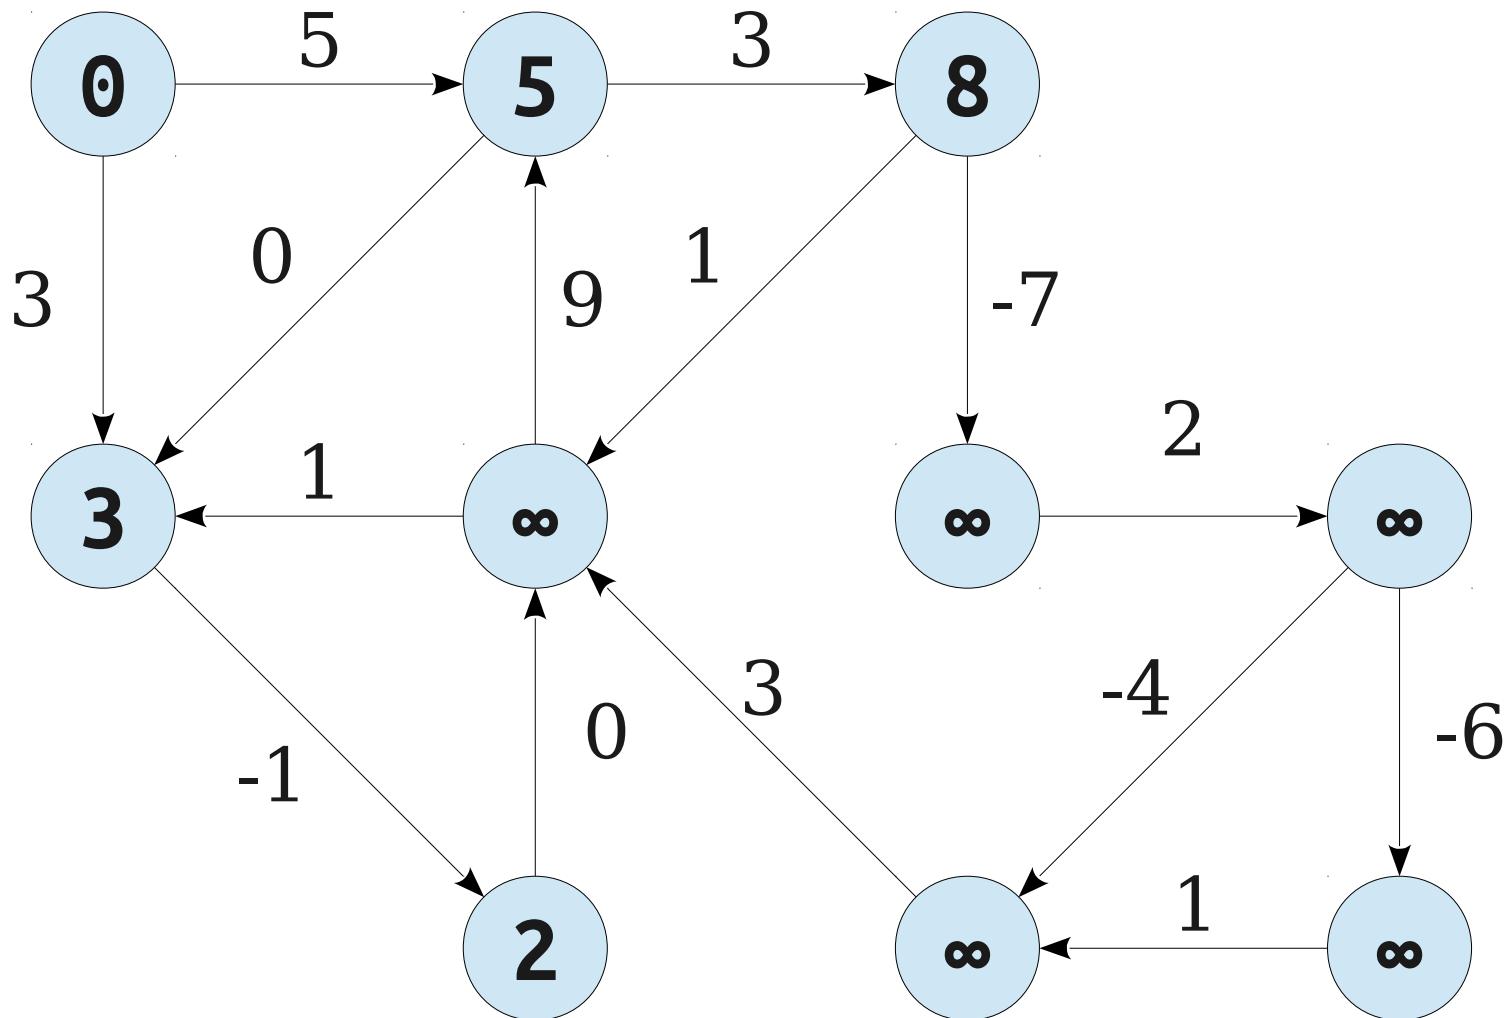
Consider paths of length **1**.



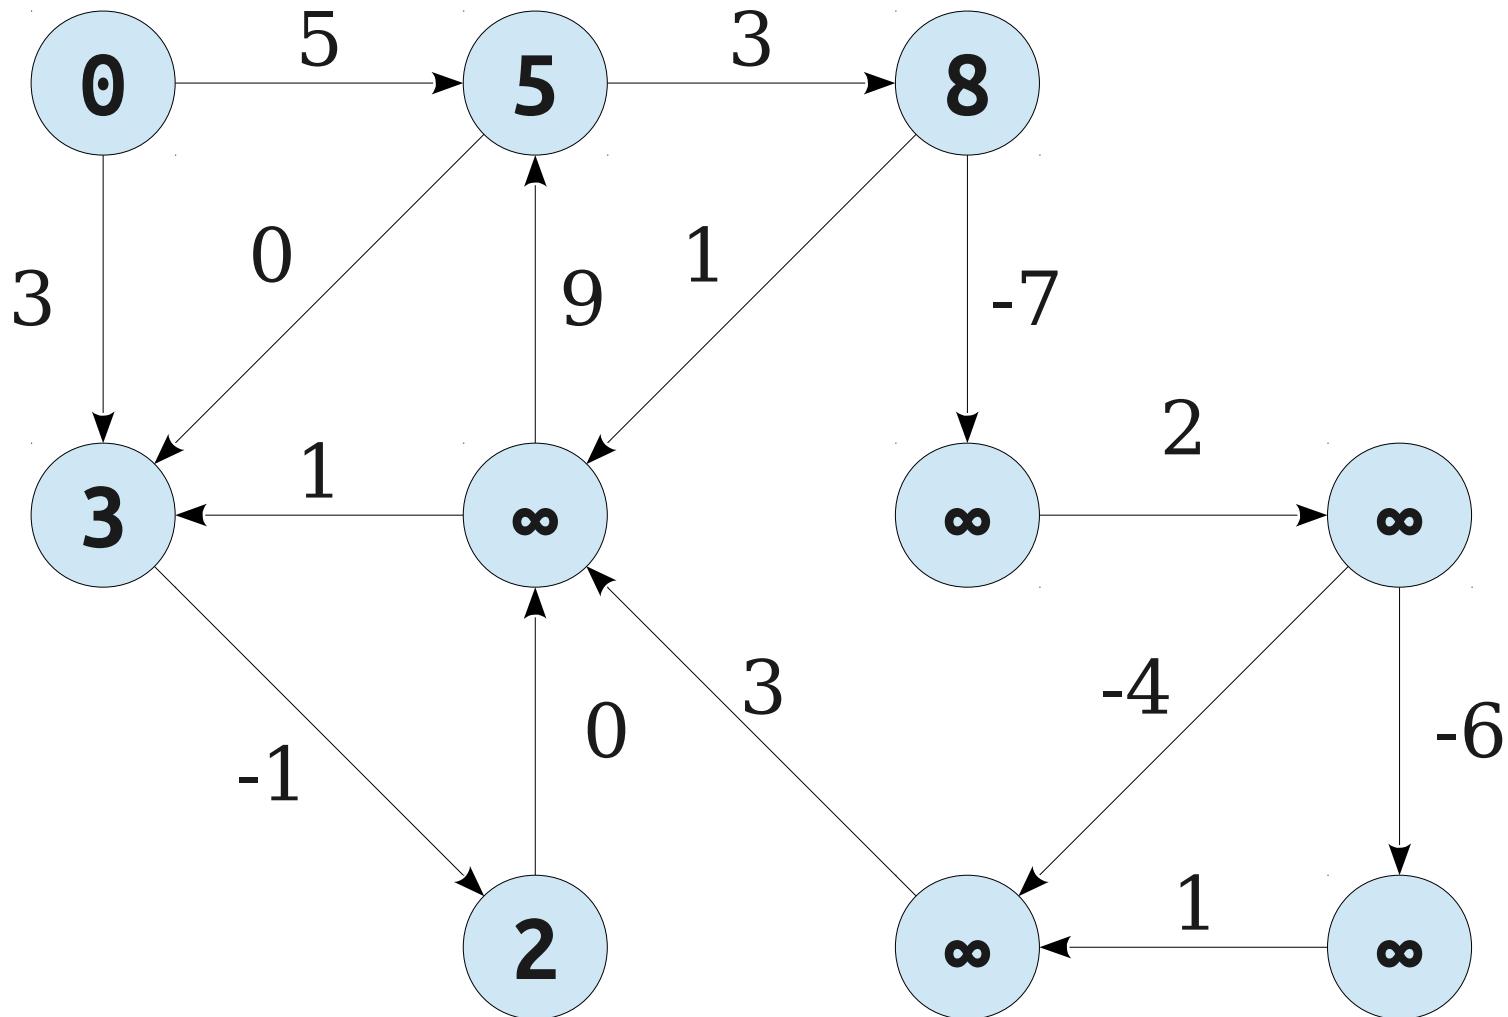
Consider paths of length **2**.



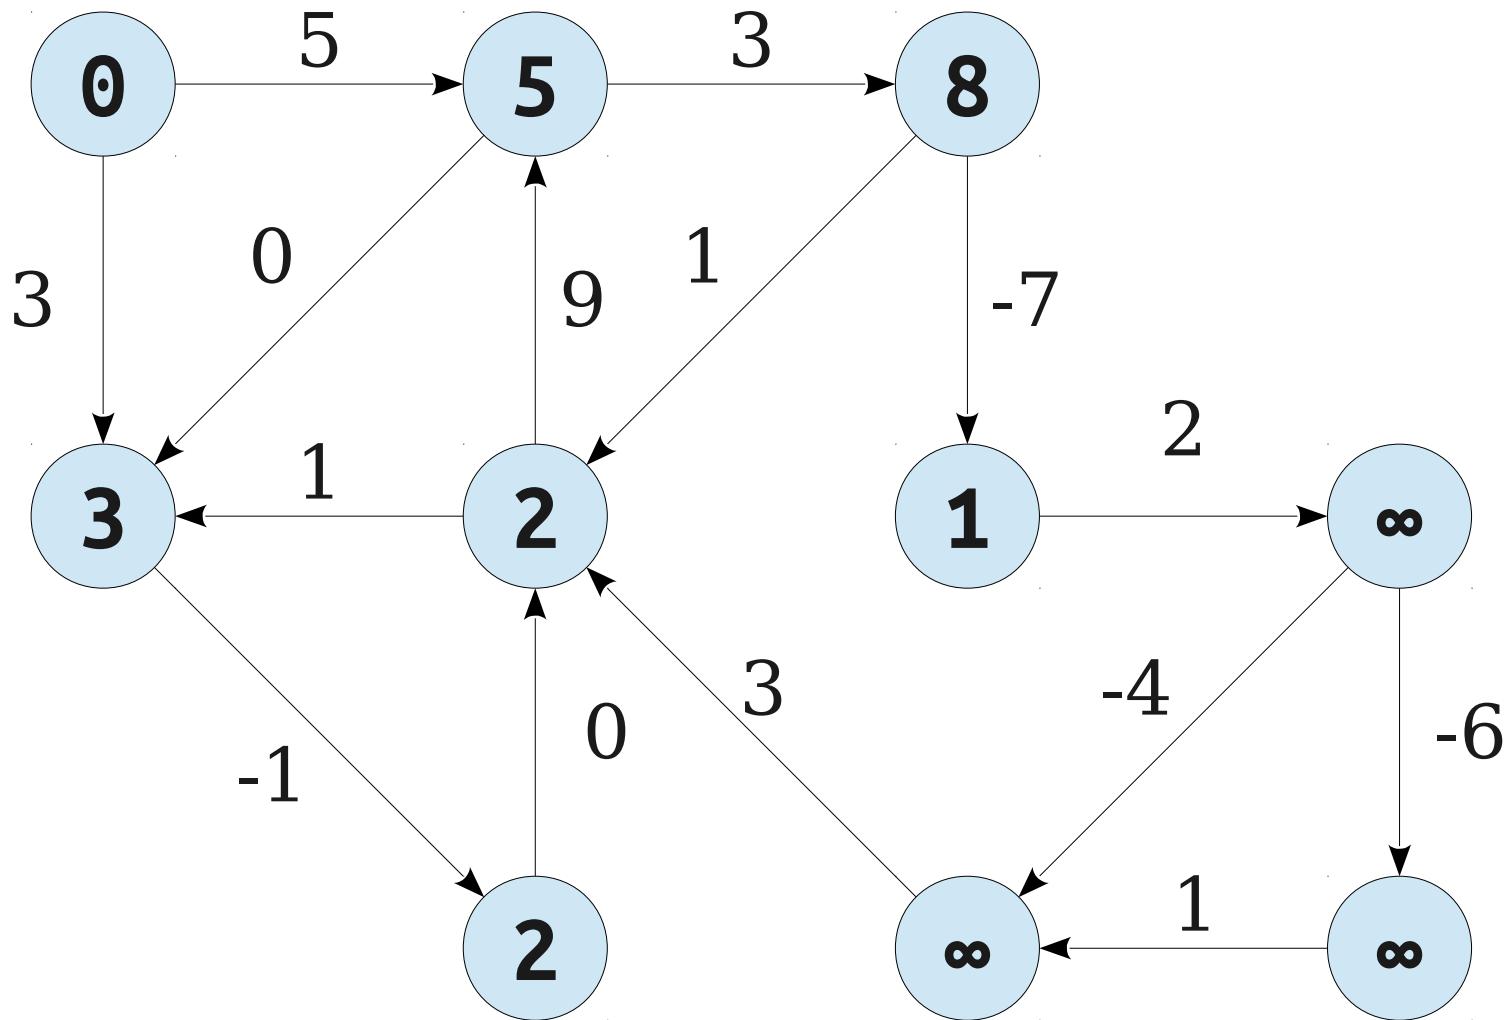
Consider paths of length **2**.



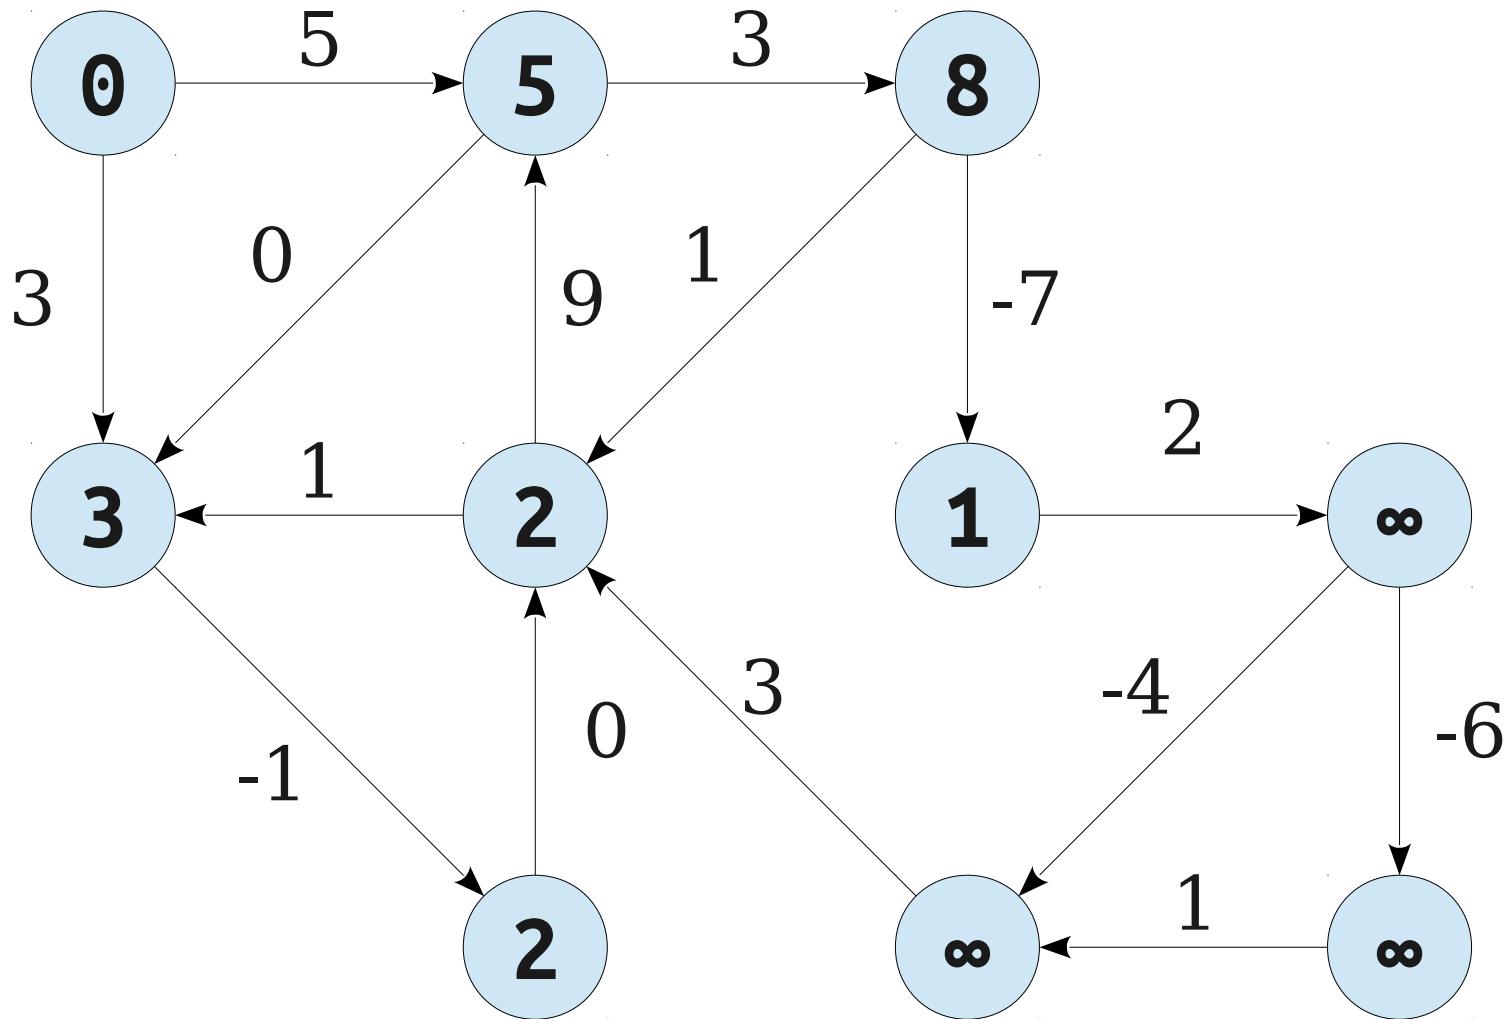
Consider paths of length **2**.



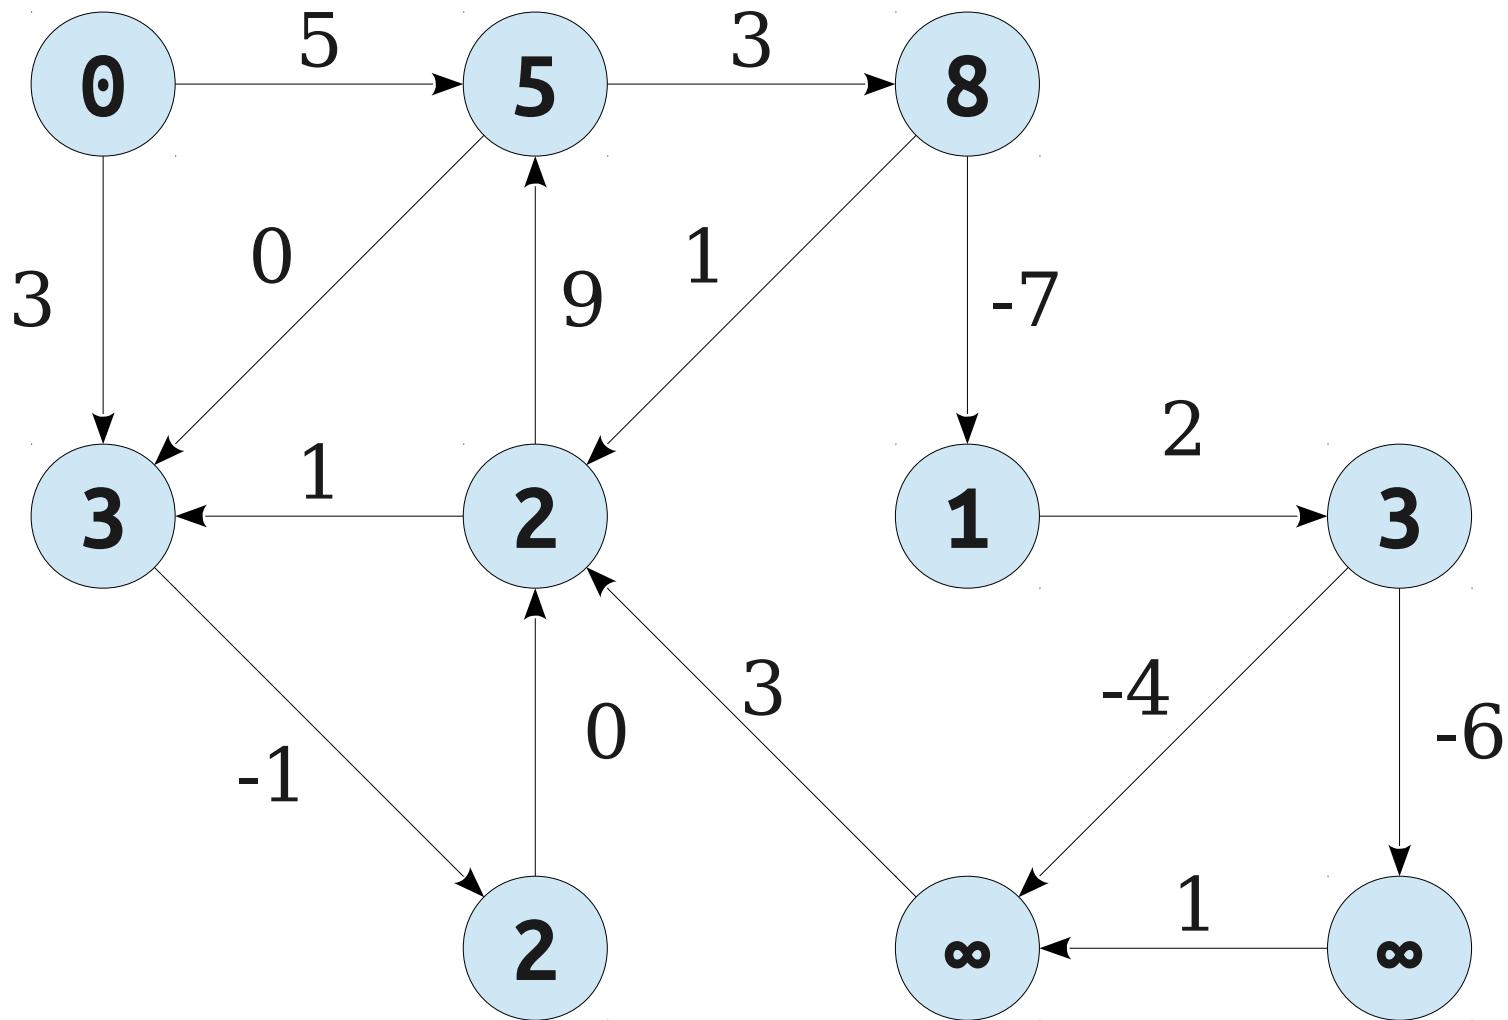
Consider paths of length **3**.



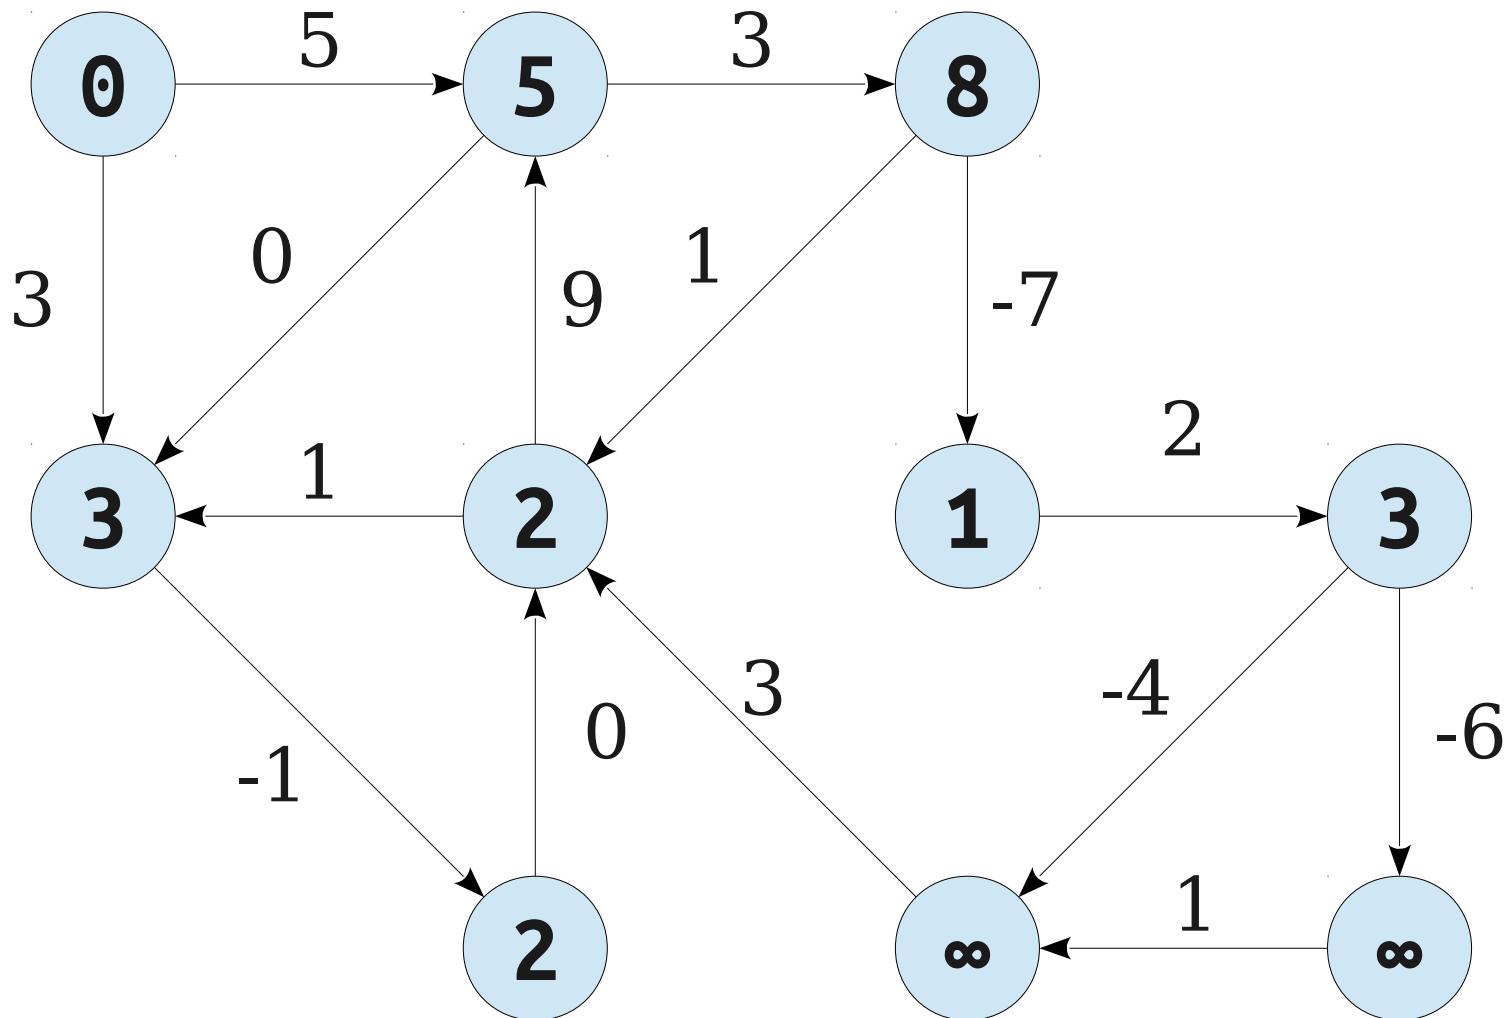
Consider paths of length **3**.



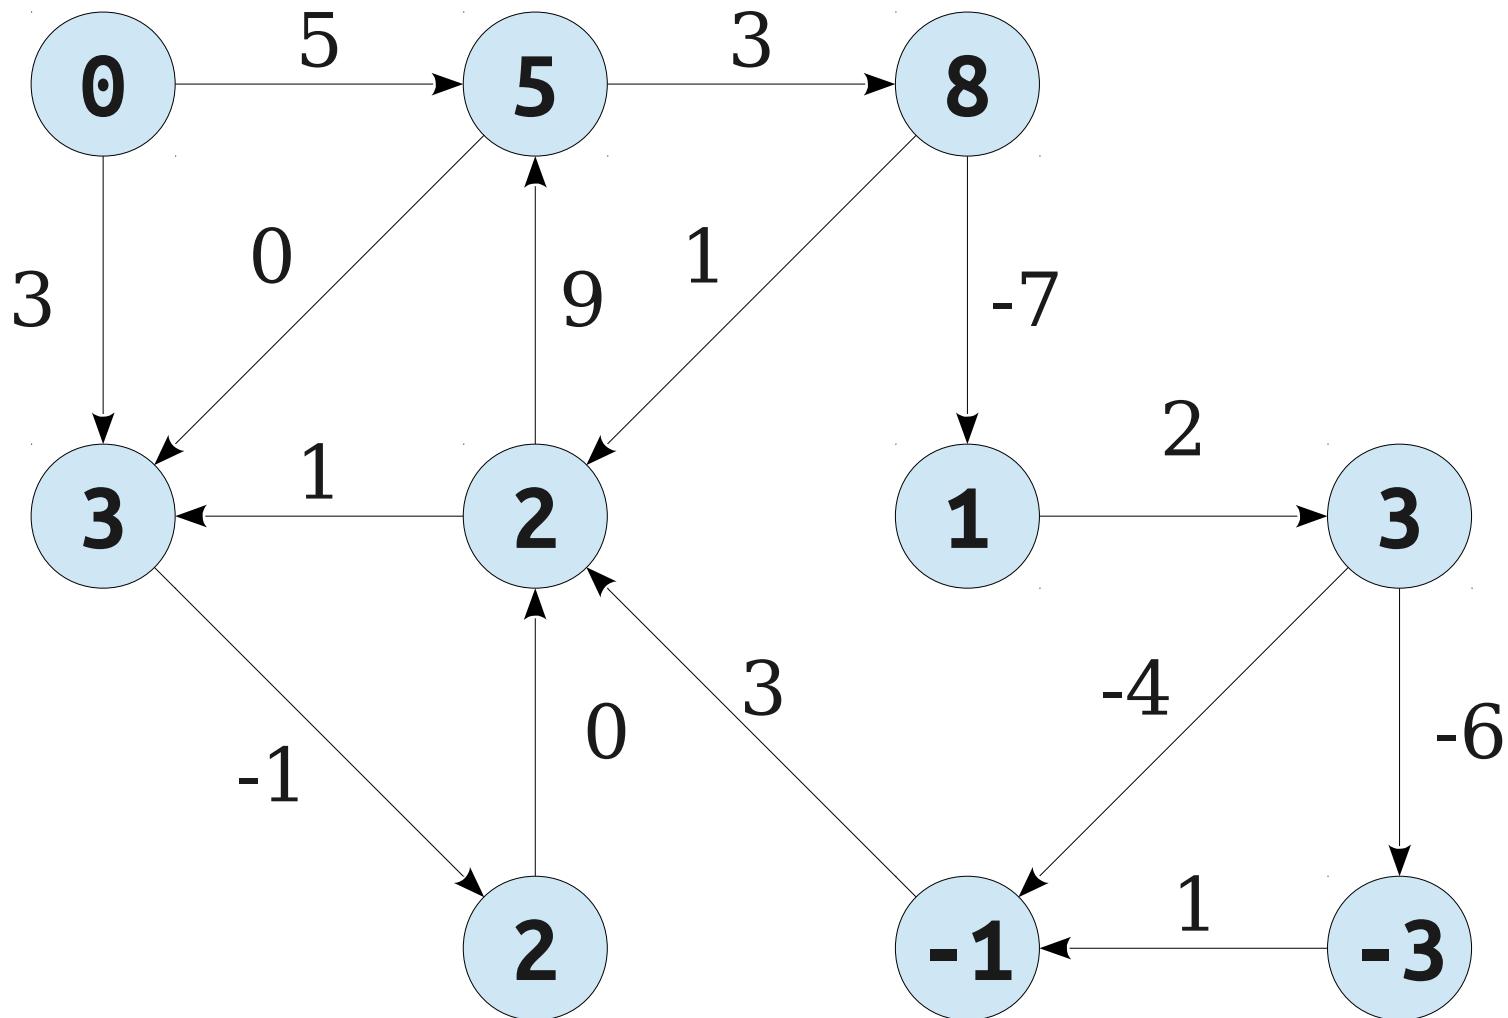
Consider paths of length **4**.



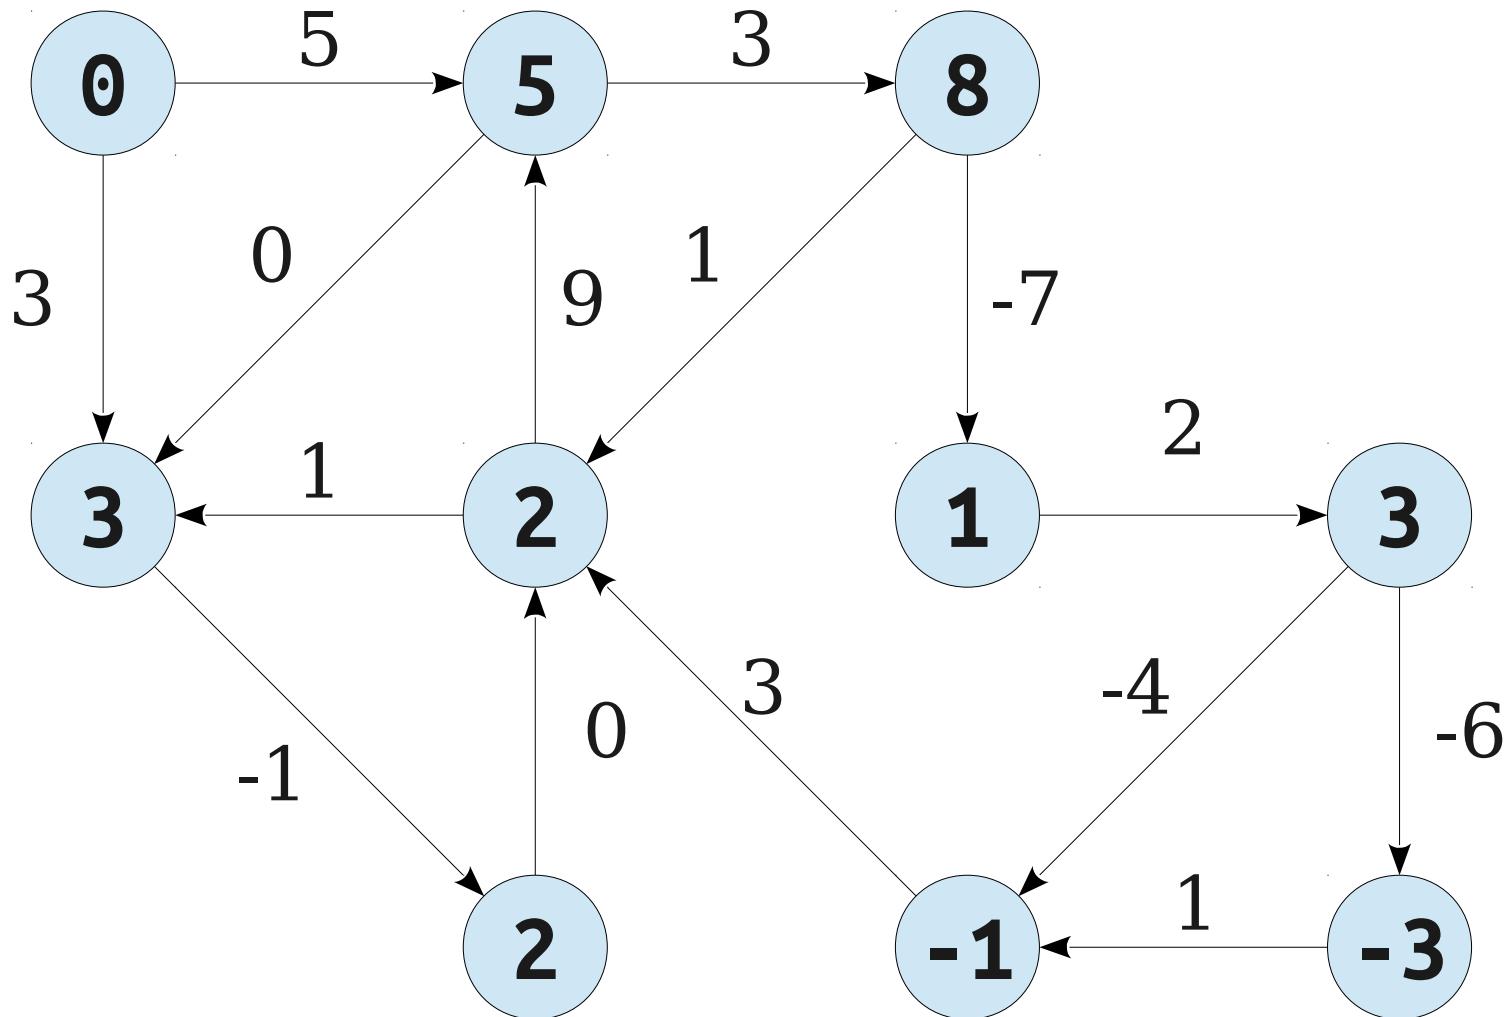
Consider paths of length **4**.



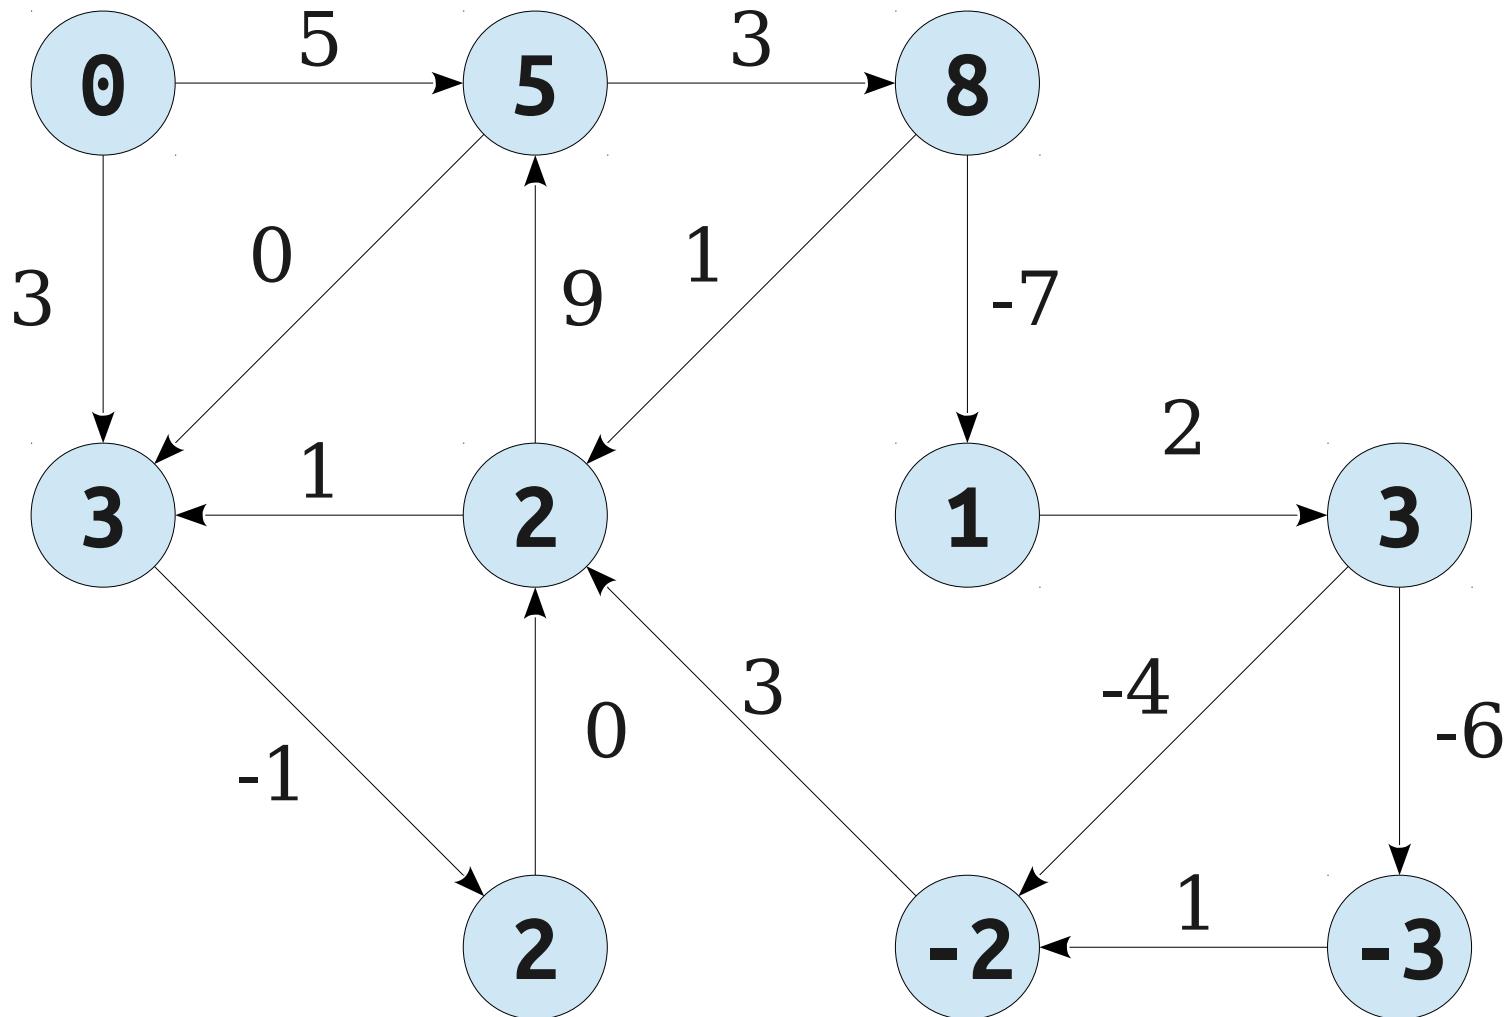
Consider paths of length **5**.



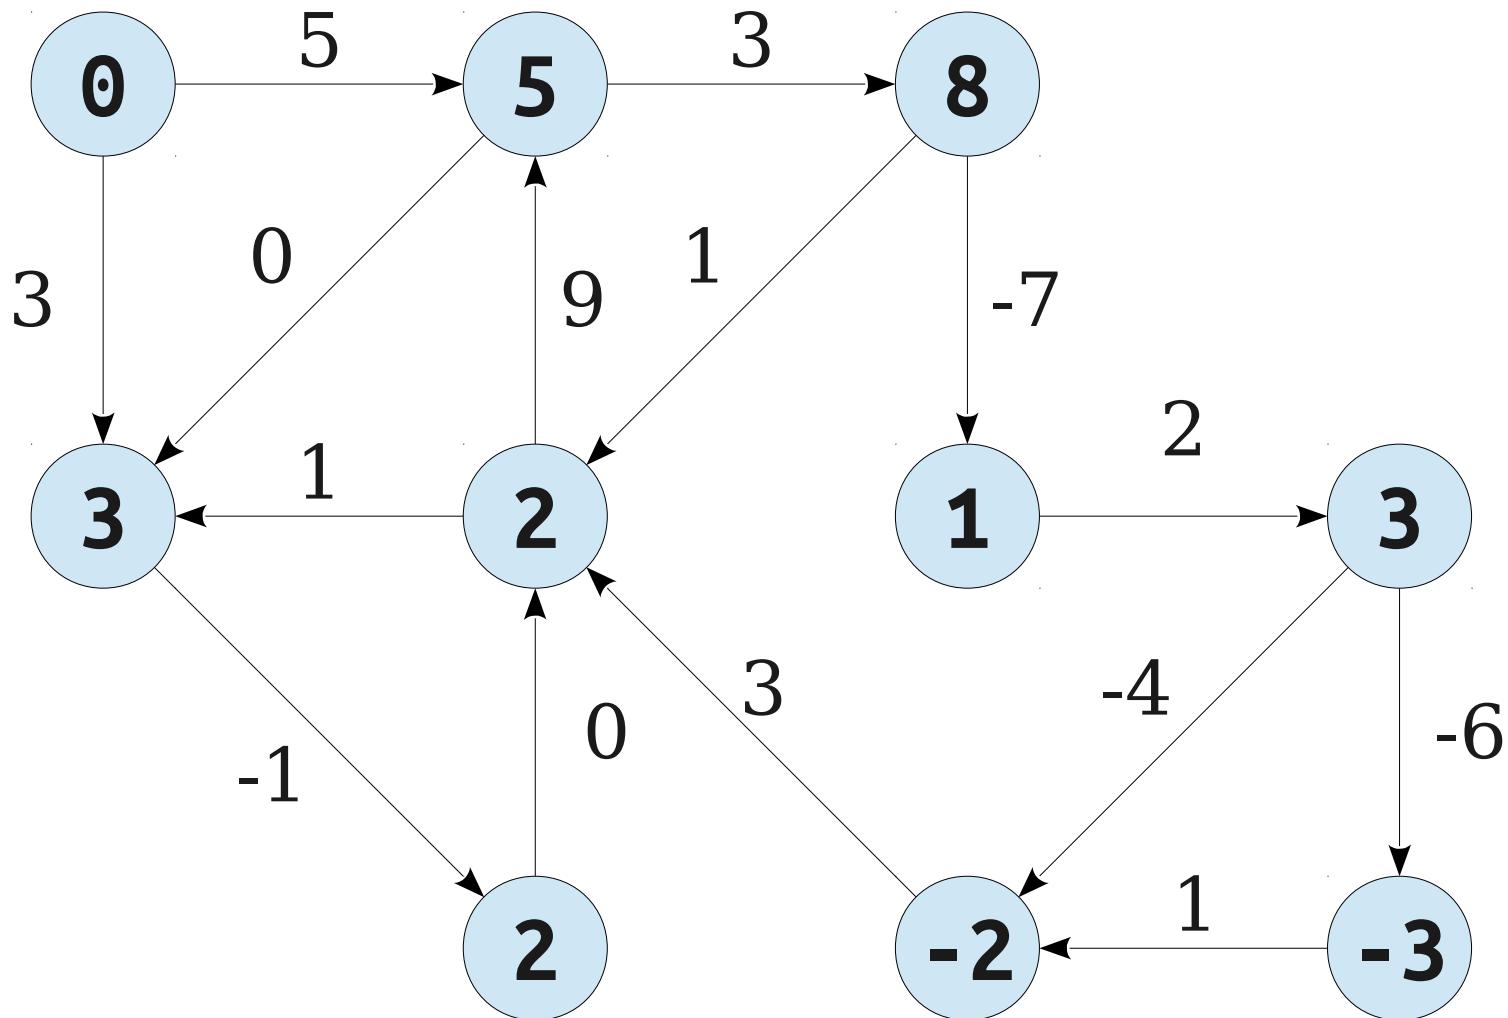
Consider paths of length **5**.



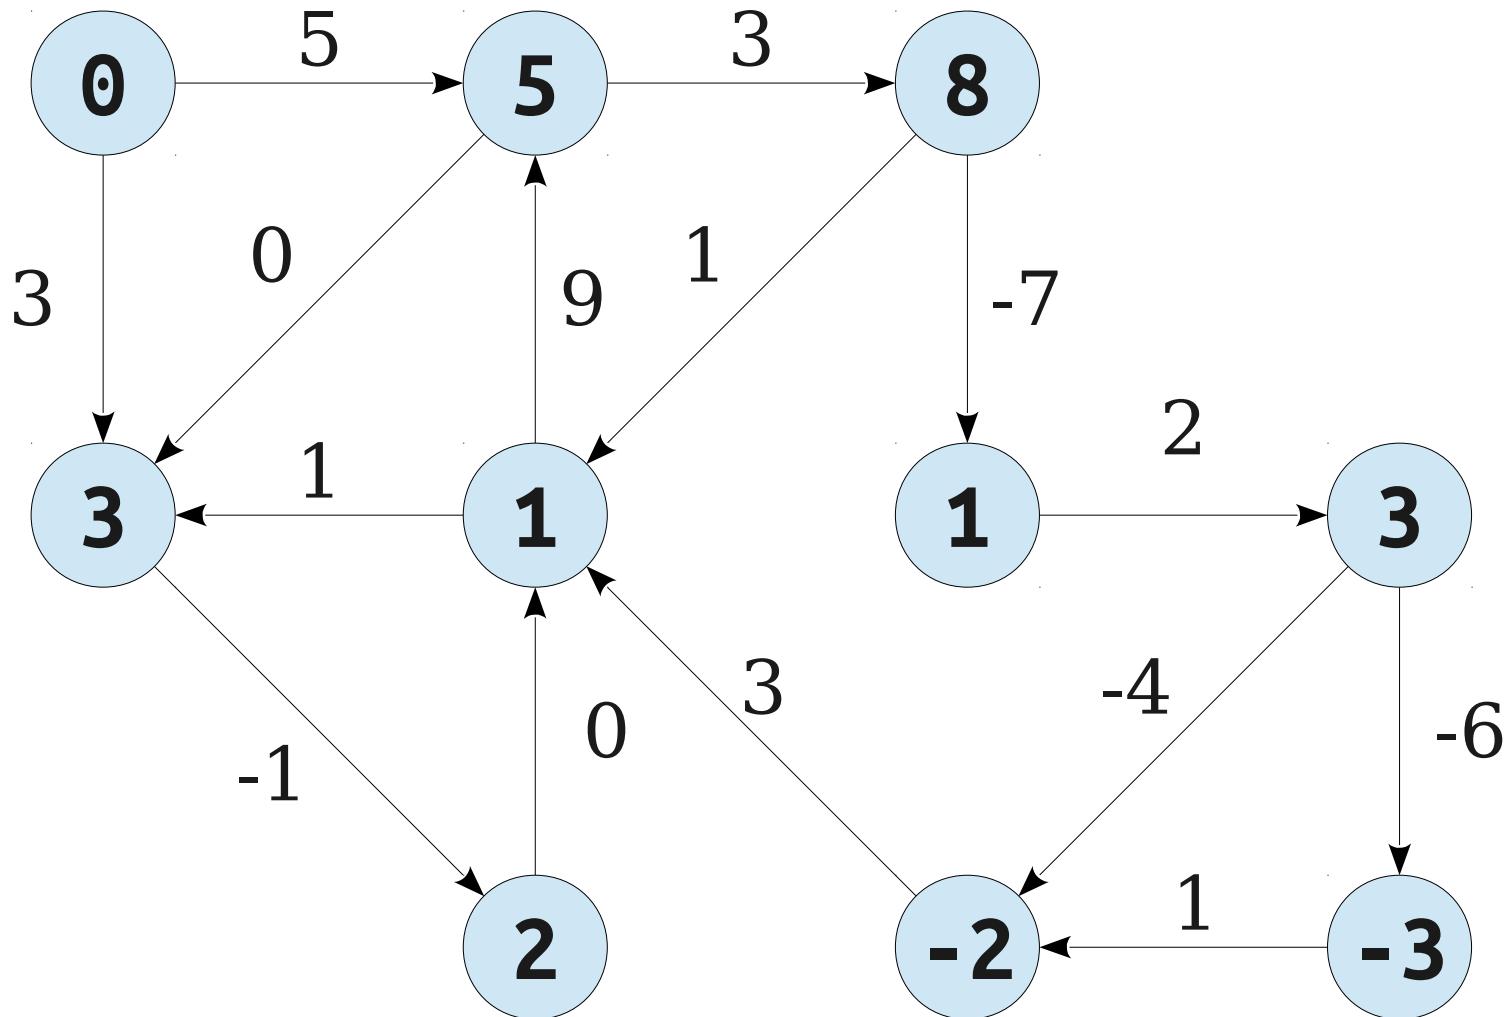
Consider paths of length **6**.



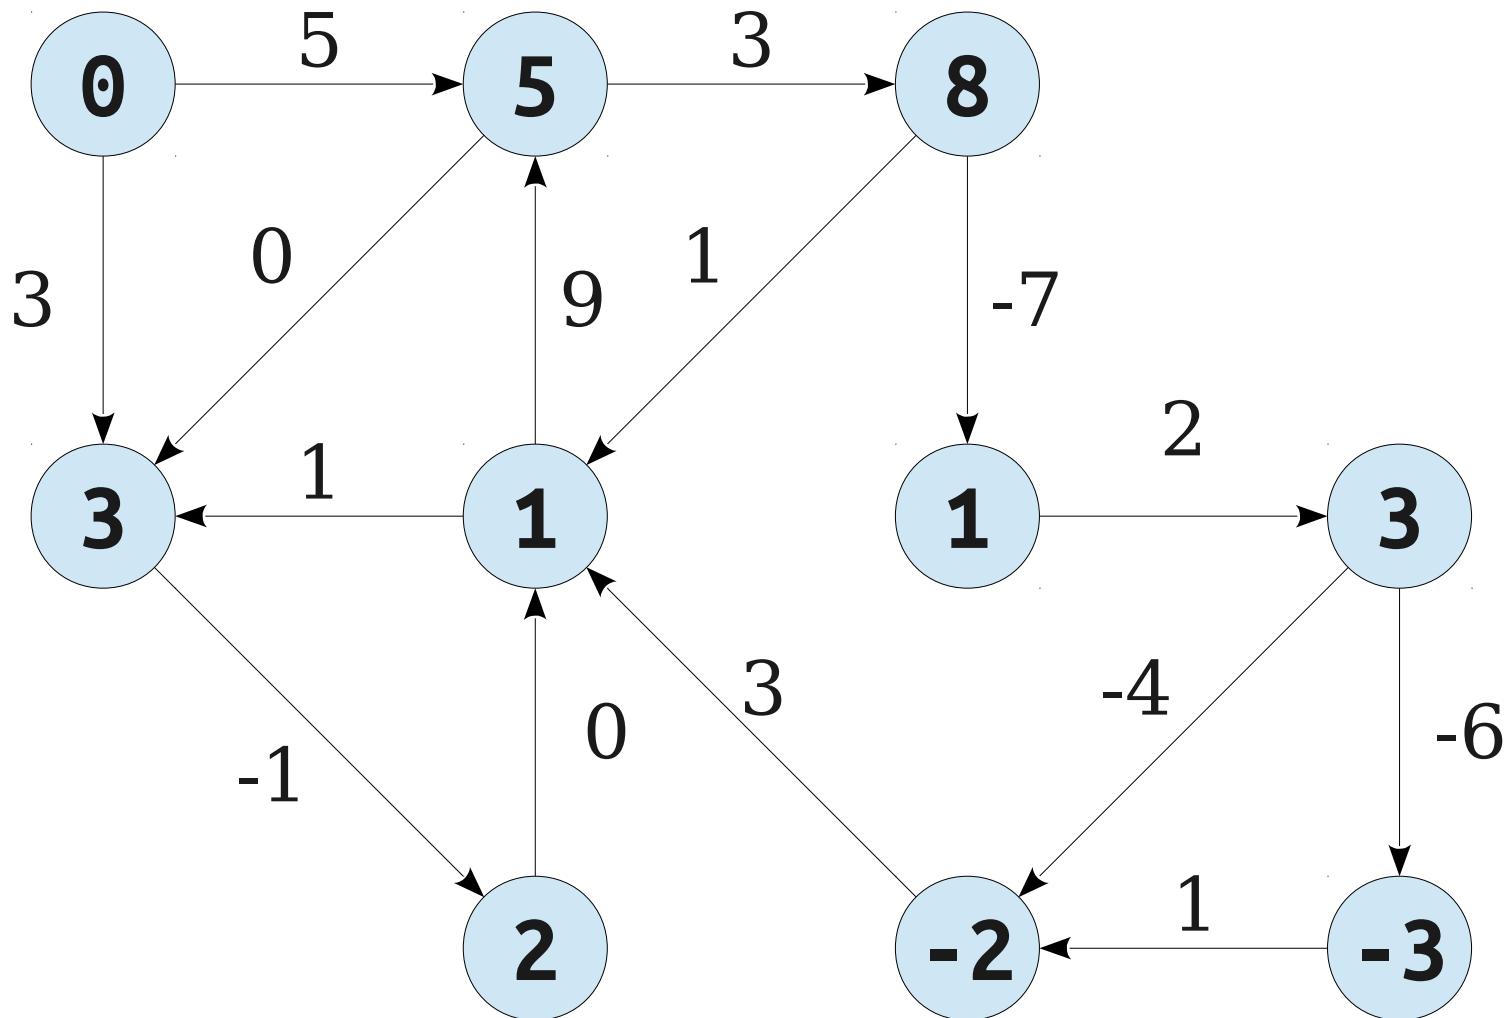
Consider paths of length **6**.



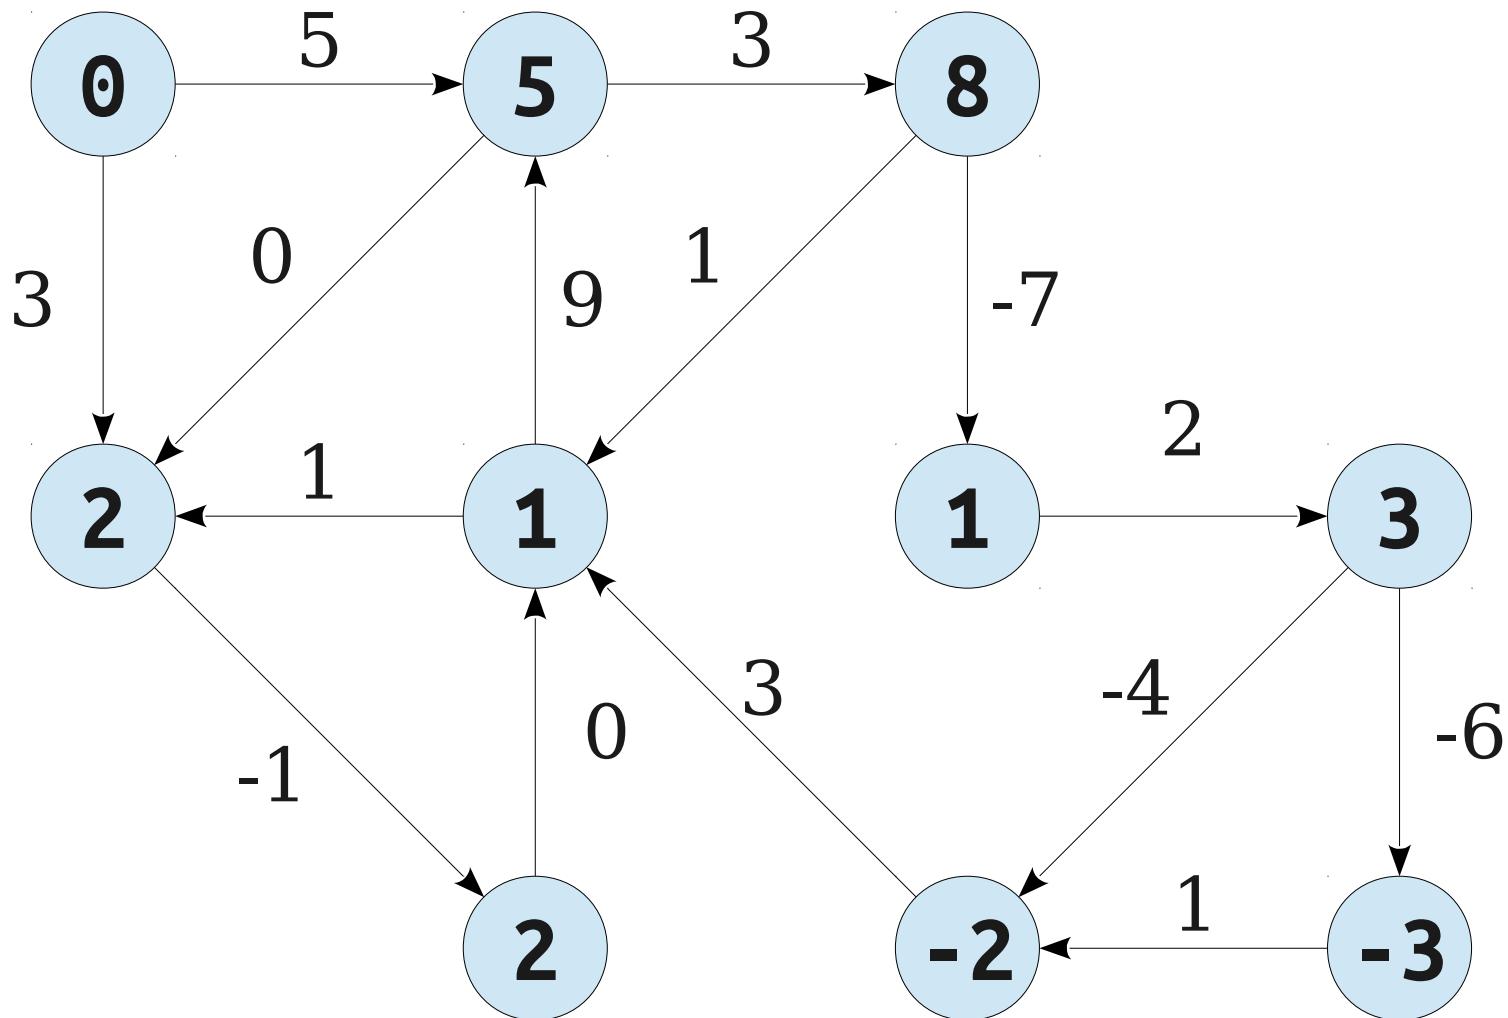
Consider paths of length 7.



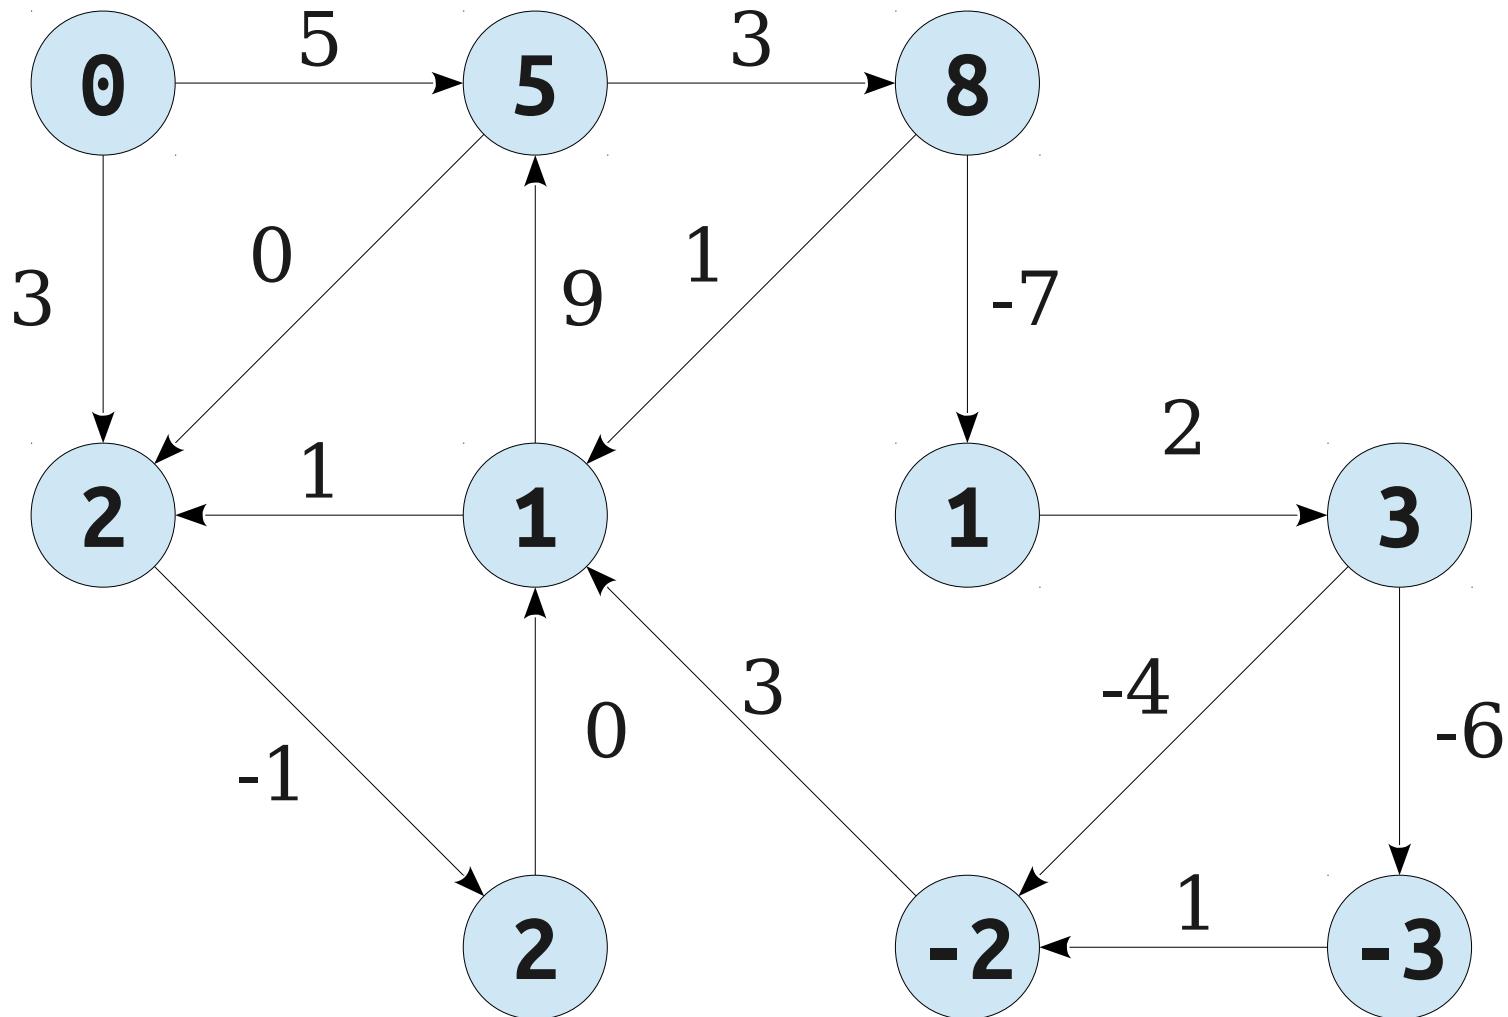
Consider paths of length 7.



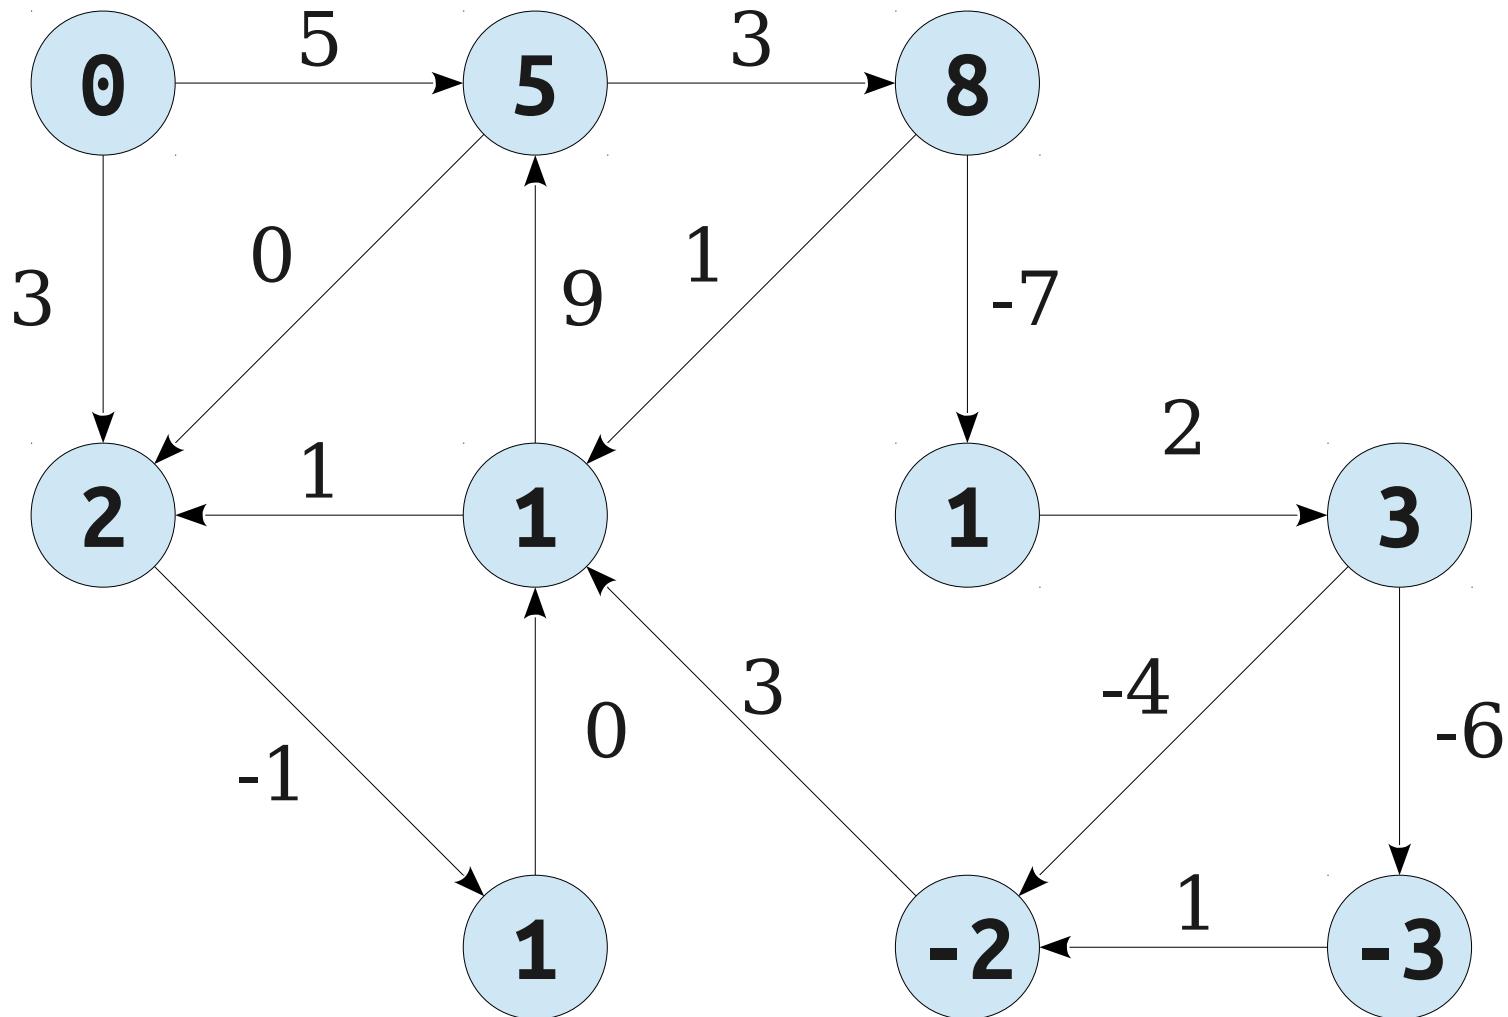
Consider paths of length **8**.



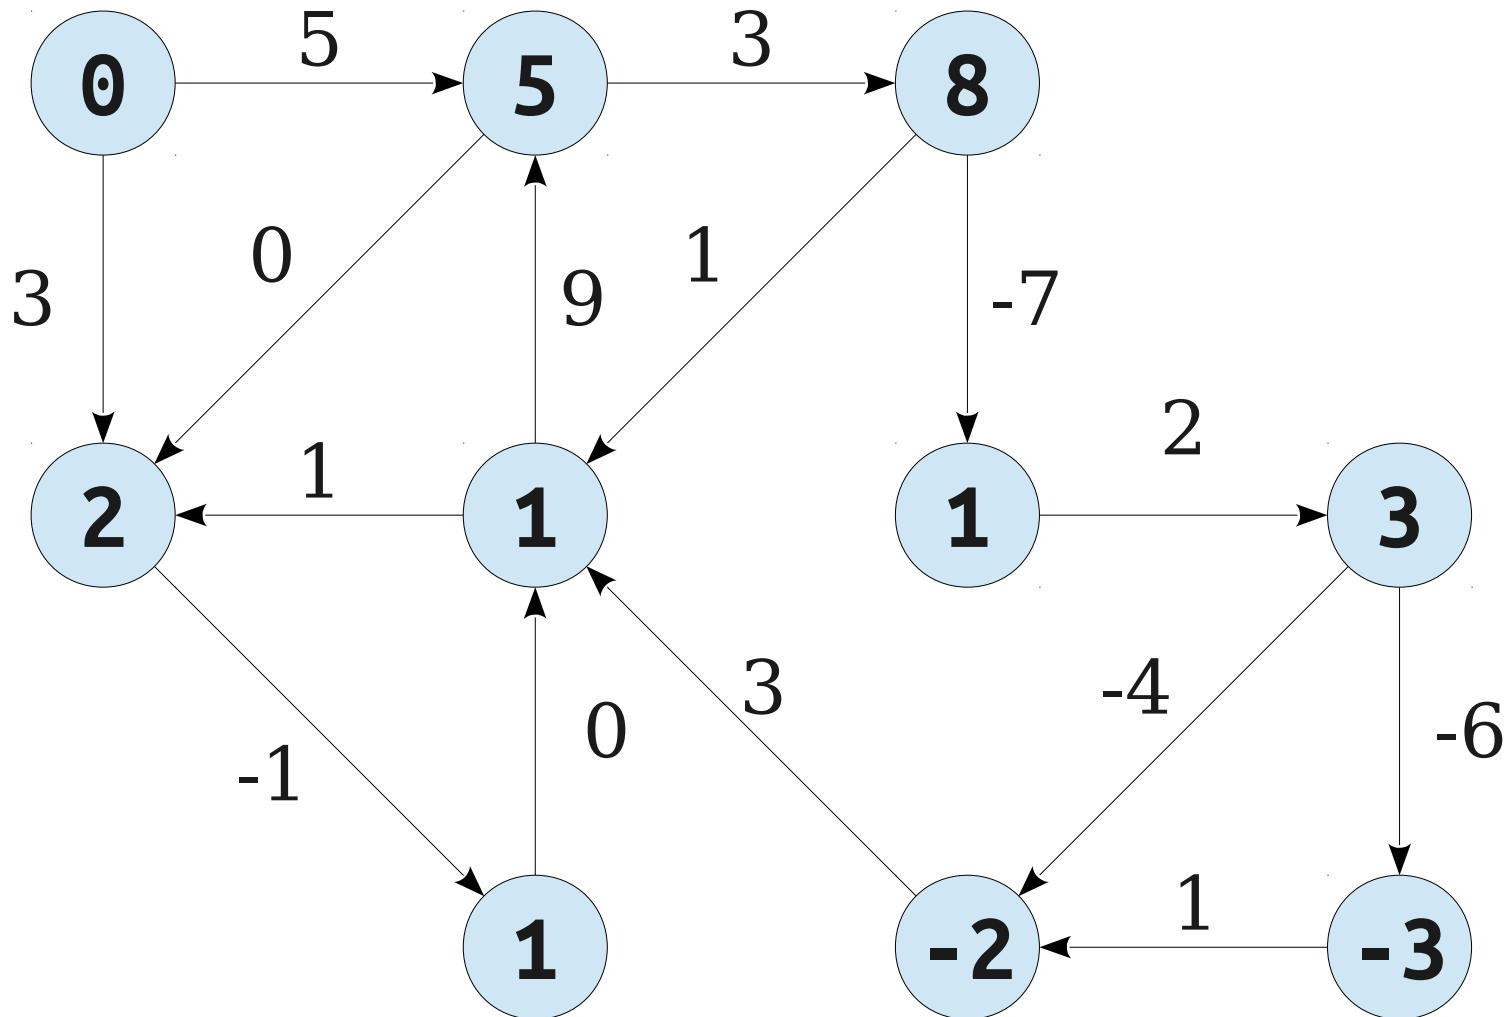
Consider paths of length **8**.



Consider paths of length **9**.



Consider paths of length **9**.



Consider paths of length **10+**.

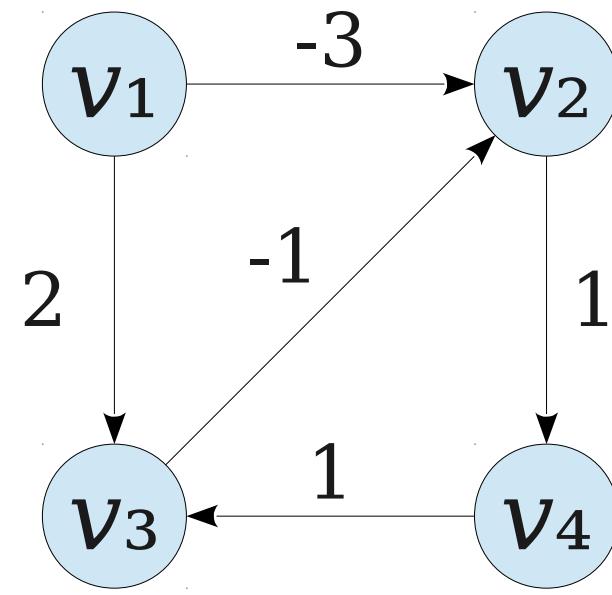
The Recurrence

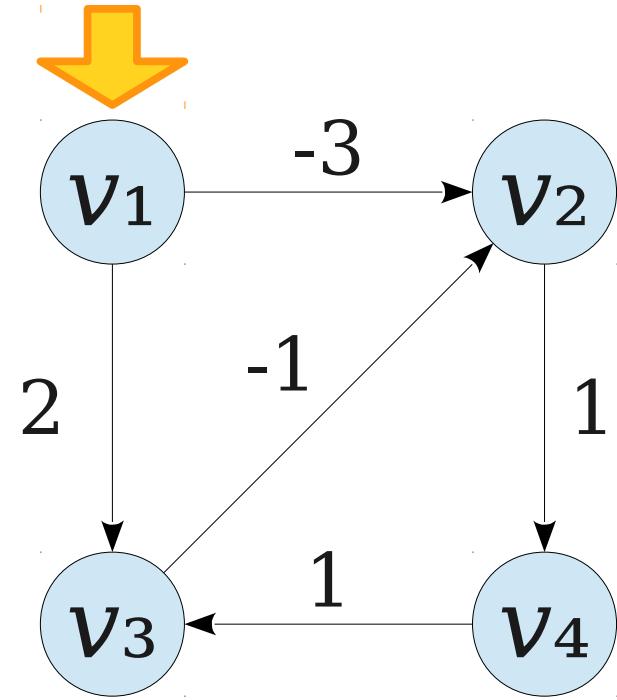
- **Idea:** Find paths of lengths at most $0, 1, 2, \dots, n$.
- Let $w(u, v)$ denote the weight of edge (u, v) .
- Let s be our start node. Let $\text{OPT}(v, i)$ be the length of the shortest $s - v$ path whose length is at most i , or ∞ if no path exists.
- **Claim:** $\text{OPT}(v, i)$ satisfies the following recurrence:

$$\text{OPT}(v, i) = \begin{cases} 0 & \text{if } i=0 \text{ and } v=s \\ \infty & \text{if } i=0 \text{ and } v \neq s \\ \min \left\{ \begin{array}{l} \text{OPT}(v, i-1), \\ \min_{(u,v) \in E} \{ \text{OPT}(u, i-1) + w(u, v) \} \end{array} \right\} & \text{otherwise} \end{cases}$$

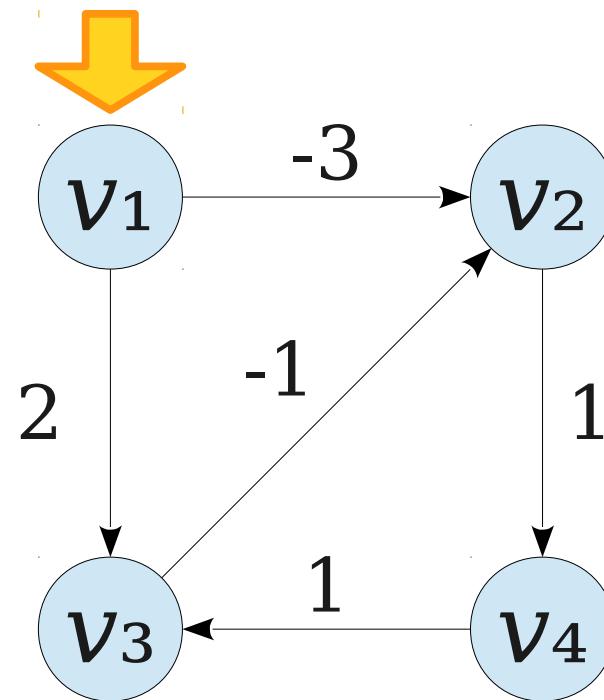
The Bellman-Ford Algorithm

- The **Bellman-Ford algorithm** evaluates this recurrence bottom-up:
 - Create a table DP of size $n \times n$.
 - Set $\text{DP}[v][0] = \infty$ for all $v \neq s$.
 - Set $\text{DP}[s][0] = 0$
 - For $i = 1$ to $n - 1$, for all $v \in V$:
 - Set $\text{DP}[v][i] = \min \{ \text{DP}[v][i - 1], \min \{ \text{DP}[u][i - 1] + w(u, v) \} \text{ (where } (u, v) \in E \}$
 - Return row n of DP.

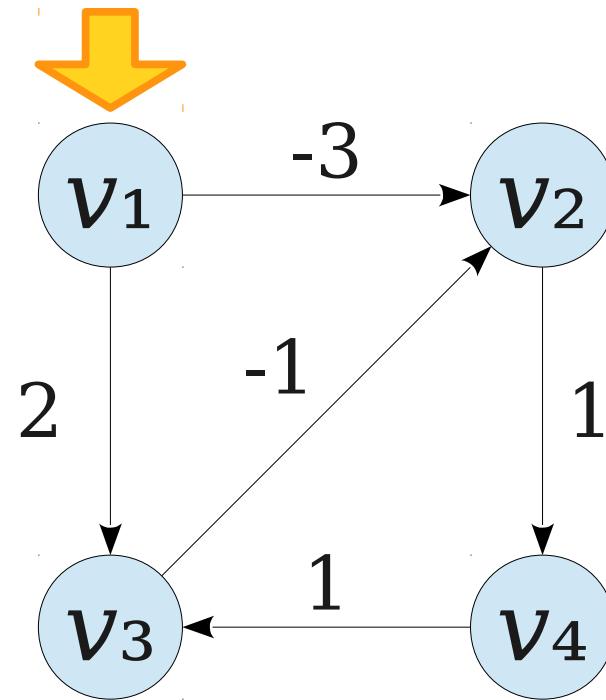




	v_1	v_2	v_3	v_4
3				
2				
1				
0				

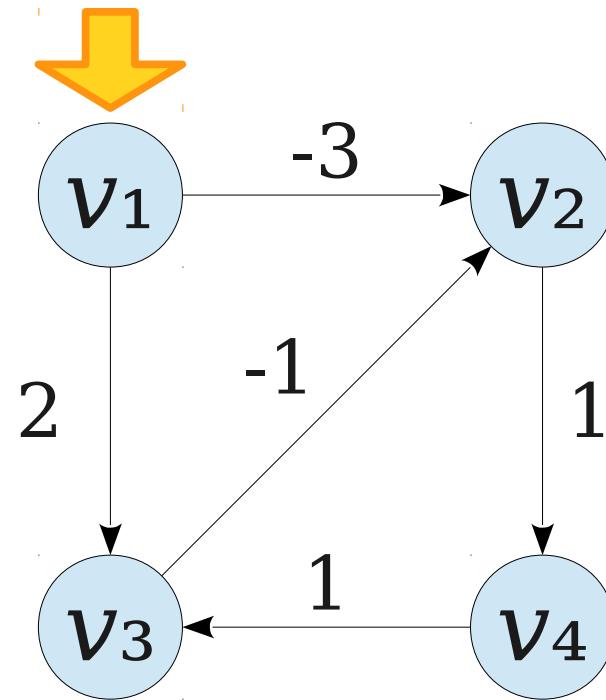


	v_1	v_2	v_3	v_4
3				
2				
1				
0				



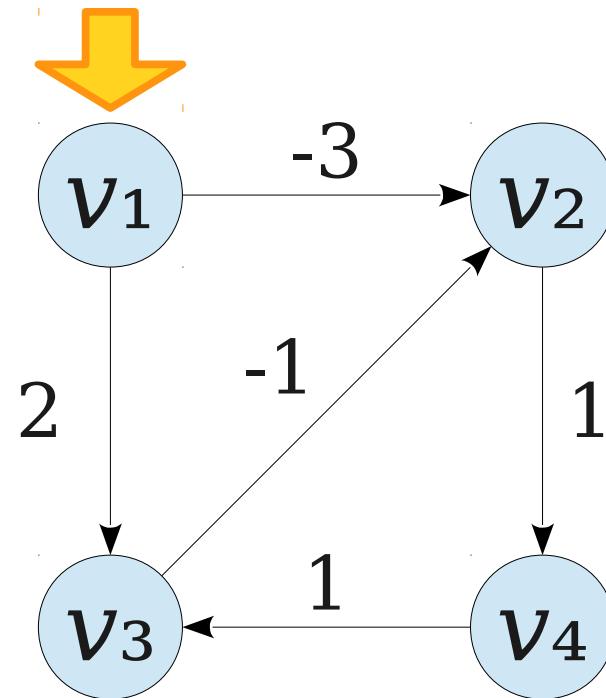
$$\text{OPT}(v, i) = \begin{cases} 0 & \text{if } i=0 \text{ and } v=s \\ \infty & \text{if } i=0 \text{ and } v \neq s \\ \min \left\{ \begin{array}{l} \text{OPT}(v, i-1), \\ \min_{(u,v) \in E} \{ \text{OPT}(u, i-1) + w(u, v) \} \end{array} \right\} & \text{otherwise} \end{cases}$$

	v_1	v_2	v_3	v_4
3				
2				
1				
0	0	∞	∞	∞



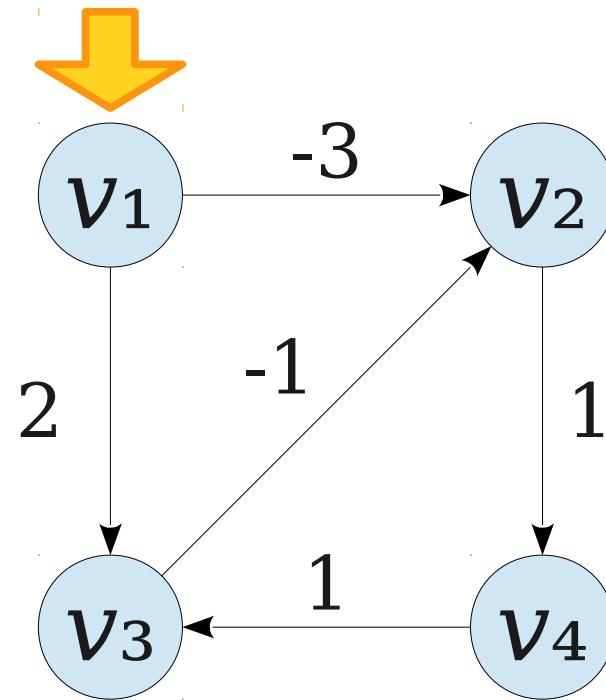
$$\text{OPT}(v, i) = \begin{cases} 0 & \text{if } i=0 \text{ and } v=s \\ \infty & \text{if } i=0 \text{ and } v \neq s \\ \min \left\{ \begin{array}{l} \text{OPT}(v, i-1), \\ \min_{(u,v) \in E} \{ \text{OPT}(u, i-1) + w(u, v) \} \end{array} \right\} & \text{otherwise} \end{cases}$$

	v_1	v_2	v_3	v_4
3				
2				
1				
0	0	∞	∞	∞



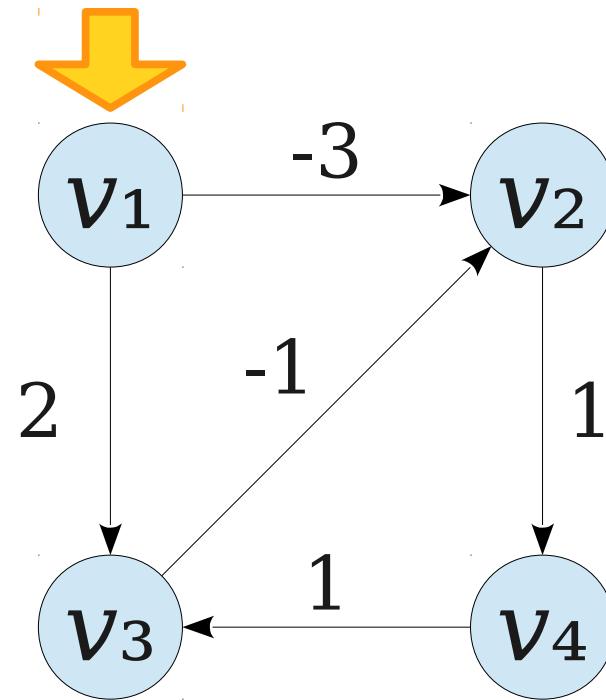
$$\text{OPT}(v, i) = \begin{cases} 0 & \text{if } i=0 \text{ and } v=s \\ \infty & \text{if } i=0 \text{ and } v \neq s \\ \min \left\{ \begin{array}{l} \text{OPT}(v, i-1), \\ \min_{(u,v) \in E} \{ \text{OPT}(u, i-1) + w(u, v) \} \end{array} \right\} & \text{otherwise} \end{cases}$$

	v_1	v_2	v_3	v_4
3				
2				
1	0			
0	0	∞	∞	∞



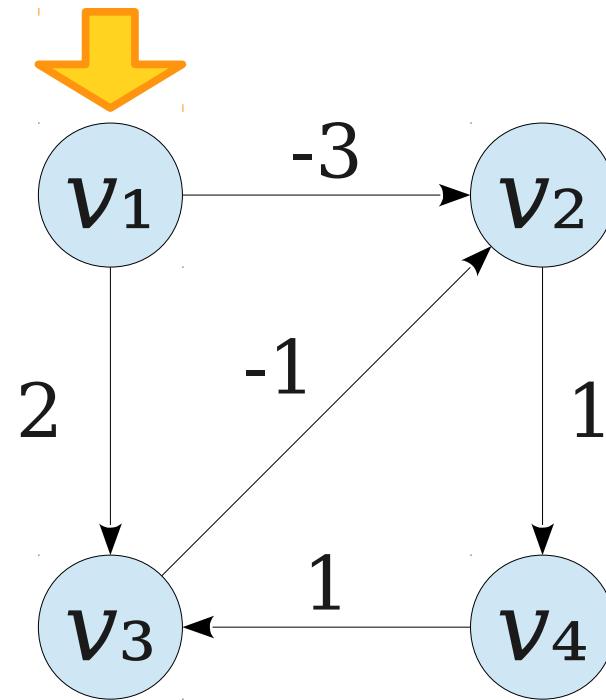
$$\text{OPT}(v, i) = \begin{cases} 0 & \text{if } i=0 \text{ and } v=s \\ \infty & \text{if } i=0 \text{ and } v \neq s \\ \min \left\{ \begin{array}{l} \text{OPT}(v, i-1), \\ \min_{(u,v) \in E} \{ \text{OPT}(u, i-1) + w(u, v) \} \end{array} \right\} & \text{otherwise} \end{cases}$$

	v_1	v_2	v_3	v_4
3				
2				
1	0	-3		
0	0	∞	∞	∞



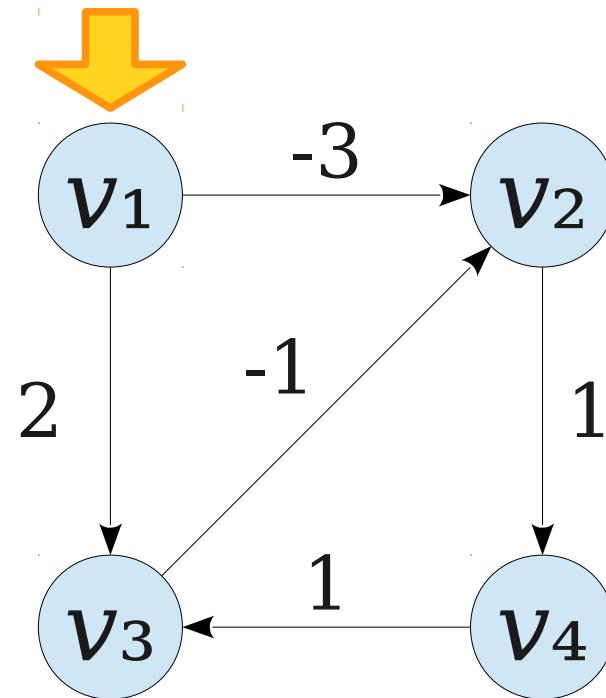
$$\text{OPT}(v, i) = \begin{cases} 0 & \text{if } i=0 \text{ and } v=s \\ \infty & \text{if } i=0 \text{ and } v \neq s \\ \min \left\{ \begin{array}{l} \text{OPT}(v, i-1), \\ \min_{(u,v) \in E} \{ \text{OPT}(u, i-1) + w(u, v) \} \end{array} \right\} & \text{otherwise} \end{cases}$$

	v_1	v_2	v_3	v_4
3				
2				
1	0	-3	2	
0	0	∞	∞	∞



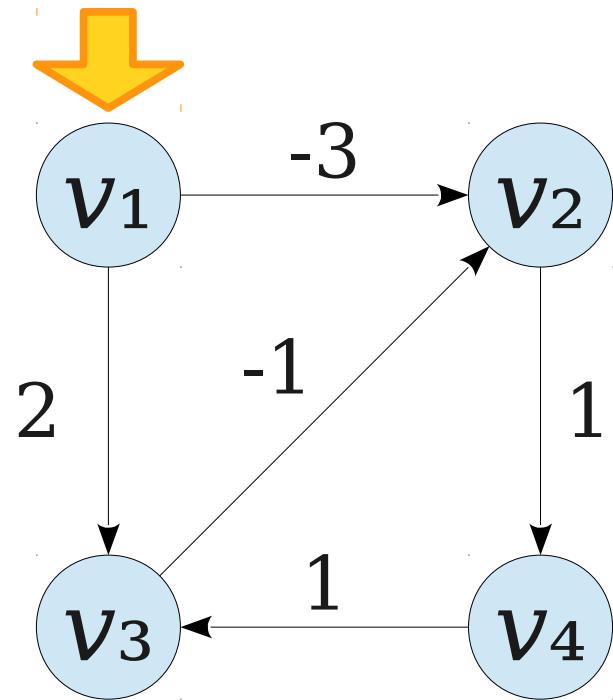
$$\text{OPT}(v, i) = \begin{cases} 0 & \text{if } i=0 \text{ and } v=s \\ \infty & \text{if } i=0 \text{ and } v \neq s \\ \min \left\{ \begin{array}{l} \text{OPT}(v, i-1), \\ \min_{(u,v) \in E} \{ \text{OPT}(u, i-1) + w(u, v) \} \end{array} \right\} & \text{otherwise} \end{cases}$$

	v_1	v_2	v_3	v_4	
3					
2					
1	0	-3	2	∞	
0	0	∞	∞	∞	



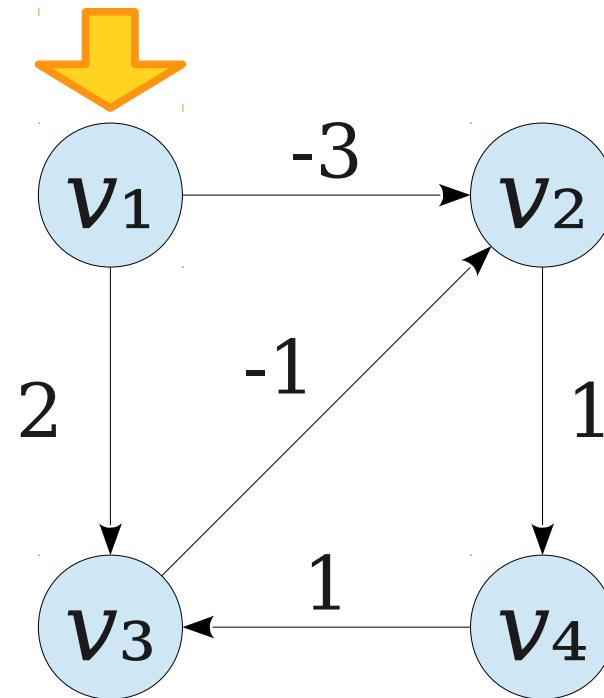
$$\text{OPT}(v, i) = \begin{cases} 0 & \text{if } i=0 \text{ and } v=s \\ \infty & \text{if } i=0 \text{ and } v \neq s \\ \min \left\{ \begin{array}{l} \text{OPT}(v, i-1), \\ \min_{(u,v) \in E} \{ \text{OPT}(u, i-1) + w(u, v) \} \end{array} \right\} & \text{otherwise} \end{cases}$$

	v_1	v_2	v_3	v_4	
3					
2					
1	0	-3	2	∞	
0	0	∞	∞	∞	



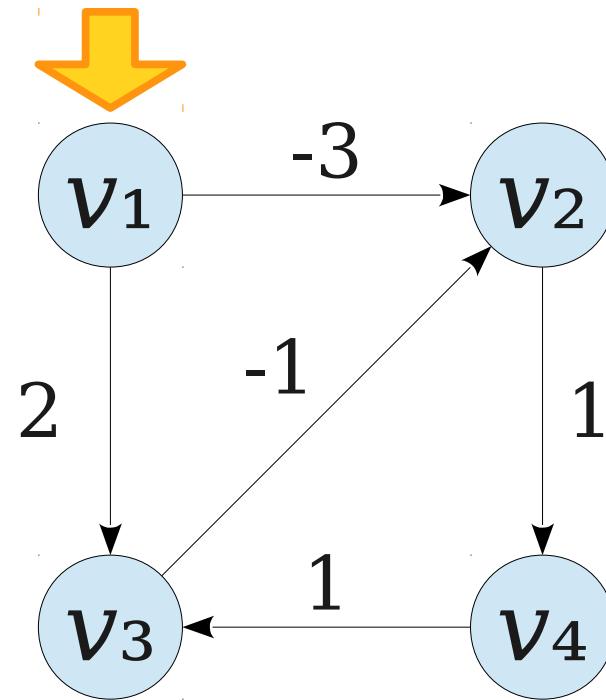
$$\text{OPT}(v, i) = \begin{cases} 0 & \text{if } i=0 \text{ and } v=s \\ \infty & \text{if } i=0 \text{ and } v \neq s \\ \min \left\{ \begin{array}{l} \text{OPT}(v, i-1), \\ \min_{(u,v) \in E} \{ \text{OPT}(u, i-1) + w(u, v) \} \end{array} \right\} & \text{otherwise} \end{cases}$$

	v_1	v_2	v_3	v_4	
3					
2	0				
1	0	-3	2	∞	
0	0	∞	∞	∞	



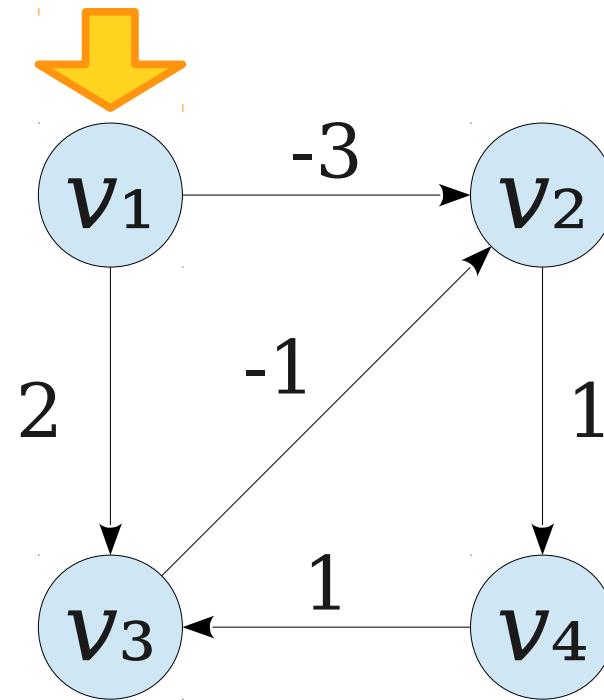
$$\text{OPT}(v, i) = \begin{cases} 0 & \text{if } i=0 \text{ and } v=s \\ \infty & \text{if } i=0 \text{ and } v \neq s \\ \min \left\{ \begin{array}{l} \text{OPT}(v, i-1), \\ \min_{(u,v) \in E} \{ \text{OPT}(u, i-1) + w(u, v) \} \end{array} \right\} & \text{otherwise} \end{cases}$$

	v_1	v_2	v_3	v_4
v_1	3			
v_2	2	0	-3	
v_3	1	0	-3	2
v_4	0	0	∞	∞



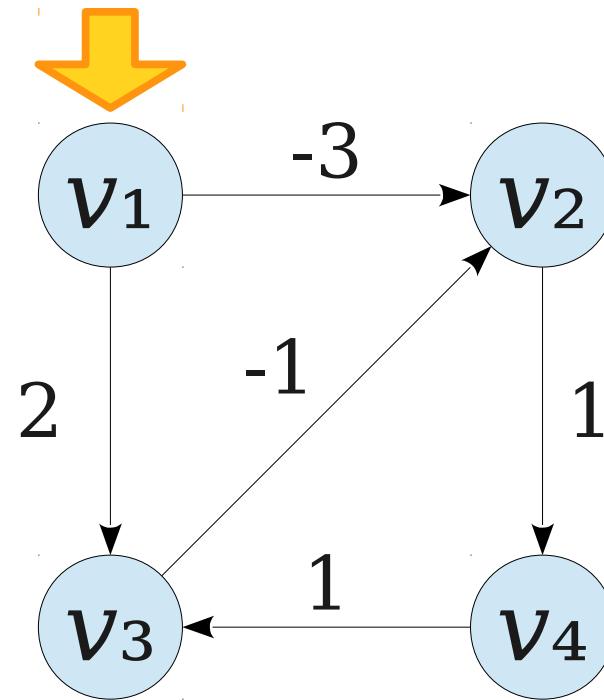
$$\text{OPT}(v, i) = \begin{cases} 0 & \text{if } i=0 \text{ and } v=s \\ \infty & \text{if } i=0 \text{ and } v \neq s \\ \min \left\{ \begin{array}{l} \text{OPT}(v, i-1), \\ \min_{(u,v) \in E} \{ \text{OPT}(u, i-1) + w(u, v) \} \end{array} \right\} & \text{otherwise} \end{cases}$$

	v_1	v_2	v_3	v_4
v_1	3			
v_2	2	0	-3	2
v_3	1	0	-3	2
v_4	0	0	∞	∞



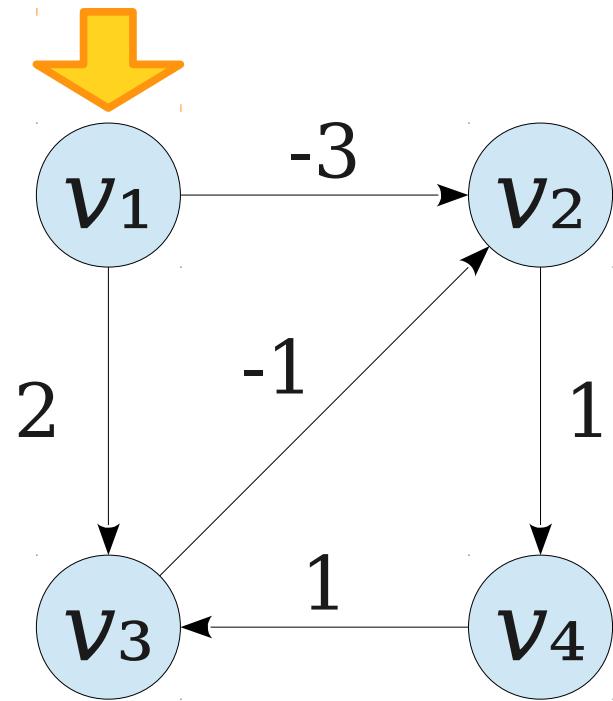
$$\text{OPT}(v, i) = \begin{cases} 0 & \text{if } i=0 \text{ and } v=s \\ \infty & \text{if } i=0 \text{ and } v \neq s \\ \min \left\{ \begin{array}{l} \text{OPT}(v, i-1), \\ \min_{(u,v) \in E} \{ \text{OPT}(u, i-1) + w(u, v) \} \end{array} \right\} & \text{otherwise} \end{cases}$$

	v_1	v_2	v_3	v_4	
3					
2	0	-3	2	-2	
1	0	-3	2	∞	
0	0	∞	∞	∞	



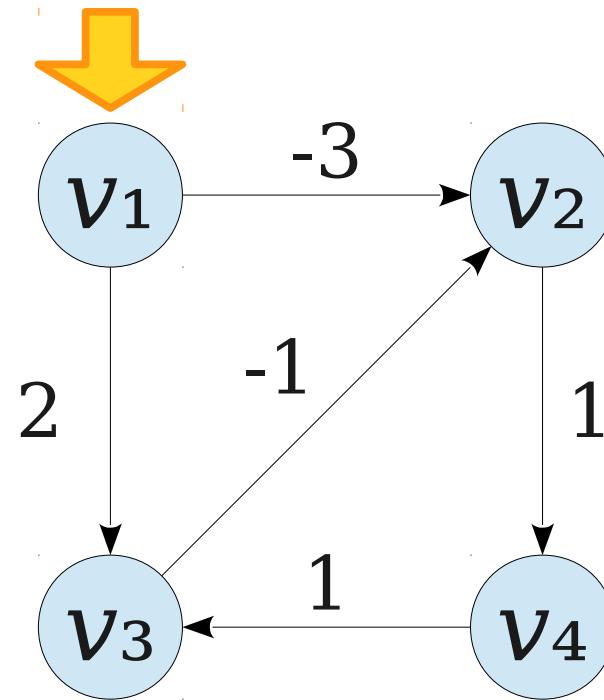
$$\text{OPT}(v, i) = \begin{cases} 0 & \text{if } i=0 \text{ and } v=s \\ \infty & \text{if } i=0 \text{ and } v \neq s \\ \min \left\{ \begin{array}{l} \text{OPT}(v, i-1), \\ \min_{(u,v) \in E} \{ \text{OPT}(u, i-1) + w(u, v) \} \end{array} \right\} & \text{otherwise} \end{cases}$$

	v_1	v_2	v_3	v_4
v_1	3			
v_2	0	-3	2	-2
v_3	0	-3	2	∞
v_4	0	∞	∞	∞



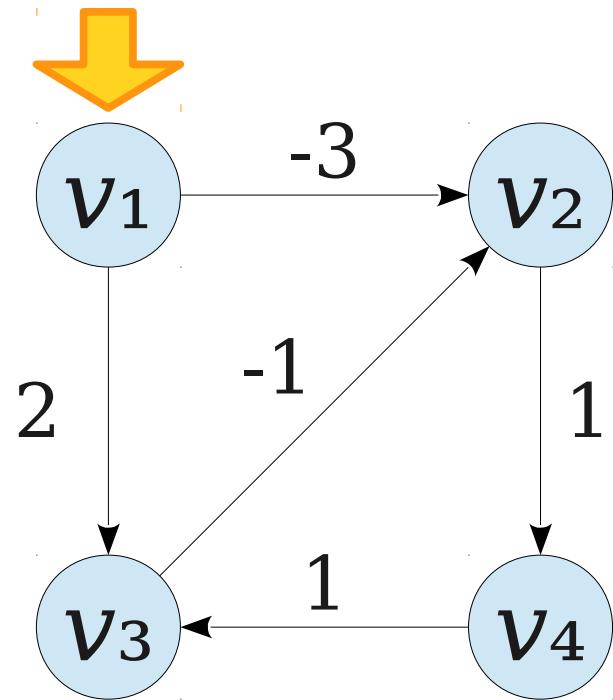
$$\text{OPT}(v, i) = \begin{cases} 0 & \text{if } i=0 \text{ and } v=s \\ \infty & \text{if } i=0 \text{ and } v \neq s \\ \min \left\{ \begin{array}{l} \text{OPT}(v, i-1), \\ \min_{(u,v) \in E} \{ \text{OPT}(u, i-1) + w(u, v) \} \end{array} \right\} & \text{otherwise} \end{cases}$$

	v_1	v_2	v_3	v_4
v_1	3	0		
v_2	2	0	-3	2
v_3	1	0	-3	2
v_4	0	0	∞	∞



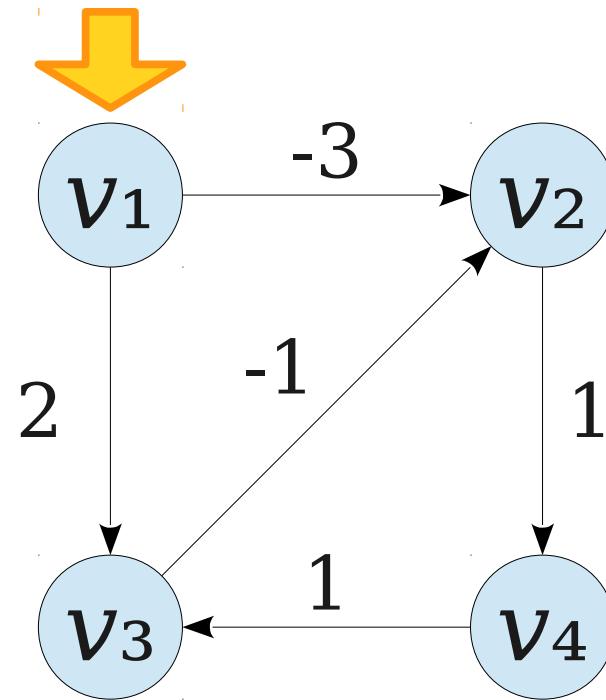
$$\text{OPT}(v, i) = \begin{cases} 0 & \text{if } i=0 \text{ and } v=s \\ \infty & \text{if } i=0 \text{ and } v \neq s \\ \min \left\{ \begin{array}{l} \text{OPT}(v, i-1), \\ \min_{(u,v) \in E} \{ \text{OPT}(u, i-1) + w(u, v) \} \end{array} \right\} & \text{otherwise} \end{cases}$$

	v_1	v_2	v_3	v_4
v_1	3	0	-3	
v_2	2	0	-3	2
v_3	1	0	-3	2
v_4	0	0	∞	∞



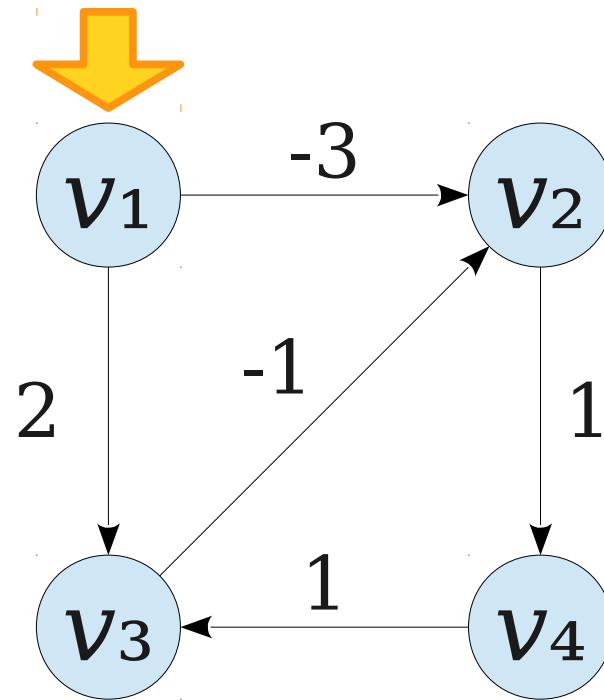
$$\text{OPT}(v, i) = \begin{cases} 0 & \text{if } i=0 \text{ and } v=s \\ \infty & \text{if } i=0 \text{ and } v \neq s \\ \min \left\{ \begin{array}{l} \text{OPT}(v, i-1), \\ \min_{(u,v) \in E} \{ \text{OPT}(u, i-1) + w(u, v) \} \end{array} \right\} & \text{otherwise} \end{cases}$$

	v_1	v_2	v_3	v_4	
3	0	-3	-1		
2	0	-3	2	-2	
1	0	-3	2	∞	
0	0	∞	∞	∞	



$$\text{OPT}(v, i) = \begin{cases} 0 & \text{if } i=0 \text{ and } v=s \\ \infty & \text{if } i=0 \text{ and } v \neq s \\ \min \left\{ \begin{array}{l} \text{OPT}(v, i-1), \\ \min_{(u,v) \in E} \{ \text{OPT}(u, i-1) + w(u, v) \} \end{array} \right\} & \text{otherwise} \end{cases}$$

	v_1	v_2	v_3	v_4
v_1	3	0	-3	-1
v_2	2	0	-3	2
v_3	1	0	-3	2
v_4	0	0	∞	∞



$$\text{OPT}(v, i) = \begin{cases} 0 & \text{if } i=0 \text{ and } v=s \\ \infty & \text{if } i=0 \text{ and } v \neq s \\ \min \left\{ \begin{array}{l} \text{OPT}(v, i-1), \\ \min_{(u,v) \in E} \{ \text{OPT}(u, i-1) + w(u, v) \} \end{array} \right\} & \text{otherwise} \end{cases}$$

Analyzing Time Complexity

- What is the time complexity of this algorithm?
 - Create a table DP of size $n \times n$.
 - Set $DP[v][0] = \infty$ for all $v \neq s$.
 - Set $DP[s][0] = 0$
 - For $i = 1$ to $n - 1$, for all $v \in V$:
 - Set $DP[v][i] = \min \{ DP[v][i - 1], \min \{ DP[u][i - 1] + w(u, v) \} \text{ (where } (u, v) \in E\}$
 - Return row n of DP.
- Answer: **O(mn)**, if you reverse G prior to running the algorithm.

Analyzing Space Complexity

- What is the *space* complexity of this algorithm?
 - Create a table DP of size $n \times n$.
 - Set $\text{DP}[v][0] = \infty$ for all $v \neq s$.
 - Set $\text{DP}[s][0] = 0$
 - For $i = 1$ to $n - 1$, for all $v \in V$:
 - Set $\text{DP}[v][i] = \min \{ \text{DP}[v][i - 1], \min \{ \text{DP}[u][i - 1] + w(u, v) \} \text{ (where } (u, v) \in E \}$
 - Return row n of DP.
- Answer: **O(n^2)**. (*Can we reduce this?*)

Shortest paths with negative weights: practical improvements

Space optimization. Maintain two 1D arrays (instead of 2D array).

- $d[v]$ = length of a shortest $v \rightsquigarrow t$ path that we have found so far.
- $\text{successor}[v]$ = next node on a $v \rightsquigarrow t$ path.

Performance optimization. If $d[w]$ was not updated in iteration $i - 1$, then no reason to consider edges entering w in iteration i .

Bellman–Ford–Moore: efficient implementation

BELLMAN–FORD–MOORE(V, E, c, t)

FOREACH node $v \in V$:

$d[v] \leftarrow \infty.$

$successor[v] \leftarrow null.$

$d[t] \leftarrow 0.$

FOR $i = 1$ TO $n - 1$

FOREACH node $w \in V$:

IF ($d[w]$ was updated in previous pass)

FOREACH edge $(v, w) \in E$:

IF ($d[v] > d[w] + \ell_{vw}$)

$d[v] \leftarrow d[w] + \ell_{vw}.$

$successor[v] \leftarrow w.$

pass i
 $O(m)$ time

IF (no $d[\cdot]$ value changed in pass i) STOP.

Bellman–Ford–Moore: analysis

Lemma 3. For each node v : $d[v]$ is the length of some $v \rightsquigarrow t$ path.

Lemma 4. For each node v : $d[v]$ is monotone non-increasing.

Lemma 5. After pass i , $d[v] \leq$ length of a shortest $v \rightsquigarrow t$ path using $\leq i$ edges.

Pf. [by induction on i]

- Base case: $i = 0$.
- Assume true after pass i .
- Let P be any $v \rightsquigarrow t$ path with $\leq i + 1$ edges.
- Let (v, w) be first edge in P and let P' be subpath from w to t .
- By inductive hypothesis, at the end of pass i , $d[w] \leq \ell(P')$ because P' is a $w \rightsquigarrow t$ path with $\leq i$ edges.
- After considering edge (v, w) in pass $i + 1$:

and by Lemma 4,
 $d[w]$ does not increase

$$\begin{aligned} d[v] &\leq \ell_{vw} + d[w] \\ &\leq \ell_{vw} + \ell(P') \\ &= \ell(P) \quad \blacksquare \end{aligned}$$

and by Lemma 4,
 $d[v]$ does not increase

Bellman–Ford–Moore: analysis

Theorem 2. Assuming no negative cycles, Bellman–Ford–Moore computes the lengths of the shortest $v \rightsquigarrow t$ paths in $O(mn)$ time and $\Theta(n)$ extra space.

Pf. Lemma 2 + Lemma 5. ■


shortest path exists and
has at most $n-1$ edges


after i passes,
 $d[v] \leq$ length of shortest path
that uses $\leq i$ edges

Remark. Bellman–Ford–Moore is typically faster in practice.

- Edge (v, w) considered in pass $i + 1$ only if $d[w]$ updated in pass i .
- If shortest path has k edges, then algorithm finds it after $\leq k$ passes.

BELLMAN-FORD(G, w, s)

```
1  INITIALIZE-SINGLE-SOURCE( $G, s$ )
2  for  $i = 1$  to  $|G.V| - 1$ 
3      for each edge  $(u, v) \in G.E$ 
4          RELAX( $u, v, w$ )
5  for each edge  $(u, v) \in G.E$ 
6      if  $v.d > u.d + w(u, v)$ 
7          return FALSE
8  return TRUE
```

Single-source shortest paths with negative weights

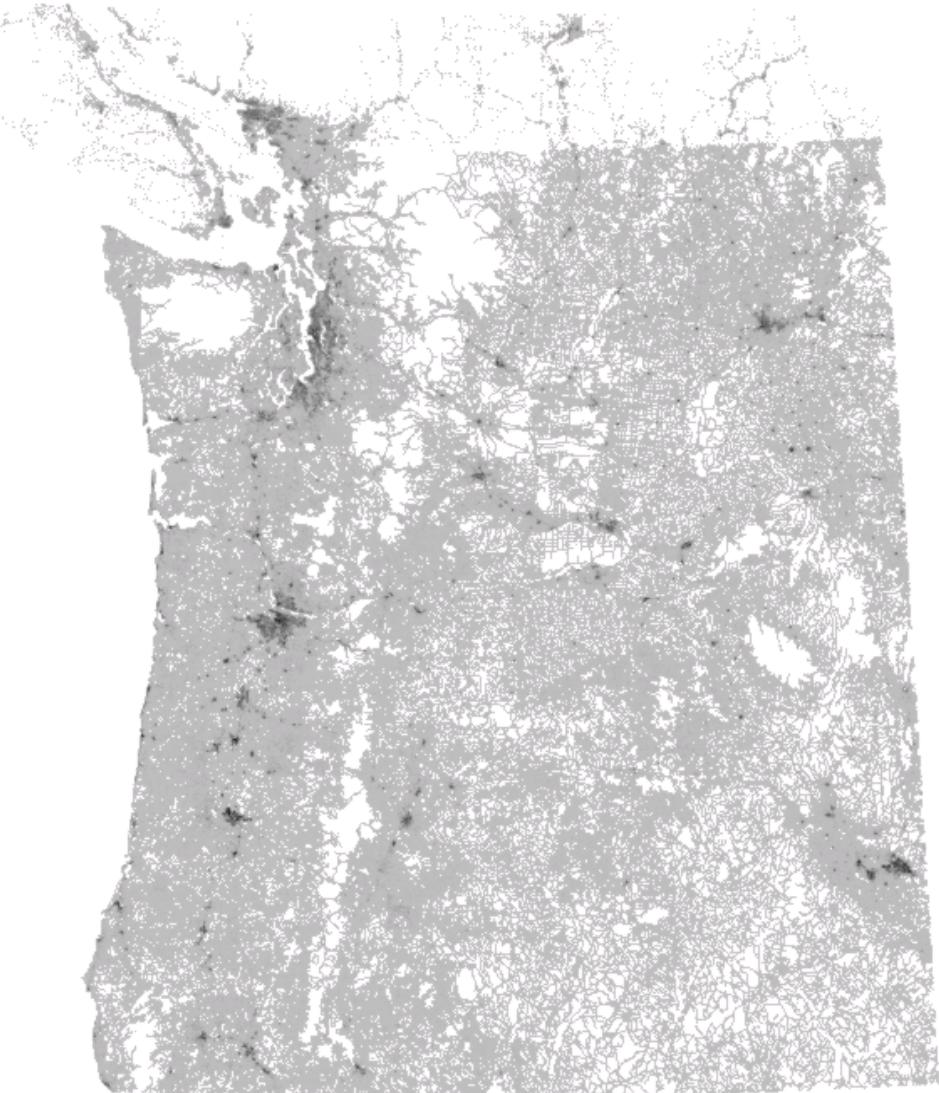
year	worst case	discovered by
1955	$O(n^4)$	Shimbel
1956	$O(m n^2 W)$	Ford
1958	$O(m n)$	Bellman, Moore
1983	$O(n^{3/4} m \log W)$	Gabow
1989	$O(m n^{1/2} \log(nW))$	Gabow–Tarjan
1993	$O(m n^{1/2} \log W)$	Goldberg
2005	$O(n^{2.38} W)$	Sankowski, Yuster–Zwick
2016	$\tilde{O}(n^{10/7} \log W)$	Cohen–Mądry–Sankowski–Vlădu
20xx	???	

single-source shortest paths with weights between $-W$ and W

Practical Concerns



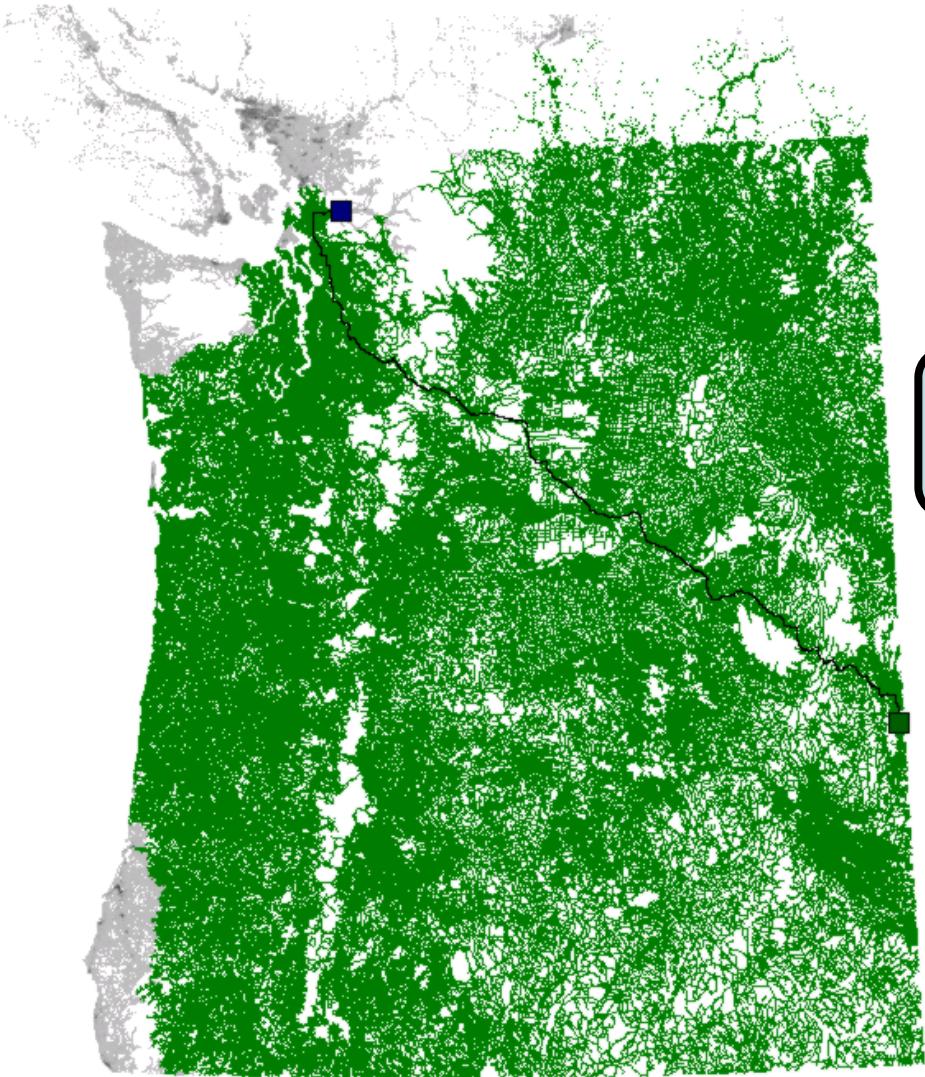
Dijkstra Example



1.6 million vertices
3.8 million edges
Distance = travel time.



Dijkstra Example

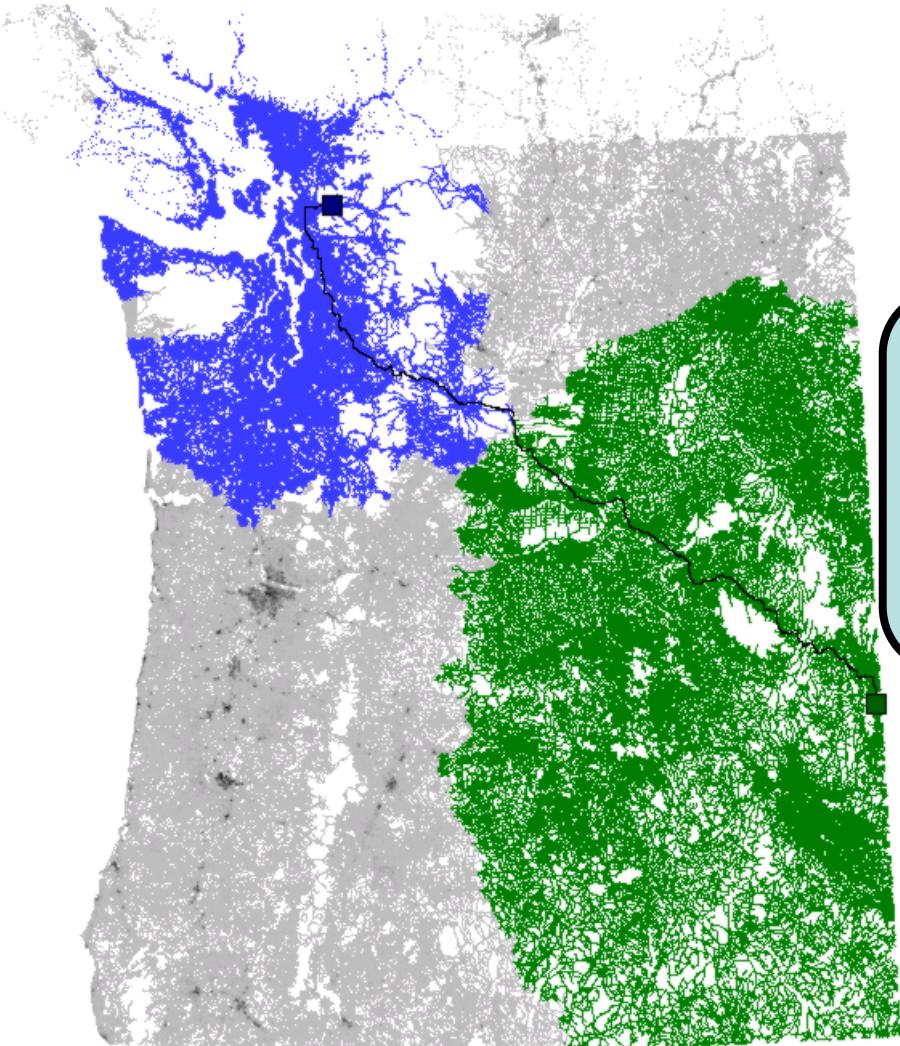


Searched Area
(starting from green point)

Problem of Dijkstra:
Didn't take account of where is t



Bidirectional Dijkstra



Forward search
Backward search

Problem of Bidirectional Dijkstra:
Forward search did not take
account of t
Backward search did not take
account of s .



A* Search

```
AStar( $G, c, s, t$ ) {
    Initialize set of explored nodes  $S \leftarrow \{s\}$ 

    // Maintain distance from  $s$  to each vertices in  $S$ 
     $d[s] \leftarrow 0$ 

    while ( $S \neq V$ )
    {
        Pick an edge  $(u, v)$  such that  $u \in S$  and  $v \notin S$  and
         $d[u] + c_{(u,v)} + h(v)$  is as small as possible.

        Add  $v$  to  $S$  and define  $d[v] = d[u] + c_{(u,v)}$ .
         $Parent(v) \leftarrow u$ .
    }
}
```



BFS



Dijkstra



A^*

- $h(v)$ is the estimate of distance from v to t
- If $h(v)$ is exactly the shortest distance from v to t , then the algorithm would go directly to t .

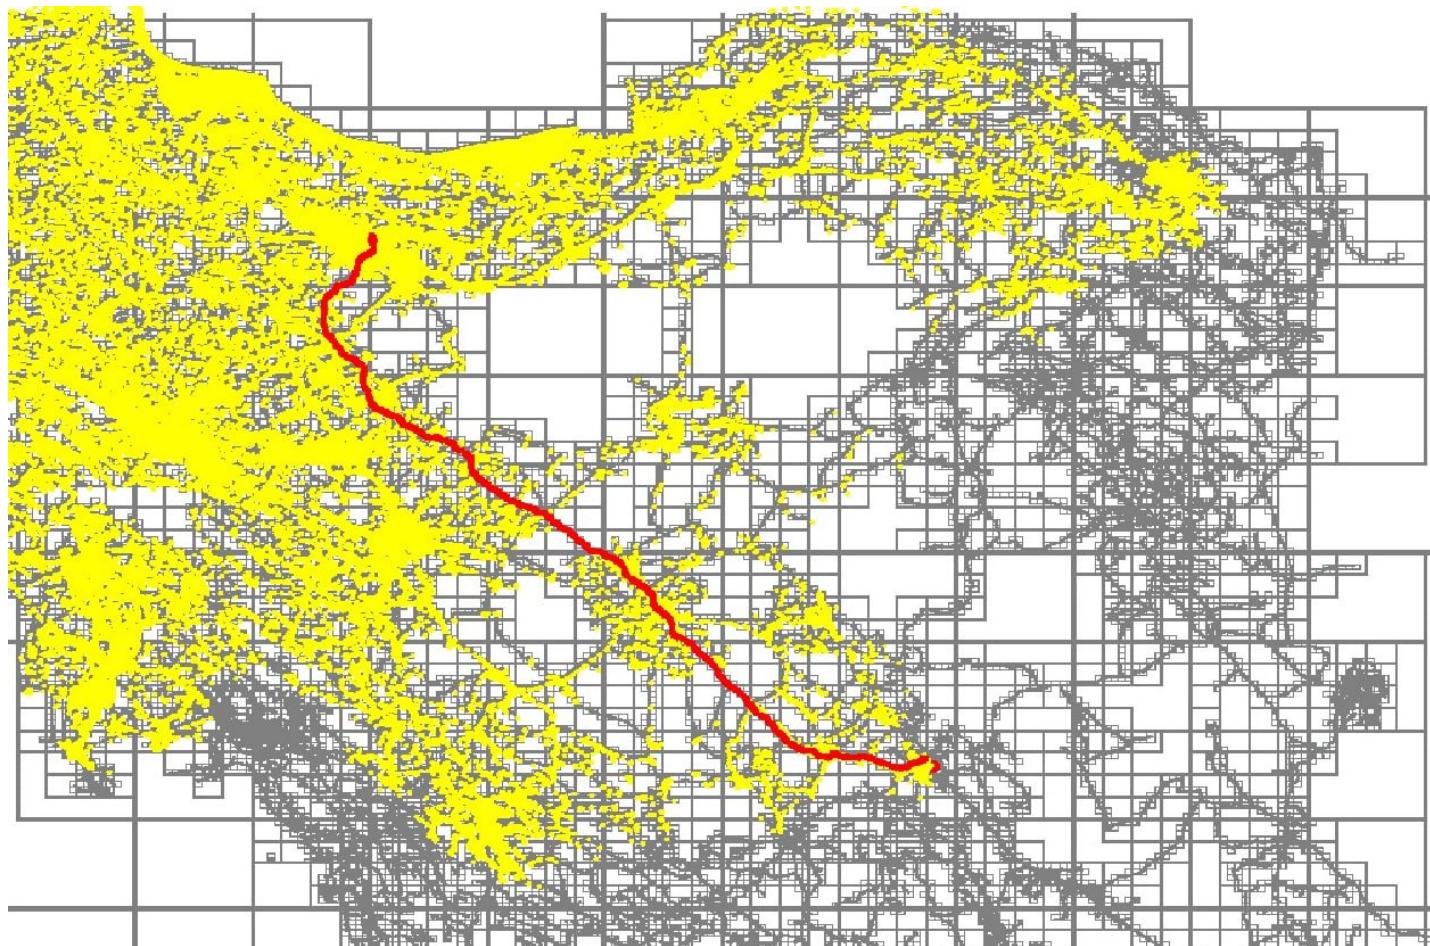
Routing Problem in Large Graphs

- Iran's Graph has 2,052,067 vertices and 2707945 edges.
- Consider Routing From Azadi Gate to Kerman
- The shortest path contains 3437 many edges.
- Dijkstra algorithm visits 1722445 many nodes!



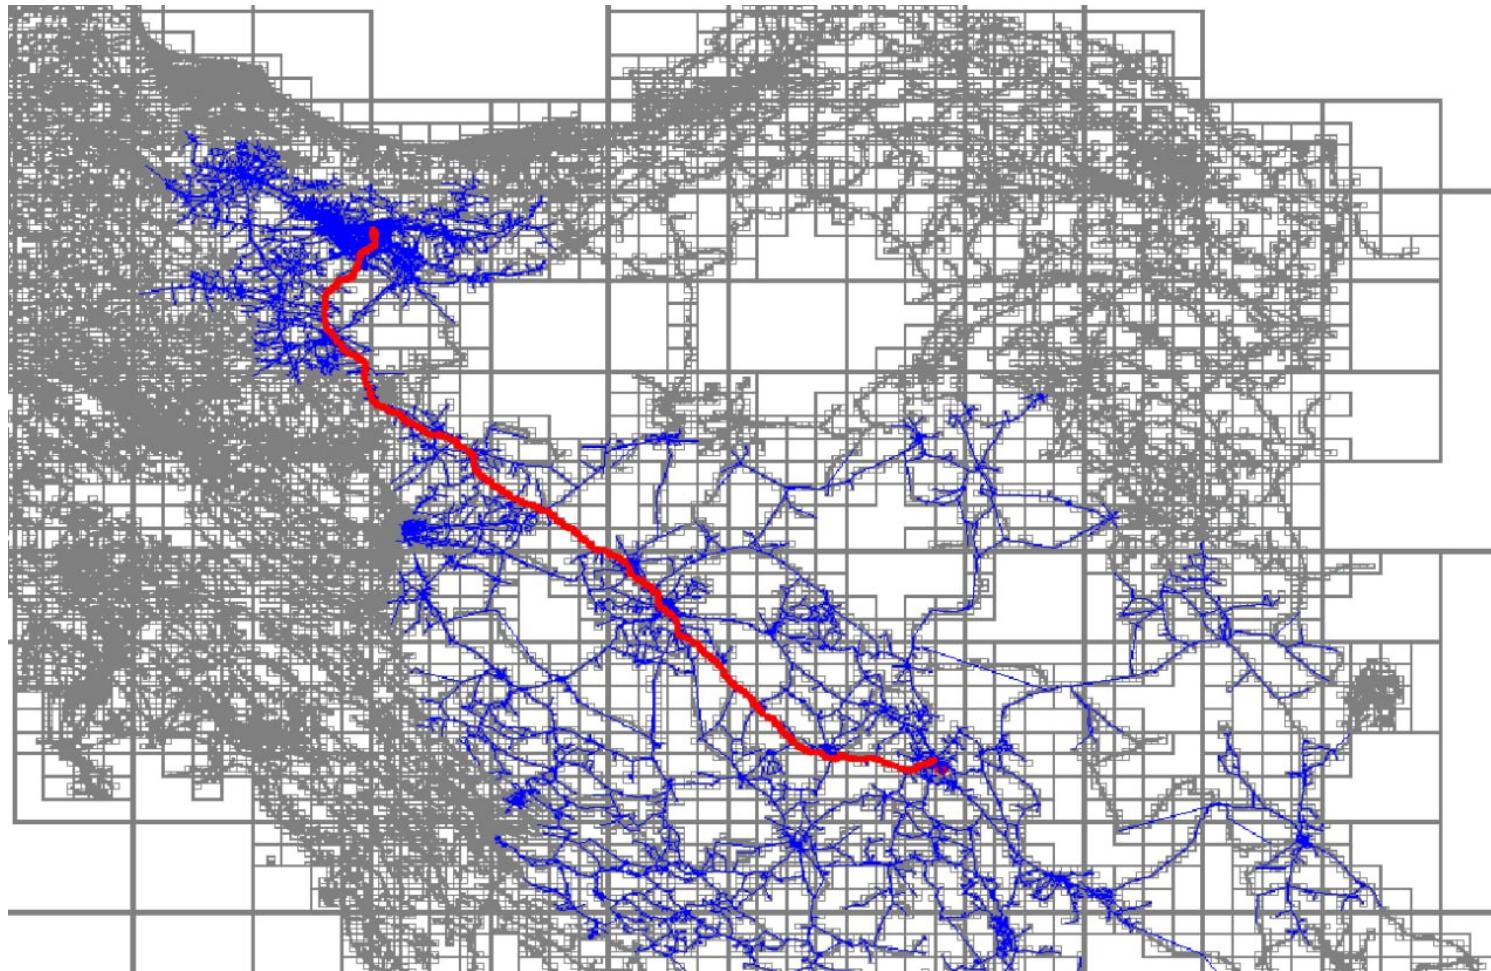
Dijkstra Performance

#Visited Nodes: 1,722,445



Bi-Directional Dijkstra

#Visited Nodes: 660,564



Landmark Selection

Preprocessing

- Random selection is fast.
- Many heuristics find better landmarks.
- Local search can find a good subset of candidate landmarks.
- We use a heuristic with local search.

Preprocessing/query trade-off.

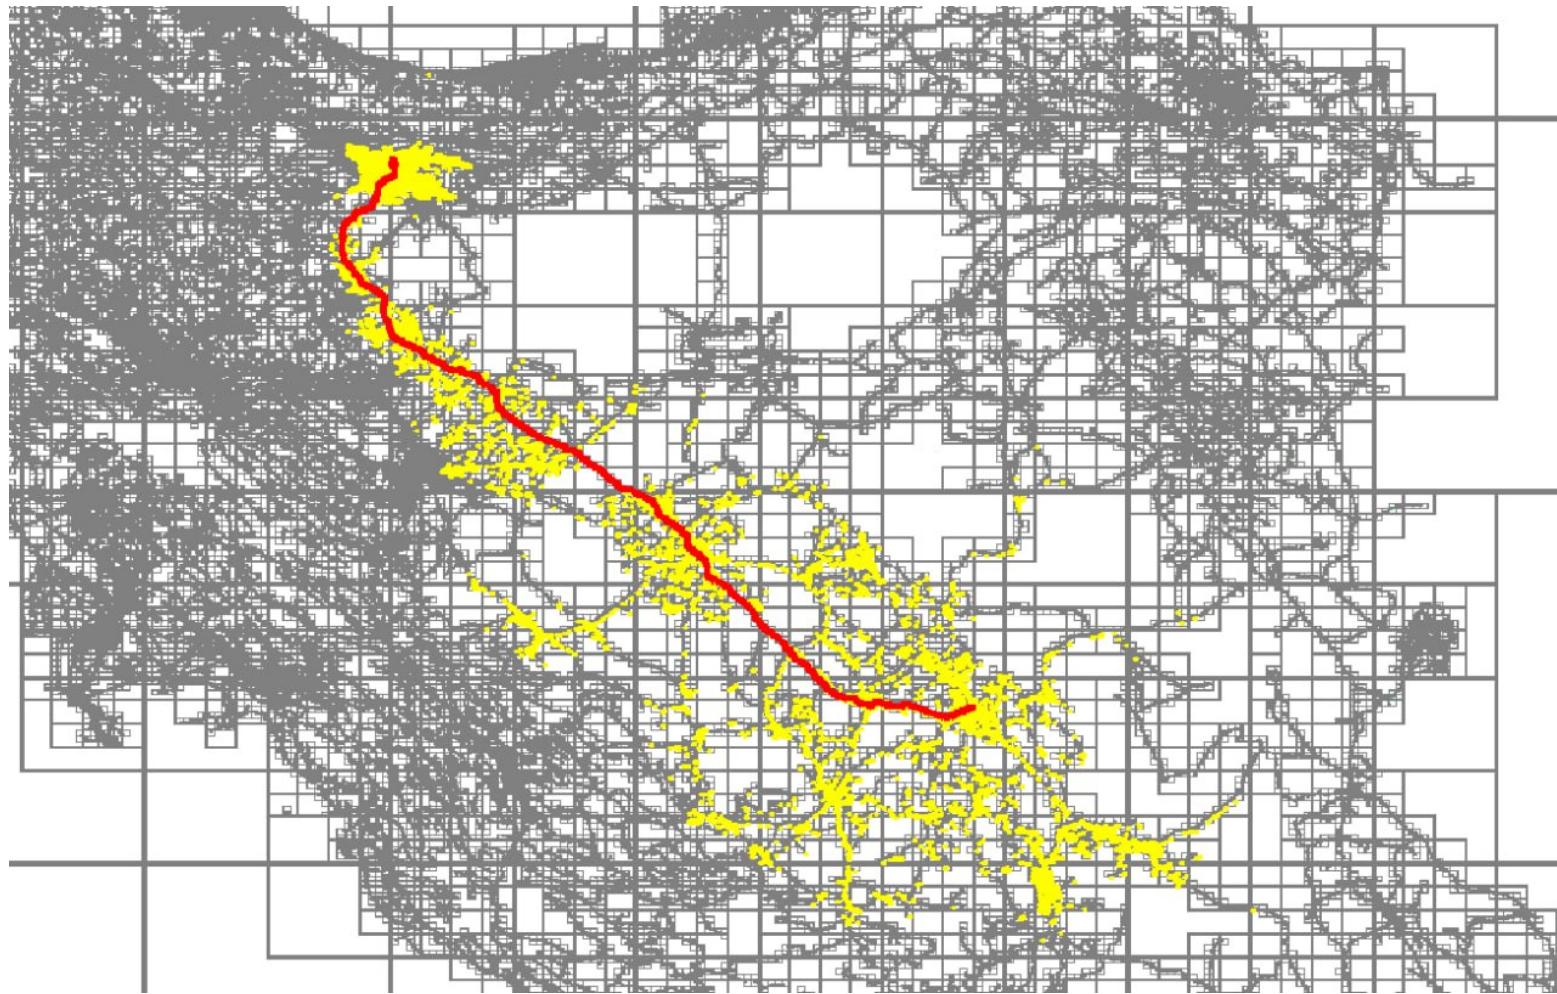
Query

- For a specific s, t pair, only some landmarks are useful.
- Use only **active landmarks** that give best bounds on $\text{dist}(s, t)$.
- If needed, **dynamically** add active landmarks (good for the search frontier).
- Only three active landmarks on the average.

Allows using many landmarks with small time overhead.

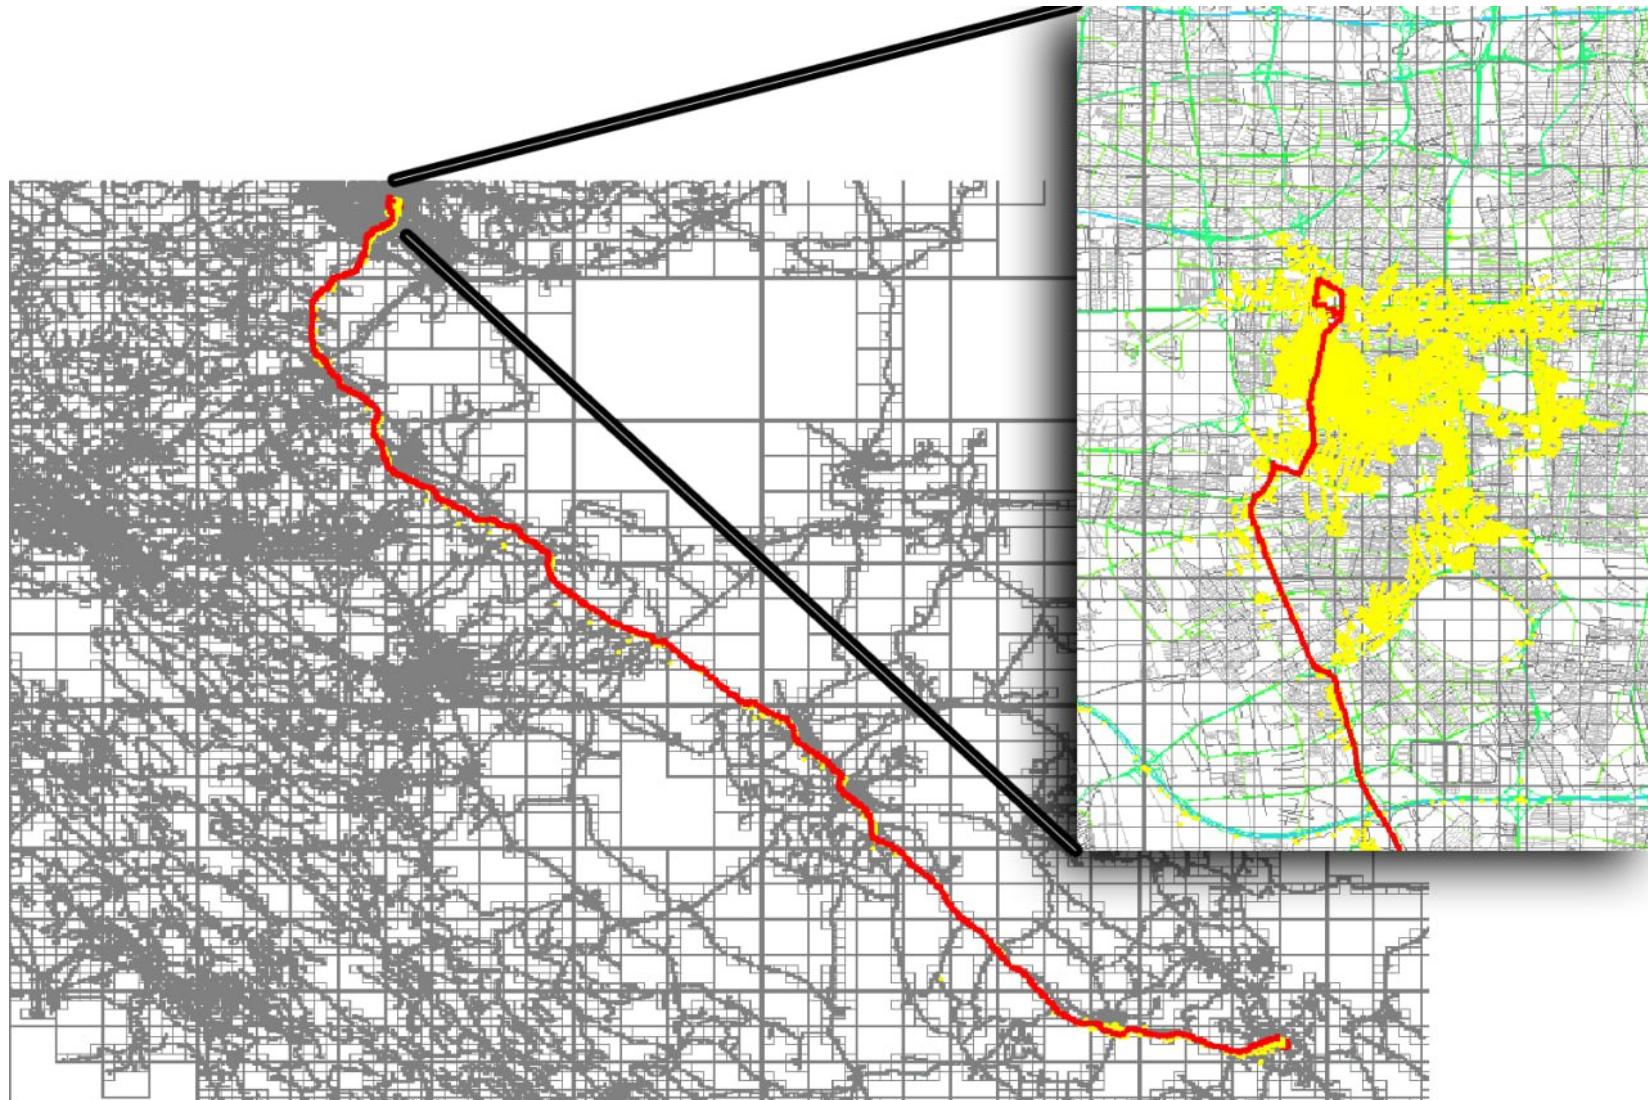
Dijkstra + Landmark

#Visited Nodes: 405,784



Bi-directional Dijkstra + Landmark

#Visited Nodes: 19,056



Algorithms & Photo By: A. R. Tofighi Mohammadi

